

D R A F T

Document for a Standard Message-Passing Interface

Message Passing Interface Forum

August 18, 2009

This work was supported in part by NSF and ARPA under NSF contract CDA-9115428 and Esprit under project HPC Standards (21111).

This is the result of a LaTeX run of a draft of a single chapter of the MPIF Final Report document.

Chapter 8

MPI Environmental Management

This chapter discusses routines for getting and, where appropriate, setting various parameters that relate to the MPI implementation and the execution environment (such as error handling). The procedures for entering and leaving the MPI execution environment are also described here.

8.1 Implementation Information

8.1.1 Version Inquiries

In order to cope with changes to the MPI Standard, there are both compile-time and run-time ways to determine which version of the standard is in use in the environment one is using.

The “version” will be represented by two separate integers, for the version and subversion: In C and C++,

```
#define MPI_VERSION      2
#define MPI_SUBVERSION [ticket101.] [1]2
```

in Fortran,

```
INTEGER MPI_VERSION, MPI_SUBVERSION
PARAMETER (MPI_VERSION      = 2)
PARAMETER (MPI_SUBVERSION = [ticket101.] [1]2)
```

For runtime determination,

`MPI_GET_VERSION(version, subversion)`

OUT	version	version number (integer)
OUT	subversion	subversion number (integer)

```
int MPI_Get_version(int *version, int *subversion)
```

```
MPI_GET_VERSION(VERSION, SUBVERSION, IERROR)
```

```
INTEGER VERSION, SUBVERSION, IERROR
```

```
{void MPI::Get_version(int& version, int& subversion) (binding deprecated, see
```

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47 ticket150.
48 ticket150.

Section 15.2) }

MPI_GET_VERSION is one of the few functions that can be called before MPI_INIT and after MPI_FINALIZE. Valid (MPI_VERSION, MPI_SUBVERSION) pairs in this and previous versions of the MPI standard are (2,2), (2,1), (2,0), and (1,2).

8.1.2 Environmental Inquiries

A set of attributes that describe the execution environment are attached to the communicator MPI_COMM_WORLD when MPI is initialized. The value of these attributes can be inquired by using the function [MPI_ATTR_GET]MPI_COMM_GET_ATTR described in Chapter 6. It is erroneous to delete these attributes, free their keys, or change their values.

The list of predefined attribute keys include

MPI_TAG_UB Upper bound for tag value.

MPI_HOST Host process rank, if such exists, MPI_PROC_NULL, otherwise.

MPI_IO rank of a node that has regular I/O facilities (possibly myrank). Nodes in the same communicator may return different values for this parameter.

MPI_WTIME_IS_GLOBAL Boolean variable that indicates whether clocks are synchronized.

Vendors may add implementation specific parameters (such as node number, real memory size, virtual memory size, etc.)

These predefined attributes do not change value between MPI initialization (MPI_INIT and MPI completion (MPI_FINALIZE), and cannot be updated or deleted by users.

Advice to users. Note that in the C binding, the value returned by these attributes is a *pointer* to an `int` containing the requested value. (*End of advice to users.*)

The required parameter values are discussed in more detail below:

Tag Values

Tag values range from 0 to the value returned for MPI_TAG_UB inclusive. These values are guaranteed to be unchanging during the execution of an MPI program. In addition, the tag upper bound value must be *at least* 32767. An MPI implementation is free to make the value of MPI_TAG_UB larger than this; for example, the value $2^{30} - 1$ is also a legal value for MPI_TAG_UB.

The attribute MPI_TAG_UB has the same value on all processes of MPI_COMM_WORLD.

Host Rank

The value returned for MPI_HOST gets the rank of the HOST process in the group associated with communicator MPI_COMM_WORLD, if there is such. MPI_PROC_NULL is returned if there is no host. MPI does not specify what it means for a process to be a HOST, nor does it require that a HOST exists.

The attribute MPI_HOST has the same value on all processes of MPI_COMM_WORLD.

IO Rank

The value returned for `MPI_IO` is the rank of a processor that can provide language-standard I/O facilities. For Fortran, this means that all of the Fortran I/O operations are supported (e.g., `OPEN`, `REWIND`, `WRITE`). For C and C++, this means that all of the ISO C and C++, I/O operations are supported (e.g., `fopen`, `fprintf`, `lseek`).

If every process can provide language-standard I/O, then the value `MPI_ANY_SOURCE` will be returned. Otherwise, if the calling process can provide language-standard I/O, then its rank will be returned. Otherwise, if some process can provide language-standard I/O then the rank of one such process will be returned. The same value need not be returned by all processes. If no process can provide language-standard I/O, then the value `MPI_PROC_NULL` will be returned.

Advice to users. Note that input is not collective, and this attribute does *not* indicate which process can or does provide input. (*End of advice to users.*)

Clock Synchronization

The value returned for `MPI_WTIME_IS_GLOBAL` is 1 if clocks at all processes in `MPI_COMM_WORLD` are synchronized, 0 otherwise. A collection of clocks is considered synchronized if explicit effort has been taken to synchronize them. The expectation is that the variation in time, as measured by calls to `MPI_WTIME`, will be less than one half the round-trip time for an MPI message of length zero. If time is measured at a process just before a send and at another process just after a matching receive, the second time should be always higher than the first one.

The attribute `MPI_WTIME_IS_GLOBAL` need not be present when the clocks are not synchronized (however, the attribute key `MPI_WTIME_IS_GLOBAL` is always valid). This attribute may be associated with communicators other than `MPI_COMM_WORLD`.

The attribute `MPI_WTIME_IS_GLOBAL` has the same value on all processes of `MPI_COMM_WORLD`.

`MPI_GET_PROCESSOR_NAME(name, resultlen)`

OUT	name	A unique specifier for the actual (as opposed to virtual) node.
OUT	resultlen	Length (in printable characters) of the result returned in name

```
int MPI_Get_processor_name(char *name, int *resultlen)
```

```
MPI_GET_PROCESSOR_NAME( NAME, RESULTLEN, IERROR)
```

```
CHARACTER*(*) NAME
```

```
INTEGER RESULTLEN, IERROR
```

```
{void MPI::Get_processor_name(char* name, int& resultlen) (binding deprecated,  
see Section 15.2) }
```

This routine returns the name of the processor on which it was called at the moment of the call. The name is a character string for maximum flexibility. From this value it must be possible to identify a specific piece of hardware; possible values include “processor

ticket150.
ticket150.

9 in rack 4 of mpp.cs.org” and “231” (where 231 is the actual processor number in the running homogeneous system). The argument `name` must represent storage that is at least `MPI_MAX_PROCESSOR_NAME` characters long. `MPI_GET_PROCESSOR_NAME` may write up to this many characters into `name`.

The number of characters actually written is returned in the output argument, `resultlen`. In C, a null character is additionally stored at `name[resultlen]`. The `resultlen` cannot be larger than `MPI_MAX_PROCESSOR_NAME-1`. In Fortran, `name` is padded on the right with blank characters. The `resultlen` cannot be larger than `MPI_MAX_PROCESSOR_NAME`.

Rationale. This function allows MPI implementations that do process migration to return the current processor. Note that nothing in MPI *requires* or defines process migration; this definition of `MPI_GET_PROCESSOR_NAME` simply allows such an implementation. (*End of rationale.*)

Advice to users. The user must provide at least `MPI_MAX_PROCESSOR_NAME` space to write the processor name — processor names can be this long. The user should examine the output argument, `resultlen`, to determine the actual length of the name. (*End of advice to users.*)

The constant `MPI_BSEND_OVERHEAD` provides an upper bound on the fixed overhead per message buffered by a call to `MPI_BSEND` (see Section 3.6.1).

8.2 Memory Allocation

In some systems, message-passing and remote-memory-access (RMA) operations run faster when accessing specially allocated memory (e.g., memory that is shared by the other processes in the communicating group on an SMP). MPI provides a mechanism for allocating and freeing such special memory. The use of such memory for message-passing or RMA is not mandatory, and this memory can be used without restrictions as any other dynamically allocated memory. However, implementations may restrict the use of the `MPI_WIN_LOCK` and `MPI_WIN_UNLOCK` functions to windows allocated in such memory (see Section 11.4.3.)

`MPI_ALLOC_MEM(size, info, baseptr)`

IN	size	size of memory segment in bytes ([nonnegative]non-negative integer)
IN	info	info argument (handle)
OUT	baseptr	pointer to beginning of memory segment allocated

`int MPI_Alloc_mem(MPI_Aint size, MPI_Info info, void *baseptr)`

`MPI_ALLOC_MEM(SIZE, INFO, BASEPTR, IERROR)`

INTEGER INFO, IERROR

INTEGER(KIND=MPI_ADDRESS_KIND) SIZE, BASEPTR

{void* MPI::Alloc_mem(MPI::Aint size, const MPI::Info& info) (*binding deprecated, see Section 15.2*) }

The `info` argument can be used to provide directives that control the desired location of the allocated memory. Such a directive does not affect the semantics of the call. Valid `info` values are implementation-dependent; a null directive value of `info = MPI_INFO_NULL` is always valid.

The function `MPI_ALLOC_MEM` may return an error code of class `MPI_ERR_NO_MEM` to indicate it failed because memory is exhausted.

`MPI_FREE_MEM(base)`

IN base initial address of memory segment allocated by
 `MPI_ALLOC_MEM` (choice)

`int MPI_Free_mem(void *base)`

`MPI_FREE_MEM(BASE, IERROR)`

 <type> BASE(*)

 INTEGER IERROR

`{void MPI::Free_mem(void *base) (binding deprecated, see Section 15.2) }`

The function `MPI_FREE_MEM` may return an error code of class `MPI_ERR_BASE` to indicate an invalid base argument.

Rationale. The C and C++ bindings of `MPI_ALLOC_MEM` and `MPI_FREE_MEM` are similar to the bindings for the `malloc` and `free` C library calls: a call to `MPI_Alloc_mem(..., &base)` should be paired with a call to `MPI_Free_mem(base)` (one less level of indirection). Both arguments are declared to be of same type `void*` so as to facilitate type casting. The Fortran binding is consistent with the C and C++ bindings: the Fortran `MPI_ALLOC_MEM` call returns in `baseptr` the (integer valued) address of the allocated memory. The `base` argument of `MPI_FREE_MEM` is a choice argument, which passes (a reference to) the variable stored at that location. (*End of rationale.*)

Advice to implementors. If `MPI_ALLOC_MEM` allocates special memory, then a design similar to the design of C `malloc` and `free` functions has to be used, in order to find out the size of a memory segment, when the segment is freed. If no special memory is used, `MPI_ALLOC_MEM` simply invokes `malloc`, and `MPI_FREE_MEM` invokes `free`.

A call to `MPI_ALLOC_MEM` can be used in shared memory systems to allocate memory in a shared memory segment. (*End of advice to implementors.*)

Example 8.1 Example of use of `MPI_ALLOC_MEM`, in Fortran with pointer support. We assume 4-byte REALs, and assume that pointers are address-sized.

```
REAL A
POINTER (P, A(100,100)) ! no memory is allocated
CALL MPI_ALLOC_MEM(4*100*100, MPI_INFO_NULL, P, IERR)
! memory is allocated
...
A(3,5) = 2.71;
```

```

1  ...
2  CALL MPI_FREE_MEM(A, IERR) ! memory is freed

```

Since standard Fortran does not support (C-like) pointers, this code is not Fortran 77 or Fortran 90 code. Some compilers (in particular, at the time of writing, g77 and Fortran compilers for Intel) do not support this code.

Example 8.2 Same example, in C

```

9  float  (* f)[100][100] ;
10 /* no memory is allocated */
11 MPI_Alloc_mem(sizeof(float)*100*100, MPI_INFO_NULL, &f);
12 /* memory allocated */
13 ...
14 (*f)[5][3] = 2.71;
15 ...
16 MPI_Free_mem(f);

```

8.3 Error Handling

An MPI implementation cannot or may choose not to handle some errors that occur during MPI calls. These can include errors that generate exceptions or traps, such as floating point errors or access violations. The set of errors that are handled by MPI is implementation-dependent. Each such error generates an **MPI exception**.

The above text takes precedence over any text on error handling within this document. Specifically, text that states that errors *will* be handled should be read as *may* be handled.

A user can associate error handlers to three types of objects: communicators, windows, and files. The specified error handling routine will be used for any MPI exception that occurs during a call to MPI for the respective object. MPI calls that are not related to any objects are considered to be attached to the communicator MPI_COMM_WORLD. The attachment of error handlers to objects is purely local: different processes may attach different error handlers to corresponding objects.

Several predefined error handlers are available in MPI:

MPI_ERRORS_ARE_FATAL The handler, when called, causes the program to abort on all executing processes. This has the same effect as if MPI_ABORT was called by the process that invoked the handler.

MPI_ERRORS_RETURN The handler has no effect other than returning the error code to the user.

Implementations may provide additional predefined error handlers and programmers can code their own error handlers.

The error handler MPI_ERRORS_ARE_FATAL is associated by default with MPI_COMM_WORLD after initialization. Thus, if the user chooses not to control error handling, every error that MPI handles is treated as fatal. Since (almost) all MPI calls return an error code, a user may choose to handle errors in its main code, by testing the return code of MPI calls and executing a suitable recovery code when the call was not successful. In this case, the error handler MPI_ERRORS_RETURN will be used. Usually it is more convenient and more

efficient not to test for errors after each MPI call, and have such error handled by a non trivial MPI error handler.

After an error is detected, the state of MPI is undefined. That is, using a user-defined error handler, or `MPI_ERRORS_RETURN`, does *not* necessarily allow the user to continue to use MPI after an error is detected. The purpose of these error handlers is to allow a user to issue user-defined error messages and to take actions unrelated to MPI (such as flushing I/O buffers) before a program exits. An MPI implementation is free to allow MPI to continue after an error but is not required to do so.

Advice to implementors. A good quality implementation will, to the greatest possible extent, circumscribe the impact of an error, so that normal processing can continue after an error handler was invoked. The implementation documentation will provide information on the possible effect of each class of errors. (*End of advice to implementors.*)

An MPI error handler is an opaque object, which is accessed by a handle. MPI calls are provided to create new error handlers, to associate error handlers with objects, and to test which error handler is associated with an object. C and C++ have distinct typedefs for user defined error handling callback functions that accept communicator, file, and window arguments. In Fortran there are three user routines.

An error handler object is created by a call to `MPI_XXX_CREATE_ERRHANDLER(function, errhandler)`, where XXX is, respectively, COMM, WIN, or FILE.

An error handler is attached to a communicator, window, or file by a call to `MPI_XXX_SET_ERRHANDLER`. The error handler must be either a predefined error handler, or an error handler that was created by a call to `MPI_XXX_CREATE_ERRHANDLER`, with matching XXX. The predefined error handlers `MPI_ERRORS_RETURN` and `MPI_ERRORS_ARE_FATAL` can be attached to communicators, windows, and files. In C++, the predefined error handler `MPI::ERRORS_THROW_EXCEPTIONS` can also be attached to communicators, windows, and files.

The error handler currently associated with a communicator, window, or file can be retrieved by a call to `MPI_XXX_GET_ERRHANDLER`.

The MPI function `MPI_ERRHANDLER_FREE` can be used to free an error handler that was created by a call to `MPI_XXX_CREATE_ERRHANDLER`.

`MPI_{COMM,WIN,FILE}_GET_ERRHANDLER` behave as if a new error handler object is created. That is, once the error handler is no longer needed, `MPI_ERRHANDLER_FREE` should be called with the error handler returned from `MPI_ERRHANDLER_GET` or `MPI_{COMM,WIN,FILE}_GET_ERRHANDLER` to mark the error handler for deallocation. This provides behavior similar to that of `MPI_COMM_GROUP` and `MPI_GROUP_FREE`.

Advice to implementors. High-quality implementation should raise an error when an error handler that was created by a call to `MPI_XXX_CREATE_ERRHANDLER` is attached to an object of the wrong type with a call to `MPI_YYY_SET_ERRHANDLER`. To do so, it is necessary to maintain, with each error handler, information on the typedef of the associated user function. (*End of advice to implementors.*)

The syntax for these calls is given below.

8.3.1 Error Handlers for Communicators

MPI_COMM_CREATE_ERRHANDLER(function, errhandler)

IN	function	user defined error handling procedure (function)
OUT	errhandler	MPI error handler (handle)

```
int MPI_Comm_create_errhandler(MPI_Comm_errhandler_ [fn]function *function,
                               MPI_Errhandler *errhandler)
```

MPI_COMM_CREATE_ERRHANDLER(FUNCTION, ERRHANDLER, IERROR)

EXTERNAL FUNCTION

INTEGER ERRHANDLER, IERROR

```
{static MPI::Errhandler
```

```
    MPI::Comm::Create_errhandler(MPI::Comm::Errhandler_ [fn]function*
    function) (binding deprecated, see Section 15.2) }
```

Creates an error handler that can be attached to communicators. This function is identical to MPI_ERRHANDLER_CREATE, whose use is deprecated.

The user routine should be, in C, a function of type [MPI_Comm_errhandler_fn] MPI_Comm_errhandler_function, which is defined as

```
typedef void MPI_Comm_errhandler_function(MPI_Comm *, int *, ...);
```

The first argument is the communicator in use. The second is the error code to be returned by the MPI routine that raised the error. If the routine would have returned MPI_ERR_IN_STATUS, it is the error code returned in the status for the request that caused the error handler to be invoked. The remaining arguments are “stdargs” arguments whose number and meaning is implementation-dependent. An implementation should clearly document these arguments. Addresses are used so that the handler may be written in Fortran. This typedef replaces MPI_Handler_function, whose use is deprecated.

In Fortran, the user routine should be of the form:

```
SUBROUTINE COMM_ERRHANDLER_FUNCTION(COMM, ERROR_CODE)
  INTEGER COMM, ERROR_CODE
```

Advice to users. Users are discouraged from using a Fortran {COMM|WIN|FILE}_ERRHANDLER_FUNCTION since the routine expects a variable number of arguments. Some Fortran systems may allow this but some may fail to give the correct result or compile/link this code. Thus, it will not, in general, be possible to create portable code with a Fortran {COMM|WIN|FILE}_ERRHANDLER_FUNCTION. *(End of advice to users.)*

In C++, the user routine should be of the form:

```
typedef void MPI::Comm::Errhandler_function(MPI::Comm &, int *, ...);
```

Rationale. The variable argument list is provided because it provides an ISO-standard hook for providing additional information to the error handler; without this hook, ISO C prohibits additional arguments. *(End of rationale.)*

Advice to users. A newly created communicator inherits the error handler that is associated with the “parent” communicator. In particular, the user can specify a “global” error handler for all communicators by associating this handler with the communicator MPI_COMM_WORLD immediately after initialization. (*End of advice to users.*)

MPI_COMM_SET_ERRHANDLER(comm, errhandler)

INOUT comm communicator (handle)

IN errhandler new error handler for communicator (handle)

int MPI_Comm_set_errhandler(MPI_Comm comm, MPI_Errhandler errhandler)

MPI_COMM_SET_ERRHANDLER(COMM, ERRHANDLER, IERROR)

INTEGER COMM, ERRHANDLER, IERROR

{void MPI::Comm::Set_errhandler(const MPI::Errhandler& errhandler) (binding deprecated, see Section 15.2) }

Attaches a new error handler to a communicator. The error handler must be either a predefined error handler, or an error handler created by a call to MPI_COMM_CREATE_ERRHANDLER. This call is identical to MPI_ERRHANDLER_SET, whose use is deprecated.

MPI_COMM_GET_ERRHANDLER(comm, errhandler)

IN comm communicator (handle)

OUT errhandler error handler currently associated with communicator (handle)

int MPI_Comm_get_errhandler(MPI_Comm comm, MPI_Errhandler *errhandler)

MPI_COMM_GET_ERRHANDLER(COMM, ERRHANDLER, IERROR)

INTEGER COMM, ERRHANDLER, IERROR

{MPI::Errhandler MPI::Comm::Get_errhandler() const (binding deprecated, see Section 15.2) }

Retrieves the error handler currently associated with a communicator. This call is identical to MPI_ERRHANDLER_GET, whose use is deprecated.

Example: A library function may register at its entry point the current error handler for a communicator, set its own private error handler for this communicator, and restore before exiting the previous error handler.

ticket150.
ticket150.

ticket150.
ticket150.

8.3.2 Error Handlers for Windows

MPI_WIN_CREATE_ERRHANDLER(function, errhandler)

IN	function	user defined error handling procedure (function)
OUT	errhandler	MPI error handler (handle)

```
int MPI_Win_create_errhandler(MPI_Win_errhandler_function *function,
                             MPI_Errhandler *errhandler)
```

MPI_WIN_CREATE_ERRHANDLER(FUNCTION, ERRHANDLER, IERROR)

EXTERNAL FUNCTION

INTEGER ERRHANDLER, IERROR

```
{static MPI::Errhandler
    MPI::Win::Create_errhandler(MPI::Win::Errhandler_function*
    function) (binding deprecated, see Section 15.2) }
```

Creates an error handler that can be attached to a window object. The user routine should be, in C, a function of type `[MPI_Win_errhandler_fn]MPI_Win_errhandler_function` which is defined as

```
typedef void MPI_Win_errhandler_function(MPI_Win *, int *, ...);
```

The first argument is the window in use, the second is the error code to be returned.

In Fortran, the user routine should be of the form:

```
SUBROUTINE WIN_ERRHANDLER_FUNCTION(WIN, ERROR_CODE)
    INTEGER WIN, ERROR_CODE
```

In C++, the user routine should be of the form:

```
typedef void MPI::Win::Errhandler_function(MPI::Win &, int *, ...);
```

MPI_WIN_SET_ERRHANDLER(win, errhandler)

INOUT	win	window (handle)
IN	errhandler	new error handler for window (handle)

int MPI_Win_set_errhandler(MPI_Win win, MPI_Errhandler errhandler)

MPI_WIN_SET_ERRHANDLER(WIN, ERRHANDLER, IERROR)

INTEGER WIN, ERRHANDLER, IERROR

```
{void MPI::Win::Set_errhandler(const MPI::Errhandler& errhandler) (binding
    deprecated, see Section 15.2) }
```

Attaches a new error handler to a window. The error handler must be either a pre-defined error handler, or an error handler created by a call to MPI_WIN_CREATE_ERRHANDLER.

MPI_WIN_GET_ERRHANDLER(win, errhandler)

IN	win	window (handle)
OUT	errhandler	error handler currently associated with window (handle)

int MPI_Win_get_errhandler(MPI_Win win, MPI_Errhandler *errhandler)

MPI_WIN_GET_ERRHANDLER(WIN, ERRHANDLER, IERROR)

INTEGER WIN, ERRHANDLER, IERROR

{MPI::Errhandler MPI::Win::Get_errhandler() const *(binding deprecated, see Section 15.2)* }

Retrieves the error handler currently associated with a window.

8.3.3 Error Handlers for Files

MPI_FILE_CREATE_ERRHANDLER(function, errhandler)

IN	function	user defined error handling procedure (function)
OUT	errhandler	MPI error handler (handle)

`int MPI_File_create_errhandler(MPI_File_errhandler_function *function, MPI_Errhandler *errhandler)`

MPI_FILE_CREATE_ERRHANDLER(FUNCTION, ERRHANDLER, IERROR)

EXTERNAL FUNCTION

INTEGER ERRHANDLER, IERROR

`{static MPI::Errhandler MPI::File::Create_errhandler(MPI::File::Errhandler_function* function) (binding deprecated, see Section 15.2) }`

Creates an error handler that can be attached to a file object. The user routine should be, in C, a function of type `[MPI_File_errhandler_fn]MPI_File_errhandler_function`, which is defined as

`typedef void MPI_File_errhandler_function(MPI_File *, int *, ...);`

The first argument is the file in use, the second is the error code to be returned.

In Fortran, the user routine should be of the form:

`SUBROUTINE FILE_ERRHANDLER_FUNCTION(FILE, ERROR_CODE)
INTEGER FILE, ERROR_CODE`

In C++, the user routine should be of the form:

`typedef void MPI::File::Errhandler_function(MPI::File &, int *, ...);`

```
1 MPI_FILE_SET_ERRHANDLER(file, errhandler)
```

```
2     INOUT    file                                file (handle)
```

```
3     IN       errhandler                          new error handler for file (handle)
```

```
5
6 int MPI_File_set_errhandler(MPI_File file, MPI_Errhandler errhandler)
```

```
7 MPI_FILE_SET_ERRHANDLER(FILE, ERRHANDLER, IERROR)
```

```
8     INTEGER FILE, ERRHANDLER, IERROR
```

```
9 {void MPI::File::Set_errhandler(const MPI::Errhandler& errhandler) (binding
10 deprecated, see Section 15.2) }
```

Attaches a new error handler to a file. The error handler must be either a predefined error handler, or an error handler created by a call to MPI_FILE_CREATE_ERRHANDLER.

```
16 MPI_FILE_GET_ERRHANDLER(file, errhandler)
```

```
17     IN       file                                file (handle)
```

```
18     OUT      errhandler                          error handler currently associated with file (handle)
```

```
20
21 int MPI_File_get_errhandler(MPI_File file, MPI_Errhandler *errhandler)
```

```
22 MPI_FILE_GET_ERRHANDLER(FILE, ERRHANDLER, IERROR)
```

```
23     INTEGER FILE, ERRHANDLER, IERROR
```

```
24 {MPI::Errhandler MPI::File::Get_errhandler() const (binding deprecated, see
25 Section 15.2) }
```

Retrieves the error handler currently associated with a file.

8.3.4 Freeing Errorhandlers and Retrieving Error Strings

```
33 MPI_ERRHANDLER_FREE( errhandler )
```

```
34     INOUT    errhandler                          MPI error handler (handle)
```

```
36
37 int MPI_Errhandler_free(MPI_Errhandler *errhandler)
```

```
38 MPI_ERRHANDLER_FREE(ERRHANDLER, IERROR)
```

```
39     INTEGER ERRHANDLER, IERROR
```

```
40 {void MPI::Errhandler::Free() (binding deprecated, see Section 15.2) }
```

Marks the error handler associated with `errhandler` for deallocation and sets `errhandler` to `MPI_ERRHANDLER_NULL`. The error handler will be deallocated after all the objects associated with it (communicator, window, or file) have been deallocated.

MPI_ERROR_STRING(errorcode, string, resultlen)

IN	errorcode	Error code returned by an MPI routine
OUT	string	Text that corresponds to the errorcode
OUT	resultlen	Length (in printable characters) of the result returned in string

int MPI_Error_string(int errorcode, char *string, int *resultlen)

MPI_ERROR_STRING(ERRORCODE, STRING, RESULTLEN, IERROR)

INTEGER ERRORCODE, RESULTLEN, IERROR

CHARACTER*(*) STRING

{void MPI::Get_error_string(int errorcode, char* name, int& resultlen)
(binding deprecated, see Section 15.2) }

Returns the error string associated with an error code or class. The argument **string** must represent storage that is at least MPI_MAX_ERROR_STRING characters long.

The number of characters actually written is returned in the output argument, **resultlen**.

Rationale. The form of this function was chosen to make the Fortran and C bindings similar. A version that returns a pointer to a string has two difficulties. First, the return string must be statically allocated and different for each error message (allowing the pointers returned by successive calls to MPI_ERROR_STRING to point to the correct message). Second, in Fortran, a function declared as returning CHARACTER*(*) can not be referenced in, for example, a PRINT statement. *(End of rationale.)*

8.4 Error Codes and Classes

The error codes returned by MPI are left entirely to the implementation (with the exception of MPI_SUCCESS). This is done to allow an implementation to provide as much information as possible in the error code (for use with MPI_ERROR_STRING).

To make it possible for an application to interpret an error code, the routine MPI_ERROR_CLASS converts any error code into one of a small set of standard error codes, called *error classes*. Valid error classes are shown in Table 8.1 and Table 8.2.

The error classes are a subset of the error codes: an MPI function may return an error class number; and the function MPI_ERROR_STRING can be used to compute the error string associated with an error class. An MPI error class is a valid MPI error code. Specifically, the values defined for MPI error classes are valid MPI error codes.

The error codes satisfy,

$$0 = \text{MPI_SUCCESS} < \text{MPI_ERR_...} \leq \text{MPI_ERR_LASTCODE}.$$

Rationale. The difference between MPI_ERR_UNKNOWN and MPI_ERR_OTHER is that MPI_ERROR_STRING can return useful information about MPI_ERR_OTHER.

Note that MPI_SUCCESS = 0 is necessary to be consistent with C practice; the separation of error classes and error codes allows us to define the error classes this way. Having a known LASTCODE is often a nice sanity check as well. *(End of rationale.)*

MPI_SUCCESS	No error
MPI_ERR_BUFFER	Invalid buffer pointer
MPI_ERR_COUNT	Invalid count argument
MPI_ERR_TYPE	Invalid datatype argument
MPI_ERR_TAG	Invalid tag argument
MPI_ERR_COMM	Invalid communicator
MPI_ERR_RANK	Invalid rank
MPI_ERR_REQUEST	Invalid request (handle)
MPI_ERR_ROOT	Invalid root
MPI_ERR_GROUP	Invalid group
MPI_ERR_OP	Invalid operation
MPI_ERR_TOPOLOGY	Invalid topology
MPI_ERR_DIMS	Invalid dimension argument
MPI_ERR_ARG	Invalid argument of some other kind
MPI_ERR_UNKNOWN	Unknown error
MPI_ERR_TRUNCATE	Message truncated on receive
MPI_ERR_OTHER	Known error not in this list
MPI_ERR_INTERN	Internal MPI (implementation) error
MPI_ERR_IN_STATUS	Error code is in status
MPI_ERR_PENDING	Pending request
MPI_ERR_KEYVAL	Invalid keyval has been passed
MPI_ERR_NO_MEM	MPI_ALLOC_MEM failed because memory is exhausted
MPI_ERR_BASE	Invalid base passed to MPI_FREE_MEM
MPI_ERR_INFO_KEY	Key longer than MPI_MAX_INFO_KEY
MPI_ERR_INFO_VALUE	Value longer than MPI_MAX_INFO_VAL
MPI_ERR_INFO_NOKEY	Invalid key passed to MPI_INFO_DELETE
MPI_ERR_SPAWN	Error in spawning processes
MPI_ERR_PORT	Invalid port name passed to MPI_COMM_CONNECT
MPI_ERR_SERVICE	Invalid service name passed to MPI_UNPUBLISH_NAME
MPI_ERR_NAME	Invalid service name passed to MPI_LOOKUP_NAME
MPI_ERR_WIN	Invalid win argument
MPI_ERR_SIZE	Invalid size argument
MPI_ERR_DISP	Invalid disp argument
MPI_ERR_INFO	Invalid info argument
MPI_ERR_LOCKTYPE	Invalid locktype argument
MPI_ERR_ASSERT	Invalid assert argument
MPI_ERR_RMA_CONFLICT	Conflicting accesses to window
MPI_ERR_RMA_SYNC	Wrong synchronization of RMA calls

Table 8.1: Error classes (Part 1)

MPI_ERR_FILE	Invalid file handle	1
MPI_ERR_NOT_SAME	Collective argument not identical on all processes, or collective routines called in a different order by different processes	2
MPI_ERR_AMODE	Error related to the <code>amode</code> passed to <code>MPI_FILE_OPEN</code>	3
MPI_ERR_UNSUPPORTED_DATAREP	Unsupported <code>datarep</code> passed to <code>MPI_FILE_SET_VIEW</code>	4
MPI_ERR_UNSUPPORTED_OPERATION	Unsupported operation, such as seeking on a file which supports sequential access only	5
MPI_ERR_NO_SUCH_FILE	File does not exist	6
MPI_ERR_FILE_EXISTS	File exists	7
MPI_ERR_BAD_FILE	Invalid file name (e.g., path name too long)	8
MPI_ERR_ACCESS	Permission denied	9
MPI_ERR_NO_SPACE	Not enough space	10
MPI_ERR_QUOTA	Quota exceeded	11
MPI_ERR_READ_ONLY	Read-only file or file system	12
MPI_ERR_FILE_IN_USE	File operation could not be completed, as the file is currently open by some process	13
MPI_ERR_DUP_DATAREP	Conversion functions could not be registered because a data representation identifier that was already defined was passed to <code>MPI_REGISTER_DATAREP</code>	14
MPI_ERR_CONVERSION	An error occurred in a user supplied data conversion function.	15
MPI_ERR_IO	Other I/O error	16
MPI_ERR_LASTCODE	Last error code	17

Table 8.2: Error classes (Part 2)

`MPI_ERROR_CLASS(errorcode, errorclass)`

IN	<code>errorcode</code>	Error code returned by an MPI routine
OUT	<code>errorclass</code>	Error class associated with <code>errorcode</code>

`int MPI_Error_class(int errorcode, int *errorclass)`

`MPI_ERROR_CLASS(ERRORCODE, ERRORCLASS, IERROR)`
`INTEGER ERRORCODE, ERRORCLASS, IERROR`

`{int MPI::Get_error_class(int errorcode) (binding deprecated, see Section 15.2) }`

The function `MPI_ERROR_CLASS` maps each standard error code (error class) onto itself.

ticket150.

ticket150.

8.5 Error Classes, Error Codes, and Error Handlers

Users may want to write a layered library on top of an existing MPI implementation, and this library may have its own set of error codes and classes. An example of such a library is an I/O library based on MPI, see Chapter 13 on page 407. For this purpose, functions are needed to:

1. add a new error class to the ones an MPI implementation already knows.
2. associate error codes with this error class, so that `MPI_ERROR_CLASS` works.
3. associate strings with these error codes, so that `MPI_ERROR_STRING` works.
4. invoke the error handler associated with a communicator, window, or object.

Several functions are provided to do this. They are all local. No functions are provided to free error classes or codes: it is not expected that an application will generate them in significant numbers.

`MPI_ADD_ERROR_CLASS(errorclass)`

OUT errorclass value for the new error class (integer)

`int MPI_Add_error_class(int *errorclass)`

`MPI_ADD_ERROR_CLASS(ERRORCLASS, IERROR)`
 INTEGER ERRORCLASS, IERROR

`{int MPI::Add_error_class() (binding deprecated, see Section 15.2) }`

Creates a new error class and returns the value for it.

Rationale. To avoid conflicts with existing error codes and classes, the value is set by the implementation and not by the user. (*End of rationale.*)

Advice to implementors. A high-quality implementation will return the value for a new errorclass in the same deterministic way on all processes. (*End of advice to implementors.*)

Advice to users. Since a call to `MPI_ADD_ERROR_CLASS` is local, the same errorclass may not be returned on all processes that make this call. Thus, it is not safe to assume that registering a new error on a set of processes at the same time will yield the same errorclass on all of the processes. However, if an implementation returns the new errorclass in a deterministic way, and they are always generated in the same order on the same set of processes (for example, all processes), then the value will be the same. However, even if a deterministic algorithm is used, the value can vary across processes. This can happen, for example, if different but overlapping groups of processes make a series of calls. As a result of these issues, getting the “same” error on multiple processes may not cause the same value of error code to be generated. (*End of advice to users.*)

The value of `MPI_ERR_LASTCODE` is a constant value and is not affected by new user-defined error codes and classes. Instead, a predefined attribute key `MPI_LASTUSED` is associated with `MPI_COMM_WORLD`. The attribute value corresponding to this key is the current maximum error class including the user-defined ones. This is a local value and may be different on different processes. The value returned by this key is always greater than or equal to `MPI_ERR_LASTCODE`.

Advice to users. The value returned by the key `MPI_LASTUSED` will not change unless the user calls a function to explicitly add an error class/code. In a multi-threaded environment, the user must take extra care in assuming this value has not changed. Note that error codes and error classes are not necessarily dense. A user may not assume that each error class below `MPI_LASTUSED` is valid. (*End of advice to users.*)

`MPI_ADD_ERROR_CODE(errorclass, errorcode)`

IN	errorclass	error class (integer)
OUT	errorcode	new error code to associated with errorclass (integer)

`int MPI_Add_error_code(int errorclass, int *errorcode)`

`MPI_ADD_ERROR_CODE(ERRORCLASS, ERRORCODE, IERROR)`

`INTEGER ERRORCLASS, ERRORCODE, IERROR`

`{int MPI::Add_error_code(int errorclass) (binding deprecated, see Section 15.2) }`

Creates new error code associated with `errorclass` and returns its value in `errorcode`.

Rationale. To avoid conflicts with existing error codes and classes, the value of the new error code is set by the implementation and not by the user. (*End of rationale.*)

Advice to implementors. A high-quality implementation will return the value for a new `errorcode` in the same deterministic way on all processes. (*End of advice to implementors.*)

`MPI_ADD_ERROR_STRING(errorcode, string)`

IN	errorcode	error code or class (integer)
IN	string	text corresponding to errorcode (string)

`int MPI_Add_error_string(int errorcode, char *string)`

`MPI_ADD_ERROR_STRING(ERRORCODE, STRING, IERROR)`

`INTEGER ERRORCODE, IERROR`

`CHARACTER*(*) STRING`

`{void MPI::Add_error_string(int errorcode, const char* string) (binding deprecated, see Section 15.2) }`

ticket150.
ticket150.

ticket150.
ticket150.

Associates an error string with an error code or class. The string must be no more than MPI_MAX_ERROR_STRING characters long. The length of the string is as defined in the calling language. The length of the string does not include the null terminator in C or C++. Trailing blanks will be stripped in Fortran. Calling MPI_ADD_ERROR_STRING for an errorcode that already has a string will replace the old string with the new string. It is erroneous to call MPI_ADD_ERROR_STRING for an error code or class with a value \leq MPI_ERR_LASTCODE.

If MPI_ERROR_STRING is called when no string has been set, it will return a empty string (all spaces in Fortran, "" in C and C++).

Section 8.3 on page 6 describes the methods for creating and associating error handlers with communicators, files, and windows.

MPI_COMM_CALL_ERRHANDLER (comm, errorcode)

IN	comm	communicator with error handler (handle)
IN	errorcode	error code (integer)

int MPI_Comm_call_errhandler(MPI_Comm comm, int errorcode)

MPI_COMM_CALL_ERRHANDLER(COMM, ERRORCODE, IERROR)

INTEGER COMM, ERRORCODE, IERROR

{void MPI::Comm::Call_errhandler(int errorcode) const *(binding deprecated, see Section 15.2)* }

This function invokes the error handler assigned to the communicator with the error code supplied. This function returns MPI_SUCCESS in C and C++ and the same value in IERROR if the error handler was successfully called (assuming the process is not aborted and the error handler returns).

Advice to users. Users should note that the default error handler is MPI_ERRORS_ARE_FATAL. Thus, calling MPI_COMM_CALL_ERRHANDLER will abort the comm processes if the default error handler has not been changed for this communicator or on the parent before the communicator was created. (*End of advice to users.*)

MPI_WIN_CALL_ERRHANDLER (win, errorcode)

IN	win	window with error handler (handle)
IN	errorcode	error code (integer)

int MPI_Win_call_errhandler(MPI_Win win, int errorcode)

MPI_WIN_CALL_ERRHANDLER(WIN, ERRORCODE, IERROR)

INTEGER WIN, ERRORCODE, IERROR

{void MPI::Win::Call_errhandler(int errorcode) const *(binding deprecated, see Section 15.2)* }

This function invokes the error handler assigned to the window with the error code supplied. This function returns `MPI_SUCCESS` in C and C++ and the same value in `IERROR` if the error handler was successfully called (assuming the process is not aborted and the error handler returns).

Advice to users. As with communicators, the default error handler for windows is `MPI_ERRORS_RETURN`. (*End of advice to users.*)

`MPI_FILE_CALL_ERRHANDLER` (fh, errorcode)

IN fh file with error handler (handle)
IN errorcode error code (integer)

`int MPI_File_call_errhandler(MPI_File fh, int errorcode)`

`MPI_FILE_CALL_ERRHANDLER(FH, ERRORCODE, IERROR)`
`INTEGER FH, ERRORCODE, IERROR`

`{void MPI::File::Call_errhandler(int errorcode) const` (*binding deprecated, see*
Section 15.2) }

This function invokes the error handler assigned to the file with the error code supplied. This function returns `MPI_SUCCESS` in C and C++ and the same value in `IERROR` if the error handler was successfully called (assuming the process is not aborted and the error handler returns).

Advice to users. Unlike errors on communicators and windows, the default behavior for files is to have `MPI_ERRORS_RETURN`. (*End of advice to users.*)

Advice to users. Users are warned that handlers should not be called recursively with `MPI_COMM_CALL_ERRHANDLER`, `MPI_FILE_CALL_ERRHANDLER`, or `MPI_WIN_CALL_ERRHANDLER`. Doing this can create a situation where an infinite recursion is created. This can occur if `MPI_COMM_CALL_ERRHANDLER`, `MPI_FILE_CALL_ERRHANDLER`, or `MPI_WIN_CALL_ERRHANDLER` is called inside an error handler.

Error codes and classes are associated with a process. As a result, they may be used in any error handler. Error handlers should be prepared to deal with any error code they are given. Furthermore, it is good practice to only call an error handler with the appropriate error codes. For example, file errors would normally be sent to the file error handler. (*End of advice to users.*)

8.6 Timers and Synchronization

MPI defines a timer. A timer is specified even though it is not “message-passing,” because timing parallel programs is important in “performance debugging” and because existing timers (both in POSIX 1003.1-1988 and 1003.4D 14.1 and in Fortran 90) are either inconvenient or do not provide adequate access to high-resolution timers. See also Section 2.6.5 on page 22.

1 MPI_WTIME()

2 double MPI_Wtime(void)

3
4 DOUBLE PRECISION MPI_WTIME()

ticket150. 5 {double MPI::Wtime() (*binding deprecated, see Section 15.2*) }

ticket150. 6
7 MPI_WTIME returns a floating-point number of seconds, representing elapsed wall-clock time since some time in the past.

8 The “time in the past” is guaranteed not to change during the life of the process. The user is responsible for converting large numbers of seconds to other units if they are preferred.

9 This function is portable (it returns seconds, not “ticks”), it allows high-resolution, and carries no unnecessary baggage. One would use it like this:

10 {
11 double starttime, endtime;
12 starttime = MPI_Wtime();
13 stuff to be timed ...
14 endtime = MPI_Wtime();
15 printf("That took %f seconds\n",endtime-starttime);
16 }
17

18 The times returned are local to the node that called them. There is no requirement that different nodes return “the same time.” (But see also the discussion of MPI_WTIME_IS_GLOBAL).

19 MPI_WTICK()

20 double MPI_Wtick(void)

21
22 DOUBLE PRECISION MPI_WTICK()

ticket150. 23 {double MPI::Wtick() (*binding deprecated, see Section 15.2*) }

ticket150. 24 MPI_WTICK returns the resolution of MPI_WTIME in seconds. That is, it returns, as a double precision value, the number of seconds between successive clock ticks. For example, if the clock is implemented by the hardware as a counter that is incremented every millisecond, the value returned by MPI_WTICK should be 10^{-3} .

25 8.7 Startup

26 One goal of MPI is to achieve *source code portability*. By this we mean that a program written using MPI and complying with the relevant language standards is portable as written, and must not require any source code changes when moved from one system to another. This explicitly does *not* say anything about how an MPI program is started or launched from the command line, nor what the user must do to set up the environment in which an MPI program will run. However, an implementation may require some setup to be performed

before other MPI routines may be called. To provide for this, MPI includes an initialization routine MPI_INIT.

MPI_INIT()

```
int MPI_Init(int *argc, char ***argv)
```

```
MPI_INIT(IERROR)
```

```
INTEGER IERROR
```

```
{void MPI::Init(int& argc, char**& argv) (binding deprecated, see Section 15.2) }
```

```
{void MPI::Init() (binding deprecated, see Section 15.2) }
```

This routine must be called before any other MPI routine. It must be called at most once; subsequent calls are erroneous (see MPI_INITIALIZED).

[All MPI programs must contain a call to MPI_INIT; this routine must be called before any other MPI routine (apart from MPI-2.1 round-two - begin of modification MPI_INITIALIZED) MPI_GET_VERSION, MPI_INITIALIZED, and MPI_FINALIZED) MPI-2.1 round-two - end of modification is called. The version for MPI-2.1 Correction due to Reviews to MPI-2.1 draft Feb.23, 2008 ANSI C ISO C MPI-2.1 End of review based correction accepts the argc and argv that are provided by the arguments to main:]All MPI programs must contain exactly one call to an MPI initialization routine: MPI_INIT or MPI_INIT_THREAD. Subsequent calls to any initialization routines are erroneous. The only MPI functions that may be invoked before the MPI initialization routines are called are MPI_GET_VERSION, MPI_INITIALIZED, and MPI_FINALIZED. The version for ISO C accepts the argc and argv that are provided by the arguments to main or NULL:

```
int main([ticket60.]int argc, [ticket60.]char **argv)
```

```
[ticket60.][int argc;char **argv;]{
```

```
    MPI_Init(&argc, &argv);
```

```
    /* parse arguments */
```

```
    /* main program */
```

```
    MPI_Finalize();    /* see below */
```

```
}
```

The Fortran version takes only IERROR.

Conforming implementations of MPI are required to allow applications to pass NULL for both the argc and argv arguments of main in C and C++. In C++, there is an alternative binding for MPI::Init that does not have these arguments at all.

Rationale. In some applications, libraries may be making the call to MPI_Init, and may not have access to argc and argv from main. It is anticipated that applications requiring special information about the environment or information supplied by mpiexec can get that information from environment variables. (*End of rationale.*)

MPI_FINALIZE()

```

1  int MPI_Finalize(void)
2  MPI_FINALIZE(IERROR)
3  INTEGER IERROR

```

```

5  {void MPI::Finalize() (binding deprecated, see Section 15.2) }

```

This routine cleans up all MPI state. Each process must call MPI_FINALIZE before it exits. Unless there has been a call to MPI_ABORT, each process must ensure that all pending **[non-blocking]nonblocking** communications are (locally) complete before calling MPI_FINALIZE. Further, at the instant at which the last process calls MPI_FINALIZE, all pending sends must be matched by a receive, and all pending receives must be matched by a send.

For example, the following program is correct:

Process 0	Process 1
-----	-----
MPI_Init();	MPI_Init();
MPI_Send(dest=1);	MPI_Recv(src=0);
MPI_Finalize();	MPI_Finalize();

Without the matching receive, the program is erroneous:

Process 0	Process 1
-----	-----
MPI_Init();	MPI_Init();
MPI_Send (dest=1);	
MPI_Finalize();	MPI_Finalize();

A successful return from a blocking communication operation or from MPI_WAIT or MPI_TEST tells the user that the buffer can be reused and means that the communication is completed by the user, but does not guarantee that the local process has no more work to do. A successful return from MPI_REQUEST_FREE with a request handle generated by an MPI_ISEND nullifies the handle but provides no assurance of operation completion. The MPI_ISEND is complete only when it is known by some means that a matching receive has completed. MPI_FINALIZE guarantees that all local actions required by communications the user has completed will, in fact, occur before it returns.

MPI_FINALIZE guarantees nothing about pending communications that have not been completed (completion is assured only by MPI_WAIT, MPI_TEST, or MPI_REQUEST_FREE combined with some other verification of completion).

Example 8.3 This program is correct:

rank 0	rank 1
=====	=====
...	...
MPI_Isend();	MPI_Recv();
MPI_Request_free();	MPI_Barrier();
MPI_Barrier();	MPI_Finalize();
MPI_Finalize();	exit();
exit();	

Example 8.4 This program is erroneous and its behavior is undefined:

```

rank 0                                     rank 1
=====
...
MPI_Isend();                             MPI_Recv();
MPI_Request_free();                      MPI_Finalize();
MPI_Finalize();                          exit();
exit();

```

If no `MPI_BUFFER_DETACH` occurs between an `MPI_BSEND` (or other buffered send) and `MPI_FINALIZE`, the `MPI_FINALIZE` implicitly supplies the `MPI_BUFFER_DETACH`.

Example 8.5 This program is correct, and after the `MPI_Finalize`, it is as if the buffer had been detached.

```

rank 0                                     rank 1
=====
...
buffer = malloc(1000000);                MPI_Recv();
MPI_Buffer_attach();                     MPI_Finalize();
MPI_Bsend();                             exit();
MPI_Finalize();
free(buffer);
exit();

```

Example 8.6 In this example, `MPI_Iprobe()` must return a `FALSE` flag. `MPI_Test_cancelled()` must return a `TRUE` flag, independent of the relative order of execution of `MPI_Cancel()` in process 0 and `MPI_Finalize()` in process 1.

The `MPI_Iprobe()` call is there to make sure the implementation knows that the “tag1” message exists at the destination, without being able to claim that the user knows about it.

```

rank 0                                     rank 1
=====
MPI_Init();                             MPI_Init();
MPI_Isend(tag1);                         MPI_Barrier();
MPI_Barrier();                          MPI_Iprobe(tag2);
MPI_Barrier();                          MPI_Barrier();
                                         MPI_Finalize();
                                         exit();

MPI_Cancel();
MPI_Wait();
MPI_Test_cancelled();
MPI_Finalize();
exit();

```

Advice to implementors. An implementation may need to delay the return from `MPI_FINALIZE` until all potential future message cancellations have been processed. One possible solution is to place a barrier inside `MPI_FINALIZE` (*End of advice to implementors.*)

Once `MPI_FINALIZE` returns, no MPI routine (not even `MPI_INIT`) may be called, except for `MPI_GET_VERSION`, `MPI_INITIALIZED`, and `MPI_FINALIZED`. Each process must complete any pending communication it initiated before it calls `MPI_FINALIZE`. If the call returns, each process may continue local computations, or exit, without participating in further MPI communication with other processes. `MPI_FINALIZE` is collective over all connected processes. If no processes were spawned, accepted or connected then this means over `MPI_COMM_WORLD`; otherwise it is collective over the union of all processes that have been and continue to be connected, as explained in Section 10.5.4 on page 346.

Advice to implementors. Even though a process has completed all the communication it initiated, such communication may not yet be completed from the viewpoint of the underlying MPI system. E.g., a blocking send may have completed, even though the data is still buffered at the sender. The MPI implementation must ensure that a process has completed any involvement in MPI communication before `MPI_FINALIZE` returns. Thus, if a process exits after the call to `MPI_FINALIZE`, this will not cause an ongoing communication to fail. (*End of advice to implementors.*)

Although it is not required that all processes return from `MPI_FINALIZE`, it is required that at least process 0 in `MPI_COMM_WORLD` return, so that users can know that the MPI portion of the computation is over. In addition, in a POSIX environment, they may desire to supply an exit code for each process that returns from `MPI_FINALIZE`.

Example 8.7 The following illustrates the use of requiring that at least one process return and that it be known that process 0 is one of the processes that return. One wants code like the following to work no matter how many processes return.

```

...
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
...
MPI_Finalize();
if (myrank == 0) {
    resultfile = fopen("outfile","w");
    dump_results(resultfile);
    fclose(resultfile);
}
exit(0);

```

`MPI_INITIALIZED(flag)`

OUT	flag	Flag is true if <code>MPI_INIT</code> has been called and false otherwise.
-----	------	----------------------------------------------------------------------------

```

int MPI_Initialized(int *flag)

```

MPI_INITIALIZED(FLAG, IERROR)
 LOGICAL FLAG
 INTEGER IERROR

{bool MPI::Is_initialized() *(binding deprecated, see Section 15.2)*}

This routine may be used to determine whether MPI_INIT has been called. MPI_INITIALIZED returns true if the calling process has called MPI_INIT. Whether MPI_FINALIZE has been called does not affect the behavior of MPI_INITIALIZED. It is one of the few routines that may be called before MPI_INIT is called.

MPI_ABORT(comm, errorcode)

IN comm communicator of tasks to abort
 IN errorcode error code to return to invoking environment

int MPI_Abort(MPI_Comm comm, int errorcode)

MPI_ABORT(COMM, ERRORCODE, IERROR)
 INTEGER COMM, ERRORCODE, IERROR

{void MPI::Comm::Abort(int errorcode) *(binding deprecated, see Section 15.2)*}

This routine makes a “best attempt” to abort all tasks in the group of comm. This function does not require that the invoking environment take any action with the error code. However, a Unix or POSIX environment should handle this as a **return errorcode** from the main program.

It may not be possible for an MPI implementation to abort only the processes represented by comm if this is a subset of the processes. In this case, the MPI implementation should attempt to abort all the connected processes but should not abort any unconnected processes. If no processes were spawned, accepted or connected then this has the effect of aborting all the processes associated with MPI_COMM_WORLD.

Rationale. The communicator argument is provided to allow for future extensions of MPI to environments with, for example, dynamic process management. In particular, it allows but does not require an MPI implementation to abort a subset of MPI_COMM_WORLD. *(End of rationale.)*

Advice to users. Whether the errorcode is returned from the executable or from the MPI process startup mechanism (e.g., **mpiexec**), is an aspect of quality of the MPI library but not mandatory. *(End of advice to users.)*

Advice to implementors. Where possible, a high-quality implementation will try to return the errorcode from the MPI process startup mechanism (e.g. **mpiexec** or singleton init). *(End of advice to implementors.)*

8.7.1 Allowing User Functions at Process Termination

There are times in which it would be convenient to have actions happen when an MPI process finishes. For example, a routine may do initializations that are useful until the MPI job (or

that part of the job that being terminated in the case of dynamically created processes) is finished. This can be accomplished in MPI by attaching an attribute to MPI_COMM_SELF with a callback function. When MPI_FINALIZE is called, it will first execute the equivalent of an MPI_COMM_FREE on MPI_COMM_SELF. This will cause the delete callback function to be executed on all keys associated with MPI_COMM_SELF, *[in an arbitrary order]in the reverse order that they were set on MPI_COMM_SELF*. If no key has been attached to MPI_COMM_SELF, then no callback is invoked. The “freeing” of MPI_COMM_SELF occurs before any other parts of MPI are affected. Thus, for example, calling MPI_FINALIZED will return false in any of these callback functions. Once done with MPI_COMM_SELF, the order and rest of the actions taken by MPI_FINALIZE is not specified.

Advice to implementors. Since attributes can be added from any supported language, the MPI implementation needs to remember the creating language so the correct callback is made. *Implementations that use the attribute delete callback on MPI_COMM_SELF internally should register their internal callbacks before returning from MPI_INIT / MPI_INIT_THREAD, so that libraries or applications will not have portions of the MPI implementation shut down before the application-level callbacks are made.* (End of advice to implementors.)

8.7.2 Determining Whether MPI Has Finished

One of the goals of MPI was to allow for layered libraries. In order for a library to do this cleanly, it needs to know if MPI is active. In MPI the function MPI_INITIALIZED was provided to tell if MPI had been initialized. The problem arises in knowing if MPI has been finalized. Once MPI has been finalized it is no longer active and cannot be restarted. A library needs to be able to determine this to act accordingly. To achieve this the following function is needed:

MPI_FINALIZED(flag)

OUT flag true if MPI was finalized (logical)

int MPI_Finalized(int *flag)

MPI_FINALIZED(FLAG, IERROR)

LOGICAL FLAG

INTEGER IERROR

{bool MPI::Is_finalized() *(binding deprecated, see Section 15.2)* }

This routine returns true if MPI_FINALIZE has completed. It is legal to call MPI_FINALIZED before MPI_INIT and after MPI_FINALIZE.

Advice to users. MPI is “active” and it is thus safe to call MPI functions if MPI_INIT has completed and MPI_FINALIZE has not completed. If a library has no other way of knowing whether MPI is active or not, then it can use MPI_INITIALIZED and MPI_FINALIZED to determine this. For example, MPI is “active” in callback functions that are invoked during MPI_FINALIZE. (End of advice to users.)

8.8 Portable MPI Process Startup

A number of implementations of MPI provide a startup command for MPI programs that is of the form

```
mpirun <mpirun arguments> <program> <program arguments>
```

Separating the command to start the program from the program itself provides flexibility, particularly for network and heterogeneous implementations. For example, the startup script need not run on one of the machines that will be executing the MPI program itself.

Having a standard startup mechanism also extends the portability of MPI programs one step further, to the command lines and scripts that manage them. For example, a validation suite script that runs hundreds of programs can be a portable script if it is written using such a standard startup mechanism. In order that the “standard” command not be confused with existing practice, which is not standard and not portable among implementations, instead of `mpirun` MPI specifies `mpiexec`.

While a standardized startup mechanism improves the usability of MPI, the range of environments is so diverse (e.g., there may not even be a command line interface) that MPI cannot mandate such a mechanism. Instead, MPI specifies an `mpiexec` startup command and recommends but does not require it, as advice to implementors. However, if an implementation does provide a command called `mpiexec`, it must be of the form described below.

It is suggested that

```
mpiexec -n <numprocs> <program>
```

be at least one way to start `<program>` with an initial `MPI_COMM_WORLD` whose group contains `<numprocs>` processes. Other arguments to `mpiexec` may be implementation-dependent.

Advice to implementors. Implementors, if they do provide a special startup command for MPI programs, are advised to give it the following form. The syntax is chosen in order that `mpiexec` be able to be viewed as a command-line version of `MPI_COMM_SPAWN` (See Section 10.3.4).

Analogous to `MPI_COMM_SPAWN`, we have

```
mpiexec -n    <maxprocs>
          -soft <      >
          -host <      >
          -arch <      >
          -wdir <      >
          -path <      >
          -file <      >
          ...
          <command line>
```

for the case where a single command line for the application program and its arguments will suffice. See Section 10.3.4 for the meanings of these arguments. For the case corresponding to `MPI_COMM_SPAWN_MULTIPLE` there are two possible formats:

Form A:

```
1      mpiexec { <above arguments> } : { ... } : { ... } : ... : { ... }
```

2
3 As with MPI_COMM_SPAWN, all the arguments are optional. (Even the `-n x` argu-
4 ment is optional; the default is implementation dependent. It might be 1, it might be
5 taken from an environment variable, or it might be specified at compile time.) The
6 names and meanings of the arguments are taken from the keys in the `info` argument
7 to MPI_COMM_SPAWN. There may be other, implementation-dependent arguments
8 as well.

9 Note that Form A, though convenient to type, prevents colons from being program
10 arguments. Therefore an alternate, file-based form is allowed:

11 Form B:

```
12  
13      mpiexec -configfile <filename>
```

14
15 where the lines of `<filename>` are of the form separated by the colons in Form A.
16 Lines beginning with `#` are comments, and lines may be continued by terminating
17 the partial line with `\`.
18

19 **Example 8.8** Start 16 instances of `myprog` on the current or default machine:

```
20  
21      mpiexec -n 16 myprog
```

22
23 **Example 8.9** Start 10 processes on the machine called `ferrari`:

```
24  
25      mpiexec -n 10 -host ferrari myprog
```

26
27 **Example 8.10** Start three copies of the same program with different command-line
28 arguments:

```
29      mpiexec myprog infile1 : myprog infile2 : myprog infile3
```

30
31 **Example 8.11** Start the `ocean` program on five Suns and the `atmos` program on 10
32 RS/6000's:

```
33  
34      mpiexec -n 5 -arch sun ocean : -n 10 -arch rs6000 atmos
```

35 It is assumed that the implementation in this case has a method for choosing hosts of
36 the appropriate type. Their ranks are in the order specified.
37

38 **Example 8.12** Start the `ocean` program on five Suns and the `atmos` program on 10
39 RS/6000's (Form B):

```
40  
41      mpiexec -configfile myfile
```

42 where `myfile` contains

```
43  
44      -n 5  -arch sun    ocean  
45      -n 10 -arch rs6000 atmos
```

46
47 (*End of advice to implementors.*)
48

Index

CONST:_WORLD, 6
CONST:argc, 21
CONST:argv, 21
CONST:CHARACTER*(*), 13
CONST:COMM, 7
CONST:errorcode, 13, 15
CONST:FILE, 7
CONST:IERROR, 21
CONST:LASTCODE, 13
CONST:MPI::Errhandler, 8, 9–12
CONST:MPI::ERRORS_THROW_EXCEPTIONS, 7
CONST:MPI::File, 11, 12, 19
CONST:MPI::Info, 4
CONST:MPI::Win, 10, 11, 18
CONST:MPI_ANY_SOURCE, 3
CONST:MPI_BSEND_OVERHEAD, 4
CONST:MPI_COMM-, 6
CONST:MPI_COMM_SELF, 26
CONST:MPI_COMM_WORLD, 2, 3, 6, 9, 17, 24, 25, 27
CONST:MPI_ERR_..., 13
CONST:MPI_ERR_ACCESS, 15
CONST:MPI_ERR_AMODE, 15
CONST:MPI_ERR_ARG, 14
CONST:MPI_ERR_ASSERT, 14
CONST:MPI_ERR_BAD_FILE, 15
CONST:MPI_ERR_BASE, 5, 14
CONST:MPI_ERR_BUFFER, 14
CONST:MPI_ERR_COMM, 14
CONST:MPI_ERR_CONVERSION, 15
CONST:MPI_ERR_COUNT, 14
CONST:MPI_ERR_DIMS, 14
CONST:MPI_ERR_DISP, 14
CONST:MPI_ERR_DUP_DATAREP, 15
CONST:MPI_ERR_FILE, 15
CONST:MPI_ERR_FILE_EXISTS, 15
CONST:MPI_ERR_FILE_IN_USE, 15
CONST:MPI_ERR_GROUP, 14
CONST:MPI_ERR_IN_STATUS, 8, 14
CONST:MPI_ERR_INFO, 14
CONST:MPI_ERR_INFO_KEY, 14
CONST:MPI_ERR_INFO_NOKEY, 14
CONST:MPI_ERR_INFO_VALUE, 14
CONST:MPI_ERR_INTERN, 14
CONST:MPI_ERR_IO, 15
CONST:MPI_ERR_KEYVAL, 14
CONST:MPI_ERR_LASTCODE, 13, 15, 17, 18
CONST:MPI_ERR_LOCKTYPE, 14
CONST:MPI_ERR_NAME, 14
CONST:MPI_ERR_NO_MEM, 5, 14
CONST:MPI_ERR_NO_SPACE, 15
CONST:MPI_ERR_NO_SUCH_FILE, 15
CONST:MPI_ERR_NOT_SAME, 15
CONST:MPI_ERR_OP, 14
CONST:MPI_ERR_OTHER, 13, 14
CONST:MPI_ERR_PENDING, 14
CONST:MPI_ERR_PORT, 14
CONST:MPI_ERR_QUOTA, 15
CONST:MPI_ERR_RANK, 14
CONST:MPI_ERR_READ_ONLY, 15
CONST:MPI_ERR_REQUEST, 14
CONST:MPI_ERR_RMA_CONFLICT, 14
CONST:MPI_ERR_RMA_SYNC, 14
CONST:MPI_ERR_ROOT, 14
CONST:MPI_ERR_SERVICE, 14
CONST:MPI_ERR_SIZE, 14
CONST:MPI_ERR_SPAWN, 14
CONST:MPI_ERR_TAG, 14
CONST:MPI_ERR_TOPOLOGY, 14
CONST:MPI_ERR_TRUNCATE, 14
CONST:MPI_ERR_TYPE, 14
CONST:MPI_ERR_UNKNOWN, 13, 14
CONST:MPI_ERR_UNSUPPORTED_DATAREP, 15
CONST:MPI_ERR_UNSUPPORTED_OPERATION, 15
CONST:MPI_ERR_WIN, 14
CONST:MPI_Errhandler, 8, 9–12

- 1 CONST:MPI_ERRHANDLER_NULL, [12](#)
- 2 CONST:MPI_ERROR_STRING, [13](#)
- 3 CONST:MPI_ERRORS_ARE_FATAL, [6](#), [7](#),
- 4 [18](#), [19](#)
- 5 CONST:MPI_ERRORS_RETURN, [6](#), [7](#), [19](#)
- 6 CONST:MPI_File, [11](#), [12](#), [19](#)
- 7 CONST:MPI_HOST, [2](#)
- 8 CONST:MPI_Info, [4](#)
- 9 CONST:MPI_IO, [2](#), [3](#)
- 10 CONST:MPI_LASTUSED_CODE, [17](#)
- 11 CONST:MPI_MAX_ERROR_STRING, [13](#),
- 12 [18](#)
- 13 CONST:MPI_MAX_INFO_KEY, [14](#)
- 14 CONST:MPI_MAX_INFO_VAL, [14](#)
- 15 CONST:MPI_MAX_PROCESSOR_NAME, [4](#)
- 16 CONST:MPI_PROC_NULL, [2](#), [3](#)
- 17 CONST:MPI_SUBVERSION, [2](#)
- 18 CONST:MPI_SUCCESS, [13](#), [14](#), [18](#), [19](#)
- 19 CONST:MPI_TAG_UB, [2](#)
- 20 CONST:MPI_VERSION, [2](#)
- 21 CONST:MPI_Win, [10](#), [11](#), [18](#)
- 22 CONST:MPI_WTIME_IS_GLOBAL, [2](#), [3](#), [20](#)
- 23 CONST:NULL, [21](#)
- 24 CONST:PRINT, [13](#)
- 25 CONST:string, [13](#)
- 26 CONST:void*, [5](#)
- 27 CONST:WIN, [7](#)
- 28 CONST:XXX, [7](#)
- 29
- 30 EXAMPLES:MPI_ALLOC_MEM, [5](#)
- 31 EXAMPLES:MPI_Alloc_mem, [6](#)
- 32 EXAMPLES:MPI_Barrier, [22](#), [23](#)
- 33 EXAMPLES:MPI_Buffer_attach, [23](#)
- 34 EXAMPLES:MPI_Cancel, [23](#)
- 35 EXAMPLES:MPI_Finalize, [22-24](#)
- 36 EXAMPLES:MPI_FREE_MEM, [5](#)
- 37 EXAMPLES:MPI_Iprobe, [23](#)
- 38 EXAMPLES:MPI_Request_free, [22](#), [23](#)
- 39 EXAMPLES:MPI_Test_cancelled, [23](#)
- 40 EXAMPLES:mpiexec, [28](#)
- 41
- 42 MPI_ABORT, [6](#), [22](#)
- 43 MPI_ABORT(comm, errorcode), [25](#)
- 44 MPI_ADD_ERROR_CLASS, [16](#)
- 45 MPI_ADD_ERROR_CLASS(errorclass), [16](#)
- 46 MPI_ADD_ERROR_CODE(errorclass, error-
- 47 code), [17](#)
- 48 MPI_ADD_ERROR_STRING, [18](#)
- MPI_ADD_ERROR_STRING(errorcode, string),
- [17](#)
- MPI_ALLOC_MEM, [5](#), [14](#)
- MPI_ALLOC_MEM(size, info, baseptr), [4](#)
- MPI_ATTR_GET, [2](#)
- MPI_BSEND, [4](#), [23](#)
- MPI_BUFFER_DETACH, [23](#)
- MPI_COMM_CALL_ERRHANDLER, [19](#)
- MPI_COMM_CALL_ERRHANDLER(comm,
- errorcode), [18](#)
- MPI_COMM_CONNECT, [14](#)
- MPI_COMM_CREATE_ERRHANDLER, [7](#),
- [9](#)
- MPI_COMM_CREATE_ERRHANDLER(function,
- errhandler), [8](#)
- MPI_COMM_FREE, [26](#)
- MPI_COMM_GET_ATTR, [2](#)
- MPI_COMM_GET_ERRHANDLER, [7](#)
- MPI_COMM_GET_ERRHANDLER(comm,
- errhandler), [9](#)
- MPI_COMM_GROUP, [7](#)
- MPI_COMM_SET_ERRHANDLER, [7](#)
- MPI_COMM_SET_ERRHANDLER(comm, er-
- rhandler), [9](#)
- MPI_COMM_SPAWN, [27](#), [28](#)
- MPI_COMM_SPAWN_MULTIPLE, [27](#)
- MPI_ERRHANDLER_CREATE, [8](#)
- MPI_ERRHANDLER_FREE, [7](#)
- MPI_ERRHANDLER_FREE(errhandler),
- [12](#)
- MPI_ERRHANDLER_GET, [7](#), [9](#)
- MPI_ERRHANDLER_SET, [9](#)
- MPI_ERROR_CLASS, [13](#), [15](#), [16](#)
- MPI_ERROR_CLASS(errorcode, errorclass
-), [15](#)
- MPI_ERROR_STRING, [13](#), [16](#), [18](#)
- MPI_ERROR_STRING(errorcode, string, re-
- sultlen), [13](#)
- MPI_FILE_CALL_ERRHANDLER, [19](#)
- MPI_FILE_CALL_ERRHANDLER(fh, er-
- rorcode), [19](#)
- MPI_FILE_CREATE_ERRHANDLER, [7](#), [12](#)
- MPI_FILE_CREATE_ERRHANDLER(function,
- errhandler), [11](#)
- MPI_FILE_GET_ERRHANDLER, [7](#)
- MPI_FILE_GET_ERRHANDLER(file, errhan-
- dler), [12](#)
- MPI_FILE_OPEN, [15](#)

MPI_FILE_SET_ERRHANDLER, 7	mpiexec, 21 , 25 , 27 , 27	1
MPI_FILE_SET_ERRHANDLER(file, errhan-	mpirun, 27	2
dler), 12		3
MPI_FILE_SET_VIEW, 15	TYPEDEF:MPI_Comm_errhandler_function(MPI_Comm	4
MPI_FINALIZE, 2 , 22–26	*, int *, ...), 8	5
MPI_FINALIZE(), 21	TYPEDEF:MPI_File_errhandler_function(MPI_File	6
MPI_FINALIZED, 21 , 24 , 26	*, int *, ...), 11	7
MPI_FINALIZED(flag), 26	TYPEDEF:MPI_Win_errhandler_function(MPI_Win	8
MPI_FREE_MEM, 5 , 14	*, int *, ...), 10	9
MPI_FREE_MEM(base), 5		10
MPI_GET_PROCESSOR_NAME, 4		11
MPI_GET_PROCESSOR_NAME(name, re-		12
sultlen), 3		13
MPI_GET_VERSION, 2 , 21 , 24		14
MPI_GET_VERSION(version, subversion),		15
1		16
MPI_GROUP_FREE, 7		17
MPI_INFO_DELETE, 14		18
MPI_INIT, 2 , 21 , 24–26		19
MPI_INIT(), 21		20
MPI_INIT_THREAD, 21 , 26		21
MPI_INITIALIZED, 21 , 24–26		22
MPI_INITIALIZED(flag), 24		23
MPI_ISEND, 22		24
MPI_LOOKUP_NAME, 14		25
MPI_REGISTER_DATAREP, 15		26
MPI_REQUEST_FREE, 22		27
MPI_TEST, 22		28
MPI_UNPUBLISH_NAME, 14		29
MPI_WAIT, 22		30
MPI_WIN_CALL_ERRHANDLER, 19		31
MPI_WIN_CALL_ERRHANDLER (win, er-		32
rorcode), 18		33
MPI_WIN_CREATE_ERRHANDLER, 7 , 10		34
MPI_WIN_CREATE_ERRHANDLER(function,		35
errhandler), 10		36
MPI_WIN_GET_ERRHANDLER, 7		37
MPI_WIN_GET_ERRHANDLER(win, errhan-		38
dler), 11		39
MPI_WIN_LOCK, 4		40
MPI_WIN_SET_ERRHANDLER, 7		41
MPI_WIN_SET_ERRHANDLER(win, errhan-		42
dler), 10		43
MPI_WIN_UNLOCK, 4		44
MPI_WTICK, 20		45
MPI_WTICK(), 20		46
MPI_WTIME, 3 , 20		47
MPI_WTIME(), 20		48