

`MPI_REDUCE_LOCAL( inbuf, inoutbuf, count, datatype, op)`

IN	inbuf	input buffer (choice)
INOUT	inoutbuf	combined input and output buffer (choice)
IN	count	number of elements in inbuf and inoutbuf buffers (non-negative integer)
IN	datatype	data type of elements of inbuf and inoutbuf buffers (handle)
IN	op	operation (handle)

```
int MPI_Reduce_local(void *inbuf, void *inoutbuf, int count,
                    MPI_Datatype datatype, MPI_Op op)
```

```
MPI_REDUCE_LOCAL(INBUF, INOUBUF, COUNT, DATATYPE, OP, IERROR)
    <type> INBUF(*), INOUBUF(*)
    INTEGER COUNT, DATATYPE, OP, IERROR
```

```
{void MPI::Op::Reduce_local(const void *inbuf, void *inoutbuf, int count,
    const MPI::Datatype& datatype) const (binding deprecated, see
    Section ??) }
```

The function applies the operation given by `op` element-wise to the elements of `inbuf` and `inoutbuf` with the result stored element-wise in `inoutbuf`, as explained for user-defined operations in Section 5.9.5. Both `inbuf` and `inoutbuf` (input as well as result) have the same number of elements given by `count` and the same datatype given by `datatype`. The `MPI_IN_PLACE` option is not allowed.

Reduction operations can be queried for their commutativity.

`MPI_OP_COMMUTATIVE( op, commute)`

IN	op	operation (handle)
OUT	commute	true if op is commutative, false otherwise (logical)

```
int MPI_Op_commutative(MPI_Op op, int *commute)
```

```
MPI_OP_COMMUTATIVE(OP, COMMUTE, IERROR)
    LOGICAL COMMUTE
    INTEGER OP, IERROR
```

```
{bool MPI::Op::Is_commutative() const (binding deprecated, see Section ??) }
```

## 5.10 Reduce-Scatter

[MPI includes a variant of the reduce operations where the result is scattered to all processes in a group on return. ]MPI includes variants of the reduce operations where the result is scattered to all processes in a group on return. One variant scatters equal-sized blocks to all processes, while another variant scatters blocks that may vary in size for each process.

## 5.10.1 MPI\_Reduce\_scatter\_block

MPI\_REDUCE\_SCATTER\_BLOCK( sendbuf, recvbuf, recvcnt, datatype, op, comm)

IN	sendbuf	starting address of send buffer (choice)
OUT	recvbuf	starting address of receive buffer (choice)
IN	recvcnt	element count per block (non-negative integer)
IN	datatype	data type of elements of send and receive buffers (handle)
IN	op	operation (handle)
IN	comm	communicator (handle)

```
int MPI_Reduce_scatter_block(void *sendbuf, void *recvbuf, int recvcnt,
                             MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
```

```
MPI_REDUCE_SCATTER_BLOCK(SENDBUF, RECVBUF, REVCOUNT, DATATYPE, OP, COMM,
                           IERROR)
```

```
<type> SENDBUF(*), RECVBUF(*)
```

```
INTEGER REVCOUNT, DATATYPE, OP, COMM, IERROR
```

```
{void MPI::Comm::Reduce_scatter_block(const void *sendbuf, void *recvbuf,
int recvcnt, const MPI::Datatype& datatype,
const MPI::Op& op) const = 0 (binding deprecated, see Section ??) }
```

If `comm` is an intracommunicator, `MPI_REDUCE_SCATTER_BLOCK` first performs a global, element-wise reduction on vectors of `count = n*recvcnt` elements in the send buffers defined by `sendbuf`, `count` and `datatype`, using the operation `op`, where `n` is the number of processes in the group of `comm`. The routine is called by all group members using the same arguments for `recvcnt`, `datatype`, `op` and `comm`. The resulting vector is treated as `n` consecutive blocks of `recvcnt` elements that are scattered to the processes of the group. The `i`-th block is sent to process `i` and stored in the receive buffer defined by `recvbuf`, `recvcnt`, and `datatype`.

*Advice to implementors.* The `MPI_REDUCE_SCATTER_BLOCK` routine is functionally equivalent to: an `MPI_REDUCE` collective operation with `count` equal to `recvcnt*n`, followed by an `MPI_SCATTER` with `sendcount` equal to `recvcnt`. However, a direct implementation may run faster. (*End of advice to implementors.*)

The “in place” option for intracommunicators is specified by passing `MPI_IN_PLACE` in the `sendbuf` argument on *all* processes. In this case, the input data is taken from the receive buffer.

If `comm` is an intercommunicator, then the result of the reduction of the data provided by processes in one group (group A) is scattered among processes in the other group (group B) and vice versa. Within each group, all processes provide the same value for the `recvcnt` argument, and provide input vectors of `count = n*recvcnt` elements stored in the send buffers, where `n` is the size of the group. The number of elements `count` must be the same for the two groups. The resulting vector from the other group is scattered in blocks of `recvcnt` elements among the processes in the group.

*Rationale.* The last restriction is needed so that the length of the send buffer of one group can be determined by the local `recvcount` argument of the other group. Otherwise, a communication is needed to figure out how many elements are reduced. (*End of rationale.*)

### 5.10.2 MPI\_Reduce\_scatter

`MPI_REDUCE_SCATTER` extends the functionality of `MPI_REDUCE_SCATTER_BLOCK` such that the scattered blocks can vary in size. Block sizes are determined by the `recvcounts` array, such that the *i*-th block contains `recvcounts[i]` elements.

`MPI_REDUCE_SCATTER( sendbuf, recvbuf, recvcounts, datatype, op, comm)`

IN	sendbuf	starting address of send buffer (choice)
OUT	recvbuf	starting address of receive buffer (choice)
IN	recvcounts	non-negative integer array (of length group size) specifying the number of elements [in]of the result distributed to each process. [Array must be identical on all calling processes.]
IN	datatype	data type of elements of [input buffer]send and receive buffers (handle)
IN	op	operation (handle)
IN	comm	communicator (handle)

```
int MPI_Reduce_scatter(void* sendbuf, void* recvbuf, int *recvcounts,
                      MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
```

```
MPI_REDUCE_SCATTER(SENDBUF, RECVBUF, RECVCOUNTS, DATATYPE, OP, COMM,
                   IERROR)
<type> SENDBUF(*), RECVBUF(*)
INTEGER REVCOUNTS(*), DATATYPE, OP, COMM, IERROR
```

```
{void MPI::Comm::Reduce_scatter(const void* sendbuf, void* recvbuf,
                               int recvcounts[], const MPI::Datatype& datatype,
                               const MPI::Op& op) const = 0 (binding deprecated, see Section ??) }
```

If `comm` is an intracommunicator, `MPI_REDUCE_SCATTER` first [does an element-wise reduction on vector of count =  $\sum_i \text{recvcounts}[i]$  elements in the send buffer defined by `sendbuf`, count and `datatype`. Next, the resulting vector of results is split into *n* disjoint segments, where *n* is the number of members in the group. Segment *i* contains `recvcounts[i]` elements. The *i*-th segment ]performs a global, element-wise reduction on vectors of count =  $\sum_{i=0}^{n-1} \text{recvcounts}[i]$  elements in the send buffers defined by `sendbuf`, count and `datatype`, using the operation `op`, where *n* is the number of processes in the group of `comm`. The routine is called by all group members using the same arguments for `recvcounts`, `datatype`, `op` and `comm`. The resulting vector is treated as *n* consecutive blocks where the number of elements of the *i*-th block is `recvcounts[i]`. The blocks are scattered to the processes of the group.