

MPI3 Hybrid Programming – Proposal for Helper Threads

June 14 Meeting notes: The following items need further discussion

- 1 Error handling/reporting: how do various errors get reported and handled – are errors thrown to MPI_COMM_WORLD? Or to the LEAVE call? Or ??
 - 1.1 The proposal has each function returning error codes, but there is also a need to define how error handlers are used with this.
 - 1.1.1 Where does one register a handler for helper threads calls?
- 2 Stacking Libraries – nesting. How would JOIN/LEAVE be used with libraries?
 - 2.1 All (non-helper) MPI calls are used the same regardless of whether inside a JOIN/LEAVE or not, the implementation detects the JOIN/LEAVE case and does things differently internally. This means that a communications-only library need not be aware of whether the caller is within a JOIN/LEAVE or not, provided the caller observes the library's restrictions on threading (i.e. only one thread calls the library if it is single-thread).
 - 2.1.1 What other types of complexity exist in current libraries?
 - 2.1.2 What do we expect library writers to do when using helper threads?
- 3 Leave semantics clarification – is it a (more restrictive) “barrier-like” operation or (less restrictive) “collective”?
 - 3.1 “barrier-like” means that all team members must call LEAVE before any will return from LEAVE.
 - 3.2 “collective” (only) would mean that simply all members must call LEAVE, but there is not requirement that all will be in the LEAVE at the same instant.
 - 3.3 Commentary: I think the distinction comes for implementations or run modes where JOIN/LEAVE are essentially not supported. I think a “productive” LEAVE will actually be barrier-like, although one could imagine a case where the “master” (thread doing MPI calls) reaches LEAVE before others and finds all communications completed, it could continue and when other threads reach LEAVE they simply pass-through. But we should probably ponder this more to be sure that “barrier-like” is not required under some circumstances, or that an implementation is free to use “barrier-like” if needed. Is it acceptable to state that the user cannot expect more than “collective”, but also must tolerate “barrier-like”?
- 4 What about a “leave but stay joined” sort of operation, like a fence.
 - 4.1 The idea being there might be a performance benefit if a program wants to “synchronize” communications at various points between JOIN and LEAVE, compared to having to do a full LEAVE/JOIN (and possible barrier) sequence.
 - 4.1.1 What would such a call be named?
 - 4.1.2 Would this be “barrier-like” or simply collective? i.e. how would it differ from LEAVE?