

MPI: A Message-Passing Interface Standard

Version **3.0**

(Draft, with MPI 3 Nonblocking Collectives

and new Fortran 2008 Interface)

Unofficial, for comment only

Message Passing Interface Forum

March 26, 2011

1 This document describes the Message-Passing Interface (MPI) standard, version 3.0.
2 The MPI standard includes point-to-point message-passing, collective communications, group
3 and communicator concepts, process topologies, environmental management, process cre-
4 ation and management, one-sided communications, extended collective operations, external
5 interfaces, I/O, some miscellaneous topics, and a profiling interface. Language bindings for
6 C, C++ and Fortran are defined.

7 Historically, the evolution of the standards is from MPI-1.0 (June 1994) to MPI-1.1
8 (June 12, 1995) to MPI-1.2 (July 18, 1997), with several clarifications and additions and
9 published as part of the MPI-2 document, to MPI-2.0 (July 18, 1997), with new functionality,
10 to MPI-1.3 (May 30, 2008), combining for historical reasons the documents 1.1 and 1.2
11 and some errata documents to one combined document, and to MPI-2.1 (June 23, 2008),
12 combining the previous documents. Version MPI-2.2 (September 2009) added additional
13 clarifications and seven new routines. This version, MPI-3.0, is an extension of MPI-2.2.

14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45 ©1993, 1994, 1995, 1996, 1997, 2008, 2009, 2010 University of Tennessee, Knoxville,
46 Tennessee. Permission to copy without fee all or part of this material is granted, provided
47 the University of Tennessee copyright notice and the title of this document appear, and
48 notice is given that copying is by permission of the University of Tennessee.

| | |
|--|------------|
| Problems Due to Strong Typing | 544 |
| Problems Due to Data Copying and Sequence Association | 545 |
| Special Constants | 547 |
| FortranDerived Types | 548 |
| A Problem with Register Optimization and Temporary Memory Mod- ifications | 549 |
| 16.2.3 Fortran Support Through the <code>mpif.h</code> Include File | 556 |
| 16.2.4 Fortran Support Through the <code>mpi</code> Module | 557 |
| 16.2.5 Fortran Support Through the <code>mpi_f08</code> Module | 558 |
| 16.2.6 Additional Support for Fortran Register-Memory-Synchronization | 560 |
| 16.2.7 Additional Support for Fortran Numeric Intrinsic Types | 560 |
| Parameterized Datatypes with Specified Precision and Exponent Range | 561 |
| Support for Size-specific MPI Datatypes | 565 |
| Communication With Size-specific Types | 568 |
| 16.3 Language Interoperability | 569 |
| 16.3.1 Introduction | 569 |
| 16.3.2 Assumptions | 570 |
| 16.3.3 Initialization | 570 |
| 16.3.4 Transfer of Handles | 571 |
| 16.3.5 Status | 574 |
| 16.3.6 MPI Opaque Objects | 576 |
| Datatypes | 576 |
| Callback Functions | 577 |
| Error Handlers | 578 |
| Reduce Operations | 578 |
| Addresses | 578 |
| 16.3.7 Attributes | 579 |
| 16.3.8 Extra State | 583 |
| 16.3.9 Constants | 583 |
| 16.3.10 Interlanguage Communication | 584 |
| A Language Bindings Summary | 585 |
| A.1 Defined Values and Handles | 585 |
| A.1.1 Defined Constants | 585 |
| A.1.2 Types | 597 |
| A.1.3 Prototype Definitions | 598 |
| A.1.4 Deprecated Prototype Definitions | 602 |
| A.1.5 Info Keys | 602 |
| A.1.6 Info Values | 603 |
| A.2 C Bindings | 604 |
| A.2.1 Point-to-Point Communication C Bindings | 604 |
| A.2.2 Datatypes C Bindings | 605 |
| A.2.3 Collective Communication C Bindings | 607 |
| A.2.4 Groups, Contexts, Communicators, and Caching C Bindings | 610 |
| A.2.5 Process Topologies C Bindings | 612 |
| A.2.6 MPI Environmental Management C Bindings | 613 |
| A.2.7 The Info Object C Bindings | 614 |
| A.2.8 Process Creation and Management C Bindings | 615 |

Chapter 1

Introduction to MPI

1.1 Overview and Goals

MPI (Message-Passing Interface) is a *message-passing library interface specification*. All parts of this definition are significant. MPI addresses primarily the message-passing parallel programming model, in which data is moved from the address space of one process to that of another process through cooperative operations on each process. Extensions to the “classical” message-passing model are provided in collective operations, remote-memory access operations, dynamic process creation, and parallel I/O. MPI is a *specification*, not an implementation; there are multiple implementations of MPI. This specification is for a *library interface*; MPI is not a language, and all MPI operations are expressed as functions, subroutines, or methods, according to the appropriate language bindings, which for C, C++, and Fortran, are part of the MPI standard. The standard has been defined through an open process by a community of parallel computing vendors, computer scientists, and application developers. The next few sections provide an overview of the history of MPI’s development.

The main advantages of establishing a message-passing standard are portability and ease of use. In a distributed memory communication environment in which the higher level routines and/or abstractions are built upon lower level message-passing routines the benefits of standardization are particularly apparent. Furthermore, the definition of a message-passing standard, such as that proposed here, provides vendors with a clearly defined base set of routines that they can implement efficiently, or in some cases **for which they can provide hardware support**, thereby enhancing scalability.

The goal of the Message-Passing Interface simply stated is to develop a widely used standard for writing message-passing programs. As such the interface should establish a practical, portable, efficient, and flexible standard for message passing.

A complete list of goals follows.

- Design an application programming interface (not necessarily for compilers or a system implementation library).
- Allow efficient communication: Avoid memory-to-memory copying, allow overlap of computation and communication, and offload to communication co-processor, where available.
- Allow for implementations that can be used in a heterogeneous environment.

- Allow convenient C, C++, and Fortran bindings for the interface.
- Assume a reliable communication interface: the user need not cope with communication failures. Such failures are dealt with by the underlying communication subsystem.
- Define an interface that can be implemented on many vendor's platforms, with no significant changes in the underlying communication and system software.
- Semantics of the interface should be language independent.
- The interface should be designed to allow for thread safety.

1.2 Background of MPI-1.0

MPI sought to make use of the most attractive features of a number of existing message-passing systems, rather than selecting one of them and adopting it as the standard. Thus, MPI was strongly influenced by work at the IBM T. J. Watson Research Center [1, 2], Intel's NX/2 [44], Express [12], nCUBE's Vertex [40], p4 [7, 8], and PARMACS [5, 9]. Other important contributions have come from Zipcode [47, 48], Chimp [16, 17], PVM [4, 14], Chameleon [25], and PICL [24].

The MPI standardization effort involved about 60 people from 40 organizations mainly from the United States and Europe. Most of the major vendors of concurrent computers were involved in MPI, along with researchers from universities, government laboratories, and industry. The standardization process began with the Workshop on Standards for Message-Passing in a Distributed Memory Environment, sponsored by the Center for Research on Parallel Computing, held April 29-30, 1992, in Williamsburg, Virginia [55]. At this workshop the basic features essential to a standard message-passing interface were discussed, and a working group established to continue the standardization process.

A preliminary draft proposal, known as MPI1, was put forward by Dongarra, Hempel, Hey, and Walker in November 1992, and a revised version was completed in February 1993 [15]. MPI1 embodied the main features that were identified at the Williamsburg workshop as being necessary in a message passing standard. Since MPI1 was primarily intended to promote discussion and "get the ball rolling," it focused mainly on point-to-point communications. MPI1 brought to the forefront a number of important standardization issues, but did not include any collective communication routines and was not thread-safe.

In November 1992, a meeting of the MPI working group was held in Minneapolis, at which it was decided to place the standardization process on a more formal footing, and to generally adopt the procedures and organization of the High Performance Fortran Forum. Subcommittees were formed for the major component areas of the standard, and an email discussion service established for each. In addition, the goal of producing a draft MPI standard by the Fall of 1993 was set. To achieve this goal the MPI working group met every 6 weeks for two days throughout the first 9 months of 1993, and presented the draft MPI standard at the Supercomputing 93 conference in November 1993. These meetings and the email discussion together constituted the MPI Forum, membership of which has been open to all members of the high performance computing community.

MPI-2.0 was released, and a second ballot was voted on May 22, 2002. Both votes were done electronically. Both ballots were combined into one document: “Errata for MPI-2”, May 15, 2002. This errata process was then interrupted, but the Forum and its e-mail reflectors kept working on new requests for clarification.

Restarting regular work of the MPI Forum was initiated in three meetings, at EuroPVM/MPI’06 in Bonn, at EuroPVM/MPI’07 in Paris, and at SC’07 in Reno. In December 2007, a steering committee started the organization of new MPI Forum meetings at regular 8-weeks intervals. At the January 14-16, 2008 meeting in Chicago, the MPI Forum decided to combine the existing and future MPI documents to one document for each version of the MPI standard. For technical and historical reasons, this series was started with MPI-1.3. Additional Ballots 3 and 4 solved old questions from the errata list started in 1995 up to new questions from the last years. After all documents (MPI-1.1, MPI-2, Errata for MPI-1.1 (Oct. 12, 1998), and MPI-2.1 Ballots 1-4) were combined into one draft document, for each chapter, a chapter author and review team were defined. They cleaned up the document to achieve a consistent MPI-2.1 document. The final MPI-2.1 standard document was finished in June 2008, and finally released with a second vote in September 2008 in the meeting at Dublin, just before EuroPVM/MPI’08. The major work of the current MPI Forum is the preparation of MPI-3.

1.5 Background of MPI-2.2

MPI-2.2 is a minor update to the MPI-2.1 standard. This version addresses additional errors and ambiguities that were not corrected in the MPI-2.1 standard as well as a small number of extensions to MPI-2.1 that met the following criteria:

- Any correct MPI-2.1 program is a correct MPI-2.2 program.
- Any extension must have significant benefit for users.
- Any extension must not require significant implementation effort. To that end, all such changes are accompanied by an open source implementation.

The discussions of MPI-2.2 proceeded concurrently with the MPI-3 discussions; in some cases, extensions were proposed for MPI-2.2 but were later moved to MPI-3.

1.6 Background of MPI-3.0

MPI-3.0 is a major update to the MPI standard. Areas of particular interest are the extension of collective operations to include nonblocking and sparse-group routines and more flexible and powerful one-sided operations. This *draft* contains the MPI Forum’s current draft of nonblocking collective routines.

A new Fortran `mpi_f08` module is introduced to provide extended compile-time argument checking and buffer handling in nonblocking routines. The existing `mpi` module provides compile-time argument checking on the basis of existing MPI-2.2 routine definitions. The use of `mpif.h` is strongly discouraged.

MPI_CLASS_ACTION. For C and Fortran we use the C++ terminology to define the **Class**. In C++, the routine is a method on **Class** and is named `MPI::Class::Action_subset`. If the routine is associated with a certain class, but does not make sense as an object method, it is a static member function of the class.

2. If the routine is not associated with a class, the name should be of the form `MPI_Action_subset` in C and `MPI_ACTION_SUBSET` in Fortran, and in C++ should be scoped in the MPI namespace, `MPI::Action_subset`.
3. The names of certain actions have been standardized. In particular, **Create** creates a new object, **Get** retrieves information about an object, **Set** sets this information, **Delete** deletes information, **Is** asks whether or not an object has a certain property.

C and Fortran names for some MPI functions (that were defined during the MPI-1 process) violate these rules in several cases. The most common exceptions are the omission of the **Class** name from the routine and the omission of the **Action** where one can be inferred.

MPI identifiers are limited to 30 characters (31 with the profiling interface). This is done to avoid exceeding the limit on some compilation systems.

2.3 Procedure Specification

MPI procedures are specified using a language-independent notation. The arguments of procedure calls are marked as IN, OUT or INOUT. The meanings of these are:

- IN: the call may use the input value but does not update the argument,
- OUT: the call may update the argument but does not use its input value,
- INOUT: the call may both use and update the argument.

There is one special case — if an argument is a handle to an opaque object (these terms are defined in Section 2.5.1), and the object is updated by the procedure call, then the argument is marked INOUT or OUT. It is marked this way even though the handle itself is not modified — we use the INOUT or OUT attribute to denote that what the handle *references* is updated. Thus, in C++, IN arguments are usually either references or pointers to `const` objects.

Rationale. The definition of MPI tries to avoid, to the largest possible extent, the use of INOUT arguments, because such use is error-prone, especially for scalar arguments. (*End of rationale.*)

MPI’s use of IN, OUT and INOUT is intended to indicate to the user how an argument is to be used, but does not provide a rigorous classification that can be translated directly into all language bindings (e.g., `INTENT` in Fortran 90 bindings or `const` in C bindings). For instance, the “constant” `MPI_BOTTOM` can usually be passed to OUT buffer arguments. Similarly, `MPI_STATUS_IGNORE` can be passed as the OUT status argument (with the `mpi module` or `mpif.h`).

A common occurrence for MPI functions is an argument that is used as IN by some processes and OUT by other processes. Such an argument is, syntactically, an INOUT argument

and is marked as such, although, semantically, it is not used in one call both for input and for output on a single process.

Another frequent situation arises when an argument value is needed only by a subset of the processes. When an argument is not significant at a process then an arbitrary value can be passed as an argument.

Unless specified otherwise, an argument of type OUT or type INOUT cannot be aliased with any other argument passed to an MPI procedure. An example of argument aliasing in C appears below. If we define a C procedure like this,

```
void copyIntBuffer( int *pin, int *pout, int len )
{
    int i;
    for (i=0; i<len; ++i) *pout++ = *pin++;
}
```

then a call to it in the following code fragment has aliased arguments.

```
int a[10];
copyIntBuffer( a, a+3, 7);
```

Although the C language allows this, such usage of MPI procedures is forbidden unless otherwise specified. Note that Fortran prohibits aliasing of arguments.

All MPI functions are first specified in the language-independent notation. Immediately below this, **language dependent bindings follow**:

- The ISO C version of the function.
- The Fortran version of the same function used with `USE mpi` or `INCLUDE 'mpif.h'`
- The Fortran version used with `USE mpi_f08`.
- The C++ binding (which is deprecated).

Fortran in this document refers to Fortran 90 and higher; see Section 2.6.

2.4 Semantic Terms

When discussing MPI procedures the following semantic terms are used.

nonblocking A procedure is nonblocking if the procedure may return before the operation completes, and before the user is allowed to reuse resources (such as buffers) specified in the call. A nonblocking request is **started** by the call that initiates it, e.g., `MPI_ISEND`. The word complete is used with respect to operations, requests, and communications. An **operation completes** when the user is allowed to reuse resources, and any output buffers have been updated; i.e. a call to `MPI_TEST` will return `flag = true`. A **request is completed** by a call to wait, which returns, or a test or get status call which returns `flag = true`. This completing call has two effects: the status is extracted from the request; in the case of test and wait, if the request was nonpersistent, it is **freed**, and becomes **inactive** if it was persistent. A **communication completes** when all participating operations complete.

blocking A procedure is blocking if return from the procedure indicates the user is allowed to reuse resources specified in the call.

local A procedure is local if completion of the procedure depends only on the local executing process.

non-local A procedure is non-local if completion of the operation may require the execution of some MPI procedure on another process. Such an operation may require communication occurring with another user process.

collective A procedure is collective if all processes in a process group need to invoke the procedure. A collective call may or may not be synchronizing. Collective calls over the same communicator must be executed in the same order by all members of the process group.

predefined A predefined datatype is a datatype with a predefined (constant) name (such as `MPI_INT`, `MPI_FLOAT_INT`, or `MPI_UB`) or a datatype constructed with `MPI_TYPE_CREATE_F90_INTEGER`, `MPI_TYPE_CREATE_F90_REAL`, or `MPI_TYPE_CREATE_F90_COMPLEX`. The former are **named** whereas the latter are **unnamed**.

derived A derived datatype is any datatype that is not predefined.

portable A datatype is portable, if it is a predefined datatype, or it is derived from a portable datatype using only the type constructors `MPI_TYPE_CONTIGUOUS`, `MPI_TYPE_VECTOR`, `MPI_TYPE_INDEXED`, `MPI_TYPE_CREATE_INDEXED_BLOCK`, `MPI_TYPE_CREATE_SUBARRAY`, `MPI_TYPE_DUP`, and `MPI_TYPE_CREATE_DARRAY`. Such a datatype is portable because all displacements in the datatype are in terms of extents of one predefined datatype. Therefore, if such a datatype fits a data layout in one memory, it will fit the corresponding data layout in another memory, if the same declarations were used, even if the two systems have different architectures. On the other hand, if a datatype was constructed using `MPI_TYPE_CREATE_HINDEXED`, `MPI_TYPE_CREATE_HVECTOR` or `MPI_TYPE_CREATE_STRUCT`, then the datatype contains explicit byte displacements (e.g., providing padding to meet alignment restrictions). These displacements are unlikely to be chosen correctly if they fit data layout on one memory, but are used for data layouts on another process, running on a processor with a different architecture.

equivalent Two datatypes are equivalent if they appear to have been created with the same sequence of calls (and arguments) and thus have the same typemap. Two equivalent datatypes do not necessarily have the same cached attributes or the same names.

2.5 Data Types

2.5.1 Opaque Objects

MPI manages **system memory** that is used for buffering messages and for storing internal representations of various MPI objects such as groups, communicators, datatypes, etc. This memory is not directly accessible to the user, and objects stored there are **opaque**: their size and shape is not visible to the user. Opaque objects are accessed via **handles**, which exist in user space. MPI procedures that operate on opaque objects are passed handle

arguments to access these objects. In addition to their use by MPI calls for object access, handles can participate in assignments and comparisons.

In Fortran with `USE mpi` or `INCLUDE 'mpif.h'`, all handles have type `INTEGER`. In Fortran with `USE mpi_f08`, and in C and C++, a different handle type is defined for each category of objects. With Fortran `USE mpi_f08`, the handles are defined as Fortran sequenced derived types that consist of only one element `INTEGER :: MPI_VAL`. The internal handle value is identical to the Fortran `INTEGER` value used in the `mpi` module and in `mpif.h`. The names are identical to the names in C, except that they are not case sensitive. For example:

```
TYPE MPI_Comm
  SEQUENCE
  INTEGER :: MPI_VAL
END TYPE MPI_Comm
```

In addition, handles themselves are distinct objects in C++. The C and C++ types must support the use of the assignment and equality operators.

Advice to implementors. In Fortran, the handle can be an index into a table of opaque objects in a system table; in C it can be such an index or a pointer to the object. C++ handles can simply “wrap up” a table index or pointer. (*End of advice to implementors.*)

Rationale. Due to the sequence attribute in the definition of handles in the `mpi_f08` module, the new Fortran handles are associated with one numerical storage unit, i.e., they have the same C binding as the `INTEGER` handles of the `mpi` module. Due to the equivalence of the integer values, applications can easily convert MPI handles between all three supported Fortran methods. For example, an integer communicator handle `COMM` can be converted directly into an exactly equivalent `mpi_f08` communicator handle named `comm_f08` by `comm_f08%MPI_VAL=COMM`, and vice versa. (*End of rationale.*)

Opaque objects are allocated and deallocated by calls that are specific to each object type. These are listed in the sections where the objects are described. The calls accept a handle argument of matching type. In an allocate call this is an OUT argument that returns a valid reference to the object. In a call to deallocate this is an INOUT argument which returns with an “invalid handle” value. MPI provides an “invalid handle” constant for each object type. Comparisons to this constant are used to test for validity of the handle.

A call to a deallocate routine invalidates the handle and marks the object for deallocation. The object is not accessible to the user after the call. However, MPI need not deallocate the object immediately. Any operation pending (at the time of the deallocate) that involves this object will complete normally; the object will be deallocated afterwards.

An opaque object and its handle are significant only at the process where the object was created and cannot be transferred to another process.

MPI provides certain predefined opaque objects and predefined, static handles to these objects. The user must not free such objects. In C++, this is enforced by declaring the handles to these predefined objects to be `static const`.

Rationale. This design hides the internal representation used for MPI data structures, thus allowing similar calls in C, C++, and Fortran. It also avoids conflicts with the typing rules in these languages, and easily allows future extensions of functionality. The mechanism for opaque objects used here loosely follows the POSIX Fortran binding standard.

The explicit separation of handles in user space and objects in system space allows space-reclaiming and deallocation calls to be made at appropriate points in the user program. If the opaque objects were in user space, one would have to be very careful not to go out of scope before any pending operation requiring that object completed. The specified design allows an object to be marked for deallocation, the user program can then go out of scope, and the object itself still persists until any pending operations are complete.

The requirement that handles support assignment/comparison is made since such operations are common. This restricts the domain of possible implementations. The alternative would have been to allow handles to have been an arbitrary, opaque type. This would force the introduction of routines to do assignment and comparison, adding complexity, and was therefore ruled out. (*End of rationale.*)

Advice to users. A user may accidentally create a dangling reference by assigning to a handle the value of another handle, and then deallocating the object associated with these handles. Conversely, if a handle variable is deallocated before the associated object is freed, then the object becomes inaccessible (this may occur, for example, if the handle is a local variable within a subroutine, and the subroutine is exited before the associated object is deallocated). It is the user's responsibility to avoid adding or deleting references to opaque objects, except as a result of MPI calls that allocate or deallocate such objects. (*End of advice to users.*)

Advice to implementors. The intended semantics of opaque objects is that opaque objects are separate from one another; each call to allocate such an object copies all the information required for the object. Implementations may avoid excessive copying by substituting referencing for copying. For example, a derived datatype may contain references to its components, rather than copies of its components; a call to `MPI_COMM_GROUP` may return a reference to the group associated with the communicator, rather than a copy of this group. In such cases, the implementation must maintain reference counts, and allocate and deallocate objects in such a way that the visible effect is as if the objects were copied. (*End of advice to implementors.*)

2.5.2 Array Arguments

An MPI call may need an argument that is an array of opaque objects, or an array of handles. The array-of-handles is a regular array with entries that are handles to objects of the same type in consecutive locations in the array. Whenever such an array is used, an additional `len` argument is required to indicate the number of valid entries (unless this number can be derived otherwise). The valid entries are at the beginning of the array; `len` indicates how many of them there are, and need not be the size of the entire array. The same approach is followed for other array arguments. In some cases `NULL` handles are considered valid entries. When a `NULL` argument is desired for an array of statuses, one uses `MPI_STATUSES_IGNORE`. **With the `mpi_f08` module, optional arguments through function overloading are used instead of**

`MPI_STATUS_IGNORE`, `MPI_STATUSES_IGNORE`, `MPI_ERRCODES_IGNORE`, and `MPI_UNWEIGHTED`. The constants `MPI_ARGV_NULL` and `MPI_ARGVS_NULL` are not substituted by function overloading.

2.5.3 State

MPI procedures use at various places arguments with *state* types. The values of such a data type are all identified by names, and no operation is defined on them. For example, the `MPI_TYPE_CREATE_SUBARRAY` routine has a state argument `order` with values `MPI_ORDER_C` and `MPI_ORDER_FORTRAN`.

2.5.4 Named Constants

MPI procedures sometimes assign a special meaning to a special value of a basic type argument; e.g., `tag` is an integer-valued argument of point-to-point communication operations, with a special wild-card value, `MPI_ANY_TAG`. Such arguments will have a range of regular values, which is a proper subrange of the range of values of the corresponding basic type; special values (such as `MPI_ANY_TAG`) will be outside the regular range. The range of regular values, such as `tag`, can be queried using environmental inquiry functions (Chapter 7 of the MPI-1 document). The range of other values, such as `source`, depends on values given by other MPI routines (in the case of `source` it is the communicator size).

MPI also provides predefined named constant handles, such as `MPI_COMM_WORLD`.

All named constants, with the exceptions noted below for Fortran, can be used in initialization expressions or assignments, but not necessarily in array declarations or as labels in C/C++ `switch` or Fortran `select/case` statements. This implies named constants to be link-time but not necessarily compile-time constants. The named constants listed below are required to be compile-time constants in both C/C++ and Fortran. These constants do not change values during execution. Opaque objects accessed by constant handles are defined and do not change value between MPI initialization (`MPI_INIT`) and MPI completion (`MPI_FINALIZE`). The handles themselves are constants and can be also used in initialization expressions or assignments.

The constants that are required to be compile-time constants (and can thus be used for array length declarations and labels in C/C++ `switch` and Fortran `case/select` statements) are:

- `MPI_MAX_PROCESSOR_NAME`
- `MPI_MAX_ERROR_STRING`
- `MPI_MAX_DATAREP_STRING`
- `MPI_MAX_INFO_KEY`
- `MPI_MAX_INFO_VAL`
- `MPI_MAX_OBJECT_NAME`
- `MPI_MAX_PORT_NAME`
- `MPI_STATUS_SIZE` (Fortran only)
- `MPI_ADDRESS_KIND` (Fortran only)
- `MPI_INTEGER_KIND` (Fortran only)
- `MPI_OFFSET_KIND` (Fortran only)

and their C++ counterparts where appropriate.

The constants that cannot be used in initialization expressions or assignments in Fortran are:

```

1      MPI_BOTTOM
2      MPI_STATUS_IGNORE
3      MPI_STATUSES_IGNORE
4      MPI_ERRCODES_IGNORE
5      MPI_IN_PLACE
6      MPI_ARGV_NULL
7      MPI_ARGVS_NULL
8      MPI_UNWEIGHTED

```

Advice to implementors. In Fortran the implementation of these special constants may require the use of language constructs that are outside the Fortran standard. Using special values for the constants (e.g., by defining them through `PARAMETER` statements) is not possible because an implementation cannot distinguish these values from legal data. Typically, these constants are implemented as predefined static variables (e.g., a variable in an MPI-declared `COMMON` block), relying on the fact that the target compiler passes data by address. Inside the subroutine, this address can be extracted by some mechanism outside the Fortran standard (e.g., by Fortran extensions or by implementing the function in C). (*End of advice to implementors.*)

2.5.5 Choice

MPI functions sometimes use arguments with a *choice* (or union) data type. Distinct calls to the same routine may pass by reference actual arguments of different types. The mechanism for providing such arguments will differ from language to language. For Fortran **with the include file `mpif.h` or the `mpi` module**, the document uses `<type>` to represent a choice variable; **with the Fortran `mpi_f08` module, such arguments are declared with the Fortran 2008 syntax `TYPE(*)`, `DIMENSION(..)`**; for C and C++, we use `void *`.

Advice to implementors. The implementor can freely choose how to implement choice arguments in the `mpi` module, e.g., with a non-standard compiler-dependent method that has the quality of the call mechanism in the implicit Fortran interfaces, or with the method defined for the `mpi_f08` module. (*End of advice to implementors.*)

2.5.6 Addresses

Some MPI procedures use *address* arguments that represent an absolute address in the calling program. The datatype of such an argument is `MPI_Aint` in C, `MPI::Aint` in C++ and `INTEGER (KIND=MPI_ADDRESS_KIND)` in Fortran. These types must have the same width and encode address values in the same manner such that address values in one language may be passed directly to another language without conversion. There is the MPI constant `MPI_BOTTOM` to indicate the start of the address range.

2.5.7 File Offsets

For I/O there is a need to give the size, displacement, and offset into a file. These quantities can easily be larger than 32 bits which can be the default size of a Fortran integer. To overcome this, these quantities are declared to be `INTEGER (KIND=MPI_OFFSET_KIND)` in Fortran. In C one uses `MPI_Offset` whereas in C++ one uses `MPI::Offset`. These types

must have the same width and encode address values in the same manner such that offset values in one language may be passed directly to another language without conversion.

2.6 Language Binding

This section defines the rules for MPI language binding in general and for Fortran, ISO C, and C++, in particular. (Note that ANSI C has been replaced by ISO C.) The C++ language bindings have been deprecated. Defined here are various object representations, as well as the naming conventions used for expressing this standard. The actual calling sequences are defined elsewhere.

MPI bindings are for Fortran 90 or later, though they were originally designed to be usable in Fortran 77 environments. With the `mpi_f08` module, the two Fortran 2008 features *assumed type* and *assumed rank* are also required, see Section 2.5.5 on page 16.

Since the word `PARAMETER` is a keyword in the Fortran language, we use the word “argument” to denote the arguments to a subroutine. These are normally referred to as parameters in C and C++, however, we expect that C and C++ programmers will understand the word “argument” (which has no specific meaning in C/C++), thus allowing us to avoid unnecessary confusion for Fortran programmers.

Since Fortran is case insensitive, linkers may use either lower case or upper case when resolving Fortran names. Users of case sensitive languages should avoid the “`mpi_`” and “`pmpi_`” prefixes.

2.6.1 Deprecated Names and Functions

A number of chapters refer to deprecated or replaced MPI-1 constructs. These are constructs that continue to be part of the MPI standard, as documented in Chapter 15, but that users are recommended not to continue using, since better solutions were provided with MPI-2. For example, the Fortran binding for MPI-1 functions that have address arguments uses `INTEGER`. This is not consistent with the C binding, and causes problems on machines with 32 bit `INTEGER`s and 64 bit addresses. In MPI-2, these functions were given new names with new bindings for the address arguments. The use of the old functions is deprecated. For consistency, here and in a few other cases, new C functions are also provided, even though the new functions are equivalent to the old functions. The old names are deprecated. Another example is provided by the MPI-1 predefined datatypes `MPI_UB` and `MPI_LB`. They are deprecated, since their use is awkward and error-prone. The MPI-2 function `MPI_TYPE_CREATE_RESIZED` provides a more convenient mechanism to achieve the same effect.

Table 2.1 shows a list of all of the deprecated constructs. Note that the constants `MPI_LB` and `MPI_UB` are replaced by the function `MPI_TYPE_CREATE_RESIZED`; this is because their principal use was as input datatypes to `MPI_TYPE_STRUCT` to create resized datatypes. Also note that some C typedefs and Fortran subroutine names are included in this list; they are the types of callback functions.

2.6.2 Fortran Binding Issues

Originally, MPI-1.1 provided bindings for Fortran 77. These bindings are retained, but they are now interpreted in the context of the Fortran 90 standard. MPI can still be used with

| Deprecated | MPI-2 Replacement |
|-----------------------|-------------------------------|
| MPI_ADDRESS | MPI_GET_ADDRESS |
| MPI_TYPE_HINDEXED | MPI_TYPE_CREATE_HINDEXED |
| MPI_TYPE_HVECTOR | MPI_TYPE_CREATE_HVECTOR |
| MPI_TYPE_STRUCT | MPI_TYPE_CREATE_STRUCT |
| MPI_TYPE_EXTENT | MPI_TYPE_GET_EXTENT |
| MPI_TYPE_UB | MPI_TYPE_GET_EXTENT |
| MPI_TYPE_LB | MPI_TYPE_GET_EXTENT |
| MPI_LB | MPI_TYPE_CREATE_RESIZED |
| MPI_UB | MPI_TYPE_CREATE_RESIZED |
| MPI_ERRHANDLER_CREATE | MPI_COMM_CREATE_ERRHANDLER |
| MPI_ERRHANDLER_GET | MPI_COMM_GET_ERRHANDLER |
| MPI_ERRHANDLER_SET | MPI_COMM_SET_ERRHANDLER |
| MPI_Handler_function | MPI_Comm_errhandler_function |
| MPI_KEYVAL_CREATE | MPI_COMM_CREATE_KEYVAL |
| MPI_KEYVAL_FREE | MPI_COMM_FREE_KEYVAL |
| MPI_DUP_FN | MPI_COMM_DUP_FN |
| MPI_NULL_COPY_FN | MPI_COMM_NULL_COPY_FN |
| MPI_NULL_DELETE_FN | MPI_COMM_NULL_DELETE_FN |
| MPI_Copy_function | MPI_Comm_copy_attr_function |
| COPY_FUNCTION | COMM_COPY_ATTR_FUNCTION |
| MPI_Delete_function | MPI_Comm_delete_attr_function |
| DELETE_FUNCTION | COMM_DELETE_ATTR_FUNCTION |
| MPI_ATTR_DELETE | MPI_COMM_DELETE_ATTR |
| MPI_ATTR_GET | MPI_COMM_GET_ATTR |
| MPI_ATTR_PUT | MPI_COMM_SET_ATTR |

Table 2.1: Deprecated constructs

most Fortran 77 compilers, as noted below. When the term Fortran is used it means Fortran 90 or later; it means Fortran 2008 and later if the `mpi_f08` module is used.

All MPI names have an `MPI_` prefix, and all characters are capitals. Programs must not declare variables, parameters, or functions with names beginning with the prefix `MPI_`. To avoid conflicting with the profiling interface, programs should also avoid functions with the prefix `PMPI_`. This is mandated to avoid possible name collisions.

All MPI Fortran subroutines have a return code in the last argument. With `USE mpi_f08`, this last argument is declared as `OPTIONAL`, except for user-defined callback functions (e.g., `COMM_COPY_ATTR_FUNCTION`) and their predefined callbacks (e.g., `MPI_NULL_COPY_FN`). A few MPI operations which are functions do not have the return code argument. The return code value for successful completion is `MPI_SUCCESS`. Other error codes are implementation dependent; see the error codes in Chapter 8 and Annex A.

Constants representing the maximum length of a string are one smaller in Fortran than in C and C++ as discussed in Section 16.3.9.

Handles are represented in Fortran as `INTEGER`s, or with the `mpi_f08` module as a derived type, see Section 2.5.1 on page 12. Binary-valued variables are of type `LOGICAL`.

Array arguments are indexed from one.

The MPI Fortran bindings **are** inconsistent with the Fortran standard in several respects. These inconsistencies, such as register optimization problems, have implications for user codes that are discussed in detail in Section 16.2.2.

2.6.3 C Binding Issues

We use the ISO C declaration format. All MPI names have an `MPI_` prefix, defined constants are in all capital letters, and defined types and functions have one capital letter after the prefix. Programs must not declare variables or functions with names beginning with the prefix `MPI_`. To support the profiling interface, programs should not declare functions with names beginning with the prefix `PMPI_`.

The definition of named constants, function prototypes, and type definitions must be supplied in an include file `mpi.h`.

Almost all C functions return an error code. The successful return code will be `MPI_SUCCESS`, but failure return codes are implementation dependent.

Type declarations are provided for handles to each category of opaque objects.

Array arguments are indexed from zero.

Logical flags are integers with value 0 meaning “false” and a non-zero value meaning “true.”

Choice arguments are pointers of type `void *`.

Address arguments are of MPI defined type `MPI_Aint`. File displacements are of type `MPI_Offset`. `MPI_Aint` is defined to be an integer of the size needed to hold any valid address on the target architecture. `MPI_Offset` is defined to be an integer of the size needed to hold any valid file size on the target architecture.

2.6.4 C++ Binding Issues

The C++ language bindings have been deprecated. There are places in the standard that give rules for C and not for C++. In these cases, the C rule should be applied to the C++ case, as appropriate. In particular, the values of constants given in the text are the ones for C and Fortran. A cross index of these with the C++ names is given in Annex A.

We use the ISO C++ declaration format. All MPI names are declared within the scope of a namespace called `MPI` and therefore are referenced with an `MPI::` prefix. Defined constants are in all capital letters, and class names, defined types, and functions have only their first letter capitalized. Programs must not declare variables or functions in the `MPI` namespace. This is mandated to avoid possible name collisions.

The definition of named constants, function prototypes, and type definitions must be supplied in an include file `mpi.h`.

Advice to implementors. The file `mpi.h` may contain both the C and C++ definitions. Usually one can simply use the defined value (generally `__cplusplus`, but not required) to see if one is using C++ to protect the C++ definitions. It is possible that a C compiler will require that the source protected this way be legal C code. In this case, all the C++ definitions can be placed in a different include file and the “`#include`” directive can be used to include the necessary C++ definitions in the `mpi.h` file. (*End of advice to implementors.*)

C++ functions that create objects or return information usually place the object or information in the return value. Since the language neutral prototypes of MPI functions

operation to select a particular message. The last three parameters of the send operation, along with the rank of the sender, specify the envelope for the message sent. Process one (`myrank = 1`) receives this message with the **receive** operation `MPI_RECV`. The message to be received is selected according to the value of its envelope, and the message data is stored into the **receive buffer**. In the example above, the receive buffer consists of the storage containing the string `message` in the memory of process one. The first three parameters of the receive operation specify the location, size and type of the receive buffer. The next three parameters are used for selecting the incoming message. The last parameter is used to return information on the message just received.

The next sections describe the blocking send and receive operations. We discuss send, receive, blocking communication semantics, type matching requirements, type conversion in heterogeneous environments, and more general communication modes. Nonblocking communication is addressed next, followed by channel-like constructs and send-receive operations. Nonblocking communication is addressed next, followed by channel-like constructs and send-receive operations, ending with a description of the “dummy” process, `MPI_PROC_NULL`.

3.2 Blocking Send and Receive Operations

3.2.1 Blocking Send

The syntax of the blocking send operation is given below.

`MPI_SEND(buf, count, datatype, dest, tag, comm)`

| | | |
|----|----------|--|
| IN | buf | initial address of send buffer (choice) |
| IN | count | number of elements in send buffer (non-negative integer) |
| IN | datatype | datatype of each send buffer element (handle) |
| IN | dest | rank of destination (integer) |
| IN | tag | message tag (integer) |
| IN | comm | communicator (handle) |

```
int MPI_Send(void* buf, int count, MPI_Datatype datatype, int dest,
             int tag, MPI_Comm comm)
```

```
MPI_SEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, IERROR)
<type> BUF(*)
INTEGER COUNT, DATATYPE, DEST, TAG, COMM, IERROR
```

```
MPI_Send(buf, count, datatype, dest, tag, comm, ierror)
TYPE(*), DIMENSION(..) :: buf
INTEGER, INTENT(IN) :: count, dest, tag
TYPE(MPI_Datatype), INTENT(IN) :: datatype
TYPE(MPI_Comm), INTENT(IN) :: comm
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

A predefined communicator `MPI_COMM_WORLD` is provided by MPI. It allows communication with all processes that are accessible after MPI initialization and processes are identified by their rank in the group of `MPI_COMM_WORLD`.

Advice to users. Users that are comfortable with the notion of a flat name space for processes, and a single communication context, as offered by most existing communication libraries, need only use the predefined variable `MPI_COMM_WORLD` as the `comm` argument. This will allow communication with all the processes available at initialization time.

Users may define new communicators, as explained in Chapter 6. Communicators provide an important encapsulation mechanism for libraries and modules. They allow modules to have their own disjoint communication universe and their own process numbering scheme. (*End of advice to users.*)

Advice to implementors. The message envelope would normally be encoded by a fixed-length message header. However, the actual encoding is implementation dependent. Some of the information (e.g., source or destination) may be implicit, and need not be explicitly carried by messages. Also, processes may be identified by relative ranks, or absolute ids, etc. (*End of advice to implementors.*)

3.2.4 Blocking Receive

The syntax of the blocking receive operation is given below.

`MPI_RECV (buf, count, datatype, source, tag, comm, status)`

| | | |
|-----|----------|---|
| OUT | buf | initial address of receive buffer (choice) |
| IN | count | number of elements in receive buffer (non-negative integer) |
| IN | datatype | datatype of each receive buffer element (handle) |
| IN | source | rank of source or <code>MPI_ANY_SOURCE</code> (integer) |
| IN | tag | message tag or <code>MPI_ANY_TAG</code> (integer) |
| IN | comm | communicator (handle) |
| OUT | status | status object (Status) |

```
int MPI_Recv(void* buf, int count, MPI_Datatype datatype, int source,
            int tag, MPI_Comm comm, MPI_Status *status)
```

```
MPI_RECV(BUF, COUNT, DATATYPE, SOURCE, TAG, COMM, STATUS, IERROR)
<type> BUF(*)
INTEGER COUNT, DATATYPE, SOURCE, TAG, COMM, STATUS(MPI_STATUS_SIZE),
IERROR
```

```
MPI_Recv(buf, count, datatype, source, tag, comm, status, ierror)
TYPE(*), DIMENSION(..) :: buf
INTEGER, INTENT(IN) :: count, source, tag
TYPE(MPI_Datatype), INTENT(IN) :: datatype
```

```

TYPE(MPI_Comm), INTENT(IN) :: comm
TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

{void MPI::Comm::Recv(void* buf, int count, const MPI::Datatype& datatype,
    int source, int tag, MPI::Status& status) const(binding
    deprecated, see Section 15.2) }

{void MPI::Comm::Recv(void* buf, int count, const MPI::Datatype& datatype,
    int source, int tag) const(binding deprecated, see Section 15.2) }

```

The blocking semantics of this call are described in Section 3.4.

The receive buffer consists of the storage containing `count` consecutive elements of the type specified by `datatype`, starting at address `buf`. The length of the received message must be less than or equal to the length of the receive buffer. An overflow error occurs if all incoming data does not fit, without truncation, into the receive buffer.

If a message that is shorter than the receive buffer arrives, then only those locations corresponding to the (shorter) message are modified.

Advice to users. The `MPI_PROBE` function described in Section 3.8 can be used to receive messages of unknown length. (*End of advice to users.*)

Advice to implementors. Even though no specific behavior is mandated by MPI for erroneous programs, the recommended handling of overflow situations is to return in `status` information about the source and tag of the incoming message. The receive operation will return an error code. A quality implementation will also ensure that no memory that is outside the receive buffer will ever be overwritten.

In the case of a message shorter than the receive buffer, MPI is quite strict in that it allows no modification of the other locations. A more lenient statement would allow for some optimizations but this is not allowed. The implementation must be ready to end a copy into the receiver memory exactly at the end of the receive buffer, even if it is an odd address. (*End of advice to implementors.*)

The selection of a message by a receive operation is governed by the value of the message envelope. A message can be received by a receive operation if its envelope matches the `source`, `tag` and `comm` values specified by the receive operation. The receiver may specify a wildcard `MPI_ANY_SOURCE` value for `source`, and/or a wildcard `MPI_ANY_TAG` value for `tag`, indicating that any source and/or tag are acceptable. It cannot specify a wildcard value for `comm`. Thus, a message can be received by a receive operation only if it is addressed to the receiving process, has a matching communicator, has matching source unless `source=MPI_ANY_SOURCE` in the pattern, and has a matching tag unless `tag=MPI_ANY_TAG` in the pattern.

The message tag is specified by the `tag` argument of the receive operation. The argument `source`, if different from `MPI_ANY_SOURCE`, is specified as a rank within the process group associated with that same communicator (remote process group, for intercommunicators). Thus, the range of valid values for the `source` argument is $\{0, \dots, n-1\} \cup \{\text{MPI_ANY_SOURCE}\}$, where n is the number of processes in this group.

Note the asymmetry between send and receive operations: A receive operation may accept messages from an arbitrary sender, on the other hand, a send operation must specify

a unique receiver. This matches a “push” communication mechanism, where data transfer is effected by the sender (rather than a “pull” mechanism, where data transfer is effected by the receiver).

Source = destination is allowed, that is, a process can send a message to itself. (However, it is unsafe to do so with the blocking send and receive operations described above, since this may lead to deadlock. See Section 3.5.)

Advice to implementors. Message context and other communicator information can be implemented as an additional tag field. It differs from the regular message tag in that wild card matching is not allowed on this field, and that value setting for this field is controlled by communicator manipulation functions. (*End of advice to implementors.*)

3.2.5 Return Status

The source or tag of a received message may not be known if wildcard values were used in the receive operation. Also, if multiple requests are completed by a single MPI function (see Section 3.7.5), a distinct error code may need to be returned for each request. The information is returned by the `status` argument of `MPI_RECV`. The type of `status` is MPI-defined. Status variables need to be explicitly allocated by the user, that is, they are not system objects.

In C, `status` is a structure that contains three fields named `MPI_SOURCE`, `MPI_TAG`, and `MPI_ERROR`; the structure may contain additional fields. Thus, `status.MPI_SOURCE`, `status.MPI_TAG` and `status.MPI_ERROR` contain the source, tag, and error code, respectively, of the received message.

In Fortran with `USE mpi` or `INCLUDE 'mpif.h'`, `status` is an array of `INTEGER`s of size `MPI_STATUS_SIZE`. The constants `MPI_SOURCE`, `MPI_TAG` and `MPI_ERROR` are the indices of the entries that store the source, tag and error fields. Thus, `status(MPI_SOURCE)`, `status(MPI_TAG)` and `status(MPI_ERROR)` contain, respectively, the source, tag and error code of the received message.

With Fortran `USE mpi_f08`, `status` is defined as the Fortran derived type `TYPE(MPI_Status)`, which contains three fields named `MPI_SOURCE`, `MPI_TAG`, and `MPI_ERROR`; the derived type may contain additional fields. Thus, `status%MPI_SOURCE`, `status%MPI_TAG` and `status%MPI_ERROR` contain the source, tag, and error code, respectively, of the received message. Additionally, within the `mpi` and the `mpi_f08` module, both, the constants `MPI_STATUS_SIZE`, `MPI_SOURCE`, `MPI_TAG`, `MPI_ERROR`, and the `TYPE(MPI_Status)` is defined to allow with both modules the conversion between both `status` declarations.

Rationale. It is allowed to have the same name (e.g., `MPI_SOURCE`) defined as a constant (e.g., Fortran parameter) and as a field of a derived type. (*End of rationale.*)

Advice to implementors. The Fortran `TYPE(MPI_Status)` may be defined as a sequence derived type to achieve the same data layout as in C. (*End of advice to implementors.*)

In C++, the `status` object is handled through the following methods:
`{int MPI::Status::Get_source() const(binding deprecated, see Section 15.2) }`

```

{void MPI::Status::Set_source(int source) (binding deprecated, see Section 15.2) }
{int MPI::Status::Get_tag() const (binding deprecated, see Section 15.2) }
{void MPI::Status::Set_tag(int tag) (binding deprecated, see Section 15.2) }
{int MPI::Status::Get_error() const (binding deprecated, see Section 15.2) }
{void MPI::Status::Set_error(int error) (binding deprecated, see Section 15.2) }

```

In general, message-passing calls do not modify the value of the error code field of status variables. This field may be updated only by the functions in Section 3.7.5 which return multiple statuses. The field is updated if and only if such function returns with an error code of MPI_ERR_IN_STATUS.

Rationale. The error field in status is not needed for calls that return only one status, such as MPI_WAIT, since that would only duplicate the information returned by the function itself. The current design avoids the additional overhead of setting it, in such cases. The field is needed for calls that return multiple statuses, since each request may have had a different failure. (*End of rationale.*)

The status argument also returns information on the length of the message received. However, this information is not directly available as a field of the status variable and a call to MPI_GET_COUNT is required to “decode” this information.

MPI_GET_COUNT(status, datatype, count)

| | | |
|-----|----------|--|
| IN | status | return status of receive operation (Status) |
| IN | datatype | datatype of each receive buffer entry (handle) |
| OUT | count | number of received entries (integer) |

```
int MPI_Get_count(MPI_Status *status, MPI_Datatype datatype, int *count)
```

```
MPI_GET_COUNT(STATUS, DATATYPE, COUNT, IERROR)
```

```
INTEGER STATUS(MPI_STATUS_SIZE), DATATYPE, COUNT, IERROR
```

```
MPI_Get_count(status, datatype, count, ierror)
```

```
TYPE(MPI_Status), INTENT(IN) :: status
```

```
TYPE(MPI_Datatype), INTENT(IN) :: datatype
```

```
INTEGER, INTENT(OUT) :: count
```

```
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

```

{int MPI::Status::Get_count(const MPI::Datatype& datatype) const (binding
    deprecated, see Section 15.2) }

```

Returns the number of entries received. (Again, we count *entries*, each of type *datatype*, not *bytes*.) The *datatype* argument should match the argument provided by the receive call that set the *status* variable. (We shall later see, in Section 4.1.11, that MPI_GET_COUNT may return, in certain situations, the value MPI_UNDEFINED.)

Rationale. Some message-passing libraries use INOUT *count*, *tag* and *source* arguments, thus using them both to specify the selection criteria for incoming

messages and return the actual envelope values of the received message. The use of a separate status argument prevents errors that are often attached with INOUT argument (e.g., using the `MPI_ANY_TAG` constant as the tag in a receive). Some libraries use calls that refer implicitly to the “last message received.” This is not thread safe.

The `datatype` argument is passed to `MPI_GET_COUNT` so as to improve performance. A message might be received without counting the number of elements it contains, and the count value is often not needed. Also, this allows the same function to be used after a call to `MPI_PROBE` or `MPI_IPROBE`. With a status from `MPI_PROBE` or `MPI_IPROBE`, the same datatypes are allowed as in a call to `MPI_RECV` to receive this message. (*End of rationale.*)

The value returned as the `count` argument of `MPI_GET_COUNT` for a datatype of length zero where zero bytes have been transferred is zero. If the number of bytes transferred is greater than zero, `MPI_UNDEFINED` is returned.

Rationale. Zero-length datatypes may be created in a number of cases. An important case is `MPI_TYPE_CREATE_DARRAY`, where the definition of the particular darray results in an empty block on some MPI process. Programs written in an SPMD style will not check for this special case and may want to use `MPI_GET_COUNT` to check the status. (*End of rationale.*)

Advice to users. The buffer size required for the receive can be affected by data conversions and by the stride of the receive datatype. In most cases, the safest approach is to use the same datatype with `MPI_GET_COUNT` and the receive. (*End of advice to users.*)

All send and receive operations use the `buf`, `count`, `datatype`, `source`, `dest`, `tag`, `comm` and `status` arguments in the same way as the blocking `MPI_SEND` and `MPI_RECV` operations described in this section.

3.2.6 Passing `MPI_STATUS_IGNORE` for Status

Every call to `MPI_RECV` includes a `status` argument, wherein the system can return details about the message received. There are also a number of other MPI calls where `status` is returned. An object of type `MPI_STATUS` is not an MPI opaque object; its structure is declared in `mpi.h` and `mpif.h`, and it exists in the user’s program. In many cases, application programs are constructed so that it is unnecessary for them to examine the `status` fields. In these cases, it is a waste for the user to allocate a status object, and it is particularly wasteful for the MPI implementation to fill in fields in this object.

To cope with this problem, there are two predefined constants, `MPI_STATUS_IGNORE` and `MPI_STATUSES_IGNORE` with the C language bindings and the Fortran bindings through the `mpi` module and the `mpif.h` include file, which when passed to a receive, wait, or test function, inform the implementation that the status fields are not to be filled in. Note that `MPI_STATUS_IGNORE` is not a special type of `MPI_STATUS` object; rather, it is a special value for the argument. In C one would expect it to be `NULL`, not the address of a special `MPI_STATUS`.

`MPI_STATUS_IGNORE`, and the array version `MPI_STATUSES_IGNORE`, can be used everywhere a status argument is passed to a receive, wait, or test function. `MPI_STATUS_IGNORE`

cannot be used when status is an IN argument. Note that in Fortran `MPI_STATUS_IGNORE` and `MPI_STATUSES_IGNORE` are objects like `MPI_BOTTOM` (not usable for initialization or assignment). See Section 2.5.4.

In general, this optimization can apply to all functions for which `status` or an array of `statuses` is an OUT argument. Note that this converts `status` into an INOUT argument. The functions that can be passed `MPI_STATUS_IGNORE` are all the various forms of `MPI_RECV`, `MPI_TEST`, and `MPI_WAIT`, as well as `MPI_REQUEST_GET_STATUS`. When an array is passed, as in the `MPI_{TEST|WAIT}{ALL|SOME}` functions, a separate constant, `MPI_STATUSES_IGNORE`, is passed for the array argument. It is possible for an MPI function to return `MPI_ERR_IN_STATUS` even when `MPI_STATUS_IGNORE` or `MPI_STATUSES_IGNORE` has been passed to that function.

`MPI_STATUS_IGNORE` and `MPI_STATUSES_IGNORE` are not required to have the same values in C and Fortran.

It is not allowed to have some of the statuses in an array of statuses for `MPI_{TEST|WAIT}{ALL|SOME}` functions set to `MPI_STATUS_IGNORE`; one either specifies ignoring *all* of the statuses in such a call with `MPI_STATUSES_IGNORE`, or *none* of them by passing normal statuses in all positions in the array of statuses.

With the Fortran bindings through the `mpi_f08` module and the C++ bindings, `MPI_STATUS_IGNORE` or `MPI_STATUSES_IGNORE` does not exist. To allow an OUT or INOUT `TYPE(MPI_Status)` or `MPI::Status` argument to be ignored, all MPI `mpi_f08` or C++ bindings that have OUT or INOUT `TYPE(MPI_Status)` or `MPI::Status` parameters are overloaded with a second version that omits the OUT or INOUT `TYPE(MPI_Status)` or `MPI::Status` parameter.

Example 3.1 The `mpi_f08` bindings for `MPI_PROBE` are:

```
SUBROUTINE MPI_Probe(source, tag, comm, status, ierror)
  INTEGER, INTENT(IN) :: source, tag
  TYPE(MPI_Comm), INTENT(IN) :: comm
  TYPE(MPI_Status), INTENT(OUT) :: status
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
END SUBROUTINE

SUBROUTINE MPI_Probe(source, tag, comm, ierror)
  INTEGER, INTENT(IN) :: source, tag
  TYPE(MPI_Comm), INTENT(IN) :: comm
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
END SUBROUTINE
```

Example 3.2 The C++ bindings for `MPI_PROBE` are:

```
void MPI::Comm::Probe(int source, int tag, MPI::Status& status) const
void MPI::Comm::Probe(int source, int tag) const
```

3.3 Data Type Matching and Data Conversion

3.3.1 Type Matching Rules

One can think of message transfer as consisting of the following three phases.

1. Data is pulled out of the send buffer and a message is assembled.


```

MPI_Irecv(buf, count, datatype, source, tag, comm, request, ierror)
    <type> BUF(*)
    INTEGER COUNT, DATATYPE, SOURCE, TAG, COMM, REQUEST, IERROR
MPI_Irecv(buf, count, datatype, source, tag, comm, request, ierror)
    TYPE(*), DIMENSION(..) :: buf
    INTEGER, INTENT(IN) :: count, source, tag
    TYPE(MPI_Datatype), INTENT(IN) :: datatype
    TYPE(MPI_Comm), INTENT(IN) :: comm
    TYPE(MPI_Request), INTENT(OUT) :: request
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
{MPI::Request MPI::Comm::Irecv(void* buf, int count, const
    MPI::Datatype& datatype, int source, int tag) const(binding
    deprecated, see Section 15.2) }
```

Start a nonblocking receive.

These calls allocate a communication request object and associate it with the request handle (the argument `request`). The request can be used later to query the status of the communication or wait for its completion.

A nonblocking send call indicates that the system may start copying data out of the send buffer. The sender should not modify any part of the send buffer after a nonblocking send operation is called, until the send completes.

A nonblocking receive call indicates that the system may start writing data into the receive buffer. The receiver should not access any part of the receive buffer after a nonblocking receive operation is called, until the receive completes.

Advice to users. To prevent problems with the argument copying and register optimization done by Fortran compilers, please note the hints in subsections “Problems Due to Data Copying and Sequence Association,” and “A Problem with Register Optimization and Temporary Memory Modifications” in Section 16.2.2 on pages 545 and 549. (*End of advice to users.*)

3.7.3 Communication Completion

The functions `MPI_WAIT` and `MPI_TEST` are used to complete a nonblocking communication. The completion of a send operation indicates that the sender is now free to update the locations in the send buffer (the send operation itself leaves the content of the send buffer unchanged). It does not indicate that the message has been received, rather, it may have been buffered by the communication subsystem. However, if a synchronous mode send was used, the completion of the send operation indicates that a matching receive was initiated, and that the message will eventually be received by this matching receive.

The completion of a receive operation indicates that the receive buffer contains the received message, the receiver is now free to access it, and that the status object is set. It does not indicate that the matching send operation has completed (but indicates, of course, that the send was initiated).

We shall use the following terminology: A **null** handle is a handle with value `MPI_REQUEST_NULL`. A persistent request and the handle to it are **inactive** if the request is not associated with any ongoing communication (see Section 3.9). A handle is **active** if it is neither null nor inactive. An **empty** status is a status which is set to return `tag`

3.9 Persistent Communication Requests

Often a communication with the same argument list is repeatedly executed within the inner loop of a parallel computation. In such a situation, it may be possible to optimize the communication by binding the list of communication arguments to a **persistent** communication request once and, then, repeatedly using the request to initiate and complete messages. The persistent request thus created can be thought of as a communication port or a “half-channel.” It does not provide the full functionality of a conventional channel, since there is no binding of the send port to the receive port. This construct allows reduction of the overhead for communication between the process and communication controller, but not of the overhead for communication between one communication controller and another. It is not necessary that messages sent with a persistent request be received by a receive operation using a persistent request, or vice versa.

A persistent communication request is created using one of the five following calls. These calls involve no communication.

`MPI_SEND_INIT(buf, count, datatype, dest, tag, comm, request)`

| | | |
|-----|----------|--|
| IN | buf | initial address of send buffer (choice) |
| IN | count | number of elements sent (non-negative integer) |
| IN | datatype | type of each element (handle) |
| IN | dest | rank of destination (integer) |
| IN | tag | message tag (integer) |
| IN | comm | communicator (handle) |
| OUT | request | communication request (handle) |

```
int MPI_Send_init(void* buf, int count, MPI_Datatype datatype, int dest,
                 int tag, MPI_Comm comm, MPI_Request *request)
```

```
MPI_SEND_INIT(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
<type> BUF(*)
INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR
```

```
MPI_Send_init(buf, count, datatype, dest, tag, comm, request, ierror)
TYPE(*), DIMENSION(..) :: buf
INTEGER, INTENT(IN) :: count, dest, tag
TYPE(MPI_Datatype), INTENT(IN) :: datatype
TYPE(MPI_Comm), INTENT(IN) :: comm
TYPE(MPI_Request), INTENT(OUT) :: request
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

```
{MPI::Prequest MPI::Comm::Send_init(const void* buf, int count, const
MPI::Datatype& datatype, int dest, int tag) const(binding
deprecated, see Section 15.2) }
```

Creates a persistent communication request for a standard mode send operation, and binds to it all the arguments of a send operation.

```

MPI_BSEND_INIT(buf, count, datatype, dest, tag, comm, request)
    IN      buf                initial address of send buffer (choice)
    IN      count              number of elements sent (non-negative integer)
    IN      datatype           type of each element (handle)
    IN      dest               rank of destination (integer)
    IN      tag                message tag (integer)
    IN      comm               communicator (handle)
    OUT     request            communication request (handle)

int MPI_Bsend_init(void* buf, int count, MPI_Datatype datatype, int dest,
                  int tag, MPI_Comm comm, MPI_Request *request)

MPI_BSEND_INIT(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
    <type> BUF(*)
    INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR

MPI_Bsend_init(buf, count, datatype, dest, tag, comm, request, ierror)
    TYPE(*), DIMENSION(..) :: buf
    INTEGER, INTENT(IN) :: count, dest, tag
    TYPE(MPI_Datatype), INTENT(IN) :: datatype
    TYPE(MPI_Comm), INTENT(IN) :: comm
    TYPE(MPI_Request), INTENT(OUT) :: request
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

{MPI::Prequest MPI::Comm::Bsend_init(const void* buf, int count, const
    MPI::Datatype& datatype, int dest, int tag) const(binding
    deprecated, see Section 15.2) }

    Creates a persistent communication request for a buffered mode send.

MPI_SSEND_INIT(buf, count, datatype, dest, tag, comm, request)
    IN      buf                initial address of send buffer (choice)
    IN      count              number of elements sent (non-negative integer)
    IN      datatype           type of each element (handle)
    IN      dest               rank of destination (integer)
    IN      tag                message tag (integer)
    IN      comm               communicator (handle)
    OUT     request            communication request (handle)

int MPI_Ssend_init(void* buf, int count, MPI_Datatype datatype, int dest,
                  int tag, MPI_Comm comm, MPI_Request *request)

MPI_SSEND_INIT(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
    <type> BUF(*)

```

where * indicates zero or more repetitions. If the same communication object is used in several concurrent threads, it is the user's responsibility to coordinate calls so that the correct sequence is obeyed.

A send operation initiated with MPI_START can be matched with any receive operation and, likewise, a receive operation initiated with MPI_START can receive messages generated by any send operation.

Advice to users. To prevent problems with the argument copying and register optimization done by Fortran compilers, please note the hints in subsections “Problems Due to Data Copying and Sequence Association,” and “A Problem with Register Optimization and Temporary Memory Modifications” in Section 16.2.2 on pages 545 and 549. (*End of advice to users.*)

3.10 Send-Receive

The **send-receive** operations combine in one call the sending of a message to one destination and the receiving of another message, from another process. The two (source and destination) are possibly the same. A send-receive operation is very useful for executing a shift operation across a chain of processes. If blocking sends and receives are used for such a shift, then one needs to order the sends and receives correctly (for example, even processes send, then receive, odd processes receive first, then send) so as to prevent cyclic dependencies that may lead to deadlock. When a send-receive operation is used, the communication subsystem takes care of these issues. The send-receive operation can be used in conjunction with the functions described in Chapter 7 in order to perform shifts on various logical topologies. Also, a send-receive operation is useful for implementing remote procedure calls.

A message sent by a send-receive operation can be received by a regular receive operation or probed by a probe operation; a send-receive operation can receive a message sent by a regular send operation.

Advice to users. C users may be tempted to avoid the usage of `MPI_GET_ADDRESS` and rely on the availability of the address operator `&`. Note, however, that `& cast-expression` is a pointer, not an address. ISO C does not require that the value of a pointer (or the pointer cast to `int`) be the absolute address of the object pointed at — although this is commonly the case. Furthermore, referencing may not have a unique definition on machines with a segmented address space. The use of `MPI_GET_ADDRESS` to “reference” C variables guarantees portability to such machines as well. (*End of advice to users.*)

Advice to users. To prevent problems with the argument copying and register optimization done by Fortran compilers, please note the hints in subsections “Problems Due to Data Copying and Sequence Association,” and “A Problem with Register Optimization and Temporary Memory Modifications” in Section 16.2.2 on pages 545 and 549. (*End of advice to users.*)

The following auxiliary function provides useful information on derived datatypes.

`MPI_TYPE_SIZE(datatype, size)`

| | | |
|-----|----------|-------------------------|
| IN | datatype | datatype (handle) |
| OUT | size | datatype size (integer) |

`int MPI_Type_size(MPI_Datatype datatype, int *size)`

`MPI_TYPE_SIZE(DATATYPE, SIZE, IERROR)`

INTEGER DATATYPE, SIZE, IERROR

`MPI_Type_size(datatype, size, ierror)`

TYPE(MPI_Datatype), INTENT(IN) :: datatype

INTEGER, INTENT(OUT) :: size

INTEGER, OPTIONAL, INTENT(OUT) :: ierror

`{int MPI::Datatype::Get_size() const(binding deprecated, see Section 15.2) }`

`MPI_TYPE_SIZE` returns the total size, in bytes, of the entries in the type signature associated with `datatype`; i.e., the total size of the data in a message that would be created with this datatype. Entries that occur multiple times in the datatype are counted with their multiplicity.

4.1.6 Lower-Bound and Upper-Bound Markers

It is often convenient to define explicitly the lower bound and upper bound of a type map, and override the definition given on page 106. This allows one to define a datatype that has “holes” at its beginning or its end, or a datatype with entries that extend above the upper bound or below the lower bound. Examples of such usage are provided in Section 4.1.14. Also, the user may want to override the alignment rules that are used to compute upper bounds and extents. E.g., a C compiler may allow the user to override default alignment rules for some of the structures within a program. The user has to specify explicitly the bounds of the datatypes that match these structures.

4.1.10 Duplicating a Datatype

`MPI_TYPE_DUP(oldtype, newtype)`

| | | |
|-----|---------|---------------------------------------|
| IN | type | datatype (handle) |
| OUT | newtype | copy of <code>oldtype</code> (handle) |

`int MPI_Type_dup(MPI_Datatype oldtype, MPI_Datatype *newtype)`

`MPI_TYPE_DUP(OLDTYPE, NEWTYPE, IERROR)`

`INTEGER OLDTYPE, NEWTYPE, IERROR`

`MPI_Type_dup(oldtype, newtype, ierror)`

`TYPE(MPI_Datatype), INTENT(IN) :: oldtype`

`TYPE(MPI_Datatype), INTENT(OUT) :: newtype`

`INTEGER, OPTIONAL, INTENT(OUT) :: ierror`

`{MPI::Datatype MPI::Datatype::Dup() const(binding deprecated, see Section 15.2) }`

`MPI_TYPE_DUP` is a type constructor which duplicates the existing type with associated key values. For each key value, the respective copy callback function determines the attribute value associated with this key in the new communicator; one particular action that a copy callback may take is to delete the attribute from the new datatype. Returns in `newtype` a new datatype with exactly the same properties as `oldtype` and any copied cached information, see Section 6.7.4 on page 275. The new datatype has identical upper bound and lower bound and yields the same net result when fully decoded with the functions in Section 4.1.13. The `newtype` has the same committed state as the old `oldtype`.

4.1.11 Use of General Datatypes in Communication

Handles to derived datatypes can be passed to a communication call wherever a datatype argument is required. A call of the form `MPI_SEND(buf, count, datatype, ...)`, where `count > 1`, is interpreted as if the call was passed a new datatype which is the concatenation of `count` copies of `datatype`. Thus, `MPI_SEND(buf, count, datatype, dest, tag, comm)` is equivalent to,

`MPI_TYPE_CONTIGUOUS(count, datatype, newtype)`

`MPI_TYPE_COMMIT(newtype)`

`MPI_SEND(buf, 1, newtype, dest, tag, comm).`

Similar statements apply to all other communication functions that have a `count` and `datatype` argument.

Suppose that a send operation `MPI_SEND(buf, count, datatype, dest, tag, comm)` is executed, where `datatype` has type map,

$$\{(type_0, disp_0), \dots, (type_{n-1}, disp_{n-1})\},$$

and extent `extent`. (Empty entries of “pseudo-type” `MPI_UB` and `MPI_LB` are not listed in the type map, but they affect the value of `extent`.) The send operation sends $n \cdot \text{count}$

Rationale. The definition of MPI_MINLOC and MPI_MAXLOC given here has the advantage that it does not require any special-case handling of these two operations: they are handled like any other reduce operation. A programmer can provide his or her own definition of MPI_MAXLOC and MPI_MINLOC, if so desired. The disadvantage is that values and indices have to be first interleaved, and that indices and values have to be coerced to the same type, in Fortran. (*End of rationale.*)

5.9.5 User-Defined Reduction Operations

MPI_OP_CREATE(*user_fn*, commute, op)

| | | |
|-----|----------------|---------------------------------------|
| IN | <i>user_fn</i> | user defined function (function) |
| IN | commute | true if commutative; false otherwise. |
| OUT | op | operation (handle) |

```
int MPI_Op_create(MPI_User_function* user_fn, int commute, MPI_Op* op)
```

```
MPI_OP_CREATE( USER_FN, COMMUTE, OP, IERROR)
```

```
EXTERNAL USER_FN
LOGICAL COMMUTE
INTEGER OP, IERROR
```

```
MPI_Op_create(user_fn, commute, op, ierror)
```

```
EXTERNAL :: user_fn
LOGICAL, INTENT(IN) :: commute
TYPE(MPI_Op), INTENT(OUT) :: op
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

```
{void MPI::Op::Init(MPI::User_function* user_fn, bool commute) (binding
    deprecated, see Section 15.2) }
```

MPI_OP_CREATE binds a user-defined reduction operation to an op handle that can subsequently be used in MPI_REDUCE, MPI_ALLREDUCE, MPI_REDUCE_SCATTER, MPI_SCAN, and MPI_EXSCAN. The user-defined operation is assumed to be associative. If *commute* = true, then the operation should be both commutative and associative. If *commute* = false, then the order of operands is fixed and is defined to be in ascending, process rank order, beginning with process zero. The order of evaluation can be changed, taking advantage of the associativity of the operation. If *commute* = true then the order of evaluation can be changed, taking advantage of commutativity and associativity.

The argument *user_fn* is the user-defined function, which must have the following four arguments: invec, inoutvec, len and datatype.

The ISO C prototype for the function is the following.

```
typedef void MPI_User_function(void* invec, void* inoutvec, int* len,
    MPI_Datatype* datatype);
```

The Fortran declaration of the user-defined function appears below.

```
SUBROUTINE USER_FUNCTION(INVEC, INOUTVEC, LEN, TYPE)
    <type> INVEC(LEN), INOUTVEC(LEN)
```

INTEGER LEN, TYPE

The C++ declaration of the user-defined function appears below.

```
{typedef void MPI::User_function(const void* invec, void* inoutvec, int
    len, const Datatype& datatype); (binding deprecated, see
    Section 15.2)}
```

The `datatype` argument is a handle to the data type that was passed into the call to `MPI_REDUCE`. The user reduce function should be written such that the following holds: Let $u[0], \dots, u[\text{len}-1]$ be the `len` elements in the communication buffer described by the arguments `invec`, `len` and `datatype` when the function is invoked; let $v[0], \dots, v[\text{len}-1]$ be `len` elements in the communication buffer described by the arguments `inoutvec`, `len` and `datatype` when the function is invoked; let $w[0], \dots, w[\text{len}-1]$ be `len` elements in the communication buffer described by the arguments `inoutvec`, `len` and `datatype` when the function returns; then $w[i] = u[i] \circ v[i]$, for $i=0, \dots, \text{len}-1$, where \circ is the reduce operation that the function computes.

Informally, we can think of `invec` and `inoutvec` as arrays of `len` elements that `user_fn` is combining. The result of the reduction over-writes values in `inoutvec`, hence the name. Each invocation of the function results in the pointwise evaluation of the reduce operator on `len` elements: i.e., the function returns in `inoutvec[i]` the value `invec[i] \circ inoutvec[i]`, for $i = 0, \dots, \text{count} - 1$, where \circ is the combining operation computed by the function.

Rationale. The `len` argument allows `MPI_REDUCE` to avoid calling the function for each element in the input buffer. Rather, the system can choose to apply the function to chunks of input. In C, it is passed in as a reference for reasons of compatibility with Fortran.

By internally comparing the value of the `datatype` argument to known, global handles, it is possible to overload the use of a single user-defined function for several, different data types. (*End of rationale.*)

General datatypes may be passed to the user function. However, use of datatypes that are not contiguous is likely to lead to inefficiencies.

No MPI communication function may be called inside the user function. `MPI_ABORT` may be called inside the function in case of an error.

Advice to users. Suppose one defines a library of user-defined reduce functions that are overloaded: the `datatype` argument is used to select the right execution path at each invocation, according to the types of the operands. The user-defined reduce function cannot “decode” the `datatype` argument that it is passed, and cannot identify, by itself, the correspondence between the datatype handles and the datatype they represent. This correspondence was established when the datatypes were created. Before the library is used, a library initialization preamble must be executed. This preamble code will define the datatypes that are used by the library, and store handles to these datatypes in global, static variables that are shared by the user code and the library code.

The Fortran version of `MPI_REDUCE` will invoke a user-defined reduce function using the Fortran calling conventions and will pass a Fortran-type datatype argument; the C version will use C calling convention and the C representation of a datatype handle.

```

1 MPI_REDUCE_LOCAL( inbuf, inoutbuf, count, datatype, op)
2     IN          inbuf                input buffer (choice)
3
4     INOUT      inoutbuf             combined input and output buffer (choice)
5
6     IN          count                number of elements in inbuf and inoutbuf buffers (non-
7                                     negative integer)
8
9     IN          datatype             data type of elements of inbuf and inoutbuf buffers
10                                     (handle)
11
12     IN          op                   operation (handle)

```

```

12 int MPI_Reduce_local(void* inbuf, void* inoutbuf, int count,
13                     MPI_Datatype datatype, MPI_Op op)
14
15 MPI_REDUCE_LOCAL(INBUF, INOUTBUF, COUNT, DATATYPE, OP, IERROR)
16     <type> INBUF(*), INOUTBUF(*)
17     INTEGER COUNT, DATATYPE, OP, IERROR

```

```

18 MPI_Reduce_local(inbuf, inoutbuf, count, datatype, op, ierror)
19     TYPE(*), DIMENSION(..) :: inbuf, inoutbuf
20     INTEGER, INTENT(IN) :: count
21     TYPE(MPI_Datatype), INTENT(IN) :: datatype
22     TYPE(MPI_Op), INTENT(IN) :: op
23     INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

```

24 {void MPI::Op::Reduce_local(const void* inbuf, void* inoutbuf, int count,
25                             const MPI::Datatype& datatype) const(binding deprecated, see
26                             Section 15.2) }
27

```

The function applies the operation given by `op` element-wise to the elements of `inbuf` and `inoutbuf` with the result stored element-wise in `inoutbuf`, as explained for user-defined operations in Section 5.9.5. Both `inbuf` and `inoutbuf` (input as well as result) have the same number of elements given by `count` and the same datatype given by `datatype`. The `MPI_IN_PLACE` option is not allowed.

Reduction operations can be queried for their commutativity.

```

36 MPI_OP_COMMUTATIVE( op, commute)
37     IN          op                   operation (handle)
38
39     OUT         commute              true if op is commutative, false otherwise (logical)

```

```

40 int MPI_Op_commutative(MPI_Op op, int *commute)

```

```

42 MPI_OP_COMMUTATIVE(OP, COMMUTE, IERROR)
43     LOGICAL COMMUTE
44     INTEGER OP, IERROR

```

```

46 MPI_Op_commutative(op, commute, ierror)
47     TYPE(MPI_Op), INTENT(IN) :: op
48     LOGICAL, INTENT(OUT) :: commute

```


no pending communication on `peer_comm` that could interfere with this communication.

Advice to users. We recommend using a dedicated peer communicator, such as a duplicate of `MPI_COMM_WORLD`, to avoid trouble with peer communicators. (*End of advice to users.*)

`MPI_INTERCOMM_MERGE(intercomm, high, newintracomm)`

| | | |
|-----|--------------|---------------------------------|
| IN | intercomm | Inter-Communicator (handle) |
| IN | high | (logical) |
| OUT | newintracomm | new intra-communicator (handle) |

```
int MPI_Intercomm_merge(MPI_Comm intercomm, int high,
                        MPI_Comm *newintracomm)
```

```
MPI_INTERCOMM_MERGE(INTERCOMM, HIGH, NEWINTRACOMM, IERROR)
    INTEGER NEWINTERCOMM, INTRACOMM, IERROR
    LOGICAL HIGH
```

```
MPI_Intercomm_merge(intercomm, high, newintracomm, ierror)
    TYPE(MPI_Comm), INTENT(IN) :: intercomm
    LOGICAL, INTENT(IN) :: high
    TYPE(MPI_Comm), INTENT(OUT) :: newintracomm
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

```
{MPI::Intracomm MPI::Intercomm::Merge(bool high) const(binding deprecated, see
    Section 15.2) }
```

This function creates an intra-communicator from the union of the two groups that are associated with `intercomm`. All processes should provide the same `high` value within each of the two groups. If processes in one group provided the value `high = false` and processes in the other group provided the value `high = true` then the union orders the “low” group before the “high” group. If all processes provided the same `high` argument then the order of the union is arbitrary. This call is blocking and collective within the union of the two groups.

The error handler on the new intercommunicator in each process is inherited from the communicator that contributes the local group. Note that this can result in different processes in the same communicator having different error handlers.

Advice to implementors. The implementation of `MPI_INTERCOMM_MERGE`, `MPI_COMM_FREE` and `MPI_COMM_DUP` are similar to the implementation of `MPI_INTERCOMM_CREATE`, except that contexts private to the input inter-communicator are used for communication between group leaders rather than contexts inside a bridge communicator. (*End of advice to implementors.*)

- obtain a key value (used to identify an attribute); the user specifies “callback” functions by which MPI informs the application when the communicator is destroyed or copied.
- store and retrieve the value of an attribute;

Advice to implementors. Caching and callback functions are only called synchronously, in response to explicit application requests. This avoids problems that result from repeated crossings between user and system space. (This synchronous calling rule is a general property of MPI.)

The choice of key values is under control of MPI. This allows MPI to optimize its implementation of attribute sets. It also avoids conflict between independent modules caching information on the same communicators.

A much smaller interface, consisting of just a callback facility, would allow the entire caching facility to be implemented by portable code. However, with the minimal callback interface, some form of table searching is implied by the need to handle arbitrary communicators. In contrast, the more complete interface defined here permits rapid access to attributes through the use of pointers in communicators (to find the attribute table) and cleverly chosen key values (to retrieve individual attributes). In light of the efficiency “hit” inherent in the minimal interface, the more complete interface defined here is seen to be superior. (*End of advice to implementors.*)

MPI provides the following services related to caching. They are all process local.

6.7.2 Communicators

Functions for caching on communicators are:

`MPI_COMM_CREATE_KEYVAL(comm_copy_attr_fn, comm_delete_attr_fn, comm_keyval, extra_state)`

| | | |
|-----|----------------------------------|--|
| IN | <code>comm_copy_attr_fn</code> | copy callback function for <code>comm_keyval</code> (function) |
| IN | <code>comm_delete_attr_fn</code> | delete callback function for <code>comm_keyval</code> (function) |
| OUT | <code>comm_keyval</code> | key value for future access (integer) |
| IN | <code>extra_state</code> | extra state for callback functions |

```
int MPI_Comm_create_keyval(MPI_Comm_copy_attr_function *comm_copy_attr_fn,
                           MPI_Comm_delete_attr_function *comm_delete_attr_fn,
                           int *comm_keyval, void *extra_state)
```

```
MPI_COMM_CREATE_KEYVAL(COMM_COPY_ATTR_FN, COMM_DELETE_ATTR_FN, COMM_KEYVAL,
                        EXTRA_STATE, IERROR)
EXTERNAL COMM_COPY_ATTR_FN, COMM_DELETE_ATTR_FN
INTEGER COMM_KEYVAL, IERROR
INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
```

```
MPI_Comm_create_keyval(comm_copy_attr_fn, comm_delete_attr_fn, comm_keyval,
                        extra_state, ierror)
```

```

EXTERNAL :: comm_copy_attr_fn, comm_delete_attr_fn
INTEGER, INTENT(OUT) :: comm_keyval
INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: extra_state
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

{static int MPI::Comm::Create_keyval(MPI::Comm::Copy_attr_function*
    comm_copy_attr_fn,
    MPI::Comm::Delete_attr_function* comm_delete_attr_fn,
    void* extra_state) (binding deprecated, see Section 15.2) }
```

Generates a new attribute key. Keys are locally unique in a process, and opaque to user, though they are explicitly stored in integers. Once allocated, the key value can be used to associate attributes and access them on any locally defined communicator.

This function replaces MPI_KEYVAL_CREATE, whose use is deprecated. The C binding is identical. The Fortran binding differs in that `extra_state` is an address-sized integer. Also, the copy and delete callback functions have Fortran bindings that are consistent with address-sized attributes.

The C callback functions are:

```

typedef int MPI_Comm_copy_attr_function(MPI_Comm oldcomm, int comm_keyval,
    void *extra_state, void *attribute_val_in,
    void *attribute_val_out, int *flag);

and

typedef int MPI_Comm_delete_attr_function(MPI_Comm comm, int comm_keyval,
    void *attribute_val, void *extra_state);
```

which are the same as the MPI-1.1 calls but with a new name. The old names are deprecated.

The Fortran callback functions are:

```

SUBROUTINE COMM_COPY_ATTR_FUNCTION(OLDCOMM, COMM_KEYVAL, EXTRA_STATE,
    ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERROR)
    INTEGER OLDCOMM, COMM_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,
        ATTRIBUTE_VAL_OUT
    LOGICAL FLAG

and

SUBROUTINE COMM_DELETE_ATTR_FUNCTION(COMM, COMM_KEYVAL, ATTRIBUTE_VAL,
    EXTRA_STATE, IERROR)
    INTEGER COMM, COMM_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE
```

The C++ callbacks are:

```

{typedef int MPI::Comm::Copy_attr_function(const MPI::Comm& oldcomm,
    int comm_keyval, void* extra_state, void* attribute_val_in,
    void* attribute_val_out, bool& flag); (binding deprecated, see
    Section 15.2)}

and

{typedef int MPI::Comm::Delete_attr_function(MPI::Comm& comm,
    int comm_keyval, void* attribute_val, void* extra_state);
    (binding deprecated, see Section 15.2)}
```

The `comm_copy_attr_fn` function is invoked when a communicator is duplicated by `MPI_COMM_DUP`. `comm_copy_attr_fn` should be of type `MPI_Comm_copy_attr_function`. The copy callback function is invoked for each key value in `oldcomm` in arbitrary order. Each call to the copy callback is made with a key value and its corresponding attribute. If it returns `flag = 0`, then the attribute is deleted in the duplicated communicator. Otherwise (`flag = 1`), the new attribute value is set to the value returned in `attribute_val_out`. The function returns `MPI_SUCCESS` on success and an error code on failure (in which case `MPI_COMM_DUP` will fail).

The argument `comm_copy_attr_fn` may be specified as `MPI_COMM_NULL_COPY_FN` or `MPI_COMM_DUP_FN` from either C, C++, or Fortran. `MPI_COMM_NULL_COPY_FN` is a function that does nothing other than returning `flag = 0` and `MPI_SUCCESS`. `MPI_COMM_DUP_FN` is a simple-minded copy function that sets `flag = 1`, returns the value of `attribute_val_in` in `attribute_val_out`, and returns `MPI_SUCCESS`. These replace the MPI-1 predefined callbacks `MPI_NULL_COPY_FN` and `MPI_DUP_FN`, whose use is deprecated.

Advice to users. Even though both formal arguments `attribute_val_in` and `attribute_val_out` are of type `void *`, their usage differs. The C copy function is passed by MPI in `attribute_val_in` the *value* of the attribute, and in `attribute_val_out` the *address* of the attribute, so as to allow the function to return the (new) attribute value. The use of type `void *` for both is to avoid messy type casts.

A valid copy function is one that completely duplicates the information by making a full duplicate copy of the data structures implied by an attribute; another might just make another reference to that data structure, while using a reference-count mechanism. Other types of attributes might not copy at all (they might be specific to `oldcomm` only). (*End of advice to users.*)

Advice to implementors. A C interface should be assumed for copy and delete functions associated with key values created in C; a Fortran calling interface should be assumed for key values created in Fortran. (*End of advice to implementors.*)

Analogous to `comm_copy_attr_fn` is a callback deletion function, defined as follows. The `comm_delete_attr_fn` function is invoked when a communicator is deleted by `MPI_COMM_FREE` or when a call is made explicitly to `MPI_COMM_DELETE_ATTR`. `comm_delete_attr_fn` should be of type `MPI_Comm_delete_attr_function`.

This function is called by `MPI_COMM_FREE`, `MPI_COMM_DELETE_ATTR`, and `MPI_COMM_SET_ATTR` to do whatever is needed to remove an attribute. The function returns `MPI_SUCCESS` on success and an error code on failure (in which case `MPI_COMM_FREE` will fail).

The argument `comm_delete_attr_fn` may be specified as `MPI_COMM_NULL_DELETE_FN` from either C, C++, or Fortran. `MPI_COMM_NULL_DELETE_FN` is a function that does nothing, other than returning `MPI_SUCCESS`. `MPI_COMM_NULL_DELETE_FN` replaces `MPI_NULL_DELETE_FN`, whose use is deprecated.

If an attribute copy function or attribute delete function returns other than `MPI_SUCCESS`, then the call that caused it to be invoked (for example, `MPI_COMM_FREE`), is erroneous.

The special key value `MPI_KEYVAL_INVALID` is never returned by `MPI_KEYVAL_CREATE`. Therefore, it can be used for static initialization of key values.

Advice to implementors. To be able to use the predefined C functions `MPI_COMM_NULL_COPY_FN` or `MPI_COMM_DUP_FN` as `comm_copy_attr_fn` argument and/or `MPI_COMM_NULL_DELETE_FN` as the `comm_delete_attr_fn` argument in a call to the C++ routine `MPI::Comm::Create_keyval`, this routine may be overloaded with 3 additional routines that accept the C functions as the first, the second, or both input arguments (instead of an argument that matches the C++ prototype). *(End of advice to implementors.)*

Advice to users. If a user wants to write a “wrapper” routine that internally calls `MPI::Comm::Create_keyval` and `comm_copy_attr_fn` and/or `comm_delete_attr_fn` are arguments of this wrapper routine, and if this wrapper routine should be callable with both user-defined C++ copy and delete functions and with the predefined C functions, then the same overloading as described above in the advice to implementors may be necessary. (*End of advice to users.*)

MPI_COMM_FREE_KEYVAL(comm_keyval)

| | | |
|-------|-------------|---------------------|
| INOUT | comm_keyval | key value (integer) |
|-------|-------------|---------------------|

```
int MPI_Comm_free_keyval(int *comm_keyval)
```

`MPI_COMM_FREE_KEYVAL(COMM_KEYVAL, IERROR)`

```
INTEGER COMM_KEYVAL, IERROR
```

```
MPI_Comm_free_keyval(comm_keyval, ierror)
```

```
INTEGER, INTENT(INOUT) :: comm_keyval
```

```
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

```
{static void MPI::Comm::Free_keyval(int& comm_keyval) (binding deprecated, see  

Section 15.2) }
```

Frees an extant attribute key. This function sets the value of `keyval` to `MPI_KEYVAL_INVALID`. Note that it is not erroneous to free an attribute key that is in use, because the actual free does not transpire until after all references (in other communicators on the process) to the key have been freed. These references need to be explicitly freed by the program, either via calls to `MPI_COMM_DELETE_ATTR` that free one attribute instance, or by calls to `MPI_COMM_FREE` that free all attribute instances associated with the freed communicator.

This call is identical to the MPI-1 call `MPI_KEYVAL_FREE` but is needed to match the new communicator-specific creation function. The use of `MPI_KEYVAL_FREE` is deprecated.

MPI_COMM_SET_ATTR(comm, comm_keyval, attribute_val)

| | | |
|-------|------|--|
| INOUT | comm | communicator from which attribute will be attached (handle) |
|-------|------|--|

| | | |
|----|-------------|---------------------|
| IN | comm_keyval | key value (integer) |
|----|-------------|---------------------|

| IN | attribute_val | attribute value |
|----|---------------|-----------------|
|----|---------------|-----------------|

6.7.3 Windows

The new functions for caching on windows are:

`MPI_WIN_CREATE_KEYVAL(win_copy_attr_fn, win_delete_attr_fn, win_keyval, extra_state)`

| | | |
|-----|---------------------------------|---|
| IN | <code>win_copy_attr_fn</code> | copy callback function for <code>win_keyval</code> (function) |
| IN | <code>win_delete_attr_fn</code> | delete callback function for <code>win_keyval</code> (function) |
| OUT | <code>win_keyval</code> | key value for future access (integer) |
| IN | <code>extra_state</code> | extra state for callback functions |

```
int MPI_Win_create_keyval(MPI_Win_copy_attr_function *win_copy_attr_fn,
                          MPI_Win_delete_attr_function *win_delete_attr_fn,
                          int *win_keyval, void *extra_state)
```

```
MPI_WIN_CREATE_KEYVAL(WIN_COPY_ATTR_FN, WIN_DELETE_ATTR_FN, WIN_KEYVAL,
                      EXTRA_STATE, IERROR)
EXTERNAL WIN_COPY_ATTR_FN, WIN_DELETE_ATTR_FN
INTEGER WIN_KEYVAL, IERROR
INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
```

```
MPI_Win_create_keyval(win_copy_attr_fn, win_delete_attr_fn, win_keyval,
                      extra_state, ierror)
EXTERNAL :: win_copy_attr_fn, win_delete_attr_fn
INTEGER, INTENT(OUT) :: win_keyval
INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: extra_state
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

```
{static int MPI::Win::Create_keyval(MPI::Win::Copy_attr_function*
    win_copy_attr_fn,
    MPI::Win::Delete_attr_function* win_delete_attr_fn,
    void* extra_state) (binding deprecated, see Section 15.2) }
```

The argument `win_copy_attr_fn` may be specified as `MPI_WIN_NULL_COPY_FN` or `MPI_WIN_DUP_FN` from either C, C++, or Fortran. `MPI_WIN_NULL_COPY_FN` is a function that does nothing other than returning `flag = 0` and `MPI_SUCCESS`. `MPI_WIN_DUP_FN` is a simple-minded copy function that sets `flag = 1`, returns the value of `attribute_val_in` in `attribute_val_out`, and returns `MPI_SUCCESS`.

The argument `win_delete_attr_fn` may be specified as `MPI_WIN_NULL_DELETE_FN` from either C, C++, or Fortran. `MPI_WIN_NULL_DELETE_FN` is a function that does nothing, other than returning `MPI_SUCCESS`.

The C callback functions are:

```
typedef int MPI_Win_copy_attr_function(MPI_Win oldwin, int win_keyval,
    void *extra_state, void *attribute_val_in,
    void *attribute_val_out, int *flag);
```

and

```
typedef int MPI_Win_delete_attr_function(MPI_Win win, int win_keyval,
                                         void *attribute_val, void *extra_state);
```

The Fortran callback functions are:

```
SUBROUTINE WIN_COPY_ATTR_FUNCTION(OLDWIN, WIN_KEYVAL, EXTRA_STATE,
                                  ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERROR)
  INTEGER OLDWIN, WIN_KEYVAL, IERROR
  INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,
    ATTRIBUTE_VAL_OUT
  LOGICAL FLAG
```

and

```
SUBROUTINE WIN_DELETE_ATTR_FUNCTION(WIN, WIN_KEYVAL, ATTRIBUTE_VAL,
                                    EXTRA_STATE, IERROR)
  INTEGER WIN, WIN_KEYVAL, IERROR
  INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE
```

The C++ callbacks are:

```
{typedef int MPI::Win::Copy_attr_function(const MPI::Win& oldwin,
                                           int win_keyval, void* extra_state, void* attribute_val_in,
                                           void* attribute_val_out, bool& flag); (binding deprecated, see
                                           Section 15.2)}
```

and

```
{typedef int MPI::Win::Delete_attr_function(MPI::Win& win, int win_keyval,
                                             void* attribute_val, void* extra_state); (binding deprecated, see
                                             Section 15.2)}
```

If an attribute copy function or attribute delete function returns other than MPI_SUCCESS, then the call that caused it to be invoked (for example, MPI_WIN_FREE), is erroneous.

```
MPI_WIN_FREE_KEYVAL(win_keyval)
```

```
  INOUT    win_keyval          key value (integer)
```

```
int MPI_Win_free_keyval(int *win_keyval)
```

```
MPI_WIN_FREE_KEYVAL(WIN_KEYVAL, IERROR)
  INTEGER WIN_KEYVAL, IERROR
```

```
MPI_Win_free_keyval(win_keyval, ierror)
  INTEGER, INTENT(INOUT) :: win_keyval
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

```
{static void MPI::Win::Free_keyval(int& win_keyval) (binding deprecated, see
  Section 15.2) }
```

```

MPI_WIN_DELETE_ATTR(win, win_keyval)
    INOUT    win                window from which the attribute is deleted (handle)
    IN       win_keyval         key value (integer)

int MPI_Win_delete_attr(MPI_Win win, int win_keyval)

MPI_WIN_DELETE_ATTR(WIN, WIN_KEYVAL, IERROR)
    INTEGER WIN, WIN_KEYVAL, IERROR

MPI_Win_delete_attr(win, win_keyval, ierror)
    TYPE(MPI_Win), INTENT(IN) :: win
    INTEGER, INTENT(IN) :: win_keyval
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

{void MPI::Win::Delete_attr(int win_keyval) (binding deprecated, see Section 15.2)
    }

```

6.7.4 Datatypes

The new functions for caching on datatypes are:

```

MPI_TYPE_CREATE_KEYVAL(type_copy_attr_fn, type_delete_attr_fn, type_keyval, extra_state)

    IN       type_copy_attr_fn    copy callback function for type_keyval (function)
    IN       type_delete_attr_fn  delete callback function for type_keyval (function)
    OUT      type_keyval          key value for future access (integer)
    IN       extra_state          extra state for callback functions

int MPI_Type_create_keyval(MPI_Type_copy_attr_function *type_copy_attr_fn,
    MPI_Type_delete_attr_function *type_delete_attr_fn,
    int *type_keyval, void *extra_state)

MPI_TYPE_CREATE_KEYVAL(TYPE_COPY_ATTR_FN, TYPE_DELETE_ATTR_FN, TYPE_KEYVAL,
    EXTRA_STATE, IERROR)
    EXTERNAL TYPE_COPY_ATTR_FN, TYPE_DELETE_ATTR_FN
    INTEGER TYPE_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE

MPI_Type_create_keyval(type_copy_attr_fn, type_delete_attr_fn, type_keyval,
    extra_state, ierror)
    EXTERNAL :: type_copy_attr_fn, type_delete_attr_fn
    INTEGER, INTENT(OUT) :: type_keyval
    INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: extra_state
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

{static int MPI::Datatype::Create_keyval(MPI::Datatype::Copy_attr_function*
    type_copy_attr_fn, MPI::Datatype::Delete_attr_function*

```



```

1 MPI_TYPE_GET_ATTR(DATATYPE, TYPE_KEYVAL, ATTRIBUTE_VAL, FLAG, IERROR)
2     INTEGER DATATYPE, TYPE_KEYVAL, IERROR
3     INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL
4     LOGICAL FLAG
5
6 MPI_Type_get_attr(datatype, type_keyval, attribute_val, flag, ierror)
7     TYPE(MPI_Datatype), INTENT(IN) :: datatype
8     INTEGER, INTENT(IN) :: type_keyval
9     INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(OUT) :: attribute_val
10    LOGICAL, INTENT(OUT) :: flag
11    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
12
13 {bool MPI::Datatype::Get_attr(int type_keyval, void* attribute_val)
14     const(binding deprecated, see Section 15.2) }
15
16 MPI_TYPE_DELETE_ATTR(datatype, type_keyval)
17
18     INOUT    datatype                datatype from which the attribute is deleted (handle)
19     IN       type_keyval             key value (integer)
20
21 int MPI_Type_delete_attr(MPI_Datatype datatype, int type_keyval)
22
23 MPI_TYPE_DELETE_ATTR(DATATYPE, TYPE_KEYVAL, IERROR)
24     INTEGER DATATYPE, TYPE_KEYVAL, IERROR
25
26 MPI_Type_delete_attr(datatype, type_keyval, ierror)
27     TYPE(MPI_Datatype), INTENT(IN) :: datatype
28     INTEGER, INTENT(IN) :: type_keyval
29     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
30
31 {void MPI::Datatype::Delete_attr(int type_keyval)(binding deprecated, see
32     Section 15.2) }
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48

```

6.7.5 Error Class for Invalid Keyval

Key values for attributes are system-allocated, by `MPI_{TYPE,COMM,WIN}_CREATE_KEYVAL`. Only such values can be passed to the functions that use key values as input arguments. In order to signal that an erroneous key value has been passed to one of these functions, there is a new MPI error class: `MPI_ERR_KEYVAL`. It can be returned by `MPI_ATTR_PUT`, `MPI_ATTR_GET`, `MPI_ATTR_DELETE`, `MPI_KEYVAL_FREE`, `MPI_{TYPE,COMM,WIN}_DELETE_ATTR`, `MPI_{TYPE,COMM,WIN}_SET_ATTR`, `MPI_{TYPE,COMM,WIN}_GET_ATTR`, `MPI_{TYPE,COMM,WIN}_FREE_KEYVAL`, `MPI_COMM_DUP`, `MPI_COMM_DISCONNECT`, and `MPI_COMM_FREE`. The last three are included because `keyval` is an argument to the copy and delete functions for attributes.

6.7.6 Attributes Example

Advice to users. This example shows how to write a collective communication operation that uses caching to be more efficient after the first call. The coding style

reasons:

- It is not, in general, possible to store a string as an attribute from Fortran.
- It is not easy to set up the delete function for a string attribute unless it is known to have been allocated from the heap.
- To make the attribute key useful additional code to call `strdup` is necessary. If this is not standardized then users have to write it. This is extra unneeded work which we can easily eliminate.
- The Fortran binding is not trivial to write (it will depend on details of the Fortran compilation system), and will not be portable. Therefore it should be in the library rather than in user code.

(End of rationale.)

Advice to users. The above definition means that it is safe simply to print the string returned by `MPI_COMM_GET_NAME`, as it is always a valid string even if there was no name.

Note that associating a name with a communicator has no effect on the semantics of an MPI program, and will (necessarily) increase the store requirement of the program, since the names must be saved. Therefore there is no requirement that users use these functions to associate names with communicators. However debugging and profiling MPI applications may be made easier if names are associated with communicators, since the debugger or profiler should then be able to present information in a less cryptic manner. *(End of advice to users.)*

The following functions are used for setting and getting names of datatypes.

`MPI_TYPE_SET_NAME` (`datatype`, `type_name`)

| | | |
|-------|------------------------|---|
| INOUT | <code>datatype</code> | datatype whose identifier is to be set (handle) |
| IN | <code>type_name</code> | the character string which is remembered as the name (string) |

`int MPI_Type_set_name(MPI_Datatype datatype, char *type_name)`

`MPI_TYPE_SET_NAME(DATATYPE, TYPE_NAME, IERROR)`

INTEGER `DATATYPE`, IERROR

CHARACTER*(*) `TYPE_NAME`

`MPI_Type_set_name(datatype, type_name, ierror)`

`TYPE(MPI_Datatype), INTENT(IN) :: datatype`

`CHARACTER(LEN=*), INTENT(IN) :: type_name`

`INTEGER, OPTIONAL, INTENT(OUT) :: ierror`

`{void MPI::Datatype::Set_name(const char* type_name) (binding deprecated, see Section 15.2) }`

```

1 MPI_TYPE_GET_NAME (datatype, type_name, resultlen)
2     IN      datatype          datatype whose name is to be returned (handle)
3
4     OUT     type_name         the name previously stored on the datatype, or a empty
5                               string if no such name exists (string)
6
7     OUT     resultlen        length of returned name (integer)

```

```

8 int MPI_Type_get_name(MPI_Datatype datatype, char *type_name, int
9                       *resultlen)

```

```

10
11 MPI_TYPE_GET_NAME(DATATYPE, TYPE_NAME, RESULTLEN, IERROR)
12     INTEGER DATATYPE, RESULTLEN, IERROR
13     CHARACTER*(*) TYPE_NAME

```

```

14 MPI_Type_get_name(datatype, type_name, resultlen, ierror)
15     TYPE(MPI_Datatype), INTENT(IN) :: datatype
16     CHARACTER(LEN=*), INTENT(OUT) :: type_name
17     INTEGER, INTENT(OUT) :: resultlen
18     INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

```

19
20 {void MPI::Datatype::Get_name(char* type_name, int& resultlen) const(binding
21                               deprecated, see Section 15.2) }

```

Named predefined datatypes have the default names of the datatype name. For example, MPI_WCHAR has the default name of MPI_WCHAR.

The following functions are used for setting and getting names of windows.

```

27 MPI_WIN_SET_NAME (win, win_name)

```

```

28     INOUT   win              window whose identifier is to be set (handle)
29
30     IN      win_name         the character string which is remembered as the name
31                               (string)

```

```

33 int MPI_Win_set_name(MPI_Win win, char *win_name)

```

```

34 MPI_WIN_SET_NAME(WIN, WIN_NAME, IERROR)
35     INTEGER WIN, IERROR
36     CHARACTER*(*) WIN_NAME

```

```

38 MPI_Win_set_name(win, win_name, ierror)
39     TYPE(MPI_Win), INTENT(IN) :: win
40     CHARACTER(LEN=*), INTENT(IN) :: win_name
41     INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

```

42
43 {void MPI::Win::Set_name(const char* win_name) (binding deprecated, see
44                                                  Section 15.2) }

```

Chapter 8

MPI Environmental Management

This chapter discusses routines for getting and, where appropriate, setting various parameters that relate to the MPI implementation and the execution environment (such as error handling). The procedures for entering and leaving the MPI execution environment are also described here.

8.1 Implementation Information

8.1.1 Version Inquiries

In order to cope with changes to the MPI Standard, there are both compile-time and run-time ways to determine which version of the standard is in use in the environment one is using.

The “version” will be represented by two separate integers, for the version and subversion: In C and C++,

```
#define MPI_VERSION      2
#define MPI_SUBVERSION 2
```

in Fortran,

```
INTEGER :: MPI_VERSION, MPI_SUBVERSION
PARAMETER (MPI_VERSION      = 2)
PARAMETER (MPI_SUBVERSION = 2)
```

For runtime determination,

`MPI_GET_VERSION(version, subversion)`

| | | |
|-----|------------|-----------------------------|
| OUT | version | version number (integer) |
| OUT | subversion | subversion number (integer) |

```
int MPI_Get_version(int *version, int *subversion)
```

```
MPI_GET_VERSION(VERSION, SUBVERSION, IERROR)
```

```
INTEGER VERSION, SUBVERSION, IERROR
```

```
MPI_Get_version(version, subversion, ierror)
```

```

1  {void MPI::Get_processor_name(char* name, int& resultlen) (binding deprecated,
2      see Section 15.2) }
3

```

This routine returns the name of the processor on which it was called at the moment of the call. The name is a character string for maximum flexibility. From this value it must be possible to identify a specific piece of hardware; possible values include “processor 9 in rack 4 of mpp.cs.org” and “231” (where 231 is the actual processor number in the running homogeneous system). The argument `name` must represent storage that is at least `MPI_MAX_PROCESSOR_NAME` characters long. `MPI_GET_PROCESSOR_NAME` may write up to this many characters into `name`.

The number of characters actually written is returned in the output argument, `resultlen`. In C, a null character is additionally stored at `name[resultlen]`. The `resultlen` cannot be larger than `MPI_MAX_PROCESSOR_NAME-1`. In Fortran, `name` is padded on the right with blank characters. The `resultlen` cannot be larger than `MPI_MAX_PROCESSOR_NAME`.

Rationale. This function allows MPI implementations that do process migration to return the current processor. Note that nothing in MPI *requires* or defines process migration; this definition of `MPI_GET_PROCESSOR_NAME` simply allows such an implementation. (*End of rationale.*)

Advice to users. The user must provide at least `MPI_MAX_PROCESSOR_NAME` space to write the processor name — processor names can be this long. The user should examine the output argument, `resultlen`, to determine the actual length of the name. (*End of advice to users.*)

The constant `MPI_BSEND_OVERHEAD` provides an upper bound on the fixed overhead per message buffered by a call to `MPI_BSEND` (see Section 3.6.1).

8.2 Memory Allocation

In some systems, message-passing and remote-memory-access (RMA) operations run faster when accessing specially allocated memory (e.g., memory that is shared by the other processes in the communicating group on an SMP). MPI provides a mechanism for allocating and freeing such special memory. The use of such memory for message-passing or RMA is not mandatory, and this memory can be used without restrictions as any other dynamically allocated memory. However, implementations may restrict the use of the `MPI_WIN_LOCK` and `MPI_WIN_UNLOCK` functions to windows allocated in such memory (see Section 11.4.3.)

```

39 MPI_ALLOC_MEM(size, info, baseptr)
40

```

| | | | |
|----|-----|---------|--|
| 41 | IN | size | size of memory segment in bytes (non-negative integer) |
| 42 | | | |
| 43 | IN | info | info argument (handle) |
| 44 | OUT | baseptr | pointer to beginning of memory segment allocated |
| 45 | | | |

```

46 int MPI_Alloc_mem(MPI_Aint size, MPI_Info info, void *baseptr)
47

```

```

48 MPI_ALLOC_MEM(SIZE, INFO, BASEPTR, IERROR)

```

```

INTEGER INFO, IERROR
INTEGER(KIND=MPI_ADDRESS_KIND) SIZE, BASEPTR

MPI_Alloc_mem(size, info, baseptr, ierror)
  USE, INTRINSIC :: ISO_C_BINDING
  INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: size
  TYPE(MPI_Info), INTENT(IN) :: info
  TYPE(C_PTR), INTENT(OUT) :: baseptr
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

{void* MPI::Alloc_mem(MPI::Aint size, const MPI::Info& info) (binding
  deprecated, see Section 15.2) }
```

The `info` argument can be used to provide directives that control the desired location of the allocated memory. Such a directive does not affect the semantics of the call. Valid `info` values are implementation-dependent; a null directive value of `info = MPI_INFO_NULL` is always valid.

The function `MPI_ALLOC_MEM` may return an error code of class `MPI_ERR_NO_MEM` to indicate it failed because memory is exhausted.

```

MPI_FREE_MEM(base)
  IN      base          initial address of memory segment allocated by
                        MPI_ALLOC_MEM (choice)

int MPI_Free_mem(void *base)

MPI_FREE_MEM(BASE, IERROR)
  <type> BASE(*)
  INTEGER IERROR

MPI_Free_mem(base, ierror)
  USE, INTRINSIC :: ISO_C_BINDING
  TYPE(C_PTR), INTENT(IN) :: base
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

{void MPI::Free_mem(void *base) (binding deprecated, see Section 15.2) }
```

The function `MPI_FREE_MEM` may return an error code of class `MPI_ERR_BASE` to indicate an invalid base argument.

Rationale. The C and C++ bindings of `MPI_ALLOC_MEM` and `MPI_FREE_MEM` are similar to the bindings for the `malloc` and `free` C library calls: a call to `MPI_Alloc_mem(..., &base)` should be paired with a call to `MPI_Free_mem(base)` (one less level of indirection). Both arguments are declared to be of same type `void*` so as to facilitate type casting. The Fortran binding is consistent with the C and C++ bindings: the Fortran `MPI_ALLOC_MEM` call returns in `baseptr` the (integer valued) address of the allocated memory. The `base` argument of `MPI_FREE_MEM` is a choice argument, which passes (a reference to) the variable stored at that location. (*End of rationale.*)

Advice to implementors. If `MPI_ALLOC_MEM` allocates special memory, then a design similar to the design of C `malloc` and `free` functions has to be used, in order to find out the size of a memory segment, when the segment is freed. If no special memory is used, `MPI_ALLOC_MEM` simply invokes `malloc`, and `MPI_FREE_MEM` invokes `free`.

A call to `MPI_ALLOC_MEM` can be used in shared memory systems to allocate memory in a shared memory segment. (*End of advice to implementors.*)

Example 8.1

Example of use of `MPI_ALLOC_MEM`, in Fortran with pointer support. We assume 4-byte REALs, and assume that pointers are address-sized.

```

REAL A
POINTER (P, A(100,100))    ! no memory is allocated
CALL MPI_ALLOC_MEM(4*100*100, MPI_INFO_NULL, P, IERR)
! memory is allocated
...
A(3,5) = 2.71;
...
CALL MPI_FREE_MEM(A, IERR) ! memory is freed

```

Since standard Fortran does not support (C-like) pointers, this code is not Fortran 77 or Fortran 90 code. Some compilers (in particular, at the time of writing, g77 and Fortran compilers for Intel) do not support this code.

Example 8.2 Same example, in C

```

float (*f)[100][100] ;
/* no memory is allocated */
MPI_Alloc_mem(sizeof(float)*100*100, MPI_INFO_NULL, &f);
/* memory allocated */
...
(*f)[5][3] = 2.71;
...
MPI_Free_mem(f);

```

8.3 Error Handling

An MPI implementation cannot or may choose not to handle some errors that occur during MPI calls. These can include errors that generate exceptions or traps, such as floating point errors or access violations. The set of errors that are handled by MPI is implementation-dependent. Each such error generates an **MPI exception**.

The above text takes precedence over any text on error handling within this document. Specifically, text that states that errors *will* be handled should be read as *may* be handled.

A user can associate error handlers to three types of objects: communicators, windows, and files. The specified error handling routine will be used for any MPI exception that occurs during a call to MPI for the respective object. MPI calls that are not related to any objects are considered to be attached to the communicator `MPI_COMM_WORLD`. The attachment

The MPI function `MPI_ERRHANDLER_FREE` can be used to free an error handler that was created by a call to `MPI_XXX_CREATE_ERRHANDLER`.

`MPI_{COMM,WIN,FILE}_GET_ERRHANDLER` behave as if a new error handler object is created. That is, once the error handler is no longer needed, `MPI_ERRHANDLER_FREE` should be called with the error handler returned from `MPI_ERRHANDLER_GET` or `MPI_{COMM,WIN,FILE}_GET_ERRHANDLER` to mark the error handler for deallocation. This provides behavior similar to that of `MPI_COMM_GROUP` and `MPI_GROUP_FREE`.

Advice to implementors. High-quality implementation should raise an error when an error handler that was created by a call to `MPI_XXX_CREATE_ERRHANDLER` is attached to an object of the wrong type with a call to `MPI_YYY_SET_ERRHANDLER`. To do so, it is necessary to maintain, with each error handler, information on the typedef of the associated user function. (*End of advice to implementors.*)

The syntax for these calls is given below.

8.3.1 Error Handlers for Communicators

```
MPI_COMM_CREATE_ERRHANDLER(comm_errhandler_fn, errhandler)
    IN      comm_errhandler_fn      user defined error handling procedure (function)
    OUT     errhandler              MPI error handler (handle)

int MPI_Comm_create_errhandler(MPI_Comm_errhandler_function
                               *comm_errhandler_fn, MPI_Errhandler *errhandler)

MPI_COMM_CREATE_ERRHANDLER(COMM_ERRHANDLER_FN, ERRHANDLER, IERROR)
    EXTERNAL COMM_ERRHANDLER_FN
    INTEGER ERRHANDLER, IERROR

MPI_Comm_create_errhandler(comm_errhandler_fn, errhandler, ierror)
    EXTERNAL :: comm_errhandler_fn
    TYPE(MPI_Errhandler), INTENT(OUT) :: errhandler
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

{static MPI::Errhandler
    MPI::Comm::Create_errhandler(MPI::Comm::Errhandler_function*
                                comm_errhandler_fn) (binding deprecated, see Section 15.2) }
```

Creates an error handler that can be attached to communicators. This function is identical to `MPI_ERRHANDLER_CREATE`, whose use is deprecated.

The user routine should be, in C, a function of type `MPI_Comm_errhandler_function`, which is defined as

```
typedef void MPI_Comm_errhandler_function(MPI_Comm *, int *, ...);
```

The first argument is the communicator in use. The second is the error code to be returned by the MPI routine that raised the error. If the routine would have returned `MPI_ERR_IN_STATUS`, it is the error code returned in the status for the request that caused the error handler to be invoked. The remaining arguments are “`stdargs`” arguments whose

number and meaning is implementation-dependent. An implementation should clearly document these arguments. Addresses are used so that the handler may be written in Fortran. This typedef replaces `MPI_Handler_function`, whose use is deprecated.

In Fortran, the user routine should be of the form:

```
SUBROUTINE COMM_ERRHANDLER_FUNCTION(COMM, ERROR_CODE)
  INTEGER COMM, ERROR_CODE
```

In C++, the user routine should be of the form:

```
{typedef void MPI::Comm::Errhandler_function(MPI::Comm &, int *, ...);
  (binding deprecated, see Section 15.2)}
```

Rationale. The variable argument list is provided because it provides an ISO-standard hook for providing additional information to the error handler; without this hook, ISO C prohibits additional arguments. (*End of rationale.*)

Advice to users. A newly created communicator inherits the error handler that is associated with the “parent” communicator. In particular, the user can specify a “global” error handler for all communicators by associating this handler with the communicator `MPI_COMM_WORLD` immediately after initialization. (*End of advice to users.*)

`MPI_COMM_SET_ERRHANDLER(comm, errhandler)`

| | | |
|-------|------------|---|
| INOUT | comm | communicator (handle) |
| IN | errhandler | new error handler for communicator (handle) |

```
int MPI_Comm_set_errhandler(MPI_Comm comm, MPI_Errhandler errhandler)
```

```
MPI_COMM_SET_ERRHANDLER(COMM, ERRHANDLER, IERROR)
  INTEGER COMM, ERRHANDLER, IERROR
```

```
MPI_Comm_set_errhandler(comm, errhandler, ierror)
  TYPE(MPI_Comm), INTENT(IN) :: comm
  TYPE(MPI_Errhandler), INTENT(IN) :: errhandler
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

```
{void MPI::Comm::Set_errhandler(const MPI::Errhandler& errhandler) (binding
  deprecated, see Section 15.2) }
```

Attaches a new error handler to a communicator. The error handler must be either a predefined error handler, or an error handler created by a call to `MPI_COMM_CREATE_ERRHANDLER`. This call is identical to `MPI_ERRHANDLER_SET`, whose use is deprecated.

```

1 MPI_COMM_GET_ERRHANDLER(comm, errhandler)
2     IN      comm      communicator (handle)
3
4     OUT     errhandler  error handler currently associated with communicator
5                        (handle)
6

```

```

7 int MPI_Comm_get_errhandler(MPI_Comm comm, MPI_Errhandler *errhandler)
8

```

```

9 MPI_COMM_GET_ERRHANDLER(COMM, ERRHANDLER, IERROR)
10    INTEGER COMM, ERRHANDLER, IERROR
11

```

```

12 MPI_Comm_get_errhandler(comm, errhandler, ierror)
13    TYPE(MPI_Comm), INTENT(IN) :: comm
14    TYPE(MPI_Errhandler), INTENT(OUT) :: errhandler
15    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
16

```

```

17 {MPI::Errhandler MPI::Comm::Get_errhandler() const(binding deprecated, see
18    Section 15.2) }
19

```

Retrieves the error handler currently associated with a communicator. This call is identical to MPI_ERRHANDLER_GET, whose use is deprecated.

Example: A library function may register at its entry point the current error handler for a communicator, set its own private error handler for this communicator, and restore before exiting the previous error handler.

8.3.2 Error Handlers for Windows

```

27 MPI_WIN_CREATE_ERRHANDLER(win_errhandler_fn, errhandler)
28
29     IN      win_errhandler_fn  user defined error handling procedure (function)
30
31     OUT     errhandler          MPI error handler (handle)
32

```

```

33 int MPI_Win_create_errhandler(MPI_Win_errhandler_function
34    *win_errhandler_fn, MPI_Errhandler *errhandler)
35

```

```

36 MPI_WIN_CREATE_ERRHANDLER(WIN_ERRHANDLER_FN, ERRHANDLER, IERROR)
37    EXTERNAL WIN_ERRHANDLER_FN
38    INTEGER ERRHANDLER, IERROR
39

```

```

40 MPI_Win_create_errhandler(win_errhandler_fn, errhandler, ierror)
41    EXTERNAL :: win_errhandler_fn
42    TYPE(MPI_Errhandler), INTENT(OUT) :: errhandler
43    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
44

```

```

45 {static MPI::Errhandler
46    MPI::Win::Create_errhandler(MPI::Win::Errhandler_function*
47    win_errhandler_fn)(binding deprecated, see Section 15.2) }
48

```

Creates an error handler that can be attached to a window object. The user routine should be, in C, a function of type MPI_Win_errhandler_function which is defined as

```
typedef void MPI_Win_errhandler_function(MPI_Win *, int *, ...);
```

The first argument is the window in use, the second is the error code to be returned.

In Fortran, the user routine should be of the form:

```
SUBROUTINE WIN_ERRHANDLER_FUNCTION(WIN, ERROR_CODE)
  INTEGER WIN, ERROR_CODE
```

In C++, the user routine should be of the form:

```
{typedef void MPI::Win::Errhandler_function(MPI::Win &, int *, ...);
  (binding deprecated, see Section 15.2)}
```

```
MPI_WIN_SET_ERRHANDLER(win, errhandler)
```

| | | |
|-------|------------|---------------------------------------|
| INOUT | win | window (handle) |
| IN | errhandler | new error handler for window (handle) |

```
int MPI_Win_set_errhandler(MPI_Win win, MPI_Errhandler errhandler)
```

```
MPI_WIN_SET_ERRHANDLER(WIN, ERRHANDLER, IERROR)
  INTEGER WIN, ERRHANDLER, IERROR
```

```
MPI_Win_set_errhandler(win, errhandler, ierror)
  TYPE(MPI_Win), INTENT(IN) :: win
  TYPE(MPI_Errhandler), INTENT(IN) :: errhandler
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

```
{void MPI::Win::Set_errhandler(const MPI::Errhandler& errhandler) (binding
  deprecated, see Section 15.2) }
```

Attaches a new error handler to a window. The error handler must be either a pre-defined error handler, or an error handler created by a call to MPI_WIN_CREATE_ERRHANDLER.

```
MPI_WIN_GET_ERRHANDLER(win, errhandler)
```

| | | |
|-----|------------|---|
| IN | win | window (handle) |
| OUT | errhandler | error handler currently associated with window (handle) |

```
int MPI_Win_get_errhandler(MPI_Win win, MPI_Errhandler *errhandler)
```

```
MPI_WIN_GET_ERRHANDLER(WIN, ERRHANDLER, IERROR)
  INTEGER WIN, ERRHANDLER, IERROR
```

```
MPI_Win_get_errhandler(win, errhandler, ierror)
  TYPE(MPI_Win), INTENT(IN) :: win
  TYPE(MPI_Errhandler), INTENT(OUT) :: errhandler
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

```

1  {MPI::Errhandler MPI::Win::Get_errhandler() const(binding deprecated, see
2      Section 15.2) }

```

Retrieves the error handler currently associated with a window.

8.3.3 Error Handlers for Files

```

9  MPI_FILE_CREATE_ERRHANDLER(file_errhandler_fn, errhandler)

```

| | | | |
|----|-----|---------------------------|--|
| 10 | IN | <i>file_errhandler_fn</i> | user defined error handling procedure (function) |
| 11 | | | |
| 12 | OUT | errhandler | MPI error handler (handle) |

```

13
14  int MPI_File_create_errhandler(MPI_File_errhandler_function
15      *file_errhandler_fn, MPI_Errhandler *errhandler)

```

```

16  MPI_FILE_CREATE_ERRHANDLER(FILE_ERRHANDLER_FN, ERRHANDLER, IERROR)
17      EXTERNAL FILE_ERRHANDLER_FN
18      INTEGER ERRHANDLER, IERROR

```

```

19
20  MPI_File_create_errhandler(file_errhandler_fn, errhandler, ierror)
21      EXTERNAL :: file_errhandler_fn
22      TYPE(MPI_Errhandler), INTENT(OUT) :: errhandler
23      INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

```

24  {static MPI::Errhandler
25      MPI::File::Create_errhandler(MPI::File::Errhandler_function*
26          file_errhandler_fn)(binding deprecated, see Section 15.2) }
27

```

Creates an error handler that can be attached to a file object. The user routine should be, in C, a function of type MPI_File_errhandler_function, which is defined as

```

30  typedef void MPI_File_errhandler_function(MPI_File *, int *, ...);

```

The first argument is the file in use, the second is the error code to be returned.

In Fortran, the user routine should be of the form:

```

33  SUBROUTINE FILE_ERRHANDLER_FUNCTION(FILE, ERROR_CODE)
34      INTEGER FILE, ERROR_CODE

```

In C++, the user routine should be of the form:

```

37  {typedef void MPI::File::Errhandler_function(MPI::File &, int *, ...);
38      (binding deprecated, see Section 15.2)}

```

```

41  MPI_FILE_SET_ERRHANDLER(file, errhandler)

```

| | | | |
|----|-------|------------|-------------------------------------|
| 42 | INOUT | file | file (handle) |
| 43 | | | |
| 44 | IN | errhandler | new error handler for file (handle) |

```

45
46  int MPI_File_set_errhandler(MPI_File file, MPI_Errhandler errhandler)

```

```

47  MPI_FILE_SET_ERRHANDLER(FILE, ERRHANDLER, IERROR)
48

```

Chapter 9

The Info Object

Many of the routines in MPI take an argument `info`. `info` is an opaque object with a handle of type `MPI_Info` in C and Fortran with the `mpi_f08` module, `MPI::Info` in C++, and `INTEGER` in Fortran with the `mpi` module or the include file `mpif.h`. It stores an unordered set of (key,value) pairs (both key and value are strings). A key can have only one value. MPI reserves several keys and requires that if an implementation uses a reserved key, it must provide the specified functionality. An implementation is not required to support these keys and may support any others not reserved by MPI.

An implementation must support info objects as caches for arbitrary (key, value) pairs, regardless of whether it recognizes the key. Each function that takes hints in the form of an `MPI_Info` must be prepared to ignore any key it does not recognize. This description of info objects does not attempt to define how a particular function should react if it recognizes a key but not the associated value. `MPI_INFO_GET_NKEYS`, `MPI_INFO_GET_NTHKEY`, `MPI_INFO_GET_VALUELEN`, and `MPI_INFO_GET` must retain all (key,value) pairs so that layered functionality can also use the `Info` object.

Keys have an implementation-defined maximum length of `MPI_MAX_INFO_KEY`, which is at least 32 and at most 255. Values have an implementation-defined maximum length of `MPI_MAX_INFO_VAL`. In Fortran, leading and trailing spaces are stripped from both. Returned values will never be larger than these maximum lengths. Both key and value are case sensitive.

Rationale. Keys have a maximum length because the set of known keys will always be finite and known to the implementation and because there is no reason for keys to be complex. The small maximum size allows applications to declare keys of size `MPI_MAX_INFO_KEY`. The limitation on value sizes is so that an implementation is not forced to deal with arbitrarily long strings. (*End of rationale.*)

Advice to users. `MPI_MAX_INFO_VAL` might be very large, so it might not be wise to declare a string of that size. (*End of advice to users.*)

When it is an argument to a nonblocking routine, `info` is parsed before that routine returns, so that it may be modified or freed immediately after return.

When the descriptions refer to a key or value as being a boolean, an integer, or a list, they mean the string representation of these types. An implementation may define its own rules for how info value strings are converted to other types, but to ensure portability, every implementation must support the following representations. Legal values for a boolean must

```

int MPI_Comm_spawn(char *command, char *argv[], int maxprocs, MPI_Info
                    info, int root, MPI_Comm comm, MPI_Comm *intercomm,
                    int array_of_errcodes[])
MPI_COMM_SPAWN(COMMAND, ARGV, MAXPROCS, INFO, ROOT, COMM, INTERCOMM,
                ARRAY_OF_ERRCODES, IERROR)
CHARACTER*(*) COMMAND, ARGV(*)
INTEGER INFO, MAXPROCS, ROOT, COMM, INTERCOMM, ARRAY_OF_ERRCODES(*),
IERROR

MPI_Comm_spawn(command, argv, maxprocs, info, root, comm, intercomm,
                array_of_errcodes, ierror)
CHARACTER(LEN=*), INTENT(IN) :: command, argv(*)
INTEGER, INTENT(IN) :: maxprocs, root
TYPE(MPI_Info), INTENT(IN) :: info
TYPE(MPI_Comm), INTENT(IN) :: comm
TYPE(MPI_Comm), INTENT(OUT) :: intercomm
INTEGER, INTENT(OUT) :: array_of_errcodes(*) ! optional by
overloading
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

{MPI::Intercomm MPI::Intracomm::Spawn(const char* command,
    const char* argv[], int maxprocs, const MPI::Info& info,
    int root, int array_of_errcodes[]) const(binding deprecated, see
    Section 15.2) }

{MPI::Intercomm MPI::Intracomm::Spawn(const char* command,
    const char* argv[], int maxprocs, const MPI::Info& info,
    int root) const(binding deprecated, see Section 15.2) }

```

MPI_COMM_SPAWN tries to start maxprocs identical copies of the MPI program specified by command, establishing communication with them and returning an intercommunicator. The spawned processes are referred to as children. The children have their own MPI_COMM_WORLD, which is separate from that of the parents. MPI_COMM_SPAWN is collective over comm, and also may not return until MPI_INIT has been called in the children. Similarly, MPI_INIT in the children may not return until all parents have called MPI_COMM_SPAWN. In this sense, MPI_COMM_SPAWN in the parents and MPI_INIT in the children form a collective operation over the union of parent and child processes. The intercommunicator returned by MPI_COMM_SPAWN contains the parent processes in the local group and the child processes in the remote group. The ordering of processes in the local and remote groups is the same as the ordering of the group of the comm in the parents and of MPI_COMM_WORLD of the children, respectively. This intercommunicator can be obtained in the children through the function MPI_COMM_GET_PARENT.

Advice to users. An implementation may automatically establish communication before MPI_INIT is called by the children. Thus, completion of MPI_COMM_SPAWN in the parent does not necessarily mean that MPI_INIT has been called in the children (although the returned intercommunicator can be used immediately). (*End of advice to users.*)

Advice to users. By default, requests are hard and MPI errors are fatal. This means that by default there will be a fatal error if MPI cannot spawn all the requested processes. If you want the behavior “spawn as many processes as possible, up to N ,” you should do a soft spawn, where the set of allowed values $\{m_i\}$ is $\{0 \dots N\}$. However, this is not completely portable, as implementations are not required to support soft spawning. (*End of advice to users.*)

The info argument The info argument to all of the routines in this chapter is an opaque handle of type `MPI_Info` in C and Fortran with the `mpi_f08` module, `MPI::Info` in C++ and `INTEGER` in Fortran with the `mpi` module or the include file `mpif.h`. It is a container for a number of user-specified (key,value) pairs. key and value are strings (null-terminated `char*` in C, `character(*)` in Fortran). Routines to create and manipulate the info argument are described in Section 9 on page 349.

For the `SPAWN` calls, info provides additional (and possibly implementation-dependent) instructions to MPI and the runtime system on how to start processes. An application may pass `MPI_INFO_NULL` in C or Fortran. Portable programs not requiring detailed control over process locations should use `MPI_INFO_NULL`.

MPI does not specify the content of the info argument, except to reserve a number of special key values (see Section 10.3.4 on page 366). The info argument is quite flexible and could even be used, for example, to specify the executable and its command-line arguments. In this case the `command` argument to `MPI_COMM_SPAWN` could be empty. The ability to do this follows from the fact that MPI does not specify how an executable is found, and the info argument can tell the runtime system where to “find” the executable “” (empty string). Of course a program that does this will not be portable across MPI implementations.

The root argument All arguments before the root argument are examined only on the process whose rank in `comm` is equal to `root`. The value of these arguments on other processes is ignored.

The array_of_errcodes argument The `array_of_errcodes` is an array of length `maxprocs` in which MPI reports the status of each process that MPI was requested to start. If all `maxprocs` processes were spawned, `array_of_errcodes` is filled in with the value `MPI_SUCCESS`. If only m ($0 \leq m < \text{maxprocs}$) processes are spawned, m of the entries will contain `MPI_SUCCESS` and the rest will contain an implementation-specific error code indicating the reason MPI could not start the process. MPI does not specify which entries correspond to failed processes. An implementation may, for instance, fill in error codes in one-to-one correspondence with a detailed specification in the info argument. These error codes all belong to the error class `MPI_ERR_SPAWN` if there was no error in the argument list. In C or in the Fortran `mpi` module or `mpif.h` include file, an application may pass `MPI_ERRCODES_IGNORE` if it is not interested in the error codes. In the Fortran `mpi_f08` module or in C++ this constant does not exist, and the `array_of_errcodes` argument may be omitted from the argument list.

Advice to implementors. In the Fortran `mpi` module or `mpif.h` include file, `MPI_ERRCODES_IGNORE` is a special type of constant, like `MPI_BOTTOM`. See the discussion in Section 2.5.4 on page 15. In the Fortran `mpi_f08` module, the optional argument has to be implemented through function overloading. See the discussion in Section 2.5.2 on page 14. (*End of advice to implementors.*)

MPI standard, process 0 must return from the complete call after a bounded delay, even if process 1 does not reach any MPI call in this period of time. According to another interpretation, the complete call may block until process 1 reaches the wait call, or reaches another MPI call. The qualitative behavior is the same, under both interpretations, unless a process is caught in an infinite compute loop, in which case the difference may not matter. However, the quantitative expectations are different. Different MPI implementations reflect these different interpretations. While this ambiguity is unfortunate, it does not seem to affect many real codes. The MPI forum decided not to decide which interpretation of the standard is the correct one, since the issue is very contentious, and a decision would have much impact on implementors but less impact on users. (*End of rationale.*)

11.7.3 Registers and Compiler Optimizations

Advice to users. All the material in this section is an advice to users. (*End of advice to users.*)

A coherence problem exists between variables kept in registers and the memory value of these variables. An RMA call may access a variable in memory (or cache), while the up-to-date value of this variable is in register. A get will not return the latest variable value, and a put may be overwritten when the register is stored back in memory.

The problem is illustrated by the following code:

| Source of Process 1 | Source of Process 2 | Executed in Process 2 |
|-------------------------|---------------------|--------------------------|
| bbbb = 777 | buff = 999 | reg_A:=999 |
| call MPI_WIN_FENCE | call MPI_WIN_FENCE | |
| call MPI_PUT(bbbb | | stop appl.thread |
| into buff of process 2) | | buff:=777 in PUT handler |
| | | continue appl.thread |
| call MPI_WIN_FENCE | call MPI_WIN_FENCE | |
| | ccc = buff | ccc:=reg_A |

In this example, variable `buff` is allocated in the register `reg_A` and therefore `ccc` will have the old value of `buff` and not the new value 777.

This problem, which also afflicts in some cases send/receive communication, is discussed more at length in Section 16.2.2.

MPI implementations will avoid this problem for standard conforming C programs. Many Fortran compilers will avoid this problem, without disabling compiler optimizations. However, in order to avoid register coherence problems in a completely portable manner, users should restrict their use of RMA windows to variables stored in **in modules or COMMON** blocks, or to variables that were declared **VOLATILE** (but this attribute may inhibit optimization of any code containing the RMA window). Further details and additional solutions are discussed in Section 16.2.2, “A Problem with Register Optimization and Temporary Memory Modifications,” on page 549. See also, “Problems Due to Data Copying and Sequence Association,” on page 545, for additional Fortran problems.

For a generalized request, the operation associated with the request is performed by the application; therefore, the application must notify MPI when the operation completes. This is done by making a call to `MPI_GREQUEST_COMPLETE`. MPI maintains the “completion” status of generalized requests. Any other request state has to be maintained by the user.

A new generalized request is started with

```
MPI_GREQUEST_START(query_fn, free_fn, cancel_fn, extra_state, request)
```

| | | |
|-----|-------------|---|
| IN | query_fn | callback function invoked when request status is queried (function) |
| IN | free_fn | callback function invoked when request is freed (function) |
| IN | cancel_fn | callback function invoked when request is cancelled (function) |
| IN | extra_state | extra state |
| OUT | request | generalized request (handle) |

```
int MPI_Grequest_start(MPI_Grequest_query_function *query_fn,
                      MPI_Grequest_free_function *free_fn,
                      MPI_Grequest_cancel_function *cancel_fn, void *extra_state,
                      MPI_Request *request)
```

```
MPI_GREQUEST_START(QUERY_FN, FREE_FN, CANCEL_FN, EXTRA_STATE, REQUEST,
                  IERROR)
```

```
INTEGER REQUEST, IERROR
EXTERNAL QUERY_FN, FREE_FN, CANCEL_FN
INTEGER (KIND=MPI_ADDRESS_KIND) EXTRA_STATE
```

```
MPI_Grequest_start(query_fn, free_fn, cancel_fn, extra_state, request,
                  ierror)
```

```
EXTERNAL :: query_fn, free_fn, cancel_fn
INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: extra_state
TYPE(MPI_Request), INTENT(OUT) :: request
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

```
{static MPI::Grequest
    MPI::Grequest::Start(const MPI::Grequest::Query_function*
                        query_fn, const MPI::Grequest::Free_function* free_fn,
                        const MPI::Grequest::Cancel_function* cancel_fn,
                        void *extra_state) (binding deprecated, see Section 15.2) }
```

Advice to users. Note that a generalized request belongs, in C++, to the class `MPI::Grequest`, which is a derived class of `MPI::Request`. It is of the same type as regular requests, in C and Fortran. (*End of advice to users.*)

The call starts a generalized request and returns a handle to it in `request`.

The syntax and meaning of the callback functions are listed below. All callback functions are passed the `extra_state` argument that was associated with the request by the

starting call `MPI_GREQUEST_START`; `extra_state` can be used to maintain user-defined state for the request.

In C, the query function is

```
typedef int MPI_Grequest_query_function(void *extra_state,
                                       MPI_Status *status);
```

in Fortran

```
SUBROUTINE GREQUEST_QUERY_FUNCTION(EXTRA_STATE, STATUS, IERROR)
  INTEGER STATUS(MPI_STATUS_SIZE), IERROR
  INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
```

and in C++

```
{typedef int MPI::Grequest::Query_function(void* extra_state,
      MPI::Status& status); (binding deprecated, see Section 15.2)}
```

The `query_fn` function computes the status that should be returned for the generalized request. The status also includes information about successful/unsuccessful cancellation of the request (result to be returned by `MPI_TEST_CANCELLED`).

The `query_fn` callback is invoked by the `MPI_{WAIT|TEST}{ANY|SOME|ALL}` call that completed the generalized request associated with this callback. The callback function is also invoked by calls to `MPI_REQUEST_GET_STATUS`, if the request is complete when the call occurs. In both cases, the callback is passed a reference to the corresponding status variable passed by the user to the MPI call; the status set by the callback function is returned by the MPI call. If the user provided `MPI_STATUS_IGNORE` or `MPI_STATUSES_IGNORE` to the MPI function that causes `query_fn` to be called or has omitted the status argument (with the `mpi_f08` Fortran module or C++), then MPI will pass a valid status object to `query_fn`, and this status will be ignored upon return of the callback function. Note that `query_fn` is invoked only after `MPI_GREQUEST_COMPLETE` is called on the request; it may be invoked several times for the same generalized request, e.g., if the user calls `MPI_REQUEST_GET_STATUS` several times for this request. Note also that a call to `MPI_{WAIT|TEST}{SOME|ALL}` may cause multiple invocations of `query_fn` callback functions, one for each generalized request that is completed by the MPI call. The order of these invocations is not specified by MPI.

In C, the free function is

```
typedef int MPI_Grequest_free_function(void *extra_state);
```

and in Fortran

```
SUBROUTINE GREQUEST_FREE_FUNCTION(EXTRA_STATE, IERROR)
  INTEGER IERROR
  INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
```

and in C++

```
{typedef int MPI::Grequest::Free_function(void* extra_state); (binding
      deprecated, see Section 15.2)}
```

The `free_fn` function is invoked to clean up user-allocated resources when the generalized request is freed.

The `free_fn` callback is invoked by the `MPI_{WAIT|TEST}{ANY|SOME|ALL}` call that completed the generalized request associated with this callback. `free_fn` is invoked after

the call to `query_fn` for the same request. However, if the MPI call completed multiple generalized requests, the order in which `free_fn` callback functions are invoked is not specified by MPI.

The `free_fn` callback is also invoked for generalized requests that are freed by a call to `MPI_REQUEST_FREE` (no call to `WAIT_{WAIT|TEST}{ANY|SOME|ALL}` will occur for such a request). In this case, the callback function will be called either in the MPI call `MPI_REQUEST_FREE(request)`, or in the MPI call `MPI_GREQUEST_COMPLETE(request)`, whichever happens last, i.e., in this case the actual freeing code is executed as soon as both calls `MPI_REQUEST_FREE` and `MPI_GREQUEST_COMPLETE` have occurred. The `request` is not deallocated until after `free_fn` completes. Note that `free_fn` will be invoked only once per request by a correct program.

Advice to users. Calling `MPI_REQUEST_FREE(request)` will cause the `request` handle to be set to `MPI_REQUEST_NULL`. This handle to the generalized request is no longer valid. However, user copies of this handle are valid until after `free_fn` completes since MPI does not deallocate the object until then. Since `free_fn` is not called until after `MPI_GREQUEST_COMPLETE`, the user copy of the handle can be used to make this call. Users should note that MPI will deallocate the object after `free_fn` executes. At this point, user copies of the `request` handle no longer point to a valid request. MPI will not set user copies to `MPI_REQUEST_NULL` in this case, so it is up to the user to avoid accessing this stale handle. This is a special case in which MPI defers deallocating the object until a later time that is known by the user. (*End of advice to users.*)

In C, the cancel function is

```
typedef int MPI_Grequest_cancel_function(void *extra_state, int complete);
```

in Fortran

```
SUBROUTINE GREQUEST_CANCEL_FUNCTION(EXTRA_STATE, COMPLETE, IERROR)
  INTEGER IERROR
  INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
  LOGICAL COMPLETE
```

and in C++

```
{typedef int MPI::Grequest::Cancel_function(void* extra_state,
      bool complete); (binding deprecated, see Section 15.2)}
```

The `cancel_fn` function is invoked to start the cancelation of a generalized request. It is called by `MPI_CANCEL(request)`. MPI passes `complete=true` to the callback function if `MPI_GREQUEST_COMPLETE` was already called on the request, and `complete=false` otherwise.

All callback functions return an error code. The code is passed back and dealt with as appropriate for the error code by the MPI function that invoked the callback function. For example, if error codes are returned then the error code returned by the callback function will be returned by the MPI function that invoked the callback function. In the case of an `MPI_{WAIT|TEST}{ANY}` call that invokes both `query_fn` and `free_fn`, the MPI call will return the error code returned by the last callback, namely `free_fn`. If one or more of the requests in a call to `MPI_{WAIT|TEST}{SOME|ALL}` failed, then the MPI call will return `MPI_ERR_IN_STATUS`. In such a case, if the MPI call was passed an array of statuses, then

MPI will return in each of the statuses that correspond to a completed generalized request the error code returned by the corresponding invocation of its `free_fn` callback function. However, if the MPI function was passed `MPI_STATUSES_IGNORE` or the `status` argument was omitted, then the individual error codes returned by each callback functions will be lost.

Advice to users. `query_fn` must **not** set the error field of `status` since `query_fn` may be called by `MPI_WAIT` or `MPI_TEST`, in which case the error field of `status` should not change. The MPI library knows the “context” in which `query_fn` is invoked and can decide correctly when to put in the error field of `status` the returned error code. (*End of advice to users.*)

MPI_GREQUEST_COMPLETE(request)

| | | |
|-------|---------|------------------------------|
| INOUT | request | generalized request (handle) |
|-------|---------|------------------------------|

```
int MPI_Grequest_complete(MPI_Request request)
```

```

MPI GREQUEST COMPLETE(REQUEST, IERROR)

```

INTEGER REQUEST, IERROR

```
MPI_Grequest_complete(request, ierror)
```

```
TYPE(MPI_Request), INTENT(IN) :: request
```

```
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

```
{void MPI::Grequest::Complete() (binding deprecated, see Section 15.2) }
```

The call informs MPI that the operations represented by the generalized request `request` are complete (see definitions in Section 2.4). A call to `MPI_WAIT(request, status)` will return and a call to `MPI_TEST(request, flag, status)` will return `flag=true` only after a call to `MPI_GREQUEST_COMPLETE` has declared that these operations are complete.

MPI imposes no restrictions on the code executed by the callback functions. However, new nonblocking operations should be defined so that the general semantic rules about MPI calls such as `MPI_TEST`, `MPI_REQUEST_FREE`, or `MPI_CANCEL` still hold. For example, all these calls are supposed to be local and nonblocking. Therefore, the callback functions `query_fn`, `free_fn`, or `cancel_fn` should invoke blocking MPI communication calls only if the context is such that these calls are guaranteed to return in finite time. Once `MPI_CANCEL` is invoked, the cancelled operation should complete in finite time, irrespective of the state of other processes (the operation has acquired “local” semantics). It should either succeed, or fail without side-effects. The user should guarantee these same properties for newly defined operations.

Advice to implementors. A call to MPI_GREQUEST_COMPLETE may unblock a blocked user process/thread. The MPI library should ensure that the blocked user computation will resume. (*End of advice to implementors.*)

It is erroneous to access the local buffer of a nonblocking data access operation, or to use that buffer as the source or target of other communications, between the initiation and completion of the operation.

The split collective routines support a restricted form of “nonblocking” operations for collective data access (see Section 13.4.5, page 480).

Coordination

Every noncollective data access routine `MPI_FILE_XXX` has a collective counterpart. For most routines, this counterpart is `MPI_FILE_XXX_ALL` or a pair of `MPI_FILE_XXX_BEGIN` and `MPI_FILE_XXX_END`. The counterparts to the `MPI_FILE_XXX_SHARED` routines are `MPI_FILE_XXX_ORDERED`.

The completion of a noncollective call only depends on the activity of the calling process. However, the completion of a collective call (which must be called by all members of the process group) may depend on the activity of the other processes participating in the collective call. See Section 13.6.4, page 502, for rules on semantics of collective calls.

Collective operations may perform much better than their noncollective counterparts, as global data accesses have significant potential for automatic optimization.

Data Access Conventions

Data is moved between files and processes by calling read and write routines. Read routines move data from a file into memory. Write routines move data from memory into a file. The file is designated by a file handle, `fh`. The location of the file data is specified by an offset into the current view. The data in memory is specified by a triple: `buf`, `count`, and `datatype`. Upon completion, the amount of data accessed by the calling process is returned in a `status`.

An offset designates the starting position in the file for an access. The offset is always in etype units relative to the current view. Explicit offset routines pass `offset` as an argument (negative values are erroneous). The file pointer routines use implicit offsets maintained by MPI.

A data access routine attempts to transfer (read or write) `count` data items of type `datatype` between the user’s buffer `buf` and the file. The `datatype` passed to the routine must be a committed datatype. The layout of data in memory corresponding to `buf`, `count`, `datatype` is interpreted the same way as in MPI communication functions; see Section 3.2.2 on page 29 and Section 4.1.11 on page 111. The data is accessed from those parts of the file specified by the current view (Section 13.3, page 456). The type signature of `datatype` must match the type signature of some number of contiguous copies of the `etype` of the current view. As in a receive, it is erroneous to specify a `datatype` for reading that contains overlapping regions (areas of memory which would be stored into more than once).

The nonblocking data access routines indicate that MPI can start a data access and associate a request handle, `request`, with the I/O operation. Nonblocking operations are completed via `MPI_TEST`, `MPI_WAIT`, or any of their variants.

Data access operations, when completed, return the amount of data accessed in `status`.

Advice to users. To prevent problems with the argument copying and register optimization done by Fortran compilers, please note the hints in subsections “Problems Due to Data Copying and Sequence Association,” and “A Problem with Register Optimization and Temporary Memory Modifications” in Section 16.2.2, pages 545 and 549. (*End of advice to users.*)

For blocking routines, `status` is returned directly. For nonblocking routines and split collective routines, `status` is returned when the operation is completed. The number of `datatype` entries and predefined elements accessed by the calling process can be extracted from `status` by using `MPI_GET_COUNT` and `MPI_GET_ELEMENTS`, respectively. The interpretation of the `MPI_ERROR` field is the same as for other operations — normally undefined, but meaningful if an MPI routine returns `MPI_ERR_IN_STATUS`. The user can pass (in C and **with the Fortran `mpi module` or `mpif.h` include file**) `MPI_STATUS_IGNORE` in the `status` argument if the return value of this argument is not needed. **With the Fortran `mpi_f08 module` or in C++**, the `status` argument is optional. The `status` can be passed to `MPI_TEST_CANCELLED` to determine if the operation was cancelled. All other fields of `status` are undefined.

When reading, a program can detect the end of file by noting that the amount of data read is less than the amount requested. Writing past the end of file increases the file size. The amount of data accessed will be the amount requested, unless an error is raised (or a read reaches the end of file).

13.4.2 Data Access with Explicit Offsets

If `MPI_MODE_SEQUENTIAL` mode was specified when the file was opened, it is erroneous to call the routines in this section.

`MPI_FILE_READ_AT(fh, offset, buf, count, datatype, status)`

| | | |
|-----|----------|--|
| IN | fh | file handle (handle) |
| IN | offset | file offset (integer) |
| OUT | buf | initial address of buffer (choice) |
| IN | count | number of elements in buffer (integer) |
| IN | datatype | datatype of each buffer element (handle) |
| OUT | status | status object (Status) |

```
int MPI_File_read_at(MPI_File fh, MPI_Offset offset, void *buf, int count,
                    MPI_Datatype datatype, MPI_Status *status)
```

```
MPI_FILE_READ_AT(FH, OFFSET, BUF, COUNT, DATATYPE, STATUS, IERROR)
```

```
<type> BUF(*)
```

```
INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
```

```
INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
```

```
MPI_File_read_at(fh, offset, buf, count, datatype, status, ierror)
```

```
TYPE(MPI_File), INTENT(IN) :: fh
```

```
INTEGER(KIND=MPI_OFFSET_KIND), INTENT(IN) :: offset
```

```
TYPE(*), DIMENSION(..) :: buf
```

```
INTEGER, INTENT(IN) :: count
```

```
TYPE(MPI_Datatype), INTENT(IN) :: datatype
```

```
TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
```

```
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- On any MPI process, each file handle may have at most one active split collective operation at any time.
- Begin calls are collective over the group of processes that participated in the collective open and follow the ordering rules for collective calls.
- End calls are collective over the group of processes that participated in the collective open and follow the ordering rules for collective calls. Each end call matches the preceding begin call for the same collective operation. When an “end” call is made, exactly one unmatched “begin” call for the same operation must precede it.
- An implementation is free to implement any split collective data access routine using the corresponding blocking collective routine when either the begin call (e.g., `MPI_FILE_READ_ALL_BEGIN`) or the end call (e.g., `MPI_FILE_READ_ALL_END`) is issued. The begin and end calls are provided to allow the user and MPI implementation to optimize the collective operation.
- Split collective operations do not match the corresponding regular collective operation. For example, in a single collective read operation, an `MPI_FILE_READ_ALL` on one process does not match an `MPI_FILE_READ_ALL_BEGIN`/`MPI_FILE_READ_ALL_END` pair on another process.
- Split collective routines must specify a buffer in both the begin and end routines. By specifying the buffer that receives data in the end routine, we can avoid many (though not all) of the problems described in “A Problem with Register Optimization and Temporary Memory Modifications,” Section 16.2.2, page 549.
- No collective I/O operations are permitted on a file handle concurrently with a split collective access on that file handle (i.e., between the begin and end of the access). That is

```

MPI_File_read_all_begin(fh, ...);
...
MPI_File_read_all(fh, ...);
...
MPI_File_read_all_end(fh, ...);

```

is erroneous.

- In a multithreaded implementation, any split collective begin and end operation called by a process must be called from the same thread. This restriction is made to simplify the implementation in the multithreaded case. (Note that we have already disallowed having two threads begin a split collective operation on the same file handle since only one split collective operation can be active on a file handle at any time.)

The arguments for these routines have the same meaning as for the equivalent collective versions (e.g., the argument definitions for `MPI_FILE_READ_ALL_BEGIN` and `MPI_FILE_READ_ALL_END` are equivalent to the arguments for `MPI_FILE_READ_ALL`). The begin routine (e.g., `MPI_FILE_READ_ALL_BEGIN`) begins a split collective operation that, when completed with the matching end routine (i.e., `MPI_FILE_READ_ALL_END`)

Advice to implementors. When converting a larger size integer to a smaller size integer, only the less significant bytes are moved. Care must be taken to preserve the sign bit value. This allows no conversion errors if the data range is within the range of the smaller size integer. (*End of advice to implementors.*)

Table 13.2 specifies the sizes of predefined datatypes in “external32” format.

13.5.3 User-Defined Data Representations

There are two situations that cannot be handled by the required representations:

1. a user wants to write a file in a representation unknown to the implementation, and
2. a user wants to read a file written in a representation unknown to the implementation.

User-defined data representations allow the user to insert a third party converter into the I/O stream to do the data representation conversion.

```
MPI_REGISTER_DATAREP(datarep, read_conversion_fn, write_conversion_fn,
                      dtype_file_extent_fn, extra_state)
```

| | | |
|----|----------------------|--|
| IN | datarep | data representation identifier (string) |
| IN | read_conversion_fn | function invoked to convert from file representation to native representation (function) |
| IN | write_conversion_fn | function invoked to convert from native representation to file representation (function) |
| IN | dtype_file_extent_fn | function invoked to get the extent of a datatype as represented in the file (function) |
| IN | extra_state | extra state |

```
int MPI_Register_datarep(char *datarep,
                        MPI_Datarep_conversion_function *read_conversion_fn,
                        MPI_Datarep_conversion_function *write_conversion_fn,
                        MPI_Datarep_extent_function *dtype_file_extent_fn,
                        void *extra_state)

MPI_REGISTER_DATAREP(DATAREP, READ_CONVERSION_FN, WRITE_CONVERSION_FN,
                     DTYPE_FILE_EXTENT_FN, EXTRA_STATE, IERROR)
CHARACTER*(*) DATAREP
EXTERNAL READ_CONVERSION_FN, WRITE_CONVERSION_FN, DTYPE_FILE_EXTENT_FN
INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
INTEGER IERROR

MPI_Register_datarep(datarep, read_conversion_fn, write_conversion_fn,
                     dtype_file_extent_fn, extra_state, ierror)
CHARACTER(LEN=*), INTENT(IN) :: datarep
EXTERNAL :: read_conversion_fn, write_conversion_fn,
dtype_file_extent_fn
INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: extra_state
```

```

    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
{void MPI::Register_datarep(const char* datarep,
    MPI::Datarep_conversion_function* read_conversion_fn,
    MPI::Datarep_conversion_function* write_conversion_fn,
    MPI::Datarep_extent_function* dtype_file_extent_fn,
    void* extra_state) (binding deprecated, see Section 15.2) }
```

The call associates `read_conversion_fn`, `write_conversion_fn`, and `dtype_file_extent_fn` with the data representation identifier `datarep`. `datarep` can then be used as an argument to `MPI_FILE_SET_VIEW`, causing subsequent data access operations to call the conversion functions to convert all data items accessed between file data representation and native representation. `MPI_REGISTER_DATAREP` is a local operation and only registers the data representation for the calling MPI process. If `datarep` is already defined, an error in the error class `MPI_ERR_DUP_DATAREP` is raised using the default file error handler (see Section 13.7, page 508). The length of a data representation string is limited to the value of `MPI_MAX_DATAREP_STRING`. `MPI_MAX_DATAREP_STRING` must have a value of at least 64. No routines are provided to delete data representations and free the associated resources; it is not expected that an application will generate them in significant numbers.

Extent Callback

```

typedef int MPI_Datarep_extent_function(MPI_Datatype datatype,
    MPI_Aint *file_extent, void *extra_state);

SUBROUTINE DATAREP_EXTENT_FUNCTION(DATATYPE, EXTENT, EXTRA_STATE, IERROR)
    INTEGER DATATYPE, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTENT, EXTRA_STATE

{typedef void MPI::Datarep_extent_function(const MPI::Datatype& datatype,
    MPI::Aint& file_extent, void* extra_state); (binding deprecated,
    see Section 15.2)}
```

The function `dtype_file_extent_fn` must return, in `file_extent`, the number of bytes required to store `datatype` in the file representation. The function is passed, in `extra_state`, the argument that was passed to the `MPI_REGISTER_DATAREP` call. MPI will only call this routine with predefined datatypes employed by the user.

Datarep Conversion Functions

```

typedef int MPI_Datarep_conversion_function(void *userbuf,
    MPI_Datatype datatype, int count, void *filebuf,
    MPI_Offset position, void *extra_state);

SUBROUTINE DATAREP_CONVERSION_FUNCTION(USERBUF, DATATYPE, COUNT, FILEBUF,
    POSITION, EXTRA_STATE, IERROR)
    <TYPE> USERBUF(*), FILEBUF(*)
    INTEGER COUNT, DATATYPE, IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) POSITION
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
```

```

1 {typedef void MPI::Datarep_conversion_function(void* userbuf,
2         MPI::Datatype& datatype, int count, void* filebuf,
3         MPI::Offset position, void* extra_state); (binding deprecated, see
4         Section 15.2)}
```

The function `read_conversion_fn` must convert from file data representation to native representation. Before calling this routine, MPI allocates and fills `filebuf` with `count` contiguous data items. The type of each data item matches the corresponding entry for the predefined datatype in the type signature of `datatype`. The function is passed, in `extra_state`, the argument that was passed to the `MPI_REGISTER_DATAREP` call. The function must copy all `count` data items from `filebuf` to `userbuf` in the distribution described by `datatype`, converting each data item from file representation to native representation. `datatype` will be equivalent to the datatype that the user passed to the read function. If the size of `datatype` is less than the size of the `count` data items, the conversion function must treat `datatype` as being contiguously tiled over the `userbuf`. The conversion function must begin storing converted data at the location in `userbuf` specified by `position` into the (tiled) `datatype`.

Advice to users. Although the conversion functions have similarities to `MPI_PACK` and `MPI_UNPACK`, one should note the differences in the use of the arguments `count` and `position`. In the conversion functions, `count` is a count of data items (i.e., count of typemap entries of `datatype`), and `position` is an index into this typemap. In `MPI_PACK`, `incount` refers to the number of whole datatypes, and `position` is a number of bytes. (*End of advice to users.*)

Advice to implementors. A converted read operation could be implemented as follows:

1. Get file extent of all data items
2. Allocate a `filebuf` large enough to hold all `count` data items
3. Read data from file into `filebuf`
4. Call `read_conversion_fn` to convert data and place it into `userbuf`
5. Deallocate `filebuf`

(*End of advice to implementors.*)

If MPI cannot allocate a buffer large enough to hold all the data to be converted from a read operation, it may call the conversion function repeatedly using the same `datatype` and `userbuf`, and reading successive chunks of data to be converted in `filebuf`. For the first call (and in the case when all the data to be converted fits into `filebuf`), MPI will call the function with `position` set to zero. Data converted during this call will be stored in the `userbuf` according to the first `count` data items in `datatype`. Then in subsequent calls to the conversion function, MPI will increment the value in `position` by the `count` of items converted in the previous call, and the `userbuf` pointer will be unchanged.

Rationale. Passing the conversion function a `position` and one `datatype` for the transfer allows the conversion function to decode the datatype only once and cache an internal representation of it on the datatype. Then on subsequent calls, the conversion function can use the `position` to quickly find its place in the datatype and continue storing converted data where it left off at the end of the previous call. (*End of rationale.*)

13.6.7 MPI_Offset Type

MPI_Offset is an integer type of size sufficient to represent the size (in bytes) of the largest file supported by MPI. Displacements and offsets are always specified as values of type MPI_Offset.

In Fortran, the corresponding integer is an integer **with kind parameter MPI_OFFSET_KIND**, **which is defined in the `mpi_f08` module, the `mpi` module and the `mpif.h` include file.**

In Fortran 77 environments that do not support KIND parameters, MPI_Offset arguments should be declared as an INTEGER of suitable size. The language interoperability implications for MPI_Offset are similar to those for addresses (see Section 16.3, page 569).

13.6.8 Logical vs. Physical File Layout

MPI specifies how the data should be laid out in a virtual file structure (the view), not how that file structure is to be stored on one or more disks. Specification of the physical file structure was avoided because it is expected that the mapping of files to disks will be system specific, and any specific control over file layout would therefore restrict program portability. However, there are still cases where some information may be necessary to optimize file layout. This information can be provided as *hints* specified via *info* when a file is created (see Section 13.2.8, page 453).

13.6.9 File Size

The size of a file may be increased by writing to the file after the current end of file. The size may also be changed by calling MPI *size changing* routines, such as MPI_FILE_SET_SIZE. A call to a size changing routine does not necessarily change the file size. For example, calling MPI_FILE_PREALLOCATE with a size less than the current size does not change the size.

Consider a set of bytes that has been written to a file since the most recent call to a size changing routine, or since MPI_FILE_OPEN if no such routine has been called. Let the *high byte* be the byte in that set with the largest displacement. The file size is the larger of

- One plus the displacement of the high byte.
- The size immediately after the size changing routine, or MPI_FILE_OPEN, returned.

When applying consistency semantics, calls to MPI_FILE_SET_SIZE and MPI_FILE_PREALLOCATE are considered writes to the file (which conflict with operations that access bytes at displacements between the old and new file sizes), and MPI_FILE_GET_SIZE is considered a read of the file (which overlaps with all accesses to the file).

Advice to users. Any sequence of operations containing the collective routines MPI_FILE_SET_SIZE and MPI_FILE_PREALLOCATE is a write sequence. As such, sequential consistency in nonatomic mode is not guaranteed unless the conditions in Section 13.6.1, page 498, are satisfied. (*End of advice to users.*)

File pointer update semantics (i.e., file pointers are updated by the amount accessed) are only guaranteed if file size changes are sequentially consistent.

```
1 MPI_TYPE_UB( datatype, displacement)
```

```
2     IN          datatype          datatype (handle)
3
4     OUT         displacement      displacement of upper bound from origin, in bytes (in-
5                                     teger)
```

```
6
7 int MPI_Type_ub(MPI_Datatype datatype, MPI_Aint* displacement)
```

```
8 MPI_TYPE_UB( DATATYPE, DISPLACEMENT, IERROR)
9     INTEGER DATATYPE, DISPLACEMENT, IERROR
```

11 The following function is deprecated and is superseded by
 12 MPI_COMM_CREATE_KEYVAL in MPI-2.0. The language independent definition of the
 13 deprecated function is the same as that of the new function, except for the function name
 14 and a different behavior in the C/Fortran language interoperability, see Section 16.3.7 on
 15 page 579. The language bindings are modified.

```
16
17 MPI_KEYVAL_CREATE(copy_fn, delete_fn, keyval, extra_state)
```

```
19     IN          copy_fn          Copy callback function for keyval
20
21     IN          delete_fn       Delete callback function for keyval
22
23     OUT         keyval          key value for future access (integer)
24
25     IN          extra_state      Extra state for callback functions
```

```
26 int MPI_Keyval_create(MPI_Copy_function *copy_fn, MPI_Delete_function
27                       *delete_fn, int *keyval, void* extra_state)
```

```
28 MPI_KEYVAL_CREATE(COPY_FN, DELETE_FN, KEYVAL, EXTRA_STATE, IERROR)
29     EXTERNAL COPY_FN, DELETE_FN
30     INTEGER KEYVAL, EXTRA_STATE, IERROR
```

31 The copy_fn function is invoked when a communicator is duplicated by
 32 MPI_COMM_DUP. copy_fn should be of type MPI_Copy_function, which is defined as follows:

```
33
34
35 typedef int MPI_Copy_function(MPI_Comm oldcomm, int keyval,
36                               void *extra_state, void *attribute_val_in,
37                               void *attribute_val_out, int *flag)
```

38 A Fortran declaration for such a function is as follows:

```
39 SUBROUTINE COPY_FUNCTION(OLDCOMM, KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,
40                           ATTRIBUTE_VAL_OUT, FLAG, IERR)
41     INTEGER OLDCOMM, KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,
42     ATTRIBUTE_VAL_OUT, IERR
43     LOGICAL FLAG
```

44 copy_fn may be specified as MPI_NULL_COPY_FN or MPI_DUP_FN from either C or
 45 FORTRAN; MPI_NULL_COPY_FN is a function that does nothing other than returning
 46 flag = 0 and MPI_SUCCESS. MPI_DUP_FN is a simple-minded copy function that sets flag =

1, returns the value of `attribute_val_in` in `attribute_val_out`, and returns `MPI_SUCCESS`. Note that `MPI_NULL_COPY_FN` and `MPI_DUP_FN` are also deprecated.

Analogous to `copy_fn` is a callback deletion function, defined as follows. The `delete_fn` function is invoked when a communicator is deleted by `MPI_COMM_FREE` or when a call is made explicitly to `MPI_ATTR_DELETE`. `delete_fn` should be of type `MPI_Delete_function`, which is defined as follows:

```
typedef int MPI_Delete_function(MPI_Comm comm, int keyval,
void *attribute_val, void *extra_state);
```

A Fortran declaration for such a function is as follows:

```
SUBROUTINE DELETE_FUNCTION(COMM, KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE, IERR)
  INTEGER COMM, KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE, IERR
```

`delete_fn` may be specified as `MPI_NULL_DELETE_FN` from either C or FORTRAN; `MPI_NULL_DELETE_FN` is a function that does nothing, other than returning `MPI_SUCCESS`. Note that `MPI_NULL_DELETE_FN` is also deprecated.

The following function is deprecated and is superseded by `MPI_COMM_FREE_KEYVAL` in MPI-2.0. The language independent definition of the deprecated function is the same as of the new function, except of the function name. The language bindings are modified.

`MPI_KEYVAL_FREE(keyval)`

| | | |
|-------|--------|---------------------------------------|
| INOUT | keyval | Frees the integer key value (integer) |
|-------|--------|---------------------------------------|

```
int MPI_Keyval_free(int *keyval)
```

```
MPI_KEYVAL_FREE(KEYVAL, IERROR)
  INTEGER KEYVAL, IERROR
```

The following function is deprecated and is superseded by `MPI_COMM_SET_ATTR` in MPI-2.0. The language independent definition of the deprecated function is the same as of the new function, except of the function name. The language bindings are modified.

`MPI_ATTR_PUT(comm, keyval, attribute_val)`

| | | |
|-------|---------------|---|
| INOUT | comm | communicator to which attribute will be attached (handle) |
| IN | keyval | key value, as returned by <code>MPI_KEYVAL_CREATE</code> (integer) |
| IN | attribute_val | attribute value |

```
int MPI_Attr_put(MPI_Comm comm, int keyval, void* attribute_val)
```

```
MPI_ATTR_PUT(COMM, KEYVAL, ATTRIBUTE_VAL, IERROR)
  INTEGER COMM, KEYVAL, ATTRIBUTE_VAL, IERROR
```

The following function is deprecated and is superseded by `MPI_COMM_GET_ATTR` in MPI-2.0. The language independent definition of the deprecated function is the same as of the new function, except of the function name. The language bindings are modified.

1 MPI_ATTR_GET(comm, keyval, attribute_val, flag)

| | | | |
|---|-----|---------------|--|
| 2 | IN | comm | communicator to which attribute is attached (handle) |
| 3 | | | |
| 4 | IN | keyval | key value (integer) |
| 5 | OUT | attribute_val | attribute value, unless flag = false |
| 6 | OUT | flag | true if an attribute value was extracted; false if no attribute is associated with the key |
| 7 | | | |
| 8 | | | |

9
10 int MPI_Attr_get(MPI_Comm comm, int keyval, void *attribute_val, int *flag)

11 MPI_ATTR_GET(COMM, KEYVAL, ATTRIBUTE_VAL, FLAG, IERROR)

12 INTEGER COMM, KEYVAL, ATTRIBUTE_VAL, IERROR

13 LOGICAL FLAG

14
15 The following function is deprecated and is superseded by MPI_COMM_DELETE_ATTR
16 in MPI-2.0. The language independent definition of the deprecated function is the same as
17 of the new function, except of the function name. The language bindings are modified.

18
19 MPI_ATTR_DELETE(comm, keyval)

| | | | |
|----|-------|--------|--|
| 20 | INOUT | comm | communicator to which attribute is attached (handle) |
| 21 | | | |
| 22 | IN | keyval | The key value of the deleted attribute (integer) |
| 23 | | | |

24 int MPI_Attr_delete(MPI_Comm comm, int keyval)

25
26 MPI_ATTR_DELETE(COMM, KEYVAL, IERROR)

27 INTEGER COMM, KEYVAL, IERROR

28
29 The following function is deprecated and is superseded by
30 MPI_COMM_CREATE_ERRHANDLER in MPI-2.0. The language independent definition
31 of the deprecated function is the same as of the new function, except of the function name.
32 The language bindings are modified.

33
34 MPI_ERRHANDLER_CREATE(handler_fn, errhandler)

| | | | |
|----|-----|------------|---------------------------------------|
| 35 | IN | handler_fn | user defined error handling procedure |
| 36 | | | |
| 37 | OUT | errhandler | MPI error handler (handle) |
| 38 | | | |

39 int MPI_Errhandler_create(MPI_Handler_function *handler_fn,
40 MPI_Errhandler *errhandler)

41 MPI_ERRHANDLER_CREATE(HANDLER_FN, ERRHANDLER, IERROR)

42 EXTERNAL HANDLER_FN

43 INTEGER ERRHANDLER, IERROR

44
45 Register the user routine handler_fn for use as an MPI exception handler. Returns in
46 errhandler a handle to the registered exception handler.

47 In the C language, the user routine should be a C function of type MPI_Handler_function,
48 which is defined as:

```
typedef void (MPI_Handler_function)(MPI_Comm *, int *, ...);
```

The first argument is the communicator in use, the second is the error code to be returned.

In the Fortran language, the user routine should be of the form:

```
SUBROUTINE HANDLER_FUNCTION(COMM, ERROR_CODE)
  INTEGER COMM, ERROR_CODE
```

The following function is deprecated and is superseded by `MPI_COMM_SET_ERRHANDLER` in MPI-2.0. The language independent definition of the deprecated function is the same as of the new function, except of the function name. The language bindings are modified.

```
MPI_ERRHANDLER_SET( comm, errhandler )
```

| | | |
|-------|------------|--|
| INOUT | comm | communicator to set the error handler for (handle) |
| IN | errhandler | new MPI error handler for communicator (handle) |

```
int MPI_Errhandler_set(MPI_Comm comm, MPI_Errhandler errhandler)
```

```
MPI_ERRHANDLER_SET(COMM, ERRHANDLER, IERROR)
  INTEGER COMM, ERRHANDLER, IERROR
```

Associates the new error handler `errorhandler` with communicator `comm` at the calling process. Note that an error handler is always associated with the communicator.

The following function is deprecated and is superseded by `MPI_COMM_GET_ERRHANDLER` in MPI-2.0. The language independent definition of the deprecated function is the same as of the new function, except of the function name. The language bindings are modified.

```
MPI_ERRHANDLER_GET( comm, errhandler )
```

| | | |
|-----|------------|---|
| IN | comm | communicator to get the error handler from (handle) |
| OUT | errhandler | MPI error handler currently associated with communicator (handle) |

```
int MPI_Errhandler_get(MPI_Comm comm, MPI_Errhandler *errhandler)
```

```
MPI_ERRHANDLER_GET(COMM, ERRHANDLER, IERROR)
  INTEGER COMM, ERRHANDLER, IERROR
```

Returns in `errhandler` (a handle to) the error handler that is currently associated with communicator `comm`.

15.2 Deprecated since MPI-2.2

The entire set of C++ language bindings have been deprecated.

Example 16.10 `mpi_profile.cc`, to be compiled into `libpmpi.a`.

```
int MPI::Comm::Get_size() const
{
    // Do profiling stuff
    int ret = pmpi_comm.Get_size();
    // More profiling stuff
    return ret;
}
```

(End of advice to implementors.)

16.2 Fortran Support

16.2.1 Overview

The Fortran **MPI** language bindings have been designed to be compatible with the Fortran 90 standard (and later).

Rationale. Fortran 90 contains numerous features designed to make it a more “modern” language than Fortran 77. It seems natural that **MPI** should be able to take advantage of these new features with a set of bindings tailored to Fortran 90. In Fortran 2008, the only new language features used, are of assumed type and assumed rank. They were defined to allow the definition of choice arguments as part of the Fortran language. *(End of rationale.)*

MPI defines **three methods** of Fortran support:

1. **INCULDE 'mpif.h'** This method is described in Section 16.2.3. The use of the include file `mpif.h` is strongly discouraged since **MPI-3.0**.
2. **USE mpi** This method is described in Section 16.2.4 and requires compile-time argument checking.
3. **USE mpi_f08** This method is described in Section 16.2.5 and requires compile-time argument checking that includes also unique handle types.

A compliant **MPI-3** implementation providing a Fortran interface must provide **all three Fortran support methods**.

Application **subroutines and functions** may use either **one of the modules** or the `mpif.h` include file. An implementation may require use of the module to prevent type mismatch errors.

Advice to users. It is recommended to use **one of the MPI modules** even if it is not necessary to use it to avoid type mismatch errors on a particular system. Using a module provides several potential advantages over using an include file. *(End of advice to users.)*

In a single application, it must be possible to link together routines some of which `USE mpi` and others of which `USE mpi_f08` or `INCLUDE mpif.h`.

The `INTEGER` compile-time constant `MPI_SUBARRAYS` is `MPI_SUBARRAYS_SUPPORTED` if all choice arguments are defined in explicit interfaces with standardized assumed type and assumed rank, otherwise it equals `MPI_SUBARRAYS_UNSUPPORTED`. This constant exists with each Fortran support method, but not in the C/C++ header files. The value may be different for each Fortran support method.

Section 16.2.2 gives an overview on known problems when using Fortran together with MPI. Section 16.2.6 and Section 16.2.7 describe additional functionality that is part of the Fortran support. `funcMPI_F_SYNC_REG` is needed for one of the methods to prevent register optimization problems. A set of functions provides additional support for Fortran intrinsic numeric types, including parameterized types. The functions are: `MPI_SIZEOF`, `MPI_TYPE_MATCH_SIZE`, `MPI_TYPE_CREATE_F90_INTEGER`, `MPI_TYPE_CREATE_F90_REAL` and `MPI_TYPE_CREATE_F90_COMPLEX`. Parameterized types are Fortran intrinsic types which are specified using `KIND` type parameters.

16.2.2 Problems With Fortran Bindings for MPI

This section discusses a number of problems that may arise when using MPI in a Fortran program. It is intended as advice to users, and clarifies how MPI interacts with Fortran. It does not add to the standard, but is intended to clarify the standard.

As noted in the original MPI specification, the interface violates the Fortran standard in several ways. While these cause few problems for Fortran 77 programs, they become more significant for Fortran 90 programs, so that users must exercise care when using new Fortran 90 features. The violations were originally adopted and have been retained because they are important for the usability of MPI. The rest of this section describes the potential problems in detail.

The following MPI features are inconsistent with Fortran 90.

1. An MPI subroutine with a choice argument may be called with different argument types. Using the `mpi_f08` module, this problem is resolved.
2. An MPI subroutine with an assumed-size dummy argument may be passed an actual scalar argument.
3. Many MPI routines assume that actual arguments are passed by address and that arguments are not copied on entrance to or exit from the subroutine.
4. An MPI implementation may read or modify user data (e.g., communication buffers used by nonblocking communications) concurrently with a user program that is executing outside of MPI calls.
5. Several named “constants,” such as `MPI_BOTTOM`, `MPI_IN_PLACE`, `MPI_STATUS_IGNORE`, `MPI_STATUSES_IGNORE`, `MPI_ERRCODES_IGNORE`, `MPI_UNWEIGHTED`, `MPI_ARGV_NULL`, and `MPI_ARGVS_NULL` are not ordinary Fortran constants and require a special implementation. See Section 2.5.4 on page 15 for more information.
6. The memory allocation routine `MPI_ALLOC_MEM` can’t be usefully used in Fortran without a language extension that allows the allocated memory to be associated with a Fortran variable.

Additionally, MPI is inconsistent with Fortran 77 in a number of ways, as noted below.

- MPI identifiers exceed 6 characters.
- MPI identifiers may contain underscores after the first character.
- MPI requires an include file, `mpif.h`. On systems that do not support include files, the implementation should specify the values of named constants.
- Many routines in MPI have KIND-parameterized integers (e.g., `MPI_ADDRESS_KIND` and `MPI_OFFSET_KIND`) that hold address information. On systems that do not support Fortran 90-style parameterized types, `INTEGER*8` or `INTEGER` should be used instead.

MPI-1 contained several routines that take address-sized information as input or return address-sized information as output. In C such arguments were of type `MPI_Aint` and in Fortran of type `INTEGER`. On machines where integers are smaller than addresses, these routines can lose information. In MPI-2 the use of these functions has been deprecated and they have been replaced by routines taking `INTEGER` arguments of `KIND=MPI_ADDRESS_KIND`. A number of new MPI-2 functions also take `INTEGER` arguments of non-default `KIND`. See Section 2.6 on page 17 and Section 4.1.1 on page 87 for more information.

Problems Due to Strong Typing

All MPI functions with choice arguments associate actual arguments of different Fortran datatypes with the same dummy argument. This is not allowed by Fortran 77, and in Fortran 90 is technically only allowed if the function is overloaded with a different function for each type. In C, the use of `void*` formal arguments avoids these problems. *Similar to C, with Fortran 2008 and later together with the `mpi_f08` module, the problem is avoided by declaring choice arguments with `TYPE(*)`, `DIMENSION(..)`, i.e., as assumed type and assumed rank dummy arguments.*

Using `INCLUDE mpif.h`, the following code fragment *might* technically be invalid and may generate a compile-time error.

```
integer i(5)
real    x(5)
...
call mpi_send(x, 5, MPI_REAL, ...)
call mpi_send(i, 5, MPI_INTEGER, ...)
```

In practice, it is rare for compilers to do more than issue a warning. *Using the `mpi_f08` or `mpi` module, the problem is usually resolved through the standardized assume-type and assume-rank declarations of the dummy arguments, or with non-standard Fortran options preventing type checking for choice arguments.*

It is also technically *invalid* in Fortran to pass a scalar actual argument to an array dummy argument. Thus, *when using the `mpi_f08` or `mpi` module, the following code fragment usually generates an error since the `dims` and `periods` arguments to `MPI_CART_CREATE` are declared as assumed size arrays `INTEGER, DIMS(*)` and `LOGICAL, PERIODS(*)`.*

```

USE mpi_f08
INTEGER size
CALL MPI_Cart_create( comm_old,1,size,.TRUE.,.TRUE.,comm_cart,ierror )

```

Using INCLUDE 'mpif.h', compiler warnings are not expected except if this include file also uses Fortran explicit interfaces.

Problems Due to Data Copying and Sequence Association

- If MPI_SUBARRAYS equals MPI_SUBARRAYS_SUPPORTED:

Choice buffer arguments are declared as TYPE(*), DIMENSION(...). For example, considering the following code fragment:

```

REAL s(100), r(100)
CALL MPI_Isend(s(1:100:5), 3, MPI_REAL, ..., rq, ierror)
CALL MPI_Wait(rq, status, ierror)
CALL MPI_Irecv(r(1:100:5), 3, MPI_REAL, ..., rq, ierror)
CALL MPI_Wait(rq, status, ierror)

```

In this case, the individual elements `s(1)`, `s(6)`, `s(11)`, etc. are sent between the start of `MPI_ISEND` and the end of `MPI_WAIT` even though the compiled code may not copy `s(1:100:5)` to a contiguous temporary scratch buffer. Instead, the compiled code may pass a descriptor to `MPI_ISEND` that allows MPI to operate directly on `s(1)`, `s(6)`, `s(11)`, ..., `s(96)`.

All nonblocking MPI communication functions behave as if the user-specified elements of choice buffers are copied to a contiguous scratch buffer in the MPI runtime environment. All datatype descriptions (in the example above, “3, MPI_REAL”) read and store data from and to this virtual contiguous scratch buffer. Displacements in MPI derived datatypes are relative to the beginning of this virtual contiguous scratch buffer. Upon completion of a nonblocking receive operation (e.g., when `MPI_WAIT` on a corresponding `MPI_Request` returns), it is as if the received data has been copied from the virtual contiguous scratch buffer back to the non-contiguous application buffer. In the example above, `r(1)`, `r(6)`, and `r(11)` will be filled with the received data when `MPI_WAIT` returns.

Advice to implementors. The Fortran descriptor for TYPE(*), DIMENSION(...) arguments contains enough information that the MPI library can make a real contiguous copy of non-contiguous user buffers. Efficient implementations may avoid such additional memory-to-memory data copying. (*End of advice to implementors.*)

Rationale. If MPI_SUBARRAYS equals MPI_SUBARRAYS_SUPPORTED, non-contiguous buffers are handled inside of the MPI library instead of by the compiled user code. Therefore the scope of scratch buffers can be from the beginning of a nonblocking operation until the completion of the operation although beginning and completion are implemented in different routines. If MPI_SUBARRAYS equals MPI_SUBARRAYS_UNSUPPORTED, such scratch buffers can be organized only by the compiler for the duration of the nonblocking call, which is too short for implementing the whole MPI operation. (*End of rationale.*)

• If `MPI_SUBARRAYS` equals `MPI_SUBARRAYS_UNSUPPORTED`:

Implicit in MPI is the idea of a contiguous chunk of memory accessible through a linear address space. MPI copies data to and from this memory. An MPI program specifies the location of data by providing memory addresses and offsets. In the C language, sequence association rules plus pointers provide all the necessary low-level structure.

In Fortran 90, user data is not necessarily stored contiguously. For example, the array section `A(1:N:2)` involves only the elements of `A` with indices 1, 3, 5, The same is true for a pointer array whose target is such a section. Most compilers ensure that an array that is a dummy argument is held in contiguous memory if it is declared with an explicit shape (e.g., `B(N)`) or is of assumed size (e.g., `B(*)`). If necessary, they do this by making a copy of the array into contiguous memory. Both Fortran 77 and Fortran 90 are carefully worded to allow such copying to occur, but few Fortran 77 compilers do it.¹

Because MPI dummy buffer arguments are assumed-size arrays, this leads to a serious problem for a nonblocking call: the compiler copies the temporary array back on return but MPI continues to copy data to the memory that held it. For example, consider the following code fragment:

```
real a(100)
call MPI_IRECV(a(1:100:2), MPI_REAL, 50, ...)
```

Since the first dummy argument to `MPI_IRECV` is an assumed-size array (`<type> buf(*)`), the array section `a(1:100:2)` is copied to a temporary before being passed to `MPI_IRECV`, so that it is contiguous in memory. `MPI_IRECV` returns immediately, and data is copied from the temporary back into the array `a`. Sometime later, MPI may write to the address of the deallocated temporary. Copying is also a problem for `MPI_ISEND` since the temporary array may be deallocated before the data has all been sent from it.

Most Fortran 90 compilers do not make a copy if the actual argument is the whole of an explicit-shape or assumed-size array or is a ‘simple’ section such as `A(1:N)` of such an array. (We define ‘simple’ more fully in the next paragraph.) Also, many compilers treat allocatable arrays the same as they treat explicit-shape arrays in this regard (though we know of one that does not). However, the same is not true for assumed-shape and pointer arrays; since they may be discontinuous, copying is often done. It is this copying that causes problems for MPI as described in the previous paragraph.

Our formal definition of a ‘simple’ array section is

```
name ( [:,]... [<subscript>]:<subscript> [,<subscript>]... )
```

That is, there are zero or more dimensions that are selected in full, then one dimension selected without a stride, then zero or more dimensions that are selected with a simple subscript. Examples are

¹Technically, the Fortran standards are worded to allow non-contiguous storage of any array data.

```
A(1:N), A(:,N), A(:,1:N,1), A(1:6,N), A(:, :, 1:N)
```

Because of Fortran’s column-major ordering, where the first index varies fastest, a simple section of a contiguous array will also be contiguous.²

The same problem can occur with a scalar argument. Some compilers, even for Fortran 77, make a copy of some scalar dummy arguments within a called procedure. That this can cause a problem is illustrated by the example

```
call user1(a,rq)
call MPI_WAIT(rq,status,ierr)
write (*,*) a

subroutine user1(buf,request)
call MPI_IRecv(buf,...,request,...)
end
```

If `a` is copied, `MPI_IRecv` will alter the copy when it completes the communication and will not alter `a` itself.

Note that copying will almost certainly occur for an argument that is a non-trivial expression (one with at least one operator or function call), a section that does not select a contiguous part of its parent (e.g., `A(1:n:2)`), a pointer whose target is such a section, or an assumed-shape array that is (directly or indirectly) associated with such a section.

If there is a compiler option that inhibits copying of arguments, in either the calling or called procedure, this should be employed.

If a compiler makes copies in the calling procedure of arguments that are explicit-shape or assumed-size arrays, simple array sections of such arrays, or scalars, and if there is no compiler option to inhibit this, then the compiler cannot be used for applications that use `MPI_GET_ADDRESS`, or any nonblocking MPI routine. If a compiler copies scalar arguments in the called procedure and there is no compiler option to inhibit this, then this compiler cannot be used for applications that use memory references across subroutine calls as in the example above.

Special Constants

MPI requires a number of special “constants” that cannot be implemented as normal Fortran constants, e.g., `MPI_BOTTOM`. The complete list can be found in Section 2.5.4 on page 15. In C, these are implemented as constant pointers, usually as `NULL` and are used where the function prototype calls for a pointer to a variable, not the variable itself.

In Fortran the implementation of these special constants may require the use of language constructs that are outside the Fortran standard. Using special values for the constants (e.g., by defining them through `parameter` statements) is not possible because an

²To keep the definition of ‘simple’ simple, we have chosen to require all but one of the section subscripts to be without bounds. A colon without bounds makes it obvious both to the compiler and to the reader that the whole of the dimension is selected. It would have been possible to allow cases where the whole dimension is selected with one or two bounds, but this means for the reader that the array declaration or most recent allocation has to be consulted and for the compiler that a run-time check may be required.

implementation cannot distinguish these values from legal data. Typically these constants are implemented as predefined static variables (e.g., a variable in an MPI-declared `COMMON` block), relying on the fact that the target compiler passes data by address. Inside the subroutine, this address can be extracted by some mechanism outside the Fortran standard (e.g., by Fortran extensions or by implementing the function in C). With `USE mpi_f08`, the attributes `INTENT(IN)`, `INTENT(OUT)`, and `INTENT(INOUT)` are used in the Fortran interface. In most cases `INTENT(IN)` is used if the C interface uses call-by-value. For all buffer arguments and for `OUT` dummy arguments that allow one of these special constants as input, an `INTENT(...)` is not specified.

Fortran Derived Types

MPI does explicitly support passing Fortran derived types to choice dummy arguments, but does not support Fortran non-sequence derived types.

The following code fragment shows one possible way to send a `sequence` derived type in Fortran. The example assumes that all data is passed by address.

```

type mytype
  SEQUENCE
  integer i
  real x
  double precision d
end type mytype

type(mytype) foo
integer blocklen(3), type(3)
integer(MPI_ADDRESS_KIND) disp(3), base

call MPI_GET_ADDRESS(foo%i, disp(1), ierr)
call MPI_GET_ADDRESS(foo%x, disp(2), ierr)
call MPI_GET_ADDRESS(foo%d, disp(3), ierr)

base = disp(1)
disp(1) = disp(1) - base
disp(2) = disp(2) - base
disp(3) = disp(3) - base

blocklen(1) = 1
blocklen(2) = 1
blocklen(3) = 1

type(1) = MPI_INTEGER
type(2) = MPI_REAL
type(3) = MPI_DOUBLE_PRECISION

call MPI_TYPE_CREATE_STRUCT(3, blocklen, disp, type, newtype, ierr)
call MPI_TYPE_COMMIT(newtype, ierr)

```

```
call MPI_SEND(foo%i, 1, newtype, ...)
```

A Problem with Register Optimization and Temporary Memory Modifications

MPI provides operations that may be hidden from the user code and run concurrently with it, accessing the same memory as user code. Examples include the data transfer for an `MPI_IRECV`. The optimizer of a compiler will assume that it can recognize periods when a copy of a variable can be kept in a register without reloading from or storing to memory. When the user code is working with a register copy of some variable while the hidden operation reads or writes the memory copy, problems occur. This section discusses register optimization pitfalls and problems with temporary memory modifications. These problems are independent of the Fortran support method, i.e., they occur with the `mpi_f08` module, the `mpi` module, and the `mif.h` include file.

This section shows four problematic usage areas (the abbreviations in parentheses are used in the table below):

- Usage of nonblocking routines (*Nonbl.*).
- Usage of one-sided routines (*1-sided*).
- Usage of MPI parallel file I/O split collective operations (*Split*).
- Use of `MPI_BOTTOM` together with absolute displacements in MPI datatypes, or relative displacements between two variables in such datatypes (*Bottom*).

The compiler is allowed to cause two optimization problems

- Register optimization problems and code movements (*Regis.*).
- Temporary memory modifications (*Memory*).

The optimization problems can occur not in all usage areas:

| | Nonbl. | 1-sided | Split | Bottom |
|----------|--------|---------|--------|--------|
| Register | occurs | occurs | -not- | occurs |
| Memory | occurs | occurs | occurs | -not- |

The application writer has several methods to circumvent parts of these problems with special declarations for the used send and receive buffers:

- Usage of the Fortran `ASYNCHRONOUS` attribute.
- Usage of the Fortran `TARGET` attribute.
- Usage of the helper routine `MPI_F_SYNC_REG` or a user-written dummy routine `DD(buf)`.
- Declaring the buffer as Fortran module data or within a Fortran common block.
- Usage of the Fortran `VOLATILE` attribute.

Each of these methods may solve only a subset of the problems, may have more or less performance drawbacks, and may be usable not in each application context. The following table shows the usability of each method:

| | Nonbl. Regis. | Nonbl. Memory | 1-sided Regis. | 1-sided Memory | Split Memo. | Bottom Regis. | overhead may be |
|----------------|------------------|------------------|-------------------|-------------------|----------------|------------------|--------------------|
| Examples | 16.11, 16.12 | 16.13 | Section 11.7.3 | | | 16.14, 16.15 | |
| ASYNCHRONOUS | solved | solved | may be | may be | solved | may be | medium |
| TARGET | solved | NOT s. | solved | NOT s. | NOT s. | solved | low-med |
| MPI_F_SYNC_REG | solved | NOT s. | solved | NOT s. | NOT s. | solved | low |
| Module Data | solved | NOT s. | solved | NOT s. | NOT s. | solved | low-med |
| VOLATILE | solved | solved | solved | solved | solved | solved | high |

The next paragraphs describe the problems in detail.

When a variable is local to a Fortran subroutine (i.e., not in a module or **COMMON block**), the compiler will assume that it cannot be modified by a called subroutine unless it is an actual argument of the call. In the most common linkage convention, the subroutine is expected to save and restore certain registers. Thus, the optimizer will assume that a register which held a valid copy of such a variable before the call will still hold a valid copy on return.

Example 16.11 shows extreme, but allowed, possibilities.

Example 16.11 Fortran 90 register optimization – extreme.

| Source | compiled as | or compiled as |
|--|--|--|
| <code>REAL :: buf, b1</code> | <code>REAL :: buf, b1</code> | <code>REAL :: buf, b1</code> |
| <code>call MPI_Irecv(buf,..req)</code> | <code>call MPI_Irecv(buf,..req)</code> | <code>call MPI_Irecv(buf,..req)</code> |
| | <code>register = buf</code> | <code>b1 = buf</code> |
| <code>call MPI_WAIT(req,..)</code> | <code>call MPI_WAIT(req,..)</code> | <code>call MPI_WAIT(req,..)</code> |
| <code>b1 = buf</code> | <code>b1 := register</code> | |

`MPI_WAIT` on a concurrent thread modifies `buf` between the invocation of `MPI_Irecv` and the finish of `MPI_WAIT`. But the compiler cannot see any possibility that `buf` can be changed after `MPI_Irecv` has returned, and may schedule the load of `buf` earlier than typed in the source. It has no reason to avoid using a register to hold `buf` across the call to `MPI_WAIT`. It also may reorder the instructions as in the case on the right.

Due to allowed code movement, the content of `buf` may be already overwritten when sending of the content of `buf` is executed. The code movement is permitted, because the compiler cannot detect a possible access to `buf` in `MPI_WAIT` (or in a second thread between the start of `MPI_ISEND` and the end of `MPI_WAIT`).

Note, that code movement can also be executed across subroutine boundaries when subroutines or functions are inlined.

This register optimization / code movement problem does not occur with MPI parallel file I/O split collective operations, because in the `..._BEGIN` and `..._END` calls, the same buffer has to be provided as actual argument.

Example 16.12 Similar example with MPI_ISEND

| Source | compiled as | or compiled as |
|---------------------------|---------------------------|---------------------|
| REAL :: buf, copy | REAL :: buf, copy | REAL :: buf, copy |
| buf = val | buf = val | buf = val |
| call MPI_ISEND(buf,..req) | call MPI_ISEND(buf,..req) | addr = &buf |
| copy = buf | copy=buf | copy = val |
| | buf = val_overwrite | buf = val_overwrite |
| call MPI_WAIT(req,..) | call MPI_WAIT(req,..) | send(*addr) |
| buf = val_overwrite | | |

Nonblocking operations and temporary memory modifications. The compiler is allowed to modify temporarily data in the memory. Normally, this problem may occur only if overlapping communication and computation. Example 16.13 shows a possibility.

Example 16.13 Overlapping Communication and Computation

```

USE mpi_f08
REAL :: buf(100,100)
CALL MPI_Irecv(buf(1,1:100),...req,...)
DO j=1,100
  DO i=2,100
    buf(i,j)=...
  END DO
END DO
CALL MPI_Wait(req,...)

```

The compiler may substitute the nested loops through loop fusion by

```

EQUIVALENCE (buf(1,1), buf_1dim(1))
DO h=1,100
  tmp(h)=buf(1,h)
END DO
DO j=1,10000
  buf_1dim(h)=...
END DO
DO h=1,100
  buf(1,h)=tmp(h)
END DO

```

In the substitution of Example 16.13, `buf_1dim(10000)` is the 1-dimensional equivalence of `buf(100,100)`. The nonblocking receive may receive the data in the boundary `buf(1,1:100)` while the fused loop is using temporarily this part of the buffer. When the `tmp` data is written back to `buf`, the old data is restored and the received data is lost.

Note, that this problem occurs also

- with one-sided communication with the local buffer at the origin process between an RMA call and the ensuing synchronization call

- and with the window buffer at the target process between two ensuing synchronization calls,
- and also with MPI parallel file I/O split collective operations with the local buffer between the `..._BEGIN` and `..._END` call.

This type of compiler optimization can be prevented when `buf` is declared with the Fortran attribute `ASYNCHRONOUS`:

```
REAL, ASYNCHRONOUS :: buf(100,100)
```

One-sided communication. An example with instruction reordering due to register optimization can be found in Section 11.7.3 on page 425.

One-sided communication. Normally users are not afflicted with this. But the user should pay attention to this section if in his/her program a buffer argument to an `MPI_SEND`, `MPI_RECV` etc., uses a name which hides the actual variables involved. `MPI_BOTTOM` with an `MPI_Datatype` containing absolute addresses is one example. Creating a datatype which uses one variable as an anchor and brings along others by using `MPI_GET_ADDRESS` to determine their offsets from the anchor is another. The anchor variable would be the only one mentioned in the call. Also attention must be paid if MPI operations are used that run in parallel with the user's application.

Example 16.14 shows what Fortran compilers are allowed to do.

Example 16.14 Fortran 90 register optimization.

| | |
|---|---|
| <p>This source ...</p> <pre>call MPI_GET_ADDRESS(buf,bufaddr, ierror) call MPI_TYPE_CREATE_STRUCT(1,1, bufaddr, MPI_REAL,type,ierror) call MPI_TYPE_COMMIT(type,ierror) val_old = buf call MPI_RECV(MPI_BOTTOM,1,type,...) val_new = buf</pre> | <p>can be compiled as:</p> <pre>call MPI_GET_ADDRESS(buf,...) call MPI_TYPE_CREATE_STRUCT(...) call MPI_TYPE_COMMIT(...) register = buf val_old = register call MPI_RECV(MPI_BOTTOM,...) val_new = register</pre> |
|---|---|

The compiler does not invalidate the register because it cannot see that `MPI_RECV` changes the value of `buf`. The access of `buf` is hidden by the use of `MPI_GET_ADDRESS` and `MPI_BOTTOM`.

Several successive assignments to the same variable can be combined in this way, but only the last assignment is executed. Successive means that no interfering read access to this variable is in between. The compiler cannot detect that the call to `MPI_SEND` statement is interfering, because the read access to `buf` is hidden by the usage of `MPI_BOTTOM`.

Example 16.15 Similar example with MPI_SEND

This source ...

```
! buf contains val_old
buf = val_new

call MPI_SEND(MPI_BOTTOM,1,type,...)
! with buf as a displacement in type
buf = val_overwrite
```

can be compiled as:

```
! buf contains val_old
! dead code:
!   buf=val_new is removed
call MPI_SEND(...)
! i.e. val_old is sent
buf = val_overwrite
```

Solutions. The following paragraphs show in detail how these problems can be solved in portabel way. Several solutions are presented, because all of these solutions have different implication on the performance. Only one solution (with **VOLATILE**) solves all problems, but it may have the most negative impact on the performance.

Fortran ASYNCHRONOUS attribute. Declaring a buffer with the Fortran **ASYNCHRONOUS** attribute in a scoping unit (or **BLOCK**) tells the compiler that any statement of the scoping unit may be executed while the buffer is affected by a pending asynchronous input/output operation. Each library call (e.g., to an MPI routine) within the scoping unit may contain a Fortran asynchronous I/O statement, e.g., the Fortran **WAIT** statement.

- In the case of nonblocking MPI communication, the send and receive buffers should be declared with the Fortran **ASYNCHRONOUS** attribute within each scoping unit (or **BLOCK**) where the buffers are declared and statements are executed between the start (e.g., **MPI_IRECV**) and completion (e.g., **MPI_WAIT**) of the nonblocking communication. Declaring **REAL, ASYNCHRONOUS :: buf** in Examples 16.11 and 16.12, and **REAL, ASYNCHRONOUS :: buf(100,100)** in Examples 16.13 solves the register optimization and temporary memory modification problem.

Rationale. A combination of a nonblocking MPI communication call with a buffer in the argument list together with a subsequent call to **MPI_WAIT** or **MPI_TEST** is similar to the combination a of Fortran asynchronous read or write together with the matching Fortran wait statement. To prevent incorrect register optimizations or code movement, the Fortran standard requires in the case of Fortran IO, that the **ASYNCHRONOUS** attribute is defined for the buffer. The **ASYNCHRONOUS** attribute also works with the asynchronous MPI routines because the compiler must expect that inside of the MPI routines such Fortran asynchronous read, write, or wait routines may be called. (*End of rationale.*)

- In the Examples 16.14 and 16.15 and also in the example in Section 11.7.3 on page 425, the **ASYNCHRONOUS** attribute may also help but the help is not guaranteed because there is not an IO counterpart to the MPI usage.

Rationale. In case of using **MPI_BOTTOM** or one-sided synchronizations (e.g., **MPI_WIN_FENCE**), the buffer is not specified, i.e., those calls can include only a Fortran **WAIT** statement (or another routine that finishes an asynchronous IO).

Additionally, with Fortran asynchronous IO, it is a clear and forbidden race-condition when storing new data into the buffer while an asynchronous IO is active. Exactly this storing of data into the buffer is done in Example 16.14 when there would have been an initialization `buf=val_init` prior to the call to `MPI_RECV`, or in Example 16.15, the statement `buf=val_new`. (*End of rationale.*)

Fortran TARGET attribute. Declaring a buffer with the Fortran `TARGET` attribute in a scoping unit (or `BLOCK`) tells the compiler that any statement of the scoping unit may be executed while some pointer to the buffer exist. Calling a library routine (e.g., an MPI routine) may imply that such a pointer is used to modify the buffer.

- The `TARGET` attribute solves problems of instruction reordering, code movement, and register optimization related to nonblocking and one-sided communication, or related to the usage of `MPI_BOTTOM` and derived datatype handles. Declaring `REAL, TARGET :: buf` solves the register optimization problem in Examples 16.11, 16.12, 16.14, and 16.15
- Unfortunately, the `TARGET` attribute has **not** any impact on problems caused by asynchronous accesses between the start and end of a nonblocking or one-sided communication, i.e., problems through temporary memory modifications are not solved. Example 16.13 can **not** be solved with the `TARGET` attribute.

The compiler may be prevented from moving a reference to a buffer across a call to an MPI subroutine by surrounding the call by calls to an external subroutine with the buffer as an actual argument. The MPI library provides `MPI_F_SYNC_REG` for this purpose, see Section 16.2.6 on page 560.

- Examples 16.11 and 16.12 can be solved by calling `MPI_F_SYNC_REG(buf)` once directly after `MPI_WAIT`.

First example

```
call MPI_IRECV(buf,..req)

call MPI_WAIT(req,..)
call MPI_F_SYNC_REG(buf)
b1 = buf
```

Second example

```
buf = val
call MPI_ISEND(buf,..req)
copy = buf
call MPI_WAIT(req,..)
call MPI_F_SYNC_REG(buf)
buf = val_overwrite
```

The call `MPI_F_SYNC_REG(buf)` prevents moving the last line before the `MPI_WAIT` call. Further calls to `MPI_F_SYNC_REG(buf)` are not needed, because it is still correct if the additional read access `copy=buf` is moved behind `MPI_WAIT` and before `buf=val_overwrite`.

- Examples 16.14 and 16.15 can be solved with two additional `call MPI_F_SYNC_REG(buf)`, one directly before `MPI_RECV/MPI_SEND`, and one directly after this communication operation.

First example

```
call MPI_F_SYNC_REG(buf)
call MPI_RECV(MPI_BOTTOM,...)
call MPI_F_SYNC_REG(buf)
```

Second example

```
call MPI_F_SYNC_REG(buf)
call MPI_SEND(MPI_BOTTOM,...)
call MPI_F_SYNC_REG(buf)
```

The first call to `MPI_F_SYNC_REG(buf)` is needed to finish all load and store references to `buf` prior to `MPI_RECV/MPI_SEND`, and the second call is needed to assure that the subsequent access to `buf` are not moved before `MPI_RECV/SEND`.

- In the example in Section 16.2.6 on page 560, two asynchronous accesses must be protected: In Process 1, the access to `bbbb` must be protected similar to Example 16.11, i.e., a call to `MPI_F_SYNC_REG(bbbb)` is needed after the second `MPI_WIN_FENCE` to guarantee that further accesses to `bbbb` are not moved ahead of the call to `MPI_WIN_FENCE`. In Process 2, both calls to `MPI_WIN_FENCE` together act as a communication call with `MPI_BOTTOM` as the buffer, i.e., before the first fence and after the second fence, a call to `MPI_F_SYNC_REG(buff)` is needed to guarantee that accesses to `buff` are not moved after or ahead of the calls to `MPI_WIN_FENCE`. Using `MPI_GET` instead of `MPI_PUT`, the same calls to `MPI_F_SYNC_REG` are necessary.

Source of Process 1

```
bbbb = 777

call MPI_WIN_FENCE
call MPI_PUT(bbbb
into buff of process 2)

call MPI_WIN_FENCE
call MPI_F_SYNC_REG(bbbb)
```

Source of Process 2

```
buff = 999
call MPI_F_SYNC_REG(buff)
call MPI_WIN_FENCE

call MPI_WIN_FENCE
call MPI_F_SYNC_REG(buff)
ccc = buff
```

- The temporary memory modification problem, i.e., Example 16.13, can **not** be solved with this method.

A user defined DD instead of `MPI_F_SYNC_REG`. Instead of `MPI_F_SYNC_REG`, one can use also a user defined external subroutine, which is separately compiled:

```
subroutine DD(buf)
  integer buf
end
```

Note that if the intent is declared in the external subroutine, it must be `OUT` or `INOUT`. The subroutine itself may have an empty body, but the compiler does not know this and has to assume that the buffer may be altered. For example, the above call of `MPI_RECV` might be replaced by


```

1      call DD(buf)
2      call MPI_RECV(MPI_BOTTOM,...)
3      call DD(buf)

```

An alternative is to put the buffer or variable into a module or a common block and access it through a `USE` or `COMMON` statement in each scope where it is referenced, defined or appears as an actual argument in a call to an MPI routine. The compiler will then have to assume that the MPI procedure (`MPI_RECV` in the above example) may alter the buffer or variable, provided that the compiler cannot analyze that the MPI procedure does not reference the module or common block.

- This method solves problems of instruction reordering, code movement, and register optimization related to nonblocking and one-sided communication, or related to the usage of `MPI_BOTTOM` and derived datatype handles.
- Unfortunately, this method has **not** any impact on problems caused by asynchronous accesses between the start and end of a nonblocking or one-sided communication, i.e., problems through temporary memory modifications are not solved.

The `VOLATILE` attribute, gives the buffer or variable the properties needed, but it may inhibit optimization of any code containing the buffer or variable.

In C, subroutines which modify variables that are not in the argument list will not cause register optimization problems. This is because taking pointers to storage objects by using the `&` operator and later referencing the objects by way of the pointer is an integral part of the language. A C compiler understands the implications, so that the problem should not occur, in general. However, some compilers do offer optional aggressive optimization levels which may not be safe.

16.2.3 Fortran Support Through the `mpif.h` Include File

The use of the `mpif.h` include file is strongly discouraged.

Because Fortran 90 is (for all practical purposes) a superset of Fortran 77, Fortran 90 (and future) programs can use the original Fortran interface. The Fortran bindings are compatible with Fortran 77 implicit-style interfaces in most cases. The include file `mpif.h` must:

- Define all named MPI constants.
- Declare MPI functions that return a value.
- Define all handles as `INTEGER`. This is reflected in the first of the two Fortran interfaces in each MPI function definition.
- Be valid and equivalent for both fixed- and free- source form.

For each MPI routine, an implementation can choose to use an implicit or explicit interface.

Advice to implementors. To make `mpif.h` compatible with both fixed- and free-source forms, to allow automatic inclusion by preprocessors, and to allow extended fixed-form line length, it is recommended that requirement two be met by constructing `mpif.h` without any continuation lines. This should be possible because `mpif.h` contains only declarations, and because common block declarations can be split among several lines. To support Fortran 77 as well as Fortran 90, it may be necessary to eliminate all comments from `mpif.h`. (*End of advice to implementors.*)

16.2.4 Fortran Support Through the `mpi` Module

An MPI implementation must provide a module named `mpi` that can be used in a Fortran program. This module must:

- Define all named MPI constants
- Declare MPI functions that return a value.
- Provide explicit interfaces for all MPI routines, i.e., this module guarantees compile-time argument checking, and allows positional and keyword-based argument lists.
- Define all handles as `INTEGER`. This is reflected in the first of the two Fortran interfaces in each MPI function definition.
- Define also all the named handle types and `MPI_Status` that are used in the `mpi_f08` module. They are needed only when the application needs to convert an old-style `INTEGER` handle into a new-style handle with a named type.

An MPI implementation may provide other features in the `mpi` module that enhance the usability of MPI while maintaining adherence to the standard. For example, it may provide `INTENT` information in these interface blocks.

Advice to implementors. The appropriate `INTENT` may be different from what is given in the MPI generic interface. Implementations must choose `INTENT` so that the function adheres to the MPI standard. (*End of advice to implementors.*)

Rationale. The intent given by the MPI generic interface is not precisely defined and does not in all cases correspond to the correct Fortran `INTENT`. For instance, receiving into a buffer specified by a datatype with absolute addresses may require associating `MPI_BOTTOM` with a dummy `OUT` argument. Moreover, “constants” such as `MPI_BOTTOM` and `MPI_STATUS_IGNORE` are not constants as defined by Fortran, but “special addresses” used in a nonstandard way. Finally, the MPI-1 generic intent is changed in several places by MPI-2. For instance, `MPI_IN_PLACE` changes the sense of an `OUT` argument to be `INOUT`. (*End of rationale.*)

Advice to implementors. In the `mpi` module with some compilers, a choice argument can be implemented with the following explicit interface:

```
!DEC$ ATTRIBUTES NO_ARG_CHECK :: BUF
!$PRAGMA IGNORE_TKR BUF
REAL, DIMENSION(*) :: BUF
```

In this case, the compile-time constant `MPI_SUBARRAYS` equals `MPI_SUBARRAYS_UNSUPPORTED`. It is explicitly allowed that the choice arguments are implemented in the same way as with the `mpi_f08` module. In the case where the compiler does not provide such functionality, a set of overloaded functions may be used. See the paper of M. Hennecke [26]. (*End of advice to implementors.*)

16.2.5 Fortran Support Through the `mpi_f08` Module

An MPI implementation must provide a module named `mpi_f08` that can be used in a Fortran program. With this module, new Fortran definitions are added for each MPI routine, except for routines that are deprecated. This module must:

- Define all named MPI constants.
- Declare MPI functions that return a value.
- Provide explicit interfaces for all MPI routines, i.e., this module guarantees compile-time argument checking.
- Define all handles with uniquely named handle types (instead of `INTEGER` handles in the `mpi` module). This is reflected in the second of the two Fortran interfaces in each MPI function definition.
- Set the `INTEGER` compile-time constant `MPI_SUBARRAYS` to `MPI_SUBARRAYS_SUPPORTED` and declare choice buffers with the Fortran 2008 feature assumed-type and assumed-rank `TYPE(*)`, `DIMENSION(..)` if the underlying Fortran compiler supports it. With this, non-contiguous sub-arrays are valid also in nonblocking routines.
- Set the `MPI_SUBARRAYS` compile-time constant to `MPI_SUBARRAYS_UNSUPPORTED` and declare choice buffers with a compiler-dependent mechanism that overrides type checking if the underlying Fortran compiler does not support the Fortran 2008 assumed-type and assumed-rank notation. In this case, the use of non-contiguous sub-arrays in nonblocking calls may be restricted as with the `mpi` module.

Advice to implementors. In this case, the choice argument may be implemented with an explicit interface with compiler directives, for example:

```
!DEC$ ATTRIBUTES NO_ARG_CHECK :: BUF
!$PRAGMA IGNORE_TKR BUF
REAL, DIMENSION(*) :: BUF
```

(End of advice to implementors.)

- Declare each argument with an `INTENT=IN`, `OUT`, or `INOUT` as appropriate.
- Declare all `status` and `array_of_statuses` output arguments as optional through function overloading, instead of using `MPI_STATUS_IGNORE`.
- Declare all `array_of_errcodes` output arguments as optional through function overloading, instead of using `MPI_ERRCODES_IGNORE`.
- Declare all `ierror` output arguments as `OPTIONAL`, except for user-defined callback functions (e.g., `fctypeCOMM_COPY_ATTR_FUNCTION`) and their predefined callbacks (e.g., `MPI_NULL_COPY_FN`).

Rationale. For user-defined callback functions (e.g., `COMM_COPY_ATTR_FUNCTION`) and their predefined callbacks (e.g., `MPI_NULL_COPY_FN`), the `ierror` argument is not optional, i.e., these user-defined functions need not to check whether the MPI library calls these routines with or without an actual `ierror` output argument. (End of rationale.)

The MPI Fortran bindings in the `mpi_f08` module are designed based on the Fortran 2008 standard [34] together with the Technical Report (TR) on Further Interoperability with C [35] of the ISO/IEC JTC1/SC22/WG5 (Fortran) working group.

Rationale. The TR on further interoperability with C was defined by WG5 to support the MPI-3.0 standardization. “It is the intention of ISO/IEC JTC1/SC22/WG5 that the semantics and syntax specified by this technical report be included in the next revision of the Fortran International Standard without change unless experience in the implementation and use of this feature identifies errors that need to be corrected, or changes are needed to achieve proper integration, in which case every reasonable effort will be made to minimize the impact of such changes on existing implementations.”³

This TR contains language features that are needed for the MPI bindings in the `mpi_f08` module: assumed type and assumed rank. Here, it is important that any possible actual argument can be used for such dummy arguments, e.g., scalars, arrays, assumed-shape arrays, assumed-size arrays, allocatable arrays, and with any element type, e.g., `REAL`, `REAL*4`, `CHARACTER*5`, `CHARACTER*(*)`, derived types.

Furthermore, the implementors of the MPI Fortran bindings can freely choose whether all bindings are defined as `BIND(C)`. This is important to implement the Fortran `mpi_f08` interface with only **one** set of **portable** wrapper routines written in C. For this implementation goal, the following additional features are used: `BIND(C)` together

³[35] page iv, sentence 7.

with OPTIONAL, and with standard Fortran types like INTEGER and CHARACTER (i.e., not only with INTEGER(C_INT)).

The MPI Forum wants to acknowledge this important effort by the Fortran WG5 committee. (*End of rationale.*)

16.2.6 Additional Support for Fortran Register-Memory-Synchronization

As described in Section 16.2.2 on page 549, a dummy call is needed to tell the compiler that registers are to be flushed for a given buffer. It is a generic Fortran routine and has only a Fortran binding.

MPI_F_SYNC_REG(buf)

INOUT buf initial address of buffer (choice)

MPI_F_SYNC_REG(buf)

<type> buf(*)

MPI_F_sync_reg(buf)

TYPE(*), DIMENSION(...) :: buf

This routine has no operation associated with. It must be compiled in the MPI library in the way that a Fortran compiler cannot detect in the module that the routine has an empty body. It is used only to tell the compiler that a cached register value of a variable or buffer should be flushed, i.e., stored back to the memory (when necessary) or invalidated.

Rationale. This function is not available in other languages because it would not be useful. This routine has not an ierror return argument because there isn't any operation that can detect an error. (*End of rationale.*)

Advice to implementors. It is recommended to bind this routine to a C routine to minimize the risk that the fortran compiler can learn that this routine is empty, i.e., that the compiler can learn that a call to this routine can be removed as part of the automated optimization. (*End of advice to implementors.*)

16.2.7 Additional Support for Fortran Numeric Intrinsic Types

The routines in this section are part of Extended Fortran Support described in Section 16.2.4.

MPI provides a small number of named datatypes that correspond to named intrinsic types supported by C and Fortran. These include MPI_INTEGER, MPI_REAL, MPI_INT, MPI_DOUBLE, etc., as well as the optional types MPI_REAL4, MPI_REAL8, etc. There is a one-to-one correspondence between language declarations and MPI types.

Fortran (starting with Fortran 90) provides so-called KIND-parameterized types. These types are declared using an intrinsic type (one of INTEGER, REAL, COMPLEX, LOGICAL and

```
MPI_TYPE_MATCH_SIZE(TYPECLASS, SIZE, DATATYPE, IERROR)
```

```
INTEGER TYPECLASS, SIZE, DATATYPE, IERROR
```

```
MPI_Type_match_size(typeclass, size, datatype, ierror)
```

```
INTEGER, INTENT(IN) :: typeclass, size
```

```
TYPE(MPI_Datatype), INTENT(OUT) :: datatype
```

```
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

```
{static MPI::Datatype MPI::Datatype::Match_size(int typeclass,  
int size) (binding deprecated, see Section 15.2) }
```

typeclass is one of MPI_TYPECLASS_REAL, MPI_TYPECLASS_INTEGER and MPI_TYPECLASS_COMPLEX, corresponding to the desired **typeclass**. The function returns an MPI datatype matching a local variable of type (**typeclass**, **size**).

This function returns a reference (handle) to one of the predefined named datatypes, not a duplicate. This type cannot be freed. MPI_TYPE_MATCH_SIZE can be used to obtain a size-specific type that matches a Fortran numeric intrinsic type by first calling MPI_SIZEOF in order to compute the variable size, and then calling MPI_TYPE_MATCH_SIZE to find a suitable datatype. In C and C++, one can use the C function sizeof(), instead of MPI_SIZEOF. In addition, for variables of default kind the variable's size can be computed by a call to MPI_TYPE_GET_EXTENT, if the typeclass is known. It is erroneous to specify a size not supported by the compiler.

Rationale. This is a convenience function. Without it, it can be tedious to find the correct named type. See note to implementors below. (*End of rationale.*)

Advice to implementors. This function could be implemented as a series of tests.

```
int MPI_Type_match_size(int typeclass, int size, MPI_Datatype *rtype)
{
    switch(typeclass) {
        case MPI_TYPECLASS_REAL: switch(size) {
            case 4: *rtype = MPI_REAL4; return MPI_SUCCESS;
            case 8: *rtype = MPI_REAL8; return MPI_SUCCESS;
            default: error(...);
        }
        case MPI_TYPECLASS_INTEGER: switch(size) {
            case 4: *rtype = MPI_INTEGER4; return MPI_SUCCESS;
            case 8: *rtype = MPI_INTEGER8; return MPI_SUCCESS;
            default: error(...);
        }
        ... etc. ...
    }
}
```

(*End of advice to implementors.*)

16.3.4 Transfer of Handles

Handles are passed between Fortran and C or C++ by using an explicit C wrapper to convert Fortran handles to C handles. There is no direct access to C or C++ handles in Fortran. Handles are passed between C and C++ using overloaded C++ operators called from C++ code. There is no direct access to C++ objects from C.

The type definition `MPI_Fint` is provided in C/C++ for an integer of the size that matches a Fortran `INTEGER`; usually, `MPI_Fint` will be equivalent to `int`. With the Fortran `mpi` module or the `mpif.h` include file, a Fortran handle is a Fortran `INTEGER` value that can be used in the following conversion functions. With the Fortran `mpi_f08` module, a Fortran handle is a derived type that contains the Fortran `INTEGER` field `MPI_VAL`, which contains the the `INTEGER` value that can be used in the following conversion functions.

The following functions are provided in C to convert from a Fortran communicator handle (which is an integer) to a C communicator handle, and vice versa. See also Section 2.6.5 on page 22.

```
MPI_Comm MPI_Comm_f2c(MPI_Fint comm)
```

If `comm` is a valid Fortran handle to a communicator, then `MPI_Comm_f2c` returns a valid C handle to that same communicator; if `comm = MPI_COMM_NULL` (Fortran value), then `MPI_Comm_f2c` returns a null C handle; if `comm` is an invalid Fortran handle, then `MPI_Comm_f2c` returns an invalid C handle.

```
MPI_Fint MPI_Comm_c2f(MPI_Comm comm)
```

The function `MPI_Comm_c2f` translates a C communicator handle into a Fortran handle to the same communicator; it maps a null handle into a null handle and an invalid handle into an invalid handle.

Similar functions are provided for the other types of opaque objects.

```
MPI_Datatype MPI_Type_f2c(MPI_Fint datatype)
```

```
MPI_Fint MPI_Type_c2f(MPI_Datatype datatype)
```

```
MPI_Group MPI_Group_f2c(MPI_Fint group)
```

```
MPI_Fint MPI_Group_c2f(MPI_Group group)
```

```
MPI_Request MPI_Request_f2c(MPI_Fint request)
```

```
MPI_Fint MPI_Request_c2f(MPI_Request request)
```

```
MPI_File MPI_File_f2c(MPI_Fint file)
```

```
MPI_Fint MPI_File_c2f(MPI_File file)
```

```
MPI_Win MPI_Win_f2c(MPI_Fint win)
```

```
MPI_Fint MPI_Win_c2f(MPI_Win win)
```

```
MPI_Op MPI_Op_f2c(MPI_Fint op)
```

```
MPI_Fint MPI_Op_c2f(MPI_Op op)
```

```
MPI_Info MPI_Info_f2c(MPI_Fint info)
```

```
MPI_Fint MPI_Info_c2f(MPI_Info info)
```

```

1 // MPI::COMM_WORLD
2 MPI::Intracomm cpp_comm(MPI::COMM_WORLD.Dup());
3 c_lib_call(cpp_comm);
4 }

```

Rationale. Providing conversion from C to C++ via constructors and from C++ to C via casting allows the compiler to make automatic conversions. Calling C from C++ becomes trivial, as does the provision of a C or Fortran interface to a C++ library. (*End of rationale.*)

Advice to users. Note that the casting and promotion operators return new handles by value. Using these new handles as INOUT parameters will affect the internal MPI object, but will *not* affect the original handle from which it was cast. (*End of advice to users.*)

It is important to note that all C++ objects with corresponding C handles can be used interchangeably by an application. For example, an application can cache an attribute on MPI_COMM_WORLD and later retrieve it from MPI::COMM_WORLD.

16.3.5 Status

The following two procedures are provided in C to convert from a Fortran (**with the mpi module or mpif.h**) status (which is an array of integers) to a C status (which is a structure), and vice versa. The conversion occurs on all the information in status, including that which is hidden. That is, no status information is lost in the conversion.

```
int MPI_Status_f2c(MPI_Fint *f_status, MPI_Status *c_status)
```

If `f_status` is a valid Fortran status, but not the Fortran value of MPI_STATUS_IGNORE or MPI_STATUSES_IGNORE, then `MPI_Status_f2c` returns in `c_status` a valid C status with the same content. If `f_status` is the Fortran value of MPI_STATUS_IGNORE or MPI_STATUSES_IGNORE, or if `f_status` is not a valid Fortran status, then the call is erroneous.

The C status has the same source, tag and error code values as the Fortran status, and returns the same answers when queried for count, elements, and cancellation. The conversion function may be called with a Fortran status argument that has an undefined error field, in which case the value of the error field in the C status argument is undefined.

Two global variables of type MPI_Fint*, MPI_F_STATUS_IGNORE and MPI_F_STATUSES_IGNORE are declared in mpi.h. They can be used to test, in C, whether `f_status` is the Fortran value of MPI_STATUS_IGNORE or MPI_STATUSES_IGNORE, respectively. These are global variables, not C constant expressions and cannot be used in places where C requires constant expressions. Their value is defined only between the calls to MPI_INIT and MPI_FINALIZE and should not be changed by user code.

To do the conversion in the other direction, we have the following:

```
int MPI_Status_c2f(MPI_Status *c_status, MPI_Fint *f_status)
```

This call converts a C status into a Fortran status, and has a behavior similar to MPI_Status_f2c. That is, the value of `c_status` must not be either MPI_STATUS_IGNORE or MPI_STATUSES_IGNORE.

Advice to users. There is not a separate conversion function for arrays of statuses, since one can simply loop through the array, converting each status. (*End of advice to users.*)

Rationale. The handling of `MPI_STATUS_IGNORE` is required in order to layer libraries with only a C wrapper: if the Fortran call has passed `MPI_STATUS_IGNORE`, then the C wrapper must handle this correctly. Note that this constant need not have the same value in Fortran and C. If `MPI_Status_f2c` were to handle `MPI_STATUS_IGNORE`, then the type of its result would have to be `MPI_Status**`, which was considered an inferior solution. (*End of rationale.*)

Using the `mpi_f08` Fortran module, a status is declared as `TYPE(MPI_Status)`. The C datatype `MPI_F_status` can be used to hand over a Fortran `TYPE(MPI_Status)` argument into a C routine.

```
int MPI_Status_f082c(MPI_F_status *f08_status, MPI_Status *c_status)
```

This C routine converts a Fortran `mpi_f08 f08_status` into a C `c_status`.

```
int MPI_Status_c2f08(MPI_Status *c_status, MPI_F_status *f08_status)
```

This C routine converts a C `c_status` into a Fortran `mpi_f08 f08_status`.

Conversion between the two Fortran versions of a status can be done with:

```
MPI_STATUS_F2F08(f_status, f08_status)
```

| | | |
|----|----------|---------------------------------|
| IN | f_status | status object declared as array |
|----|----------|---------------------------------|

| | | |
|-----|------------|--------------------------------------|
| OUT | f08_status | status object declared as named type |
|-----|------------|--------------------------------------|

```
int MPI_Status_f2f08(MPI_Fint *f_status, MPI_F_status *f08_status)
```

```
MPI_STATUS_F2F08(F_STATUS, F08_STATUS, IERROR)
```

```
INTEGER F_STATUS(MPI_STATUS_SIZE)
```

```
TYPE(MPI_Status) :: F08_STATUS
```

```
INTEGER IERROR
```

```
MPI_Status_f2f08(f_status, f08_status, ierror)
```

```
INTEGER f_status(MPI_STATUS_SIZE)
```

```
TYPE(MPI_Status) :: f08_status
```

```
INTEGER, OPTIONAL :: ierror
```

This routine converts a Fortran `mpi module status` into a Fortran `mpi_f08 f08_status`.

```
MPI_STATUS_F082F(f08_status, f_status)
```

| | | |
|----|------------|--------------------------------------|
| IN | f08_status | status object declared as named type |
|----|------------|--------------------------------------|

| | | |
|-----|----------|---------------------------------|
| OUT | f_status | status object declared as array |
|-----|----------|---------------------------------|

```
int MPI_Status_f082f(MPI_F_status *f08_status, MPI_Fint *f_status)
```

```
MPI_STATUS_F082F(F08_STATUS, F_STATUS, IERROR)
```

```

1      TYPE(MPI_Status) :: F08_STATUS
2      INTEGER F_STATUS(MPI_STATUS_SIZE)
3      INTEGER IERROR
4
5      MPI_Status_f082f(f08_status, f_status, ierror)
6      TYPE(MPI_Status) :: f08_status
7      INTEGER :: f_status(MPI_STATUS_SIZE)
8      INTEGER, OPTIONAL :: ierror

```

This routine converts a Fortran `mpi_f08` module `f08_status` into a Fortran `mpi` status.

16.3.6 MPI Opaque Objects

Unless said otherwise, opaque objects are “the same” in all languages: they carry the same information, and have the same meaning in both languages. The mechanism described in the previous section can be used to pass references to MPI objects from language to language. An object created in one language can be accessed, modified or freed in another language.

We examine below in more detail, issues that arise for each type of MPI object.

Datatypes

Datatypes encode the same information in all languages. E.g., a datatype accessor like `MPI_TYPE_GET_EXTENT` will return the same information in all languages. If a datatype defined in one language is used for a communication call in another language, then the message sent will be identical to the message that would be sent from the first language: the same communication buffer is accessed, and the same representation conversion is performed, if needed. All predefined datatypes can be used in datatype constructors in any language. If a datatype is committed, it can be used for communication in any language.

The function `MPI_GET_ADDRESS` returns the same value in all languages. Note that we do not require that the constant `MPI_BOTTOM` have the same value in all languages (see [16.3.9](#), page 583).

Example 16.19

```

34      ! FORTRAN CODE
35      REAL R(5)
36      INTEGER TYPE, IERR, AOBLLEN(1), AOTYPE(1)
37      INTEGER (KIND=MPI_ADDRESS_KIND) AODISP(1)
38
39      ! create an absolute datatype for array R
40      AOBLLEN(1) = 5
41      CALL MPI_GET_ADDRESS( R, AODISP(1), IERR)
42      AOTYPE(1) = MPI_REAL
43      CALL MPI_TYPE_CREATE_STRUCT(1, AOBLLEN,AODISP,AOTYPE, TYPE, IERR)
44      CALL C_ROUTINE(TYPE)

```

Assorted Constants

| | |
|--|--|
| C type: <code>const int</code> (or unnamed <code>enum</code>) | C++ type: |
| Fortran type: <code>INTEGER</code> | <code>const int</code> (or unnamed <code>enum</code>) |
| <code>MPI_PROC_NULL</code> | <code>MPI::PROC_NULL</code> |
| <code>MPI_ANY_SOURCE</code> | <code>MPI::ANY_SOURCE</code> |
| <code>MPI_ANY_TAG</code> | <code>MPI::ANY_TAG</code> |
| <code>MPI_UNDEFINED</code> | <code>MPI::UNDEFINED</code> |
| <code>MPI_BSEND_OVERHEAD</code> | <code>MPI::BSEND_OVERHEAD</code> |
| <code>MPI_KEYVAL_INVALID</code> | <code>MPI::KEYVAL_INVALID</code> |
| <code>MPI_LOCK_EXCLUSIVE</code> | <code>MPI::LOCK_EXCLUSIVE</code> |
| <code>MPI_LOCK_SHARED</code> | <code>MPI::LOCK_SHARED</code> |
| <code>MPI_ROOT</code> | <code>MPI::ROOT</code> |

Status size and reserved index values (Fortran only)

| | |
|------------------------------------|---------------------|
| Fortran type: <code>INTEGER</code> | |
| <code>MPI_STATUS_SIZE</code> | Not defined for C++ |
| <code>MPI_SOURCE</code> | Not defined for C++ |
| <code>MPI_TAG</code> | Not defined for C++ |
| <code>MPI_ERROR</code> | Not defined for C++ |

Variable Address Size (Fortran only)

| | |
|------------------------------------|---------------------|
| Fortran type: <code>INTEGER</code> | |
| <code>MPI_ADDRESS_KIND</code> | Not defined for C++ |
| <code>MPI_INTEGER_KIND</code> | Not defined for C++ |
| <code>MPI_OFFSET_KIND</code> | Not defined for C++ |

Error-handling specifiers

| | |
|--------------------------------------|---|
| C type: <code>MPI_Errhandler</code> | C++ type: <code>MPI::Errhandler</code> |
| Fortran type: <code>INTEGER</code> | |
| or <code>TYPE(MPI_Errhandler)</code> | |
| <code>MPI_ERRORS_ARE_FATAL</code> | <code>MPI::ERRORS_ARE_FATAL</code> |
| <code>MPI_ERRORS_RETURN</code> | <code>MPI::ERRORS_RETURN</code> |
| | <code>MPI::ERRORS_THROW_EXCEPTIONS</code> |

Maximum Sizes for Strings

| | |
|--|--|
| C type: <code>const int</code> (or unnamed <code>enum</code>) | C++ type: |
| Fortran type: <code>INTEGER</code> | <code>const int</code> (or unnamed <code>enum</code>) |
| <code>MPI_MAX_PROCESSOR_NAME</code> | <code>MPI::MAX_PROCESSOR_NAME</code> |
| <code>MPI_MAX_ERROR_STRING</code> | <code>MPI::MAX_ERROR_STRING</code> |
| <code>MPI_MAX_DATAREP_STRING</code> | <code>MPI::MAX_DATAREP_STRING</code> |
| <code>MPI_MAX_INFO_KEY</code> | <code>MPI::MAX_INFO_KEY</code> |
| <code>MPI_MAX_INFO_VAL</code> | <code>MPI::MAX_INFO_VAL</code> |
| <code>MPI_MAX_OBJECT_NAME</code> | <code>MPI::MAX_OBJECT_NAME</code> |
| <code>MPI_MAX_PORT_NAME</code> | <code>MPI::MAX_PORT_NAME</code> |

| Named Predefined Datatypes | | C/C++ types |
|----------------------------|-------------------------|---|
| C type: MPI_Datatype | C++ type: MPI::Datatype | |
| Fortran type: INTEGER | | |
| or TYPE(MPI_Datatype) | | |
| MPI_CHAR | MPI::CHAR | char (treated as printable character) |
| MPI_SHORT | MPI::SHORT | signed short int |
| MPI_INT | MPI::INT | signed int |
| MPI_LONG | MPI::LONG | signed long |
| MPI_LONG_LONG_INT | MPI::LONG_LONG_INT | signed long long |
| MPI_LONG_LONG | MPI::LONG_LONG | long long (synonym) |
| MPI_SIGNED_CHAR | MPI::SIGNED_CHAR | signed char (treated as integral value) |
| MPI_UNSIGNED_CHAR | MPI::UNSIGNED_CHAR | unsigned char (treated as integral value) |
| MPI_UNSIGNED_SHORT | MPI::UNSIGNED_SHORT | unsigned short |
| MPI_UNSIGNED | MPI::UNSIGNED | unsigned int |
| MPI_UNSIGNED_LONG | MPI::UNSIGNED_LONG | unsigned long |
| MPI_UNSIGNED_LONG_LONG | MPI::UNSIGNED_LONG_LONG | unsigned long long |
| MPI_FLOAT | MPI::FLOAT | float |
| MPI_DOUBLE | MPI::DOUBLE | double |
| MPI_LONG_DOUBLE | MPI::LONG_DOUBLE | long double |
| MPI_WCHAR | MPI::WCHAR | wchar_t (defined in <stddef.h> (treated as printable character) |
| MPI_C_BOOL | (use C datatype handle) | _Bool |
| MPI_INT8_T | (use C datatype handle) | int8_t |
| MPI_INT16_T | (use C datatype handle) | int16_t |
| MPI_INT32_T | (use C datatype handle) | int32_t |
| MPI_INT64_T | (use C datatype handle) | int64_t |
| MPI_UINT8_T | (use C datatype handle) | uint8_t |
| MPI_UINT16_T | (use C datatype handle) | uint16_t |
| MPI_UINT32_T | (use C datatype handle) | uint32_t |
| MPI_UINT64_T | (use C datatype handle) | uint64_t |
| MPI_AINT | (use C datatype handle) | MPI_Aint |
| MPI_OFFSET | (use C datatype handle) | MPI_Offset |
| MPI_C_COMPLEX | (use C datatype handle) | float _Complex |
| MPI_C_FLOAT_COMPLEX | (use C datatype handle) | float _Complex |
| MPI_C_DOUBLE_COMPLEX | (use C datatype handle) | double _Complex |
| MPI_C_LONG_DOUBLE_COMPLEX | (use C datatype handle) | long double _Complex |
| MPI_BYTE | MPI::BYTE | (any C/C++ type) |
| MPI_PACKED | MPI::PACKED | (any C/C++ type) |

| Named Predefined Datatypes | | Fortran types |
|----------------------------|-------------------------|---------------------------------|
| C type: MPI_Datatype | C++ type: MPI::Datatype | |
| Fortran type: INTEGER | | |
| or TYPE(MPI_Datatype) | | |
| MPI_INTEGER | MPI::INTEGER | INTEGER |
| MPI_REAL | MPI::REAL | REAL |
| MPI_DOUBLE_PRECISION | MPI::DOUBLE_PRECISION | DOUBLE PRECISION |
| MPI_COMPLEX | MPI::F_COMPLEX | COMPLEX |
| MPI_LOGICAL | MPI::LOGICAL | LOGICAL |
| MPI_CHARACTER | MPI::CHARACTER | CHARACTER(1) |
| MPI_AINT | (use C datatype handle) | INTEGER (KIND=MPI_ADDRESS_KIND) |
| MPI_OFFSET | (use C datatype handle) | INTEGER (KIND=MPI_OFFSET_KIND) |
| MPI_BYTE | MPI::BYTE | (any Fortran type) |
| MPI_PACKED | MPI::PACKED | (any Fortran type) |

| C++-Only Named Predefined Datatypes | C++ types |
|-------------------------------------|----------------------|
| C++ type: MPI::Datatype | |
| MPI::BOOL | bool |
| MPI::COMPLEX | Complex<float> |
| MPI::DOUBLE_COMPLEX | Complex<double> |
| MPI::LONG_DOUBLE_COMPLEX | Complex<long double> |

| Optional datatypes (Fortran) | | Fortran types |
|------------------------------|-------------------------|----------------|
| C type: MPI_Datatype | C++ type: MPI::Datatype | |
| Fortran type: INTEGER | | |
| or TYPE(MPI_Datatype) | | |
| MPI_DOUBLE_COMPLEX | MPI::F_DOUBLE_COMPLEX | DOUBLE COMPLEX |
| MPI_INTEGER1 | MPI::INTEGER1 | INTEGER*1 |
| MPI_INTEGER2 | MPI::INTEGER2 | INTEGER*8 |
| MPI_INTEGER4 | MPI::INTEGER4 | INTEGER*4 |
| MPI_INTEGER8 | MPI::INTEGER8 | INTEGER*8 |
| MPI_INTEGER16 | | INTEGER*16 |
| MPI_REAL2 | MPI::REAL2 | REAL*2 |
| MPI_REAL4 | MPI::REAL4 | REAL*4 |
| MPI_REAL8 | MPI::REAL8 | REAL*8 |
| MPI_REAL16 | | REAL*16 |
| MPI_COMPLEX4 | | COMPLEX*4 |
| MPI_COMPLEX8 | | COMPLEX*8 |
| MPI_COMPLEX16 | | COMPLEX*16 |
| MPI_COMPLEX32 | | COMPLEX*32 |

Datatypes for reduction functions (C and C++)

| | |
|-----------------------|-------------------------|
| C type: MPI_Datatype | C++ type: MPI::Datatype |
| Fortran type: INTEGER | |
| or TYPE(MPI_Datatype) | |

| | |
|---------------------|----------------------|
| MPI_FLOAT_INT | MPI::FLOAT_INT |
| MPI_DOUBLE_INT | MPI::DOUBLE_INT |
| MPI_LONG_INT | MPI::LONG_INT |
| MPI_2INT | MPI::TWOINT |
| MPI_SHORT_INT | MPI::SHORT_INT |
| MPI_LONG_DOUBLE_INT | MPI::LONG_DOUBLE_INT |

Datatypes for reduction functions (Fortran)

| | |
|-----------------------|-------------------------|
| C type: MPI_Datatype | C++ type: MPI::Datatype |
| Fortran type: INTEGER | |
| or TYPE(MPI_Datatype) | |

| | |
|-----------------------|--------------------------|
| MPI_2REAL | MPI::TWOREAL |
| MPI_2DOUBLE_PRECISION | MPI::TWODOUBLE_PRECISION |
| MPI_2INTEGER | MPI::TWOINTEGER |

Special datatypes for constructing derived datatypes

| | |
|-----------------------|-------------------------|
| C type: MPI_Datatype | C++ type: MPI::Datatype |
| Fortran type: INTEGER | |
| or TYPE(MPI_Datatype) | |

| | |
|--------|---------|
| MPI_UB | MPI::UB |
| MPI_LB | MPI::LB |

Reserved communicators

| | |
|-----------------------|--------------------------|
| C type: MPI_Comm | C++ type: MPI::Intracomm |
| Fortran type: INTEGER | |
| or TYPE(MPI_Comm) | |

| | |
|----------------|-----------------|
| MPI_COMM_WORLD | MPI::COMM_WORLD |
| MPI_COMM_SELF | MPI::COMM_SELF |

Results of communicator and group comparisons

| | |
|-------------------------------------|---------------------|
| C type: const int (or unnamed enum) | C++ type: const int |
| Fortran type: INTEGER | (or unnamed enum) |
| MPI_IDENT | MPI::IDENT |
| MPI_CONGRUENT | MPI::CONGRUENT |
| MPI_SIMILAR | MPI::SIMILAR |
| MPI_UNEQUAL | MPI::UNEQUAL |

Environmental inquiry keys

| | |
|--|-----------------------------------|
| C type: <code>const int</code> (or unnamed <code>enum</code>) | C++ type: <code>const int</code> |
| Fortran type: <code>INTEGER</code> | (or unnamed <code>enum</code>) |
| <code>MPI_TAG_UB</code> | <code>MPI::TAG_UB</code> |
| <code>MPI_IO</code> | <code>MPI::IO</code> |
| <code>MPI_HOST</code> | <code>MPI::HOST</code> |
| <code>MPI_WTIME_IS_GLOBAL</code> | <code>MPI::WTIME_IS_GLOBAL</code> |

Collective Operations

| | |
|------------------------------------|--------------------------------------|
| C type: <code>MPI_Op</code> | C++ type: <code>const MPI::Op</code> |
| Fortran type: <code>INTEGER</code> | |
| or <code>TYPE(MPI_Op)</code> | |
| <code>MPI_MAX</code> | <code>MPI::MAX</code> |
| <code>MPI_MIN</code> | <code>MPI::MIN</code> |
| <code>MPI_SUM</code> | <code>MPI::SUM</code> |
| <code>MPI_PROD</code> | <code>MPI::PROD</code> |
| <code>MPI_MAXLOC</code> | <code>MPI::MAXLOC</code> |
| <code>MPI_MINLOC</code> | <code>MPI::MINLOC</code> |
| <code>MPI_BAND</code> | <code>MPI::BAND</code> |
| <code>MPI_BOR</code> | <code>MPI::BOR</code> |
| <code>MPI_BXOR</code> | <code>MPI::BXOR</code> |
| <code>MPI_LAND</code> | <code>MPI::LAND</code> |
| <code>MPI_LOR</code> | <code>MPI::LOR</code> |
| <code>MPI_LXOR</code> | <code>MPI::LXOR</code> |
| <code>MPI_REPLACE</code> | <code>MPI::REPLACE</code> |

Null Handles

| C/Fortran name | C++ name |
|---|-----------------------|
| C type / Fortran type | C++ type |
| MPI_GROUP_NULL | MPI::GROUP_NULL |
| MPI_Group / INTEGER or TYPE(MPI_Group) | const MPI::Group |
| MPI_COMM_NULL | MPI::COMM_NULL |
| MPI_Comm / INTEGER or TYPE(MPI_Comm) | ¹⁾ |
| MPI_DATATYPE_NULL | MPI::DATATYPE_NULL |
| MPI_Datatype / INTEGER or TYPE(MPI_Datatype) | const MPI::Datatype |
| MPI_REQUEST_NULL | MPI::REQUEST_NULL |
| MPI_Request / INTEGER or TYPE(MPI_Request) | const MPI::Request |
| MPI_OP_NULL | MPI::OP_NULL |
| MPI_Op / INTEGER or TYPE(MPI_Op) | const MPI::Op |
| MPI_ERRHANDLER_NULL | MPI::ERRHANDLER_NULL |
| MPI_Errhandler / INTEGER or TYPE(MPI_Errhandler) | const MPI::Errhandler |
| MPI_FILE_NULL | MPI::FILE_NULL |
| MPI_File / INTEGER or TYPE(MPI_File) | |
| MPI_INFO_NULL | MPI::INFO_NULL |
| MPI_Info / INTEGER or TYPE(MPI_Info) | const MPI::Info |
| MPI_WIN_NULL | MPI::WIN_NULL |
| MPI_Win / INTEGER or TYPE(MPI_Win) | |

¹⁾ C++ type: See Section 16.1.7 on page 536 regarding class hierarchy and the specific type of MPI::COMM_NULL

Empty group

| | |
|---|----------------------------|
| C type: MPI_Group | C++ type: const MPI::Group |
| Fortran type: INTEGER or TYPE(MPI_Group) | |
| MPI_GROUP_EMPTY | MPI::GROUP_EMPTY |

Topologies

| | |
|-------------------------------------|---------------------|
| C type: const int (or unnamed enum) | C++ type: const int |
| Fortran type: INTEGER | (or unnamed enum) |
| MPI_GRAPH | MPI::GRAPH |
| MPI_CART | MPI::CART |
| MPI_DIST_GRAPH | MPI::DIST_GRAPH |

| Predefined functions | |
|--|-----------------------------|
| C/Fortran name | C++ name |
| C type / Fortran type | C++ type |
| MPI_COMM_NULL_COPY_FN | MPI_COMM_NULL_COPY_FN |
| MPI_Comm_copy_attr_function | same as in C ¹) |
| / COMM_COPY_ATTR_FUNCTION | |
| MPI_COMM_DUP_FN | MPI_COMM_DUP_FN |
| MPI_Comm_copy_attr_function | same as in C ¹) |
| / COMM_COPY_ATTR_FUNCTION | |
| MPI_COMM_NULL_DELETE_FN | MPI_COMM_NULL_DELETE_FN |
| MPI_Comm_delete_attr_function | same as in C ¹) |
| / COMM_DELETE_ATTR_FUNCTION | |
| MPI_WIN_NULL_COPY_FN | MPI_WIN_NULL_COPY_FN |
| MPI_Win_copy_attr_function | same as in C ¹) |
| / WIN_COPY_ATTR_FUNCTION | |
| MPI_WIN_DUP_FN | MPI_WIN_DUP_FN |
| MPI_Win_copy_attr_function | same as in C ¹) |
| / WIN_COPY_ATTR_FUNCTION | |
| MPI_WIN_NULL_DELETE_FN | MPI_WIN_NULL_DELETE_FN |
| MPI_Win_delete_attr_function | same as in C ¹) |
| / WIN_DELETE_ATTR_FUNCTION | |
| MPI_TYPE_NULL_COPY_FN | MPI_TYPE_NULL_COPY_FN |
| MPI_Type_copy_attr_function | same as in C ¹) |
| / TYPE_COPY_ATTR_FUNCTION | |
| MPI_TYPE_DUP_FN | MPI_TYPE_DUP_FN |
| MPI_Type_copy_attr_function | same as in C ¹) |
| / TYPE_COPY_ATTR_FUNCTION | |
| MPI_TYPE_NULL_DELETE_FN | MPI_TYPE_NULL_DELETE_FN |
| MPI_Type_delete_attr_function | same as in C ¹) |
| / TYPE_DELETE_ATTR_FUNCTION | |
| ¹ See the advice to implementors on MPI_COMM_NULL_COPY_FN, ... in Section 6.7.2 on page 266 | |

| Deprecated predefined functions | |
|---------------------------------------|----------------------|
| C/Fortran name | C++ name |
| C type / Fortran type | C++ type |
| MPI_NULL_COPY_FN | MPI::NULL_COPY_FN |
| MPI_Copy_function / COPY_FUNCTION | MPI::Copy_function |
| MPI_DUP_FN | MPI::DUP_FN |
| MPI_Copy_function / COPY_FUNCTION | MPI::Copy_function |
| MPI_NULL_DELETE_FN | MPI::NULL_DELETE_FN |
| MPI_Delete_function / DELETE_FUNCTION | MPI::Delete_function |

File Operation Constants, Part 2

| C type: <code>const int</code> (or unnamed <code>enum</code>) | C++ type: |
|--|--|
| Fortran type: <code>INTEGER</code> | <code>const int</code> (or unnamed <code>enum</code>) |
| <code>MPI_DISTRIBUTE_BLOCK</code> | <code>MPI::DISTRIBUTE_BLOCK</code> |
| <code>MPI_DISTRIBUTE_CYCLIC</code> | <code>MPI::DISTRIBUTE_CYCLIC</code> |
| <code>MPI_DISTRIBUTE_DFLT_DARG</code> | <code>MPI::DISTRIBUTE_DFLT_DARG</code> |
| <code>MPI_DISTRIBUTE_NONE</code> | <code>MPI::DISTRIBUTE_NONE</code> |
| <code>MPI_ORDER_C</code> | <code>MPI::ORDER_C</code> |
| <code>MPI_ORDER_FORTRAN</code> | <code>MPI::ORDER_FORTRAN</code> |
| <code>MPI_SEEK_CUR</code> | <code>MPI::SEEK_CUR</code> |
| <code>MPI_SEEK_END</code> | <code>MPI::SEEK_END</code> |
| <code>MPI_SEEK_SET</code> | <code>MPI::SEEK_SET</code> |

F90 Datatype Matching Constants

| C type: <code>const int</code> (or unnamed <code>enum</code>) | C++ type: |
|--|--|
| Fortran type: <code>INTEGER</code> | <code>const int</code> (or unnamed <code>enum</code>) |
| <code>MPI_TYPECLASS_COMPLEX</code> | <code>MPI::TYPECLASS_COMPLEX</code> |
| <code>MPI_TYPECLASS_INTEGER</code> | <code>MPI::TYPECLASS_INTEGER</code> |
| <code>MPI_TYPECLASS_REAL</code> | <code>MPI::TYPECLASS_REAL</code> |

Constants Specifying Empty or Ignored Input

| C/Fortran name | C++ name |
|--|---|
| C type / Fortran type with <code>mpi</code> module / Fortran type with <code>mpi_f08</code> module | C++ type |
| <code>MPI_ARGVS_NULL</code> <code>char***</code> / 2-dim. array of <code>CHARACTER*(*)</code> / 2-dim. array of <code>CHARACTER*(*)</code> | <code>MPI::ARGVS_NULL</code> <code>const char ***</code> |
| <code>MPI_ARGV_NULL</code> <code>char**</code> / array of <code>CHARACTER*(*)</code> / array of <code>CHARACTER*(*)</code> | <code>MPI::ARGV_NULL</code> <code>const char **</code> |
| <code>MPI_ERRCODES_IGNORE</code> <code>int*</code> / <code>INTEGER</code> array / not defined | Not defined for C++ |
| <code>MPI_STATUSES_IGNORE</code> <code>MPI_Status*</code> / <code>INTEGER</code> , <code>DIMENSION(MPI_STATUS_SIZE,*)</code> / not defined | Not defined for C++ |
| <code>MPI_STATUS_IGNORE</code> <code>MPI_Status*</code> / <code>INTEGER</code> , <code>DIMENSION(MPI_STATUS_SIZE)</code> / not defined | Not defined for C++ |
| <code>MPI_UNWEIGHTED</code> <code>int*</code> / <code>INTEGER</code> array / not defined | Not defined for C++ |

```

1  MPI::File
2  MPI::Group
3  MPI::Info
4  MPI::Op
5  MPI::Request
6  MPI::Prequest
7  MPI::Grequest
8  MPI::Win

```

The following are defined Fortran type definitions, included in the `mpi_f08` and `mpi` module.

```

12  ! Fortran opaque types in the mpi_f08 and mpi module
13  TYPE(MPI_Status)
14
15  ! Fortran handles in the mpi_f08 and mpi module
16  TYPE(MPI_Comm)
17  TYPE(MPI_Datatype)
18  TYPE(MPI_Errhandler)
19  TYPE(MPI_File)
20  TYPE(MPI_Group)
21  TYPE(MPI_Info)
22  TYPE(MPI_Op)
23  TYPE(MPI_Request)
24  TYPE(MPI_Win)

```

A.1.3 Prototype Definitions

The following are defined C typedefs for user-defined functions, also included in the file `mpi.h`.

```

31  /* prototypes for user-defined functions */
32  typedef void MPI_User_function(void *invec, void *inoutvec, int *len,
33                                MPI_Datatype *datatype);
34
35  typedef int MPI_Comm_copy_attr_function(MPI_Comm oldcomm,
36                                          int comm_keyval, void *extra_state, void *attribute_val_in,
37                                          void *attribute_val_out, int*flag);
38  typedef int MPI_Comm_delete_attr_function(MPI_Comm comm,
39                                          int comm_keyval, void *attribute_val, void *extra_state);
40
41  typedef int MPI_Win_copy_attr_function(MPI_Win oldwin, int win_keyval,
42                                          void *extra_state, void *attribute_val_in,
43                                          void *attribute_val_out, int *flag);
44  typedef int MPI_Win_delete_attr_function(MPI_Win win, int win_keyval,
45                                          void *attribute_val, void *extra_state);
46
47  typedef int MPI_Type_copy_attr_function(MPI_Datatype oldtype,
48                                          int type_keyval, void *extra_state,

```

```

        void *attribute_val_in, void *attribute_val_out, int *flag);
typedef int MPI_Type_delete_attr_function(MPI_Datatype type,
        int type_keyval, void *attribute_val, void *extra_state);

typedef void MPI_Comm_errhandler_function(MPI_Comm *, int *, ...);
typedef void MPI_Win_errhandler_function(MPI_Win *, int *, ...);
typedef void MPI_File_errhandler_function(MPI_File *, int *, ...);

typedef int MPI_Grequest_query_function(void *extra_state,
        MPI_Status *status);
typedef int MPI_Grequest_free_function(void *extra_state);
typedef int MPI_Grequest_cancel_function(void *extra_state, int complete);

typedef int MPI_Datarep_extent_function(MPI_Datatype datatype,
        MPI_Aint *file_extent, void *extra_state);
typedef int MPI_Datarep_conversion_function(void *userbuf,
        MPI_Datatype datatype, int count, void *filebuf,
        MPI_Offset position, void *extra_state);

```

For Fortran, here are examples of how each of the user-defined subroutines should be declared.

The user-function argument to MPI_OP_CREATE should be declared like this:

```

SUBROUTINE USER_FUNCTION(INVEC, INOUTVEC, LEN, TYPE)
    <type> INVEC(LEN), INOUTVEC(LEN)
    INTEGER LEN, TYPE

```

The copy and delete function arguments to MPI_COMM_CREATE_KEYVAL should be declared like these:

```

SUBROUTINE COMM_COPY_ATTR_FUNCTION(OLDCOMM, COMM_KEYVAL, EXTRA_STATE,
        ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERROR)
    INTEGER OLDCOMM, COMM_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,
        ATTRIBUTE_VAL_OUT
    LOGICAL FLAG

SUBROUTINE COMM_DELETE_ATTR_FUNCTION(COMM, COMM_KEYVAL, ATTRIBUTE_VAL,
        EXTRA_STATE, IERROR)
    INTEGER COMM, COMM_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE

```

The copy and delete function arguments to MPI_WIN_CREATE_KEYVAL should be declared like these:

```

SUBROUTINE WIN_COPY_ATTR_FUNCTION(OLDWIN, WIN_KEYVAL, EXTRA_STATE,
        ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERROR)
    INTEGER OLDWIN, WIN_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,

```

```

1      ATTRIBUTE_VAL_OUT
2      LOGICAL FLAG

```

```

3
4      SUBROUTINE WIN_DELETE_ATTR_FUNCTION(WIN, WIN_KEYVAL, ATTRIBUTE_VAL,
5          EXTRA_STATE, IERROR)
6          INTEGER WIN, WIN_KEYVAL, IERROR
7          INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE
8

```

9 The copy and delete function arguments to MPI_TYPE_CREATE_KEYVAL should be
10 declared like these:

```

11
12      SUBROUTINE TYPE_COPY_ATTR_FUNCTION(OLDTYPE, TYPE_KEYVAL, EXTRA_STATE,
13          ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERROR)
14          INTEGER OLDTYPE, TYPE_KEYVAL, IERROR
15          INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE,
16          ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT
17          LOGICAL FLAG
18
19      SUBROUTINE TYPE_DELETE_ATTR_FUNCTION(DATATYPE, TYPE_KEYVAL, ATTRIBUTE_VAL,
20          EXTRA_STATE, IERROR)
21          INTEGER DATATYPE, TYPE_KEYVAL, IERROR
22          INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE
23

```

24 The handler-function argument to MPI_COMM_CREATE_ERRHANDLER should be de-
25 clared like this:

```

26
27      SUBROUTINE COMM_ERRHANDLER_FUNCTION(COMM, ERROR_CODE)
28          INTEGER COMM, ERROR_CODE
29

```

30 The handler-function argument to MPI_WIN_CREATE_ERRHANDLER should be de-
31 clared like this:

```

32      SUBROUTINE WIN_ERRHANDLER_FUNCTION(WIN, ERROR_CODE)
33          INTEGER WIN, ERROR_CODE
34

```

35 The handler-function argument to MPI_FILE_CREATE_ERRHANDLER should be de-
36 clared like this:

```

37
38      SUBROUTINE FILE_ERRHANDLER_FUNCTION(FILE, ERROR_CODE)
39          INTEGER FILE, ERROR_CODE
40

```

41 The query, free, and cancel function arguments to MPI_GREQUEST_START should be
42 declared like these:

```

43
44      SUBROUTINE GREQUEST_QUERY_FUNCTION(EXTRA_STATE, STATUS, IERROR)
45          INTEGER STATUS(MPI_STATUS_SIZE), IERROR
46          INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
47
48      SUBROUTINE GREQUEST_FREE_FUNCTION(EXTRA_STATE, IERROR)

```

```

INTEGER IERROR
INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE

SUBROUTINE GREQUEST_CANCEL_FUNCTION(EXTRA_STATE, COMPLETE, IERROR)
  INTEGER IERROR
  INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
  LOGICAL COMPLETE

```

The extend and conversion function arguments to MPI_REGISTER_DATAREP should be declared like these:

```

SUBROUTINE DATAREP_EXTENT_FUNCTION(DATATYPE, EXTENT, EXTRA_STATE, IERROR)
  INTEGER DATATYPE, IERROR
  INTEGER(KIND=MPI_ADDRESS_KIND) EXTENT, EXTRA_STATE

SUBROUTINE DATAREP_CONVERSION_FUNCTION(USERBUF, DATATYPE, COUNT, FILEBUF,
  POSITION, EXTRA_STATE, IERROR)
  <TYPE> USERBUF(*), FILEBUF(*)
  INTEGER COUNT, DATATYPE, IERROR
  INTEGER(KIND=MPI_OFFSET_KIND) POSITION
  INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE

```

The following are defined C++ typedefs, also included in the file `mpi.h`.

```

namespace MPI {
  typedef void User_function(const void* invec, void *inoutvec,
    int len, const Datatype& datatype);

  typedef int Comm::Copy_attr_function(const Comm& oldcomm,
    int comm_keyval, void* extra_state, void* attribute_val_in,
    void* attribute_val_out, bool& flag);
  typedef int Comm::Delete_attr_function(Comm& comm, int
    comm_keyval, void* attribute_val, void* extra_state);

  typedef int Win::Copy_attr_function(const Win& oldwin,
    int win_keyval, void* extra_state, void* attribute_val_in,
    void* attribute_val_out, bool& flag);
  typedef int Win::Delete_attr_function(Win& win, int
    win_keyval, void* attribute_val, void* extra_state);

  typedef int Datatype::Copy_attr_function(const Datatype& oldtype,
    int type_keyval, void* extra_state,
    const void* attribute_val_in, void* attribute_val_out,
    bool& flag);
  typedef int Datatype::Delete_attr_function(Datatype& type,
    int type_keyval, void* attribute_val, void* extra_state);

  typedef void Comm::Errhandler_function(Comm &, int *, ...);
  typedef void Win::Errhandler_function(Win &, int *, ...);

```

```

1  int MPI_Pack(void* inbuf, int incount, MPI_Datatype datatype, void *outbuf,
2      int outsize, int *position, MPI_Comm comm)
3
4  int MPI_Pack_external(char *datarep, void *inbuf, int incount,
5      MPI_Datatype datatype, void *outbuf, MPI_Aint outsize,
6      MPI_Aint *position)
7
8  int MPI_Pack_external_size(char *datarep, int incount,
9      MPI_Datatype datatype, MPI_Aint *size)
10
11 int MPI_Pack_size(int incount, MPI_Datatype datatype, MPI_Comm comm,
12     int *size)
13
14 int MPI_Type_commit(MPI_Datatype *datatype)
15
16 int MPI_Type_contiguous(int count, MPI_Datatype oldtype,
17     MPI_Datatype *newtype)
18
19 int MPI_Type_create_darray(int size, int rank, int ndims,
20     int array_of_gsizes[], int array_of_distribs[], int
21     array_of_dargs[], int array_of_psize[], int order,
22     MPI_Datatype oldtype, MPI_Datatype *newtype)
23
24 int MPI_Type_create_hindexed(int count, int array_of_blocklengths[],
25     MPI_Aint array_of_displacements[], MPI_Datatype oldtype,
26     MPI_Datatype *newtype)
27
28 int MPI_Type_create_hvector(int count, int blocklength, MPI_Aint stride,
29     MPI_Datatype oldtype, MPI_Datatype *newtype)
30
31 int MPI_Type_create_indexed_block(int count, int blocklength,
32     int array_of_displacements[], MPI_Datatype oldtype,
33     MPI_Datatype *newtype)
34
35 int MPI_Type_create_resized(MPI_Datatype oldtype, MPI_Aint lb, MPI_Aint
36     extent, MPI_Datatype *newtype)
37
38 int MPI_Type_create_struct(int count, int array_of_blocklengths[],
39     MPI_Aint array_of_displacements[],
40     MPI_Datatype array_of_types[], MPI_Datatype *newtype)
41
42 int MPI_Type_create_subarray(int ndims, int array_of_sizes[],
43     int array_of_subsizes[], int array_of_starts[], int order,
44     MPI_Datatype oldtype, MPI_Datatype *newtype)
45
46 int MPI_Type_dup(MPI_Datatype oldtype, MPI_Datatype *newtype)
47
48 int MPI_Type_free(MPI_Datatype *datatype)
49
50 int MPI_Type_get_contents(MPI_Datatype datatype, int max_integers,
51     int max_addresses, int max_datatypes, int array_of_integers[],
52     MPI_Aint array_of_addresses[],
53     MPI_Datatype array_of_datatypes[])
54
55 int MPI_Type_get_envelope(MPI_Datatype datatype, int *num_integers,

```

```

int MPI_Ireduce(void* sendbuf, void* recvbuf, int count,
                MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm comm,
                MPI_Request* request)
int MPI_Ireduce_scatter(void* sendbuf, void* recvbuf, int recvcounts[],
                        MPI_Datatype datatype, MPI_Op op, MPI_Comm comm,
                        MPI_Request* request)
int MPI_Ireduce_scatter_block(void* sendbuf, void* recvbuf, int recvcount,
                              MPI_Datatype datatype, MPI_Op op, MPI_Comm comm,
                              MPI_Request* request)
int MPI_Iscan(void* sendbuf, void* recvbuf, int count,
              MPI_Datatype datatype, MPI_Op op, MPI_Comm comm,
              MPI_Request* request)
int MPI_Iscatter(void* sendbuf, int sendcount, MPI_Datatype sendtype,
                void* recvbuf, int recvcount, MPI_Datatype recvtype, int root,
                MPI_Comm comm, MPI_Request* request)
int MPI_Iscatterv(void* sendbuf, int sendcounts[], int displs[],
                  MPI_Datatype sendtype, void* recvbuf, int recvcount,
                  MPI_Datatype recvtype, int root, MPI_Comm comm,
                  MPI_Request* request)
int MPI_Op_commutative(MPI_Op op, int *commute)
int MPI_Op_create(MPI_User_function* user_fn, int commute, MPI_Op* op)
int MPI_Reduce(void* sendbuf, void* recvbuf, int count,
               MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm comm)
int MPI_Reduce_local(void* inbuf, void* inoutbuf, int count,
                     MPI_Datatype datatype, MPI_Op op)
int MPI_Reduce_scatter(void* sendbuf, void* recvbuf, int recvcounts[],
                       MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
int MPI_Reduce_scatter_block(void* sendbuf, void* recvbuf, int recvcount,
                              MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
int MPI_Scan(void* sendbuf, void* recvbuf, int count,
             MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
int MPI_Scatter(void* sendbuf, int sendcount, MPI_Datatype sendtype,
                void* recvbuf, int recvcount, MPI_Datatype recvtype, int root,
                MPI_Comm comm)
int MPI_Scatterv(void* sendbuf, int sendcounts[], int displs[],
                 MPI_Datatype sendtype, void* recvbuf, int recvcount,
                 MPI_Datatype recvtype, int root, MPI_Comm comm)
int MPI_op_free(MPI_Op *op)

```



```

int MPI_Group_incl(MPI_Group group, int n, int *ranks, MPI_Group *newgroup) 1
int MPI_Group_intersection(MPI_Group group1, MPI_Group group2, 2
    MPI_Group *newgroup) 3
int MPI_Group_range_excl(MPI_Group group, int n, int ranges[][3], 4
    MPI_Group *newgroup) 5
int MPI_Group_range_incl(MPI_Group group, int n, int ranges[][3], 6
    MPI_Group *newgroup) 7
int MPI_Group_rank(MPI_Group group, int *rank) 8
int MPI_Group_size(MPI_Group group, int *size) 9
int MPI_Group_translate_ranks (MPI_Group group1, int n, int *ranks1, 10
    MPI_Group group2, int *ranks2) 11
int MPI_Group_union(MPI_Group group1, MPI_Group group2, 12
    MPI_Group *newgroup) 13
int MPI_Intercomm_create(MPI_Comm local_comm, int local_leader, 14
    MPI_Comm peer_comm, int remote_leader, int tag, 15
    MPI_Comm *newintercomm) 16
int MPI_Intercomm_merge(MPI_Comm intercomm, int high, 17
    MPI_Comm *newintracomm) 18
int MPI_TYPE_DUP_FN(MPI_Datatype oldtype, int type_keyval, 19
    void *extra_state, void *attribute_val_in, 20
    void *attribute_val_out, int *flag) 21
int MPI_TYPE_NULL_COPY_FN(MPI_Datatype oldtype, int type_keyval, 22
    void *extra_state, void *attribute_val_in, 23
    void *attribute_val_out, int *flag) 24
int MPI_TYPE_NULL_DELETE_FN(MPI_Datatype datatype, int type_keyval, void 25
    *attribute_val, void *extra_state) 26
int MPI_Type_create_keyval(MPI_Type_copy_attr_function *type_copy_attr_fn, 27
    MPI_Type_delete_attr_function *type_delete_attr_fn, 28
    int *type_keyval, void *extra_state) 29
int MPI_Type_delete_attr(MPI_Datatype datatype, int type_keyval) 30
int MPI_Type_free_keyval(int *type_keyval) 31
int MPI_Type_get_attr(MPI_Datatype datatype, int type_keyval, void 32
    *attribute_val, int *flag) 33
int MPI_Type_get_name(MPI_Datatype datatype, char *type_name, int 34
    *resultlen) 35
int MPI_Type_set_attr(MPI_Datatype datatype, int type_keyval, 36
    void *attribute_val) 37

```

```

1  int MPI_Type_set_name(MPI_Datatype datatype, char *type_name)
2
3  int MPI_WIN_DUP_FN(MPI_Win oldwin, int win_keyval, void *extra_state,
4                      void *attribute_val_in, void *attribute_val_out, int *flag)
5
6  int MPI_WIN_NULL_COPY_FN(MPI_Win oldwin, int win_keyval, void *extra_state,
7                          void *attribute_val_in, void *attribute_val_out, int *flag)
8
9  int MPI_WIN_NULL_DELETE_FN(MPI_Win win, int win_keyval, void
10                           *attribute_val, void *extra_state)
11
12 int MPI_Win_create_keyval(MPI_Win_copy_attr_function *win_copy_attr_fn,
13                          MPI_Win_delete_attr_function *win_delete_attr_fn,
14                          int *win_keyval, void *extra_state)
15
16 int MPI_Win_delete_attr(MPI_Win win, int win_keyval)
17
18 int MPI_Win_free_keyval(int *win_keyval)
19
20 int MPI_Win_get_attr(MPI_Win win, int win_keyval, void *attribute_val,
21                     int *flag)
22
23 int MPI_Win_get_name(MPI_Win win, char *win_name, int *resultlen)
24
25 int MPI_Win_set_attr(MPI_Win win, int win_keyval, void *attribute_val)
26
27 int MPI_Win_set_name(MPI_Win win, char *win_name)

```

A.2.5 Process Topologies C Bindings

```

28 int MPI_Cart_coords(MPI_Comm comm, int rank, int maxdims, int *coords)
29
30 int MPI_Cart_create(MPI_Comm comm_old, int ndims, int *dims, int *periods,
31                   int reorder, MPI_Comm *comm_cart)
32
33 int MPI_Cart_get(MPI_Comm comm, int maxdims, int *dims, int *periods,
34                 int *coords)
35
36 int MPI_Cart_map(MPI_Comm comm, int ndims, int *dims, int *periods,
37                 int *newrank)
38
39 int MPI_Cart_rank(MPI_Comm comm, int *coords, int *rank)
40
41 int MPI_Cart_shift(MPI_Comm comm, int direction, int disp,
42                   int *rank_source, int *rank_dest)
43
44 int MPI_Cart_sub(MPI_Comm comm, int *remain_dims, MPI_Comm *newcomm)
45
46 int MPI_Cartdim_get(MPI_Comm comm, int *ndims)
47
48 int MPI_Dims_create(int nnodes, int ndims, int *dims)
49
50 int MPI_Dist_graph_create(MPI_Comm comm_old, int n, int sources[],
51                          int degrees[], int destinations[], int weights[],
52                          MPI_Info info, int reorder, MPI_Comm *comm_dist_graph)

```

```

int MPI_Dist_graph_create_adjacent(MPI_Comm comm_old, int indegree,
    int sources[], int sourceweights[], int outdegree,
    int destinations[], int destweights[], MPI_Info info,
    int reorder, MPI_Comm *comm_dist_graph)

int MPI_Dist_graph_neighbors(MPI_Comm comm, int maxindegree, int sources[],
    int sourceweights[], int maxoutdegree, int destinations[],
    int destweights[])

int MPI_Dist_graph_neighbors_count(MPI_Comm comm, int *indegree,
    int *outdegree, int *weighted)

int MPI_Graph_create(MPI_Comm comm_old, int nnodes, int *index, int *edges,
    int reorder, MPI_Comm *comm_graph)

int MPI_Graph_get(MPI_Comm comm, int maxindex, int maxedges, int *index,
    int *edges)

int MPI_Graph_map(MPI_Comm comm, int nnodes, int *index, int *edges,
    int *newrank)

int MPI_Graph_neighbors(MPI_Comm comm, int rank, int maxneighbors,
    int *neighbors)

int MPI_Graph_neighbors_count(MPI_Comm comm, int rank, int *nneighbors)

int MPI_Graphdims_get(MPI_Comm comm, int *nnodes, int *nedges)

int MPI_Topo_test(MPI_Comm comm, int *status)

```

A.2.6 MPI Environmental Management C Bindings

```

double MPI_Wtick(void)

double MPI_Wtime(void)

int MPI_Abort(MPI_Comm comm, int errorcode)

int MPI_Add_error_class(int *errorclass)

int MPI_Add_error_code(int errorclass, int *errorcode)

int MPI_Add_error_string(int errorcode, char *string)

int MPI_Alloc_mem(MPI_Aint size, MPI_Info info, void *baseptr)

int MPI_Comm_call_errhandler(MPI_Comm comm, int errorcode)

int MPI_Comm_create_errhandler(MPI_Comm_errhandler_function
    *comm_errhandler_fn, MPI_Errhandler *errhandler)

int MPI_Comm_get_errhandler(MPI_Comm comm, MPI_Errhandler *errhandler)

int MPI_Comm_set_errhandler(MPI_Comm comm, MPI_Errhandler errhandler)

int MPI_Errhandler_free(MPI_Errhandler *errhandler)

```

```

1  int MPI_Error_class(int errorcode, int *errorclass)
2
3  int MPI_Error_string(int errorcode, char *string, int *resultlen)
4
5  int MPI_File_call_errhandler(MPI_File fh, int errorcode)
6
7  int MPI_File_create_errhandler(MPI_File_errhandler_function
8      *file_errhandler_fn, MPI_Errhandler *errhandler)
9
10 int MPI_File_get_errhandler(MPI_File file, MPI_Errhandler *errhandler)
11
12 int MPI_File_set_errhandler(MPI_File file, MPI_Errhandler errhandler)
13
14 int MPI_Finalize(void)
15
16 int MPI_Finalized(int *flag)
17
18 int MPI_Free_mem(void *base)
19
20 int MPI_Get_processor_name(char *name, int *resultlen)
21
22 int MPI_Get_version(int *version, int *subversion)
23
24 int MPI_Init(int *argc, char ***argv)
25
26 int MPI_Initialized(int *flag)
27
28 int MPI_Win_call_errhandler(MPI_Win win, int errorcode)
29
30 int MPI_Win_create_errhandler(MPI_Win_errhandler_function
31     *win_errhandler_fn, MPI_Errhandler *errhandler)
32
33 int MPI_Win_get_errhandler(MPI_Win win, MPI_Errhandler *errhandler)
34
35 int MPI_Win_set_errhandler(MPI_Win win, MPI_Errhandler errhandler)

```

A.2.7 The Info Object C Bindings

```

31 int MPI_Info_create(MPI_Info *info)
32
33 int MPI_Info_delete(MPI_Info info, char *key)
34
35 int MPI_Info_dup(MPI_Info info, MPI_Info *newinfo)
36
37 int MPI_Info_free(MPI_Info *info)
38
39 int MPI_Info_get(MPI_Info info, char *key, int valuelen, char *value,
40     int *flag)
41
42 int MPI_Info_get_nkeys(MPI_Info info, int *nkeys)
43
44 int MPI_Info_get_nthkey(MPI_Info info, int n, char *key)
45
46 int MPI_Info_get_valuelen(MPI_Info info, char *key, int *valuelen,
47     int *flag)
48
49 int MPI_Info_set(MPI_Info info, char *key, char *value)

```

```
void *extra_state)
```

A.2.12 Language Bindings C Bindings

```
int MPI_Status_f082f(MPI_F_status *f08_status, MPI_Fint *f_status)
int MPI_Status_f2f08(MPI_Fint *f_status, MPI_F_status *f08_status)
int MPI_Type_create_f90_complex(int p, int r, MPI_Datatype *newtype)
int MPI_Type_create_f90_integer(int r, MPI_Datatype *newtype)
int MPI_Type_create_f90_real(int p, int r, MPI_Datatype *newtype)
int MPI_Type_match_size(int typeclass, int size, MPI_Datatype *datatype)
MPI_Fint MPI_Comm_c2f(MPI_Comm comm)
MPI_Comm MPI_Comm_f2c(MPI_Fint comm)
MPI_Fint MPI_Errhandler_c2f(MPI_Errhandler errhandler)
MPI_Errhandler MPI_Errhandler_f2c(MPI_Fint errhandler)
MPI_Fint MPI_File_c2f(MPI_File file)
MPI_File MPI_File_f2c(MPI_Fint file)
MPI_Fint MPI_Group_c2f(MPI_Group group)
MPI_Group MPI_Group_f2c(MPI_Fint group)
MPI_Fint MPI_Info_c2f(MPI_Info info)
MPI_Info MPI_Info_f2c(MPI_Fint info)
MPI_Fint MPI_Op_c2f(MPI_Op op)
MPI_Op MPI_Op_f2c(MPI_Fint op)
MPI_Fint MPI_Request_c2f(MPI_Request request)
MPI_Request MPI_Request_f2c(MPI_Fint request)
int MPI_Status_c2f(MPI_Status *c_status, MPI_Fint *f_status)
int MPI_Status_c2f08(MPI_Status *c_status, MPI_F_status *f08_status)
int MPI_Status_f082c(MPI_F_status *f08_status, MPI_Status *c_status)
int MPI_Status_f2c(MPI_Fint *f_status, MPI_Status *c_status)
MPI_Fint MPI_Type_c2f(MPI_Datatype datatype)
MPI_Datatype MPI_Type_f2c(MPI_Fint datatype)
MPI_Fint MPI_Win_c2f(MPI_Win win)
MPI_Win MPI_Win_f2c(MPI_Fint win)
```

A.2.13 Profiling Interface C Bindings

```
int MPI_Pcontrol(const int level, ...)
```

A.2.14 Deprecated C Bindings

```
int MPI_Address(void* location, MPI_Aint *address)

int MPI_Attr_delete(MPI_Comm comm, int keyval)

int MPI_Attr_get(MPI_Comm comm, int keyval, void *attribute_val, int *flag)

int MPI_Attr_put(MPI_Comm comm, int keyval, void* attribute_val)

int MPI_DUP_FN(MPI_Comm oldcomm, int keyval, void *extra_state,
               void *attribute_val_in, void *attribute_val_out, int *flag)

int MPI_Errhandler_create(MPI_Handler_function *handler_fn,
                          MPI_Errhandler *errhandler)

int MPI_Errhandler_get(MPI_Comm comm, MPI_Errhandler *errhandler)

int MPI_Errhandler_set(MPI_Comm comm, MPI_Errhandler errhandler)

int MPI_Keyval_create(MPI_Copy_function *copy_fn, MPI_Delete_function
                     *delete_fn, int *keyval, void* extra_state)

int MPI_Keyval_free(int *keyval)

int MPI_NULL_COPY_FN(MPI_Comm oldcomm, int keyval, void *extra_state,
                    void *attribute_val_in, void *attribute_val_out, int *flag)

int MPI_NULL_DELETE_FN(MPI_Comm comm, int keyval, void *attribute_val,
                      void *extra_state)

int MPI_Type_extent(MPI_Datatype datatype, MPI_Aint *extent)

int MPI_Type_hindexed(int count, int *array_of_blocklengths,
                     MPI_Aint *array_of_displacements, MPI_Datatype oldtype,
                     MPI_Datatype *newtype)

int MPI_Type_hvector(int count, int blocklength, MPI_Aint stride,
                    MPI_Datatype oldtype, MPI_Datatype *newtype)

int MPI_Type_lb(MPI_Datatype datatype, MPI_Aint* displacement)

int MPI_Type_struct(int count, int *array_of_blocklengths,
                   MPI_Aint *array_of_displacements,
                   MPI_Datatype *array_of_types, MPI_Datatype *newtype)

int MPI_Type_ub(MPI_Datatype datatype, MPI_Aint* displacement)
```

A.3 Fortran Bindings with mpif.h or the mpi Module

A.3.1 Point-to-Point Communication Fortran Bindings

```

MPI_BSEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, IERROR)
    <type> BUF(*)
    INTEGER COUNT, DATATYPE, DEST, TAG, COMM, IERROR

MPI_BSEND_INIT(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
    <type> BUF(*)
    INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR

MPI_BUFFER_ATTACH(BUFFER, SIZE, IERROR)
    <type> BUFFER(*)
    INTEGER SIZE, IERROR

MPI_BUFFER_DETACH(BUFFER_ADDR, SIZE, IERROR)
    <type> BUFFER_ADDR(*)
    INTEGER SIZE, IERROR

MPI_CANCEL(REQUEST, IERROR)
    INTEGER REQUEST, IERROR

MPI_GET_COUNT(STATUS, DATATYPE, COUNT, IERROR)
    INTEGER STATUS(MPI_STATUS_SIZE), DATATYPE, COUNT, IERROR

MPI_IBSEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
    <type> BUF(*)
    INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR

MPI_IPROBE(SOURCE, TAG, COMM, FLAG, STATUS, IERROR)
    LOGICAL FLAG
    INTEGER SOURCE, TAG, COMM, STATUS(MPI_STATUS_SIZE), IERROR

MPI_IRECV(BUF, COUNT, DATATYPE, SOURCE, TAG, COMM, REQUEST, IERROR)
    <type> BUF(*)
    INTEGER COUNT, DATATYPE, SOURCE, TAG, COMM, REQUEST, IERROR

MPI_IRSEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
    <type> BUF(*)
    INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR

MPI_ISEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
    <type> BUF(*)
    INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR

MPI_ISSEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
    <type> BUF(*)
    INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR

MPI_PROBE(SOURCE, TAG, COMM, STATUS, IERROR)
    INTEGER SOURCE, TAG, COMM, STATUS(MPI_STATUS_SIZE), IERROR

MPI_RECV(BUF, COUNT, DATATYPE, SOURCE, TAG, COMM, STATUS, IERROR)

```

```

1      <type> BUF(*)
2      INTEGER COUNT, DATATYPE, SOURCE, TAG, COMM, STATUS(MPI_STATUS_SIZE),
3      IERROR
4
5      MPI_RECV_INIT(BUF, COUNT, DATATYPE, SOURCE, TAG, COMM, REQUEST, IERROR)
6      <type> BUF(*)
7      INTEGER COUNT, DATATYPE, SOURCE, TAG, COMM, REQUEST, IERROR
8
9      MPI_REQUEST_FREE(REQUEST, IERROR)
10     INTEGER REQUEST, IERROR
11
12     MPI_REQUEST_GET_STATUS( REQUEST, FLAG, STATUS, IERROR)
13     INTEGER REQUEST, STATUS(MPI_STATUS_SIZE), IERROR
14     LOGICAL FLAG
15
16     MPI_RSEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, IERROR)
17     <type> BUF(*)
18     INTEGER COUNT, DATATYPE, DEST, TAG, COMM, IERROR
19
20     MPI_RSEND_INIT(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
21     <type> BUF(*)
22     INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR
23
24     MPI_SEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, IERROR)
25     <type> BUF(*)
26     INTEGER COUNT, DATATYPE, DEST, TAG, COMM, IERROR
27
28     MPI_SENDRECV(SENDBUF, SENDCOUNT, SENDTYPE, DEST, SENDTAG, RECVBUF,
29     RECVCOUNT, RECVMODE, SOURCE, RECVTAG, COMM, STATUS, IERROR)
30     <type> SENDBUF(*), RECVBUF(*)
31     INTEGER SENDCOUNT, SENDTYPE, DEST, SENDTAG, RECVCOUNT, RECVMODE,
32     SOURCE, RECVTAG, COMM, STATUS(MPI_STATUS_SIZE), IERROR
33
34     MPI_SENDRECV_REPLACE(BUF, COUNT, DATATYPE, DEST, SENDTAG, SOURCE, RECVTAG,
35     COMM, STATUS, IERROR)
36     <type> BUF(*)
37     INTEGER COUNT, DATATYPE, DEST, SENDTAG, SOURCE, RECVTAG, COMM,
38     STATUS(MPI_STATUS_SIZE), IERROR
39
40     MPI_SEND_INIT(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
41     <type> BUF(*)
42     INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR
43
44     MPI_SSEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, IERROR)
45     <type> BUF(*)
46     INTEGER COUNT, DATATYPE, DEST, TAG, COMM, IERROR
47
48     MPI_SSEND_INIT(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
49     <type> BUF(*)
50     INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR
51
52     MPI_START(REQUEST, IERROR)
53     INTEGER REQUEST, IERROR

```



```

    INTEGER(KIND=MPI_ADDRESS_KIND) ARRAY_OF_DISPLACEMENTS(*)
1
MPI_TYPE_CREATE_SUBARRAY(NDIMS, ARRAY_OF_SIZES, ARRAY_OF_SUBSIZES,
2
    ARRAY_OF_STARTS, ORDER, OLDTYPE, NEWTYPE, IERROR)
3
    INTEGER NDIMS, ARRAY_OF_SIZES(*), ARRAY_OF_SUBSIZES(*),
4
    ARRAY_OF_STARTS(*), ORDER, OLDTYPE, NEWTYPE, IERROR
5
MPI_TYPE_DUP(OLDTYPE, NEWTYPE, IERROR)
6
    INTEGER OLDTYPE, NEWTYPE, IERROR
7
MPI_TYPE_FREE(DATATYPE, IERROR)
8
    INTEGER DATATYPE, IERROR
9
MPI_TYPE_GET_CONTENTS(DATATYPE, MAX_INTEGERS, MAX_ADDRESSES, MAX_DATATYPES,
10
    ARRAY_OF_INTEGERS, ARRAY_OF_ADDRESSES, ARRAY_OF_DATATYPES,
11
    IERROR)
12
    INTEGER DATATYPE, MAX_INTEGERS, MAX_ADDRESSES, MAX_DATATYPES,
13
    ARRAY_OF_INTEGERS(*), ARRAY_OF_DATATYPES(*), IERROR
14
    INTEGER(KIND=MPI_ADDRESS_KIND) ARRAY_OF_ADDRESSES(*)
15
MPI_TYPE_GET_ENVELOPE(DATATYPE, NUM_INTEGERS, NUM_ADDRESSES, NUM_DATATYPES,
16
    COMBINER, IERROR)
17
    INTEGER DATATYPE, NUM_INTEGERS, NUM_ADDRESSES, NUM_DATATYPES, COMBINER,
18
    IERROR
19
MPI_TYPE_GET_EXTENT(DATATYPE, LB, EXTENT, IERROR)
20
    INTEGER DATATYPE, IERROR
21
    INTEGER(KIND = MPI_ADDRESS_KIND) LB, EXTENT
22
MPI_TYPE_GET_TRUE_EXTENT(DATATYPE, TRUE_LB, TRUE_EXTENT, IERROR)
23
    INTEGER DATATYPE, IERROR
24
    INTEGER(KIND = MPI_ADDRESS_KIND) TRUE_LB, TRUE_EXTENT
25
MPI_TYPE_INDEXED(COUNT, ARRAY_OF_BLOCKLENGTHS, ARRAY_OF_DISPLACEMENTS,
26
    OLDTYPE, NEWTYPE, IERROR)
27
    INTEGER COUNT, ARRAY_OF_BLOCKLENGTHS(*), ARRAY_OF_DISPLACEMENTS(*),
28
    OLDTYPE, NEWTYPE, IERROR
29
MPI_TYPE_SIZE(DATATYPE, SIZE, IERROR)
30
    INTEGER DATATYPE, SIZE, IERROR
31
MPI_TYPE_VECTOR(COUNT, BLOCKLENGTH, STRIDE, OLDTYPE, NEWTYPE, IERROR)
32
    INTEGER COUNT, BLOCKLENGTH, STRIDE, OLDTYPE, NEWTYPE, IERROR
33
MPI_UNPACK(INBUF, INSIZE, POSITION, OUTBUF, OUTCOUNT, DATATYPE, COMM,
34
    IERROR)
35
    <type> INBUF(*), OUTBUF(*)
36
    INTEGER INSIZE, POSITION, OUTCOUNT, DATATYPE, COMM, IERROR
37
MPI_UNPACK_EXTERNAL(DATAREP, INBUF, INSIZE, POSITION, OUTBUF, OUTCOUNT,
38
    DATATYPE, IERROR)
39
    INTEGER OUTCOUNT, DATATYPE, IERROR
40
    INTEGER(KIND=MPI_ADDRESS_KIND) INSIZE, POSITION
41
42
43
44
45
46
47
48

```

```

1      <type> SENDBUF(*), RECVBUF(*)
2      INTEGER SENDCOUNT, SENDTYPE, RECVCOUNT, RECVTYPE, ROOT, COMM, REQUEST,
3      IERROR
4
5      MPI_IGATHERV(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNTS, DISPLS,
6                  RECVTYPE, ROOT, COMM, REQUEST, IERROR)
7      <type> SENDBUF(*), RECVBUF(*)
8      INTEGER SENDCOUNT, SENDTYPE, RECVCOUNTS(*), DISPLS(*), RECVTYPE, ROOT,
9      COMM, REQUEST, IERROR
10
11     MPI_IREDUCE(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, ROOT, COMM, REQUEST,
12                IERROR)
13     <type> SENDBUF(*), RECVBUF(*)
14     INTEGER COUNT, DATATYPE, OP, ROOT, COMM, REQUEST, IERROR
15
16     MPI_IREDUCE_SCATTER(SENDBUF, RECVBUF, RECVCOUNTS, DATATYPE, OP, COMM,
17                        REQUEST, IERROR)
18     <type> SENDBUF(*), RECVBUF(*)
19     INTEGER RECVCOUNTS(*), DATATYPE, OP, COMM, REQUEST, IERROR
20
21     MPI_IREDUCE_SCATTER_BLOCK(SENDBUF, RECVBUF, RECVCOUNT, DATATYPE, OP, COMM,
22                              REQUEST, IERROR)
23     <type> SENDBUF(*), RECVBUF(*)
24     INTEGER RECVCOUNT, DATATYPE, OP, COMM, REQUEST, IERROR
25
26     MPI_ISCAN(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, COMM, REQUEST, IERROR)
27     <type> SENDBUF(*), RECVBUF(*)
28     INTEGER COUNT, DATATYPE, OP, COMM, REQUEST, IERROR
29
30     MPI_ISCATTER(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNT, RECVTYPE,
31                 ROOT, COMM, REQUEST, IERROR)
32     <type> SENDBUF(*), RECVBUF(*)
33     INTEGER SENDCOUNT, SENDTYPE, RECVCOUNT, RECVTYPE, ROOT, COMM, REQUEST,
34     IERROR
35
36     MPI_ISCATTERV(SENDBUF, SENDCOUNTS, DISPLS, SENDTYPE, RECVBUF, RECVCOUNT,
37                  RECVTYPE, ROOT, COMM, REQUEST, IERROR)
38     <type> SENDBUF(*), RECVBUF(*)
39     INTEGER SENDCOUNTS(*), DISPLS(*), SENDTYPE, RECVCOUNT, RECVTYPE, ROOT,
40     COMM, REQUEST, IERROR
41
42     MPI_OP_COMMUTATIVE(OP, COMMUTE, IERROR)
43     LOGICAL COMMUTE
44     INTEGER OP, IERROR
45
46     MPI_OP_CREATE( USER_FN, COMMUTE, OP, IERROR)
47     EXTERNAL USER_FN
48     LOGICAL COMMUTE
49     INTEGER OP, IERROR
50
51     MPI_OP_FREE(OP, IERROR)
52     INTEGER OP, IERROR

```

```

MPI_REDUCE(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, ROOT, COMM, IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER COUNT, DATATYPE, OP, ROOT, COMM, IERROR
MPI_REDUCE_LOCAL(INBUF, INOUTBUF, COUNT, DATATYPE, OP, IERROR)
    <type> INBUF(*), INOUTBUF(*)
    INTEGER COUNT, DATATYPE, OP, IERROR
MPI_REDUCE_SCATTER(SENDBUF, RECVBUF, RECVCOUNTS, DATATYPE, OP, COMM,
    IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER RECVCOUNTS(*), DATATYPE, OP, COMM, IERROR
MPI_REDUCE_SCATTER_BLOCK(SENDBUF, RECVBUF, RECVCOUNT, DATATYPE, OP, COMM,
    IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER RECVCOUNT, DATATYPE, OP, COMM, IERROR
MPI_SCAN(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, COMM, IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER COUNT, DATATYPE, OP, COMM, IERROR
MPI_SCATTER(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNT, RECVTYPE,
    ROOT, COMM, IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER SENDCOUNT, SENDTYPE, RECVCOUNT, RECVTYPE, ROOT, COMM, IERROR
MPI_SCATTERV(SENDBUF, SENDCOUNTS, DISPLS, SENDTYPE, RECVBUF, RECVCOUNT,
    RECVTYPE, ROOT, COMM, IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER SENDCOUNTS(*), DISPLS(*), SENDTYPE, RECVCOUNT, RECVTYPE, ROOT,
    COMM, IERROR

```

A.3.4 Groups, Contexts, Communicators, and Caching Fortran Bindings

```

MPI_COMM_COMPARE(COMM1, COMM2, RESULT, IERROR)
    INTEGER COMM1, COMM2, RESULT, IERROR
MPI_COMM_CREATE(COMM, GROUP, NEWCOMM, IERROR)
    INTEGER COMM, GROUP, NEWCOMM, IERROR
MPI_COMM_CREATE_KEYVAL(COMM_COPY_ATTR_FN, COMM_DELETE_ATTR_FN, COMM_KEYVAL,
    EXTRA_STATE, IERROR)
    EXTERNAL COMM_COPY_ATTR_FN, COMM_DELETE_ATTR_FN
    INTEGER COMM_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
MPI_COMM_DELETE_ATTR(COMM, COMM_KEYVAL, IERROR)
    INTEGER COMM, COMM_KEYVAL, IERROR
MPI_COMM_DUP(COMM, NEWCOMM, IERROR)
    INTEGER COMM, NEWCOMM, IERROR

```

| | |
|--|----|
| MPI_COMM_SIZE(COMM, SIZE, IERROR) | 1 |
| INTEGER COMM, SIZE, IERROR | 2 |
| | 3 |
| MPI_COMM_SPLIT(COMM, COLOR, KEY, NEWCOMM, IERROR) | 4 |
| INTEGER COMM, COLOR, KEY, NEWCOMM, IERROR | 5 |
| | 6 |
| MPI_COMM_TEST_INTER(COMM, FLAG, IERROR) | 7 |
| INTEGER COMM, IERROR | 8 |
| LOGICAL FLAG | 9 |
| | 10 |
| MPI_GROUP_COMPARE(GROUP1, GROUP2, RESULT, IERROR) | 11 |
| INTEGER GROUP1, GROUP2, RESULT, IERROR | 12 |
| | 13 |
| MPI_GROUP_DIFFERENCE(GROUP1, GROUP2, NEWGROUP, IERROR) | 14 |
| INTEGER GROUP1, GROUP2, NEWGROUP, IERROR | 15 |
| | 16 |
| MPI_GROUP_EXCL(GROUP, N, RANKS, NEWGROUP, IERROR) | 17 |
| INTEGER GROUP, N, RANKS(*), NEWGROUP, IERROR | 18 |
| | 19 |
| MPI_GROUP_FREE(GROUP, IERROR) | 20 |
| INTEGER GROUP, IERROR | 21 |
| | 22 |
| MPI_GROUP_INCL(GROUP, N, RANKS, NEWGROUP, IERROR) | 23 |
| INTEGER GROUP, N, RANKS(*), NEWGROUP, IERROR | 24 |
| | 25 |
| MPI_GROUP_INTERSECTION(GROUP1, GROUP2, NEWGROUP, IERROR) | 26 |
| INTEGER GROUP1, GROUP2, NEWGROUP, IERROR | 27 |
| | 28 |
| MPI_GROUP_RANGE_EXCL(GROUP, N, RANGES, NEWGROUP, IERROR) | 29 |
| INTEGER GROUP, N, RANGES(3,*), NEWGROUP, IERROR | 30 |
| | 31 |
| MPI_GROUP_RANGE_INCL(GROUP, N, RANGES, NEWGROUP, IERROR) | 32 |
| INTEGER GROUP, N, RANGES(3,*), NEWGROUP, IERROR | 33 |
| | 34 |
| MPI_GROUP_RANK(GROUP, RANK, IERROR) | 35 |
| INTEGER GROUP, RANK, IERROR | 36 |
| | 37 |
| MPI_GROUP_SIZE(GROUP, SIZE, IERROR) | 38 |
| INTEGER GROUP, SIZE, IERROR | 39 |
| | 40 |
| MPI_GROUP_TRANSLATE_RANKS(GROUP1, N, RANKS1, GROUP2, RANKS2, IERROR) | 41 |
| INTEGER GROUP1, N, RANKS1(*), GROUP2, RANKS2(*), IERROR | 42 |
| | 43 |
| MPI_GROUP_UNION(GROUP1, GROUP2, NEWGROUP, IERROR) | 44 |
| INTEGER GROUP1, GROUP2, NEWGROUP, IERROR | 45 |
| | 46 |
| MPI_INTERCOMM_CREATE(LOCAL_COMM, LOCAL_LEADER, PEER_COMM, REMOTE_LEADER, | 47 |
| TAG, NEWINTERCOMM, IERROR) | 48 |
| INTEGER LOCAL_COMM, LOCAL_LEADER, PEER_COMM, REMOTE_LEADER, TAG, | |
| NEWINTERCOMM, IERROR | |
| | |
| MPI_INTERCOMM_MERGE(INTERCOMM, HIGH, NEWINTRACOMM, IERROR) | |
| INTEGER NEWINTERCOMM, INTRACOMM, IERROR | |
| LOGICAL HIGH | |

```

1  MPI_TYPE_CREATE_KEYVAL(TYPE_COPY_ATTR_FN, TYPE_DELETE_ATTR_FN, TYPE_KEYVAL,
2      EXTRA_STATE, IERROR)
3      EXTERNAL TYPE_COPY_ATTR_FN, TYPE_DELETE_ATTR_FN
4      INTEGER TYPE_KEYVAL, IERROR
5      INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
6
7  MPI_TYPE_DELETE_ATTR(DATATYPE, TYPE_KEYVAL, IERROR)
8      INTEGER DATATYPE, TYPE_KEYVAL, IERROR
9
10 MPI_TYPE_DUP_FN(OLDTYPE, TYPE_KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,
11     ATTRIBUTE_VAL_OUT, FLAG, IERROR)
12     INTEGER OLDTYPE, TYPE_KEYVAL, IERROR
13     INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,
14     ATTRIBUTE_VAL_OUT
15     LOGICAL FLAG
16
17 MPI_TYPE_FREE_KEYVAL(TYPE_KEYVAL, IERROR)
18     INTEGER TYPE_KEYVAL, IERROR
19
20 MPI_TYPE_GET_ATTR(DATATYPE, TYPE_KEYVAL, ATTRIBUTE_VAL, FLAG, IERROR)
21     INTEGER DATATYPE, TYPE_KEYVAL, IERROR
22     INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL
23     LOGICAL FLAG
24
25 MPI_TYPE_GET_NAME(DATATYPE, TYPE_NAME, RESULTLEN, IERROR)
26     INTEGER DATATYPE, RESULTLEN, IERROR
27     CHARACTER*(*) TYPE_NAME
28
29 MPI_TYPE_NULL_COPY_FN(OLDTYPE, TYPE_KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,
30     ATTRIBUTE_VAL_OUT, FLAG, IERROR)
31     INTEGER OLDTYPE, TYPE_KEYVAL, IERROR
32     INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,
33     ATTRIBUTE_VAL_OUT
34     LOGICAL FLAG
35
36 MPI_TYPE_NULL_DELETE_FN(DATATYPE, TYPE_KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE,
37     IERROR)
38     INTEGER DATATYPE, TYPE_KEYVAL, IERROR
39     INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE
40
41 MPI_TYPE_SET_ATTR(DATATYPE, TYPE_KEYVAL, ATTRIBUTE_VAL, IERROR)
42     INTEGER DATATYPE, TYPE_KEYVAL, IERROR
43     INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL
44
45 MPI_TYPE_SET_NAME(DATATYPE, TYPE_NAME, IERROR)
46     INTEGER DATATYPE, IERROR
47     CHARACTER*(*) TYPE_NAME
48
49 MPI_WIN_CREATE_KEYVAL(WIN_COPY_ATTR_FN, WIN_DELETE_ATTR_FN, WIN_KEYVAL,
50     EXTRA_STATE, IERROR)
51     EXTERNAL WIN_COPY_ATTR_FN, WIN_DELETE_ATTR_FN
52     INTEGER WIN_KEYVAL, IERROR

```

```

    INTEGER COMM, MAXINDEX, MAXEDGES, INDEX(*), EDGES(*), IERROR
MPI_GRAPH_MAP(COMM, NNODES, INDEX, EDGES, NEWRANK, IERROR)
    INTEGER COMM, NNODES, INDEX(*), EDGES(*), NEWRANK, IERROR
MPI_GRAPH_NEIGHBORS(COMM, RANK, MAXNEIGHBORS, NEIGHBORS, IERROR)
    INTEGER COMM, RANK, MAXNEIGHBORS, NEIGHBORS(*), IERROR
MPI_GRAPH_NEIGHBORS_COUNT(COMM, RANK, NNEIGHBORS, IERROR)
    INTEGER COMM, RANK, NNEIGHBORS, IERROR
MPI_TOPO_TEST(COMM, STATUS, IERROR)
    INTEGER COMM, STATUS, IERROR

```

A.3.6 MPI Environmental Management Fortran Bindings

```

DOUBLE PRECISION MPI_WTICK()
DOUBLE PRECISION MPI_WTIME()
MPI_ABORT(COMM, ERRORCODE, IERROR)
    INTEGER COMM, ERRORCODE, IERROR
MPI_ADD_ERROR_CLASS(ERRORCLASS, IERROR)
    INTEGER ERRORCLASS, IERROR
MPI_ADD_ERROR_CODE(ERRORCLASS, ERRORCODE, IERROR)
    INTEGER ERRORCLASS, ERRORCODE, IERROR
MPI_ADD_ERROR_STRING(ERRORCODE, STRING, IERROR)
    INTEGER ERRORCODE, IERROR
    CHARACTER*(*) STRING
MPI_ALLOC_MEM(SIZE, INFO, BASEPTR, IERROR)
    INTEGER INFO, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) SIZE, BASEPTR
MPI_COMM_CALL_ERRHANDLER(COMM, ERRORCODE, IERROR)
    INTEGER COMM, ERRORCODE, IERROR
MPI_COMM_CREATE_ERRHANDLER(COMM_ERRHANDLER_FN, ERRHANDLER, IERROR)
    EXTERNAL COMM_ERRHANDLER_FN
    INTEGER ERRHANDLER, IERROR
MPI_COMM_GET_ERRHANDLER(COMM, ERRHANDLER, IERROR)
    INTEGER COMM, ERRHANDLER, IERROR
MPI_COMM_SET_ERRHANDLER(COMM, ERRHANDLER, IERROR)
    INTEGER COMM, ERRHANDLER, IERROR
MPI_ERRHANDLER_FREE(ERRHANDLER, IERROR)
    INTEGER ERRHANDLER, IERROR
MPI_ERROR_CLASS(ERRORCODE, ERRORCLASS, IERROR)
    INTEGER ERRORCODE, ERRORCLASS, IERROR

```

```

1  MPI_ERROR_STRING(ERRORCODE, STRING, RESULTLEN, IERROR)
2      INTEGER ERRORCODE, RESULTLEN, IERROR
3      CHARACTER*(*) STRING
4
5  MPI_FILE_CALL_ERRHANDLER(FH, ERRORCODE, IERROR)
6      INTEGER FH, ERRORCODE, IERROR
7
8  MPI_FILE_CREATE_ERRHANDLER(FILE_ERRHANDLER_FN, ERRHANDLER, IERROR)
9      EXTERNAL FILE_ERRHANDLER_FN
10     INTEGER ERRHANDLER, IERROR
11
12 MPI_FILE_GET_ERRHANDLER(FILE, ERRHANDLER, IERROR)
13     INTEGER FILE, ERRHANDLER, IERROR
14
15 MPI_FILE_SET_ERRHANDLER(FILE, ERRHANDLER, IERROR)
16     INTEGER FILE, ERRHANDLER, IERROR
17
18 MPI_FINALIZE(IERROR)
19     INTEGER IERROR
20
21 MPI_FINALIZED(FLAG, IERROR)
22     LOGICAL FLAG
23     INTEGER IERROR
24
25 MPI_FREE_MEM(BASE, IERROR)
26     <type> BASE(*)
27     INTEGER IERROR
28
29 MPI_GET_PROCESSOR_NAME( NAME, RESULTLEN, IERROR)
30     CHARACTER*(*) NAME
31     INTEGER RESULTLEN, IERROR
32
33 MPI_GET_VERSION(VERSION, SUBVERSION, IERROR)
34     INTEGER VERSION, SUBVERSION, IERROR
35
36 MPI_INIT(IERROR)
37     INTEGER IERROR
38
39 MPI_INITIALIZED(FLAG, IERROR)
40     LOGICAL FLAG
41     INTEGER IERROR
42
43 MPI_WIN_CALL_ERRHANDLER(WIN, ERRORCODE, IERROR)
44     INTEGER WIN, ERRORCODE, IERROR
45
46 MPI_WIN_CREATE_ERRHANDLER(WIN_ERRHANDLER_FN, ERRHANDLER, IERROR)
47     EXTERNAL WIN_ERRHANDLER_FN
48     INTEGER ERRHANDLER, IERROR
49
50 MPI_WIN_GET_ERRHANDLER(WIN, ERRHANDLER, IERROR)
51     INTEGER WIN, ERRHANDLER, IERROR
52
53 MPI_WIN_SET_ERRHANDLER(WIN, ERRHANDLER, IERROR)
54     INTEGER WIN, ERRHANDLER, IERROR
55
56

```

```

1  MPI_FILE_WRITE_ORDERED_BEGIN(FH, BUF, COUNT, DATATYPE, IERROR)
2      <type> BUF(*)
3      INTEGER FH, COUNT, DATATYPE, IERROR
4
5  MPI_FILE_WRITE_ORDERED_END(FH, BUF, STATUS, IERROR)
6      <type> BUF(*)
7      INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR
8
9  MPI_FILE_WRITE_SHARED(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
10     <type> BUF(*)
11     INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
12
13 MPI_REGISTER_DATAREP(DATAREP, READ_CONVERSION_FN, WRITE_CONVERSION_FN,
14     DTYPE_FILE_EXTENT_FN, EXTRA_STATE, IERROR)
15     CHARACTER*(*) DATAREP
16     EXTERNAL READ_CONVERSION_FN, WRITE_CONVERSION_FN, DTYPE_FILE_EXTENT_FN
17     INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
18     INTEGER IERROR

```

A.3.12 Language Bindings Fortran Bindings

```

19
20
21 MPI_F_SYNC_REG(buf)
22     <type> buf(*)
23
24 MPI_SIZEOF(X, SIZE, IERROR)
25     <type> X
26     INTEGER SIZE, IERROR
27
28 MPI_STATUS_F082F(F08_STATUS, F_STATUS, IERROR)
29     TYPE(MPI_Status) :: F08_STATUS
30     INTEGER F_STATUS(MPI_STATUS_SIZE)
31     INTEGER IERROR
32
33 MPI_STATUS_F2F08(F_STATUS, F08_STATUS, IERROR)
34     INTEGER F_STATUS(MPI_STATUS_SIZE)
35     TYPE(MPI_Status) :: F08_STATUS
36     INTEGER IERROR
37
38 MPI_TYPE_CREATE_F90_COMPLEX(P, R, NEWTYPE, IERROR)
39     INTEGER P, R, NEWTYPE, IERROR
40
41 MPI_TYPE_CREATE_F90_INTEGER(R, NEWTYPE, IERROR)
42     INTEGER R, NEWTYPE, IERROR
43
44 MPI_TYPE_CREATE_F90_REAL(P, R, NEWTYPE, IERROR)
45     INTEGER P, R, NEWTYPE, IERROR
46
47 MPI_TYPE_MATCH_SIZE(TYPECLASS, SIZE, DATATYPE, IERROR)
48     INTEGER TYPECLASS, SIZE, DATATYPE, IERROR

```


A.3.13 Profiling Interface Fortran Bindings

```

MPI_PCONTROL(LEVEL)
    INTEGER LEVEL

```

A.3.14 Deprecated Fortran Bindings

```

MPI_ADDRESS(LOCATION, ADDRESS, IERROR)
    <type> LOCATION(*)
    INTEGER ADDRESS, IERROR

MPI_ATTR_DELETE(COMM, KEYVAL, IERROR)
    INTEGER COMM, KEYVAL, IERROR

MPI_ATTR_GET(COMM, KEYVAL, ATTRIBUTE_VAL, FLAG, IERROR)
    INTEGER COMM, KEYVAL, ATTRIBUTE_VAL, IERROR
    LOGICAL FLAG

MPI_ATTR_PUT(COMM, KEYVAL, ATTRIBUTE_VAL, IERROR)
    INTEGER COMM, KEYVAL, ATTRIBUTE_VAL, IERROR

MPI_DUP_FN(OLDCOMM, KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,
           ATTRIBUTE_VAL_OUT, FLAG, IERR)
    INTEGER OLDCOMM, KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,
    ATTRIBUTE_VAL_OUT, IERR
    LOGICAL FLAG

MPI_ERRHANDLER_CREATE(HANDLER_FN, ERRHANDLER, IERROR)
    EXTERNAL HANDLER_FN
    INTEGER ERRHANDLER, IERROR

MPI_ERRHANDLER_GET(COMM, ERRHANDLER, IERROR)
    INTEGER COMM, ERRHANDLER, IERROR

MPI_ERRHANDLER_SET(COMM, ERRHANDLER, IERROR)
    INTEGER COMM, ERRHANDLER, IERROR

MPI_KEYVAL_CREATE(COPY_FN, DELETE_FN, KEYVAL, EXTRA_STATE, IERROR)
    EXTERNAL COPY_FN, DELETE_FN
    INTEGER KEYVAL, EXTRA_STATE, IERROR

MPI_KEYVAL_FREE(KEYVAL, IERROR)
    INTEGER KEYVAL, IERROR

MPI_NULL_COPY_FN(OLDCOMM, KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,
                ATTRIBUTE_VAL_OUT, FLAG, IERR)
    INTEGER OLDCOMM, KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,
    ATTRIBUTE_VAL_OUT, IERR
    LOGICAL FLAG

MPI_NULL_DELETE_FN(COMM, KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE, IERROR)
    INTEGER COMM, KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE, IERROR

```

A.4 Fortran 2008 Bindings with the mpi_f08 Module

A.4.1 Point-to-Point Communication Fortran 2008 Bindings

```

MPI_Bsend(buf, count, datatype, dest, tag, comm, ierror)
  TYPE(*), DIMENSION(..) :: buf
  INTEGER, INTENT(IN) :: count, dest, tag
  TYPE(MPI_Datatype), INTENT(IN) :: datatype
  TYPE(MPI_Comm), INTENT(IN) :: comm
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Bsend_init(buf, count, datatype, dest, tag, comm, request, ierror)
  TYPE(*), DIMENSION(..) :: buf
  INTEGER, INTENT(IN) :: count, dest, tag
  TYPE(MPI_Datatype), INTENT(IN) :: datatype
  TYPE(MPI_Comm), INTENT(IN) :: comm
  TYPE(MPI_Request), INTENT(OUT) :: request
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Buffer_attach(buffer, size, ierror)
  TYPE(*), DIMENSION(..) :: buffer
  INTEGER, INTENT(IN) :: size
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Buffer_detach(buffer_addr, size, ierror)
  TYPE(*), DIMENSION(..) :: buffer_addr
  INTEGER, INTENT(OUT) :: size
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Cancel(request, ierror)
  TYPE(MPI_Request), INTENT(IN) :: request
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Get_count(status, datatype, count, ierror)
  TYPE(MPI_Status), INTENT(IN) :: status
  TYPE(MPI_Datatype), INTENT(IN) :: datatype
  INTEGER, INTENT(OUT) :: count
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Ibsend(buf, count, datatype, dest, tag, comm, request, ierror)
  TYPE(*), DIMENSION(..) :: buf
  INTEGER, INTENT(IN) :: count, dest, tag
  TYPE(MPI_Datatype), INTENT(IN) :: datatype
  TYPE(MPI_Comm), INTENT(IN) :: comm
  TYPE(MPI_Request), INTENT(OUT) :: request
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Iprobe(source, tag, comm, flag, status, ierror)
  INTEGER, INTENT(IN) :: source, tag
  TYPE(MPI_Comm), INTENT(IN) :: comm
  LOGICAL, INTENT(OUT) :: flag

```

```

1      TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
2      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
3
4      MPI_Irecv(buf, count, datatype, source, tag, comm, request, ierror)
5          TYPE(*), DIMENSION(..) :: buf
6          INTEGER, INTENT(IN) :: count, source, tag
7          TYPE(MPI_Datatype), INTENT(IN) :: datatype
8          TYPE(MPI_Comm), INTENT(IN) :: comm
9          TYPE(MPI_Request), INTENT(OUT) :: request
10         INTEGER, OPTIONAL, INTENT(OUT) :: ierror
11
12      MPI_Irsend(buf, count, datatype, dest, tag, comm, request, ierror)
13          TYPE(*), DIMENSION(..) :: buf
14          INTEGER, INTENT(IN) :: count, dest, tag
15          TYPE(MPI_Datatype), INTENT(IN) :: datatype
16          TYPE(MPI_Comm), INTENT(IN) :: comm
17          TYPE(MPI_Request), INTENT(OUT) :: request
18          INTEGER, OPTIONAL, INTENT(OUT) :: ierror
19
20      MPI_Isend(buf, count, datatype, dest, tag, comm, request, ierror)
21          TYPE(*), DIMENSION(..) :: buf
22          INTEGER, INTENT(IN) :: count, dest, tag
23          TYPE(MPI_Datatype), INTENT(IN) :: datatype
24          TYPE(MPI_Comm), INTENT(IN) :: comm
25          TYPE(MPI_Request), INTENT(OUT) :: request
26          INTEGER, OPTIONAL, INTENT(OUT) :: ierror
27
28      MPI_Issend(buf, count, datatype, dest, tag, comm, request, ierror)
29          TYPE(*), DIMENSION(..) :: buf
30          INTEGER, INTENT(IN) :: count, dest, tag
31          TYPE(MPI_Datatype), INTENT(IN) :: datatype
32          TYPE(MPI_Comm), INTENT(IN) :: comm
33          TYPE(MPI_Request), INTENT(OUT) :: request
34          INTEGER, OPTIONAL, INTENT(OUT) :: ierror
35
36      MPI_Probe(source, tag, comm, status, ierror)
37          INTEGER, INTENT(IN) :: source, tag
38          TYPE(MPI_Comm), INTENT(IN) :: comm
39          TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
40          INTEGER, OPTIONAL, INTENT(OUT) :: ierror
41
42      MPI_Recv(buf, count, datatype, source, tag, comm, status, ierror)
43          TYPE(*), DIMENSION(..) :: buf
44          INTEGER, INTENT(IN) :: count, source, tag
45          TYPE(MPI_Datatype), INTENT(IN) :: datatype
46          TYPE(MPI_Comm), INTENT(IN) :: comm
47          TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
48          INTEGER, OPTIONAL, INTENT(OUT) :: ierror
49
50      MPI_Recv_init(buf, count, datatype, source, tag, comm, request, ierror)
51          TYPE(*), DIMENSION(..) :: buf

```

```

    INTEGER, INTENT(IN) :: count, source, tag
    TYPE(MPI_Datatype), INTENT(IN) :: datatype
    TYPE(MPI_Comm), INTENT(IN) :: comm
    TYPE(MPI_Request), INTENT(OUT) :: request
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Request_free(request, ierror)
    TYPE(MPI_Request), INTENT(INOUT) :: request
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Request_get_status( request, flag, status, ierror)
    TYPE(MPI_Request), INTENT(IN) :: request
    LOGICAL, INTENT(OUT) :: flag
    TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Rsend(buf, count, datatype, dest, tag, comm, ierror)
    TYPE(*), DIMENSION(..) :: buf
    INTEGER, INTENT(IN) :: count, dest, tag
    TYPE(MPI_Datatype), INTENT(IN) :: datatype
    TYPE(MPI_Comm), INTENT(IN) :: comm
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Rsend_init(buf, count, datatype, dest, tag, comm, request, ierror)
    TYPE(*), DIMENSION(..) :: buf
    INTEGER, INTENT(IN) :: count, dest, tag
    TYPE(MPI_Datatype), INTENT(IN) :: datatype
    TYPE(MPI_Comm), INTENT(IN) :: comm
    TYPE(MPI_Request), INTENT(OUT) :: request
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Send(buf, count, datatype, dest, tag, comm, ierror)
    TYPE(*), DIMENSION(..) :: buf
    INTEGER, INTENT(IN) :: count, dest, tag
    TYPE(MPI_Datatype), INTENT(IN) :: datatype
    TYPE(MPI_Comm), INTENT(IN) :: comm
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Send_init(buf, count, datatype, dest, tag, comm, request, ierror)
    TYPE(*), DIMENSION(..) :: buf
    INTEGER, INTENT(IN) :: count, dest, tag
    TYPE(MPI_Datatype), INTENT(IN) :: datatype
    TYPE(MPI_Comm), INTENT(IN) :: comm
    TYPE(MPI_Request), INTENT(OUT) :: request
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Sendrecv(sendbuf, sendcount, sendtype, dest, sendtag, recvbuf,
    recvcount, recvtype, source, recvtag, comm, status, ierror)
    TYPE(*), DIMENSION(..) :: sendbuf, recvbuf
    INTEGER, INTENT(IN) :: sendcount, dest, sendtag, recvcount, source,
    recvtag

```

```

1      TYPE(MPI_Datatype), INTENT(IN) :: sendtype, recvtype
2      TYPE(MPI_Comm), INTENT(IN) :: comm
3      TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
4      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
5
6      MPI_Sendrecv_replace(buf, count, datatype, dest, sendtag, source, recvtag,
7                          comm, status, ierror)
8      TYPE(*), DIMENSION(..) :: buf
9      INTEGER, INTENT(IN) :: count, dest, sendtag, source, recvtag
10     TYPE(MPI_Datatype), INTENT(IN) :: datatype
11     TYPE(MPI_Comm), INTENT(IN) :: comm
12     TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
13     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
14
15     MPI_Ssend(buf, count, datatype, dest, tag, comm, ierror)
16     TYPE(*), DIMENSION(..) :: buf
17     INTEGER, INTENT(IN) :: count, dest, tag
18     TYPE(MPI_Datatype), INTENT(IN) :: datatype
19     TYPE(MPI_Comm), INTENT(IN) :: comm
20     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
21
22     MPI_Ssend_init(buf, count, datatype, dest, tag, comm, request, ierror)
23     TYPE(*), DIMENSION(..) :: buf
24     INTEGER, INTENT(IN) :: count, dest, tag
25     TYPE(MPI_Datatype), INTENT(IN) :: datatype
26     TYPE(MPI_Comm), INTENT(IN) :: comm
27     TYPE(MPI_Request), INTENT(OUT) :: request
28     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
29
30     MPI_Start(request, ierror)
31     TYPE(MPI_Request), INTENT(INOUT) :: request
32     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
33
34     MPI_Startall(count, array_of_requests, ierror)
35     INTEGER, INTENT(IN) :: count
36     TYPE(MPI_Request), INTENT(INOUT) :: array_of_requests(*)
37     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
38
39     MPI_Test(request, flag, status, ierror)
40     LOGICAL, INTENT(OUT) :: flag
41     TYPE(MPI_Request), INTENT(INOUT) :: request
42     TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
43     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
44
45     MPI_Test_cancelled(status, flag, ierror)
46     TYPE(MPI_Status), INTENT(IN) :: status
47     LOGICAL, INTENT(OUT) :: flag
48     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
49
50     MPI_Testall(count, array_of_requests, flag, array_of_statuses, ierror)
51     INTEGER, INTENT(IN) :: count

```

```

TYPE(MPI_Request), INTENT(INOUT) :: array_of_requests(*)
LOGICAL, INTENT(OUT) :: flag
TYPE(MPI_Status), INTENT(OUT) :: array_of_statuses(*) ! optional by
overloading
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Testany(count, array_of_requests, index, flag, status, ierror)
INTEGER, INTENT(IN) :: count
TYPE(MPI_Request), INTENT(INOUT) :: array_of_requests(*)
INTEGER, INTENT(OUT) :: index
LOGICAL, INTENT(OUT) :: flag
TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Testsome(incount, array_of_requests, outcount, array_of_indices,
array_of_statuses, ierror)
INTEGER, INTENT(IN) :: incount
TYPE(MPI_Request), INTENT(INOUT) :: array_of_requests(*)
INTEGER, INTENT(OUT) :: outcount, array_of_indices(*)
TYPE(MPI_Status), INTENT(OUT) :: array_of_statuses(*) ! optional by
overloading
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Wait(request, status, ierror)
TYPE(MPI_Request), INTENT(INOUT) :: request
TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Waitall(count, array_of_requests, array_of_statuses, ierror)
INTEGER, INTENT(IN) :: count
INTEGER, INTENT(INOUT) :: array_of_requests(*)
TYPE(MPI_Status), INTENT(OUT) :: array_of_statuses(*) ! optional by
overloading
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Waitany(count, array_of_requests, index, status, ierror)
INTEGER, INTENT(IN) :: count
TYPE(MPI_Request), INTENT(INOUT) :: array_of_requests(*)
INTEGER, INTENT(OUT) :: index
TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Waitsome(incount, array_of_requests, outcount, array_of_indices,
array_of_statuses, ierror)
INTEGER, INTENT(IN) :: incount
TYPE(MPI_Request), INTENT(INOUT) :: array_of_requests(*)
INTEGER, INTENT(OUT) :: outcount, array_of_indices(*)
TYPE(MPI_Status), INTENT(OUT) :: array_of_statuses(*) ! optional by
overloading
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

A.4.2 Datatypes Fortran 2008 Bindings

```

1  MPI_Get_address(location, address, ierror)
2
3  TYPE(*), DIMENSION(..) :: location
4  INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(OUT) :: address
5  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
6
7  MPI_Get_elements(status, datatype, count, ierror)
8  TYPE(MPI_Status), INTENT(IN) :: status
9  TYPE(MPI_Datatype), INTENT(IN) :: datatype
10 INTEGER, INTENT(OUT) :: count
11 INTEGER, OPTIONAL, INTENT(OUT) :: ierror
12
13 MPI_Pack(inbuf, incount, datatype, outbuf, outsize, position, comm, ierror)
14 TYPE(*), DIMENSION(..) :: inbuf, outbuf
15 INTEGER, INTENT(IN) :: incount, outsize
16 TYPE(MPI_Datatype), INTENT(IN) :: datatype
17 INTEGER, INTENT(INOUT) :: position
18 TYPE(MPI_Comm), INTENT(IN) :: comm
19 INTEGER, OPTIONAL, INTENT(OUT) :: ierror
20
21 MPI_Pack_external(datarep, inbuf, incount, datatype, outbuf, outsize,
22                  position, ierror)
23 CHARACTER(LEN=*), INTENT(IN) :: datarep
24 TYPE(*), DIMENSION(..) :: inbuf, outbuf
25 INTEGER, INTENT(IN) :: incount
26 TYPE(MPI_Datatype), INTENT(IN) :: datatype
27 INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: outsize
28 INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(INOUT) :: position
29 INTEGER, OPTIONAL, INTENT(OUT) :: ierror
30
31 MPI_Pack_external_size(datarep, incount, datatype, size, ierror)
32 TYPE(MPI_Datatype), INTENT(IN) :: datatype
33 INTEGER, INTENT(IN) :: incount
34 CHARACTER(LEN=*), INTENT(IN) :: datarep
35 INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(OUT) :: size
36 INTEGER, OPTIONAL, INTENT(OUT) :: ierror
37
38 MPI_Pack_size(incount, datatype, comm, size, ierror)
39 INTEGER, INTENT(IN) :: incount
40 TYPE(MPI_Datatype), INTENT(IN) :: datatype
41 TYPE(MPI_Comm), INTENT(IN) :: comm
42 INTEGER, INTENT(OUT) :: size
43 INTEGER, OPTIONAL, INTENT(OUT) :: ierror
44
45 MPI_Type_commit(datatype, ierror)
46 TYPE(MPI_Datatype), INTENT(INOUT) :: datatype
47 INTEGER, OPTIONAL, INTENT(OUT) :: ierror
48
49 MPI_Type_contiguous(count, oldtype, newtype, ierror)
50 INTEGER, INTENT(IN) :: count

```

```

TYPE(MPI_Datatype), INTENT(IN) :: oldtype
TYPE(MPI_Datatype), INTENT(OUT) :: newtype
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Type_create_darray(size, rank, ndims, array_of_gsizes,
    array_of_distribs, array_of_dargs, array_of_psize, order,
    oldtype, newtype, ierror)
INTEGER, INTENT(IN) :: size, rank, ndims, array_of_gsizes(*),
array_of_distribs(*), array_of_dargs(*), array_of_psize(*), order
TYPE(MPI_Datatype), INTENT(IN) :: oldtype
TYPE(MPI_Datatype), INTENT(OUT) :: newtype
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Type_create_hindexed(count, array_of_blocklengths,
    array_of_displacements, oldtype, newtype, ierror)
INTEGER, INTENT(IN) :: count, array_of_blocklengths(*)
INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) ::
array_of_displacements(*)
TYPE(MPI_Datatype), INTENT(IN) :: oldtype
TYPE(MPI_Datatype), INTENT(OUT) :: newtype
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Type_create_hvector(count, blocklength, stride, oldtype, newtype,
    ierror)
INTEGER, INTENT(IN) :: count, blocklength
INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: stride
TYPE(MPI_Datatype), INTENT(IN) :: oldtype
TYPE(MPI_Datatype), INTENT(OUT) :: newtype
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Type_create_indexed_block(count, blocklength, array_of_displacements,
    oldtype, newtype, ierror)
INTEGER, INTENT(IN) :: count, blocklength, array_of_displacements(*)
TYPE(MPI_Datatype), INTENT(IN) :: oldtype
TYPE(MPI_Datatype), INTENT(OUT) :: newtype
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Type_create_resized(oldtype, lb, extent, newtype, ierror)
INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: lb, extent
TYPE(MPI_Datatype), INTENT(IN) :: oldtype
TYPE(MPI_Datatype), INTENT(OUT) :: newtype
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Type_create_struct(count, array_of_blocklengths,
    array_of_displacements, array_of_types, newtype, ierror)
INTEGER, INTENT(IN) :: count, array_of_blocklengths(*)
INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) ::
array_of_displacements(*)
TYPE(MPI_Datatype), INTENT(IN) :: array_of_types(*)
TYPE(MPI_Datatype), INTENT(OUT) :: newtype

```



```

1      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
2
3      MPI_Type_create_subarray(ndims, array_of_sizes, array_of_subsizes,
4          array_of_starts, order, oldtype, newtype, ierror)
5      INTEGER, INTENT(IN) :: ndims, array_of_sizes(*), array_of_subsizes(*),
6      array_of_starts(*), order
7      TYPE(MPI_Datatype), INTENT(IN) :: oldtype
8      TYPE(MPI_Datatype), INTENT(OUT) :: newtype
9      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
10
11     MPI_Type_dup(oldtype, newtype, ierror)
12     TYPE(MPI_Datatype), INTENT(IN) :: oldtype
13     TYPE(MPI_Datatype), INTENT(OUT) :: newtype
14     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
15
16     MPI_Type_free(datatype, ierror)
17     TYPE(MPI_Datatype), INTENT(INOUT) :: datatype
18     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
19
20     MPI_Type_get_contents(datatype, max_integers, max_addresses, max_datatypes,
21         array_of_integers, array_of_addresses, array_of_datatypes,
22         ierror)
23     TYPE(MPI_Datatype), INTENT(IN) :: datatype
24     INTEGER, INTENT(IN) :: max_integers, max_addresses, max_datatypes
25     INTEGER, INTENT(OUT) :: array_of_integers(*)
26     INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(OUT) :: array_of_addresses(*)
27     TYPE(MPI_Datatype), INTENT(OUT) :: array_of_datatypes(*)
28     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
29
30     MPI_Type_get_envelope(datatype, num_integers, num_addresses, num_datatypes,
31         combiner, ierror)
32     TYPE(MPI_Datatype), INTENT(IN) :: datatype
33     INTEGER, INTENT(OUT) :: num_integers, num_addresses, num_datatypes,
34     combiner
35     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
36
37     MPI_Type_get_extent(datatype, lb, extent, ierror)
38     TYPE(MPI_Datatype), INTENT(IN) :: datatype
39     INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(OUT) :: lb, extent
40     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
41
42     MPI_Type_get_true_extent(datatype, true_lb, true_extent, ierror)
43     TYPE(MPI_Datatype), INTENT(IN) :: datatype
44     INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(OUT) :: true_lb, true_extent
45     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
46
47     MPI_Type_indexed(count, array_of_blocklengths, array_of_displacements,
48         oldtype, newtype, ierror)
49     INTEGER, INTENT(IN) :: count, array_of_blocklengths(*),
50     array_of_displacements(*)
51     TYPE(MPI_Datatype), INTENT(IN) :: oldtype

```

```

TYPE(MPI_Datatype), INTENT(OUT) :: newtype
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Type_size(datatype, size, ierror)
TYPE(MPI_Datatype), INTENT(IN) :: datatype
INTEGER, INTENT(OUT) :: size
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Type_vector(count, blocklength, stride, oldtype, newtype, ierror)
INTEGER, INTENT(IN) :: count, blocklength, stride
TYPE(MPI_Datatype), INTENT(IN) :: oldtype
TYPE(MPI_Datatype), INTENT(OUT) :: newtype
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Unpack(inbuf, insize, position, outbuf, outcount, datatype, comm,
           ierror)
TYPE(*), DIMENSION(..) :: inbuf, outbuf
INTEGER, INTENT(IN) :: insize, outcount
INTEGER, INTENT(INOUT) :: position
TYPE(MPI_Datatype), INTENT(IN) :: datatype
TYPE(MPI_Comm), INTENT(IN) :: comm
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Unpack_external(datarep, inbuf, insize, position, outbuf, outcount,
                    datatype, ierror)
CHARACTER(LEN=*), INTENT(IN) :: datarep
TYPE(*), DIMENSION(..) :: inbuf, outbuf
INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: insize
INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(INOUT) :: position
INTEGER, INTENT(IN) :: outcount
TYPE(MPI_Datatype), INTENT(IN) :: datatype
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

A.4.3 Collective Communication Fortran 2008 Bindings

```

MPI_Allgather(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype,
              comm, ierror)
TYPE(*), DIMENSION(..) :: sendbuf, recvbuf
INTEGER, INTENT(IN) :: sendcount, recvcount
TYPE(MPI_Datatype), INTENT(IN) :: sendtype
TYPE(MPI_Datatype), INTENT(IN) :: recvtype
TYPE(MPI_Comm), INTENT(IN) :: comm
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Allgatherv(sendbuf, sendcount, sendtype, recvbuf, recvcounts, displs,
               recvtype, comm, ierror)
TYPE(*), DIMENSION(..) :: sendbuf, recvbuf
INTEGER, INTENT(IN) :: sendcount, recvcounts(*), displs(*)
TYPE(MPI_Datatype), INTENT(IN) :: sendtype
TYPE(MPI_Datatype), INTENT(IN) :: recvtype

```

```

1      TYPE(MPI_Comm), INTENT(IN) :: comm
2      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
3
4      MPI_Allreduce(sendbuf, recvbuf, count, datatype, op, comm, ierror)
5          TYPE(*), DIMENSION(..) :: sendbuf, recvbuf
6          INTEGER, INTENT(IN) :: count
7          TYPE(MPI_Datatype), INTENT(IN) :: datatype
8          TYPE(MPI_Op), INTENT(IN) :: op
9          TYPE(MPI_Comm), INTENT(IN) :: comm
10         INTEGER, OPTIONAL, INTENT(OUT) :: ierror
11
12     MPI_Alltoall(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype,
13                 comm, ierror)
14         TYPE(*), DIMENSION(..) :: sendbuf, recvbuf
15         INTEGER, INTENT(IN) :: sendcount, recvcount
16         TYPE(MPI_Datatype), INTENT(IN) :: sendtype
17         TYPE(MPI_Datatype), INTENT(IN) :: recvtype
18         TYPE(MPI_Comm), INTENT(IN) :: comm
19         INTEGER, OPTIONAL, INTENT(OUT) :: ierror
20
21     MPI_Alltoallv(sendbuf, sendcounts, sdispls, sendtype, recvbuf, recvcounts,
22                  rdispls, recvtype, comm, ierror)
23         TYPE(*), DIMENSION(..) :: sendbuf, recvbuf
24         INTEGER, INTENT(IN) :: sendcounts(*), sdispls(*), recvcounts(*),
25         rdispls(*)
26         TYPE(MPI_Datatype), INTENT(IN) :: sendtype
27         TYPE(MPI_Datatype), INTENT(IN) :: recvtype
28         TYPE(MPI_Comm), INTENT(IN) :: comm
29         INTEGER, OPTIONAL, INTENT(OUT) :: ierror
30
31     MPI_Alltoallw(sendbuf, sendcounts, sdispls, sendtypes, recvbuf, recvcounts,
32                  rdispls, recvtypes, comm, ierror)
33         TYPE(*), DIMENSION(..) :: sendbuf, recvbuf
34         INTEGER, INTENT(IN) :: sendcounts(*), sdispls(*), recvcounts(*),
35         rdispls(*)
36         TYPE(MPI_Datatype), INTENT(IN) :: sendtypes(*)
37         TYPE(MPI_Datatype), INTENT(IN) :: recvtypes(*)
38         TYPE(MPI_Comm), INTENT(IN) :: comm
39         INTEGER, OPTIONAL, INTENT(OUT) :: ierror
40
41     MPI_Barrier(comm, ierror)
42         TYPE(MPI_Comm), INTENT(IN) :: comm
43         INTEGER, OPTIONAL, INTENT(OUT) :: ierror
44
45     MPI_Bcast(buffer, count, datatype, root, comm, ierror)
46         TYPE(*), DIMENSION(..) :: buffer
47         INTEGER, INTENT(IN) :: count, root
48         TYPE(MPI_Datatype), INTENT(IN) :: datatype
49         TYPE(MPI_Comm), INTENT(IN) :: comm
50         INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

```

MPI_Exscan(sendbuf, recvbuf, count, datatype, op, comm, ierror)
    TYPE(*), DIMENSION(..) :: sendbuf, recvbuf
    INTEGER, INTENT(IN) :: count
    TYPE(MPI_Datatype), INTENT(IN) :: datatype
    TYPE(MPI_Op), INTENT(IN) :: op
    TYPE(MPI_Comm), INTENT(IN) :: comm
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Gather(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvttype,
    root, comm, ierror)
    TYPE(*), DIMENSION(..) :: sendbuf, recvbuf
    INTEGER, INTENT(IN) :: sendcount, recvcount, root
    TYPE(MPI_Datatype), INTENT(IN) :: sendtype
    TYPE(MPI_Datatype), INTENT(IN) :: recvttype
    TYPE(MPI_Comm), INTENT(IN) :: comm
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Gatherv(sendbuf, sendcount, sendtype, recvbuf, recvcounts, displs,
    recvttype, root, comm, ierror)
    TYPE(*), DIMENSION(..) :: sendbuf, recvbuf
    INTEGER, INTENT(IN) :: sendcount, recvcounts(*), displs(*), root
    TYPE(MPI_Datatype), INTENT(IN) :: sendtype
    TYPE(MPI_Datatype), INTENT(IN) :: recvttype
    TYPE(MPI_Comm), INTENT(IN) :: comm
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Iallgather( -- Fortran 2008 -- to be done -- )

MPI_Iallgatherv( -- Fortran 2008 -- to be done -- )

MPI_Iallreduce( -- Fortran 2008 -- to be done -- )

MPI_Ialltoall( -- Fortran 2008 -- to be done -- )

MPI_Ialltoallv( -- Fortran 2008 -- to be done -- )

MPI_Ialltoallw( -- Fortran 2008 -- to be done -- )

MPI_Ibarrier( -- Fortran 2008 -- to be done -- )

MPI_Ibcast( -- Fortran 2008 -- to be done -- )

MPI_Iexscan( -- Fortran 2008 -- to be done -- )

MPI_Igather( -- Fortran 2008 -- to be done -- )

MPI_Igatherv( -- Fortran 2008 -- to be done -- )

MPI_Ireduce( -- Fortran 2008 -- to be done -- )

MPI_Ireduce_scatter( -- Fortran 2008 -- to be done -- )

MPI_Ireduce_scatter_block( -- Fortran 2008 -- to be done -- )

MPI_Iscan( -- Fortran 2008 -- to be done -- )

```

```

1  MPI_Iscatter( -- Fortran 2008 -- to be done -- )
2
3  MPI_Iscatterv( -- Fortran 2008 -- to be done -- )
4
5  MPI_Op_commutative(op, commute, ierror)
6      TYPE(MPI_Op), INTENT(IN) :: op
7      LOGICAL, INTENT(OUT) :: commute
8      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
9
10 MPI_Op_create(user_fn, commute, op, ierror)
11     EXTERNAL :: user_fn
12     LOGICAL, INTENT(IN) :: commute
13     TYPE(MPI_Op), INTENT(OUT) :: op
14     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
15
16 MPI_Op_free(op, ierror)
17     TYPE(MPI_Op), INTENT(INOUT) :: op
18     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
19
20 MPI_Reduce(sendbuf, recvbuf, count, datatype, op, root, comm, ierror)
21     TYPE(*), DIMENSION(..) :: sendbuf, recvbuf
22     INTEGER, INTENT(IN) :: count, root
23     TYPE(MPI_Datatype), INTENT(IN) :: datatype
24     TYPE(MPI_Op), INTENT(IN) :: op
25     TYPE(MPI_Comm), INTENT(IN) :: comm
26     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
27
28 MPI_Reduce_local(inbuf, inoutbuf, count, datatype, op, ierror)
29     TYPE(*), DIMENSION(..) :: inbuf, inoutbuf
30     INTEGER, INTENT(IN) :: count
31     TYPE(MPI_Datatype), INTENT(IN) :: datatype
32     TYPE(MPI_Op), INTENT(IN) :: op
33     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
34
35 MPI_Reduce_scatter(sendbuf, recvbuf, recvcounts, datatype, op, comm,
36     ierror)
37     TYPE(*), DIMENSION(..) :: sendbuf, recvbuf
38     INTEGER, INTENT(IN) :: recvcounts(*)
39     TYPE(MPI_Datatype), INTENT(IN) :: datatype
40     TYPE(MPI_Op), INTENT(IN) :: op
41     TYPE(MPI_Comm), INTENT(IN) :: comm
42     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
43
44 MPI_Reduce_scatter_block(sendbuf, recvbuf, recvcount, datatype, op, comm,
45     ierror)
46     TYPE(*), DIMENSION(..) :: sendbuf, recvbuf
47     INTEGER, INTENT(IN) :: recvcount
48     TYPE(MPI_Datatype), INTENT(IN) :: datatype
49     TYPE(MPI_Op), INTENT(IN) :: op
50     TYPE(MPI_Comm), INTENT(IN) :: comm
51     INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

```

MPI_Scan(sendbuf, recvbuf, count, datatype, op, comm, ierror)
    TYPE(*), DIMENSION(..) :: sendbuf, recvbuf
    INTEGER, INTENT(IN) :: count
    TYPE(MPI_Datatype), INTENT(IN) :: datatype
    TYPE(MPI_Op), INTENT(IN) :: op
    TYPE(MPI_Comm), INTENT(IN) :: comm
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Scatter(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvttype,
    root, comm, ierror)
    TYPE(*), DIMENSION(..) :: sendbuf, recvbuf
    INTEGER, INTENT(IN) :: sendcount, recvcount, root
    TYPE(MPI_Datatype), INTENT(IN) :: sendtype
    TYPE(MPI_Datatype), INTENT(IN) :: recvttype
    TYPE(MPI_Comm), INTENT(IN) :: comm
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Scatterv(sendbuf, sendcounts, displs, sendtype, recvbuf, recvcount,
    recvttype, root, comm, ierror)
    TYPE(*), DIMENSION(..) :: sendbuf, recvbuf
    INTEGER, INTENT(IN) :: sendcounts(*), displs(*), recvcount, root
    TYPE(MPI_Datatype), INTENT(IN) :: sendtype
    TYPE(MPI_Datatype), INTENT(IN) :: recvttype
    TYPE(MPI_Comm), INTENT(IN) :: comm
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

A.4.4 Groups, Contexts, Communicators, and Caching Fortran 2008 Bindings

```

MPI_COMM_DUP_FN(oldcomm, comm_keyval, extra_state, attribute_val_in,
    attribute_val_out, flag, ierror)
    TYPE(MPI_Comm), INTENT(IN) :: oldcomm
    INTEGER, INTENT(IN) :: comm_keyval
    INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: extra_state,
    attribute_val_in
    INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(OUT) :: attribute_val_out
    LOGICAL, INTENT(OUT) :: flag
    INTEGER, INTENT(OUT) :: ierror

MPI_COMM_NULL_COPY_FN(oldcomm, comm_keyval, extra_state, attribute_val_in,
    attribute_val_out, flag, ierror)
    TYPE(MPI_Comm), INTENT(IN) :: oldcomm
    INTEGER, INTENT(IN) :: comm_keyval
    INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: extra_state,
    attribute_val_in
    INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(OUT) :: attribute_val_out
    LOGICAL, INTENT(OUT) :: flag
    INTEGER, INTENT(OUT) :: ierror

MPI_COMM_NULL_DELETE_FN(comm, comm_keyval, attribute_val, extra_state,

```

```

1         ierror)
2     TYPE(MPI_Comm), INTENT(IN) :: comm
3     INTEGER, INTENT(IN) :: comm_keyval
4     INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: attribute_val,
5     extra_state
6     INTEGER, INTENT(OUT) :: ierror
7
8 MPI_Comm_compare(comm1, comm2, result, ierror)
9     TYPE(MPI_Comm), INTENT(IN) :: comm1
10    TYPE(MPI_Comm), INTENT(IN) :: comm2
11    INTEGER, INTENT(OUT) :: result
12    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
13
14 MPI_Comm_create(comm, group, newcomm, ierror)
15    TYPE(MPI_Comm), INTENT(IN) :: comm
16    TYPE(MPI_Group), INTENT(IN) :: group
17    TYPE(MPI_Comm), INTENT(OUT) :: newcomm
18    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
19
20 MPI_Comm_create_keyval(comm_copy_attr_fn, comm_delete_attr_fn, comm_keyval,
21    extra_state, ierror)
22    EXTERNAL :: comm_copy_attr_fn, comm_delete_attr_fn
23    INTEGER, INTENT(OUT) :: comm_keyval
24    INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: extra_state
25    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
26
27 MPI_Comm_delete_attr(comm, comm_keyval, ierror)
28    TYPE(MPI_Comm), INTENT(IN) :: comm
29    INTEGER, INTENT(IN) :: comm_keyval
30    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
31
32 MPI_Comm_dup(comm, newcomm, ierror)
33    TYPE(MPI_Comm), INTENT(IN) :: comm
34    TYPE(MPI_Comm), INTENT(OUT) :: newcomm
35    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
36
37 MPI_Comm_free(comm, ierror)
38    TYPE(MPI_Comm), INTENT(INOUT) :: comm
39    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
40
41 MPI_Comm_free_keyval(comm_keyval, ierror)
42    INTEGER, INTENT(INOUT) :: comm_keyval
43    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
44
45 MPI_Comm_get_attr(comm, comm_keyval, attribute_val, flag, ierror)
46    TYPE(MPI_Comm), INTENT(IN) :: comm
47    INTEGER, INTENT(IN) :: comm_keyval
48    INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(OUT) :: attribute_val
49    LOGICAL, INTENT(OUT) :: flag
50    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
51
52 MPI_Comm_get_name(comm, comm_name, resultlen, ierror)

```

```

TYPE(MPI_Comm), INTENT(IN) :: comm
CHARACTER(LEN=*), INTENT(OUT) :: comm_name
INTEGER, INTENT(OUT) :: resultlen
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Comm_group(comm, group, ierror)
TYPE(MPI_Comm), INTENT(IN) :: comm
TYPE(MPI_Group), INTENT(OUT) :: group
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Comm_rank(comm, rank, ierror)
TYPE(MPI_Comm), INTENT(IN) :: comm
INTEGER, INTENT(OUT) :: rank
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Comm_remote_group(comm, group, ierror)
TYPE(MPI_Comm), INTENT(IN) :: comm
TYPE(MPI_Group), INTENT(OUT) :: group
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Comm_remote_size(comm, size, ierror)
TYPE(MPI_Comm), INTENT(IN) :: comm
INTEGER, INTENT(OUT) :: size
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Comm_set_attr(comm, comm_keyval, attribute_val, ierror)
TYPE(MPI_Comm), INTENT(IN) :: comm
INTEGER, INTENT(IN) :: comm_keyval
INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: attribute_val
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Comm_set_name(comm, comm_name, ierror)
TYPE(MPI_Comm), INTENT(IN) :: comm
CHARACTER(LEN=*), INTENT(IN) :: comm_name
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Comm_size(comm, size, ierror)
TYPE(MPI_Comm), INTENT(IN) :: comm
INTEGER, INTENT(OUT) :: size
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Comm_split(comm, color, key, newcomm, ierror)
TYPE(MPI_Comm), INTENT(IN) :: comm
INTEGER, INTENT(IN) :: color, key
TYPE(MPI_Comm), INTENT(OUT) :: newcomm
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Comm_test_inter(comm, flag, ierror)
TYPE(MPI_Comm), INTENT(IN) :: comm
LOGICAL, INTENT(OUT) :: flag
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```



```

1  MPI_Group_compare(group1, group2, result, ierror)
2      TYPE(MPI_Group), INTENT(IN) :: group1, group2
3      INTEGER, INTENT(OUT) :: result
4      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
5
6  MPI_Group_difference(group1, group2, newgroup, ierror)
7      TYPE(MPI_Group), INTENT(IN) :: group1, group2
8      TYPE(MPI_Group), INTENT(OUT) :: newgroup
9      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
10
11 MPI_Group_excl(group, n, ranks, newgroup, ierror)
12     TYPE(MPI_Group), INTENT(IN) :: group
13     INTEGER, INTENT(IN) :: n, ranks(*)
14     TYPE(MPI_Group), INTENT(OUT) :: newgroup
15     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
16
17 MPI_Group_free(group, ierror)
18     TYPE(MPI_Group), INTENT(INOUT) :: group
19     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
20
21 MPI_Group_incl(group, n, ranks, newgroup, ierror)
22     INTEGER, INTENT(IN) :: n, ranks(*)
23     TYPE(MPI_Group), INTENT(IN) :: group
24     TYPE(MPI_Group), INTENT(OUT) :: newgroup
25     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
26
27 MPI_Group_intersection(group1, group2, newgroup, ierror)
28     TYPE(MPI_Group), INTENT(IN) :: group1, group2
29     TYPE(MPI_Group), INTENT(OUT) :: newgroup
30     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
31
32 MPI_Group_range_excl(group, n, ranges, newgroup, ierror)
33     TYPE(MPI_Group), INTENT(IN) :: group
34     INTEGER, INTENT(IN) :: n, ranges(3,*)
35     TYPE(MPI_Group), INTENT(OUT) :: newgroup
36     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
37
38 MPI_Group_range_incl(group, n, ranges, newgroup, ierror)
39     TYPE(MPI_Group), INTENT(IN) :: group
40     INTEGER, INTENT(IN) :: n, ranges(3,*)
41     TYPE(MPI_Group), INTENT(OUT) :: newgroup
42     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
43
44 MPI_Group_rank(group, rank, ierror)
45     TYPE(MPI_Group), INTENT(IN) :: group
46     INTEGER, INTENT(OUT) :: rank
47     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
48
49 MPI_Group_size(group, size, ierror)
50     TYPE(MPI_Group), INTENT(IN) :: group
51     INTEGER, INTENT(OUT) :: size
52     INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

```

MPI_Group_translate_ranks(group1, n, ranks1, group2, ranks2, ierror)
    TYPE(MPI_Group), INTENT(IN) :: group1, group2
    INTEGER, INTENT(IN) :: n, ranks1(*)
    INTEGER, INTENT(OUT) :: ranks2(*)
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Group_union(group1, group2, newgroup, ierror)
    TYPE(MPI_Group), INTENT(IN) :: group1, group2
    TYPE(MPI_Group), INTENT(OUT) :: newgroup
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Intercomm_create(local_comm, local_leader, peer_comm, remote_leader,
    tag, newintercomm, ierror)
    TYPE(MPI_Comm), INTENT(IN) :: local_comm, peer_comm
    INTEGER, INTENT(IN) :: local_leader, remote_leader, tag
    TYPE(MPI_Comm), INTENT(OUT) :: newintercomm
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Intercomm_merge(intercomm, high, newintracomm, ierror)
    TYPE(MPI_Comm), INTENT(IN) :: intercomm
    LOGICAL, INTENT(IN) :: high
    TYPE(MPI_Comm), INTENT(OUT) :: newintracomm
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_TYPE_DUP_FN(oldtype, type_keyval, extra_state, attribute_val_in,
    attribute_val_out, flag, ierror)
    TYPE(MPI_Datatype), INTENT(IN) :: oldtype
    INTEGER, INTENT(IN) :: type_keyval
    INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: extra_state,
    attribute_val_in
    INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(OUT) :: attribute_val_out
    LOGICAL, INTENT(OUT) :: flag
    INTEGER, INTENT(OUT) :: ierror

MPI_TYPE_NULL_COPY_FN(oldtype, type_keyval, extra_state, attribute_val_in,
    attribute_val_out, flag, ierror)
    TYPE(MPI_Datatype), INTENT(IN) :: oldtype
    INTEGER, INTENT(IN) :: type_keyval
    INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: extra_state,
    attribute_val_in
    INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(OUT) :: attribute_val_out
    LOGICAL, INTENT(OUT) :: flag
    INTEGER, INTENT(OUT) :: ierror

MPI_TYPE_NULL_DELETE_FN(datatype, type_keyval, attribute_val, extra_state,
    ierror)
    TYPE(MPI_Datatype), INTENT(IN) :: datatype
    INTEGER, INTENT(IN) :: type_keyval
    INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: attribute_val,
    extra_state

```

```

1      INTEGER, INTENT(OUT) :: ierror
2
3      MPI_Type_create_keyval(type_copy_attr_fn, type_delete_attr_fn, type_keyval,
4          extra_state, ierror)
5      EXTERNAL :: type_copy_attr_fn, type_delete_attr_fn
6      INTEGER, INTENT(OUT) :: type_keyval
7      INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: extra_state
8      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
9
10     MPI_Type_delete_attr(datatype, type_keyval, ierror)
11     TYPE(MPI_Datatype), INTENT(IN) :: datatype
12     INTEGER, INTENT(IN) :: type_keyval
13     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
14
15     MPI_Type_free_keyval(type_keyval, ierror)
16     INTEGER, INTENT(INOUT) :: type_keyval
17     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
18
19     MPI_Type_get_attr(datatype, type_keyval, attribute_val, flag, ierror)
20     TYPE(MPI_Datatype), INTENT(IN) :: datatype
21     INTEGER, INTENT(IN) :: type_keyval
22     INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(OUT) :: attribute_val
23     LOGICAL, INTENT(OUT) :: flag
24     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
25
26     MPI_Type_get_name(datatype, type_name, resultlen, ierror)
27     TYPE(MPI_Datatype), INTENT(IN) :: datatype
28     CHARACTER(LEN=*), INTENT(OUT) :: type_name
29     INTEGER, INTENT(OUT) :: resultlen
30     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
31
32     MPI_Type_set_attr(datatype, type_keyval, attribute_val, ierror)
33     TYPE(MPI_Datatype), INTENT(IN) :: datatype
34     INTEGER, INTENT(IN) :: type_keyval
35     INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: attribute_val
36     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
37
38     MPI_Type_set_name(datatype, type_name, ierror)
39     TYPE(MPI_Datatype), INTENT(IN) :: datatype
40     CHARACTER(LEN=*), INTENT(IN) :: type_name
41     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
42
43     MPI_WIN_DUP_FN(oldwin, win_keyval, extra_state, attribute_val_in,
44         attribute_val_out, flag, ierror)
45     INTEGER, INTENT(IN) :: oldwin, win_keyval
46     INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: extra_state,
47         attribute_val_in
48     INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(OUT) :: attribute_val_out
49     LOGICAL, INTENT(OUT) :: flag
50     INTEGER, INTENT(OUT) :: ierror
51
52     MPI_WIN_NULL_COPY_FN(oldwin, win_keyval, extra_state, attribute_val_in,

```

```

        attribute_val_out, flag, ierror)
INTEGER, INTENT(IN) :: oldwin, win_keyval
INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: extra_state,
attribute_val_in
INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(OUT) :: attribute_val_out
LOGICAL, INTENT(OUT) :: flag
INTEGER, INTENT(OUT) :: ierror

MPI_WIN_NULL_DELETE_FN(win, win_keyval, attribute_val, extra_state, ierror)
TYPE(MPI_Win), INTENT(IN) :: win
INTEGER, INTENT(IN) :: win_keyval
INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: attribute_val,
extra_state
INTEGER, INTENT(OUT) :: ierror

MPI_Win_create_keyval(win_copy_attr_fn, win_delete_attr_fn, win_keyval,
        extra_state, ierror)
EXTERNAL :: win_copy_attr_fn, win_delete_attr_fn
INTEGER, INTENT(OUT) :: win_keyval
INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: extra_state
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Win_delete_attr(win, win_keyval, ierror)
TYPE(MPI_Win), INTENT(IN) :: win
INTEGER, INTENT(IN) :: win_keyval
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Win_free_keyval(win_keyval, ierror)
INTEGER, INTENT(INOUT) :: win_keyval
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Win_get_attr(win, win_keyval, attribute_val, flag, ierror)
TYPE(MPI_Win), INTENT(IN) :: win
INTEGER, INTENT(IN) :: win_keyval
INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(OUT) :: attribute_val
LOGICAL, INTENT(OUT) :: flag
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Win_get_name(win, win_name, resultlen, ierror)
TYPE(MPI_Win), INTENT(IN) :: win
CHARACTER(LEN=*), INTENT(OUT) :: win_name
INTEGER, INTENT(OUT) :: resultlen
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Win_set_attr(win, win_keyval, attribute_val, ierror)
TYPE(MPI_Win), INTENT(IN) :: win
INTEGER, INTENT(IN) :: win_keyval
INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: attribute_val
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Win_set_name(win, win_name, ierror)

```

```

1      TYPE(MPI_Win), INTENT(IN) :: win
2      CHARACTER(LEN=*), INTENT(IN) :: win_name
3      INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

A.4.5 Process Topologies Fortran 2008 Bindings

```

7  MPI_Cart_coords(comm, rank, maxdims, coords, ierror)
8      TYPE(MPI_Comm), INTENT(IN) :: comm
9      INTEGER, INTENT(IN) :: rank, maxdims
10     INTEGER, INTENT(OUT) :: coords(*)
11     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
12
13 MPI_Cart_create(comm_old, ndims, dims, periods, reorder, comm_cart, ierror)
14     TYPE(MPI_Comm), INTENT(IN) :: comm_old
15     INTEGER, INTENT(IN) :: ndims, dims(*)
16     LOGICAL, INTENT(IN) :: periods(*), reorder
17     TYPE(MPI_Comm), INTENT(OUT) :: comm_cart
18     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
19
20 MPI_Cart_get(comm, maxdims, dims, periods, coords, ierror)
21     TYPE(MPI_Comm), INTENT(IN) :: comm
22     INTEGER, INTENT(IN) :: maxdims
23     INTEGER, INTENT(OUT) :: dims(*), coords(*)
24     LOGICAL, INTENT(OUT) :: periods(*)
25     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
26
27 MPI_Cart_map(comm, ndims, dims, periods, newrank, ierror)
28     TYPE(MPI_Comm), INTENT(IN) :: comm
29     INTEGER, INTENT(IN) :: ndims, dims(*)
30     LOGICAL, INTENT(IN) :: periods(*)
31     INTEGER, INTENT(OUT) :: newrank
32     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
33
34 MPI_Cart_rank(comm, coords, rank, ierror)
35     TYPE(MPI_Comm), INTENT(IN) :: comm
36     INTEGER, INTENT(IN) :: coords(*)
37     INTEGER, INTENT(OUT) :: rank
38     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
39
40 MPI_Cart_shift(comm, direction, disp, rank_source, rank_dest, ierror)
41     TYPE(MPI_Comm), INTENT(IN) :: comm
42     INTEGER, INTENT(IN) :: direction, disp
43     INTEGER, INTENT(OUT) :: rank_source, rank_dest
44     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
45
46 MPI_Cart_sub(comm, remain_dims, newcomm, ierror)
47     TYPE(MPI_Comm), INTENT(IN) :: comm
48     LOGICAL, INTENT(IN) :: remain_dims(*)
49     TYPE(MPI_Comm), INTENT(OUT) :: newcomm
50     INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

```

MPI_Cartdim_get(comm, ndims, ierror)
  TYPE(MPI_Comm), INTENT(IN) :: comm
  INTEGER, INTENT(OUT) :: ndims
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Dims_create(nnodes, ndims, dims, ierror)
  INTEGER, INTENT(IN) :: nnodes, ndims
  INTEGER, INTENT(INOUT) :: dims(*)
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Dist_graph_create(comm_old, n, sources, degrees, destinations, weights,
  info, reorder, comm_dist_graph, ierror)
  TYPE(MPI_Comm), INTENT(IN) :: comm_old
  INTEGER, INTENT(IN) :: n, sources(*), degrees(*), destinations(*)
  INTEGER, INTENT(IN) :: weights(*) ! optional by overloading
  TYPE(MPI_Info), INTENT(IN) :: info
  LOGICAL, INTENT(IN) :: reorder
  TYPE(MPI_Comm), INTENT(OUT) :: comm_dist_graph
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Dist_graph_create_adjacent(comm_old, indegree, sources, sourceweights,
  outdegree, destinations, destweights, info, reorder,
  comm_dist_graph, ierror)
  TYPE(MPI_Comm), INTENT(IN) :: comm_old
  INTEGER, INTENT(IN) :: indegree, sources(*), outdegree,
  destinations(*)
  INTEGER, INTENT(IN) :: sourceweights(*), destweights(*) ! optional by
  overloading
  TYPE(MPI_Info), INTENT(IN) :: info
  LOGICAL, INTENT(IN) :: reorder
  TYPE(MPI_Comm), INTENT(OUT) :: comm_dist_graph
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Dist_graph_neighbors(comm, maxindegree, sources, sourceweights,
  maxoutdegree, destinations, destweights, ierror)
  TYPE(MPI_Comm), INTENT(IN) :: comm
  INTEGER, INTENT(IN) :: maxindegree, maxoutdegree
  INTEGER, INTENT(OUT) :: sources(*), destinations(*)
  INTEGER :: sourceweights(*), destweights(*) ! optional by overloading
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Dist_graph_neighbors_count(comm, indegree, outdegree, weighted, ierror)
  TYPE(MPI_Comm), INTENT(IN) :: comm
  INTEGER, INTENT(OUT) :: indegree, outdegree
  LOGICAL, INTENT(OUT) :: weighted
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Graph_create(comm_old, nnodes, index, edges, reorder, comm_graph,
  ierror)
  TYPE(MPI_Comm), INTENT(IN) :: comm_old

```

```

1      INTEGER, INTENT(IN) :: nnodes, index(*), edges(*)
2      LOGICAL, INTENT(IN) :: reorder
3      TYPE(MPI_Comm), INTENT(OUT) :: comm_graph
4      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
5
6      MPI_Graph_get(comm, maxindex, maxedges, index, edges, ierror)
7      TYPE(MPI_Comm), INTENT(IN) :: comm
8      INTEGER, INTENT(IN) :: maxindex, maxedges
9      INTEGER, INTENT(OUT) :: index(*), edges(*)
10     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
11
12     MPI_Graph_map(comm, nnodes, index, edges, newrank, ierror)
13     TYPE(MPI_Comm), INTENT(IN) :: comm
14     INTEGER, INTENT(IN) :: nnodes, index(*), edges(*)
15     INTEGER, INTENT(OUT) :: newrank
16     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
17
18     MPI_Graph_neighbors(comm, rank, maxneighbors, neighbors, ierror)
19     TYPE(MPI_Comm), INTENT(IN) :: comm
20     INTEGER, INTENT(IN) :: rank, maxneighbors
21     INTEGER, INTENT(OUT) :: neighbors(*)
22     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
23
24     MPI_Graph_neighbors_count(comm, rank, nneighbors, ierror)
25     TYPE(MPI_Comm), INTENT(IN) :: comm
26     INTEGER, INTENT(IN) :: rank
27     INTEGER, INTENT(OUT) :: nneighbors
28     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
29
30     MPI_Graphdims_get(comm, nnodes, nedges, ierror)
31     TYPE(MPI_Comm), INTENT(IN) :: comm
32     INTEGER, INTENT(OUT) :: nnodes, nedges
33     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
34
35     MPI_Topo_test(comm, status, ierror)
36     TYPE(MPI_Comm), INTENT(IN) :: comm
37     INTEGER, INTENT(OUT) :: status
38     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
39
40
41

```

A.4.6 MPI Environmental Management Fortran 2008 Bindings

```

39     DOUBLE PRECISION MPI_Wtick()
40
41     DOUBLE PRECISION MPI_Wtime()
42
43     MPI_Abort(comm, errorcode, ierror)
44     TYPE(MPI_Comm), INTENT(IN) :: comm
45     INTEGER, INTENT(IN) :: errorcode
46     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
47
48     MPI_Add_error_class(errorclass, ierror)
49     INTEGER, INTENT(OUT) :: errorclass

```

```

    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
1
2
MPI_Add_error_code(errorclass, errorcode, ierror)
3
    INTEGER, INTENT(IN) :: errorclass
4
    INTEGER, INTENT(OUT) :: errorcode
5
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
6
MPI_Add_error_string(errorcode, string, ierror)
7
    INTEGER, INTENT(IN) :: errorcode
8
    CHARACTER(LEN=*), INTENT(IN) :: string
9
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
10
11
MPI_Alloc_mem(size, info, baseptr, ierror)
12
    USE, INTRINSIC :: ISO_C_BINDING
13
    INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: size
14
    TYPE(MPI_Info), INTENT(IN) :: info
15
    TYPE(C_PTR), INTENT(OUT) :: baseptr
16
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
17
MPI_Comm_call_errhandler(comm, errorcode, ierror)
18
    TYPE(MPI_Comm), INTENT(IN) :: comm
19
    INTEGER, INTENT(IN) :: errorcode
20
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
21
22
MPI_Comm_create_errhandler(comm_errhandler_fn, errhandler, ierror)
23
    EXTERNAL :: comm_errhandler_fn
24
    TYPE(MPI_Errhandler), INTENT(OUT) :: errhandler
25
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
26
MPI_Comm_get_errhandler(comm, errhandler, ierror)
27
    TYPE(MPI_Comm), INTENT(IN) :: comm
28
    TYPE(MPI_Errhandler), INTENT(OUT) :: errhandler
29
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
30
31
MPI_Comm_set_errhandler(comm, errhandler, ierror)
32
    TYPE(MPI_Comm), INTENT(IN) :: comm
33
    TYPE(MPI_Errhandler), INTENT(IN) :: errhandler
34
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
35
36
MPI_Errhandler_free(errhandler, ierror)
37
    TYPE(MPI_Errhandler), INTENT(INOUT) :: errhandler
38
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
39
MPI_Error_class(errorcode, errorclass, ierror)
40
    INTEGER, INTENT(IN) :: errorcode
41
    INTEGER, INTENT(OUT) :: errorclass
42
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
43
44
MPI_Error_string(errorcode, string, resultlen, ierror)
45
    INTEGER, INTENT(IN) :: errorcode
46
    CHARACTER(LEN=*), INTENT(OUT) :: string
47
    INTEGER, INTENT(OUT) :: resultlen
48

```



```

1      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
2
3      MPI_File_call_errhandler(fh, errorcode, ierror)
4      TYPE(MPI_File), INTENT(IN) :: fh
5      INTEGER, INTENT(IN) :: errorcode
6      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
7
8      MPI_File_create_errhandler(file_errhandler_fn, errhandler, ierror)
9      EXTERNAL :: file_errhandler_fn
10     TYPE(MPI_Errhandler), INTENT(OUT) :: errhandler
11     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
12
13     MPI_File_get_errhandler(file, errhandler, ierror)
14     TYPE(MPI_File), INTENT(IN) :: file
15     TYPE(MPI_Errhandler), INTENT(OUT) :: errhandler
16     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
17
18     MPI_File_set_errhandler(file, errhandler, ierror)
19     TYPE(MPI_File), INTENT(IN) :: file
20     TYPE(MPI_Errhandler), INTENT(IN) :: errhandler
21     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
22
23     MPI_Finalize(ierr)
24     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
25
26     MPI_Finalized(flag, ierror)
27     LOGICAL, INTENT(OUT) :: flag
28     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
29
30     MPI_Free_mem(base, ierror)
31     USE, INTRINSIC :: ISO_C_BINDING
32     TYPE(C_PTR), INTENT(IN) :: base
33     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
34
35     MPI_Get_processor_name( name, resultlen, ierror)
36     CHARACTER(LEN=*), INTENT(OUT) :: name
37     INTEGER, INTENT(OUT) :: resultlen
38     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
39
40     MPI_Get_version(version, subversion, ierror)
41     INTEGER, INTENT(OUT) :: version, subversion
42     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
43
44     MPI_Init(ierr)
45     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
46
47     MPI_Initialized(flag, ierror)
48     LOGICAL, INTENT(OUT) :: flag
49     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
50
51     MPI_Win_call_errhandler(win, errorcode, ierror)
52     TYPE(MPI_Win), INTENT(IN) :: win
53     INTEGER, INTENT(IN) :: errorcode

```

```

    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
MPI_Win_create_errhandler(win_errhandler_fn, errhandler, ierror)
    EXTERNAL :: win_errhandler_fn
    TYPE(MPI_Errhandler), INTENT(OUT) :: errhandler
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
MPI_Win_get_errhandler(win, errhandler, ierror)
    TYPE(MPI_Win), INTENT(IN) :: win
    TYPE(MPI_Errhandler), INTENT(OUT) :: errhandler
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
MPI_Win_set_errhandler(win, errhandler, ierror)
    TYPE(MPI_Win), INTENT(IN) :: win
    TYPE(MPI_Errhandler), INTENT(IN) :: errhandler
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

A.4.7 The Info Object Fortran 2008 Bindings

```

MPI_Info_create(info, ierror)
    TYPE(MPI_Info), INTENT(OUT) :: info
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
MPI_Info_delete(info, key, ierror)
    TYPE(MPI_Info), INTENT(IN) :: info
    CHARACTER(LEN=*), INTENT(IN) :: key
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
MPI_Info_dup(info, newinfo, ierror)
    TYPE(MPI_Info), INTENT(IN) :: info
    TYPE(MPI_Info), INTENT(OUT) :: newinfo
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
MPI_Info_free(info, ierror)
    TYPE(MPI_Info), INTENT(INOUT) :: info
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
MPI_Info_get(info, key, valuelen, value, flag, ierror)
    TYPE(MPI_Info), INTENT(IN) :: info
    CHARACTER(LEN=*), INTENT(IN) :: key
    INTEGER, INTENT(IN) :: valuelen
    CHARACTER(LEN=*), INTENT(OUT) :: value
    LOGICAL, INTENT(OUT) :: flag
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
MPI_Info_get_nkeys(info, nkeys, ierror)
    TYPE(MPI_Info), INTENT(IN) :: info
    INTEGER, INTENT(OUT) :: nkeys
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
MPI_Info_get_nthkey(info, n, key, ierror)
    TYPE(MPI_Info), INTENT(IN) :: info

```

```

1      INTEGER, INTENT(IN) :: n
2      CHARACTER(LEN=*), INTENT(OUT) :: key
3      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
4
5      MPI_Info_get_valuelen(info, key, valuelen, flag, ierror)
6      TYPE(MPI_Info), INTENT(IN) :: info
7      CHARACTER(LEN=*), INTENT(IN) :: key
8      INTEGER, INTENT(OUT) :: valuelen
9      LOGICAL, INTENT(OUT) :: flag
10     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
11
12     MPI_Info_set(info, key, value, ierror)
13     TYPE(MPI_Info), INTENT(IN) :: info
14     CHARACTER(LEN=*), INTENT(IN) :: key, value
15     INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

A.4.8 Process Creation and Management Fortran 2008 Bindings

```

18     MPI_Close_port(port_name, ierror)
19     CHARACTER(LEN=*), INTENT(IN) :: port_name
20     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
21
22     MPI_Comm_accept(port_name, info, root, comm, newcomm, ierror)
23     CHARACTER(LEN=*), INTENT(IN) :: port_name
24     TYPE(MPI_Info), INTENT(IN) :: info
25     INTEGER, INTENT(IN) :: root
26     TYPE(MPI_Comm), INTENT(IN) :: comm
27     TYPE(MPI_Comm), INTENT(OUT) :: newcomm
28     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
29
30     MPI_Comm_connect(port_name, info, root, comm, newcomm, ierror)
31     CHARACTER(LEN=*), INTENT(IN) :: port_name
32     TYPE(MPI_Info), INTENT(IN) :: info
33     INTEGER, INTENT(IN) :: root
34     TYPE(MPI_Comm), INTENT(IN) :: comm
35     TYPE(MPI_Comm), INTENT(OUT) :: newcomm
36     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
37
38     MPI_Comm_disconnect(comm, ierror)
39     TYPE(MPI_Comm), INTENT(INOUT) :: comm
40     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
41
42     MPI_Comm_get_parent(parent, ierror)
43     TYPE(MPI_Comm), INTENT(OUT) :: parent
44     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
45
46     MPI_Comm_join(fd, intercomm, ierror)
47     INTEGER, INTENT(IN) :: fd
48     TYPE(MPI_Comm), INTENT(OUT) :: intercomm
49     INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

```

MPI_Comm_spawn(command, argv, maxprocs, info, root, comm, intercomm,
               array_of_errcodes, ierror)
CHARACTER(LEN=*), INTENT(IN) :: command, argv(*)
INTEGER, INTENT(IN) :: maxprocs, root
TYPE(MPI_Info), INTENT(IN) :: info
TYPE(MPI_Comm), INTENT(IN) :: comm
TYPE(MPI_Comm), INTENT(OUT) :: intercomm
INTEGER, INTENT(OUT) :: array_of_errcodes(*) ! optional by
overloading
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Comm_spawn_multiple(count, array_of_commands, array_of_argv,
                       array_of_maxprocs, array_of_info, root, comm, intercomm,
                       array_of_errcodes, ierror)
INTEGER, INTENT(IN) :: count, array_of_maxprocs(*), root
CHARACTER(LEN=*), INTENT(IN) :: array_of_commands(*),
array_of_argv(count, *)
TYPE(MPI_Info), INTENT(IN) :: array_of_info(*)
TYPE(MPI_Comm), INTENT(IN) :: comm
TYPE(MPI_Comm), INTENT(OUT) :: intercomm
INTEGER, INTENT(OUT) :: array_of_errcodes(*) ! optional by
overloading
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Lookup_name(service_name, info, port_name, ierror)
CHARACTER(LEN=*), INTENT(IN) :: service_name
TYPE(MPI_Info), INTENT(IN) :: info
CHARACTER(LEN=*), INTENT(OUT) :: port_name
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Open_port(info, port_name, ierror)
TYPE(MPI_Info), INTENT(IN) :: info
CHARACTER(LEN=*), INTENT(OUT) :: port_name
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Publish_name(service_name, info, port_name, ierror)
TYPE(MPI_Info), INTENT(IN) :: info
CHARACTER(LEN=*), INTENT(IN) :: service_name, port_name
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Unpublish_name(service_name, info, port_name, ierror)
CHARACTER(LEN=*), INTENT(IN) :: service_name, port_name
TYPE(MPI_Info), INTENT(IN) :: info
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

A.4.9 One-Sided Communications Fortran 2008 Bindings

MPI_Accumulate(origin_addr, origin_count, origin_datatype, target_rank,
               target_disp, target_count, target_datatype, op, win, ierror)
TYPE(*), DIMENSION(..) :: origin_addr

```

```

1      INTEGER, INTENT(IN) :: origin_count, target_rank, target_count
2      TYPE(MPI_Datatype), INTENT(IN) :: origin_datatype
3      INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: target_disp
4      TYPE(MPI_Datatype), INTENT(IN) :: target_datatype
5      TYPE(MPI_Op), INTENT(IN) :: op
6      TYPE(MPI_Win), INTENT(IN) :: win
7      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
8
9      MPI_Get(origin_addr, origin_count, origin_datatype, target_rank,
10             target_disp, target_count, target_datatype, win, ierror)
11      TYPE(*), DIMENSION(..) :: origin_addr
12      INTEGER, INTENT(IN) :: origin_count, target_rank, target_count
13      TYPE(MPI_Datatype), INTENT(IN) :: origin_datatype
14      INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: target_disp
15      TYPE(MPI_Datatype), INTENT(IN) :: target_datatype
16      TYPE(MPI_Win), INTENT(IN) :: win
17      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
18
19      MPI_Put(origin_addr, origin_count, origin_datatype, target_rank,
20             target_disp, target_count, target_datatype, win, ierror)
21      TYPE(*), DIMENSION(..) :: origin_addr
22      INTEGER, INTENT(IN) :: origin_count, target_rank, target_count
23      TYPE(MPI_Datatype), INTENT(IN) :: origin_datatype
24      INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: target_disp
25      TYPE(MPI_Datatype), INTENT(IN) :: target_datatype
26      TYPE(MPI_Win), INTENT(IN) :: win
27      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
28
29      MPI_Win_complete(win, ierror)
30      TYPE(MPI_Win), INTENT(IN) :: win
31      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
32
33      MPI_Win_create(base, size, disp_unit, info, comm, win, ierror)
34      TYPE(*), DIMENSION(..) :: base
35      INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: size
36      INTEGER, INTENT(IN) :: disp_unit
37      TYPE(MPI_Info), INTENT(IN) :: info
38      TYPE(MPI_Comm), INTENT(IN) :: comm
39      TYPE(MPI_Win), INTENT(OUT) :: win
40      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
41
42      MPI_Win_fence(assert, win, ierror)
43      INTEGER, INTENT(IN) :: assert
44      TYPE(MPI_Win), INTENT(IN) :: win
45      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
46
47      MPI_Win_free(win, ierror)
48      TYPE(MPI_Win), INTENT(INOUT) :: win
49      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
50
51      MPI_Win_get_group(win, group, ierror)

```

```

TYPE(MPI_Win), INTENT(IN) :: win
TYPE(MPI_Group), INTENT(OUT) :: group
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Win_lock(lock_type, rank, assert, win, ierror)
  INTEGER, INTENT(IN) :: lock_type, rank, assert
  TYPE(MPI_Win), INTENT(IN) :: win
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Win_post(group, assert, win, ierror)
  TYPE(MPI_Group), INTENT(IN) :: group
  INTEGER, INTENT(IN) :: assert
  TYPE(MPI_Win), INTENT(IN) :: win
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Win_start(group, assert, win, ierror)
  TYPE(MPI_Group), INTENT(IN) :: group
  INTEGER, INTENT(IN) :: assert
  TYPE(MPI_Win), INTENT(IN) :: win
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Win_test(win, flag, ierror)
  LOGICAL, INTENT(OUT) :: flag
  TYPE(MPI_Win), INTENT(IN) :: win
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Win_unlock(rank, win, ierror)
  INTEGER, INTENT(IN) :: rank
  TYPE(MPI_Win), INTENT(IN) :: win
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Win_wait(win, ierror)
  TYPE(MPI_Win), INTENT(IN) :: win
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

A.4.10 External Interfaces Fortran 2008 Bindings

```

MPI_Grequest_complete(request, ierror)
  TYPE(MPI_Request), INTENT(IN) :: request
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Grequest_start(query_fn, free_fn, cancel_fn, extra_state, request,
  ierror)
  EXTERNAL :: query_fn, free_fn, cancel_fn
  INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: extra_state
  TYPE(MPI_Request), INTENT(OUT) :: request
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Init_thread(required, provided, ierror)
  INTEGER, INTENT(IN) :: required
  INTEGER, INTENT(OUT) :: provided

```

```

1      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
2
3      MPI_Is_thread_main(flag, ierror)
4      LOGICAL, INTENT(OUT) :: flag
5      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
6
7      MPI_Query_thread(provided, ierror)
8      INTEGER, INTENT(OUT) :: provided
9      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
10
11     MPI_Status_set_cancelled(status, flag, ierror)
12     TYPE(MPI_Status), INTENT(INOUT) :: status
13     LOGICAL, INTENT(OUT) :: flag
14     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
15
16     MPI_Status_set_elements(status, datatype, count, ierror)
17     TYPE(MPI_Status), INTENT(INOUT) :: status
18     TYPE(MPI_Datatype), INTENT(IN) :: datatype
19     INTEGER, INTENT(IN) :: count
20     INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

A.4.11 I/O Fortran 2008 Bindings

```

22     MPI_File_close(fh, ierror)
23     TYPE(MPI_File), INTENT(INOUT) :: fh
24     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
25
26     MPI_File_delete(filename, info, ierror)
27     CHARACTER(LEN=*), INTENT(IN) :: filename
28     TYPE(MPI_Info), INTENT(IN) :: info
29     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
30
31     MPI_File_get_amode(fh, amode, ierror)
32     TYPE(MPI_File), INTENT(IN) :: fh
33     INTEGER, INTENT(OUT) :: amode
34     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
35
36     MPI_File_get_atomicity(fh, flag, ierror)
37     TYPE(MPI_File), INTENT(IN) :: fh
38     LOGICAL, INTENT(OUT) :: flag
39     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
40
41     MPI_File_get_byte_offset(fh, offset, disp, ierror)
42     TYPE(MPI_File), INTENT(IN) :: fh
43     INTEGER(KIND=MPI_OFFSET_KIND), INTENT(IN) :: offset
44     INTEGER(KIND=MPI_OFFSET_KIND), INTENT(OUT) :: disp
45     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
46
47     MPI_File_get_group(fh, group, ierror)
48     TYPE(MPI_File), INTENT(IN) :: fh
49     TYPE(MPI_Group), INTENT(OUT) :: group
50     INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

```

MPI_File_get_info(fh, info_used, ierror)
  TYPE(MPI_File), INTENT(IN) :: fh
  TYPE(MPI_Info), INTENT(OUT) :: info_used
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_get_position(fh, offset, ierror)
  TYPE(MPI_File), INTENT(IN) :: fh
  INTEGER(KIND=MPI_OFFSET_KIND), INTENT(OUT) :: offset
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_get_position_shared(fh, offset, ierror)
  TYPE(MPI_File), INTENT(IN) :: fh
  INTEGER(KIND=MPI_OFFSET_KIND), INTENT(OUT) :: offset
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_get_size(fh, size, ierror)
  TYPE(MPI_File), INTENT(IN) :: fh
  INTEGER(KIND=MPI_OFFSET_KIND), INTENT(OUT) :: size
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_get_type_extent(fh, datatype, extent, ierror)
  TYPE(MPI_File), INTENT(IN) :: fh
  TYPE(MPI_Datatype), INTENT(IN) :: datatype
  INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(OUT) :: extent
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_get_view(fh, disp, etype, filetype, datarep, ierror)
  TYPE(MPI_File), INTENT(IN) :: fh
  INTEGER(KIND=MPI_OFFSET_KIND), INTENT(OUT) :: disp
  TYPE(MPI_Datatype), INTENT(OUT) :: etype
  TYPE(MPI_Datatype), INTENT(OUT) :: filetype
  CHARACTER(LEN=*), INTENT(OUT) :: datarep
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_iread(fh, buf, count, datatype, request, ierror)
  TYPE(MPI_File), INTENT(IN) :: fh
  TYPE(*), DIMENSION(..) :: buf
  INTEGER, INTENT(IN) :: count
  TYPE(MPI_Datatype), INTENT(IN) :: datatype
  TYPE(MPI_Request), INTENT(OUT) :: request
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_iread_at(fh, offset, buf, count, datatype, request, ierror)
  TYPE(MPI_File), INTENT(IN) :: fh
  INTEGER(KIND=MPI_OFFSET_KIND), INTENT(IN) :: offset
  TYPE(*), DIMENSION(..) :: buf
  INTEGER, INTENT(IN) :: count
  TYPE(MPI_Datatype), INTENT(IN) :: datatype
  TYPE(MPI_Request), INTENT(OUT) :: request
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```



```

1  MPI_File_iread_shared(fh, buf, count, datatype, request, ierror)
2      TYPE(MPI_File), INTENT(IN) :: fh
3      TYPE(*), DIMENSION(..) :: buf
4      INTEGER, INTENT(IN) :: count
5      TYPE(MPI_Datatype), INTENT(IN) :: datatype
6      TYPE(MPI_Request), INTENT(OUT) :: request
7      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
8
9  MPI_File_fwrite(fh, buf, count, datatype, request, ierror)
10     TYPE(MPI_File), INTENT(IN) :: fh
11     TYPE(*), DIMENSION(..) :: buf
12     INTEGER, INTENT(IN) :: count
13     TYPE(MPI_Datatype), INTENT(IN) :: datatype
14     TYPE(MPI_Request), INTENT(OUT) :: request
15     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
16
17 MPI_File_fwrite_at(fh, offset, buf, count, datatype, request, ierror)
18     TYPE(MPI_File), INTENT(IN) :: fh
19     INTEGER(KIND=MPI_OFFSET_KIND), INTENT(IN) :: offset
20     TYPE(*), DIMENSION(..) :: buf
21     INTEGER, INTENT(IN) :: count
22     TYPE(MPI_Datatype), INTENT(IN) :: datatype
23     TYPE(MPI_Request), INTENT(OUT) :: request
24     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
25
26 MPI_File_fwrite_shared(fh, buf, count, datatype, request, ierror)
27     TYPE(*), DIMENSION(..) :: buf
28     TYPE(MPI_File), INTENT(IN) :: fh
29     INTEGER, INTENT(IN) :: count
30     TYPE(MPI_Datatype), INTENT(IN) :: datatype
31     TYPE(MPI_Request), INTENT(OUT) :: request
32     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
33
34 MPI_File_open(comm, filename, amode, info, fh, ierror)
35     TYPE(MPI_Comm), INTENT(IN) :: comm
36     CHARACTER(LEN=*), INTENT(IN) :: filename
37     INTEGER, INTENT(IN) :: amode
38     TYPE(MPI_Info), INTENT(IN) :: info
39     TYPE(MPI_File), INTENT(OUT) :: fh
40     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
41
42 MPI_File_preallocate(fh, size, ierror)
43     TYPE(MPI_File), INTENT(IN) :: fh
44     INTEGER(KIND=MPI_OFFSET_KIND), INTENT(IN) :: size
45     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
46
47 MPI_File_read(fh, buf, count, datatype, status, ierror)
48     TYPE(MPI_File), INTENT(IN) :: fh
49     TYPE(*), DIMENSION(..) :: buf
50     INTEGER, INTENT(IN) :: count

```

```

TYPE(MPI_Datatype), INTENT(IN) :: datatype
TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_read_all(fh, buf, count, datatype, status, ierror)
TYPE(MPI_File), INTENT(IN) :: fh
TYPE(*), DIMENSION(..) :: buf
INTEGER, INTENT(IN) :: count
TYPE(MPI_Datatype), INTENT(IN) :: datatype
TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_read_all_begin(fh, buf, count, datatype, ierror)
TYPE(MPI_File), INTENT(IN) :: fh
TYPE(*), DIMENSION(..) :: buf
INTEGER, INTENT(IN) :: count
TYPE(MPI_Datatype), INTENT(IN) :: datatype
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_read_all_end(fh, buf, status, ierror)
TYPE(MPI_File), INTENT(IN) :: fh
TYPE(*), DIMENSION(..) :: buf
TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_read_at(fh, offset, buf, count, datatype, status, ierror)
TYPE(MPI_File), INTENT(IN) :: fh
INTEGER(KIND=MPI_OFFSET_KIND), INTENT(IN) :: offset
TYPE(*), DIMENSION(..) :: buf
INTEGER, INTENT(IN) :: count
TYPE(MPI_Datatype), INTENT(IN) :: datatype
TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_read_at_all(fh, offset, buf, count, datatype, status, ierror)
TYPE(MPI_File), INTENT(IN) :: fh
INTEGER(KIND=MPI_OFFSET_KIND), INTENT(IN) :: offset
TYPE(*), DIMENSION(..) :: buf
INTEGER, INTENT(IN) :: count
TYPE(MPI_Datatype), INTENT(IN) :: datatype
TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_read_at_all_begin(fh, offset, buf, count, datatype, ierror)
TYPE(MPI_File), INTENT(IN) :: fh
INTEGER(KIND=MPI_OFFSET_KIND), INTENT(IN) :: offset
TYPE(*), DIMENSION(..) :: buf
INTEGER, INTENT(IN) :: count
TYPE(MPI_Datatype), INTENT(IN) :: datatype
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

```

1  MPI_File_read_at_all_end(fh, buf, status, ierror)
2      TYPE(MPI_File), INTENT(IN) :: fh
3      TYPE(*), DIMENSION(..) :: buf
4      TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
5      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
6
7  MPI_File_read_ordered(fh, buf, count, datatype, status, ierror)
8      TYPE(MPI_File), INTENT(IN) :: fh
9      TYPE(*), DIMENSION(..) :: buf
10     INTEGER, INTENT(IN) :: count
11     TYPE(MPI_Datatype), INTENT(IN) :: datatype
12     TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
13     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
14
15 MPI_File_read_ordered_begin(fh, buf, count, datatype, ierror)
16     TYPE(MPI_File), INTENT(IN) :: fh
17     TYPE(*), DIMENSION(..) :: buf
18     INTEGER, INTENT(IN) :: count
19     TYPE(MPI_Datatype), INTENT(IN) :: datatype
20     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
21
22 MPI_File_read_ordered_end(fh, buf, status, ierror)
23     TYPE(MPI_File), INTENT(IN) :: fh
24     TYPE(*), DIMENSION(..) :: buf
25     TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
26     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
27
28 MPI_File_read_shared(fh, buf, count, datatype, status, ierror)
29     TYPE(MPI_File), INTENT(IN) :: fh
30     TYPE(*), DIMENSION(..) :: buf
31     INTEGER, INTENT(IN) :: count
32     TYPE(MPI_Datatype), INTENT(IN) :: datatype
33     TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
34     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
35
36 MPI_File_seek(fh, offset, whence, ierror)
37     TYPE(MPI_File), INTENT(IN) :: fh
38     INTEGER(KIND=MPI_OFFSET_KIND), INTENT(IN) :: offset
39     INTEGER, INTENT(IN) :: whence
40     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
41
42 MPI_File_seek_shared(fh, offset, whence, ierror)
43     TYPE(MPI_File), INTENT(IN) :: fh
44     INTEGER(KIND=MPI_OFFSET_KIND), INTENT(IN) :: offset
45     INTEGER, INTENT(IN) :: whence
46     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
47
48 MPI_File_set_atomicity(fh, flag, ierror)
49     TYPE(MPI_File), INTENT(IN) :: fh
50     LOGICAL, INTENT(IN) :: flag
51     INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

```

MPI_File_set_info(fh, info, ierror)
  TYPE(MPI_File), INTENT(IN) :: fh
  TYPE(MPI_Info), INTENT(IN) :: info
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_set_size(fh, size, ierror)
  TYPE(MPI_File), INTENT(IN) :: fh
  INTEGER(KIND=MPI_OFFSET_KIND), INTENT(IN) :: size
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_set_view(fh, disp, etype, filetype, datarep, info, ierror)
  TYPE(MPI_File), INTENT(IN) :: fh
  INTEGER(KIND=MPI_OFFSET_KIND), INTENT(IN) :: disp
  TYPE(MPI_Datatype), INTENT(IN) :: etype
  TYPE(MPI_Datatype), INTENT(IN) :: filetype
  CHARACTER(LEN=*), INTENT(IN) :: datarep
  TYPE(MPI_Info), INTENT(IN) :: info
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_sync(fh, ierror)
  TYPE(MPI_File), INTENT(IN) :: fh
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_write(fh, buf, count, datatype, status, ierror)
  TYPE(MPI_File), INTENT(IN) :: fh
  TYPE(*), DIMENSION(..) :: buf
  INTEGER, INTENT(IN) :: count
  TYPE(MPI_Datatype), INTENT(IN) :: datatype
  TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_write_all(fh, buf, count, datatype, status, ierror)
  TYPE(MPI_File), INTENT(IN) :: fh
  TYPE(*), DIMENSION(..) :: buf
  INTEGER, INTENT(IN) :: count
  TYPE(MPI_Datatype), INTENT(IN) :: datatype
  TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_write_all_begin(fh, buf, count, datatype, ierror)
  TYPE(MPI_File), INTENT(IN) :: fh
  TYPE(*), DIMENSION(..) :: buf
  INTEGER, INTENT(IN) :: count
  TYPE(MPI_Datatype), INTENT(IN) :: datatype
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_write_all_end(fh, buf, status, ierror)
  TYPE(MPI_File), INTENT(IN) :: fh
  TYPE(*), DIMENSION(..) :: buf
  TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

```

1  MPI_File_write_at(fh, offset, buf, count, datatype, status, ierror)
2      TYPE(MPI_File), INTENT(IN) :: fh
3      INTEGER(KIND=MPI_OFFSET_KIND), INTENT(IN) :: offset
4      TYPE(*), DIMENSION(..) :: buf
5      INTEGER, INTENT(IN) :: count
6      TYPE(MPI_Datatype), INTENT(IN) :: datatype
7      TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
8      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
9
10 MPI_File_write_at_all(fh, offset, buf, count, datatype, status, ierror)
11     TYPE(MPI_File), INTENT(IN) :: fh
12     INTEGER(KIND=MPI_OFFSET_KIND), INTENT(IN) :: offset
13     TYPE(*), DIMENSION(..) :: buf
14     INTEGER, INTENT(IN) :: count
15     TYPE(MPI_Datatype), INTENT(IN) :: datatype
16     TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
17     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
18
19 MPI_File_write_at_all_begin(fh, offset, buf, count, datatype, ierror)
20     TYPE(MPI_File), INTENT(IN) :: fh
21     INTEGER(KIND=MPI_OFFSET_KIND), INTENT(IN) :: offset
22     TYPE(*), DIMENSION(..) :: buf
23     INTEGER, INTENT(IN) :: count
24     TYPE(MPI_Datatype), INTENT(IN) :: datatype
25     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
26
27 MPI_File_write_at_all_end(fh, buf, status, ierror)
28     TYPE(MPI_File), INTENT(IN) :: fh
29     TYPE(*), DIMENSION(..) :: buf
30     TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
31     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
32
33 MPI_File_write_ordered(fh, buf, count, datatype, status, ierror)
34     TYPE(MPI_File), INTENT(IN) :: fh
35     TYPE(*), DIMENSION(..) :: buf
36     INTEGER, INTENT(IN) :: count
37     TYPE(MPI_Datatype), INTENT(IN) :: datatype
38     TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
39     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
40
41 MPI_File_write_ordered_begin(fh, buf, count, datatype, ierror)
42     TYPE(MPI_File), INTENT(IN) :: fh
43     TYPE(*), DIMENSION(..) :: buf
44     INTEGER, INTENT(IN) :: count
45     TYPE(MPI_Datatype), INTENT(IN) :: datatype
46     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
47
48 MPI_File_write_ordered_end(fh, buf, status, ierror)
49     TYPE(MPI_File), INTENT(IN) :: fh
50     TYPE(*), DIMENSION(..) :: buf

```

```

TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_File_write_shared(fh, buf, count, datatype, status, ierror)
TYPE(MPI_File), INTENT(IN) :: fh
TYPE(*), DIMENSION(..) :: buf
INTEGER, INTENT(IN) :: count
TYPE(MPI_Datatype), INTENT(IN) :: datatype
TYPE(MPI_Status), INTENT(OUT) :: status ! optional by overloading
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Register_datarep(datarep, read_conversion_fn, write_conversion_fn,
dtype_file_extent_fn, extra_state, ierror)
CHARACTER(LEN=*), INTENT(IN) :: datarep
EXTERNAL :: read_conversion_fn, write_conversion_fn,
dtype_file_extent_fn
INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: extra_state
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

A.4.12 Language Bindings Fortran 2008 Bindings

```

MPI_F_sync_reg(buf)
TYPE(*), DIMENSION(..) :: buf

MPI_Sizeof(x, size, ierror)
TYPE(*) :: x
INTEGER, INTENT(OUT) :: size
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Status_f082f(f08_status, f_status, ierror)
TYPE(MPI_Status) :: f08_status
INTEGER :: f_status(MPI_STATUS_SIZE)
INTEGER, OPTIONAL :: ierror

MPI_Status_f2f08(f_status, f08_status, ierror)
INTEGER f_status(MPI_STATUS_SIZE)
TYPE(MPI_Status) :: f08_status
INTEGER, OPTIONAL :: ierror

MPI_Type_create_f90_complex(p, r, newtype, ierror)
INTEGER, INTENT(IN) :: p, r
TYPE(MPI_Datatype), INTENT(OUT) :: newtype
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Type_create_f90_integer(r, newtype, ierror)
INTEGER, INTENT(IN) :: r
TYPE(MPI_Datatype), INTENT(OUT) :: newtype
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_Type_create_f90_real(p, r, newtype, ierror)
INTEGER, INTENT(IN) :: p, r

```

```
1      TYPE(MPI_Datatype), INTENT(OUT) :: newtype
2      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
3
4      MPI_Type_match_size(typeclass, size, datatype, ierror)
5      INTEGER, INTENT(IN) :: typeclass, size
6      TYPE(MPI_Datatype), INTENT(OUT) :: datatype
7      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

A.4.13 Profiling Interface Fortran 2008 Bindings

```
11     MPI_Pcontrol(level)
12     INTEGER, INTENT(IN) :: level
```

A.4.14 Deprecated Fortran 2008 Bindings

```

1  {void Op::Init(User_function* user_fn, bool commute) (binding deprecated, see
2      Section 15.2) }
3
4  {Request Comm::Ireduce(const void* sendbuf, void* recvbuf, int count,
5      const Datatype& datatype, const Op& op, int root)
6      const = 0 (binding deprecated, see Section 15.2) }
7
8  {Request Comm::Ireduce_scatter(const void* sendbuf, void* recvbuf,
9      int recvcnts[], const Datatype& datatype, const Op& op)
10     const = 0 (binding deprecated, see Section 15.2) }
11
12 {Request Comm::Ireduce_scatter_block(const void* sendbuf, void* recvbuf,
13     int recvcnt, const Datatype& datatype, const Op& op)
14     const = 0 (binding deprecated, see Section 15.2) }
15
16 {bool Op::Is_commutative() const (binding deprecated, see Section 15.2) }
17
18 {Request Intracomm::Isca(const void* sendbuf, void* recvbuf, int count,
19     const Datatype& datatype, const Op& op) const (binding deprecated,
20     see Section 15.2) }
21
22 {Request Comm::Isca(const void* sendbuf, int sendcount, const
23     Datatype& sendtype, void* recvbuf, int recvcnt,
24     const Datatype& recvtype, int root) const = 0 (binding deprecated,
25     see Section 15.2) }
26
27 {Request Comm::Iscaerv(const void* sendbuf, const int sendcounts[],
28     const int displs[], const Datatype& sendtype, void* recvbuf,
29     int recvcnt, const Datatype& recvtype, int root)
30     const = 0 (binding deprecated, see Section 15.2) }
31
32 {void Comm::Reduce(const void* sendbuf, void* recvbuf, int count,
33     const Datatype& datatype, const Op& op, int root)
34     const = 0 (binding deprecated, see Section 15.2) }
35
36 {void Op::Reduce_local(const void* inbuf, void* inoutbuf, int count,
37     const Datatype& datatype) const (binding deprecated, see
38     Section 15.2) }
39
40 {void Comm::Reduce_scatter(const void* sendbuf, void* recvbuf,
41     int recvcnts[], const Datatype& datatype, const Op& op)
42     const = 0 (binding deprecated, see Section 15.2) }
43
44 {void Comm::Reduce_scatter_block(const void* sendbuf, void* recvbuf,
45     int recvcnt, const Datatype& datatype, const Op& op)
46     const = 0 (binding deprecated, see Section 15.2) }
47
48 {void Intracomm::Scan(const void* sendbuf, void* recvbuf, int count,
49     const Datatype& datatype, const Op& op) const (binding deprecated,
50     see Section 15.2) }
51
52 {void Comm::Scatter(const void* sendbuf, int sendcount, const
53     Datatype& sendtype, void* recvbuf, int recvcnt,
54     const Datatype& recvtype, int root) const = 0 (binding deprecated,

```



```

1  {static Errhandler Comm::Create_errhandler(Comm::Errhandler_function*
2      comm_errhandler_fn) (binding deprecated, see Section 15.2) }
3
4  {static Errhandler File::Create_errhandler(File::Errhandler_function*
5      file_errhandler_fn) (binding deprecated, see Section 15.2) }
6
7  {static Errhandler Win::Create_errhandler(Win::Errhandler_function*
8      win_errhandler_fn) (binding deprecated, see Section 15.2) }
9
10 {void Finalize() (binding deprecated, see Section 15.2) }
11
12 {void Errhandler::Free() (binding deprecated, see Section 15.2) }
13
14 {void Free_mem(void *base) (binding deprecated, see Section 15.2) }
15
16 {Errhandler Comm::Get_errhandler() const (binding deprecated, see Section 15.2) }
17
18 {Errhandler File::Get_errhandler() const (binding deprecated, see Section 15.2) }
19
20 {Errhandler Win::Get_errhandler() const (binding deprecated, see Section 15.2) }
21
22 {int Get_error_class(int errorcode) (binding deprecated, see Section 15.2) }
23
24 {void Get_error_string(int errorcode, char* name, int& resultlen) (binding
25     deprecated, see Section 15.2) }
26
27 {void Get_processor_name(char* name, int& resultlen) (binding deprecated, see
28     Section 15.2) }
29
30 {void Get_version(int& version, int& subversion) (binding deprecated, see
31     Section 15.2) }
32
33 {void Init() (binding deprecated, see Section 15.2) }
34
35 {void Init(int& argc, char**& argv) (binding deprecated, see Section 15.2) }
36
37 {bool Is_finalized() (binding deprecated, see Section 15.2) }
38
39 {bool Is_initialized() (binding deprecated, see Section 15.2) }
40
41 {void Comm::Set_errhandler(const Errhandler& errhandler) (binding deprecated,
42     see Section 15.2) }
43
44 {void File::Set_errhandler(const Errhandler& errhandler) (binding deprecated,
45     see Section 15.2) }
46
47 {void Win::Set_errhandler(const Errhandler& errhandler) (binding deprecated,
48     see Section 15.2) }
49
50 {double Wtick() (binding deprecated, see Section 15.2) }
51
52 {double Wtime() (binding deprecated, see Section 15.2) }
53
54 };

```

A.5.7 The Info Object C++ Bindings

```
namespace MPI {
```