

# MPI: A Message-Passing Interface Standard

Version 3.0

ticket0.

(Draft, with MPI 3 Nonblocking Collectives)

Unofficial, for comment only

Message Passing Interface Forum

April 25, 2011

## Chapter 2

# MPI Terms and Conventions

This chapter explains notational terms and conventions used throughout the MPI document, some of the choices that have been made, and the rationale behind those choices. It is similar to the MPI-1 Terms and Conventions chapter but differs in some major and minor ways. Some of the major areas of difference are the naming conventions, some semantic definitions, file objects, Fortran 90 *vs* Fortran 77, C++, processes, and interaction with signals.

### 2.1 Document Notation

*Rationale.* Throughout this document, the rationale for the design choices made in the interface specification is set off in this format. Some readers may wish to skip these sections, while readers interested in interface design may want to read them carefully. (*End of rationale.*)

*Advice to users.* Throughout this document, material aimed at users and that illustrates usage is set off in this format. Some readers may wish to skip these sections, while readers interested in programming in MPI may want to read them carefully. (*End of advice to users.*)

*Advice to implementors.* Throughout this document, material that is primarily commentary to implementors is set off in this format. Some readers may wish to skip these sections, while readers interested in MPI implementations may want to read them carefully. (*End of advice to implementors.*)

### 2.2 Naming Conventions

In many cases MPI names for C functions are of the form `MPI_Class_action_subset`. This convention originated with MPI-1. Since MPI-2 an attempt has been made to standardize the names of MPI functions according to the following rules. The C++ bindings in particular follow these rules (see Section 2.6.4 on page 19).

1. In C, all routines associated with a particular type of MPI object should be of the form `MPI_Class_action_subset` or, if no subset exists, of the form `MPI_Class_action`. In Fortran, all routines associated with a particular type of MPI object should be of the form `MPI_CLASS_ACTION_SUBSET` or, if no subset exists, of the form

MPI\_CLASS\_ACTION. For C and Fortran we use the C++ terminology to define the **Class**. In C++, the routine is a method on **Class** and is named `MPI::Class::Action_subset`. If the routine is associated with a certain class, but does not make sense as an object method, it is a static member function of the class.

2. If the routine is not associated with a class, the name should be of the form `MPI_Action_subset` in C and `MPI_ACTION_SUBSET` in Fortran, and in C++ should be scoped in the MPI namespace, `MPI::Action_subset`.
3. The names of certain actions have been standardized. In particular, **Create** creates a new object, **Get** retrieves information about an object, **Set** sets this information, **Delete** deletes information, **Is** asks whether or not an object has a certain property.

C and Fortran names for some MPI functions (that were defined during the MPI-1 process) violate these rules in several cases. The most common exceptions are the omission of the **Class** name from the routine and the omission of the **Action** where one can be inferred.

MPI identifiers are limited to 30 characters (31 with the profiling interface). This is done to avoid exceeding the limit on some compilation systems.

## 2.3 Procedure Specification

MPI procedures are specified using a language-independent notation. The arguments of procedure calls are marked as IN, OUT or INOUT. The meanings of these are:

- IN: the call may use the input value but does not update the argument,
- OUT: the call may update the argument but does not use its input value,
- INOUT: the call may both use and update the argument.

There is one special case — if an argument is a handle to an opaque object (these terms are defined in Section 2.5.1), and the object is updated by the procedure call, then the argument is marked INOUT or OUT. It is marked this way even though the handle itself is not modified — we use the INOUT or OUT attribute to denote that what the handle *references* is updated. Thus, in C++, IN arguments are usually either references or pointers to `const` objects.

*Rationale.* The definition of MPI tries to avoid, to the largest possible extent, the use of INOUT arguments, because such use is error-prone, especially for scalar arguments. (*End of rationale.*)

MPI's use of IN, OUT and INOUT is intended to indicate to the user how an argument is to be used, but does not provide a rigorous classification that can be translated directly into all language bindings (e.g., `INTENT` in Fortran 90 bindings or `const` in C bindings). For instance, the “constant” `MPI_BOTTOM` can usually be passed to OUT buffer arguments. Similarly, `MPI_STATUS_IGNORE` can be passed as the OUT status argument.

A common occurrence for MPI functions is an argument that is used as IN by some processes and OUT by other processes. Such an argument is, syntactically, an INOUT argument

and is marked as such, although, semantically, it is not used in one call both for input and for output on a single process.

Another frequent situation arises when an argument value is needed only by a subset of the processes. When an argument is not significant at a process then an arbitrary value can be passed as an argument.

Unless specified otherwise, an argument of type OUT or type INOUT cannot be aliased with any other argument passed to an MPI procedure. An example of argument aliasing in C appears below. If we define a C procedure like this,

```
void copyIntBuffer( int *pin, int *pout, int len )
{
    int i;
    for (i=0; i<len; ++i) *pout++ = *pin++;
}
```

then a call to it in the following code fragment has aliased arguments.

```
int a[10];
copyIntBuffer( a, a+3, 7);
```

Although the C language allows this, such usage of MPI procedures is forbidden unless otherwise specified. Note that Fortran prohibits aliasing of arguments.

All MPI functions are first specified in the language-independent notation. Immediately below this, the ISO C version of the function is shown followed by a version of the same function in Fortran and then the C++ binding. Fortran in this document refers to Fortran 90; see Section 2.6.

## 2.4 Semantic Terms

When discussing MPI procedures the following semantic terms are used.

**nonblocking** A procedure is nonblocking if the procedure may return before the operation completes, and before the user is allowed to reuse resources (such as buffers) specified in the call. A nonblocking request is **started** by the call that initiates it, e.g., MPI\_ISEND. The word complete is used with respect to operations, requests, and communications. An **operation completes** when the user is allowed to reuse resources, and any output buffers have been updated; i.e. a call to MPI\_TEST will return `flag = true`. A **request is completed** by a call to wait, which returns, or a test or get status call which returns `flag = true`. This completing call has two effects: the status is extracted from the request; in the case of test and wait, if the request was nonpersistent, it is **freed**, and becomes **inactive** if it was persistent. A **communication completes** when all participating operations complete.

**blocking** A procedure is blocking if return from the procedure indicates the user is allowed to reuse resources specified in the call.

**local** A procedure is local if completion of the procedure depends only on the local executing process.

**non-local** A procedure is non-local if completion of the operation may require the execution of some MPI procedure on another process. Such an operation may require communication occurring with another user process.

**collective** A procedure is collective if all processes in a process group need to invoke the procedure. A collective call may or may not be synchronizing. Collective calls over the same communicator must be executed in the same order by all members of the process group.

**predefined** A predefined datatype is a datatype with a predefined (constant) name (such as `MPI_INT`, `MPI_FLOAT_INT`, or `MPI_UB`) or a datatype constructed with `MPI_TYPE_CREATE_F90_INTEGER`, `MPI_TYPE_CREATE_F90_REAL`, or `MPI_TYPE_CREATE_F90_COMPLEX`. The former are **named** whereas the latter are **unnamed**.

**derived** A derived datatype is any datatype that is not predefined.

**portable** A datatype is portable, if it is a predefined datatype, or it is derived from a portable datatype using only the type constructors `MPI_TYPE_CONTIGUOUS`, `MPI_TYPE_VECTOR`, `MPI_TYPE_INDEXED`, `MPI_TYPE_CREATE_INDEXED_BLOCK`, `MPI_TYPE_CREATE_SUBARRAY`, `MPI_TYPE_DUP`, and `MPI_TYPE_CREATE_DARRAY`. Such a datatype is portable because all displacements in the datatype are in terms of extents of one predefined datatype. Therefore, if such a datatype fits a data layout in one memory, it will fit the corresponding data layout in another memory, if the same declarations were used, even if the two systems have different architectures. On the other hand, if a datatype was constructed using `MPI_TYPE_CREATE_HINDEXED`, `MPI_TYPE_CREATE_HVECTOR` or `MPI_TYPE_CREATE_STRUCT`, then the datatype contains explicit byte displacements (e.g., providing padding to meet alignment restrictions). These displacements are unlikely to be chosen correctly if they fit data layout on one memory, but are used for data layouts on another process, running on a processor with a different architecture.

**equivalent** Two datatypes are equivalent if they appear to have been created with the same sequence of calls (and arguments) and thus have the same typemap. Two equivalent datatypes do not necessarily have the same cached attributes or the same names.

## 2.5 Data Types

### 2.5.1 Opaque Objects

MPI manages **system memory** that is used for buffering messages and for storing internal representations of various MPI objects such as groups, communicators, datatypes, etc. This memory is not directly accessible to the user, and objects stored there are **opaque**: their size and shape is not visible to the user. Opaque objects are accessed via **handles**, which exist in user space. MPI procedures that operate on opaque objects are passed handle arguments to access these objects. In addition to their use by MPI calls for object access, handles can participate in assignments and comparisons.

In Fortran, all handles have type `INTEGER`. In C and C++, a different handle type is defined for each category of objects. In addition, handles themselves are distinct objects in C++. The C and C++ types must support the use of the assignment and equality operators.

*Advice to implementors.* In Fortran, the handle can be an index into a table of opaque objects in a system table; in C it can be such an index or a pointer to the object. C++ handles can simply “wrap up” a table index or pointer.

*(End of advice to implementors.)*

Opaque objects are allocated and deallocated by calls that are specific to each object type. These are listed in the sections where the objects are described. The calls accept a handle argument of matching type. In an allocate call this is an OUT argument that returns a valid reference to the object. In a call to deallocate this is an INOUT argument which returns with an “invalid handle” value. MPI provides an “invalid handle” constant for each object type. Comparisons to this constant are used to test for validity of the handle.

A call to a deallocate routine invalidates the handle and marks the object for deallocation. The object is not accessible to the user after the call. However, MPI need not deallocate the object immediately. Any operation pending (at the time of the deallocate) that involves this object will complete normally; the object will be deallocated afterwards.

An opaque object and its handle are significant only at the process where the object was created and cannot be transferred to another process.

MPI provides certain predefined opaque objects and predefined, static handles to these objects. The user must not free such objects. In C++, this is enforced by declaring the handles to these predefined objects to be `static const`.

*Rationale.* This design hides the internal representation used for MPI data structures, thus allowing similar calls in C, C++, and Fortran. It also avoids conflicts with the typing rules in these languages, and easily allows future extensions of functionality. The mechanism for opaque objects used here loosely follows the POSIX Fortran binding standard.

The explicit separation of handles in user space and objects in system space allows space-reclaiming and deallocation calls to be made at appropriate points in the user program. If the opaque objects were in user space, one would have to be very careful not to go out of scope before any pending operation requiring that object completed. The specified design allows an object to be marked for deallocation, the user program can then go out of scope, and the object itself still persists until any pending operations are complete.

The requirement that handles support assignment/comparison is made since such operations are common. This restricts the domain of possible implementations. The alternative would have been to allow handles to have been an arbitrary, opaque type. This would force the introduction of routines to do assignment and comparison, adding complexity, and was therefore ruled out. *(End of rationale.)*

*Advice to users.* A user may accidentally create a dangling reference by assigning to a handle the value of another handle, and then deallocating the object associated with these handles. Conversely, if a handle variable is deallocated before the associated object is freed, then the object becomes inaccessible (this may occur, for example, if the handle is a local variable within a subroutine, and the subroutine is exited before the associated object is deallocated). It is the user’s responsibility to avoid adding or deleting references to opaque objects, except as a result of MPI calls that allocate or deallocate such objects. *(End of advice to users.)*

*Advice to implementors.* The intended semantics of opaque objects is that opaque objects are separate from one another; each call to allocate such an object copies all the information required for the object. Implementations may avoid excessive copying by substituting referencing for copying. For example, a derived datatype may contain references to its components, rather than copies of its components; a call to `MPI_COMM_GROUP` may return a reference to the group associated with the communicator, rather than a copy of this group. In such cases, the implementation must maintain reference counts, and allocate and deallocate objects in such a way that the visible effect is as if the objects were copied. (*End of advice to implementors.*)

## 2.5.2 Array Arguments

An MPI call may need an argument that is an array of opaque objects, or an array of handles. The array-of-handles is a regular array with entries that are handles to objects of the same type in consecutive locations in the array. Whenever such an array is used, an additional `len` argument is required to indicate the number of valid entries (unless this number can be derived otherwise). The valid entries are at the beginning of the array; `len` indicates how many of them there are, and need not be the size of the entire array. The same approach is followed for other array arguments. In some cases `NULL` handles are considered valid entries. When a `NULL` argument is desired for an array of statuses, one uses `MPI_STATUSES_IGNORE`.

## 2.5.3 State

MPI procedures use at various places arguments with *state* types. The values of such a data type are all identified by names, and no operation is defined on them. For example, the `MPI_TYPE_CREATE_SUBARRAY` routine has a state argument `order` with values `MPI_ORDER_C` and `MPI_ORDER_FORTRAN`.

## 2.5.4 Named Constants

MPI procedures sometimes assign a special meaning to a special value of a basic type argument; e.g., `tag` is an integer-valued argument of point-to-point communication operations, with a special wild-card value, `MPI_ANY_TAG`. Such arguments will have a range of regular values, which is a proper subrange of the range of values of the corresponding basic type; special values (such as `MPI_ANY_TAG`) will be outside the regular range. The range of regular values, such as `tag`, can be queried using environmental inquiry functions (Chapter 7 of the MPI-1 document). The range of other values, such as `source`, depends on values given by other MPI routines (in the case of `source` it is the communicator size).

MPI also provides predefined named constant handles, such as `MPI_COMM_WORLD`.

All named constants, with the exceptions noted below for Fortran, can be used in initialization expressions or assignments, but not necessarily in array declarations or as labels in C/C++ `switch` or Fortran `select/case` statements. This implies named constants to be link-time but not necessarily compile-time constants. The named constants listed below are required to be compile-time constants in both C/C++ and Fortran. These constants do not change values during execution. Opaque objects accessed by constant handles are defined and do not change value between MPI initialization (`MPI_INIT`) and MPI completion (`MPI_FINALIZE`). The handles themselves are constants and can be also used in initialization expressions or assignments.

The constants that are required to be compile-time constants (and can thus be used for array length declarations and labels in C/C++ `switch` and Fortran `case/select` statements) are:

`MPI_MAX_PROCESSOR_NAME`

`MPI_MAX_ERROR_STRING`

`MPI_MAX_DATAREP_STRING`

`MPI_MAX_INFO_KEY`

`MPI_MAX_INFO_VAL`

`MPI_MAX_OBJECT_NAME`

`MPI_MAX_PORT_NAME`

`MPI_STATUS_SIZE` (Fortran only)

`MPI_ADDRESS_KIND` (Fortran only)

`MPI_INTEGER_KIND` (Fortran only)

`MPI_OFFSET_KIND` (Fortran only)

and their C++ counterparts where appropriate.

The constants that cannot be used in initialization expressions or assignments in Fortran are:

`MPI_BOTTOM`

`MPI_STATUS_IGNORE`

`MPI_STATUSES_IGNORE`

`MPI_ERRCODES_IGNORE`

`MPI_IN_PLACE`

`MPI_ARGV_NULL`

`MPI_ARGVS_NULL`

`MPI_UNWEIGHTED`

*Advice to implementors.* In Fortran the implementation of these special constants may require the use of language constructs that are outside the Fortran standard. Using special values for the constants (e.g., by defining them through `PARAMETER` statements) is not possible because an implementation cannot distinguish these values from legal data. Typically, these constants are implemented as predefined static variables (e.g., a variable in an MPI-declared `COMMON` block), relying on the fact that the target compiler passes data by address. Inside the subroutine, this address can be extracted by some mechanism outside the Fortran standard (e.g., by Fortran extensions or by implementing the function in C). (*End of advice to implementors.*)

### 2.5.5 Choice

MPI functions sometimes use arguments with a *choice* (or union) data type. Distinct calls to the same routine may pass by reference actual arguments of different types. The mechanism for providing such arguments will differ from language to language. For Fortran, the document uses `<type>` to represent a choice variable; for C and C++, we use `void *`.

### 2.5.6 Addresses

Some MPI procedures use *address* arguments that represent an absolute address in the calling program. The datatype of such an argument is `MPI_Aint` in C, `MPI::Aint` in C++ and `INTEGER (KIND=MPI_ADDRESS_KIND)` in Fortran. These types must have the same



width and encode address values in the same manner such that address values in one language may be passed directly to another language without conversion. There is the MPI constant MPI\_BOTTOM to indicate the start of the address range.

### 2.5.7 File Offsets

For I/O there is a need to give the size, displacement, and offset into a file. These quantities can easily be larger than 32 bits which can be the default size of a Fortran integer. To overcome this, these quantities are declared to be INTEGER (KIND=MPI\_OFFSET\_KIND) in Fortran. In C one uses MPI\_Offset whereas in C++ one uses MPI::Offset. These types must have the same width and encode address values in the same manner such that offset values in one language may be passed directly to another language without conversion.

### 2.5.8 Counts

As described above, MPI defines types (e.g., MPI\_Aint) to address locations within memory and other types (e.g., MPI\_Offset) to address locations within files. In addition, some MPI procedures use *count* arguments that represent a number of MPI datatypes on which to operate. At times, one needs a single type that can be used to address locations within either memory or files as well as express *count* values, and that type is MPI\_Count in C and INTEGER (KIND=MPI\_COUNT\_KIND) in Fortran. These types must have the same width and encode values in the same manner such that count values in one language may be passed directly to another language without conversion. The size of the MPI\_Count type is determined by the MPI implementation with the restriction that it must be minimally capable of encoding any value that may be stored in a variable of type int, MPI\_Aint, or MPI\_Offset in C and of type INTEGER, INTEGER (KIND=MPI\_ADDRESS\_KIND), or INTEGER (KIND=MPI\_OFFSET\_KIND) in Fortran.

*Rationale.* Count values logically need to be large enough to encode any value used for expressing element counts, type maps in memory, type maps in file views, etc. For backward compatibility reasons, many MPI routines still use int in C and INTEGER in Fortran as the type of count arguments. (*End of rationale.*)

## 2.6 Language Binding

This section defines the rules for MPI language binding in general and for Fortran, ISO C, and C++, in particular. (Note that ANSI C has been replaced by ISO C.) The C++ language bindings have been deprecated. Defined here are various object representations, as well as the naming conventions used for expressing this standard. The actual calling sequences are defined elsewhere.

MPI bindings are for Fortran 90, though they are designed to be usable in Fortran 77 environments.

Since the word PARAMETER is a keyword in the Fortran language, we use the word “argument” to denote the arguments to a subroutine. These are normally referred to as parameters in C and C++, however, we expect that C and C++ programmers will understand the word “argument” (which has no specific meaning in C/C++), thus allowing us to avoid unnecessary confusion for Fortran programmers.

Since Fortran is case insensitive, linkers may use either lower case or upper case when resolving Fortran names. Users of case sensitive languages should avoid the “mpi\_” and “pmpi\_” prefixes.

### 2.6.1 Deprecated Names and Functions

A number of chapters refer to deprecated or replaced MPI-1 constructs. These are constructs that continue to be part of the MPI standard, as documented in Chapter 15, but that users are recommended not to continue using, since better solutions were provided with MPI-2. For example, the Fortran binding for MPI-1 functions that have address arguments uses `INTEGER`. This is not consistent with the C binding, and causes problems on machines with 32 bit `INTEGER`s and 64 bit addresses. In MPI-2, these functions were given new names with new bindings for the address arguments. The use of the old functions is deprecated. For consistency, here and in a few other cases, new C functions are also provided, even though the new functions are equivalent to the old functions. The old names are deprecated. Another example is provided by the MPI-1 predefined datatypes `MPI_UB` and `MPI_LB`. They are deprecated, since their use is awkward and error-prone. The MPI-2 function `MPI_TYPE_CREATE_RESIZED` provides a more convenient mechanism to achieve the same effect.

Table 2.1 shows a list of all of the deprecated constructs. Note that the constants `MPI_LB` and `MPI_UB` are replaced by the function `MPI_TYPE_CREATE_RESIZED`; this is because their principal use was as input datatypes to `MPI_TYPE_STRUCT` to create resized datatypes. Also note that some C typedefs and Fortran subroutine names are included in this list; they are the types of callback functions.

### 2.6.2 Fortran Binding Issues

Originally, MPI-1.1 provided bindings for Fortran 77. These bindings are retained, but they are now interpreted in the context of the Fortran 90 standard. MPI can still be used with most Fortran 77 compilers, as noted below. When the term Fortran is used it means Fortran 90.

All MPI names have an `MPI_` prefix, and all characters are capitals. Programs must not declare variables, parameters, or functions with names beginning with the prefix `MPI_`. To avoid conflicting with the profiling interface, programs should also avoid functions with the prefix `PMPI_`. This is mandated to avoid possible name collisions.

All MPI Fortran subroutines have a return code in the last argument. A few MPI operations which are functions do not have the return code argument. The return code value for successful completion is `MPI_SUCCESS`. Other error codes are implementation dependent; see the error codes in Chapter 8 and Annex A.

Constants representing the maximum length of a string are one smaller in Fortran than in C and C++ as discussed in Section 16.3.9.

Handles are represented in Fortran as `INTEGER`s. Binary-valued variables are of type `LOGICAL`.

Array arguments are indexed from one.

The MPI Fortran binding is inconsistent with the Fortran 90 standard in several respects. These inconsistencies, such as register optimization problems, have implications for user codes that are discussed in detail in Section 16.2.2. They are also inconsistent with Fortran 77.

Deprecated	MPI-2 Replacement
MPI_ADDRESS	MPI_GET_ADDRESS
MPI_TYPE_HINDEXED	MPI_TYPE_CREATE_HINDEXED
MPI_TYPE_HVECTOR	MPI_TYPE_CREATE_HVECTOR
MPI_TYPE_STRUCT	MPI_TYPE_CREATE_STRUCT
MPI_TYPE_EXTENT	MPI_TYPE_GET_EXTENT
MPI_TYPE_UB	MPI_TYPE_GET_EXTENT
MPI_TYPE_LB	MPI_TYPE_GET_EXTENT
MPI_LB	MPI_TYPE_CREATE_RESIZED
MPI_UB	MPI_TYPE_CREATE_RESIZED
MPI_ERRHANDLER_CREATE	MPI_COMM_CREATE_ERRHANDLER
MPI_ERRHANDLER_GET	MPI_COMM_GET_ERRHANDLER
MPI_ERRHANDLER_SET	MPI_COMM_SET_ERRHANDLER
MPI_Handler_function	MPI_Comm_errhandler_function
MPI_KEYVAL_CREATE	MPI_COMM_CREATE_KEYVAL
MPI_KEYVAL_FREE	MPI_COMM_FREE_KEYVAL
MPI_DUP_FN	MPI_COMM_DUP_FN
MPI_NULL_COPY_FN	MPI_COMM_NULL_COPY_FN
MPI_NULL_DELETE_FN	MPI_COMM_NULL_DELETE_FN
MPI_Copy_function	MPI_Comm_copy_attr_function
COPY_FUNCTION	COMM_COPY_ATTR_FN
MPI_Delete_function	MPI_Comm_delete_attr_function
DELETE_FUNCTION	COMM_DELETE_ATTR_FN
MPI_ATTR_DELETE	MPI_COMM_DELETE_ATTR
MPI_ATTR_GET	MPI_COMM_GET_ATTR
MPI_ATTR_PUT	MPI_COMM_SET_ATTR

Table 2.1: Deprecated constructs

### 2.6.3 C Binding Issues

We use the ISO C declaration format. All MPI names have an `MPI_` prefix, defined constants are in all capital letters, and defined types and functions have one capital letter after the prefix. Programs must not declare variables or functions with names beginning with the prefix `MPI_`. To support the profiling interface, programs should not declare functions with names beginning with the prefix `PMPI_`.

The definition of named constants, function prototypes, and type definitions must be supplied in an include file `mpi.h`.

Almost all C functions return an error code. The successful return code will be `MPI_SUCCESS`, but failure return codes are implementation dependent.

Type declarations are provided for handles to each category of opaque objects.

Array arguments are indexed from zero.

Logical flags are integers with value 0 meaning “false” and a non-zero value meaning “true.”

Choice arguments are pointers of type `void *`.

Address arguments are of MPI defined type `MPI_Aint`. File displacements are of type `MPI_Offset`. `MPI_Aint` is defined to be an integer of the size needed to hold any valid address

on the target architecture. `MPI_Offset` is defined to be an integer of the size needed to hold any valid file size on the target architecture.

#### 2.6.4 C++ Binding Issues

The C++ language bindings have been deprecated. There are places in the standard that give rules for C and not for C++. In these cases, the C rule should be applied to the C++ case, as appropriate. In particular, the values of constants given in the text are the ones for C and Fortran. A cross index of these with the C++ names is given in Annex A.

We use the ISO C++ declaration format. All MPI names are declared within the scope of a namespace called `MPI` and therefore are referenced with an `MPI::` prefix. Defined constants are in all capital letters, and class names, defined types, and functions have only their first letter capitalized. Programs must not declare variables or functions in the `MPI` namespace. This is mandated to avoid possible name collisions.

The definition of named constants, function prototypes, and type definitions must be supplied in an include file `mpi.h`.

*Advice to implementors.* The file `mpi.h` may contain both the C and C++ definitions. Usually one can simply use the defined value (generally `__cplusplus`, but not required) to see if one is using C++ to protect the C++ definitions. It is possible that a C compiler will require that the source protected this way be legal C code. In this case, all the C++ definitions can be placed in a different include file and the “`#include`” directive can be used to include the necessary C++ definitions in the `mpi.h` file. (*End of advice to implementors.*)

C++ functions that create objects or return information usually place the object or information in the return value. Since the language neutral prototypes of MPI functions include the C++ return value as an OUT parameter, semantic descriptions of MPI functions refer to the C++ return value by that parameter name. The remaining C++ functions return `void`.

In some circumstances, MPI permits users to indicate that they do not want a return value. For example, the user may indicate that the status is not filled in. Unlike C and Fortran where this is achieved through a special input value, in C++ this is done by having two bindings where one has the optional argument and one does not.

C++ functions do not return error codes. If the default error handler has been set to `MPI::ERRORS_THROW_EXCEPTIONS`, the C++ exception mechanism is used to signal an error by throwing an `MPI::Exception` object.

It should be noted that the default error handler (i.e., `MPI::ERRORS_FATAL`) on a given type has not changed. User error handlers are also permitted. `MPI::ERRORS_RETURN` simply returns control to the calling function; there is no provision for the user to retrieve the error code.

User callback functions that return integer error codes should not throw exceptions; the returned error will be handled by the MPI implementation by invoking the appropriate error handler.

*Advice to users.* C++ programmers that want to handle MPI errors on their own should use the `MPI::ERRORS_THROW_EXCEPTIONS` error handler, rather than `MPI::ERRORS_RETURN`, that is used for that purpose in C. Care should be taken using exceptions in mixed language situations. (*End of advice to users.*)

1       Opaque object handles must be objects in themselves, and have the assignment and  
2 equality operators overridden to perform semantically like their C and Fortran counterparts.

3       Array arguments are indexed from zero.

4       Logical flags are of type `bool`.

5       Choice arguments are pointers of type `void *`.

6       Address arguments are of MPI-defined integer type `MPI::Aint`, defined to be an integer  
7 of the size needed to hold any valid address on the target architecture. Analogously,  
8 `MPI::Offset` is an integer to hold file offsets.

9       Most MPI functions are methods of MPI C++ classes. MPI class names are generated  
10 from the language neutral MPI types by dropping the `MPI_` prefix and scoping the type  
11 within the MPI namespace. For example, `MPI_DATATYPE` becomes `MPI::Datatype`.

12       The names of MPI functions generally follow the naming rules given. In some circum-  
13 stances, the MPI function is related to a function defined already for MPI-1 with a name  
14 that does not follow the naming conventions. In this circumstance, the language neutral  
15 name is in analogy to the MPI name even though this gives an MPI-2 name that violates the  
16 naming conventions. The C and Fortran names are the same as the language neutral name  
17 in this case. However, the C++ names do reflect the naming rules and can differ from the C  
18 and Fortran names. Thus, the analogous name in C++ to the MPI name may be different  
19 than the language neutral name. This results in the C++ name differing from the language  
20 neutral name. An example of this is the language neutral name of `MPI_FINALIZED` and a  
21 C++ name of `MPI::Is_finalized`.

22       In C++, function `typedefs` are made publicly within appropriate classes. However,  
23 these declarations then become somewhat cumbersome, as with the following:

24 `{typedef MPI::Grequest::Query_function(); (binding deprecated, see Section 15.2)}`

25 would look like the following:

```
27 namespace MPI {
28     class Request {
29         // ...
30     };
31
32     class Grequest : public MPI::Request {
33         // ...
34         typedef Query_function(void* extra_state, MPI::Status& status);
35     };
36 };
37
```

38 Rather than including this scaffolding when declaring C++ `typedefs`, we use an abbreviated  
39 form. In particular, we explicitly indicate the class and namespace scope for the `typedef`  
40 of the function. Thus, the example above is shown in the text as follows:

```
41 typedef int MPI::Grequest::Query_function(void* extra_state,
42                                           MPI::Status& status)
```

44       The C++ bindings presented in Annex A.4 and throughout this document were gener-  
45 ated by applying a simple set of name generation rules to the MPI function specifications.  
46 While these guidelines may be sufficient in most cases, they may not be suitable for all  
47 situations. In cases of ambiguity or where a specific semantic statement is desired, these  
48 guidelines may be superseded as the situation dictates.

1. All functions, types, and constants are declared within the scope of a `namespace` called `MPI`.
2. Arrays of `MPI` handles are always left in the argument list (whether they are `IN` or `OUT` arguments).
3. If the argument list of an `MPI` function contains a scalar `IN` handle, and it makes sense to define the function as a method of the object corresponding to that handle, the function is made a member function of the corresponding `MPI` class. The member functions are named according to the corresponding `MPI` function name, but without the “`MPI_`” prefix and without the object name prefix (if applicable). In addition:
  - (a) The scalar `IN` handle is dropped from the argument list, and `this` corresponds to the dropped argument.
  - (b) The function is declared `const`.
4. `MPI` functions are made into class functions (static) when they belong on a class but do not have a unique scalar `IN` or `INOUT` parameter of that class.
5. If the argument list contains a single `OUT` argument that is not of type `MPI_STATUS` (or an array), that argument is dropped from the list and the function returns that value.

**Example 2.1** The C++ binding for `MPI_COMM_SIZE` is  
`int MPI::Comm::Get_size(void) const.`

6. If there are multiple `OUT` arguments in the argument list, one is chosen as the return value and is removed from the list.
7. If the argument list does not contain any `OUT` arguments, the function returns `void`.

**Example 2.2** The C++ binding for `MPI_REQUEST_FREE` is  
`void MPI::Request::Free(void)`

8. `MPI` functions to which the above rules do not apply are not members of any class, but are defined in the `MPI` namespace.

**Example 2.3** The C++ binding for `MPI_BUFFER_ATTACH` is  
`void MPI::Attach_buffer(void* buffer, int size).`

9. All class names, defined types, and function names have only their first letter capitalized. Defined constants are in all capital letters.
10. Any `IN` pointer, reference, or array argument must be declared `const`.
11. Handles are passed by reference.
12. Array arguments are denoted with square brackets (`[]`), not pointers, as this is more semantically precise.

### 2.6.5 Functions and Macros

An implementation is allowed to implement `MPI_WTIME`, `MPI_WTICK`, `PMPI_WTIME`, `PMPI_WTICK`, and the handle-conversion functions (`MPI_Group_f2c`, etc.) in Section 16.3.4, and no others, as macros in C.

*Advice to implementors.* Implementors should document which routines are implemented as macros. (*End of advice to implementors.*)

*Advice to users.* If these routines are implemented as macros, they will not work with the MPI profiling interface. (*End of advice to users.*)

## 2.7 Processes

An MPI program consists of autonomous processes, executing their own code, in an MIMD style. The codes executed by each process need not be identical. The processes communicate via calls to MPI communication primitives. Typically, each process executes in its own address space, although shared-memory implementations of MPI are possible.

This document specifies the behavior of a parallel program assuming that only MPI calls are used. The interaction of an MPI program with other possible means of communication, I/O, and process management is not specified. Unless otherwise stated in the specification of the standard, MPI places no requirements on the result of its interaction with external mechanisms that provide similar or equivalent functionality. This includes, but is not limited to, interactions with external mechanisms for process control, shared and remote memory access, file system access and control, interprocess communication, process signaling, and terminal I/O. High quality implementations should strive to make the results of such interactions intuitive to users, and attempt to document restrictions where deemed necessary.

*Advice to implementors.* Implementations that support such additional mechanisms for functionality supported within MPI are expected to document how these interact with MPI. (*End of advice to implementors.*)

The interaction of MPI and threads is defined in Section 12.4.

## 2.8 Error Handling

MPI provides the user with reliable message transmission. A message sent is always received correctly, and the user does not need to check for transmission errors, time-outs, or other error conditions. In other words, MPI does not provide mechanisms for dealing with failures in the communication system. If the MPI implementation is built on an unreliable underlying mechanism, then it is the job of the implementor of the MPI subsystem to insulate the user from this unreliability, or to reflect unrecoverable errors as failures. Whenever possible, such failures will be reflected as errors in the relevant communication call. Similarly, MPI itself provides no mechanisms for handling processor failures.

Of course, MPI programs may still be erroneous. A **program error** can occur when an MPI call is made with an incorrect argument (non-existing destination in a send operation, buffer too small in a receive operation, etc.). This type of error would occur in any

implementation. In addition, a **resource error** may occur when a program exceeds the amount of available system resources (number of pending messages, system buffers, etc.). The occurrence of this type of error depends on the amount of available resources in the system and the resource allocation mechanism used; this may differ from system to system. A high-quality implementation will provide generous limits on the important resources so as to alleviate the portability problem this represents.

In C and Fortran, almost all MPI calls return a code that indicates successful completion of the operation. Whenever possible, MPI calls return an error code if an error occurred during the call. By default, an error detected during the execution of the MPI library causes the parallel computation to abort, except for file operations. However, MPI provides mechanisms for users to change this default and to handle recoverable errors. The user may specify that no error is fatal, and handle error codes returned by MPI calls by himself or herself. Also, the user may provide his or her own error-handling routines, which will be invoked whenever an MPI call returns abnormally. The MPI error handling facilities are described in Section 8.3. The return values of C++ functions are not error codes. If the default error handler has been set to `MPI::ERRORS_THROW_EXCEPTIONS`, the C++ exception mechanism is used to signal an error by throwing an `MPI::Exception` object. See also Section 16.1.8 on page 508.

Several factors limit the ability of MPI calls to return with meaningful error codes when an error occurs. MPI may not be able to detect some errors; other errors may be too expensive to detect in normal execution mode; finally some errors may be “catastrophic” and may prevent MPI from returning control to the caller in a consistent state.

Another subtle issue arises because of the nature of asynchronous communications: MPI calls may initiate operations that continue asynchronously after the call returned. Thus, the operation may return with a code indicating successful completion, yet later cause an error exception to be raised. If there is a subsequent call that relates to the same operation (e.g., a call that verifies that an asynchronous operation has completed) then the error argument associated with this call will be used to indicate the nature of the error. In a few cases, the error may occur after all calls that relate to the operation have completed, so that no error value can be used to indicate the nature of the error (e.g., an error on the receiver in a send with the ready mode). Such an error must be treated as fatal, since information cannot be returned for the user to recover from it.

This document does not specify the state of a computation after an erroneous MPI call has occurred. The desired behavior is that a relevant error code be returned, and the effect of the error be localized to the greatest possible extent. E.g., it is highly desirable that an erroneous receive call will not cause any part of the receiver’s memory to be overwritten, beyond the area specified for receiving the message.

Implementations may go beyond this document in supporting in a meaningful manner MPI calls that are defined here to be erroneous. For example, MPI specifies strict type matching rules between matching send and receive operations: it is erroneous to send a floating point variable and receive an integer. Implementations may go beyond these type matching rules, and provide automatic type conversion in such situations. It will be helpful to generate warnings for such non-conforming behavior.

MPI defines a way for users to create new error codes as defined in Section 8.5.



## 2.9 Implementation Issues

There are a number of areas where an MPI implementation may interact with the operating environment and system. While MPI does not mandate that any services (such as signal handling) be provided, it does strongly suggest the behavior to be provided if those services are available. This is an important point in achieving portability across platforms that provide the same set of services.

### 2.9.1 Independence of Basic Runtime Routines

MPI programs require that library routines that are part of the basic language environment (such as `write` in Fortran and `printf` and `malloc` in ISO C) and are executed after `MPI_INIT` and before `MPI_FINALIZE` operate independently and that their *completion* is independent of the action of other processes in an MPI program.

Note that this in no way prevents the creation of library routines that provide parallel services whose operation is collective. However, the following program is expected to complete in an ISO C environment regardless of the size of `MPI_COMM_WORLD` (assuming that `printf` is available at the executing nodes).

```
int rank;
MPI_Init((void *)0, (void *)0);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
if (rank == 0) printf("Starting program\n");
MPI_Finalize();
```

The corresponding Fortran and C++ programs are also expected to complete.

An example of what is *not* required is any particular ordering of the action of these routines when called by several tasks. For example, MPI makes neither requirements nor recommendations for the output from the following program (again assuming that I/O is available at the executing nodes).

```
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
printf("Output from task rank %d\n", rank);
```

In addition, calls that fail because of resource exhaustion or other error are not considered a violation of the requirements here (however, they are required to complete, just not to complete successfully).

### 2.9.2 Interaction with Signals

MPI does not specify the interaction of processes with signals and does not require that MPI be signal safe. The implementation may reserve some signals for its own use. It is required that the implementation document which signals it uses, and it is strongly recommended that it not use `SIGALRM`, `SIGFPE`, or `SIGIO`. Implementations may also prohibit the use of MPI calls from within signal handlers.

In multithreaded environments, users can avoid conflicts between signals and the MPI library by catching signals only on threads that do not execute MPI calls. High quality single-threaded implementations will be signal safe: an MPI call suspended by a signal will resume and complete normally after the signal is handled.

## 2.10 Examples

The examples in this document are for illustration purposes only. They are not intended to specify the standard. Furthermore, the examples have not been carefully checked or verified.

1  
2  
3  
4  
5  
6  
7  
8  
9  
10  
11  
12  
13  
14  
15  
16  
17  
18  
19  
20  
21  
22  
23  
24  
25  
26  
27  
28  
29  
30  
31  
32  
33  
34  
35  
36  
37  
38  
39  
40  
41  
42  
43  
44  
45  
46  
47  
48

1  
2  
3  
4  
5  
6  
7  
8  
9  
10  
11  
12  
13  
14  
15  
16  
17  
18  
19  
20  
21  
22  
23  
24  
25  
26  
27  
28  
29  
30  
31  
32  
33  
34  
35  
36  
37  
38  
39  
40  
41  
42  
43  
44  
45  
46  
47  
48

## Chapter 3

# Point-to-Point Communication

### 3.1 Introduction

Sending and receiving of messages by processes is the basic MPI communication mechanism. The basic point-to-point communication operations are **send** and **receive**. Their use is illustrated in the example below.

```
#include "mpi.h"
int main( int argc, char **argv )
{
    char message[20];
    int myrank;
    MPI_Status status;
    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &myrank );
    if (myrank == 0) /* code for process zero */
    {
        strcpy(message,"Hello, there");
        MPI_Send(message, strlen(message)+1, MPI_CHAR, 1, 99, MPI_COMM_WORLD);
    }
    else if (myrank == 1) /* code for process one */
    {
        MPI_Recv(message, 20, MPI_CHAR, 0, 99, MPI_COMM_WORLD, &status);
        printf("received :%s:\n", message);
    }
    MPI_Finalize();
}
```

In this example, process zero (`myrank = 0`) sends a message to process one using the **send** operation `MPI_SEND`. The operation specifies a **send buffer** in the sender memory from which the message data is taken. In the example above, the send buffer consists of the storage containing the variable `message` in the memory of process zero. The location, size and type of the send buffer are specified by the first three parameters of the send operation. The message sent will contain the 13 characters of this variable. In addition, the send operation associates an **envelope** with the message. This envelope specifies the message destination and contains distinguishing information that can be used by the **receive**

operation to select a particular message. The last three parameters of the send operation, along with the rank of the sender, specify the envelope for the message sent. Process one (`myrank = 1`) receives this message with the **receive** operation `MPI_RECV`. The message to be received is selected according to the value of its envelope, and the message data is stored into the **receive buffer**. In the example above, the receive buffer consists of the storage containing the string `message` in the memory of process one. The first three parameters of the receive operation specify the location, size and type of the receive buffer. The next three parameters are used for selecting the incoming message. The last parameter is used to return information on the message just received.

The next sections describe the blocking send and receive operations. We discuss send, receive, blocking communication semantics, type matching requirements, type conversion in heterogeneous environments, and more general communication modes. Nonblocking communication is addressed next, followed by channel-like constructs and send-receive operations. Nonblocking communication is addressed next, followed by channel-like constructs and send-receive operations, ending with a description of the “dummy” process, `MPI_PROC_NULL`.

## 3.2 Blocking Send and Receive Operations

### 3.2.1 Blocking Send

The syntax of the blocking send operation is given below.

`MPI_SEND(buf, count, datatype, dest, tag, comm)`

IN	buf	initial address of send buffer (choice)
IN	count	number of elements in send buffer (non-negative integer)
IN	datatype	datatype of each send buffer element (handle)
IN	dest	rank of destination (integer)
IN	tag	message tag (integer)
IN	comm	communicator (handle)

```
int MPI_Send(void* buf, int count, MPI_Datatype datatype, int dest,
             int tag, MPI_Comm comm)
```

```
MPI_SEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, IERROR)
```

```
<type> BUF(*)
```

```
INTEGER COUNT, DATATYPE, DEST, TAG, COMM, IERROR
```

```
{void MPI::Comm::Send(const void* buf, int count, const
                      MPI::Datatype& datatype, int dest, int tag) const(binding
                      deprecated, see Section 15.2) }
```

The blocking semantics of this call are described in Section 3.4.

### 3.2.2 Message Data

The send buffer specified by the `MPI_SEND` operation consists of `count` successive entries of the type indicated by `datatype`, starting with the entry at address `buf`. Note that we specify the message length in terms of number of *elements*, not number of *bytes*. The former is machine independent and closer to the application level.

The data part of the message consists of a sequence of `count` values, each of the type indicated by `datatype`. `count` may be zero, in which case the data part of the message is empty. The basic datatypes that can be specified for message data values correspond to the basic datatypes of the host language. Possible values of this argument for Fortran and the corresponding Fortran types are listed in Table 3.1.

MPI datatype	Fortran datatype
<code>MPI_INTEGER</code>	<code>INTEGER</code>
<code>MPI_REAL</code>	<code>REAL</code>
<code>MPI_DOUBLE_PRECISION</code>	<code>DOUBLE PRECISION</code>
<code>MPI_COMPLEX</code>	<code>COMPLEX</code>
<code>MPI_LOGICAL</code>	<code>LOGICAL</code>
<code>MPI_CHARACTER</code>	<code>CHARACTER(1)</code>
<code>MPI_BYTE</code>	
<code>MPI_PACKED</code>	

Table 3.1: Predefined MPI datatypes corresponding to Fortran datatypes

Possible values for this argument for C and the corresponding C types are listed in Table 3.2.

The datatypes `MPI_BYTE` and `MPI_PACKED` do not correspond to a Fortran or C datatype. A value of type `MPI_BYTE` consists of a byte (8 binary digits). A byte is uninterpreted and is different from a character. Different machines may have different representations for characters, or may use more than one byte to represent characters. On the other hand, a byte has the same binary value on all machines. The use of the type `MPI_PACKED` is explained in Section 4.2.

MPI requires support of these datatypes, which match the basic datatypes of Fortran and ISO C. Additional MPI datatypes should be provided if the host language has additional data types: `MPI_DOUBLE_COMPLEX` for double precision complex in Fortran declared to be of type `DOUBLE COMPLEX`; `MPI_REAL2`, `MPI_REAL4` and `MPI_REAL8` for Fortran reals, declared to be of type `REAL*2`, `REAL*4` and `REAL*8`, respectively; `MPI_INTEGER1`, `MPI_INTEGER2` and `MPI_INTEGER4` for Fortran integers, declared to be of type `INTEGER*1`, `INTEGER*2` and `INTEGER*4`, respectively; etc.

*Rationale.* One goal of the design is to allow for MPI to be implemented as a library, with no need for additional preprocessing or compilation. Thus, one cannot assume that a communication call has information on the datatype of variables in the communication buffer; this information must be supplied by an explicit argument. The need for such datatype information will become clear in Section 3.3.2. (*End of rationale.*)

*Rationale.* The datatypes `MPI_C_BOOL`, `MPI_INT8_T`, `MPI_INT16_T`, `MPI_INT32_T`, `MPI_UINT8_T`, `MPI_UINT16_T`, `MPI_UINT32_T`, `MPI_C_COMPLEX`,

MPI datatype	C datatype
MPI_CHAR	char (treated as printable character)
MPI_SHORT	signed short int
MPI_INT	signed int
MPI_LONG	signed long int
MPI_LONG_LONG_INT	signed long long int
MPI_LONG_LONG (as a synonym)	signed long long int
MPI_SIGNED_CHAR	signed char (treated as integral value)
MPI_UNSIGNED_CHAR	unsigned char (treated as integral value)
MPI_UNSIGNED_SHORT	unsigned short int
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_UNSIGNED_LONG_LONG	unsigned long long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_WCHAR	wchar_t (defined in <stddef.h>) (treated as printable character)
MPI_C_BOOL	_Bool
MPI_INT8_T	int8_t
MPI_INT16_T	int16_t
MPI_INT32_T	int32_t
MPI_INT64_T	int64_t
MPI_UINT8_T	uint8_t
MPI_UINT16_T	uint16_t
MPI_UINT32_T	uint32_t
MPI_UINT64_T	uint64_t
MPI_C_COMPLEX	float _Complex
MPI_C_FLOAT_COMPLEX (as a synonym)	float _Complex
MPI_C_DOUBLE_COMPLEX	double _Complex
MPI_C_LONG_DOUBLE_COMPLEX	long double _Complex
MPI_BYTE	
MPI_PACKED	

Table 3.2: Predefined MPI datatypes corresponding to C datatypes

MPI\_C\_FLOAT\_COMPLEX, MPI\_C\_DOUBLE\_COMPLEX, and MPI\_C\_LONG\_DOUBLE\_COMPLEX have no corresponding C++ bindings. This was intentionally done to avoid potential collisions with the C preprocessor and namespaced C++ names. C++ applications can use the C bindings with no loss of functionality. (*End of rationale.*)

The datatypes MPI\_AINT [and ], MPI\_OFFSET , and MPI\_COUNT correspond to the

MPI datatype	C datatype	Fortran datatype	
MPI_AINT	MPI_Aint	INTEGER (KIND=MPI_ADDRESS_KIND)	1
MPI_OFFSET	MPI_Offset	INTEGER (KIND=MPI_OFFSET_KIND)	2
[ticket265.] MPI_COUNT	[ticket265.]MPI_Count	[ticket265.]INTEGER (KIND=MPI_COUNT_KIND)	3

Table 3.3: Predefined MPI datatypes corresponding to both C and Fortran datatypes

MPI-defined C types MPI\_Aint [and ], MPI\_Offset , and MPI\_Count and their Fortran equivalents INTEGER (KIND=MPI\_ADDRESS\_KIND) [and ], INTEGER (KIND=MPI\_OFFSET\_KIND) , and INTEGER (KIND=MPI\_COUNT\_KIND) . This is described in Table 3.3. See Section 16.3.10 for information on interlanguage communication with these types.

### 3.2.3 Message Envelope

In addition to the data part, messages carry information that can be used to distinguish messages and selectively receive them. This information consists of a fixed number of fields, which we collectively call the **message envelope**. These fields are

source  
destination  
tag  
communicator

The message source is implicitly determined by the identity of the message sender. The other fields are specified by arguments in the send operation.

The message destination is specified by the **dest** argument.

The integer-valued message tag is specified by the **tag** argument. This integer can be used by the program to distinguish different types of messages. The range of valid tag values is 0,...,UB, where the value of UB is implementation dependent. It can be found by querying the value of the attribute MPI\_TAG\_UB, as described in Chapter 8. MPI requires that UB be no less than 32767.

The **comm** argument specifies the **communicator** that is used for the send operation. Communicators are explained in Chapter 6; below is a brief summary of their usage.

A communicator specifies the communication context for a communication operation. Each communication context provides a separate “communication universe:” messages are always received within the context they were sent, and messages sent in different contexts do not interfere.

The communicator also specifies the set of processes that share this communication context. This **process group** is ordered and processes are identified by their rank within this group. Thus, the range of valid values for **dest** is 0, ... , n-1, where n is the number of processes in the group. (If the communicator is an inter-communicator, then destinations are identified by their rank in the remote group. See Chapter 6.)

A predefined communicator MPI\_COMM\_WORLD is provided by MPI. It allows communication with all processes that are accessible after MPI initialization and processes are identified by their rank in the group of MPI\_COMM\_WORLD.

*Advice to users.* Users that are comfortable with the notion of a flat name space for processes, and a single communication context, as offered by most existing communication libraries, need only use the predefined variable MPI\_COMM\_WORLD as the



comm argument. This will allow communication with all the processes available at initialization time.

Users may define new communicators, as explained in Chapter 6. Communicators provide an important encapsulation mechanism for libraries and modules. They allow modules to have their own disjoint communication universe and their own process numbering scheme. (*End of advice to users.*)

*Advice to implementors.* The message envelope would normally be encoded by a fixed-length message header. However, the actual encoding is implementation dependent. Some of the information (e.g., source or destination) may be implicit, and need not be explicitly carried by messages. Also, processes may be identified by relative ranks, or absolute ids, etc. (*End of advice to implementors.*)

### 3.2.4 Blocking Receive

The syntax of the blocking receive operation is given below.

**MPI\_RECV** (buf, count, datatype, source, tag, comm, status)

OUT	buf	initial address of receive buffer (choice)
IN	count	number of elements in receive buffer (non-negative integer)
IN	datatype	datatype of each receive buffer element (handle)
IN	source	rank of source or MPI_ANY_SOURCE (integer)
IN	tag	message tag or MPI_ANY_TAG (integer)
IN	comm	communicator (handle)
OUT	status	status object (Status)

```
int MPI_Recv(void* buf, int count, MPI_Datatype datatype, int source,
             int tag, MPI_Comm comm, MPI_Status *status)
```

```
MPI_RECV(BUF, COUNT, DATATYPE, SOURCE, TAG, COMM, STATUS, IERROR)
<type> BUF(*)
INTEGER COUNT, DATATYPE, SOURCE, TAG, COMM, STATUS(MPI_STATUS_SIZE),
IERROR
```

```
{void MPI::Comm::Recv(void* buf, int count, const MPI::Datatype& datatype,
                      int source, int tag, MPI::Status& status) const(binding
                      deprecated, see Section 15.2) }
```

```
{void MPI::Comm::Recv(void* buf, int count, const MPI::Datatype& datatype,
                      int source, int tag) const(binding deprecated, see Section 15.2) }
```

The blocking semantics of this call are described in Section 3.4.

The receive buffer consists of the storage containing count consecutive elements of the type specified by datatype, starting at address buf. The length of the received message must be less than or equal to the length of the receive buffer. An overflow error occurs if all incoming data does not fit, without truncation, into the receive buffer.

If a message that is shorter than the receive buffer arrives, then only those locations corresponding to the (shorter) message are modified.

*Advice to users.* The `MPI_PROBE` function described in Section 3.8 can be used to receive messages of unknown length. (*End of advice to users.*)

*Advice to implementors.* Even though no specific behavior is mandated by MPI for erroneous programs, the recommended handling of overflow situations is to return in status information about the source and tag of the incoming message. The receive operation will return an error code. A quality implementation will also ensure that no memory that is outside the receive buffer will ever be overwritten.

In the case of a message shorter than the receive buffer, MPI is quite strict in that it allows no modification of the other locations. A more lenient statement would allow for some optimizations but this is not allowed. The implementation must be ready to end a copy into the receiver memory exactly at the end of the receive buffer, even if it is an odd address. (*End of advice to implementors.*)

The selection of a message by a receive operation is governed by the value of the message envelope. A message can be received by a receive operation if its envelope matches the source, tag and comm values specified by the receive operation. The receiver may specify a wildcard `MPI_ANY_SOURCE` value for source, and/or a wildcard `MPI_ANY_TAG` value for tag, indicating that any source and/or tag are acceptable. It cannot specify a wildcard value for comm. Thus, a message can be received by a receive operation only if it is addressed to the receiving process, has a matching communicator, has matching source unless source=`MPI_ANY_SOURCE` in the pattern, and has a matching tag unless tag=`MPI_ANY_TAG` in the pattern.

The message tag is specified by the tag argument of the receive operation. The argument source, if different from `MPI_ANY_SOURCE`, is specified as a rank within the process group associated with that same communicator (remote process group, for intercommunicators). Thus, the range of valid values for the source argument is  $\{0, \dots, n-1\} \cup \{\text{MPI\_ANY\_SOURCE}\}$ , where n is the number of processes in this group.

Note the asymmetry between send and receive operations: A receive operation may accept messages from an arbitrary sender, on the other hand, a send operation must specify a unique receiver. This matches a “push” communication mechanism, where data transfer is effected by the sender (rather than a “pull” mechanism, where data transfer is effected by the receiver).

Source = destination is allowed, that is, a process can send a message to itself. (However, it is unsafe to do so with the blocking send and receive operations described above, since this may lead to deadlock. See Section 3.5.)

*Advice to implementors.* Message context and other communicator information can be implemented as an additional tag field. It differs from the regular message tag in that wild card matching is not allowed on this field, and that value setting for this field is controlled by communicator manipulation functions. (*End of advice to implementors.*)

### 3.2.5 Return Status

The source or tag of a received message may not be known if wildcard values were used in the receive operation. Also, if multiple requests are completed by a single MPI function

(see Section 3.7.5), a distinct error code may need to be returned for each request. The information is returned by the `status` argument of `MPI_RECV`. The type of `status` is MPI-defined. Status variables need to be explicitly allocated by the user, that is, they are not system objects.

In C, `status` is a structure that contains three fields named `MPI_SOURCE`, `MPI_TAG`, and `MPI_ERROR`; the structure may contain additional fields. Thus, `status.MPI_SOURCE`, `status.MPI_TAG` and `status.MPI_ERROR` contain the source, tag, and error code, respectively, of the received message.

In Fortran, `status` is an array of `INTEGER`s of size `MPI_STATUS_SIZE`. The constants `MPI_SOURCE`, `MPI_TAG` and `MPI_ERROR` are the indices of the entries that store the source, tag and error fields. Thus, `status(MPI_SOURCE)`, `status(MPI_TAG)` and `status(MPI_ERROR)` contain, respectively, the source, tag and error code of the received message.

In C++, the `status` object is handled through the following methods:

```
{int MPI::Status::Get_source() const(binding deprecated, see Section 15.2) }
{void MPI::Status::Set_source(int source)(binding deprecated, see Section 15.2) }
{int MPI::Status::Get_tag() const(binding deprecated, see Section 15.2) }
{void MPI::Status::Set_tag(int tag)(binding deprecated, see Section 15.2) }
{int MPI::Status::Get_error() const(binding deprecated, see Section 15.2) }
{void MPI::Status::Set_error(int error)(binding deprecated, see Section 15.2) }
```

In general, message-passing calls do not modify the value of the error code field of `status` variables. This field may be updated only by the functions in Section 3.7.5 which return multiple statuses. The field is updated if and only if such function returns with an error code of `MPI_ERR_IN_STATUS`.

*Rationale.* The error field in `status` is not needed for calls that return only one status, such as `MPI_WAIT`, since that would only duplicate the information returned by the function itself. The current design avoids the additional overhead of setting it, in such cases. The field is needed for calls that return multiple statuses, since each request may have had a different failure. (*End of rationale.*)

The `status` argument also returns information on the length of the message received. However, this information is not directly available as a field of the `status` variable and a call to `MPI_GET_COUNT` is required to “decode” this information.

`MPI_GET_COUNT(status, datatype, count)`

IN	<code>status</code>	return status of receive operation (Status)
IN	<code>datatype</code>	datatype of each receive buffer entry (handle)
OUT	<code>count</code>	number of received entries (integer)

```
int MPI_Get_count(MPI_Status *status, MPI_Datatype datatype, int *count)
```

```
MPI_GET_COUNT(STATUS, DATATYPE, COUNT, IERROR)
```

```

    INTEGER STATUS(MPI_STATUS_SIZE), DATATYPE, COUNT, IERROR
    {int MPI::Status::Get_count(const MPI::Datatype& datatype) const(binding
        deprecated, see Section 15.2) }
```

Returns the number of entries received. (Again, we count *entries*, each of type *datatype*, not *bytes*.) The *datatype* argument should match the argument provided by the receive call that set the *status* variable. If the number of entries received exceeds the limits of the *count* parameter, then `MPI_GET_COUNT` sets the value of *count* to `MPI_UNDEFINED`. [(We shall later see, in Section 4.1.11, that `MPI_GET_COUNT` may return, in certain situations, the value `MPI_UNDEFINED`.)] There are other situations where the value of *count* can be set to `MPI_UNDEFINED`; see Section 4.1.11.

*Rationale.* Some message-passing libraries use `INOUT count`, `tag` and `source` arguments, thus using them both to specify the selection criteria for incoming messages and return the actual envelope values of the received message. The use of a separate status argument prevents errors that are often attached with `INOUT` argument (e.g., using the `MPI_ANY_TAG` constant as the tag in a receive). Some libraries use calls that refer implicitly to the “last message received.” This is not thread safe.

The *datatype* argument is passed to `MPI_GET_COUNT` so as to improve performance. A message might be received without counting the number of elements it contains, and the count value is often not needed. Also, this allows the same function to be used after a call to `MPI_PROBE` or `MPI_IPROBE`. With a status from `MPI_PROBE` or `MPI_IPROBE`, the same datatypes are allowed as in a call to `MPI_RECV` to receive this message. (*End of rationale.*)

The value returned as the *count* argument of `MPI_GET_COUNT` for a datatype of length zero where zero bytes have been transferred is zero. If the number of bytes transferred is greater than zero, `MPI_UNDEFINED` is returned.

*Rationale.* Zero-length datatypes may be created in a number of cases. An important case is `MPI_TYPE_CREATE_DARRAY`, where the definition of the particular darray results in an empty block on some MPI process. Programs written in an SPMD style will not check for this special case and may want to use `MPI_GET_COUNT` to check the status. (*End of rationale.*)

*Advice to users.* The buffer size required for the receive can be affected by data conversions and by the stride of the receive datatype. In most cases, the safest approach is to use the same datatype with `MPI_GET_COUNT` and the receive. (*End of advice to users.*)

All send and receive operations use the *buf*, *count*, *datatype*, *source*, *dest*, *tag*, *comm* and *status* arguments in the same way as the blocking `MPI_SEND` and `MPI_RECV` operations described in this section.

### 3.2.6 Passing `MPI_STATUS_IGNORE` for Status

Every call to `MPI_RECV` includes a *status* argument, wherein the system can return details about the message received. There are also a number of other MPI calls where *status* is returned. An object of type `MPI_STATUS` is not an MPI opaque object; its structure

is declared in `mpi.h` and `mpif.h`, and it exists in the user's program. In many cases, application programs are constructed so that it is unnecessary for them to examine the `status` fields. In these cases, it is a waste for the user to allocate a status object, and it is particularly wasteful for the MPI implementation to fill in fields in this object.

To cope with this problem, there are two predefined constants, `MPI_STATUS_IGNORE` and `MPI_STATUSES_IGNORE`, which when passed to a receive, wait, or test function, inform the implementation that the status fields are not to be filled in. Note that `MPI_STATUS_IGNORE` is not a special type of `MPI_STATUS` object; rather, it is a special value for the argument. In C one would expect it to be `NULL`, not the address of a special `MPI_STATUS`.

`MPI_STATUS_IGNORE`, and the array version `MPI_STATUSES_IGNORE`, can be used everywhere a status argument is passed to a receive, wait, or test function. `MPI_STATUS_IGNORE` cannot be used when status is an `IN` argument. Note that in Fortran `MPI_STATUS_IGNORE` and `MPI_STATUSES_IGNORE` are objects like `MPI_BOTTOM` (not usable for initialization or assignment). See Section 2.5.4.

In general, this optimization can apply to all functions for which `status` or an array of `statuses` is an `OUT` argument. Note that this converts `status` into an `INOUT` argument. The functions that can be passed `MPI_STATUS_IGNORE` are all the various forms of `MPI_RECV`, `MPI_TEST`, and `MPI_WAIT`, as well as `MPI_REQUEST_GET_STATUS`. When an array is passed, as in the `MPI_{TEST|WAIT}{ALL|SOME}` functions, a separate constant, `MPI_STATUSES_IGNORE`, is passed for the array argument. It is possible for an MPI function to return `MPI_ERR_IN_STATUS` even when `MPI_STATUS_IGNORE` or `MPI_STATUSES_IGNORE` has been passed to that function.

`MPI_STATUS_IGNORE` and `MPI_STATUSES_IGNORE` are not required to have the same values in C and Fortran.

It is not allowed to have some of the statuses in an array of statuses for `MPI_{TEST|WAIT}{ALL|SOME}` functions set to `MPI_STATUS_IGNORE`; one either specifies ignoring *all* of the statuses in such a call with `MPI_STATUSES_IGNORE`, or *none* of them by passing normal statuses in all positions in the array of statuses.

There are no C++ bindings for `MPI_STATUS_IGNORE` or `MPI_STATUSES_IGNORE`. To allow an `OUT` or `INOUT` `MPI::Status` argument to be ignored, all MPI C++ bindings that have `OUT` or `INOUT` `MPI::Status` parameters are overloaded with a second version that omits the `OUT` or `INOUT` `MPI::Status` parameter.

**Example 3.1** The C++ bindings for `MPI_PROBE` are:

```
void MPI::Comm::Probe(int source, int tag, MPI::Status& status) const
void MPI::Comm::Probe(int source, int tag) const
```

## 3.3 Data Type Matching and Data Conversion

### 3.3.1 Type Matching Rules

One can think of message transfer as consisting of the following three phases.

1. Data is pulled out of the send buffer and a message is assembled.
2. A message is transferred from sender to receiver.
3. Data is pulled from the incoming message and disassembled into the receive buffer.

Type matching has to be observed at each of these three phases: The type of each variable in the sender buffer has to match the type specified for that entry by the send operation; the type specified by the send operation has to match the type specified by the receive operation; and the type of each variable in the receive buffer has to match the type specified for that entry by the receive operation. A program that fails to observe these three rules is erroneous.

To define type matching more precisely, we need to deal with two issues: matching of types of the host language with types specified in communication operations; and matching of types at sender and receiver.

The types of a send and receive match (phase two) if both operations use identical names. That is, `MPI_INTEGER` matches `MPI_INTEGER`, `MPI_REAL` matches `MPI_REAL`, and so on. There is one exception to this rule, discussed in Section 4.2, the type `MPI_PACKED` can match any other type.

The type of a variable in a host program matches the type specified in the communication operation if the datatype name used by that operation corresponds to the basic type of the host program variable. For example, an entry with type name `MPI_INTEGER` matches a Fortran variable of type `INTEGER`. A table giving this correspondence for Fortran and C appears in Section 3.2.2. There are two exceptions to this last rule: an entry with type name `MPI_BYTE` or `MPI_PACKED` can be used to match any byte of storage (on a byte-addressable machine), irrespective of the datatype of the variable that contains this byte. The type `MPI_PACKED` is used to send data that has been explicitly packed, or receive data that will be explicitly unpacked, see Section 4.2. The type `MPI_BYTE` allows one to transfer the binary value of a byte in memory unchanged.

To summarize, the type matching rules fall into the three categories below.

- Communication of typed values (e.g., with datatype different from `MPI_BYTE`), where the datatypes of the corresponding entries in the sender program, in the send call, in the receive call and in the receiver program must all match.
- Communication of untyped values (e.g., of datatype `MPI_BYTE`), where both sender and receiver use the datatype `MPI_BYTE`. In this case, there are no requirements on the types of the corresponding entries in the sender and the receiver programs, nor is it required that they be the same.
- Communication involving packed data, where `MPI_PACKED` is used.

The following examples illustrate the first two cases.

**Example 3.2** Sender and receiver specify matching types.

```
CALL MPI_COMM_RANK(comm, rank, ierr)
IF (rank.EQ.0) THEN
    CALL MPI_SEND(a(1), 10, MPI_REAL, 1, tag, comm, ierr)
ELSE IF (rank.EQ.1) THEN
    CALL MPI_RECV(b(1), 15, MPI_REAL, 0, tag, comm, status, ierr)
END IF
```

This code is correct if both `a` and `b` are real arrays of size  $\geq 10$ . (In Fortran, it might be correct to use this code even if `a` or `b` have size  $< 10$ : e.g., when `a(1)` can be equivalenced to an array with ten reals.)

**Example 3.3** Sender and receiver do not specify matching types.

```

CALL MPI_COMM_RANK(comm, rank, ierr)
IF (rank.EQ.0) THEN
    CALL MPI_SEND(a(1), 10, MPI_REAL, 1, tag, comm, ierr)
ELSE IF (rank.EQ.1) THEN
    CALL MPI_RECV(b(1), 40, MPI_BYTE, 0, tag, comm, status, ierr)
END IF

```

This code is erroneous, since sender and receiver do not provide matching datatype arguments.

**Example 3.4** Sender and receiver specify communication of untyped values.

```

CALL MPI_COMM_RANK(comm, rank, ierr)
IF (rank.EQ.0) THEN
    CALL MPI_SEND(a(1), 40, MPI_BYTE, 1, tag, comm, ierr)
ELSE IF (rank.EQ.1) THEN
    CALL MPI_RECV(b(1), 60, MPI_BYTE, 0, tag, comm, status, ierr)
END IF

```

This code is correct, irrespective of the type and size of *a* and *b* (unless this results in an out of bound memory access).

*Advice to users.* If a buffer of type `MPI_BYTE` is passed as an argument to `MPI_SEND`, then MPI will send the data stored at contiguous locations, starting from the address indicated by the *buf* argument. This may have unexpected results when the data layout is not as a casual user would expect it to be. For example, some Fortran compilers implement variables of type `CHARACTER` as a structure that contains the character length and a pointer to the actual string. In such an environment, sending and receiving a Fortran `CHARACTER` variable using the `MPI_BYTE` type will not have the anticipated result of transferring the character string. For this reason, the user is advised to use typed communications whenever possible. (*End of advice to users.*)

#### Type `MPI_CHARACTER`

The type `MPI_CHARACTER` matches one character of a Fortran variable of type `CHARACTER`, rather than the entire character string stored in the variable. Fortran variables of type `CHARACTER` or substrings are transferred as if they were arrays of characters. This is illustrated in the example below.

#### Example 3.5

Transfer of Fortran `CHARACTER`s.



```

CHARACTER*10 a
CHARACTER*10 b

CALL MPI_COMM_RANK(comm, rank, ierr)
IF (rank.EQ.0) THEN
    CALL MPI_SEND(a, 5, MPI_CHARACTER, 1, tag, comm, ierr)
ELSE IF (rank.EQ.1) THEN
    CALL MPI_RECV(b(6:10), 5, MPI_CHARACTER, 0, tag, comm, status, ierr)
END IF

```

The last five characters of string `b` at process 1 are replaced by the first five characters of string `a` at process 0.

*Rationale.* The alternative choice would be for `MPI_CHARACTER` to match a character of arbitrary length. This runs into problems.

A Fortran character variable is a constant length string, with no special termination symbol. There is no fixed convention on how to represent characters, and how to store their length. Some compilers pass a character argument to a routine as a pair of arguments, one holding the address of the string and the other holding the length of string. Consider the case of an MPI communication call that is passed a communication buffer with type defined by a derived datatype (Section 4.1). If this communicator buffer contains variables of type `CHARACTER` then the information on their length will not be passed to the MPI routine.

This problem forces us to provide explicit information on character length with the MPI call. One could add a length parameter to the type `MPI_CHARACTER`, but this does not add much convenience and the same functionality can be achieved by defining a suitable derived datatype. (*End of rationale.*)

*Advice to implementors.* Some compilers pass Fortran `CHARACTER` arguments as a structure with a length and a pointer to the actual string. In such an environment, the MPI call needs to dereference the pointer in order to reach the string. (*End of advice to implementors.*)

### 3.3.2 Data Conversion

One of the goals of MPI is to support parallel computations across heterogeneous environments. Communication in a heterogeneous environment may require data conversions. We use the following terminology.

**type conversion** changes the datatype of a value, e.g., by rounding a `REAL` to an `INTEGER`.

**representation conversion** changes the binary representation of a value, e.g., from Hex floating point to IEEE floating point.

The type matching rules imply that MPI communication never entails type conversion. On the other hand, MPI requires that a representation conversion be performed when a typed value is transferred across environments that use different representations for the datatype of this value. MPI does not specify rules for representation conversion. Such



conversion is expected to preserve integer, logical or character values, and to convert a floating point value to the nearest value that can be represented on the target system.

Overflow and underflow exceptions may occur during floating point conversions. Conversion of integers or characters may also lead to exceptions when a value that can be represented in one system cannot be represented in the other system. An exception occurring during representation conversion results in a failure of the communication. An error occurs either in the send operation, or the receive operation, or both.

If a value sent in a message is untyped (i.e., of type `MPI_BYTE`), then the binary representation of the byte stored at the receiver is identical to the binary representation of the byte loaded at the sender. This holds true, whether sender and receiver run in the same or in distinct environments. No representation conversion is required. (Note that representation conversion may occur when values of type `MPI_CHARACTER` or `MPI_CHAR` are transferred, for example, from an EBCDIC encoding to an ASCII encoding.)

No conversion need occur when an MPI program executes in a homogeneous system, where all processes run in the same environment.

Consider the three examples, 3.2–3.4. The first program is correct, assuming that `a` and `b` are `REAL` arrays of size  $\geq 10$ . If the sender and receiver execute in different environments, then the ten real values that are fetched from the send buffer will be converted to the representation for reals on the receiver site before they are stored in the receive buffer. While the number of real elements fetched from the send buffer equal the number of real elements stored in the receive buffer, the number of bytes stored need not equal the number of bytes loaded. For example, the sender may use a four byte representation and the receiver an eight byte representation for reals.

The second program is erroneous, and its behavior is undefined.

The third program is correct. The exact same sequence of forty bytes that were loaded from the send buffer will be stored in the receive buffer, even if sender and receiver run in a different environment. The message sent has exactly the same length (in bytes) and the same binary representation as the message received. If `a` and `b` are of different types, or if they are of the same type but different data representations are used, then the bits stored in the receive buffer may encode values that are different from the values they encoded in the send buffer.

Data representation conversion also applies to the envelope of a message: source, destination and tag are all integers that may need to be converted.

*Advice to implementors.* The current definition does not require messages to carry data type information. Both sender and receiver provide complete data type information. In a heterogeneous environment, one can either use a machine independent encoding such as XDR, or have the receiver convert from the sender representation to its own, or even have the sender do the conversion.

Additional type information might be added to messages in order to allow the system to detect mismatches between datatype at sender and receiver. This might be particularly useful in a slower but safer debug mode. (*End of advice to implementors.*)

MPI requires support for inter-language communication, i.e., if messages are sent by a C or C++ process and received by a Fortran process, or vice-versa. The behavior is defined in Section 16.3 on page 529.

## 3.4 Communication Modes

The send call described in Section 3.2.1 is **blocking**: it does not return until the message data and envelope have been safely stored away so that the sender is free to modify the send buffer. The message might be copied directly into the matching receive buffer, or it might be copied into a temporary system buffer.

Message buffering decouples the send and receive operations. A blocking send can complete as soon as the message was buffered, even if no matching receive has been executed by the receiver. On the other hand, message buffering can be expensive, as it entails additional memory-to-memory copying, and it requires the allocation of memory for buffering. MPI offers the choice of several communication modes that allow one to control the choice of the communication protocol.

The send call described in Section 3.2.1 uses the **standard** communication mode. In this mode, it is up to MPI to decide whether outgoing messages will be buffered. MPI may buffer outgoing messages. In such a case, the send call may complete before a matching receive is invoked. On the other hand, buffer space may be unavailable, or MPI may choose not to buffer outgoing messages, for performance reasons. In this case, the send call will not complete until a matching receive has been posted, and the data has been moved to the receiver.

Thus, a send in standard mode can be started whether or not a matching receive has been posted. It may complete before a matching receive is posted. The standard mode send is **non-local**: successful completion of the send operation may depend on the occurrence of a matching receive.

*Rationale.* The reluctance of MPI to mandate whether standard sends are buffering or not stems from the desire to achieve portable programs. Since any system will run out of buffer resources as message sizes are increased, and some implementations may want to provide little buffering, MPI takes the position that correct (and therefore, portable) programs do not rely on system buffering in standard mode. Buffering may improve the performance of a correct program, but it doesn't affect the result of the program. If the user wishes to guarantee a certain amount of buffering, the user-provided buffer system of Section 3.6 should be used, along with the buffered-mode send. (*End of rationale.*)

There are three additional communication modes.

A **buffered** mode send operation can be started whether or not a matching receive has been posted. It may complete before a matching receive is posted. However, unlike the standard send, this operation is **local**, and its completion does not depend on the occurrence of a matching receive. Thus, if a send is executed and no matching receive is posted, then MPI must buffer the outgoing message, so as to allow the send call to complete. An error will occur if there is insufficient buffer space. The amount of available buffer space is controlled by the user — see Section 3.6. Buffer allocation by the user may be required for the buffered mode to be effective.

A send that uses the **synchronous** mode can be started whether or not a matching receive was posted. However, the send will complete successfully only if a matching receive is posted, and the receive operation has started to receive the message sent by the synchronous send. Thus, the completion of a synchronous send not only indicates that the send buffer can be reused, but it also indicates that the receiver has reached a certain point in its

execution, namely that it has started executing the matching receive. If both sends and receives are blocking operations then the use of the synchronous mode provides synchronous communication semantics: a communication does not complete at either end before both processes rendezvous at the communication. A send executed in this mode is **non-local**.

A send that uses the **ready** communication mode may be started *only* if the matching receive is already posted. Otherwise, the operation is erroneous and its outcome is undefined. On some systems, this allows the removal of a hand-shake operation that is otherwise required and results in improved performance. The completion of the send operation does not depend on the status of a matching receive, and merely indicates that the send buffer can be reused. A send operation that uses the ready mode has the same semantics as a standard send operation, or a synchronous send operation; it is merely that the sender provides additional information to the system (namely that a matching receive is already posted), that can save some overhead. In a correct program, therefore, a ready send could be replaced by a standard send with no effect on the behavior of the program other than performance.

Three additional send functions are provided for the three additional communication modes. The communication mode is indicated by a one letter prefix: B for buffered, S for synchronous, and R for ready.

MPI\_BSEND (buf, count, datatype, dest, tag, comm)

IN	buf	initial address of send buffer (choice)
IN	count	number of elements in send buffer (non-negative integer)
IN	datatype	datatype of each send buffer element (handle)
IN	dest	rank of destination (integer)
IN	tag	message tag (integer)
IN	comm	communicator (handle)

```
int MPI_Bsend(void* buf, int count, MPI_Datatype datatype, int dest,
              int tag, MPI_Comm comm)
```

```
MPI_BSEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, IERROR)
```

```
<type> BUF(*)
```

```
INTEGER COUNT, DATATYPE, DEST, TAG, COMM, IERROR
```

```
{void MPI::Comm::Bsend(const void* buf, int count, const
                       MPI::Datatype& datatype, int dest, int tag) const(binding
                       deprecated, see Section 15.2) }
```

Send in buffered mode.

MPI\_SSEND (buf, count, datatype, dest, tag, comm)

IN	buf	initial address of send buffer (choice)
IN	count	number of elements in send buffer (non-negative integer)
IN	datatype	datatype of each send buffer element (handle)
IN	dest	rank of destination (integer)
IN	tag	message tag (integer)
IN	comm	communicator (handle)

```
int MPI_Ssend(void* buf, int count, MPI_Datatype datatype, int dest,
              int tag, MPI_Comm comm)
```

```
MPI_SSEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, IERROR)
<type> BUF(*)
INTEGER COUNT, DATATYPE, DEST, TAG, COMM, IERROR
```

```
{void MPI::Comm::Ssend(const void* buf, int count, const
                       MPI::Datatype& datatype, int dest, int tag) const(binding
depreciated, see Section 15.2) }
```

Send in synchronous mode.

MPI\_RSEND (buf, count, datatype, dest, tag, comm)

IN	buf	initial address of send buffer (choice)
IN	count	number of elements in send buffer (non-negative integer)
IN	datatype	datatype of each send buffer element (handle)
IN	dest	rank of destination (integer)
IN	tag	message tag (integer)
IN	comm	communicator (handle)

```
int MPI_Rsend(void* buf, int count, MPI_Datatype datatype, int dest,
              int tag, MPI_Comm comm)
```

```
MPI_RSEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, IERROR)
<type> BUF(*)
INTEGER COUNT, DATATYPE, DEST, TAG, COMM, IERROR
```

```
{void MPI::Comm::Rsend(const void* buf, int count, const
                       MPI::Datatype& datatype, int dest, int tag) const(binding
depreciated, see Section 15.2) }
```

Send in ready mode.

There is only one receive operation, but it matches any of the send modes. The receive operation described in the last section is **blocking**: it returns only after the receive buffer

contains the newly received message. A receive can complete before the matching send has completed (of course, it can complete only after the matching send has started).

In a multi-threaded implementation of MPI, the system may de-schedule a thread that is blocked on a send or receive operation, and schedule another thread for execution in the same address space. In such a case it is the user's responsibility not to modify a communication buffer until the communication completes. Otherwise, the outcome of the computation is undefined.

*Advice to implementors.* Since a synchronous send cannot complete before a matching receive is posted, one will not normally buffer messages sent by such an operation.

It is recommended to choose buffering over blocking the sender, whenever possible, for standard sends. The programmer can signal his or her preference for blocking the sender until a matching receive occurs by using the synchronous send mode.

A possible communication protocol for the various communication modes is outlined below.

**ready send:** The message is sent as soon as possible.

**synchronous send:** The sender sends a request-to-send message. The receiver stores this request. When a matching receive is posted, the receiver sends back a permission-to-send message, and the sender now sends the message.

**standard send:** First protocol may be used for short messages, and second protocol for long messages.

**buffered send:** The sender copies the message into a buffer and then sends it with a nonblocking send (using the same protocol as for standard send).

Additional control messages might be needed for flow control and error recovery. Of course, there are many other possible protocols.

Ready send can be implemented as a standard send. In this case there will be no performance advantage (or disadvantage) for the use of ready send.

A standard send can be implemented as a synchronous send. In such a case, no data buffering is needed. However, users may expect some buffering.

In a multi-threaded environment, the execution of a blocking communication should block only the executing thread, allowing the thread scheduler to de-schedule this thread and schedule another thread for execution. (*End of advice to implementors.*)

### 3.5 Semantics of Point-to-Point Communication

A valid MPI implementation guarantees certain general properties of point-to-point communication, which are described in this section.

**Order** Messages are *non-overtaking*: If a sender sends two messages in succession to the same destination, and both match the same receive, then this operation cannot receive the second message if the first one is still pending. If a receiver posts two receives in succession, and both match the same message, then the second receive operation cannot be satisfied by this message, if the first one is still pending. This requirement facilitates matching of sends to receives. It guarantees that message-passing code is deterministic, if processes are

single-threaded and the wildcard `MPI_ANY_SOURCE` is not used in receives. (Some of the calls described later, such as `MPI_CANCEL` or `MPI_WAITANY`, are additional sources of nondeterminism.)

If a process has a single thread of execution, then any two communications executed by this process are ordered. On the other hand, if the process is multi-threaded, then the semantics of thread execution may not define a relative order between two send operations executed by two distinct threads. The operations are logically concurrent, even if one physically precedes the other. In such a case, the two messages sent can be received in any order. Similarly, if two receive operations that are logically concurrent receive two successively sent messages, then the two messages can match the two receives in either order.

**Example 3.6** An example of non-overtaking messages.

```
CALL MPI_COMM_RANK(comm, rank, ierr)
IF (rank.EQ.0) THEN
    CALL MPI_BSEND(buf1, count, MPI_REAL, 1, tag, comm, ierr)
    CALL MPI_BSEND(buf2, count, MPI_REAL, 1, tag, comm, ierr)
ELSE IF (rank.EQ.1) THEN
    CALL MPI_RECV(buf1, count, MPI_REAL, 0, MPI_ANY_TAG, comm, status, ierr)
    CALL MPI_RECV(buf2, count, MPI_REAL, 0, tag, comm, status, ierr)
END IF
```

The message sent by the first send must be received by the first receive, and the message sent by the second send must be received by the second receive.

**Progress** If a pair of matching send and receives have been initiated on two processes, then at least one of these two operations will complete, independently of other actions in the system: the send operation will complete, unless the receive is satisfied by another message, and completes; the receive operation will complete, unless the message sent is consumed by another matching receive that was posted at the same destination process.

**Example 3.7** An example of two, intertwined matching pairs.

```
CALL MPI_COMM_RANK(comm, rank, ierr)
IF (rank.EQ.0) THEN
    CALL MPI_BSEND(buf1, count, MPI_REAL, 1, tag1, comm, ierr)
    CALL MPI_SSEND(buf2, count, MPI_REAL, 1, tag2, comm, ierr)
ELSE IF (rank.EQ.1) THEN
    CALL MPI_RECV(buf1, count, MPI_REAL, 0, tag2, comm, status, ierr)
    CALL MPI_RECV(buf2, count, MPI_REAL, 0, tag1, comm, status, ierr)
END IF
```

Both processes invoke their first communication call. Since the first send of process zero uses the buffered mode, it must complete, irrespective of the state of process one. Since no matching receive is posted, the message will be copied into buffer space. (If insufficient buffer space is available, then the program will fail.) The second send is then invoked. At that point, a matching pair of send and receive operation is enabled, and both operations must complete. Process one next invokes its second receive call, which will be satisfied by

the buffered message. Note that process one received the messages in the reverse order they were sent.

**Fairness** MPI makes no guarantee of *fairness* in the handling of communication. Suppose that a send is posted. Then it is possible that the destination process repeatedly posts a receive that matches this send, yet the message is never received, because it is each time overtaken by another message, sent from another source. Similarly, suppose that a receive was posted by a multi-threaded process. Then it is possible that messages that match this receive are repeatedly received, yet the receive is never satisfied, because it is overtaken by other receives posted at this node (by other executing threads). It is the programmer's responsibility to prevent starvation in such situations.

**Resource limitations** Any pending communication operation consumes system resources that are limited. Errors may occur when lack of resources prevent the execution of an MPI call. A quality implementation will use a (small) fixed amount of resources for each pending send in the ready or synchronous mode and for each pending receive. However, buffer space may be consumed to store messages sent in standard mode, and must be consumed to store messages sent in buffered mode, when no matching receive is available. The amount of space available for buffering will be much smaller than program data memory on many systems. Then, it will be easy to write programs that overrun available buffer space.

MPI allows the user to provide buffer memory for messages sent in the buffered mode. Furthermore, MPI specifies a detailed operational model for the use of this buffer. An MPI implementation is required to do no worse than implied by this model. This allows users to avoid buffer overflows when they use buffered sends. Buffer allocation and use is described in Section 3.6.

A buffered send operation that cannot complete because of a lack of buffer space is erroneous. When such a situation is detected, an error is signalled that may cause the program to terminate abnormally. On the other hand, a standard send operation that cannot complete because of lack of buffer space will merely block, waiting for buffer space to become available or for a matching receive to be posted. This behavior is preferable in many situations. Consider a situation where a producer repeatedly produces new values and sends them to a consumer. Assume that the producer produces new values faster than the consumer can consume them. If buffered sends are used, then a buffer overflow will result. Additional synchronization has to be added to the program so as to prevent this from occurring. If standard sends are used, then the producer will be automatically throttled, as its send operations will block when buffer space is unavailable.

In some situations, a lack of buffer space leads to deadlock situations. This is illustrated by the examples below.

**Example 3.8** An exchange of messages.

```
CALL MPI_COMM_RANK(comm, rank, ierr)
IF (rank.EQ.0) THEN
    CALL MPI_SEND(sendbuf, count, MPI_REAL, 1, tag, comm, ierr)
    CALL MPI_RECV(recvbuf, count, MPI_REAL, 1, tag, comm, status, ierr)
ELSE IF (rank.EQ.1) THEN
    CALL MPI_RECV(recvbuf, count, MPI_REAL, 0, tag, comm, status, ierr)
    CALL MPI_SEND(sendbuf, count, MPI_REAL, 0, tag, comm, ierr)
```

END IF

This program will succeed even if no buffer space for data is available. The standard send operation can be replaced, in this example, with a synchronous send.

**Example 3.9** An errant attempt to exchange messages.

```
CALL MPI_COMM_RANK(comm, rank, ierr)
IF (rank.EQ.0) THEN
    CALL MPI_RECV(recvbuf, count, MPI_REAL, 1, tag, comm, status, ierr)
    CALL MPI_SEND(sendbuf, count, MPI_REAL, 1, tag, comm, ierr)
ELSE IF (rank.EQ.1) THEN
    CALL MPI_RECV(recvbuf, count, MPI_REAL, 0, tag, comm, status, ierr)
    CALL MPI_SEND(sendbuf, count, MPI_REAL, 0, tag, comm, ierr)
END IF
```

The receive operation of the first process must complete before its send, and can complete only if the matching send of the second processor is executed. The receive operation of the second process must complete before its send and can complete only if the matching send of the first process is executed. This program will always deadlock. The same holds for any other send mode.

**Example 3.10** An exchange that relies on buffering.

```
CALL MPI_COMM_RANK(comm, rank, ierr)
IF (rank.EQ.0) THEN
    CALL MPI_SEND(sendbuf, count, MPI_REAL, 1, tag, comm, ierr)
    CALL MPI_RECV(recvbuf, count, MPI_REAL, 1, tag, comm, status, ierr)
ELSE IF (rank.EQ.1) THEN
    CALL MPI_SEND(sendbuf, count, MPI_REAL, 0, tag, comm, ierr)
    CALL MPI_RECV(recvbuf, count, MPI_REAL, 0, tag, comm, status, ierr)
END IF
```

The message sent by each process has to be copied out before the send operation returns and the receive operation starts. For the program to complete, it is necessary that at least one of the two messages sent be buffered. Thus, this program can succeed only if the communication system can buffer at least `count` words of data.

*Advice to users.* When standard send operations are used, then a deadlock situation may occur where both processes are blocked because buffer space is not available. The same will certainly happen, if the synchronous mode is used. If the buffered mode is used, and not enough buffer space is available, then the program will not complete either. However, rather than a deadlock situation, we shall have a buffer overflow error.

A program is “safe” if no message buffering is required for the program to complete. One can replace all sends in such program with synchronous sends, and the program will still run correctly. This conservative programming style provides the best portability, since program completion does not depend on the amount of buffer space available or on the communication protocol used.



Many programmers prefer to have more leeway and opt to use the “unsafe” programming style shown in Example 3.10. In such cases, the use of standard sends is likely to provide the best compromise between performance and robustness: quality implementations will provide sufficient buffering so that “common practice” programs will not deadlock. The buffered send mode can be used for programs that require more buffering, or in situations where the programmer wants more control. This mode might also be used for debugging purposes, as buffer overflow conditions are easier to diagnose than deadlock conditions.

Nonblocking message-passing operations, as described in Section 3.7, can be used to avoid the need for buffering outgoing messages. This prevents deadlocks due to lack of buffer space, and improves performance, by allowing overlap of computation and communication, and avoiding the overheads of allocating buffers and copying messages into buffers. (*End of advice to users.*)

### 3.6 Buffer Allocation and Usage

A user may specify a buffer to be used for buffering messages sent in buffered mode. Buffering is done by the sender.

`MPI_BUFFER_ATTACH(buffer, size)`

IN	buffer	initial buffer address (choice)
IN	size	buffer size, in bytes (non-negative integer)

`int MPI_Buffer_attach(void* buffer, int size)`

`MPI_BUFFER_ATTACH(BUFFER, SIZE, IERROR)`

`<type> BUFFER(*)`

`INTEGER SIZE, IERROR`

`{void MPI::Attach_buffer(void* buffer, int size) (binding deprecated, see Section 15.2) }`

Provides to MPI a buffer in the user’s memory to be used for buffering outgoing messages. The buffer is used only by messages sent in buffered mode. Only one buffer can be attached to a process at a time.

`MPI_BUFFER_DETACH(buffer_addr, size)`

OUT	buffer_addr	initial buffer address (choice)
OUT	size	buffer size, in bytes (non-negative integer)

`int MPI_Buffer_detach(void* buffer_addr, int* size)`

`MPI_BUFFER_DETACH(BUFFER_ADDR, SIZE, IERROR)`

`<type> BUFFER_ADDR(*)`

`INTEGER SIZE, IERROR`

```
{int MPI::Detach_buffer(void*& buffer) (binding deprecated, see Section 15.2) }
```

Detach the buffer currently associated with MPI. The call returns the address and the size of the detached buffer. This operation will block until all messages currently in the buffer have been transmitted. Upon return of this function, the user may reuse or deallocate the space taken by the buffer.

**Example 3.11** Calls to attach and detach buffers.

```
#define BUFFSIZE 10000
int size;
char *buff;
MPI_Buffer_attach( malloc(BUFFSIZE), BUFFSIZE);
/* a buffer of 10000 bytes can now be used by MPI_Bsend */
MPI_Buffer_detach( &buff, &size);
/* Buffer size reduced to zero */
MPI_Buffer_attach( buff, size);
/* Buffer of 10000 bytes available again */
```

*Advice to users.* Even though the C functions `MPI_Buffer_attach` and `MPI_Buffer_detach` both have a first argument of type `void*`, these arguments are used differently: A pointer to the buffer is passed to `MPI_Buffer_attach`; the address of the pointer is passed to `MPI_Buffer_detach`, so that this call can return the pointer value. (*End of advice to users.*)

*Rationale.* Both arguments are defined to be of type `void*` (rather than `void*` and `void**`, respectively), so as to avoid complex type casts. E.g., in the last example, `&buff`, which is of type `char**`, can be passed as argument to `MPI_Buffer_detach` without type casting. If the formal parameter had type `void**` then we would need a type cast before and after the call. (*End of rationale.*)

The statements made in this section describe the behavior of MPI for buffered-mode sends. When no buffer is currently associated, MPI behaves as if a zero-sized buffer is associated with the process.

MPI must provide as much buffering for outgoing messages *as if* outgoing message data were buffered by the sending process, in the specified buffer space, using a circular, contiguous-space allocation policy. We outline below a model implementation that defines this policy. MPI may provide more buffering, and may use a better buffer allocation algorithm than described below. On the other hand, MPI may signal an error whenever the simple buffering allocator described below would run out of space. In particular, if no buffer is explicitly associated with the process, then any buffered send may cause an error.

MPI does not provide mechanisms for querying or controlling buffering done by standard mode sends. It is expected that vendors will provide such information for their implementations.

*Rationale.* There is a wide spectrum of possible implementations of buffered communication: buffering can be done at sender, at receiver, or both; buffers can be dedicated to one sender-receiver pair, or be shared by all communications; buffering can be done in real or in virtual memory; it can use dedicated memory, or memory

shared by other processes; buffer space may be allocated statically or be changed dynamically; etc. It does not seem feasible to provide a portable mechanism for querying or controlling buffering that would be compatible with all these choices, yet provide meaningful information. (*End of rationale.*)

### 3.6.1 Model Implementation of Buffered Mode

The model implementation uses the packing and unpacking functions described in Section 4.2 and the nonblocking communication functions described in Section 3.7.

We assume that a circular queue of pending message entries (PME) is maintained. Each entry contains a communication request handle that identifies a pending nonblocking send, a pointer to the next entry and the packed message data. The entries are stored in successive locations in the buffer. Free space is available between the queue tail and the queue head.

A buffered send call results in the execution of the following code.

- Traverse sequentially the PME queue from head towards the tail, deleting all entries for communications that have completed, up to the first entry with an uncompleted request; update queue head to point to that entry.
- Compute the number, *n*, of bytes needed to store an entry for the new message. An upper bound on *n* can be computed as follows: A call to the function `MPI_PACK_SIZE(count, datatype, comm, size)`, with the `count`, `datatype` and `comm` arguments used in the `MPI_BSEND` call, returns an upper bound on the amount of space needed to buffer the message data (see Section 4.2). The MPI constant `MPI_BSEND_OVERHEAD` provides an upper bound on the additional space consumed by the entry (e.g., for pointers or envelope information).
- Find the next contiguous empty space of *n* bytes in buffer (space following queue tail, or space at start of buffer if queue tail is too close to end of buffer). If space is not found then raise buffer overflow error.
- Append to end of PME queue in contiguous space the new entry that contains request handle, next pointer and packed message data; `MPI_PACK` is used to pack data.
- Post nonblocking send (standard mode) for packed data.
- Return

## 3.7 Nonblocking Communication

One can improve performance on many systems by overlapping communication and computation. This is especially true on systems where communication can be executed autonomously by an intelligent communication controller. Light-weight threads are one mechanism for achieving such overlap. An alternative mechanism that often leads to better performance is to use **nonblocking communication**. A nonblocking **send start** call initiates the send operation, but does not complete it. The send start call can return before the message was copied out of the send buffer. A separate **send complete** call is needed to complete the communication, i.e., to verify that the data has been copied out of the send buffer. With suitable hardware, the transfer of data out of the sender memory may proceed

concurrently with computations done at the sender after the send was initiated and before it completed. Similarly, a nonblocking **receive start call** initiates the receive operation, but does not complete it. The call can return before a message is stored into the receive buffer. A separate **receive complete** call is needed to complete the receive operation and verify that the data has been received into the receive buffer. With suitable hardware, the transfer of data into the receiver memory may proceed concurrently with computations done after the receive was initiated and before it completed. The use of nonblocking receives may also avoid system buffering and memory-to-memory copying, as information is provided early on the location of the receive buffer.

Nonblocking send start calls can use the same four modes as blocking sends: **standard**, **buffered**, **synchronous** and **ready**. These carry the same meaning. Sends of all modes, **ready** excepted, can be started whether a matching receive has been posted or not; a nonblocking **ready** send can be started only if a matching receive is posted. In all cases, the send start call is local: it returns immediately, irrespective of the status of other processes. If the call causes some system resource to be exhausted, then it will fail and return an error code. Quality implementations of MPI should ensure that this happens only in “pathological” cases. That is, an MPI implementation should be able to support a large number of pending nonblocking operations.

The send-complete call returns when data has been copied out of the send buffer. It may carry additional meaning, depending on the send mode.

If the send mode is **synchronous**, then the send can complete only if a matching receive has started. That is, a receive has been posted, and has been matched with the send. In this case, the send-complete call is non-local. Note that a synchronous, nonblocking send may complete, if matched by a nonblocking receive, before the receive complete call occurs. (It can complete as soon as the sender “knows” the transfer will complete, but before the receiver “knows” the transfer will complete.)

If the send mode is **buffered** then the message must be buffered if there is no pending receive. In this case, the send-complete call is local, and must succeed irrespective of the status of a matching receive.

If the send mode is **standard** then the send-complete call may return before a matching receive is posted, if the message is buffered. On the other hand, the send-complete may not complete until a matching receive is posted, and the message was copied into the receive buffer.

Nonblocking sends can be matched with blocking receives, and vice-versa.

*Advice to users.* The completion of a send operation may be delayed, for standard mode, and must be delayed, for synchronous mode, until a matching receive is posted. The use of nonblocking sends in these two cases allows the sender to proceed ahead of the receiver, so that the computation is more tolerant of fluctuations in the speeds of the two processes.

Nonblocking sends in the buffered and ready modes have a more limited impact, e.g., the blocking version of buffered send is capable of completing regardless of when a matching receive call is made. However, separating the start from the completion of these sends still gives some opportunity for optimization within the MPI library. For example, starting a buffered send gives an implementation more flexibility in determining if and how the message is buffered. There are also advantages for both nonblocking buffered and ready modes when data copying can be done concurrently with computation.

The message-passing model implies that communication is initiated by the sender. The communication will generally have lower overhead if a receive is already posted when the sender initiates the communication (data can be moved directly to the receive buffer, and there is no need to queue a pending send request). However, a receive operation can complete only after the matching send has occurred. The use of nonblocking receives allows one to achieve lower communication overheads without blocking the receiver while it waits for the send. (*End of advice to users.*)

### 3.7.1 Communication Request Objects

Nonblocking communications use opaque **request** objects to identify communication operations and match the operation that initiates the communication with the operation that terminates it. These are system objects that are accessed via a handle. A request object identifies various properties of a communication operation, such as the send mode, the communication buffer that is associated with it, its context, the tag and destination arguments to be used for a send, or the tag and source arguments to be used for a receive. In addition, this object stores information about the status of the pending communication operation.

### 3.7.2 Communication Initiation

We use the same naming conventions as for blocking communication: a prefix of B, S, or R is used for buffered, synchronous or ready mode. In addition a prefix of l (for immediate) indicates that the call is nonblocking.

**MPI\_ISEND**(buf, count, datatype, dest, tag, comm, request)

IN	buf	initial address of send buffer (choice)
IN	count	number of elements in send buffer (non-negative integer)
IN	datatype	datatype of each send buffer element (handle)
IN	dest	rank of destination (integer)
IN	tag	message tag (integer)
IN	comm	communicator (handle)
OUT	request	communication request (handle)

```
int MPI_Isend(void* buf, int count, MPI_Datatype datatype, int dest,
              int tag, MPI_Comm comm, MPI_Request *request)
```

```
MPI_ISEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
```

```
<type> BUF(*)
```

```
INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR
```

```
{MPI::Request MPI::Comm::Isend(const void* buf, int count, const
                               MPI::Datatype& datatype, int dest, int tag) const(binding
                               deprecated, see Section 15.2) }
```

Start a standard mode, nonblocking send.

```

MPI_IBSEND(buf, count, datatype, dest, tag, comm, request) 1
    IN      buf      initial address of send buffer (choice) 2
    IN      count    number of elements in send buffer (non-negative inte- 3
                        ger) 4
    IN      datatype  datatype of each send buffer element (handle) 5
    IN      dest      rank of destination (integer) 6
    IN      tag        message tag (integer) 7
    IN      comm       communicator (handle) 8
    OUT     request    communication request (handle) 9
10
11
12
13
int MPI_Ibsend(void* buf, int count, MPI_Datatype datatype, int dest, 14
               int tag, MPI_Comm comm, MPI_Request *request) 15
MPI_IBSEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR) 16
    <type> BUF(*) 17
    INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR 18
19
{MPI::Request MPI::Comm::Ibsend(const void* buf, int count, const 20
    MPI::Datatype& datatype, int dest, int tag) const(binding 21
    deprecated, see Section 15.2) } 22
23
    Start a buffered mode, nonblocking send. 24
25
MPI_ISSEND(buf, count, datatype, dest, tag, comm, request) 26
    IN      buf      initial address of send buffer (choice) 27
    IN      count    number of elements in send buffer (non-negative inte- 28
                        ger) 29
    IN      datatype  datatype of each send buffer element (handle) 30
    IN      dest      rank of destination (integer) 31
    IN      tag        message tag (integer) 32
    IN      comm       communicator (handle) 33
    OUT     request    communication request (handle) 34
35
36
37
38
int MPI_Issend(void* buf, int count, MPI_Datatype datatype, int dest, 39
               int tag, MPI_Comm comm, MPI_Request *request) 40
MPI_ISSEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR) 41
    <type> BUF(*) 42
    INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR 43
44
{MPI::Request MPI::Comm::Issend(const void* buf, int count, const 45
    MPI::Datatype& datatype, int dest, int tag) const(binding 46
    deprecated, see Section 15.2) } 47
48
    Start a synchronous mode, nonblocking send.

```

```

1 MPI_IRSEND(buf, count, datatype, dest, tag, comm, request)
2     IN      buf                      initial address of send buffer (choice)
3
4     IN      count                    number of elements in send buffer (non-negative integer)
5
6     IN      datatype                 datatype of each send buffer element (handle)
7
8     IN      dest                     rank of destination (integer)
9
10    IN      tag                      message tag (integer)
11
12    IN      comm                     communicator (handle)
13
14    OUT     request                   communication request (handle)

```

```

14 int MPI_Irsend(void* buf, int count, MPI_Datatype datatype, int dest,
15               int tag, MPI_Comm comm, MPI_Request *request)

```

```

16 MPI_IRSEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)

```

```

17     <type> BUF(*)

```

```

18     INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR

```

```

19 {MPI::Request MPI::Comm::Irsend(const void* buf, int count, const
20                               MPI::Datatype& datatype, int dest, int tag) const(binding
21                               deprecated, see Section 15.2) }

```

Start a ready mode nonblocking send.

```

26 MPI_IRECV (buf, count, datatype, source, tag, comm, request)
27     OUT     buf                      initial address of receive buffer (choice)
28
29     IN      count                    number of elements in receive buffer (non-negative integer)
30
31     IN      datatype                 datatype of each receive buffer element (handle)
32
33     IN      source                   rank of source or MPI_ANY_SOURCE (integer)
34
35     IN      tag                      message tag or MPI_ANY_TAG (integer)
36
37     IN      comm                     communicator (handle)
38
39     OUT     request                   communication request (handle)

```

```

38 int MPI_Irecv(void* buf, int count, MPI_Datatype datatype, int source,
39              int tag, MPI_Comm comm, MPI_Request *request)

```

```

40 MPI_IRECV(BUF, COUNT, DATATYPE, SOURCE, TAG, COMM, REQUEST, IERROR)

```

```

41     <type> BUF(*)

```

```

42     INTEGER COUNT, DATATYPE, SOURCE, TAG, COMM, REQUEST, IERROR

```

```

43 {MPI::Request MPI::Comm::Irecv(void* buf, int count, const
44                               MPI::Datatype& datatype, int source, int tag) const(binding
45                               deprecated, see Section 15.2) }

```

Start a nonblocking receive.

These calls allocate a communication request object and associate it with the request handle (the argument `request`). The request can be used later to query the status of the communication or wait for its completion.

A nonblocking send call indicates that the system may start copying data out of the send buffer. The sender should not modify any part of the send buffer after a nonblocking send operation is called, until the send completes.

A nonblocking receive call indicates that the system may start writing data into the receive buffer. The receiver should not access any part of the receive buffer after a nonblocking receive operation is called, until the receive completes.

*Advice to users.* To prevent problems with the argument copying and register optimization done by Fortran compilers, please note the hints in subsections “Problems Due to Data Copying and Sequence Association,” and “A Problem with Register Optimization” in Section 16.2.2 on pages 514 and 517. (*End of advice to users.*)

### 3.7.3 Communication Completion

The functions `MPI_WAIT` and `MPI_TEST` are used to complete a nonblocking communication. The completion of a send operation indicates that the sender is now free to update the locations in the send buffer (the send operation itself leaves the content of the send buffer unchanged). It does not indicate that the message has been received, rather, it may have been buffered by the communication subsystem. However, if a **synchronous** mode send was used, the completion of the send operation indicates that a matching receive was initiated, and that the message will eventually be received by this matching receive.

The completion of a receive operation indicates that the receive buffer contains the received message, the receiver is now free to access it, and that the status object is set. It does not indicate that the matching send operation has completed (but indicates, of course, that the send was initiated).

We shall use the following terminology: A **null** handle is a handle with value `MPI_REQUEST_NULL`. A persistent request and the handle to it are **inactive** if the request is not associated with any ongoing communication (see Section 3.9). A handle is **active** if it is neither null nor inactive. An **empty** status is a status which is set to return `tag = MPI_ANY_TAG`, `source = MPI_ANY_SOURCE`, `error = MPI_SUCCESS`, and is also internally configured so that calls to `MPI_GET_COUNT` and `MPI_GET_ELEMENTS` return `count = 0` and `MPI_TEST_CANCELLED` returns false. We set a status variable to empty when the value returned by it is not significant. Status is set in this way so as to prevent errors due to accesses of stale information.

The fields in a **status** object returned by a call to `MPI_WAIT`, `MPI_TEST`, or any of the other derived functions (`MPI_{TEST|WAIT}{ALL|SOME|ANY}`), where the **request** corresponds to a send call, are undefined, with two exceptions: The error status field will contain valid information if the wait or test call returned with `MPI_ERR_IN_STATUS`; and the returned status can be queried by the call `MPI_TEST_CANCELLED`.

Error codes belonging to the error class `MPI_ERR_IN_STATUS` should be returned only by the MPI completion functions that take arrays of `MPI_STATUS`. For the functions `MPI_TEST`, `MPI_TESTANY`, `MPI_WAIT`, and `MPI_WAITANY`, which return a single `MPI_STATUS` value, the normal MPI error return process should be used (not the `MPI_ERROR` field in the `MPI_STATUS` argument).



```
1 MPI_WAIT(request, status)
```

```
2     INOUT    request                request (handle)
3
4     OUT     status                status object (Status)
```

```
5
6 int MPI_Wait(MPI_Request *request, MPI_Status *status)
```

```
7 MPI_WAIT(REQUEST, STATUS, IERROR)
8     INTEGER REQUEST, STATUS(MPI_STATUS_SIZE), IERROR
```

```
9
10 {void MPI::Request::Wait(MPI::Status& status) (binding deprecated, see
11     Section 15.2) }
```

```
12
13 {void MPI::Request::Wait() (binding deprecated, see Section 15.2) }
```

14 A call to MPI\_WAIT returns when the operation identified by `request` is complete. If  
15 the communication object associated with this request was created by a nonblocking send  
16 or receive call, then the object is deallocated by the call to MPI\_WAIT and the request  
17 handle is set to MPI\_REQUEST\_NULL. MPI\_WAIT is a non-local operation.

18 The call returns, in `status`, information on the completed operation. The content of  
19 the status object for a receive operation can be accessed as described in Section 3.2.5. The  
20 status object for a send operation may be queried by a call to MPI\_TEST\_CANCELLED  
21 (see Section 3.8).

22 One is allowed to call MPI\_WAIT with a null or inactive `request` argument. In this case  
23 the operation returns immediately with empty `status`.

24  
25 *Advice to users.* Successful return of MPI\_WAIT after a MPI\_IBSEND implies that  
26 the user send buffer can be reused — i.e., data has been sent out or copied into  
27 a buffer attached with MPI\_BUFFER\_ATTACH. Note that, at this point, we can no  
28 longer cancel the send (see Section 3.8). If a matching receive is never posted, then the  
29 buffer cannot be freed. This runs somewhat counter to the stated goal of MPI\_CANCEL  
30 (always being able to free program space that was committed to the communication  
31 subsystem). (*End of advice to users.*)

32  
33 *Advice to implementors.* In a multi-threaded environment, a call to MPI\_WAIT should  
34 block only the calling thread, allowing the thread scheduler to schedule another thread  
35 for execution. (*End of advice to implementors.*)

```
36
37
38 MPI_TEST(request, flag, status)
```

```
39     INOUT    request                communication request (handle)
40
41     OUT     flag                    true if operation completed (logical)
42
43     OUT     status                status object (Status)
```

```
44
45 int MPI_Test(MPI_Request *request, int *flag, MPI_Status *status)
```

```
46 MPI_TEST(REQUEST, FLAG, STATUS, IERROR)
47     LOGICAL FLAG
48     INTEGER REQUEST, STATUS(MPI_STATUS_SIZE), IERROR
```

```
{bool MPI::Request::Test(MPI::Status& status) (binding deprecated, see
    Section 15.2) }
```

```
{bool MPI::Request::Test() (binding deprecated, see Section 15.2) }
```

A call to `MPI_TEST` returns `flag = true` if the operation identified by `request` is complete. In such a case, the status object is set to contain information on the completed operation; if the communication object was created by a nonblocking send or receive, then it is deallocated and the request handle is set to `MPI_REQUEST_NULL`. The call returns `flag = false`, otherwise. In this case, the value of the status object is undefined. `MPI_TEST` is a local operation.

The return status object for a receive operation carries information that can be accessed as described in Section 3.2.5. The status object for a send operation carries information that can be accessed by a call to `MPI_TEST_CANCELLED` (see Section 3.8).

One is allowed to call `MPI_TEST` with a null or inactive `request` argument. In such a case the operation returns with `flag = true` and empty `status`.

The functions `MPI_WAIT` and `MPI_TEST` can be used to complete both sends and receives.

*Advice to users.* The use of the nonblocking `MPI_TEST` call allows the user to schedule alternative activities within a single thread of execution. An event-driven thread scheduler can be emulated with periodic calls to `MPI_TEST`. (*End of advice to users.*)

**Example 3.12** Simple usage of nonblocking operations and `MPI_WAIT`.

```
CALL MPI_COMM_RANK(comm, rank, ierr)
IF (rank.EQ.0) THEN
    CALL MPI_ISEND(a(1), 10, MPI_REAL, 1, tag, comm, request, ierr)
    **** do some computation to mask latency ****
    CALL MPI_WAIT(request, status, ierr)
ELSE IF (rank.EQ.1) THEN
    CALL MPI_IRECV(a(1), 15, MPI_REAL, 0, tag, comm, request, ierr)
    **** do some computation to mask latency ****
    CALL MPI_WAIT(request, status, ierr)
END IF
```

A request object can be deallocated without waiting for the associated communication to complete, by using the following operation.

```
MPI_REQUEST_FREE(request)
```

```
INOUT    request                communication request (handle)
```

```
int MPI_Request_free(MPI_Request *request)
```

```
MPI_REQUEST_FREE(REQUEST, IERROR)
```

```
INTEGER REQUEST, IERROR
```

```
{void MPI::Request::Free() (binding deprecated, see Section 15.2) }
```

Mark the request object for deallocation and set `request` to `MPI_REQUEST_NULL`. An ongoing communication that is associated with the request will be allowed to complete. The request will be deallocated only after its completion.

*Rationale.* The `MPI_REQUEST_FREE` mechanism is provided for reasons of performance and convenience on the sending side. (*End of rationale.*)

*Advice to users.* Once a request is freed by a call to `MPI_REQUEST_FREE`, it is not possible to check for the successful completion of the associated communication with calls to `MPI_WAIT` or `MPI_TEST`. Also, if an error occurs subsequently during the communication, an error code cannot be returned to the user — such an error must be treated as fatal. An active receive request should never be freed as the receiver will have no way to verify that the receive has completed and the receive buffer can be reused. (*End of advice to users.*)

**Example 3.13** An example using `MPI_REQUEST_FREE`.

```
CALL MPI_COMM_RANK(MPI_COMM_WORLD, rank, ierr)
IF (rank.EQ.0) THEN
  DO i=1, n
    CALL MPI_ISEND(outval, 1, MPI_REAL, 1, 0, MPI_COMM_WORLD, req, ierr)
    CALL MPI_REQUEST_FREE(req, ierr)
    CALL MPI_Irecv(inval, 1, MPI_REAL, 1, 0, MPI_COMM_WORLD, req, ierr)
    CALL MPI_WAIT(req, status, ierr)
  END DO
ELSE IF (rank.EQ.1) THEN
  CALL MPI_Irecv(inval, 1, MPI_REAL, 0, 0, MPI_COMM_WORLD, req, ierr)
  CALL MPI_WAIT(req, status, ierr)
  DO I=1, n-1
    CALL MPI_ISEND(outval, 1, MPI_REAL, 0, 0, MPI_COMM_WORLD, req, ierr)
    CALL MPI_REQUEST_FREE(req, ierr)
    CALL MPI_Irecv(inval, 1, MPI_REAL, 0, 0, MPI_COMM_WORLD, req, ierr)
    CALL MPI_WAIT(req, status, ierr)
  END DO
  CALL MPI_ISEND(outval, 1, MPI_REAL, 0, 0, MPI_COMM_WORLD, req, ierr)
  CALL MPI_WAIT(req, status, ierr)
END IF
```

### 3.7.4 Semantics of Nonblocking Communications

The semantics of nonblocking communication is defined by suitably extending the definitions in Section 3.5.

**Order** Nonblocking communication operations are ordered according to the execution order of the calls that initiate the communication. The non-overtaking requirement of Section 3.5 is extended to nonblocking communication, with this definition of order being used.

**Example 3.14** Message ordering for nonblocking operations.

```

CALL MPI_COMM_RANK(comm, rank, ierr)
IF (RANK.EQ.0) THEN
    CALL MPI_ISEND(a, 1, MPI_REAL, 1, 0, comm, r1, ierr)
    CALL MPI_ISEND(b, 1, MPI_REAL, 1, 0, comm, r2, ierr)
ELSE IF (rank.EQ.1) THEN
    CALL MPI_Irecv(a, 1, MPI_REAL, 0, MPI_ANY_TAG, comm, r1, ierr)
    CALL MPI_Irecv(b, 1, MPI_REAL, 0, 0, comm, r2, ierr)
END IF
CALL MPI_WAIT(r1, status, ierr)
CALL MPI_WAIT(r2, status, ierr)

```

The first send of process zero will match the first receive of process one, even if both messages are sent before process one executes either receive.

**Progress** A call to `MPI_WAIT` that completes a receive will eventually terminate and return if a matching send has been started, unless the send is satisfied by another receive. In particular, if the matching send is nonblocking, then the receive should complete even if no call is executed by the sender to complete the send. Similarly, a call to `MPI_WAIT` that completes a send will eventually return if a matching receive has been started, unless the receive is satisfied by another send, and even if no call is executed to complete the receive.

**Example 3.15** An illustration of progress semantics.

```

CALL MPI_COMM_RANK(comm, rank, ierr)
IF (RANK.EQ.0) THEN
    CALL MPI_SSEND(a, 1, MPI_REAL, 1, 0, comm, ierr)
    CALL MPI_SEND(b, 1, MPI_REAL, 1, 1, comm, ierr)
ELSE IF (rank.EQ.1) THEN
    CALL MPI_Irecv(a, 1, MPI_REAL, 0, 0, comm, r, ierr)
    CALL MPI_RECV(b, 1, MPI_REAL, 0, 1, comm, status, ierr)
    CALL MPI_WAIT(r, status, ierr)
END IF

```

This code should not deadlock in a correct MPI implementation. The first synchronous send of process zero must complete after process one posts the matching (nonblocking) receive even if process one has not yet reached the completing wait call. Thus, process zero will continue and execute the second send, allowing process one to complete execution.

If an `MPI_TEST` that completes a receive is repeatedly called with the same arguments, and a matching send has been started, then the call will eventually return `flag = true`, unless the send is satisfied by another receive. If an `MPI_TEST` that completes a send is repeatedly called with the same arguments, and a matching receive has been started, then the call will eventually return `flag = true`, unless the receive is satisfied by another send.

### 3.7.5 Multiple Completions

It is convenient to be able to wait for the completion of any, some, or all the operations in a list, rather than having to wait for a specific message. A call to `MPI_WAITANY` or `MPI_TESTANY` can be used to wait for the completion of one out of several operations. A call to `MPI_WAITALL` or `MPI_TESTALL` can be used to wait for all pending operations in

a list. A call to `MPI_WAITSOME` or `MPI_TESTSOME` can be used to complete all enabled operations in a list.

`MPI_WAITANY` (count, array\_of\_requests, index, status)

IN	count	list length (non-negative integer)
INOUT	array_of_requests	array of requests (array of handles)
OUT	index	index of handle for operation that completed (integer)
OUT	status	status object (Status)

```
int MPI_Waitany(int count, MPI_Request *array_of_requests, int *index,
               MPI_Status *status)
```

```
MPI_WAITANY(COUNT, ARRAY_OF_REQUESTS, INDEX, STATUS, IERROR)
  INTEGER COUNT, ARRAY_OF_REQUESTS(*), INDEX, STATUS(MPI_STATUS_SIZE),
  IERROR
```

```
{static int MPI::Request::Waitany(int count,
    MPI::Request array_of_requests[], MPI::Status& status) (binding
    deprecated, see Section 15.2) }
```

```
{static int MPI::Request::Waitany(int count,
    MPI::Request array_of_requests[]) (binding deprecated, see
    Section 15.2) }
```

Blocks until one of the operations associated with the active requests in the array has completed. If more than one operation is enabled and can terminate, one is arbitrarily chosen. Returns in `index` the index of that request in the array and returns in `status` the status of the completing communication. (The array is indexed from zero in C, and from one in Fortran.) If the request was allocated by a nonblocking communication operation, then it is deallocated and the request handle is set to `MPI_REQUEST_NULL`.

The `array_of_requests` list may contain null or inactive handles. If the list contains no active handles (list has length zero or all entries are null or inactive), then the call returns immediately with `index = MPI_UNDEFINED`, and an empty status.

The execution of `MPI_WAITANY(count, array_of_requests, index, status)` has the same effect as the execution of `MPI_WAIT(&array_of_requests[i], status)`, where `i` is the value returned by `index` (unless the value of `index` is `MPI_UNDEFINED`). `MPI_WAITANY` with an array containing one active entry is equivalent to `MPI_WAIT`.

<b>MPI_TESTANY(count, array_of_requests, index, flag, status)</b>			1
IN	count	list length (non-negative integer)	2
INOUT	array_of_requests	array of requests (array of handles)	3
OUT	index	index of operation that completed, or MPI_UNDEFINED if none completed (integer)	4
OUT	flag	true if one of the operations is complete (logical)	5
OUT	status	status object (Status)	6

```

int MPI_Testany(int count, MPI_Request *array_of_requests, int *index,
                int *flag, MPI_Status *status)
MPI_TESTANY(COUNT, ARRAY_OF_REQUESTS, INDEX, FLAG, STATUS, IERROR)
    LOGICAL FLAG
    INTEGER COUNT, ARRAY_OF_REQUESTS(*), INDEX, STATUS(MPI_STATUS_SIZE),
    IERROR
{static bool MPI::Request::Testany(int count,
    MPI::Request array_of_requests[], int& index,
    MPI::Status& status) (binding deprecated, see Section 15.2) }
{static bool MPI::Request::Testany(int count,
    MPI::Request array_of_requests[], int& index) (binding deprecated,
    see Section 15.2) }

```

Tests for completion of either one or none of the operations associated with active handles. In the former case, it returns **flag** = true, returns in **index** the index of this request in the array, and returns in **status** the status of that operation; if the request was allocated by a nonblocking communication call then the request is deallocated and the handle is set to MPI\_REQUEST\_NULL. (The array is indexed from zero in C, and from one in Fortran.) In the latter case (no operation completed), it returns **flag** = false, returns a value of MPI\_UNDEFINED in **index** and **status** is undefined.

The array may contain null or inactive handles. If the array contains no active handles then the call returns immediately with **flag** = true, **index** = MPI\_UNDEFINED, and an empty **status**.

If the array of requests contains active handles then the execution of MPI\_TESTANY(count, array\_of\_requests, index, status) has the same effect as the execution of MPI\_TEST( &array\_of\_requests[i], flag, status), for i=0, 1, ..., count-1, in some arbitrary order, until one call returns **flag** = true, or all fail. In the former case, **index** is set to the last value of i, and in the latter case, it is set to MPI\_UNDEFINED. MPI\_TESTANY with an array containing one active entry is equivalent to MPI\_TEST.

<b>MPI_WAITALL( count, array_of_requests, array_of_statuses)</b>			43
IN	count	lists length (non-negative integer)	44
INOUT	array_of_requests	array of requests (array of handles)	45
OUT	array_of_statuses	array of status objects (array of Status)	46

```

1  int MPI_Waitall(int count, MPI_Request *array_of_requests,
2                  MPI_Status *array_of_statuses)
3
4  MPI_WAITALL(COUNT, ARRAY_OF_REQUESTS, ARRAY_OF_STATUSES, IERROR)
5      INTEGER COUNT, ARRAY_OF_REQUESTS(*)
6      INTEGER ARRAY_OF_STATUSES(MPI_STATUS_SIZE,*), IERROR
7
8  {static void MPI::Request::Waitall(int count,
9      MPI::Request array_of_requests[],
10     MPI::Status array_of_statuses[]) (binding deprecated, see
11     Section 15.2) }
12
13 {static void MPI::Request::Waitall(int count,
14     MPI::Request array_of_requests[]) (binding deprecated, see
15     Section 15.2) }

```

Blocks until all communication operations associated with active handles in the list complete, and return the status of all these operations (this includes the case where no handle in the list is active). Both arrays have the same number of valid entries. The *i*-th entry in `array_of_statuses` is set to the return status of the *i*-th operation. Requests that were created by nonblocking communication operations are deallocated and the corresponding handles in the array are set to `MPI_REQUEST_NULL`. The list may contain null or inactive handles. The call sets to empty the status of each such entry.

The error-free execution of `MPI_WAITALL(count, array_of_requests, array_of_statuses)` has the same effect as the execution of `MPI_WAIT(&array_of_request[i], &array_of_statuses[i])`, for *i*=0 ,..., count-1, in some arbitrary order. `MPI_WAITALL` with an array of length one is equivalent to `MPI_WAIT`.

When one or more of the communications completed by a call to `MPI_WAITALL` fail, it is desirable to return specific information on each communication. The function `MPI_WAITALL` will return in such case the error code `MPI_ERR_IN_STATUS` and will set the error field of each status to a specific error code. This code will be `MPI_SUCCESS`, if the specific communication completed; it will be another specific error code, if it failed; or it can be `MPI_ERR_PENDING` if it has neither failed nor completed. The function `MPI_WAITALL` will return `MPI_SUCCESS` if no request had an error, or will return another error code if it failed for other reasons (such as invalid arguments). In such cases, it will not update the error fields of the statuses.

*Rationale.* This design streamlines error handling in the application. The application code need only test the (single) function result to determine if an error has occurred. It needs to check each individual status only when an error occurred. (*End of rationale.*)

```

41 MPI_TESTALL(count, array_of_requests, flag, array_of_statuses)
42
43 IN      count                lists length (non-negative integer)
44 INOUT   array_of_requests    array of requests (array of handles)
45 OUT     flag                 (logical)
46 OUT     array_of_statuses    array of status objects (array of Status)

```

```

int MPI_Testall(int count, MPI_Request *array_of_requests, int *flag,
                MPI_Status *array_of_statuses)
MPI_TESTALL(COUNT, ARRAY_OF_REQUESTS, FLAG, ARRAY_OF_STATUSES, IERROR)
    LOGICAL FLAG
    INTEGER COUNT, ARRAY_OF_REQUESTS(*),
    ARRAY_OF_STATUSES(MPI_STATUS_SIZE,*), IERROR
{static bool MPI::Request::Testall(int count,
    MPI::Request array_of_requests[],
    MPI::Status array_of_statuses[]) (binding deprecated, see
    Section 15.2) }
{static bool MPI::Request::Testall(int count,
    MPI::Request array_of_requests[]) (binding deprecated, see
    Section 15.2) }

```

Returns `flag = true` if all communications associated with active handles in the array have completed (this includes the case where no handle in the list is active). In this case, each status entry that corresponds to an active handle request is set to the status of the corresponding communication; if the request was allocated by a nonblocking communication call then it is deallocated, and the handle is set to `MPI_REQUEST_NULL`. Each status entry that corresponds to a null or inactive handle is set to empty.

Otherwise, `flag = false` is returned, no request is modified and the values of the status entries are undefined. This is a local operation.

Errors that occurred during the execution of `MPI_TESTALL` are handled as errors in `MPI_WAITALL`.

```

MPI_WAITSOME(incount, array_of_requests, outcount, array_of_indices, array_of_statuses)

```

IN	incount	length of <code>array_of_requests</code> (non-negative integer)
INOUT	<code>array_of_requests</code>	array of requests (array of handles)
OUT	outcount	number of completed requests (integer)
OUT	<code>array_of_indices</code>	array of indices of operations that completed (array of integers)
OUT	<code>array_of_statuses</code>	array of status objects for operations that completed (array of Status)

```

int MPI_Waitsome(int incount, MPI_Request *array_of_requests,
                int *outcount, int *array_of_indices,
                MPI_Status *array_of_statuses)
MPI_WAITSOME(INCOUNT, ARRAY_OF_REQUESTS, OUTCOUNT, ARRAY_OF_INDICES,
    ARRAY_OF_STATUSES, IERROR)
    INTEGER INCOUNT, ARRAY_OF_REQUESTS(*), OUTCOUNT, ARRAY_OF_INDICES(*),
    ARRAY_OF_STATUSES(MPI_STATUS_SIZE,*), IERROR

```



```

1  {static int MPI::Request::Waitsome(int incout,
2      MPI::Request array_of_requests[], int array_of_indices[],
3      MPI::Status array_of_statuses[]) (binding deprecated, see
4      Section 15.2) }
5
6  {static int MPI::Request::Waitsome(int incout,
7      MPI::Request array_of_requests[],
8      int array_of_indices[]) (binding deprecated, see Section 15.2) }

```

Waits until at least one of the operations associated with active handles in the list have completed. Returns in `outcount` the number of requests from the list `array_of_requests` that have completed. Returns in the first `outcount` locations of the array `array_of_indices` the indices of these operations (index within the array `array_of_requests`; the array is indexed from zero in C and from one in Fortran). Returns in the first `outcount` locations of the array `array_of_status` the status for these completed operations. If a request that completed was allocated by a nonblocking communication call, then it is deallocated, and the associated handle is set to `MPI_REQUEST_NULL`.

If the list contains no active handles, then the call returns immediately with `outcount = MPI_UNDEFINED`.

When one or more of the communications completed by `MPI_WAITSOME` fails, then it is desirable to return specific information on each communication. The arguments `outcount`, `array_of_indices` and `array_of_statuses` will be adjusted to indicate completion of all communications that have succeeded or failed. The call will return the error code `MPI_ERR_IN_STATUS` and the error field of each status returned will be set to indicate success or to indicate the specific error that occurred. The call will return `MPI_SUCCESS` if no request resulted in an error, and will return another error code if it failed for other reasons (such as invalid arguments). In such cases, it will not update the error fields of the statuses.

`MPI_TESTSOME(incout, array_of_requests, outcount, array_of_indices, array_of_statuses)`

IN	<code>incout</code>	length of <code>array_of_requests</code> (non-negative integer)
INOUT	<code>array_of_requests</code>	array of requests (array of handles)
OUT	<code>outcount</code>	number of completed requests (integer)
OUT	<code>array_of_indices</code>	array of indices of operations that completed (array of integers)
OUT	<code>array_of_statuses</code>	array of status objects for operations that completed (array of Status)

```

41 int MPI_Testsome(int incout, MPI_Request *array_of_requests,
42     int *outcount, int *array_of_indices,
43     MPI_Status *array_of_statuses)
44
45 MPI_TESTSOME(INCOUNT, ARRAY_OF_REQUESTS, OUTCOUNT, ARRAY_OF_INDICES,
46     ARRAY_OF_STATUSES, IERROR)
47     INTEGER INCOUNT, ARRAY_OF_REQUESTS(*), OUTCOUNT, ARRAY_OF_INDICES(*),
48     ARRAY_OF_STATUSES(MPI_STATUS_SIZE,*), IERROR

```

```

{static int MPI::Request::Testsome(int incout,
    MPI::Request array_of_requests[], int array_of_indices[],
    MPI::Status array_of_statuses[]) (binding deprecated, see
    Section 15.2) }
{static int MPI::Request::Testsome(int incout,
    MPI::Request array_of_requests[],
    int array_of_indices[]) (binding deprecated, see Section 15.2) }

```

Behaves like MPI\_WAITSSOME, except that it returns immediately. If no operation has completed it returns `outcount = 0`. If there is no active handle in the list it returns `outcount = MPI_UNDEFINED`.

MPI\_TESTSSOME is a local operation, which returns immediately, whereas MPI\_WAITSSOME will block until a communication completes, if it was passed a list that contains at least one active handle. Both calls fulfill a fairness requirement: If a request for a receive repeatedly appears in a list of requests passed to MPI\_WAITSSOME or MPI\_TESTSSOME, and a matching send has been posted, then the receive will eventually succeed, unless the send is satisfied by another receive; and similarly for send requests.

Errors that occur during the execution of MPI\_TESTSSOME are handled as for MPI\_WAITSSOME.

*Advice to users.* The use of MPI\_TESTSSOME is likely to be more efficient than the use of MPI\_TESTANY. The former returns information on all completed communications, with the latter, a new call is required for each communication that completes.

A server with multiple clients can use MPI\_WAITSSOME so as not to starve any client. Clients send messages to the server with service requests. The server calls MPI\_WAITSSOME with one receive request for each client, and then handles all receives that completed. If a call to MPI\_WAITANY is used instead, then one client could starve while requests from another client always sneak in first. (*End of advice to users.*)

*Advice to implementors.* MPI\_TESTSSOME should complete as many pending communications as possible. (*End of advice to implementors.*)

**Example 3.16** Client-server code (starvation can occur).

```

CALL MPI_COMM_SIZE(comm, size, ierr)
CALL MPI_COMM_RANK(comm, rank, ierr)
IF(rank .GT. 0) THEN      ! client code
  DO WHILE(.TRUE.)
    CALL MPI_ISEND(a, n, MPI_REAL, 0, tag, comm, request, ierr)
    CALL MPI_WAIT(request, status, ierr)
  END DO
ELSE      ! rank=0 -- server code
  DO i=1, size-1
    CALL MPI_Irecv(a(1,i), n, MPI_REAL, i, tag,
      comm, request_list(i), ierr)
  END DO
  DO WHILE(.TRUE.)

```

```

1      CALL MPI_WAITANY(size-1, request_list, index, status, ierr)
2      CALL DO_SERVICE(a(1,index)) ! handle one message
3      CALL MPI_Irecv(a(1, index), n, MPI_REAL, index, tag,
4                    comm, request_list(index), ierr)
5  END DO
6 END IF

```

**Example 3.17** Same code, using MPI\_WAITSSOME.

```

11 CALL MPI_COMM_SIZE(comm, size, ierr)
12 CALL MPI_COMM_RANK(comm, rank, ierr)
13 IF(rank .GT. 0) THEN ! client code
14   DO WHILE(.TRUE.)
15     CALL MPI_ISEND(a, n, MPI_REAL, 0, tag, comm, request, ierr)
16     CALL MPI_WAIT(request, status, ierr)
17   END DO
18 ELSE ! rank=0 -- server code
19   DO i=1, size-1
20     CALL MPI_Irecv(a(1,i), n, MPI_REAL, i, tag,
21                   comm, request_list(i), ierr)
22   END DO
23   DO WHILE(.TRUE.)
24     CALL MPI_WAITSSOME(size, request_list, numdone,
25                       indices, statuses, ierr)
26     DO i=1, numdone
27       CALL DO_SERVICE(a(1, indices(i)))
28       CALL MPI_Irecv(a(1, indices(i)), n, MPI_REAL, 0, tag,
29                     comm, request_list(indices(i)), ierr)
30     END DO
31   END DO
32 END IF

```

### 3.7.6 Non-destructive Test of status

This call is useful for accessing the information associated with a request, without freeing the request (in case the user is expected to access it later). It allows one to layer libraries more conveniently, since multiple layers of software may access the same completed request and extract from it the status information.

**MPI\_REQUEST\_GET\_STATUS( request, flag, status )**

IN	request	request (handle)
OUT	flag	boolean flag, same as from MPI_TEST (logical)
OUT	status	MPI_STATUS object if flag is true (Status)

```

int MPI_Request_get_status(MPI_Request request, int *flag,
                           MPI_Status *status)
MPI_REQUEST_GET_STATUS( REQUEST, FLAG, STATUS, IERROR)
    INTEGER REQUEST, STATUS(MPI_STATUS_SIZE), IERROR
    LOGICAL FLAG
{bool MPI::Request::Get_status(MPI::Status& status) const(binding deprecated,
    see Section 15.2) }
{bool MPI::Request::Get_status() const(binding deprecated, see Section 15.2) }

```

Sets `flag=true` if the operation is complete, and, if so, returns in status the request status. However, unlike `test` or `wait`, it does not deallocate or inactivate the request; a subsequent call to `test`, `wait` or `free` should be executed with that request. It sets `flag=false` if the operation is not complete.

One is allowed to call `MPI_REQUEST_GET_STATUS` with a null or inactive request argument. In such a case the operation returns with `flag=true` and empty status.

### 3.8 Probe and Cancel

The `MPI_PROBE` and `MPI_IPROBE` operations allow incoming messages to be checked for, without actually receiving them. The user can then decide how to receive them, based on the information returned by the probe (basically, the information returned by `status`). In particular, the user may allocate memory for the receive buffer, according to the length of the probed message.

The `MPI_CANCEL` operation allows pending communications to be canceled. This is required for cleanup. Posting a send or a receive ties up user resources (send or receive buffers), and a cancel may be needed to free these resources gracefully.

```

MPI_IPROBE(source, tag, comm, flag, status)
    IN      source      rank of source or MPI_ANY_SOURCE (integer)
    IN      tag         message tag or MPI_ANY_TAG (integer)
    IN      comm        communicator (handle)
    OUT     flag         (logical)
    OUT     status      status object (Status)

```

```

int MPI_Iprobe(int source, int tag, MPI_Comm comm, int *flag,
               MPI_Status *status)
MPI_IPROBE(SOURCE, TAG, COMM, FLAG, STATUS, IERROR)
    LOGICAL FLAG
    INTEGER SOURCE, TAG, COMM, STATUS(MPI_STATUS_SIZE), IERROR
{bool MPI::Comm::Iprobe(int source, int tag, MPI::Status& status)
    const(binding deprecated, see Section 15.2) }

```

```

1 {bool MPI::Comm::Iprobe(int source, int tag) const(binding deprecated, see
2 Section 15.2) }

```

MPI\_IPROBE(source, tag, comm, flag, status) returns flag = true if there is a message that can be received and that matches the pattern specified by the arguments source, tag, and comm. The call matches the same message that would have been received by a call to MPI\_RECV(..., source, tag, comm, status) executed at the same point in the program, and returns in status the same value that would have been returned by MPI\_RECV(). Otherwise, the call returns flag = false, and leaves status undefined.

If MPI\_IPROBE returns flag = true, then the content of the status object can be subsequently accessed as described in Section 3.2.5 to find the source, tag and length of the probed message.

A subsequent receive executed with the same communicator, and the source and tag returned in status by MPI\_IPROBE will receive the message that was matched by the probe, if no other intervening receive occurs after the probe, and the send is not successfully cancelled before the receive. If the receiving process is multi-threaded, it is the user's responsibility to ensure that the last condition holds.

The source argument of MPI\_PROBE can be MPI\_ANY\_SOURCE, and the tag argument can be MPI\_ANY\_TAG, so that one can probe for messages from an arbitrary source and/or with an arbitrary tag. However, a specific communication context must be provided with the comm argument.

It is not necessary to receive a message immediately after it has been probed for, and the same message may be probed for several times before it is received.

```

25 MPI_PROBE(source, tag, comm, status)

```

27	IN	source	rank of source or MPI_ANY_SOURCE (integer)
28	IN	tag	message tag or MPI_ANY_TAG (integer)
29	IN	comm	communicator (handle)
31	OUT	status	status object (Status)

```

33 int MPI_Probe(int source, int tag, MPI_Comm comm, MPI_Status *status)

```

```

34 MPI_PROBE(SOURCE, TAG, COMM, STATUS, IERROR)

```

```

35     INTEGER SOURCE, TAG, COMM, STATUS(MPI_STATUS_SIZE), IERROR

```

```

37 {void MPI::Comm::Probe(int source, int tag, MPI::Status& status)
38     const(binding deprecated, see Section 15.2) }

```

```

39 {void MPI::Comm::Probe(int source, int tag) const(binding deprecated, see
40 Section 15.2) }

```

MPI\_PROBE behaves like MPI\_IPROBE except that it is a blocking call that returns only after a matching message has been found.

The MPI implementation of MPI\_PROBE and MPI\_IPROBE needs to guarantee progress: if a call to MPI\_PROBE has been issued by a process, and a send that matches the probe has been initiated by some process, then the call to MPI\_PROBE will return, unless the message is received by another concurrent receive operation (that is executed by another thread at the probing process). Similarly, if a process busy waits with MPI\_IPROBE and

a matching message has been issued, then the call to `MPI_Iprobe` will eventually return `flag = true` unless the message is received by another concurrent receive operation.

### Example 3.18

Use blocking probe to wait for an incoming message.

```

CALL MPI_COMM_RANK(comm, rank, ierr)
IF (rank.EQ.0) THEN
    CALL MPI_SEND(i, 1, MPI_INTEGER, 2, 0, comm, ierr)
ELSE IF (rank.EQ.1) THEN
    CALL MPI_SEND(x, 1, MPI_REAL, 2, 0, comm, ierr)
ELSE IF (rank.EQ.2) THEN
    DO i=1, 2
        CALL MPI_PROBE(MPI_ANY_SOURCE, 0,
                        comm, status, ierr)
        IF (status(MPI_SOURCE) .EQ. 0) THEN
100      CALL MPI_RECV(i, 1, MPI_INTEGER, 0, 0, comm, status, ierr)
        ELSE
200      CALL MPI_RECV(x, 1, MPI_REAL, 1, 0, comm, status, ierr)
        END IF
    END DO
END IF

```

Each message is received with the right type.

**Example 3.19** A similar program to the previous example, but now it has a problem.

```

CALL MPI_COMM_RANK(comm, rank, ierr)
IF (rank.EQ.0) THEN
    CALL MPI_SEND(i, 1, MPI_INTEGER, 2, 0, comm, ierr)
ELSE IF (rank.EQ.1) THEN
    CALL MPI_SEND(x, 1, MPI_REAL, 2, 0, comm, ierr)
ELSE IF (rank.EQ.2) THEN
    DO i=1, 2
        CALL MPI_PROBE(MPI_ANY_SOURCE, 0,
                        comm, status, ierr)
        IF (status(MPI_SOURCE) .EQ. 0) THEN
100      CALL MPI_RECV(i, 1, MPI_INTEGER, MPI_ANY_SOURCE,
                        0, comm, status, ierr)
        ELSE
200      CALL MPI_RECV(x, 1, MPI_REAL, MPI_ANY_SOURCE,
                        0, comm, status, ierr)
        END IF
    END DO
END IF

```

We slightly modified Example 3.18, using `MPI_ANY_SOURCE` as the source argument in the two receive calls in statements labeled 100 and 200. The program is now incorrect: the receive operation may receive a message that is distinct from the message probed by the preceding call to `MPI_PROBE`.

## MPI\_CANCEL(request)

```
int MPI_Cancel(MPI_Request *request)
```

INTEGER REQUEST, IERROR

Either the cancellation succeeds, or the communication succeeds, but not both. If a send is marked for cancellation, then it must be the case that either the send completes normally, in which case the message sent was received at the destination process, or that the send is successfully canceled, in which case no part of the message was received at the destination. Then, any matching receive has to be satisfied by another send. If a receive is marked for cancellation, then it must be the case that either the receive completes normally,

or that the receive is successfully canceled, in which case no part of the receive buffer is altered. Then, any matching send has to be satisfied by another receive.

If the operation has been canceled, then information to that effect will be returned in the status argument of the operation that completes the communication.

*Rationale.* Although the IN request handle parameter should not need to be passed by reference, the C binding has listed the argument type as `MPI_Request*` since MPI-1.0. This function signature therefore cannot be changed without breaking existing MPI applications. (*End of rationale.*)

`MPI_TEST_CANCELLED(status, flag)`

IN	status	status object (Status)
OUT	flag	(logical)

`int MPI_Test_cancelled(MPI_Status *status, int *flag)`

`MPI_TEST_CANCELLED(STATUS, FLAG, IERROR)`

LOGICAL FLAG

INTEGER STATUS(MPI\_STATUS\_SIZE), IERROR

`{bool MPI::Status::Is_cancelled() const(binding deprecated, see Section 15.2) }`

Returns `flag = true` if the communication associated with the status object was canceled successfully. In such a case, all other fields of `status` (such as `count` or `tag`) are undefined. Returns `flag = false`, otherwise. If a receive operation might be canceled then one should call `MPI_TEST_CANCELLED` first, to check whether the operation was canceled, before checking on the other fields of the return status.

*Advice to users.* Cancel can be an expensive operation that should be used only exceptionally. (*End of advice to users.*)

*Advice to implementors.* If a send operation uses an “eager” protocol (data is transferred to the receiver before a matching receive is posted), then the cancellation of this send may require communication with the intended receiver in order to free allocated buffers. On some systems this may require an interrupt to the intended receiver. Note that, while communication may be needed to implement `MPI_CANCEL`, this is still a local operation, since its completion does not depend on the code executed by other processes. If processing is required on another process, this should be transparent to the application (hence the need for an interrupt and an interrupt handler). (*End of advice to implementors.*)

### 3.9 Persistent Communication Requests

Often a communication with the same argument list is repeatedly executed within the inner loop of a parallel computation. In such a situation, it may be possible to optimize the communication by binding the list of communication arguments to a **persistent** communication request once and, then, repeatedly using the request to initiate and complete



messages. The persistent request thus created can be thought of as a communication port or a “half-channel.” It does not provide the full functionality of a conventional channel, since there is no binding of the send port to the receive port. This construct allows reduction of the overhead for communication between the process and communication controller, but not of the overhead for communication between one communication controller and another. It is not necessary that messages sent with a persistent request be received by a receive operation using a persistent request, or vice versa.

A persistent communication request is created using one of the five following calls. These calls involve no communication.

**MPI\_SEND\_INIT**(buf, count, datatype, dest, tag, comm, request)

IN	buf	initial address of send buffer (choice)
IN	count	number of elements sent (non-negative integer)
IN	datatype	type of each element (handle)
IN	dest	rank of destination (integer)
IN	tag	message tag (integer)
IN	comm	communicator (handle)
OUT	request	communication request (handle)

```
int MPI_Send_init(void* buf, int count, MPI_Datatype datatype, int dest,
                  int tag, MPI_Comm comm, MPI_Request *request)
```

```
MPI_SEND_INIT(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
<type> BUF(*)
INTEGER REQUEST, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR
```

```
{MPI::Prequest MPI::Comm::Send_init(const void* buf, int count, const
MPI::Datatype& datatype, int dest, int tag) const(binding
deprecated, see Section 15.2) }
```

Creates a persistent communication request for a standard mode send operation, and binds to it all the arguments of a send operation.

**MPI\_BSEND\_INIT**(buf, count, datatype, dest, tag, comm, request)

IN	buf	initial address of send buffer (choice)
IN	count	number of elements sent (non-negative integer)
IN	datatype	type of each element (handle)
IN	dest	rank of destination (integer)
IN	tag	message tag (integer)
IN	comm	communicator (handle)
OUT	request	communication request (handle)

```

int MPI_Bsend_init(void* buf, int count, MPI_Datatype datatype, int dest,
                  int tag, MPI_Comm comm, MPI_Request *request)
MPI_BSEND_INIT(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
<type> BUF(*)
INTEGER REQUEST, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR
{MPI::Prequest MPI::Comm::Bsend_init(const void* buf, int count, const
MPI::Datatype& datatype, int dest, int tag) const(binding
deprecated, see Section 15.2) }
```

Creates a persistent communication request for a buffered mode send.

```

MPI_SSEND_INIT(buf, count, datatype, dest, tag, comm, request)
IN      buf                initial address of send buffer (choice)
IN      count              number of elements sent (non-negative integer)
IN      datatype           type of each element (handle)
IN      dest               rank of destination (integer)
IN      tag                message tag (integer)
IN      comm               communicator (handle)
OUT     request            communication request (handle)
```

```

int MPI_Ssend_init(void* buf, int count, MPI_Datatype datatype, int dest,
                  int tag, MPI_Comm comm, MPI_Request *request)
MPI_SSEND_INIT(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
<type> BUF(*)
INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR
{MPI::Prequest MPI::Comm::Ssend_init(const void* buf, int count, const
MPI::Datatype& datatype, int dest, int tag) const(binding
deprecated, see Section 15.2) }
```

Creates a persistent communication object for a synchronous mode send operation.

```

MPI_RSEND_INIT(buf, count, datatype, dest, tag, comm, request)
IN      buf                initial address of send buffer (choice)
IN      count              number of elements sent (non-negative integer)
IN      datatype           type of each element (handle)
IN      dest               rank of destination (integer)
IN      tag                message tag (integer)
IN      comm               communicator (handle)
OUT     request            communication request (handle)
```

```

1  int MPI_Rsend_init(void* buf, int count, MPI_Datatype datatype, int dest,
2      int tag, MPI_Comm comm, MPI_Request *request)
3
4  MPI_RSEND_INIT(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
5      <type> BUF(*)
6      INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR
7
8  {MPI::Prequest MPI::Comm::Rsend_init(const void* buf, int count, const
9      MPI::Datatype& datatype, int dest, int tag) const(binding
10      deprecated, see Section 15.2) }

```

Creates a persistent communication object for a ready mode send operation.

```

13 MPI_RECV_INIT(buf, count, datatype, source, tag, comm, request)
14
15 OUT    buf                initial address of receive buffer (choice)
16 IN     count              number of elements received (non-negative integer)
17 IN     datatype           type of each element (handle)
18 IN     source             rank of source or MPI_ANY_SOURCE (integer)
19 IN     tag                message tag or MPI_ANY_TAG (integer)
20 IN     comm              communicator (handle)
21 IN     request            communication request (handle)
22
23
24
25 int MPI_Recv_init(void* buf, int count, MPI_Datatype datatype, int source,
26     int tag, MPI_Comm comm, MPI_Request *request)
27
28 MPI_RECV_INIT(BUF, COUNT, DATATYPE, SOURCE, TAG, COMM, REQUEST, IERROR)
29     <type> BUF(*)
30     INTEGER COUNT, DATATYPE, SOURCE, TAG, COMM, REQUEST, IERROR
31
32 {MPI::Prequest MPI::Comm::Recv_init(void* buf, int count, const
33     MPI::Datatype& datatype, int source, int tag) const(binding
34     deprecated, see Section 15.2) }

```

Creates a persistent communication request for a receive operation. The argument `buf` is marked as OUT because the user gives permission to write on the receive buffer by passing the argument to `MPI_RECV_INIT`.

A persistent communication request is inactive after it was created — no active communication is attached to the request.

A communication (send or receive) that uses a persistent request is initiated by the function `MPI_START`.

```

43 MPI_START(request)
44 INOUT request            communication request (handle)
45
46
47 int MPI_Start(MPI_Request *request)
48
49 MPI_START(REQUEST, IERROR)

```

INTEGER REQUEST, IERROR

```
{void MPI::Prequest::Start() (binding deprecated, see Section 15.2) }
```

The argument, `request`, is a handle returned by one of the previous five calls. The associated request should be inactive. The request becomes active once the call is made.

If the request is for a send with ready mode, then a matching receive should be posted before the call is made. The communication buffer should not be modified after the call, and until the operation completes.

The call is local, with similar semantics to the nonblocking communication operations described in Section 3.7. That is, a call to `MPI_START` with a request created by `MPI_SEND_INIT` starts a communication in the same manner as a call to `MPI_ISEND`; a call to `MPI_START` with a request created by `MPI_BSEND_INIT` starts a communication in the same manner as a call to `MPI_IBSEND`; and so on.

`MPI_STARTALL(count, array_of_requests)`

IN	count	list length (non-negative integer)
INOUT	array_of_requests	array of requests (array of handle)

```
int MPI_Startall(int count, MPI_Request *array_of_requests)
```

```
MPI_STARTALL(COUNT, ARRAY_OF_REQUESTS, IERROR)
    INTEGER COUNT, ARRAY_OF_REQUESTS(*), IERROR
```

```
{static void MPI::Prequest::Startall(int count,
    MPI::Prequest array_of_requests[]) (binding deprecated, see
    Section 15.2) }
```

Start all communications associated with requests in `array_of_requests`. A call to `MPI_STARTALL(count, array_of_requests)` has the same effect as calls to `MPI_START (&array_of_requests[i])`, executed for  $i=0, \dots, \text{count}-1$ , in some arbitrary order.

A communication started with a call to `MPI_START` or `MPI_STARTALL` is completed by a call to `MPI_WAIT`, `MPI_TEST`, or one of the derived functions described in Section 3.7.5. The request becomes inactive after successful completion of such call. The request is not deallocated and it can be activated anew by an `MPI_START` or `MPI_STARTALL` call.

A persistent request is deallocated by a call to `MPI_REQUEST_FREE` (Section 3.7.3).

The call to `MPI_REQUEST_FREE` can occur at any point in the program after the persistent request was created. However, the request will be deallocated only after it becomes inactive. Active receive requests should not be freed. Otherwise, it will not be possible to check that the receive has completed. It is preferable, in general, to free requests when they are inactive. If this rule is followed, then the functions described in this section will be invoked in a sequence of the form,

**Create (Start Complete)\* Free**

where  $*$  indicates zero or more repetitions. If the same communication object is used in several concurrent threads, it is the user's responsibility to coordinate calls so that the

correct sequence is obeyed.

A send operation initiated with `MPI_START` can be matched with any receive operation and, likewise, a receive operation initiated with `MPI_START` can receive messages generated by any send operation.

*Advice to users.* To prevent problems with the argument copying and register optimization done by Fortran compilers, please note the hints in subsections “Problems Due to Data Copying and Sequence Association,” and “A Problem with Register Optimization” in Section 16.2.2 on pages 514 and 517. (*End of advice to users.*)

### 3.10 Send-Receive

The **send-receive** operations combine in one call the sending of a message to one destination and the receiving of another message, from another process. The two (source and destination) are possibly the same. A send-receive operation is very useful for executing a shift operation across a chain of processes. If blocking sends and receives are used for such a shift, then one needs to order the sends and receives correctly (for example, even processes send, then receive, odd processes receive first, then send) so as to prevent cyclic dependencies that may lead to deadlock. When a send-receive operation is used, the communication subsystem takes care of these issues. The send-receive operation can be used in conjunction with the functions described in Chapter 7 in order to perform shifts on various logical topologies. Also, a send-receive operation is useful for implementing remote procedure calls.

A message sent by a send-receive operation can be received by a regular receive operation or probed by a probe operation; a send-receive operation can receive a message sent by a regular send operation.

`MPI_SENDRECV(sendbuf, sendcount, sendtype, dest, sendtag, recvbuf, recvcount, recvtype, source, recvtag, comm, status)`

IN	sendbuf	initial address of send buffer (choice)
IN	sendcount	number of elements in send buffer (non-negative integer)
IN	sendtype	type of elements in send buffer (handle)
IN	dest	rank of destination (integer)
IN	sendtag	send tag (integer)
OUT	recvbuf	initial address of receive buffer (choice)
IN	recvcount	number of elements in receive buffer (non-negative integer)
IN	recvtype	type of elements in receive buffer (handle)
IN	source	rank of source or <code>MPI_ANY_SOURCE</code> (integer)
IN	recvtag	receive tag or <code>MPI_ANY_TAG</code> (integer)
IN	comm	communicator (handle)
OUT	status	status object (Status)

```

int MPI_Sendrecv(void *sendbuf, int sendcount, MPI_Datatype sendtype,
                 int dest, int sendtag, void *recvbuf, int recvcount,
                 MPI_Datatype recvtype, int source, int recvtag, MPI_Comm comm,
                 MPI_Status *status)
MPI_SENDRECV(SENDBUF, SENDCOUNT, SENDTYPE, DEST, SENDTAG, RECVBUF,
             RECVCOUNT, RECVTYPE, SOURCE, RECVTAG, COMM, STATUS, IERROR)
<type> SENDBUF(*), RECVBUF(*)
INTEGER SENDCOUNT, SENDTYPE, DEST, SENDTAG, RECVCOUNT, RECVTYPE,
SOURCE, RECVTAG, COMM, STATUS(MPI_STATUS_SIZE), IERROR
{void MPI::Comm::Sendrecv(const void *sendbuf, int sendcount, const
                          MPI::Datatype& sendtype, int dest, int sendtag, void *recvbuf,
                          int recvcount, const MPI::Datatype& recvtype, int source,
                          int recvtag, MPI::Status& status) const(binding deprecated, see
                          Section 15.2) }
{void MPI::Comm::Sendrecv(const void *sendbuf, int sendcount, const
                          MPI::Datatype& sendtype, int dest, int sendtag, void *recvbuf,
                          int recvcount, const MPI::Datatype& recvtype, int source,
                          int recvtag) const(binding deprecated, see Section 15.2) }

```

Execute a blocking send and receive operation. Both send and receive use the same communicator, but possibly different tags. The send buffer and receive buffers must be disjoint, and may have different lengths and datatypes.

The semantics of a send-receive operation is what would be obtained if the caller forked two concurrent threads, one to execute the send, and one to execute the receive, followed by a join of these two threads.

MPI_SENDRECV_REPLACE(buf, count, datatype, dest, sendtag, source, recvtag, comm, status)			29
			30
INOUT	buf	initial address of send and receive buffer (choice)	31
IN	count	number of elements in send and receive buffer (non-negative integer)	32
			33
			34
IN	datatype	type of elements in send and receive buffer (handle)	35
IN	dest	rank of destination (integer)	36
IN	sendtag	send message tag (integer)	37
IN	source	rank of source or MPI_ANY_SOURCE (integer)	38
IN	recvtag	receive message tag or MPI_ANY_TAG (integer)	39
IN	comm	communicator (handle)	40
OUT	status	status object (Status)	41
			42
			43
			44

```

int MPI_Sendrecv_replace(void* buf, int count, MPI_Datatype datatype,
                        int dest, int sendtag, int source, int recvtag, MPI_Comm comm,
                        MPI_Status *status)

```

```

1 MPI_SENDRECV_REPLACE(BUF, COUNT, DATATYPE, DEST, SENDTAG, SOURCE, RECVTAG,
2   COMM, STATUS, IERROR)
3   <type> BUF(*)
4   INTEGER COUNT, DATATYPE, DEST, SENDTAG, SOURCE, RECVTAG, COMM,
5   STATUS(MPI_STATUS_SIZE), IERROR

```

```

6 {void MPI::Comm::Sendrecv_replace(void* buf, int count, const
7   MPI::Datatype& datatype, int dest, int sendtag, int source,
8   int recvtag, MPI::Status& status) const(binding deprecated, see
9   Section 15.2) }

```

```

11 {void MPI::Comm::Sendrecv_replace(void* buf, int count, const
12   MPI::Datatype& datatype, int dest, int sendtag, int source,
13   int recvtag) const(binding deprecated, see Section 15.2) }

```

Execute a blocking send and receive. The same buffer is used both for the send and for the receive, so that the message sent is replaced by the message received.

*Advice to implementors.* Additional intermediate buffering is needed for the “replace” variant. (*End of advice to implementors.*)

### 3.11 Null Processes

In many instances, it is convenient to specify a “dummy” source or destination for communication. This simplifies the code that is needed for dealing with boundaries, for example, in the case of a non-circular shift done with calls to send-receive.

The special value `MPI_PROC_NULL` can be used instead of a rank wherever a source or a destination argument is required in a call. A communication with process `MPI_PROC_NULL` has no effect. A send to `MPI_PROC_NULL` succeeds and returns as soon as possible. A receive from `MPI_PROC_NULL` succeeds and returns as soon as possible with no modifications to the receive buffer. When a receive with `source = MPI_PROC_NULL` is executed then the status object returns `source = MPI_PROC_NULL`, `tag = MPI_ANY_TAG` and `count = 0`.

# Chapter 4

## Datatypes

Basic datatypes were introduced in Section 3.2.2 Message Data on page 29 and in Section 3.3 Data Type Matching and Data Conversion on page 36. In this chapter, this model is extended to describe any data layout. We consider general datatypes that allow one to transfer efficiently heterogeneous and noncontiguous data. We conclude with the description of calls for explicit packing and unpacking of messages.

### 4.1 Derived Datatypes

Up to here, all point to point communication have involved only buffers containing a sequence of identical basic datatypes. This is too constraining on two accounts. One often wants to pass messages that contain values with different datatypes (e.g., an integer count, followed by a sequence of real numbers); and one often wants to send noncontiguous data (e.g., a sub-block of a matrix). One solution is to pack noncontiguous data into a contiguous buffer at the sender site and unpack it at the receiver site. This has the disadvantage of requiring additional memory-to-memory copy operations at both sites, even when the communication subsystem has scatter-gather capabilities. Instead, MPI provides mechanisms to specify more general, mixed, and noncontiguous communication buffers. It is up to the implementation to decide whether data should be first packed in a contiguous buffer before being transmitted, or whether it can be collected directly from where it resides.

The general mechanisms provided here allow one to transfer directly, without copying, objects of various shape and size. It is not assumed that the MPI library is cognizant of the objects declared in the host language. Thus, if one wants to transfer a structure, or an array section, it will be necessary to provide in MPI a definition of a communication buffer that mimics the definition of the structure or array section in question. These facilities can be used by library designers to define communication functions that can transfer objects defined in the host language — by decoding their definitions as available in a symbol table or a dope vector. Such higher-level communication functions are not part of MPI.

More general communication buffers are specified by replacing the basic datatypes that have been used so far with derived datatypes that are constructed from basic datatypes using the constructors described in this section. These methods of constructing derived datatypes can be applied recursively.

A **general datatype** is an opaque object that specifies two things:

- A sequence of basic datatypes



- A sequence of integer (byte) displacements

The displacements are not required to be positive, distinct, or in increasing order. Therefore, the order of items need not coincide with their order in store, and an item may appear more than once. We call such a pair of sequences (or sequence of pairs) a **type map**. The sequence of basic datatypes (displacements ignored) is the **type signature** of the datatype.

Let

$$Typemap = \{(type_0, disp_0), \dots, (type_{n-1}, disp_{n-1})\},$$

be such a type map, where  $type_i$  are basic types, and  $disp_i$  are displacements. Let

$$Typesig = \{type_0, \dots, type_{n-1}\}$$

be the associated type signature. This type map, together with a base address  $buf$ , specifies a communication buffer: the communication buffer that consists of  $n$  entries, where the  $i$ -th entry is at address  $buf + disp_i$  and has type  $type_i$ . A message assembled from such a communication buffer will consist of  $n$  values, of the types defined by  $Typesig$ .

Most datatype constructors have replication count or block length arguments. Allowed values are non-negative integers. If the value is zero, no elements are generated in the type map and there is no effect on datatype bounds or extent.

We can use a handle to a general datatype as an argument in a send or receive operation, instead of a basic datatype argument. The operation `MPI_SEND(buf, 1, datatype,...)` will use the send buffer defined by the base address `buf` and the general datatype associated with `datatype`; it will generate a message with the type signature determined by the `datatype` argument. `MPI_RECV(buf, 1, datatype,...)` will use the receive buffer defined by the base address `buf` and the general datatype associated with `datatype`.

General datatypes can be used in all send and receive operations. We discuss, in Section 4.1.11, the case where the second argument `count` has value  $> 1$ .

The basic datatypes presented in Section 3.2.2 are particular cases of a general datatype, and are predefined. Thus, `MPI_INT` is a predefined handle to a datatype with type map  $\{(int, 0)\}$ , with one entry of type `int` and displacement zero. The other basic datatypes are similar.

The **extent** of a datatype is defined to be the span from the first byte to the last byte occupied by entries in this datatype, rounded up to satisfy alignment requirements. That is, if

$$Typemap = \{(type_0, disp_0), \dots, (type_{n-1}, disp_{n-1})\},$$

then

$$\begin{aligned} lb(Typemap) &= \min_j disp_j, \\ ub(Typemap) &= \max_j (disp_j + sizeof(type_j)) + \epsilon, \text{ and} \\ extent(Typemap) &= ub(Typemap) - lb(Typemap). \end{aligned} \tag{4.1}$$

If  $type_i$  requires alignment to a byte address that is a multiple of  $k_i$ , then  $\epsilon$  is the least non-negative increment needed to round  $extent(Typemap)$  to the next multiple of  $\max_i k_i$ . The complete definition of **extent** is given on page 99.

**Example 4.1** Assume that  $Type = \{(double, 0), (char, 8)\}$  (a double at displacement zero, followed by a char at displacement eight). Assume, furthermore, that doubles have to be strictly aligned at addresses that are multiples of eight. Then, the extent of this datatype is 16 (9 rounded to the next multiple of 8). A datatype that consists of a character immediately followed by a double will also have an extent of 16.

*Rationale.* The definition of extent is motivated by the assumption that the amount of padding added at the end of each structure in an array of structures is the least needed to fulfill alignment constraints. More explicit control of the extent is provided in Section 4.1.6. Such explicit control is needed in cases where the assumption does not hold, for example, where union types are used. (*End of rationale.*)

#### 4.1.1 Type Constructors with Explicit Addresses

In Fortran, the functions `MPI_TYPE_CREATE_HVECTOR`, `MPI_TYPE_CREATE_HINDEXED`, `MPI_TYPE_CREATE_STRUCT`, and `MPI_GET_ADDRESS` accept arguments of type `INTEGER(KIND=MPI_ADDRESS_KIND)`, wherever arguments of type `MPI_Aint` and `MPI::Aint` are used in C and C++. On Fortran 77 systems that do not support the Fortran 90 `KIND` notation, and where addresses are 64 bits whereas default `INTEGER`s are 32 bits, these arguments will be of type `INTEGER*8`.

#### 4.1.2 Datatype Constructors

**Contiguous** The simplest datatype constructor is `MPI_TYPE_CONTIGUOUS` which allows replication of a datatype into contiguous locations.

`MPI_TYPE_CONTIGUOUS(count, oldtype, newtype)`

IN	count	replication count (non-negative integer)
IN	oldtype	old datatype (handle)
OUT	newtype	new datatype (handle)

```
int MPI_Type_contiguous(int count, MPI_Datatype oldtype,
                        MPI_Datatype *newtype)
```

```
MPI_TYPE_CONTIGUOUS(COUNT, OLDTYPE, NEWTYPE, IERROR)
INTEGER COUNT, OLDTYPE, NEWTYPE, IERROR
```

```
{MPI::Datatype MPI::Datatype::Create_contiguous(int count) const(binding
deprecated, see Section 15.2) }
```

`newtype` is the datatype obtained by concatenating `count` copies of `oldtype`. Concatenation is defined using *extent* as the size of the concatenated copies.

**Example 4.2** Let `oldtype` have type map  $\{(double, 0), (char, 8)\}$ , with extent 16, and let `count = 3`. The type map of the datatype returned by `newtype` is

$\{(double, 0), (char, 8), (double, 16), (char, 24), (double, 32), (char, 40)\};$

i.e., alternating double and char elements, with displacements 0, 8, 16, 24, 32, 40.

In general, assume that the type map of `oldtype` is

$$\{(type_0, disp_0), \dots, (type_{n-1}, disp_{n-1})\},$$

with extent  $ex$ . Then `newtype` has a type map with `count · n` entries defined by:

$$\{(type_0, disp_0), \dots, (type_{n-1}, disp_{n-1}), (type_0, disp_0 + ex), \dots, (type_{n-1}, disp_{n-1} + ex), \\ \dots, (type_0, disp_0 + ex \cdot (\text{count} - 1)), \dots, (type_{n-1}, disp_{n-1} + ex \cdot (\text{count} - 1))\}.$$

**Vector** The function `MPI_TYPE_VECTOR` is a more general constructor that allows replication of a datatype into locations that consist of equally spaced blocks. Each block is obtained by concatenating the same number of copies of the old datatype. The spacing between blocks is a multiple of the extent of the old datatype.

`MPI_TYPE_VECTOR( count, blocklength, stride, oldtype, newtype)`

IN	count	number of blocks (non-negative integer)
IN	blocklength	number of elements in each block (non-negative integer)
IN	stride	number of elements between start of each block (integer)
IN	oldtype	old datatype (handle)
OUT	newtype	new datatype (handle)

```
int MPI_Type_vector(int count, int blocklength, int stride,
    MPI_Datatype oldtype, MPI_Datatype *newtype)
```

```
MPI_TYPE_VECTOR(COUNT, BLOCKLENGTH, STRIDE, OLDTYPE, NEWTYPE, IERROR)
    INTEGER COUNT, BLOCKLENGTH, STRIDE, OLDTYPE, NEWTYPE, IERROR
```

```
{MPI::Datatype MPI::Datatype::Create_vector(int count, int blocklength,
    int stride) const(binding deprecated, see Section 15.2) }
```

**Example 4.3** Assume, again, that `oldtype` has type map  $\{(double, 0), (char, 8)\}$ , with extent 16. A call to `MPI_TYPE_VECTOR( 2, 3, 4, oldtype, newtype)` will create the datatype with type map,

$$\{(double, 0), (char, 8), (double, 16), (char, 24), (double, 32), (char, 40), \\ (double, 64), (char, 72), (double, 80), (char, 88), (double, 96), (char, 104)\}.$$

That is, two blocks with three copies each of the old type, with a stride of 4 elements (4 · 16 bytes) between the blocks.

**Example 4.4** A call to `MPI_TYPE_VECTOR(3, 1, -2, oldtype, newtype)` will create the datatype,

$$\{(\text{double}, 0), (\text{char}, 8), (\text{double}, -32), (\text{char}, -24), (\text{double}, -64), (\text{char}, -56)\}.$$

In general, assume that `oldtype` has type map,

$$\{(type_0, disp_0), \dots, (type_{n-1}, disp_{n-1})\},$$

with extent  $ex$ . Let `bl` be the blocklength. The newly created datatype has a type map with `count · bl · n` entries:

$$\begin{aligned} &\{(type_0, disp_0), \dots, (type_{n-1}, disp_{n-1}), \\ &(type_0, disp_0 + ex), \dots, (type_{n-1}, disp_{n-1} + ex), \dots, \\ &(type_0, disp_0 + (bl - 1) \cdot ex), \dots, (type_{n-1}, disp_{n-1} + (bl - 1) \cdot ex), \\ &(type_0, disp_0 + stride \cdot ex), \dots, (type_{n-1}, disp_{n-1} + stride \cdot ex), \dots, \\ &(type_0, disp_0 + (stride + bl - 1) \cdot ex), \dots, (type_{n-1}, disp_{n-1} + (stride + bl - 1) \cdot ex), \dots, \\ &(type_0, disp_0 + stride \cdot (count - 1) \cdot ex), \dots, \\ &(type_{n-1}, disp_{n-1} + stride \cdot (count - 1) \cdot ex), \dots, \\ &(type_0, disp_0 + (stride \cdot (count - 1) + bl - 1) \cdot ex), \dots, \\ &(type_{n-1}, disp_{n-1} + (stride \cdot (count - 1) + bl - 1) \cdot ex)\}. \end{aligned}$$

A call to `MPI_TYPE_CONTIGUOUS(count, oldtype, newtype)` is equivalent to a call to `MPI_TYPE_VECTOR(count, 1, 1, oldtype, newtype)`, or to a call to `MPI_TYPE_VECTOR(1, count, n, oldtype, newtype)`,  $n$  arbitrary.

**Hvector** The function `MPI_TYPE_CREATE_HVECTOR` is identical to `MPI_TYPE_VECTOR`, except that `stride` is given in bytes, rather than in elements. The use for both types of vector constructors is illustrated in Section 4.1.14. (H stands for “heterogeneous”).

`MPI_TYPE_CREATE_HVECTOR( count, blocklength, stride, oldtype, newtype)`

IN	count	number of blocks (non-negative integer)
IN	blocklength	number of elements in each block (non-negative integer)
IN	stride	number of bytes between start of each block (integer)
IN	oldtype	old datatype (handle)
OUT	newtype	new datatype (handle)

```

1  int MPI_Type_create_hvector(int count, int blocklength, MPI_Aint stride,
2      MPI_Datatype oldtype, MPI_Datatype *newtype)
3
4  MPI_TYPE_CREATE_HVECTOR(COUNT, BLOCKLENGTH, STRIDE, OLDTYPE, NEWTYPE,
5      IERROR)
6      INTEGER COUNT, BLOCKLENGTH, OLDTYPE, NEWTYPE, IERROR
7      INTEGER(KIND=MPI_ADDRESS_KIND) STRIDE
8
9  {MPI::Datatype MPI::Datatype::Create_hvector(int count, int blocklength,
10      MPI::Aint stride) const(binding deprecated, see Section 15.2) }

```

This function replaces MPI\_TYPE\_HVECTOR, whose use is deprecated. See also Chapter 15.

Assume that `oldtype` has type map,

$$\{(type_0, disp_0), \dots, (type_{n-1}, disp_{n-1})\},$$

with extent  $ex$ . Let  $bl$  be the `blocklength`. The newly created datatype has a type map with  $count \cdot bl \cdot n$  entries:

$$\begin{aligned}
 &\{(type_0, disp_0), \dots, (type_{n-1}, disp_{n-1}), \\
 &(type_0, disp_0 + ex), \dots, (type_{n-1}, disp_{n-1} + ex), \dots, \\
 &(type_0, disp_0 + (bl - 1) \cdot ex), \dots, (type_{n-1}, disp_{n-1} + (bl - 1) \cdot ex), \\
 &(type_0, disp_0 + stride), \dots, (type_{n-1}, disp_{n-1} + stride), \dots, \\
 &(type_0, disp_0 + stride + (bl - 1) \cdot ex), \dots, \\
 &(type_{n-1}, disp_{n-1} + stride + (bl - 1) \cdot ex), \dots, \\
 &(type_0, disp_0 + stride \cdot (count - 1)), \dots, (type_{n-1}, disp_{n-1} + stride \cdot (count - 1)), \dots, \\
 &(type_0, disp_0 + stride \cdot (count - 1) + (bl - 1) \cdot ex), \dots, \\
 &(type_{n-1}, disp_{n-1} + stride \cdot (count - 1) + (bl - 1) \cdot ex)\}.
 \end{aligned}$$

**Indexed** The function MPI\_TYPE\_INDEXED allows replication of an old datatype into a sequence of blocks (each block is a concatenation of the old datatype), where each block can contain a different number of copies and have a different displacement. All block displacements are multiples of the old type extent.

```

MPI_TYPE_INDEXED( count, array_of_blocklengths, array_of_displacements, oldtype, new-
    type)
    IN      count      number of blocks – also number of entries in
                        array_of_displacements and array_of_blocklengths (non-
                        negative integer)
    IN      array_of_blocklengths  number of elements per block (array of non-negative
                        integers)
    IN      array_of_displacements  displacement for each block, in multiples of oldtype
                        extent (array of integer)
    IN      oldtype      old datatype (handle)
    OUT     newtype      new datatype (handle)

int MPI_Type_indexed(int count, int *array_of_blocklengths,
    int *array_of_displacements, MPI_Datatype oldtype,
    MPI_Datatype *newtype)

MPI_TYPE_INDEXED(COUNT, ARRAY_OF_BLOCKLENGTHS, ARRAY_OF_DISPLACEMENTS,
    OLDTYPE, NEWTYPE, IERROR)
    INTEGER COUNT, ARRAY_OF_BLOCKLENGTHS(*), ARRAY_OF_DISPLACEMENTS(*),
    OLDTYPE, NEWTYPE, IERROR

{MPI::Datatype MPI::Datatype::Create_indexed(int count,
    const int array_of_blocklengths[],
    const int array_of_displacements[]) const(binding deprecated, see
    Section 15.2) }

```

**Example 4.5**

Let `oldtype` have type map  $\{(\text{double}, 0), (\text{char}, 8)\}$ , with extent 16. Let  $B = (3, 1)$  and let  $D = (4, 0)$ . A call to `MPI_TYPE_INDEXED(2, B, D, oldtype, newtype)` returns a datatype with type map,

$$\{(\text{double}, 64), (\text{char}, 72), (\text{double}, 80), (\text{char}, 88), (\text{double}, 96), (\text{char}, 104),$$

$$(\text{double}, 0), (\text{char}, 8)\}.$$

That is, three copies of the old type starting at displacement 64, and one copy starting at displacement 0.

In general, assume that `oldtype` has type map,

$$\{(type_0, disp_0), \dots, (type_{n-1}, disp_{n-1})\},$$

with extent  $ex$ . Let  $B$  be the `array_of_blocklength` argument and  $D$  be the `array_of_displacements` argument. The newly created datatype has  $n \cdot \sum_{i=0}^{\text{count}-1} B[i]$  entries:

$$\{(type_0, disp_0 + D[0] \cdot ex), \dots, (type_{n-1}, disp_{n-1} + D[0] \cdot ex), \dots,$$

$$(type_0, disp_0 + (D[0] + B[0] - 1) \cdot ex), \dots, (type_{n-1}, disp_{n-1} + (D[0] + B[0] - 1) \cdot ex), \dots,$$

$$\begin{aligned}
& (type_0, disp_0 + D[count-1] \cdot ex), \dots, (type_{n-1}, disp_{n-1} + D[count-1] \cdot ex), \dots, \\
& (type_0, disp_0 + (D[count-1] + B[count-1] - 1) \cdot ex), \dots, \\
& (type_{n-1}, disp_{n-1} + (D[count-1] + B[count-1] - 1) \cdot ex)\}.
\end{aligned}$$

A call to `MPI_TYPE_VECTOR(count, blocklength, stride, oldtype, newtype)` is equivalent to a call to `MPI_TYPE_INDEXED(count, B, D, oldtype, newtype)` where

$$D[j] = j \cdot \text{stride}, \quad j = 0, \dots, \text{count} - 1,$$

and

$$B[j] = \text{blocklength}, \quad j = 0, \dots, \text{count} - 1.$$

**Hindexed** The function `MPI_TYPE_CREATE_HINDEXED` is identical to `MPI_TYPE_INDEXED`, except that block displacements in `array_of_displacements` are specified in bytes, rather than in multiples of the `oldtype` extent.

`MPI_TYPE_CREATE_HINDEXED( count, array_of_blocklengths, array_of_displacements, oldtype, newtype)`

IN	count	number of blocks — also number of entries in <code>array_of_displacements</code> and <code>array_of_blocklengths</code> (non-negative integer)
IN	array_of_blocklengths	number of elements in each block (array of non-negative integers)
IN	array_of_displacements	byte displacement of each block (array of integer)
IN	oldtype	old datatype (handle)
OUT	newtype	new datatype (handle)

```

int MPI_Type_create_hindexed(int count, int array_of_blocklengths[],
                             MPI_Aint array_of_displacements[], MPI_Datatype oldtype,
                             MPI_Datatype *newtype)

```

```

MPI_TYPE_CREATE_HINDEXED(COUNT, ARRAY_OF_BLOCKLENGTHS,
                          ARRAY_OF_DISPLACEMENTS, OLDTYPE, NEWTYPE, IERROR)
INTEGER COUNT, ARRAY_OF_BLOCKLENGTHS(*), OLDTYPE, NEWTYPE, IERROR
INTEGER(KIND=MPI_ADDRESS_KIND) ARRAY_OF_DISPLACEMENTS(*)

```

```

{MPI::Datatype MPI::Datatype::Create_hindexed(int count,
        const int array_of_blocklengths[],
        const MPI::Aint array_of_displacements[]) const (binding deprecated, see Section 15.2) }

```

This function replaces `MPI_TYPE_HINDEXED`, whose use is deprecated. See also Chapter 15.

Assume that `oldtype` has type map,

$$\{(type_0, disp_0), \dots, (type_{n-1}, disp_{n-1})\},$$

with extent  $ex$ . Let `B` be the `array_of_blocklength` argument and `D` be the `array_of_displacements` argument. The newly created datatype has a type map with  $n \cdot \sum_{i=0}^{count-1} B[i]$  entries:

$$\begin{aligned} &\{(type_0, disp_0 + D[0]), \dots, (type_{n-1}, disp_{n-1} + D[0]), \dots, \\ &(type_0, disp_0 + D[0] + (B[0] - 1) \cdot ex), \dots, \\ &(type_{n-1}, disp_{n-1} + D[0] + (B[0] - 1) \cdot ex), \dots, \\ &(type_0, disp_0 + D[count-1]), \dots, (type_{n-1}, disp_{n-1} + D[count-1]), \dots, \\ &(type_0, disp_0 + D[count-1] + (B[count-1] - 1) \cdot ex), \dots, \\ &(type_{n-1}, disp_{n-1} + D[count-1] + (B[count-1] - 1) \cdot ex)\}. \end{aligned}$$

**Indexed\_block** This function is the same as `MPI_TYPE_INDEXED` except that the block-length is the same for all blocks. There are many codes using indirect addressing arising from unstructured grids where the blocksize is always 1 (gather/scatter). The following convenience function allows for constant blocksize and arbitrary displacements.

`MPI_TYPE_CREATE_INDEXED_BLOCK(count, blocklength, array_of_displacements, oldtype, newtype)`

IN	count	length of array of displacements (non-negative integer)
IN	blocklength	size of block (non-negative integer)
IN	array_of_displacements	array of displacements (array of integer)
IN	oldtype	old datatype (handle)
OUT	newtype	new datatype (handle)

```
int MPI_Type_create_indexed_block(int count, int blocklength,
    int array_of_displacements[], MPI_Datatype oldtype,
    MPI_Datatype *newtype)
```

```
MPI_TYPE_CREATE_INDEXED_BLOCK(COUNT, BLOCKLENGTH, ARRAY_OF_DISPLACEMENTS,
    OLDTYPE, NEWTYPE, IERROR)
INTEGER COUNT, BLOCKLENGTH, ARRAY_OF_DISPLACEMENTS(*), OLDTYPE,
NEWTYPE, IERROR
```

```
{MPI::Datatype MPI::Datatype::Create_indexed_block(int count,
    int blocklength,
    const int array_of_displacements[]) const (binding deprecated, see
    Section 15.2) }
```



Struct `MPI_TYPE_STRUCT` is the most general type constructor. It further generalizes `MPI_TYPE_CREATE_HINDEXED` in that it allows each block to consist of replications of different datatypes.

`MPI_TYPE_CREATE_STRUCT(count, array_of_blocklengths, array_of_displacements, array_of_types, newtype)`

IN	count	number of blocks (non-negative integer) — also number of entries in arrays <code>array_of_types</code> , <code>array_of_displacements</code> and <code>array_of_blocklengths</code>
IN	array_of_blocklength	number of elements in each block (array of non-negative integer)
IN	array_of_displacements	byte displacement of each block (array of integer)
IN	array_of_types	type of elements in each block (array of handles to datatype objects)
OUT	newtype	new datatype (handle)

```
int MPI_Type_create_struct(int count, int array_of_blocklengths[],
                          MPI_Aint array_of_displacements[],
                          MPI_Datatype array_of_types[], MPI_Datatype *newtype)
```

```
MPI_TYPE_CREATE_STRUCT(COUNT, ARRAY_OF_BLOCKLENGTHS,
                       ARRAY_OF_DISPLACEMENTS, ARRAY_OF_TYPES, NEWTYPE, IERROR)
INTEGER COUNT, ARRAY_OF_BLOCKLENGTHS(*), ARRAY_OF_TYPES(*), NEWTYPE,
IERROR
INTEGER(KIND=MPI_ADDRESS_KIND) ARRAY_OF_DISPLACEMENTS(*)
```

```
{static MPI::Datatype MPI::Datatype::Create_struct(int count,
            const int array_of_blocklengths[], const MPI::Aint
            array_of_displacements[],
            const MPI::Datatype array_of_types[]) (binding deprecated, see
            Section 15.2) }
```

This function replaces `MPI_TYPE_STRUCT`, whose use is deprecated. See also Chapter 15.

**Example 4.6** Let `type1` have type map,

`{(double, 0), (char, 8)}`,

with extent 16. Let `B = (2, 1, 3)`, `D = (0, 16, 26)`, and `T = (MPI_FLOAT, type1, MPI_CHAR)`. Then a call to `MPI_TYPE_STRUCT(3, B, D, T, newtype)` returns a datatype with type map,

`{(float, 0), (float, 4), (double, 16), (char, 24), (char, 26), (char, 27), (char, 28)}`.

That is, two copies of `MPI_FLOAT` starting at 0, followed by one copy of `type1` starting at 16, followed by three copies of `MPI_CHAR`, starting at 26. (We assume that a float occupies four bytes.)

In general, let  $T$  be the `array_of_types` argument, where  $T[i]$  is a handle to,

$$typemap_i = \{(type_0^i, disp_0^i), \dots, (type_{n_i-1}^i, disp_{n_i-1}^i)\},$$

with extent  $ex_i$ . Let  $B$  be the `array_of_blocklength` argument and  $D$  be the `array_of_displacements` argument. Let  $c$  be the count argument. Then the newly created datatype has a type map with  $\sum_{i=0}^{c-1} B[i] \cdot n_i$  entries:

$$\begin{aligned} &\{(type_0^0, disp_0^0 + D[0]), \dots, (type_{n_0}^0, disp_{n_0}^0 + D[0]), \dots, \\ &(type_0^0, disp_0^0 + D[0] + (B[0] - 1) \cdot ex_0), \dots, (type_{n_0}^0, disp_{n_0}^0 + D[0] + (B[0]-1) \cdot ex_0), \dots, \\ &(type_0^{c-1}, disp_0^{c-1} + D[c-1]), \dots, (type_{n_{c-1}-1}^{c-1}, disp_{n_{c-1}-1}^{c-1} + D[c-1]), \dots, \\ &(type_0^{c-1}, disp_0^{c-1} + D[c-1] + (B[c-1] - 1) \cdot ex_{c-1}), \dots, \\ &(type_{n_{c-1}-1}^{c-1}, disp_{n_{c-1}-1}^{c-1} + D[c-1] + (B[c-1]-1) \cdot ex_{c-1})\}. \end{aligned}$$

A call to `MPI_TYPE_CREATE_HINDEXED(count, B, D, oldtype, newtype)` is equivalent to a call to `MPI_TYPE_CREATE_STRUCT(count, B, D, T, newtype)`, where each entry of  $T$  is equal to `oldtype`.

#### 4.1.3 Subarray Datatype Constructor

`MPI_TYPE_CREATE_SUBARRAY(ndims, array_of_sizes, array_of_subsizes, array_of_starts, order, oldtype, newtype)`

IN	ndims	number of array dimensions (positive integer)
IN	array_of_sizes	number of elements of type <code>oldtype</code> in each dimension of the full array (array of positive integers)
IN	array_of_subsizes	number of elements of type <code>oldtype</code> in each dimension of the subarray (array of positive integers)
IN	array_of_starts	starting coordinates of the subarray in each dimension (array of non-negative integers)
IN	order	array storage order flag (state)
IN	oldtype	array element datatype (handle)
OUT	newtype	new datatype (handle)

```
int MPI_Type_create_subarray(int ndims, int array_of_sizes[],
                           int array_of_subsizes[], int array_of_starts[], int order,
                           MPI_Datatype oldtype, MPI_Datatype *newtype)
```

```
MPI_TYPE_CREATE_SUBARRAY(NDIMS, ARRAY_OF_SIZES, ARRAY_OF_SUBSIZES,
                          ARRAY_OF_STARTS, ORDER, OLDTYPE, NEWTYPE, IERROR)
INTEGER NDIMS, ARRAY_OF_SIZES(*), ARRAY_OF_SUBSIZES(*),
ARRAY_OF_STARTS(*), ORDER, OLDTYPE, NEWTYPE, IERROR
```

```

1 {MPI::Datatype MPI::Datatype::Create_subarray(int ndims,
2         const int array_of_sizes[], const int array_of_subsizes[],
3         const int array_of_starts[], int order) const(binding deprecated,
4         see Section 15.2) }
```

The subarray type constructor creates an MPI datatype describing an  $n$ -dimensional subarray of an  $n$ -dimensional array. The subarray may be situated anywhere within the full array, and may be of any nonzero size up to the size of the larger array as long as it is confined within this array. This type constructor facilitates creating filetypes to access arrays distributed in blocks among processes to a single file that contains the global array, see MPI I/O, especially Section 13.1.1 on page 421.

This type constructor can handle arrays with an arbitrary number of dimensions and works for both C and Fortran ordered matrices (i.e., row-major or column-major). Note that a C program may use Fortran order and a Fortran program may use C order.

The `ndims` parameter specifies the number of dimensions in the full data array and gives the number of elements in `array_of_sizes`, `array_of_subsizes`, and `array_of_starts`.

The number of elements of type `oldtype` in each dimension of the  $n$ -dimensional array and the requested subarray are specified by `array_of_sizes` and `array_of_subsizes`, respectively. For any dimension  $i$ , it is erroneous to specify `array_of_subsizes[i] < 1` or `array_of_subsizes[i] > array_of_sizes[i]`.

The `array_of_starts` contains the starting coordinates of each dimension of the subarray. Arrays are assumed to be indexed starting from zero. For any dimension  $i$ , it is erroneous to specify `array_of_starts[i] < 0` or `array_of_starts[i] > (array_of_sizes[i] - array_of_subsizes[i])`.

*Advice to users.* In a Fortran program with arrays indexed starting from 1, if the starting coordinate of a particular dimension of the subarray is  $n$ , then the entry in `array_of_starts` for that dimension is  $n-1$ . (*End of advice to users.*)

The `order` argument specifies the storage order for the subarray as well as the full array. It must be set to one of the following:

**MPI\_ORDER\_C** The ordering used by C arrays, (i.e., row-major order)

**MPI\_ORDER\_FORTRAN** The ordering used by Fortran arrays, (i.e., column-major order)

A  $ndims$ -dimensional subarray (`newtype`) with no extra padding can be defined by the function `Subarray()` as follows:

```

newtype = Subarray(ndims, {size0, size1, ..., sizendims-1},
                    {subsize0, subsize1, ..., subsizendims-1},
                    {start0, start1, ..., startndims-1}, oldtype)
```

Let the `typemap` of `oldtype` have the form:

$$\{(type_0, disp_0), (type_1, disp_1), \dots, (type_{n-1}, disp_{n-1})\}$$

where  $type_i$  is a predefined MPI datatype, and let  $ex$  be the extent of `oldtype`. Then we define the `Subarray()` function recursively using the following three equations. Equation 4.2 defines the base step. Equation 4.3 defines the recursion step when `order = MPI_ORDER_FORTRAN`, and Equation 4.4 defines the recursion step when `order = MPI_ORDER_C`.

$$\begin{aligned}
& \text{Subarray}(1, \{size_0\}, \{subsize_0\}, \{start_0\}, \\
& \quad \{(type_0, disp_0), (type_1, disp_1), \dots, (type_{n-1}, disp_{n-1})\}) \\
& = \{(\text{MPI\_LB}, 0), \\
& \quad (type_0, disp_0 + start_0 \times ex), \dots, (type_{n-1}, disp_{n-1} + start_0 \times ex), \\
& \quad (type_0, disp_0 + (start_0 + 1) \times ex), \dots, (type_{n-1}, \\
& \quad \quad disp_{n-1} + (start_0 + 1) \times ex), \dots \\
& \quad (type_0, disp_0 + (start_0 + subsize_0 - 1) \times ex), \dots, \\
& \quad \quad (type_{n-1}, disp_{n-1} + (start_0 + subsize_0 - 1) \times ex), \\
& \quad (\text{MPI\_UB}, size_0 \times ex)\}
\end{aligned} \tag{4.2}$$

$$\begin{aligned}
& \text{Subarray}(ndims, \{size_0, size_1, \dots, size_{ndims-1}\}, \\
& \quad \{subsize_0, subsize_1, \dots, subsize_{ndims-1}\}, \\
& \quad \{start_0, start_1, \dots, start_{ndims-1}\}, \text{oldtype}) \\
& = \text{Subarray}(ndims - 1, \{size_1, size_2, \dots, size_{ndims-1}\}, \\
& \quad \{subsize_1, subsize_2, \dots, subsize_{ndims-1}\}, \\
& \quad \{start_1, start_2, \dots, start_{ndims-1}\}, \\
& \quad \text{Subarray}(1, \{size_0\}, \{subsize_0\}, \{start_0\}, \text{oldtype}))
\end{aligned} \tag{4.3}$$

$$\begin{aligned}
& \text{Subarray}(ndims, \{size_0, size_1, \dots, size_{ndims-1}\}, \\
& \quad \{subsize_0, subsize_1, \dots, subsize_{ndims-1}\}, \\
& \quad \{start_0, start_1, \dots, start_{ndims-1}\}, \text{oldtype}) \\
& = \text{Subarray}(ndims - 1, \{size_0, size_1, \dots, size_{ndims-2}\}, \\
& \quad \{subsize_0, subsize_1, \dots, subsize_{ndims-2}\}, \\
& \quad \{start_0, start_1, \dots, start_{ndims-2}\}, \\
& \quad \text{Subarray}(1, \{size_{ndims-1}\}, \{subsize_{ndims-1}\}, \{start_{ndims-1}\}, \text{oldtype}))
\end{aligned} \tag{4.4}$$

For an example use of `MPI_TYPE_CREATE_SUBARRAY` in the context of I/O see Section 13.9.2.

#### 4.1.4 Distributed Array Datatype Constructor

The distributed array type constructor supports HPF-like [34] data distributions. However, unlike in HPF, the storage order may be specified for C arrays as well as for Fortran arrays.

*Advice to users.* One can create an HPF-like file view using this type constructor as follows. Complementary filetypes are created by having every process of a group call this constructor with identical arguments (with the exception of `rank` which should be set appropriately). These filetypes (along with identical `disp` and `etype`) are then used to define the view (via `MPI_FILE_SET_VIEW`), see MPI I/O, especially Section 13.1.1 on page 421 and Section 13.3 on page 433. Using this view, a collective data access operation (with identical offsets) will yield an HPF-like distribution pattern. (*End of advice to users.*)

```

1 MPI_TYPE_CREATE_DARRAY(size, rank, ndims, array_of_gsizes, array_of_distrib,
2   array_of_dargs, array_of_psize, order, oldtype, newtype)
3
4   IN      size                size of process group (positive integer)
5   IN      rank                rank in process group (non-negative integer)
6   IN      ndims               number of array dimensions as well as process grid
7   dimensions (positive integer)
8   IN      array_of_gsizes      number of elements of type oldtype in each dimension
9   of global array (array of positive integers)
10  IN      array_of_distrib      distribution of array in each dimension (array of state)
11  IN      array_of_dargs        distribution argument in each dimension (array of positive
12  integers)
13  IN      array_of_psize        size of process grid in each dimension (array of positive
14  integers)
15  IN      order                array storage order flag (state)
16  IN      oldtype              old datatype (handle)
17  OUT     newtype              new datatype (handle)

```

```

21 int MPI_Type_create_darray(int size, int rank, int ndims,
22   int array_of_gsizes[], int array_of_distrib[], int
23   array_of_dargs[], int array_of_psize[], int order,
24   MPI_Datatype oldtype, MPI_Datatype *newtype)
25
26 MPI_TYPE_CREATE_DARRAY(SIZE, RANK, NDIMS, ARRAY_OF_GSIZES,
27   ARRAY_OF_DISTRIBS, ARRAY_OF_DARGS, ARRAY_OF_PSIZE, ORDER,
28   OLDTYPE, NEWTYPE, IERROR)
29
30 INTEGER SIZE, RANK, NDIMS, ARRAY_OF_GSIZES(*), ARRAY_OF_DISTRIBS(*),
31   ARRAY_OF_DARGS(*), ARRAY_OF_PSIZE(*), ORDER, OLDTYPE, NEWTYPE, IERROR
32
33 {MPI::Datatype MPI::Datatype::Create_darray(int size, int rank, int ndims,
34   const int array_of_gsizes[], const int array_of_distrib[],
35   const int array_of_dargs[], const int array_of_psize[],
36   int order) const(binding deprecated, see Section 15.2) }

```

MPI\_TYPE\_CREATE\_DARRAY can be used to generate the datatypes corresponding to the distribution of an `ndims`-dimensional array of `oldtype` elements onto an `ndims`-dimensional grid of logical processes. Unused dimensions of `array_of_psize` should be set to 1. (See Example 4.7, page 95.) For a call to MPI\_TYPE\_CREATE\_DARRAY to be correct, the equation  $\prod_{i=0}^{ndims-1} array\_of\_psize[i] = size$  must be satisfied. The ordering of processes in the process grid is assumed to be row-major, as in the case of virtual Cartesian process topologies.

*Advice to users.* For both Fortran and C arrays, the ordering of processes in the process grid is assumed to be row-major. This is consistent with the ordering used in virtual Cartesian process topologies in MPI. To create such virtual process topologies, or to find the coordinates of a process in the process grid, etc., users may use the corresponding process topology functions, see Chapter 7 on page 277. (*End of advice to users.*)

Each dimension of the array can be distributed in one of three ways:

- MPI\_DISTRIBUTE\_BLOCK - Block distribution
- MPI\_DISTRIBUTE\_CYCLIC - Cyclic distribution
- MPI\_DISTRIBUTE\_NONE - Dimension not distributed.

The constant MPI\_DISTRIBUTE\_DFLT\_DARG specifies a default distribution argument. The distribution argument for a dimension that is not distributed is ignored. For any dimension  $i$  in which the distribution is MPI\_DISTRIBUTE\_BLOCK, it is erroneous to specify  $\text{array\_of\_dargs}[i] * \text{array\_of\_psizes}[i] < \text{array\_of\_gsizes}[i]$ .

For example, the HPF layout `ARRAY(CYCLIC(15))` corresponds to MPI\_DISTRIBUTE\_CYCLIC with a distribution argument of 15, and the HPF layout `ARRAY(BLOCK)` corresponds to MPI\_DISTRIBUTE\_BLOCK with a distribution argument of MPI\_DISTRIBUTE\_DFLT\_DARG.

The `order` argument is used as in MPI\_TYPE\_CREATE\_SUBARRAY to specify the storage order. Therefore, arrays described by this type constructor may be stored in Fortran (column-major) or C (row-major) order. Valid values for `order` are MPI\_ORDER\_FORTRAN and MPI\_ORDER\_C.

This routine creates a new MPI datatype with a typemap defined in terms of a function called “cyclic()” (see below).

Without loss of generality, it suffices to define the typemap for the MPI\_DISTRIBUTE\_CYCLIC case where MPI\_DISTRIBUTE\_DFLT\_DARG is not used.

MPI\_DISTRIBUTE\_BLOCK and MPI\_DISTRIBUTE\_NONE can be reduced to the MPI\_DISTRIBUTE\_CYCLIC case for dimension  $i$  as follows.

MPI\_DISTRIBUTE\_BLOCK with  $\text{array\_of\_dargs}[i]$  equal to MPI\_DISTRIBUTE\_DFLT\_DARG is equivalent to MPI\_DISTRIBUTE\_CYCLIC with  $\text{array\_of\_dargs}[i]$  set to

$$(\text{array\_of\_gsizes}[i] + \text{array\_of\_psizes}[i] - 1) / \text{array\_of\_psizes}[i].$$

If  $\text{array\_of\_dargs}[i]$  is not MPI\_DISTRIBUTE\_DFLT\_DARG, then MPI\_DISTRIBUTE\_BLOCK and MPI\_DISTRIBUTE\_CYCLIC are equivalent.

MPI\_DISTRIBUTE\_NONE is equivalent to MPI\_DISTRIBUTE\_CYCLIC with  $\text{array\_of\_dargs}[i]$  set to  $\text{array\_of\_gsizes}[i]$ .

Finally, MPI\_DISTRIBUTE\_CYCLIC with  $\text{array\_of\_dargs}[i]$  equal to MPI\_DISTRIBUTE\_DFLT\_DARG is equivalent to MPI\_DISTRIBUTE\_CYCLIC with  $\text{array\_of\_dargs}[i]$  set to 1.

For MPI\_ORDER\_FORTRAN, an  $\text{ndims}$ -dimensional distributed array (`newtype`) is defined by the following code fragment:

```
oldtype[0] = oldtype;
for ( i = 0; i < ndims; i++ ) {
    oldtype[i+1] = cyclic(array_of_dargs[i],
                        array_of_gsizes[i],
                        r[i],
                        array_of_psizes[i],
                        oldtype[i]);
}
newtype = oldtype[ndims];
```

For MPI\_ORDER\_C, the code is:

```

oldtype[0] = oldtype;
for ( i = 0; i < ndims; i++ ) {
    oldtype[i + 1] = cyclic(array_of_dargs[ndims - i - 1],
                           array_of_gsizes[ndims - i - 1],
                           r[ndims - i - 1],
                           array_of_psize[ndims - i - 1],
                           oldtype[i]);
}
newtype = oldtype[ndims];

```

where  $r[i]$  is the position of the process (with rank  $rank$ ) in the process grid at dimension  $i$ . The values of  $r[i]$  are given by the following code fragment:

```

t_rank = rank;
t_size = 1;
for (i = 0; i < ndims; i++)
    t_size *= array_of_psize[i];
for (i = 0; i < ndims; i++) {
    t_size = t_size / array_of_psize[i];
    r[i] = t_rank / t_size;
    t_rank = t_rank % t_size;
}

```

Let the typemap of `oldtype` have the form:

$$\{(type_0, disp_0), (type_1, disp_1), \dots, (type_{n-1}, disp_{n-1})\}$$

where  $type_i$  is a predefined MPI datatype, and let  $ex$  be the extent of `oldtype`.

Given the above, the function `cyclic()` is defined as follows:

```

cyclic(darg, gsize, r, psize, oldtype)
= { (MPI_LB, 0),
    (type0, disp0 + r × darg × ex), ...,
    (typen-1, dispn-1 + r × darg × ex),
    (type0, disp0 + (r × darg + 1) × ex), ...,
    (typen-1, dispn-1 + (r × darg + 1) × ex),
    ...
    (type0, disp0 + ((r + 1) × darg - 1) × ex), ...,
    (typen-1, dispn-1 + ((r + 1) × darg - 1) × ex),
    (type0, disp0 + r × darg × ex + psize × darg × ex), ...,
    (typen-1, dispn-1 + r × darg × ex + psize × darg × ex),
    (type0, disp0 + (r × darg + 1) × ex + psize × darg × ex), ...,
    (typen-1, dispn-1 + (r × darg + 1) × ex + psize × darg × ex),

```

```

...
(type0, disp0 + ((r + 1) × darg - 1) × ex + psize × darg × ex), ...,
(typen-1, dispn-1 + ((r + 1) × darg - 1) × ex + psize × darg × ex),
:
(type0, disp0 + r × darg × ex + psize × darg × ex × (count - 1)), ...,
(typen-1, dispn-1 + r × darg × ex + psize × darg × ex × (count - 1)),
(type0, disp0 + (r × darg + 1) × ex + psize × darg × ex × (count - 1)), ...,
(typen-1, dispn-1 + (r × darg + 1) × ex
+ psize × darg × ex × (count - 1)),
...
(type0, disp0 + (r × darg + darglast - 1) × ex
+ psize × darg × ex × (count - 1)), ...,
(typen-1, dispn-1 + (r × darg + darglast - 1) × ex
+ psize × darg × ex × (count - 1)),
(MPI_UB, gsize × ex)

```

where *count* is defined by this code fragment:

```

nblocks = (gsize + (darg - 1)) / darg;
count = nblocks / psize;
left_over = nblocks - count * psize;
if (r < left_over)
    count = count + 1;

```

Here, *nblocks* is the number of blocks that must be distributed among the processors. Finally, *darg<sub>last</sub>* is defined by this code fragment:

```

if ((num_in_last_cyclic = gsize % (psize * darg)) == 0)
    darg_last = darg;
else
    darg_last = num_in_last_cyclic - darg * r;
    if (darg_last > darg)
        darg_last = darg;
    if (darg_last <= 0)
        darg_last = darg;

```

**Example 4.7** Consider generating the filetypes corresponding to the HPF distribution:

```

<oldtype> FILEARRAY(100, 200, 300)
!HPF$ PROCESSORS PROCESSES(2, 3)
!HPF$ DISTRIBUTE FILEARRAY(CYCLIC(10), *, BLOCK) ONTO PROCESSES

```

This can be achieved by the following Fortran code, assuming there will be six processes attached to the run:



```

1      ndims = 3
2      array_of_gsizes(1) = 100
3      array_of_distribs(1) = MPI_DISTRIBUTE_CYCLIC
4      array_of_dargs(1) = 10
5      array_of_gsizes(2) = 200
6      array_of_distribs(2) = MPI_DISTRIBUTE_NONE
7      array_of_dargs(2) = 0
8      array_of_gsizes(3) = 300
9      array_of_distribs(3) = MPI_DISTRIBUTE_BLOCK
10     array_of_dargs(3) = MPI_DISTRIBUTE_DFLT_DARG
11     array_of_psize(1) = 2
12     array_of_psize(2) = 1
13     array_of_psize(3) = 3
14     call MPI_COMM_SIZE(MPI_COMM_WORLD, size, ierr)
15     call MPI_COMM_RANK(MPI_COMM_WORLD, rank, ierr)
16     call MPI_TYPE_CREATE_DARRAY(size, rank, ndims, array_of_gsizes, &
17                                array_of_distribs, array_of_dargs, array_of_psize, &
18                                MPI_ORDER_FORTRAN, oldtype, newtype, ierr)
19

```

#### 4.1.5 Address and Size Functions

The displacements in a general datatype are relative to some initial buffer address. **Absolute addresses** can be substituted for these displacements: we treat them as displacements relative to “address zero,” the start of the address space. This initial address zero is indicated by the constant `MPI_BOTTOM`. Thus, a datatype can specify the absolute address of the entries in the communication buffer, in which case the `buf` argument is passed the value `MPI_BOTTOM`.

The address of a location in memory can be found by invoking the function `MPI_GET_ADDRESS`.

```

31 MPI_GET_ADDRESS(location, address)
32

```

IN	location	location in caller memory (choice)
OUT	address	address of location (integer)

```

36 int MPI_Get_address(void *location, MPI_Aint *address)
37

```

```

38 MPI_GET_ADDRESS(LOCATION, ADDRESS, IERROR)
39     <type> LOCATION(*)
40     INTEGER IERROR
41     INTEGER(KIND=MPI_ADDRESS_KIND) ADDRESS

```

```

42 {MPI::Aint MPI::Get_address(void* location) (binding deprecated, see Section 15.2)
43     }
44

```

This function replaces `MPI_ADDRESS`, whose use is deprecated. See also Chapter 15. Returns the (byte) address of `location`.

*Advice to users.* Current Fortran MPI codes will run unmodified, and will port

to any system. However, they may fail if addresses larger than  $2^{32} - 1$  are used in the program. New codes should be written so that they use the new functions. This provides compatibility with C/C++ and avoids errors on 64 bit architectures. However, such newly written codes may need to be (slightly) rewritten to port to old Fortran 77 environments that do not support KIND declarations. (*End of advice to users.*)

**Example 4.8** Using MPI\_GET\_ADDRESS for an array.

```
REAL A(100,100)
INTEGER(KIND=MPI_ADDRESS_KIND) I1, I2, DIFF
CALL MPI_GET_ADDRESS(A(1,1), I1, IERROR)
CALL MPI_GET_ADDRESS(A(10,10), I2, IERROR)
DIFF = I2 - I1
! The value of DIFF is 909*sizeofreal; the values of I1 and I2 are
! implementation dependent.
```

*Advice to users.* C users may be tempted to avoid the usage of MPI\_GET\_ADDRESS and rely on the availability of the address operator &. Note, however, that & *cast-expression* is a pointer, not an address. ISO C does not require that the value of a pointer (or the pointer cast to int) be the absolute address of the object pointed at — although this is commonly the case. Furthermore, referencing may not have a unique definition on machines with a segmented address space. The use of MPI\_GET\_ADDRESS to “reference” C variables guarantees portability to such machines as well. (*End of advice to users.*)

*Advice to users.* To prevent problems with the argument copying and register optimization done by Fortran compilers, please note the hints in subsections “Problems Due to Data Copying and Sequence Association,” and “A Problem with Register Optimization” in Section 16.2.2 on pages 514 and 517. (*End of advice to users.*)

The following auxillary [\[function provides\]](#)[functions provide](#) useful information on derived datatypes.

MPI\_TYPE\_SIZE(datatype, size)

IN	datatype	datatype (handle)
OUT	size	datatype size (integer)

```
int MPI_Type_size(MPI_Datatype datatype, int *size)
```

```
MPI_TYPE_SIZE(DATATYPE, SIZE, IERROR)
```

```
INTEGER DATATYPE, SIZE, IERROR
```

```
{int MPI::Datatype::Get_size() const(binding deprecated, see Section 15.2) }
```

```

1 MPI_TYPE_SIZE_X(datatype, size)
2     IN      datatype      datatype (handle)
3     OUT     size          datatype size (integer)
4
5
6 int MPI_Type_size_x(MPI_Datatype datatype, MPI_Count *size)
7
8 MPI_TYPE_SIZE_X(DATATYPE, SIZE, IERROR)
9     INTEGER DATATYPE, IERROR
10    INTEGER (KIND=MPI_COUNT_KIND) SIZE

```

`MPI_TYPE_SIZE` and `MPI_TYPE_SIZE_X` [returns] set the value of `size` to the total size, in bytes, of the entries in the type signature associated with `datatype`; i.e., the total size of the data in a message that would be created with this datatype. Entries that occur multiple times in the datatype are counted with their multiplicity.

If the total size of the datatype can not be expressed by the `size` parameter, then `MPI_TYPE_SIZE` and `MPI_TYPE_SIZE_X` set the value of `size` to `MPI_UNDEFINED`.

#### 4.1.6 Lower-Bound and Upper-Bound Markers

It is often convenient to define explicitly the lower bound and upper bound of a type map, and override the definition given on page 99. This allows one to define a datatype that has “holes” at its beginning or its end, or a datatype with entries that extend above the upper bound or below the lower bound. Examples of such usage are provided in Section 4.1.14. Also, the user may want to override the alignment rules that are used to compute upper bounds and extents. E.g., a C compiler may allow the user to override default alignment rules for some of the structures within a program. The user has to specify explicitly the bounds of the datatypes that match these structures.

To achieve this, we add two additional “pseudo-datatypes,” `MPI_LB` and `MPI_UB`, that can be used, respectively, to mark the lower bound or the upper bound of a datatype. These pseudo-datatypes occupy no space ( $extent(MPI\_LB) = extent(MPI\_UB) = 0$ ). They do not affect the size or count of a datatype, and do not affect the content of a message created with this datatype. However, they do affect the definition of the extent of a datatype and, therefore, affect the outcome of a replication of this datatype by a datatype constructor.

**Example 4.9** Let  $D = (-3, 0, 6)$ ;  $T = (MPI\_LB, MPI\_INT, MPI\_UB)$ , and  $B = (1, 1, 1)$ . Then a call to `MPI_TYPE_STRUCT(3, B, D, T, type1)` creates a new datatype that has an extent of 9 (from -3 to 5, 5 included), and contains an integer at displacement 0. This is the datatype defined by the sequence  $\{(lb, -3), (int, 0), (ub, 6)\}$ . If this type is replicated twice by a call to `MPI_TYPE_CONTIGUOUS(2, type1, type2)` then the newly created type can be described by the sequence  $\{(lb, -3), (int, 0), (int, 9), (ub, 15)\}$ . (An entry of type `ub` can be deleted if there is another entry of type `ub` with a higher displacement; an entry of type `lb` can be deleted if there is another entry of type `lb` with a lower displacement.)

In general, if

$$Typemap = \{(type_0, disp_0), \dots, (type_{n-1}, disp_{n-1})\},$$

then the **lower bound** of *Typemap* is defined to be

$$lb(Typemap) = \begin{cases} \min_j disp_j & \text{if no entry has basic type lb} \\ \min_j \{disp_j \text{ such that } type_j = lb\} & \text{otherwise} \end{cases}$$

Similarly, the **upper bound** of *Typemap* is defined to be

$$ub(Typemap) = \begin{cases} \max_j disp_j + sizeof(type_j) + \epsilon & \text{if no entry has basic type } ub \\ \max_j \{ disp_j \text{ such that } type_j = ub \} & \text{otherwise} \end{cases}$$

Then

$$extent(Typemap) = ub(Typemap) - lb(Typemap)$$

If  $type_i$  requires alignment to a byte address that is a multiple of  $k_i$ , then  $\epsilon$  is the least non-negative increment needed to round  $extent(Typemap)$  to the next multiple of  $\max_i k_i$ .

The formal definitions given for the various datatype constructors apply now, with the amended definition of **extent**.

#### 4.1.7 Extent and Bounds of Datatypes

The [following function replaces]functions **MPI\_TYPE\_GET\_EXTENT** and **MPI\_TYPE\_GET\_EXTENT\_X** replace the three functions **MPI\_TYPE\_UB**, **MPI\_TYPE\_LB** and **MPI\_TYPE\_EXTENT** [. It also returns]and also return address sized integers[,] in the Fortran binding. The use of **MPI\_TYPE\_UB**, **MPI\_TYPE\_LB** and **MPI\_TYPE\_EXTENT** is deprecated.

**MPI\_TYPE\_GET\_EXTENT**(datatype, lb, extent)

IN	datatype	datatype to get information on (handle)
OUT	lb	lower bound of datatype (integer)
OUT	extent	extent of datatype (integer)

```
int MPI_Type_get_extent(MPI_Datatype datatype, MPI_Aint *lb,
    MPI_Aint *extent)
```

```
MPI_TYPE_GET_EXTENT(DATATYPE, LB, EXTENT, IERROR)
    INTEGER DATATYPE, IERROR
    INTEGER(KIND = MPI_ADDRESS_KIND) LB, EXTENT
```

```
{void MPI::Datatype::Get_extent(MPI::Aint& lb, MPI::Aint& extent)
    const(binding deprecated, see Section 15.2) }
```

**MPI\_TYPE\_GET\_EXTENT\_X**(datatype, lb, extent)

IN	datatype	datatype to get information on (handle)
OUT	lb	lower bound of datatype (integer)
OUT	extent	extent of datatype (integer)

```
int MPI_Type_get_extent_x(MPI_Datatype datatype, MPI_Count *lb,
    MPI_Count *extent)
```

```
MPI_TYPE_GET_EXTENT_X(DATATYPE, LB, EXTENT, IERROR)
    INTEGER DATATYPE, IERROR
```

```
1      INTEGER(KIND = MPI_COUNT_KIND) LB, EXTENT
```

2 Returns the lower bound and the extent of `datatype` (as defined in Section 4.1.6 on page 98).

3 MPI allows one to change the extent of a datatype, using lower bound and upper  
4 bound markers (`MPI_LB` and `MPI_UB`). This is useful, as it allows to control the stride of  
5 successive datatypes that are replicated by datatype constructors, or are replicated by the  
6 `count` argument in a send or receive call. However, the current mechanism for achieving  
7 it is painful; also it is restrictive. `MPI_LB` and `MPI_UB` are “sticky”: once present in a  
8 datatype, they cannot be overridden (e.g., the upper bound can be moved up, by adding  
9 a new `MPI_UB` marker, but cannot be moved down below an existing `MPI_UB` marker). A  
10 new type constructor is provided to facilitate these changes. The use of `MPI_LB` and `MPI_UB`  
11 is deprecated.

```
12  
13  
14  
15 MPI_TYPE_CREATE_RESIZED(oldtype, lb, extent, newtype)
```

16	IN	oldtype	input datatype (handle)
17	IN	lb	new lower bound of datatype (integer)
18	IN	extent	new extent of datatype (integer)
19	OUT	newtype	output datatype (handle)

```
20  
21  
22  
23 int MPI_Type_create_resized(MPI_Datatype oldtype, MPI_Aint lb, MPI_Aint  
24 extent, MPI_Datatype *newtype)
```

```
25 MPI_TYPE_CREATE_RESIZED(OLDTYPE, LB, EXTENT, NEWTYPE, IERROR)  
26     INTEGER OLDTYPE, NEWTYPE, IERROR  
27     INTEGER(KIND=MPI_ADDRESS_KIND) LB, EXTENT
```

```
28  
29 {MPI::Datatype MPI::Datatype::Create_resized(const MPI::Aint lb,  
30 const MPI::Aint extent) const(binding deprecated, see Section 15.2) }
```

31 Returns in `newtype` a handle to a new datatype that is identical to `oldtype`, except that  
32 the lower bound of this new datatype is set to be `lb`, and its upper bound is set to be `lb`  
33 + `extent`. Any previous `lb` and `ub` markers are erased, and a new pair of lower bound and  
34 upper bound markers are put in the positions indicated by the `lb` and `extent` arguments.  
35 This affects the behavior of the datatype when used in communication operations, with  
36 `count > 1`, and when used in the construction of new derived datatypes.

37  
38 *Advice to users.* It is strongly recommended that users use these two new functions,  
39 rather than the old MPI-1 functions to set and access lower bound, upper bound and  
40 extent of datatypes. (*End of advice to users.*)

#### 41 4.1.8 True Extent of Datatypes

42 Suppose we implement gather (see also Section 5.5 on page 142) as a spanning tree im-  
43 plemented on top of point-to-point routines. Since the receive buffer is only valid on the  
44 root process, one will need to allocate some temporary space for receiving data on interme-  
45 diate nodes. However, the datatype extent cannot be used as an estimate of the amount  
46 of space that needs to be allocated, if the user has modified the extent using the `MPI_UB`  
47

and MPI\_LB values. [ A function is]The functions MPI\_TYPE\_GET\_TRUE\_EXTENT and MPI\_TYPE\_GET\_TRUE\_EXTENT\_X are provided which [returns]return the true extent of the datatype.

MPI\_TYPE\_GET\_TRUE\_EXTENT(datatype, true\_lb, true\_extent)

IN	datatype	datatype to get information on (handle)
OUT	true_lb	true lower bound of datatype (integer)
OUT	true_extent	true size of datatype (integer)

```
int MPI_Type_get_true_extent(MPI_Datatype datatype, MPI_Aint *true_lb,
                             MPI_Aint *true_extent)
```

```
MPI_TYPE_GET_TRUE_EXTENT(DATATYPE, TRUE_LB, TRUE_EXTENT, IERROR)
    INTEGER DATATYPE, IERROR
    INTEGER(KIND = MPI_ADDRESS_KIND) TRUE_LB, TRUE_EXTENT
```

```
{void MPI::Datatype::Get_true_extent(MPI::Aint& true_lb,
                                     MPI::Aint& true_extent) const(binding deprecated, see Section 15.2) }
```

MPI\_TYPE\_GET\_TRUE\_EXTENT\_X(datatype, true\_lb, true\_extent)

IN	datatype	datatype to get information on (handle)
OUT	true_lb	true lower bound of datatype (integer)
OUT	true_extent	true size of datatype (integer)

```
int MPI_Type_get_true_extent_x(MPI_Datatype datatype, MPI_Count *true_lb,
                               MPI_Count *true_extent)
```

```
MPI_TYPE_GET_TRUE_EXTENT_X(DATATYPE, TRUE_LB, TRUE_EXTENT, IERROR)
    INTEGER DATATYPE, IERROR
    INTEGER(KIND = MPI_COUNT_KIND) TRUE_LB, TRUE_EXTENT
```

true\_lb returns the offset of the lowest unit of store which is addressed by the datatype, i.e., the lower bound of the corresponding typemap, ignoring MPI\_LB markers. true\_extent returns the true size of the datatype, i.e., the extent of the corresponding typemap, ignoring MPI\_LB and MPI\_UB markers, and performing no rounding for alignment. If the typemap associated with datatype is

$$Typemap = \{(type_0, disp_0), \dots, (type_{n-1}, disp_{n-1})\}$$

Then

$$true\_lb(Typemap) = \min_j \{disp_j : type_j \neq lb, ub\},$$

$$true\_ub(Typemap) = \max_j \{disp_j + sizeof(type_j) : type_j \neq lb, ub\},$$

and

$$true\_extent(Typemap) = true\_ub(Typemap) - true\_lb(Typemap).$$

(Readers should compare this with the definitions in Section 4.1.6 on page 98 and Section 4.1.7 on page 99, which describe the function `MPI_TYPE_GET_EXTENT`.)

The `true_extent` is the minimum number of bytes of memory necessary to hold a datatype, uncompressed.

#### 4.1.9 Commit and Free

A datatype object has to be **committed** before it can be used in a communication. As an argument in datatype constructors, uncommitted and also committed datatypes can be used. There is no need to commit basic datatypes. They are “pre-committed.”

`MPI_TYPE_COMMIT(datatype)`

INOUT     datatype                             datatype that is committed (handle)

`int MPI_Type_commit(MPI_Datatype *datatype)`

`MPI_TYPE_COMMIT(DATATYPE, IERROR)`

INTEGER DATATYPE, IERROR

`{void MPI::Datatype::Commit() (binding deprecated, see Section 15.2) }`

The commit operation commits the datatype, that is, the formal description of a communication buffer, not the content of that buffer. Thus, after a datatype has been committed, it can be repeatedly reused to communicate the changing content of a buffer or, indeed, the content of different buffers, with different starting addresses.

*Advice to implementors.* The system may “compile” at commit time an internal representation for the datatype that facilitates communication, e.g. change from a compacted representation to a flat representation of the datatype, and select the most convenient transfer mechanism. (*End of advice to implementors.*)

`MPI_TYPE_COMMIT` will accept a committed datatype; in this case, it is equivalent to a no-op.

**Example 4.10** The following code fragment gives examples of using `MPI_TYPE_COMMIT`.

```
INTEGER type1, type2
CALL MPI_TYPE_CONTIGUOUS(5, MPI_REAL, type1, ierr)
      ! new type object created
CALL MPI_TYPE_COMMIT(type1, ierr)
      ! now type1 can be used for communication
type2 = type1
      ! type2 can be used for communication
      ! (it is a handle to same object as type1)
CALL MPI_TYPE_VECTOR(3, 5, 4, MPI_REAL, type1, ierr)
      ! new uncommitted type object created
CALL MPI_TYPE_COMMIT(type1, ierr)
      ! now type1 can be used anew for communication
```

MPI\_TYPE\_FREE(datatype)

INOUT datatype datatype that is freed (handle)

int MPI\_Type\_free(MPI\_Datatype \*datatype)

MPI\_TYPE\_FREE(DATATYPE, IERROR)

INTEGER DATATYPE, IERROR

{void MPI::Datatype::Free() (*binding deprecated, see Section 15.2*) }

Marks the datatype object associated with `datatype` for deallocation and sets `datatype` to `MPI_DATATYPE_NULL`. Any communication that is currently using this datatype will complete normally. Freeing a datatype does not affect any other datatype that was built from the freed datatype. The system behaves as if input datatype arguments to derived datatype constructors are passed by value.

*Advice to implementors.* The implementation may keep a reference count of active communications that use the datatype, in order to decide when to free it. Also, one may implement constructors of derived datatypes so that they keep pointers to their datatype arguments, rather than copying them. In this case, one needs to keep track of active datatype definition references in order to know when a datatype object can be freed. (*End of advice to implementors.*)

#### 4.1.10 Duplicating a Datatype

MPI\_TYPE\_DUP(type, newtype)

IN type datatype (handle)  
OUT newtype copy of type (handle)

int MPI\_Type\_dup(MPI\_Datatype type, MPI\_Datatype \*newtype)

MPI\_TYPE\_DUP(TYPE, NEWTYPE, IERROR)

INTEGER TYPE, NEWTYPE, IERROR

{MPI::Datatype MPI::Datatype::Dup() const (*binding deprecated, see Section 15.2*) }

`MPI_TYPE_DUP` is a type constructor which duplicates the existing `type` with associated key values. For each key value, the respective copy callback function determines the attribute value associated with this key in the new communicator; one particular action that a copy callback may take is to delete the attribute from the new datatype. Returns in `newtype` a new datatype with exactly the same properties as `type` and any copied cached information, see Section 6.7.4 on page 265. The new datatype has identical upper bound and lower bound and yields the same net result when fully decoded with the functions in Section 4.1.13. The `newtype` has the same committed state as the old `type`.



#### 4.1.11 Use of General Datatypes in Communication

Handles to derived datatypes can be passed to a communication call wherever a datatype argument is required. A call of the form `MPI_SEND(buf, count, datatype, ...)`, where `count > 1`, is interpreted as if the call was passed a new datatype which is the concatenation of `count` copies of `datatype`. Thus, `MPI_SEND(buf, count, datatype, dest, tag, comm)` is equivalent to,

```
MPI_TYPE_CONTIGUOUS(count, datatype, newtype)
MPI_TYPE_COMMIT(newtype)
MPI_SEND(buf, 1, newtype, dest, tag, comm).
```

Similar statements apply to all other communication functions that have a `count` and `datatype` argument.

Suppose that a send operation `MPI_SEND(buf, count, datatype, dest, tag, comm)` is executed, where `datatype` has type map,

$$\{(type_0, disp_0), \dots, (type_{n-1}, disp_{n-1})\},$$

and extent *extent*. (Empty entries of “pseudo-type” `MPI_UB` and `MPI_LB` are not listed in the type map, but they affect the value of *extent*.) The send operation sends  $n \cdot \text{count}$  entries, where entry  $i \cdot n + j$  is at location  $addr_{i,j} = \text{buf} + \text{extent} \cdot i + disp_j$  and has type  $type_j$ , for  $i = 0, \dots, \text{count} - 1$  and  $j = 0, \dots, n - 1$ . These entries need not be contiguous, nor distinct; their order can be arbitrary.

The variable stored at address  $addr_{i,j}$  in the calling program should be of a type that matches  $type_j$ , where type matching is defined as in Section 3.3.1. The message sent contains  $n \cdot \text{count}$  entries, where entry  $i \cdot n + j$  has type  $type_j$ .

Similarly, suppose that a receive operation `MPI_RECV(buf, count, datatype, source, tag, comm, status)` is executed, where `datatype` has type map,

$$\{(type_0, disp_0), \dots, (type_{n-1}, disp_{n-1})\},$$

with extent *extent*. (Again, empty entries of “pseudo-type” `MPI_UB` and `MPI_LB` are not listed in the type map, but they affect the value of *extent*.) This receive operation receives  $n \cdot \text{count}$  entries, where entry  $i \cdot n + j$  is at location  $\text{buf} + \text{extent} \cdot i + disp_j$  and has type  $type_j$ . If the incoming message consists of  $k$  elements, then we must have  $k \leq n \cdot \text{count}$ ; the  $i \cdot n + j$ -th element of the message should have a type that matches  $type_j$ .

Type matching is defined according to the type signature of the corresponding datatypes, that is, the sequence of basic type components. Type matching does not depend on some aspects of the datatype definition, such as the displacements (layout in memory) or the intermediate types used.

**Example 4.11** This example shows that type matching is defined in terms of the basic types that a derived type consists of.

```

...
CALL MPI_TYPE_CONTIGUOUS( 2, MPI_REAL, type2, ...)
CALL MPI_TYPE_CONTIGUOUS( 4, MPI_REAL, type4, ...)
CALL MPI_TYPE_CONTIGUOUS( 2, type2, type22, ...)
...
CALL MPI_SEND( a, 4, MPI_REAL, ...)
CALL MPI_SEND( a, 2, type2, ...)
CALL MPI_SEND( a, 1, type22, ...)
CALL MPI_SEND( a, 1, type4, ...)
...
CALL MPI_RECV( a, 4, MPI_REAL, ...)
CALL MPI_RECV( a, 2, type2, ...)
CALL MPI_RECV( a, 1, type22, ...)
CALL MPI_RECV( a, 1, type4, ...)

```

Each of the sends matches any of the receives.

A datatype may specify overlapping entries. The use of such a datatype in a receive operation is erroneous. (This is erroneous even if the actual message received is short enough not to write any entry more than once.)

Suppose that `MPI_RECV(buf, count, datatype, dest, tag, comm, status)` is executed, where `datatype` has type map,

$$\{(type_0, disp_0), \dots, (type_{n-1}, disp_{n-1})\}.$$

The received message need not fill all the receive buffer, nor does it need to fill a number of locations which is a multiple of  $n$ . Any number,  $k$ , of basic elements can be received, where  $0 \leq k \leq \text{count} \cdot n$ . The number of basic elements received can be retrieved from `status` using the query `[function]functions MPI_GET_ELEMENTS[ or MPI_GET_ELEMENTS_X]`.

`MPI_GET_ELEMENTS( status, datatype, count)`

IN	status	return status of receive operation (Status)
IN	datatype	datatype used by receive operation (handle)
OUT	count	number of received basic elements (integer)

```
int MPI_Get_elements(MPI_Status *status, MPI_Datatype datatype, int *count)
```

```

MPI_GET_ELEMENTS(STATUS, DATATYPE, COUNT, IERROR)
    INTEGER STATUS(MPI_STATUS_SIZE), DATATYPE, COUNT, IERROR

```

```

{int MPI::Status::Get_elements(const MPI::Datatype& datatype) const(binding
    deprecated, see Section 15.2) }

```

```

1  MPI_GET_ELEMENTS_X( status, datatype, count)
2      IN          status          return status of receive operation (Status)
3
4      IN          datatype        datatype used by receive operation (handle)
5
6      OUT         count           number of received basic elements (integer)
7
8  int MPI_Get_elements_x(MPI_Status *status, MPI_Datatype datatype,
9                          MPI_Count *count)
10
11 MPI_GET_ELEMENTS_X(STATUS, DATATYPE, COUNT, IERROR)
12     INTEGER STATUS(MPI_STATUS_SIZE), DATATYPE, IERROR
13     INTEGER (KIND=MPI_COUNT_KIND) COUNT

```

The previously defined function `MPI_GET_COUNT` (Section 3.2.5), has a different behavior. It returns the number of “top-level entries” received, i.e. the number of “copies” of type `datatype`. In the previous example, `MPI_GET_COUNT` may return any integer value  $k$ , where  $0 \leq k \leq \text{count}$ . If `MPI_GET_COUNT` returns  $k$ , then the number of basic elements received (and the value returned by `MPI_GET_ELEMENTS` or `MPI_GET_ELEMENTS_X`) is  $n \cdot k$ . If the number of basic elements received is not a multiple of  $n$ , that is, if the receive operation has not received an integral number of `datatype` “copies,” then `MPI_GET_COUNT` [returns]sets the value of `count` to `MPI_UNDEFINED`. The `datatype` argument should match the argument provided by the receive call that set the status variable.

If the number of basic elements received exceeds the limits of the `count` parameter, then `MPI_GET_ELEMENTS` and `MPI_GET_ELEMENTS_X` set the value of `count` to `MPI_UNDEFINED`.

**Example 4.12** Usage of `MPI_GET_COUNT` and `MPI_GET_ELEMENTS`.

```

...
CALL MPI_TYPE_CONTIGUOUS(2, MPI_REAL, Type2, ierr)
CALL MPI_TYPE_COMMIT(Type2, ierr)
...
CALL MPI_COMM_RANK(comm, rank, ierr)
IF (rank.EQ.0) THEN
    CALL MPI_SEND(a, 2, MPI_REAL, 1, 0, comm, ierr)
    CALL MPI_SEND(a, 3, MPI_REAL, 1, 0, comm, ierr)
ELSE IF (rank.EQ.1) THEN
    CALL MPI_RECV(a, 2, Type2, 0, 0, comm, stat, ierr)
    CALL MPI_GET_COUNT(stat, Type2, i, ierr)      ! returns i=1
    CALL MPI_GET_ELEMENTS(stat, Type2, i, ierr)   ! returns i=2
    CALL MPI_RECV(a, 2, Type2, 0, 0, comm, stat, ierr)
    CALL MPI_GET_COUNT(stat, Type2, i, ierr)      ! returns i=MPI_UNDEFINED
    CALL MPI_GET_ELEMENTS(stat, Type2, i, ierr)   ! returns i=3
END IF

```

The [function]functions `MPI_GET_ELEMENTS` and `MPI_GET_ELEMENTS_X` can also be used after a probe to find the number of elements in the probed message. Note that the [two] functions `MPI_GET_COUNT` [and], `MPI_GET_ELEMENTS` and `MPI_GET_ELEMENTS_X` return the same values when they are used with basic datatypes.

*Rationale.* The extension given to the definition of `MPI_GET_COUNT` seems natural: one would expect this function to return the value of the `count` argument, when the receive buffer is filled. Sometimes `datatype` represents a basic unit of data one wants to transfer, for example, a record in an array of records (structures). One should be able to find out how many components were received without bothering to divide by the number of elements in each component. However, on other occasions, `datatype` is used to define a complex layout of data in the receiver memory, and does not represent a basic unit of data for transfers. In such cases, one needs to use the function `MPI_GET_ELEMENTS` or `MPI_GET_ELEMENTS_X`. (*End of rationale.*)

*Advice to implementors.* The definition implies that a receive cannot change the value of storage outside the entries defined to compose the communication buffer. In particular, the definition implies that padding space in a structure should not be modified when such a structure is copied from one process to another. This would prevent the obvious optimization of copying the structure, together with the padding, as one contiguous block. The implementation is free to do this optimization when it does not impact the outcome of the computation. The user can “force” this optimization by explicitly including padding as part of the message. (*End of advice to implementors.*)

#### 4.1.12 Correct Use of Addresses

Successively declared variables in C or Fortran are not necessarily stored at contiguous locations. Thus, care must be exercised that displacements do not cross from one variable to another. Also, in machines with a segmented address space, addresses are not unique and address arithmetic has some peculiar properties. Thus, the use of **addresses**, that is, displacements relative to the start address `MPI_BOTTOM`, has to be restricted.

Variables belong to the same **sequential storage** if they belong to the same array, to the same `COMMON` block in Fortran, or to the same structure in C. Valid addresses are defined recursively as follows:

1. The function `MPI_GET_ADDRESS` returns a valid address, when passed as argument a variable of the calling program.
2. The `buf` argument of a communication function evaluates to a valid address, when passed as argument a variable of the calling program.
3. If `v` is a valid address, and `i` is an integer, then `v+i` is a valid address, provided `v` and `v+i` are in the same sequential storage.
4. If `v` is a valid address then `MPI_BOTTOM + v` is a valid address.

A correct program uses only valid addresses to identify the locations of entries in communication buffers. Furthermore, if `u` and `v` are two valid addresses, then the (integer) difference `u - v` can be computed only if both `u` and `v` are in the same sequential storage. No other arithmetic operations can be meaningfully executed on addresses.

The rules above impose no constraints on the use of derived datatypes, as long as they are used to define a communication buffer that is wholly contained within the same sequential storage. However, the construction of a communication buffer that contains variables that are not within the same sequential storage must obey certain restrictions. Basically, a communication buffer with variables that are not within the same sequential

storage can be used only by specifying in the communication call `buf = MPI_BOTTOM`, `count = 1`, and using a `datatype` argument where all displacements are valid (absolute) addresses.

*Advice to users.* It is not expected that MPI implementations will be able to detect erroneous, “out of bound” displacements — unless those overflow the user address space — since the MPI call may not know the extent of the arrays and records in the host program. (*End of advice to users.*)

*Advice to implementors.* There is no need to distinguish (absolute) addresses and (relative) displacements on a machine with contiguous address space: `MPI_BOTTOM` is zero, and both addresses and displacements are integers. On machines where the distinction is required, addresses are recognized as expressions that involve `MPI_BOTTOM`. (*End of advice to implementors.*)

#### 4.1.13 Decoding a Datatype

MPI datatype objects allow users to specify an arbitrary layout of data in memory. There are several cases where accessing the layout information in opaque datatype objects would be useful. The opaque datatype object has found a number of uses outside MPI. Furthermore, a number of tools wish to display internal information about a datatype. To achieve this, datatype decoding functions are provided. The two functions in this section are used together to decode datatypes to recreate the calling sequence used in their initial definition. These can be used to allow a user to determine the type map and type signature of a datatype.

`MPI_TYPE_GET_ENVELOPE(datatype, num_integers, num_addresses, num_datatypes, combiner)`

IN	<code>datatype</code>	datatype to access (handle)
OUT	<code>num_integers</code>	number of input integers used in the call constructing combiner (non-negative integer)
OUT	<code>num_addresses</code>	number of input addresses used in the call constructing combiner (non-negative integer)
OUT	<code>num_datatypes</code>	number of input datatypes used in the call constructing combiner (non-negative integer)
OUT	<code>combiner</code>	combiner (state)

```
int MPI_Type_get_envelope(MPI_Datatype datatype, int *num_integers,
                          int *num_addresses, int *num_datatypes, int *combiner)
MPI_TYPE_GET_ENVELOPE(DATATYPE, NUM_INTEGERS, NUM_ADDRESSES, NUM_DATATYPES,
                      COMBINER, IERROR)
    INTEGER DATATYPE, NUM_INTEGERS, NUM_ADDRESSES, NUM_DATATYPES, COMBINER,
    IERROR
{void MPI::Datatype::Get_envelope(int& num_integers, int& num_addresses,
                                  int& num_datatypes, int& combiner) const(binding deprecated, see
Section 15.2) }
```

For the given datatype, `MPI_TYPE_GET_ENVELOPE` returns information on the number and type of input arguments used in the call that created the `datatype`. The number-of-arguments values returned can be used to provide sufficiently large arrays in the decoding routine `MPI_TYPE_GET_CONTENTS`. This call and the meaning of the returned values is described below. The combiner reflects the MPI datatype constructor call that was used in creating `datatype`.

*Rationale.* By requiring that the combiner reflect the constructor used in the creation of the `datatype`, the decoded information can be used to effectively recreate the calling sequence used in the original creation. One call is effectively the same as another when the information obtained from `MPI_TYPE_GET_CONTENTS` may be used with either to produce the same outcome. C calls `MPI_Type_hindexed` and `MPI_Type_create_hindexed` are always effectively the same while the Fortran call `MPI_TYPE_HINDEXED` will be different than either of these in some MPI implementations. This is the most useful information and was felt to be reasonable even though it constrains implementations to remember the original constructor sequence even if the internal representation is different.

The decoded information keeps track of datatype duplications. This is important as one needs to distinguish between a predefined datatype and a dup of a predefined datatype. The former is a constant object that cannot be freed, while the latter is a derived datatype that can be freed. (*End of rationale.*)

The list below has the values that can be returned in `combiner` on the left and the call associated with them on the right.

If `combiner` is `MPI_COMBINER_NAMED` then `datatype` is a named predefined datatype.

For deprecated calls with address arguments, we sometimes need to differentiate whether the call used an integer or an address size argument. For example, there are two combiners for `hvector`: `MPI_COMBINER_HVECTOR_INTEGER` and `MPI_COMBINER_HVECTOR`. The former is used if it was the MPI-1 call from Fortran, and the latter is used if it was the MPI-1 call from C or C++. However, on systems where `MPI_ADDRESS_KIND = MPI_INTEGER_KIND` (i.e., where integer arguments and address size arguments are the same), the combiner `MPI_COMBINER_HVECTOR` may be returned for a datatype constructed by a call to `MPI_TYPE_HVECTOR` from Fortran. Similarly, `MPI_COMBINER_HINDEXED` may be returned for a datatype constructed by a call to `MPI_TYPE_HINDEXED` from Fortran, and `MPI_COMBINER_STRUCT` may be returned for a datatype constructed by a call to `MPI_TYPE_STRUCT` from Fortran. On such systems, one need not differentiate constructors that take address size arguments from constructors that take integer arguments, since these are the same. The preferred calls all use address sized arguments so two combiners are not required for them.

*Rationale.* For recreating the original call, it is important to know if address information may have been truncated. The deprecated calls from Fortran for a few routines could be subject to truncation in the case where the default `INTEGER` size is smaller than the size of an address. (*End of rationale.*)

The actual arguments used in the creation call for a `datatype` can be obtained from the call:

MPI_COMBINER_NAMED	a named predefined datatype
MPI_COMBINER_DUP	MPI_TYPE_DUP
MPI_COMBINER_CONTIGUOUS	MPI_TYPE_CONTIGUOUS
MPI_COMBINER_VECTOR	MPI_TYPE_VECTOR
MPI_COMBINER_HVECTOR_INTEGER	MPI_TYPE_HVECTOR from Fortran
MPI_COMBINER_HVECTOR	MPI_TYPE_HVECTOR from C or C++ and in some case Fortran or MPI_TYPE_CREATE_HVECTOR
MPI_COMBINER_INDEXED	MPI_TYPE_INDEXED
MPI_COMBINER_HINDEXED_INTEGER	MPI_TYPE_HINDEXED from Fortran
MPI_COMBINER_HINDEXED	MPI_TYPE_HINDEXED from C or C++ and in some case Fortran or MPI_TYPE_CREATE_HINDEXED
MPI_COMBINER_INDEXED_BLOCK	MPI_TYPE_CREATE_INDEXED_BLOCK
MPI_COMBINER_STRUCT_INTEGER	MPI_TYPE_STRUCT from Fortran
MPI_COMBINER_STRUCT	MPI_TYPE_STRUCT from C or C++ and in some case Fortran or MPI_TYPE_CREATE_STRUCT
MPI_COMBINER_SUBARRAY	MPI_TYPE_CREATE_SUBARRAY
MPI_COMBINER_DARRAY	MPI_TYPE_CREATE_DARRAY
MPI_COMBINER_F90_REAL	MPI_TYPE_CREATE_F90_REAL
MPI_COMBINER_F90_COMPLEX	MPI_TYPE_CREATE_F90_COMPLEX
MPI_COMBINER_F90_INTEGER	MPI_TYPE_CREATE_F90_INTEGER
MPI_COMBINER_RESIZED	MPI_TYPE_CREATE_RESIZED

Table 4.1: combiner values returned from MPI\_TYPE\_GET\_ENVELOPE

MPI_TYPE_GET_CONTENTS(datatype, max_integers, max_addresses, max_datatypes, array_of_integers, array_of_addresses, array_of_datatypes)			
IN	datatype	datatype to access (handle)	
IN	max_integers	number of elements in array_of_integers (non-negative integer)	
IN	max_addresses	number of elements in array_of_addresses (non-negative integer)	
IN	max_datatypes	number of elements in array_of_datatypes (non-negative integer)	
OUT	array_of_integers	contains integer arguments used in constructing datatype (array of integers)	
OUT	array_of_addresses	contains address arguments used in constructing datatype (array of integers)	
OUT	array_of_datatypes	contains datatype arguments used in constructing datatype (array of handles)	

```

int MPI_Type_get_contents(MPI_Datatype datatype, int max_integers,
                          int max_addresses, int max_datatypes, int array_of_integers[],
                          MPI_Aint array_of_addresses[],
                          MPI_Datatype array_of_datatypes[])
MPI_TYPE_GET_CONTENTS(DATATYPE, MAX_INTEGERS, MAX_ADDRESSES, MAX_DATATYPES,
                      ARRAY_OF_INTEGERS, ARRAY_OF_ADDRESSES, ARRAY_OF_DATATYPES,
                      IERROR)
    INTEGER DATATYPE, MAX_INTEGERS, MAX_ADDRESSES, MAX_DATATYPES,
    ARRAY_OF_INTEGERS(*), ARRAY_OF_DATATYPES(*), IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) ARRAY_OF_ADDRESSES(*)
{void MPI::Datatype::Get_contents(int max_integers, int max_addresses,
    int max_datatypes, int array_of_integers[],
    MPI::Aint array_of_addresses[],
    MPI::Datatype array_of_datatypes[]) const(binding deprecated, see
    Section 15.2) }
```

`datatype` must be a predefined unnamed or a derived datatype; the call is erroneous if `datatype` is a predefined named datatype.

The values given for `max_integers`, `max_addresses`, and `max_datatypes` must be at least as large as the value returned in `num_integers`, `num_addresses`, and `num_datatypes`, respectively, in the call `MPI_TYPE_GET_ENVELOPE` for the same `datatype` argument.

*Rationale.* The arguments `max_integers`, `max_addresses`, and `max_datatypes` allow for error checking in the call. (*End of rationale.*)

The datatypes returned in `array_of_datatypes` are handles to datatype objects that are equivalent to the datatypes used in the original construction call. If these were derived datatypes, then the returned datatypes are new datatype objects, and the user is responsible for freeing these datatypes with `MPI_TYPE_FREE`. If these were predefined datatypes, then the returned datatype is equal to that (constant) predefined datatype and cannot be freed.

The committed state of returned derived datatypes is undefined, i.e., the datatypes may or may not be committed. Furthermore, the content of attributes of returned datatypes is undefined.

Note that `MPI_TYPE_GET_CONTENTS` can be invoked with a `datatype` argument that was constructed using `MPI_TYPE_CREATE_F90_REAL`, `MPI_TYPE_CREATE_F90_INTEGER`, or `MPI_TYPE_CREATE_F90_COMPLEX` (an unnamed predefined datatype). In such a case, an empty `array_of_datatypes` is returned.

*Rationale.* The definition of datatype equivalence implies that equivalent predefined datatypes are equal. By requiring the same handle for named predefined datatypes, it is possible to use the `==` or `.EQ.` comparison operator to determine the datatype involved. (*End of rationale.*)

*Advice to implementors.* The datatypes returned in `array_of_datatypes` must appear to the user as if each is an equivalent copy of the datatype used in the type constructor call. Whether this is done by creating a new datatype or via another mechanism such as a reference count mechanism is up to the implementation as long as the semantics are preserved. (*End of advice to implementors.*)



*Rationale.* The committed state and attributes of the returned datatype is deliberately left vague. The datatype used in the original construction may have been modified since its use in the constructor call. Attributes can be added, removed, or modified as well as having the datatype committed. The semantics given allow for a reference count implementation without having to track these changes. (*End of rationale.*)

In the deprecated datatype constructor calls, the address arguments in Fortran are of type `INTEGER`. In the preferred calls, the address arguments are of type `INTEGER(KIND=MPI_ADDRESS_KIND)`. The call `MPI_TYPE_GET_CONTENTS` returns all addresses in an argument of type `INTEGER(KIND=MPI_ADDRESS_KIND)`. This is true even if the deprecated calls were used. Thus, the location of values returned can be thought of as being returned by the C bindings. It can also be determined by examining the preferred calls for datatype constructors for the deprecated calls that involve addresses.

*Rationale.* By having all address arguments returned in the `array_of_addresses` argument, the result from a C and Fortran decoding of a `datatype` gives the result in the same argument. It is assumed that an integer of type `INTEGER(KIND=MPI_ADDRESS_KIND)` will be at least as large as the `INTEGER` argument used in datatype construction with the old MPI-1 calls so no loss of information will occur. (*End of rationale.*)

The following defines what values are placed in each entry of the returned arrays depending on the datatype constructor used for `datatype`. It also specifies the size of the arrays needed which is the values returned by `MPI_TYPE_GET_ENVELOPE`. In Fortran, the following calls were made:

```

PARAMETER (LARGE = 1000)
INTEGER TYPE, NI, NA, ND, COMBINER, I(LARGE), D(LARGE), IERROR
INTEGER(KIND=MPI_ADDRESS_KIND) A(LARGE)
! CONSTRUCT DATATYPE TYPE (NOT SHOWN)
CALL MPI_TYPE_GET_ENVELOPE(TYPE, NI, NA, ND, COMBINER, IERROR)
IF ((NI .GT. LARGE) .OR. (NA .GT. LARGE) .OR. (ND .GT. LARGE)) THEN
    WRITE (*, *) "NI, NA, OR ND = ", NI, NA, ND, &
    " RETURNED BY MPI_TYPE_GET_ENVELOPE IS LARGER THAN LARGE = ", LARGE
    CALL MPI_ABORT(MPI_COMM_WORLD, 99, IERROR)
ENDIF
CALL MPI_TYPE_GET_CONTENTS(TYPE, NI, NA, ND, I, A, D, IERROR)

```

or in C the analogous calls of:

```

#define LARGE 1000
int ni, na, nd, combiner, i[LARGE];
MPI_Aint a[LARGE];
MPI_Datatype type, d[LARGE];
/* construct datatype type (not shown) */
MPI_Type_get_envelope(type, &ni, &na, &nd, &combiner);
if ((ni > LARGE) || (na > LARGE) || (nd > LARGE)) {
    fprintf(stderr, "ni, na, or nd = %d %d %d returned by ", ni, na, nd);

```

```

    fprintf(stderr, "MPI_Type_get_envelope is larger than LARGE = %d\n",
            LARGE);
    MPI_Abort(MPI_COMM_WORLD, 99);
};
MPI_Type_get_contents(type, ni, na, nd, i, a, d);

```

The C++ code is in analogy to the C code above with the same values returned.

In the descriptions that follow, the lower case name of arguments is used.

If combiner is MPI\_COMBINER\_NAMED then it is erroneous to call

MPI\_TYPE\_GET\_CONTENTS.

If combiner is MPI\_COMBINER\_DUP then

Constructor argument	C & C++ location	Fortran location
oldtype	d[0]	D(1)

and ni = 0, na = 0, nd = 1.

If combiner is MPI\_COMBINER\_CONTIGUOUS then

Constructor argument	C & C++ location	Fortran location
count	i[0]	I(1)
oldtype	d[0]	D(1)

and ni = 1, na = 0, nd = 1.

If combiner is MPI\_COMBINER\_VECTOR then

Constructor argument	C & C++ location	Fortran location
count	i[0]	I(1)
blocklength	i[1]	I(2)
stride	i[2]	I(3)
oldtype	d[0]	D(1)

and ni = 3, na = 0, nd = 1.

If combiner is MPI\_COMBINER\_HVECTOR\_INTEGER or MPI\_COMBINER\_HVECTOR then

Constructor argument	C & C++ location	Fortran location
count	i[0]	I(1)
blocklength	i[1]	I(2)
stride	a[0]	A(1)
oldtype	d[0]	D(1)

and ni = 2, na = 1, nd = 1.

If combiner is MPI\_COMBINER\_INDEXED then

Constructor argument	C & C++ location	Fortran location
count	i[0]	I(1)
array_of_blocklengths	i[1] to i[i[0]]	I(2) to I(I(1)+1)
array_of_displacements	i[i[0]+1] to i[2*i[0]]	I(I(1)+2) to I(2*I(1)+1)
oldtype	d[0]	D(1)

and ni = 2\*count+1, na = 0, nd = 1.

If combiner is MPI\_COMBINER\_HINDEXED\_INTEGER or MPI\_COMBINER\_HINDEXED then

Constructor argument	C & C++ location	Fortran location
count	i[0]	I(1)
array_of_blocklengths	i[1] to i[i[0]]	I(2) to I(I(1)+1)
array_of_displacements	a[0] to a[i[0]-1]	A(1) to A(I(1))
oldtype	d[0]	D(1)

and ni = count+1, na = count, nd = 1.

If combiner is MPI\_COMBINER\_INDEXED\_BLOCK then

Constructor argument	C & C++ location	Fortran location
count	i[0]	I(1)
blocklength	i[1]	I(2)
array_of_displacements	i[2] to i[i[0]+1]	I(3) to I(I(1)+2)
oldtype	d[0]	D(1)

and ni = count+2, na = 0, nd = 1.

If combiner is MPI\_COMBINER\_STRUCT\_INTEGER or MPI\_COMBINER\_STRUCT then

Constructor argument	C & C++ location	Fortran location
count	i[0]	I(1)
array_of_blocklengths	i[1] to i[i[0]]	I(2) to I(I(1)+1)
array_of_displacements	a[0] to a[i[0]-1]	A(1) to A(I(1))
array_of_types	d[0] to d[i[0]-1]	D(1) to D(I(1))

and ni = count+1, na = count, nd = count.

If combiner is MPI\_COMBINER\_SUBARRAY then

Constructor argument	C & C++ location	Fortran location
ndims	i[0]	I(1)
array_of_sizes	i[1] to i[i[0]]	I(2) to I(I(1)+1)
array_of_subsizes	i[i[0]+1] to i[2*i[0]]	I(I(1)+2) to I(2*I(1)+1)
array_of_starts	i[2*i[0]+1] to i[3*i[0]]	I(2*I(1)+2) to I(3*I(1)+1)
order	i[3*i[0]+1]	I(3*I(1)+2)
oldtype	d[0]	D(1)

and ni = 3\*ndims+2, na = 0, nd = 1.

If combiner is MPI\_COMBINER\_DARRAY then

Constructor argument	C & C++ location	Fortran location
size	i[0]	I(1)
rank	i[1]	I(2)
ndims	i[2]	I(3)
array_of_gsizes	i[3] to i[i[2]+2]	I(4) to I(I(3)+3)
array_of_distribs	i[i[2]+3] to i[2*i[2]+2]	I(I(3)+4) to I(2*I(3)+3)
array_of_dargs	i[2*i[2]+3] to i[3*i[2]+2]	I(2*I(3)+4) to I(3*I(3)+3)
array_of_psize	i[3*i[2]+3] to i[4*i[2]+2]	I(3*I(3)+4) to I(4*I(3)+3)
order	i[4*i[2]+3]	I(4*I(3)+4)
oldtype	d[0]	D(1)

and ni = 4\*ndims+4, na = 0, nd = 1.

If combiner is MPI\_COMBINER\_F90\_REAL then

Constructor argument	C & C++ location	Fortran location
p	i[0]	I(1)
r	i[1]	I(2)

and ni = 2, na = 0, nd = 0.

If combiner is MPI\_COMBINER\_F90\_COMPLEX then

Constructor argument	C & C++ location	Fortran location
p	i[0]	I(1)
r	i[1]	I(2)

and ni = 2, na = 0, nd = 0.

If combiner is MPI\_COMBINER\_F90\_INTEGER then

Constructor argument	C & C++ location	Fortran location
r	i[0]	I(1)

and ni = 1, na = 0, nd = 0.

If combiner is MPI\_COMBINER\_RESIZED then

Constructor argument	C & C++ location	Fortran location
lb	a[0]	A(1)
extent	a[1]	A(2)
oldtype	d[0]	D(1)

and ni = 0, na = 2, nd = 1.

#### 4.1.14 Examples

The following examples illustrate the use of derived datatypes.

**Example 4.13** Send and receive a section of a 3D array.

```

REAL a(100,100,100), e(9,9,9)
INTEGER oneslice, twoslice, threeslice, sizeofreal, myrank, ierr
INTEGER status(MPI_STATUS_SIZE)

C      extract the section a(1:17:2, 3:11, 2:10)
C      and store it in e(:, :, :).

CALL MPI_COMM_RANK(MPI_COMM_WORLD, myrank, ierr)

CALL MPI_TYPE_EXTENT( MPI_REAL, sizeofreal, ierr)

C      create datatype for a 1D section
CALL MPI_TYPE_VECTOR( 9, 1, 2, MPI_REAL, oneslice, ierr)

C      create datatype for a 2D section
CALL MPI_TYPE_HVECTOR(9, 1, 100*sizeofreal, oneslice, twoslice, ierr)

C      create datatype for the entire section
CALL MPI_TYPE_HVECTOR( 9, 1, 100*100*sizeofreal, twoslice,
                      threeslice, ierr)

```

```

1
2     CALL MPI_TYPE_COMMIT( threeslice, ierr)
3     CALL MPI_SENDRECV(a(1,3,2), 1, threeslice, myrank, 0, e, 9*9*9,
4                     MPI_REAL, myrank, 0, MPI_COMM_WORLD, status, ierr)
5
6

```

**Example 4.14** Copy the (strictly) lower triangular part of a matrix.

```

8     REAL a(100,100), b(100,100)
9     INTEGER disp(100), blocklen(100), ltype, myrank, ierr
10    INTEGER status(MPI_STATUS_SIZE)
11
12    C    copy lower triangular part of array a
13    C    onto lower triangular part of array b
14
15    CALL MPI_COMM_RANK(MPI_COMM_WORLD, myrank, ierr)
16
17    C    compute start and size of each column
18    DO i=1, 100
19        disp(i) = 100*(i-1) + i
20        blocklen(i) = 100-i
21    END DO
22
23    C    create datatype for lower triangular part
24    CALL MPI_TYPE_INDEXED( 100, blocklen, disp, MPI_REAL, ltype, ierr)
25
26    CALL MPI_TYPE_COMMIT(ltype, ierr)
27    CALL MPI_SENDRECV( a, 1, ltype, myrank, 0, b, 1,
28                    ltype, myrank, 0, MPI_COMM_WORLD, status, ierr)
29
30

```

**Example 4.15** Transpose a matrix.

```

32    REAL a(100,100), b(100,100)
33    INTEGER row, xpose, sizeofreal, myrank, ierr
34    INTEGER status(MPI_STATUS_SIZE)
35
36    C    transpose matrix a onto b
37
38    CALL MPI_COMM_RANK(MPI_COMM_WORLD, myrank, ierr)
39
40    CALL MPI_TYPE_EXTENT( MPI_REAL, sizeofreal, ierr)
41
42    C    create datatype for one row
43    CALL MPI_TYPE_VECTOR( 100, 1, 100, MPI_REAL, row, ierr)
44
45    C    create datatype for matrix in row-major order
46    CALL MPI_TYPE_HVECTOR( 100, 1, sizeofreal, row, xpose, ierr)
47
48

```

```

CALL MPI_TYPE_COMMIT( xpose, ierr)
C      send matrix in row-major order and receive in column major order
CALL MPI_SENDRECV( a, 1, xpose, myrank, 0, b, 100*100,
                  MPI_REAL, myrank, 0, MPI_COMM_WORLD, status, ierr)

```

**Example 4.16** Another approach to the transpose problem:

```

REAL a(100,100), b(100,100)
INTEGER disp(2), blocklen(2), type(2), row, row1, sizeofreal
INTEGER myrank, ierr
INTEGER status(MPI_STATUS_SIZE)

CALL MPI_COMM_RANK(MPI_COMM_WORLD, myrank, ierr)

C      transpose matrix a onto b

CALL MPI_TYPE_EXTENT( MPI_REAL, sizeofreal, ierr)

C      create datatype for one row
CALL MPI_TYPE_VECTOR( 100, 1, 100, MPI_REAL, row, ierr)

C      create datatype for one row, with the extent of one real number
disp(1) = 0
disp(2) = sizeofreal
type(1) = row
type(2) = MPI_UB
blocklen(1) = 1
blocklen(2) = 1
CALL MPI_TYPE_STRUCT( 2, blocklen, disp, type, row1, ierr)

CALL MPI_TYPE_COMMIT( row1, ierr)

C      send 100 rows and receive in column major order
CALL MPI_SENDRECV( a, 100, row1, myrank, 0, b, 100*100,
                  MPI_REAL, myrank, 0, MPI_COMM_WORLD, status, ierr)

```

**Example 4.17** We manipulate an array of structures.

```

struct Partstruct
{
    int    class; /* particle class */
    double d[6];  /* particle coordinates */
    char   b[7];  /* some additional information */
};

struct Partstruct    particle[1000];

```

```

1
2  int          i, dest, rank, tag;
3  MPI_Comm     comm;
4
5
6  /* build datatype describing structure */
7
8  MPI_Datatype Particletype;
9  MPI_Datatype type[3] = {MPI_INT, MPI_DOUBLE, MPI_CHAR};
10 int          blocklen[3] = {1, 6, 7};
11 MPI_Aint      disp[3];
12 MPI_Aint      base;
13
14
15 /* compute displacements of structure components */
16
17 MPI_Address( particle, disp);
18 MPI_Address( particle[0].d, disp+1);
19 MPI_Address( particle[0].b, disp+2);
20 base = disp[0];
21 for (i=0; i < 3; i++) disp[i] -= base;
22
23 MPI_Type_struct( 3, blocklen, disp, type, &Particletype);
24
25 /* If compiler does padding in mysterious ways,
26    the following may be safer */
27
28 MPI_Datatype type1[4] = {MPI_INT, MPI_DOUBLE, MPI_CHAR, MPI_UB};
29 int          blocklen1[4] = {1, 6, 7, 1};
30 MPI_Aint      disp1[4];
31
32 /* compute displacements of structure components */
33
34 MPI_Address( particle, disp1);
35 MPI_Address( particle[0].d, disp1+1);
36 MPI_Address( particle[0].b, disp1+2);
37 MPI_Address( particle+1, disp1+3);
38 base = disp1[0];
39 for (i=0; i < 4; i++) disp1[i] -= base;
40
41 /* build datatype describing structure */
42
43 MPI_Type_struct( 4, blocklen1, disp1, type1, &Particletype);
44
45
46 /* 4.1:
47    send the entire array */
48

```

```

MPI_Type_commit( &Particletype);
MPI_Send( particle, 1000, Particletype, dest, tag, comm);

/* 4.2:
send only the entries of class zero particles,
preceded by the number of such entries */

MPI_Datatype Zparticles; /* datatype describing all particles
                           with class zero (needs to be recomputed
                           if classes change) */

MPI_Datatype Ztype;

MPI_Aint      zdisp[1000];
int           zblock[1000], j, k;
int           zzblock[2] = {1,1};
MPI_Aint      zzdisp[2];
MPI_Datatype  zztype[2];

/* compute displacements of class zero particles */
j = 0;
for(i=0; i < 1000; i++)
    if (particle[i].class == 0)
    {
        zdisp[j] = i;
        zblock[j] = 1;
        j++;
    }

/* create datatype for class zero particles */
MPI_Type_indexed( j, zblock, zdisp, Particletype, &Zparticles);

/* prepend particle count */
MPI_Address(&j, zzdisp);
MPI_Address(particle, zzdisp+1);
zztype[0] = MPI_INT;
zztype[1] = Zparticles;
MPI_Type_struct(2, zzblock, zzdisp, zztype, &Ztype);

MPI_Type_commit( &Ztype);
MPI_Send( MPI_BOTTOM, 1, Ztype, dest, tag, comm);

/* A probably more efficient way of defining Zparticles */

/* consecutive particles with index zero are handled as one block */
j=0;
for (i=0; i < 1000; i++)

```



```

1      if (particle[i].index == 0)
2          {
3              for (k=i+1; (k < 1000)&&(particle[k].index == 0) ; k++);
4              zdisp[j] = i;
5              zblock[j] = k-i;
6              j++;
7              i = k;
8          }
9      MPI_Type_indexed( j, zblock, zdisp, Particletype, &Zparticles);
10
11
12          /* 4.3:
13             send the first two coordinates of all entries */
14
15      MPI_Datatype Allpairs;      /* datatype for all pairs of coordinates */
16
17      MPI_Aint sizeofentry;
18
19      MPI_Type_extent( Particletype, &sizeofentry);
20
21          /* sizeofentry can also be computed by subtracting the address
22             of particle[0] from the address of particle[1] */
23
24      MPI_Type_hvector( 1000, 2, sizeofentry, MPI_DOUBLE, &Allpairs);
25      MPI_Type_commit( &Allpairs);
26      MPI_Send( particle[0].d, 1, Allpairs, dest, tag, comm);
27
28          /* an alternative solution to 4.3 */
29
30      MPI_Datatype Onepair;      /* datatype for one pair of coordinates, with
31                                 the extent of one particle entry */
32      MPI_Aint disp2[3];
33      MPI_Datatype type2[3] = {MPI_LB, MPI_DOUBLE, MPI_UB};
34      int blocklen2[3] = {1, 2, 1};
35
36      MPI_Address( particle, disp2);
37      MPI_Address( particle[0].d, disp2+1);
38      MPI_Address( particle+1, disp2+2);
39      base = disp2[0];
40      for (i=0; i<2; i++) disp2[i] -= base;
41
42      MPI_Type_struct( 3, blocklen2, disp2, type2, &Onepair);
43      MPI_Type_commit( &Onepair);
44      MPI_Send( particle[0].d, 1000, Onepair, dest, tag, comm);
45
46
47

```

**Example 4.18** The same manipulations as in the previous example, but use absolute

addresses in datatypes.

```

struct Partstruct
{
    int class;
    double d[6];
    char b[7];
};

struct Partstruct particle[1000];

/* build datatype describing first array entry */

MPI_Datatype Particletype;
MPI_Datatype type[3] = {MPI_INT, MPI_DOUBLE, MPI_CHAR};
int          block[3] = {1, 6, 7};
MPI_Aint     disp[3];

MPI_Address( particle, disp);
MPI_Address( particle[0].d, disp+1);
MPI_Address( particle[0].b, disp+2);
MPI_Type_struct( 3, block, disp, type, &Particletype);

/* Particletype describes first array entry -- using absolute
   addresses */

/* 5.1:
   send the entire array */

MPI_Type_commit( &Particletype);
MPI_Send( MPI_BOTTOM, 1000, Particletype, dest, tag, comm);

/* 5.2:
   send the entries of class zero,
   preceded by the number of such entries */

MPI_Datatype Zparticles, Ztype;

MPI_Aint     zdisp[1000];
int          zblock[1000], i, j, k;
int          zzblock[2] = {1,1};
MPI_Datatype zztype[2];
MPI_Aint     zzdisp[2];

j=0;
for (i=0; i < 1000; i++)
    if (particle[i].index == 0)

```

```

1      {
2          for (k=i+1; (k < 1000)&&(particle[k].index == 0) ; k++);
3          zdisp[j] = i;
4          zblock[j] = k-i;
5          j++;
6          i = k;
7      }
8      MPI_Type_indexed( j, zblock, zdisp, Particletype, &Zparticles);
9      /* Zparticles describe particles with class zero, using
10         their absolute addresses*/
11
12     /* prepend particle count */
13     MPI_Address(&j, zzdisp);
14     zzdisp[1] = MPI_BOTTOM;
15     zztype[0] = MPI_INT;
16     zztype[1] = Zparticles;
17     MPI_Type_struct(2, zzbblock, zzdisp, zztype, &Ztype);
18
19     MPI_Type_commit( &Ztype);
20     MPI_Send( MPI_BOTTOM, 1, Ztype, dest, tag, comm);
21
22
23

```

#### Example 4.19 Handling of unions.

```

24
25     union {
26         int      ival;
27         float    fval;
28         } u[1000];
29
30     int      utype;
31
32     /* All entries of u have identical type; variable
33        utype keeps track of their current type */
34
35     MPI_Datatype  type[2];
36     int          blocklen[2] = {1,1};
37     MPI_Aint      disp[2];
38     MPI_Datatype  mpi_utype[2];
39     MPI_Aint      i,j;
40
41     /* compute an MPI datatype for each possible union type;
42        assume values are left-aligned in union storage. */
43
44     MPI_Address( u, &i);
45     MPI_Address( u+1, &j);
46     disp[0] = 0; disp[1] = j-i;
47     type[1] = MPI_UB;
48

```

```

type[0] = MPI_INT;
MPI_Type_struct(2, blocklen, disp, type, &mpi_utype[0]);

type[0] = MPI_FLOAT;
MPI_Type_struct(2, blocklen, disp, type, &mpi_utype[1]);

for(i=0; i<2; i++) MPI_Type_commit(&mpi_utype[i]);

/* actual communication */

MPI_Send(u, 1000, mpi_utype[utype], dest, tag, comm);

```

**Example 4.20** This example shows how a datatype can be decoded. The routine `printdatatype` prints out the elements of the datatype. Note the use of `MPI_Type_free` for datatypes that are not predefined.

```

/*
   Example of decoding a datatype.

   Returns 0 if the datatype is predefined, 1 otherwise
*/
#include <stdio.h>
#include <stdlib.h>
#include "mpi.h"
int printdatatype( MPI_Datatype datatype )
{
    int *array_of_ints;
    MPI_Aint *array_of_adds;
    MPI_Datatype *array_of_dtypes;
    int num_ints, num_adds, num_dtypes, combiner;
    int i;

    MPI_Type_get_envelope( datatype,
                          &num_ints, &num_adds, &num_dtypes, &combiner );
    switch (combiner) {
    case MPI_COMBINER_NAMED:
        printf( "Datatype is named:" );
        /* To print the specific type, we can match against the
           predefined forms. We can NOT use a switch statement here
           We could also use MPI_TYPE_GET_NAME if we preferred to use
           names that the user may have changed.
        */
        if (datatype == MPI_INT) printf( "MPI_INT\n" );
        else if (datatype == MPI_DOUBLE) printf( "MPI_DOUBLE\n" );
        ... else test for other types ...
        return 0;
        break;

```

```

1  case MPI_COMBINER_STRUCT:
2  case MPI_COMBINER_STRUCT_INTEGER:
3      printf( "Datatype is struct containing" );
4      array_of_ints = (int *)malloc( num_ints * sizeof(int) );
5      array_of_adds =
6          (MPI_Aint *) malloc( num_adds * sizeof(MPI_Aint) );
7      array_of_dtypes = (MPI_Datatype *)
8          malloc( num_dtypes * sizeof(MPI_Datatype) );
9      MPI_Type_get_contents( datatype, num_ints, num_adds, num_dtypes,
10                          array_of_ints, array_of_adds, array_of_dtypes );
11     printf( " %d datatypes:\n", array_of_ints[0] );
12     for (i=0; i<array_of_ints[0]; i++) {
13         printf( "blocklength %d, displacement %ld, type:\n",
14             array_of_ints[i+1], array_of_adds[i] );
15         if (printdatatype( array_of_dtypes[i] )) {
16             /* Note that we free the type ONLY if it
17              is not predefined */
18             MPI_Type_free( &array_of_dtypes[i] );
19         }
20     }
21     free( array_of_ints );
22     free( array_of_adds );
23     free( array_of_dtypes );
24     break;
25     ... other combiner values ...
26 default:
27     printf( "Unrecognized combiner type\n" );
28 }
29 return 1;
30 }

```

## 4.2 Pack and Unpack

Some existing communication libraries provide pack/unpack functions for sending noncontiguous data. In these, the user explicitly packs data into a contiguous buffer before sending it, and unpacks it from a contiguous buffer after receiving it. Derived datatypes, which are described in Section 4.1, allow one, in most cases, to avoid explicit packing and unpacking. The user specifies the layout of the data to be sent or received, and the communication library directly accesses a noncontiguous buffer. The pack/unpack routines are provided for compatibility with previous libraries. Also, they provide some functionality that is not otherwise available in MPI. For instance, a message can be received in several parts, where the receive operation done on a later part may depend on the content of a former part. Another use is that outgoing messages may be explicitly buffered in user supplied space, thus overriding the system buffering policy. Finally, the availability of pack and unpack operations facilitates the development of additional communication libraries layered on top of MPI.

```

MPI_PACK(inbuf, incount, datatype, outbuf, outsize, position, comm)
    IN      inbuf      input buffer start (choice)
    IN      incount    number of input data items (non-negative integer)
    IN      datatype   datatype of each input data item (handle)
    OUT     outbuf     output buffer start (choice)
    IN      outsize    output buffer size, in bytes (non-negative integer)
    INOUT   position   current position in buffer, in bytes (integer)
    IN      comm       communicator for packed message (handle)

int MPI_Pack(void* inbuf, int incount, MPI_Datatype datatype, void *outbuf,
             int outsize, int *position, MPI_Comm comm)

MPI_PACK(INBUF, INCOUNT, DATATYPE, OUTBUF, OUTSIZE, POSITION, COMM, IERROR)
    <type> INBUF(*), OUTBUF(*)
    INTEGER INCOUNT, DATATYPE, OUTSIZE, POSITION, COMM, IERROR

{void MPI::Datatype::Pack(const void* inbuf, int incount, void *outbuf,
                          int outsize, int& position, const MPI::Comm &comm)
                          const(binding deprecated, see Section 15.2) }
```

Packs the message in the send buffer specified by `inbuf`, `incount`, `datatype` into the buffer space specified by `outbuf` and `outsize`. The input buffer can be any communication buffer allowed in `MPI_SEND`. The output buffer is a contiguous storage area containing `outsize` bytes, starting at the address `outbuf` (length is counted in bytes, not elements, as if it were a communication buffer for a message of type `MPI_PACKED`).

The input value of `position` is the first location in the output buffer to be used for packing. `position` is incremented by the size of the packed message, and the output value of `position` is the first location in the output buffer following the locations occupied by the packed message. The `comm` argument is the communicator that will be subsequently used for sending the packed message.

```

MPI_UNPACK(inbuf, insize, position, outbuf, outcount, datatype, comm)
    IN      inbuf      input buffer start (choice)
    IN      insize     size of input buffer, in bytes (non-negative integer)
    INOUT   position   current position in bytes (integer)
    OUT     outbuf     output buffer start (choice)
    IN      outcount   number of items to be unpacked (integer)
    IN      datatype   datatype of each output data item (handle)
    IN      comm       communicator for packed message (handle)

int MPI_Unpack(void* inbuf, int insize, int *position, void *outbuf,
               int outcount, MPI_Datatype datatype, MPI_Comm comm)
```

```

1 MPI_UNPACK(INBUF, INSIZE, POSITION, OUTBUF, OUTCOUNT, DATATYPE, COMM,
2           IERROR)
3     <type> INBUF(*), OUTBUF(*)
4     INTEGER INSIZE, POSITION, OUTCOUNT, DATATYPE, COMM, IERROR
5
6 {void MPI::Datatype::Unpack(const void* inbuf, int insize, void *outbuf,
7     int outcount, int& position, const MPI::Comm& comm)
8     const(binding deprecated, see Section 15.2) }

```

Unpacks a message into the receive buffer specified by `outbuf`, `outcount`, `datatype` from the buffer space specified by `inbuf` and `insize`. The output buffer can be any communication buffer allowed in `MPI_RECV`. The input buffer is a contiguous storage area containing `insize` bytes, starting at address `inbuf`. The input value of `position` is the first location in the input buffer occupied by the packed message. `position` is incremented by the size of the packed message, so that the output value of `position` is the first location in the input buffer after the locations occupied by the message that was unpacked. `comm` is the communicator used to receive the packed message.

*Advice to users.* Note the difference between `MPI_RECV` and `MPI_UNPACK`: in `MPI_RECV`, the `count` argument specifies the maximum number of items that can be received. The actual number of items received is determined by the length of the incoming message. In `MPI_UNPACK`, the `count` argument specifies the actual number of items that are unpacked; the “size” of the corresponding message is the increment in `position`. The reason for this change is that the “incoming message size” is not predetermined since the user decides how much to unpack; nor is it easy to determine the “message size” from the number of items to be unpacked. In fact, in a heterogeneous system, this number may not be determined *a priori*. (*End of advice to users.*)

To understand the behavior of pack and unpack, it is convenient to think of the data part of a message as being the sequence obtained by concatenating the successive values sent in that message. The pack operation stores this sequence in the buffer space, as if sending the message to that buffer. The unpack operation retrieves this sequence from buffer space, as if receiving a message from that buffer. (It is helpful to think of internal Fortran files or `sscanf` in C, for a similar function.)

Several messages can be successively packed into one **packing unit**. This is effected by several successive **related** calls to `MPI_PACK`, where the first call provides `position = 0`, and each successive call inputs the value of `position` that was output by the previous call, and the same values for `outbuf`, `outcount` and `comm`. This packing unit now contains the equivalent information that would have been stored in a message by one send call with a send buffer that is the “concatenation” of the individual send buffers.

A packing unit can be sent using type `MPI_PACKED`. Any point to point or collective communication function can be used to move the sequence of bytes that forms the packing unit from one process to another. This packing unit can now be received using any receive operation, with any datatype: the type matching rules are relaxed for messages sent with type `MPI_PACKED`.

A message sent with any type (including `MPI_PACKED`) can be received using the type `MPI_PACKED`. Such a message can then be unpacked by calls to `MPI_UNPACK`.

A packing unit (or a message created by a regular, “typed” send) can be unpacked into several successive messages. This is effected by several successive related calls to

MPI\_UNPACK, where the first call provides `position = 0`, and each successive call inputs the value of `position` that was output by the previous call, and the same values for `inbuf`, `insize` and `comm`.

The concatenation of two packing units is not necessarily a packing unit; nor is a substring of a packing unit necessarily a packing unit. Thus, one cannot concatenate two packing units and then unpack the result as one packing unit; nor can one unpack a substring of a packing unit as a separate packing unit. Each packing unit, that was created by a related sequence of pack calls, or by a regular send, must be unpacked as a unit, by a sequence of related unpack calls.

*Rationale.* The restriction on “atomic” packing and unpacking of packing units allows the implementation to add at the head of packing units additional information, such as a description of the sender architecture (to be used for type conversion, in a heterogeneous environment) (*End of rationale.*)

The following call allows the user to find out how much space is needed to pack a message and, thus, manage space allocation for buffers.

MPI\_PACK\_SIZE(`incount`, `datatype`, `comm`, `size`)

IN	<code>incount</code>	count argument to packing call (non-negative integer)
IN	<code>datatype</code>	datatype argument to packing call (handle)
IN	<code>comm</code>	communicator argument to packing call (handle)
OUT	<code>size</code>	upper bound on size of packed message, in bytes (non-negative integer)

```
int MPI_Pack_size(int incount, MPI_Datatype datatype, MPI_Comm comm,
                  int *size)
```

```
MPI_PACK_SIZE(INCOUNT, DATATYPE, COMM, SIZE, IERROR)
    INTEGER INCOUNT, DATATYPE, COMM, SIZE, IERROR
```

```
{int MPI::Datatype::Pack_size(int incount, const MPI::Comm& comm)
    const(binding deprecated, see Section 15.2) }
```

A call to MPI\_PACK\_SIZE(`incount`, `datatype`, `comm`, `size`) returns in `size` an upper bound on the increment in `position` that is effected by a call to MPI\_PACK(`inbuf`, `incount`, `datatype`, `outbuf`, `outcount`, `position`, `comm`). If the packed size of the datatype can not be expressed by the `size` parameter, then MPI\_PACK\_SIZE sets the value of `size` to MPI\_UNDEFINED.

*Rationale.* The call returns an upper bound, rather than an exact bound, since the exact amount of space needed to pack the message may depend on the context (e.g., first message packed in a packing unit may take more space). (*End of rationale.*)

**Example 4.21** An example using MPI\_PACK.



```

1  int      position, i, j, a[2];
2  char      buff[1000];
3
4  MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
5  if (myrank == 0)
6  {
7      /* SENDER CODE */
8
9      position = 0;
10     MPI_Pack(&i, 1, MPI_INT, buff, 1000, &position, MPI_COMM_WORLD);
11     MPI_Pack(&j, 1, MPI_INT, buff, 1000, &position, MPI_COMM_WORLD);
12     MPI_Send( buff, position, MPI_PACKED, 1, 0, MPI_COMM_WORLD);
13 }
14 else /* RECEIVER CODE */
15     MPI_Recv( a, 2, MPI_INT, 0, 0, MPI_COMM_WORLD);
16
17

```

**Example 4.22** An elaborate example.

```

19  int  position, i;
20  float a[1000];
21  char  buff[1000];
22
23  MPI_Comm_rank(MPI_Comm_world, &myrank);
24  if (myrank == 0)
25  {
26      /* SENDER CODE */
27
28      int len[2];
29      MPI_Aint disp[2];
30      MPI_Datatype type[2], newtype;
31
32      /* build datatype for i followed by a[0]...a[i-1] */
33
34      len[0] = 1;
35      len[1] = i;
36      MPI_Address( &i, disp);
37      MPI_Address( a, disp+1);
38      type[0] = MPI_INT;
39      type[1] = MPI_FLOAT;
40      MPI_Type_struct( 2, len, disp, type, &newtype);
41      MPI_Type_commit( &newtype);
42
43      /* Pack i followed by a[0]...a[i-1]*/
44
45      position = 0;
46      MPI_Pack( MPI_BOTTOM, 1, newtype, buff, 1000, &position, MPI_COMM_WORLD);
47
48      /* Send */

```

```

MPI_Send( buff, position, MPI_PACKED, 1, 0,
          MPI_COMM_WORLD);

/* *****
   One can replace the last three lines with
   MPI_Send( MPI_BOTTOM, 1, newtype, 1, 0, MPI_COMM_WORLD);
   ***** */
}
else if (myrank == 1)
{
    /* RECEIVER CODE */

    MPI_Status status;

    /* Receive */

    MPI_Recv( buff, 1000, MPI_PACKED, 0, 0, MPI_COMM_WORLD, &status);

    /* Unpack i */

    position = 0;
    MPI_Unpack(buff, 1000, &position, &i, 1, MPI_INT, MPI_COMM_WORLD);

    /* Unpack a[0]...a[i-1] */
    MPI_Unpack(buff, 1000, &position, a, i, MPI_FLOAT, MPI_COMM_WORLD);
}

```

**Example 4.23** Each process sends a count, followed by count characters to the root; the root concatenates all characters into one string.

```

int  count, gsize, counts[64], totalcount, k1, k2, k,
     displs[64], position, concat_pos;
char chr[100], *lbuf, *rbuf, *cbuf;

MPI_Comm_size(comm, &gsize);
MPI_Comm_rank(comm, &myrank);

    /* allocate local pack buffer */
MPI_Pack_size(1, MPI_INT, comm, &k1);
MPI_Pack_size(count, MPI_CHAR, comm, &k2);
k = k1+k2;
lbuf = (char *)malloc(k);

    /* pack count, followed by count characters */
position = 0;
MPI_Pack(&count, 1, MPI_INT, lbuf, k, &position, comm);
MPI_Pack(chr, count, MPI_CHAR, lbuf, k, &position, comm);

```

```

1
2  if (myrank != root) {
3      /* gather at root sizes of all packed messages */
4      MPI_Gather( &position, 1, MPI_INT, NULL, 0,
5                  MPI_DATATYPE_NULL, root, comm);
6
7      /* gather at root packed messages */
8      MPI_Gatherv( lbuf, position, MPI_PACKED, NULL,
9                  NULL, NULL, NULL, root, comm);
10
11 } else { /* root code */
12     /* gather sizes of all packed messages */
13     MPI_Gather( &position, 1, MPI_INT, counts, 1,
14                 MPI_INT, root, comm);
15
16     /* gather all packed messages */
17     displs[0] = 0;
18     for (i=1; i < gsize; i++)
19         displs[i] = displs[i-1] + counts[i-1];
20     totalcount = displs[gsiz-1] + counts[gsiz-1];
21     rbuf = (char *)malloc(totalcount);
22     cbuf = (char *)malloc(totalcount);
23     MPI_Gatherv( lbuf, position, MPI_PACKED, rbuf,
24                 counts, displs, MPI_PACKED, root, comm);
25
26     /* unpack all messages and concatenate strings */
27     concat_pos = 0;
28     for (i=0; i < gsize; i++) {
29         position = 0;
30         MPI_Unpack( rbuf+displs[i], totalcount-displs[i],
31                     &position, &count, 1, MPI_INT, comm);
32         MPI_Unpack( rbuf+displs[i], totalcount-displs[i],
33                     &position, cbuf+concat_pos, count, MPI_CHAR, comm);
34         concat_pos += count;
35     }
36     cbuf[concat_pos] = '\0';
37 }

```

### 4.3 Canonical MPI\_PACK and MPI\_UNPACK

These functions read/write data to/from the buffer in the “external32” data format specified in Section 13.5.2, and calculate the size needed for packing. Their first arguments specify the data format, for future extensibility, but currently the only valid value of the `datarep` argument is “external32.”

*Advice to users.* These functions could be used, for example, to send typed data in a portable format from one MPI implementation to another. (*End of advice to users.*)

The buffer will contain exactly the packed data, without headers. MPI\_BYTE should be used to send and receive data that is packed using MPI\_PACK\_EXTERNAL.

*Rationale.* MPI\_PACK\_EXTERNAL specifies that there is no header on the message and further specifies the exact format of the data. Since MPI\_PACK may (and is allowed to) use a header, the datatype MPI\_PACKED cannot be used for data packed with MPI\_PACK\_EXTERNAL. (*End of rationale.*)

MPI\_PACK\_EXTERNAL(datarep, inbuf, incount, datatype, outbuf, outsize, position )

IN	datarep	data representation (string)
IN	inbuf	input buffer start (choice)
IN	incount	number of input data items (integer)
IN	datatype	datatype of each input data item (handle)
OUT	outbuf	output buffer start (choice)
IN	outsize	output buffer size, in bytes (integer)
INOUT	position	current position in buffer, in bytes (integer)

```
int MPI_Pack_external(char *datarep, void *inbuf, int incount,
                     MPI_Datatype datatype, void *outbuf, MPI_Aint outsize,
                     MPI_Aint *position)
```

```
MPI_PACK_EXTERNAL(DATAREP, INBUF, INCOUNT, DATATYPE, OUTBUF, OUTSIZE,
                  POSITION, IERROR)
```

```
INTEGER INCOUNT, DATATYPE, IERROR
INTEGER(KIND=MPI_ADDRESS_KIND) OUTSIZE, POSITION
CHARACTER*(*) DATAREP
<type> INBUF(*), OUTBUF(*)
```

```
{void MPI::Datatype::Pack_external(const char* datarep, const void* inbuf,
                                   int incount, void* outbuf, MPI::Aint outsize,
                                   MPI::Aint& position) const(binding deprecated, see Section 15.2) }
```

MPI\_UNPACK\_EXTERNAL(datarep, inbuf, insize, position, outbuf, outsize, position )

IN	datarep	data representation (string)
IN	inbuf	input buffer start (choice)
IN	insize	input buffer size, in bytes (integer)
INOUT	position	current position in buffer, in bytes (integer)
OUT	outbuf	output buffer start (choice)
IN	outcount	number of output data items (integer)
IN	datatype	datatype of output data item (handle)

```

1  int MPI_Unpack_external(char *datarep, void *inbuf, MPI_Aint insize,
2      MPI_Aint *position, void *outbuf, int outcount,
3      MPI_Datatype datatype)
4
5  MPI_UNPACK_EXTERNAL(DATAREP, INBUF, INSIZE, POSITION, OUTBUF, OUTCOUNT,
6      DATATYPE, IERROR)
7      INTEGER OUTCOUNT, DATATYPE, IERROR
8      INTEGER(KIND=MPI_ADDRESS_KIND) INSIZE, POSITION
9      CHARACTER*(*) DATAREP
10     <type> INBUF(*), OUTBUF(*)
11
12 {void MPI::Datatype::Unpack_external(const char* datarep,
13     const void* inbuf, MPI::Aint insize, MPI::Aint& position,
14     void* outbuf, int outcount) const (binding deprecated, see
15     Section 15.2) }
16
17 MPI_PACK_EXTERNAL_SIZE( datarep, incount, datatype, size )
18
19     IN      datarep      data representation (string)
20     IN      incount      number of input data items (integer)
21     IN      datatype      datatype of each input data item (handle)
22     OUT     size          output buffer size, in bytes (integer)
23
24
25 int MPI_Pack_external_size(char *datarep, int incount,
26     MPI_Datatype datatype, MPI_Aint *size)
27
28 MPI_PACK_EXTERNAL_SIZE(DATAREP, INCOUNT, DATATYPE, SIZE, IERROR)
29     INTEGER INCOUNT, DATATYPE, IERROR
30     INTEGER(KIND=MPI_ADDRESS_KIND) SIZE
31     CHARACTER*(*) DATAREP
32
33 {MPI::Aint MPI::Datatype::Pack_external_size(const char* datarep,
34     int incount) const(binding deprecated, see Section 15.2) }
35
36
37
38
39
40
41
42
43
44
45
46
47
48

```

# Chapter 5

## Collective Communication

### 5.1 Introduction and Overview

Collective communication is defined as communication that involves a group or groups of processes. The functions of this type provided by MPI are the following:

- `MPI_BARRIER`, `MPI_IBARRIER`: Barrier synchronization across all members of a group (Section 5.3 and Section 5.12.1 ).
- `MPI_BCAST` , `MPI_IBCAST` : Broadcast from one member to all members of a group (Section 5.4 and Section 5.12.2). This is shown as “broadcast” in Figure 5.1.
- `MPI_GATHER`, `MPI_IGATHER`, `MPI_GATHERV`, `MPI_IGATHERV` , `MPI_GATHERW`, `MPI_IGATHERW`: Gather data from all members of a group to one member (Section 5.5 and Section 5.12.3). This is shown as “gather” in Figure 5.1.
- `MPI_SCATTER`, `MPI_ISCATTER` , `MPI_SCATTERV`, `MPI_ISCATTERV` , `MPI_SCATTERW`, `MPI_ISCATTERW`: Scatter data from one member to all members of a group (Section 5.6 and Section 5.12.4). This is shown as “scatter” in Figure 5.1.
- `MPI_ALLGATHER`, `MPI_IALLGATHER` , `MPI_ALLGATHERV`, `MPI_IALLGATHERV` , `MPI_ALLGATHERW`, `MPI_IALLGATHERW`: A variation on Gather where all members of a group receive the result (Section 5.7 and Section 5.12.5). This is shown as “allgather” in Figure 5.1.
- `MPI_ALLTOALL`, `MPI_IALLTOALL` , `MPI_ALLTOALLV`, `MPI_IALLTOALLV` , `MPI_ALLTOALLW`, `MPI_IALLTOALLW` : Scatter/Gather data from all members to all members of a group (also called complete exchange) (Section 5.8 and Section 5.12.6). This is shown as “complete exchange” in Figure 5.1.
- `MPI_ALLREDUCE`, `MPI_IALLREDUCE` , `MPI_REDUCE`, `MPI_IREDUCE` : Global reduction operations such as sum, max, min, or user-defined functions, where the result is returned to all members of a group (Section 5.9.6 and Section 5.12.8) and a variation where the result is returned to only one member (Section 5.9 and Section 5.12.7).
- `MPI_REDUCE_SCATTER_BLOCK`, `MPI_IREDUCE_SCATTER_BLOCK`, `MPI_REDUCE_SCATTER` , `MPI_IREDUCE_SCATTER` : A combined reduction and scatter operation (Section 5.10, Section 5.12.9, and Section 5.12.10).

- `MPI_SCAN`, `MPI_ISCAN`, `MPI_EXSCAN`, `MPI_IEXSCAN`: Scan across all members of a group (also called prefix) (Section 5.11, Section 5.11.2, Section 5.12.11, and Section 5.12.12).

One of the key arguments in a call to a collective routine is a communicator that defines the group or groups of participating processes and provides a context for the operation. This is discussed further in Section 5.2. The syntax and semantics of the collective operations are defined to be consistent with the syntax and semantics of the point-to-point operations. Thus, general datatypes are allowed and must match between sending and receiving processes as specified in Chapter 4. Several collective routines such as broadcast and gather have a single originating or receiving process. Such a process is called the *root*. Some arguments in the collective functions are specified as “significant only at root,” and are ignored for all participants except the root. The reader is referred to Chapter 4 for information concerning communication buffers, general datatypes and type matching rules, and to Chapter 6 for information on how to define groups and create communicators.

The type-matching conditions for the collective operations are more strict than the corresponding conditions between sender and receiver in point-to-point. Namely, for collective operations, the amount of data sent must exactly match the amount of data specified by the receiver. Different type maps (the layout in memory, see Section 4.1) between sender and receiver are still allowed.

Collective [routine calls]operations can (but are not required to) [return]complete as soon as [their]the caller’s participation in the collective communication is [complete]finished. A blocking operation is complete as soon as the call returns. A nonblocking (immediate) call requires a separate completion call (cf. Section 3.7). The completion of a [call]collective operation indicates that the caller is [now] free to modify locations in the communication buffer. It does not indicate that other processes in the group have completed or even started the operation (unless otherwise implied by the description of the operation). [Thus, a collective communication call may, or may not, have the effect of synchronizing all calling processes. This statement excludes, of course, the barrier function]Thus, a collective communication operation may, or may not, have the effect of synchronizing all calling processes. This statement excludes, of course, the barrier operation.

Collective communication calls may use the same communicators as point-to-point communication; MPI guarantees that messages generated on behalf of collective communication calls will not be confused with messages generated by point-to-point communication. The collective operations do not have a message tag argument. A more detailed discussion of correct use of collective routines is found in Section 5.13.

*Rationale.* The equal-data restriction (on type matching) was made so as to avoid the complexity of providing a facility analogous to the status argument of `MPI_RECV` for discovering the amount of data sent. Some of the collective routines would require an array of status values.

The statements about synchronization are made so as to allow a variety of implementations of the collective functions.

[The collective operations do not accept a message tag argument. If future revisions of MPI define nonblocking collective functions, then tags (or a similar mechanism) might need to be added so as to allow the dis-ambiguation of multiple, pending, collective operations.] (*End of rationale.*)




Figure 5.1: Collective move functions illustrated for a group of six processes. In each case, each row of boxes represents data locations in one process. Thus, in the broadcast, initially just the first process contains the data  $A_0$ , but after the broadcast all processes contain it.



*Advice to users.* It is dangerous to rely on synchronization side-effects of the collective operations for program correctness. For example, even though a particular implementation may provide a broadcast routine with a side-effect of synchronization, the standard does not require this, and a program that relies on this will not be portable.

On the other hand, a correct, portable program must allow for the fact that a collective call *may* be synchronizing. Though one cannot rely on any synchronization side-effect, one must program so as to allow it. These issues are discussed further in Section 5.13. (*End of advice to users.*)

*Advice to implementors.* While vendors may write optimized collective routines matched to their architectures, a complete library of the collective communication routines can be written entirely using the MPI point-to-point communication functions and a few auxiliary functions. If implementing on top of point-to-point, a hidden, special communicator might be created for the collective operation so as to avoid interference with any on-going point-to-point communication at the time of the collective call. This is discussed further in Section 5.13. (*End of advice to implementors.*)

Many of the descriptions of the collective routines provide illustrations in terms of blocking MPI point-to-point routines. These are intended solely to indicate what data is sent or received by what process. Many of these examples are *not* correct MPI programs; for purposes of simplicity, they often assume infinite buffering.

## 5.2 Communicator Argument

The key concept of the collective functions is to have a group or groups of participating processes. The routines do not have group identifiers as explicit arguments. Instead, there is a communicator argument. Groups and communicators are discussed in full detail in Chapter 6. For the purposes of this chapter, it is sufficient to know that there are two types of communicators: *intra-communicators* and *inter-communicators*. An intracommunicator can be thought of as an identifier for a single group of processes linked with a context. An intercommunicator identifies two distinct groups of processes linked with a context.

### 5.2.1 Specifics for Intracommunicator Collective Operations

All processes in the group identified by the intracommunicator must call the collective routine.

In many cases, collective communication can occur “in place” for intracommunicators, with the output buffer being identical to the input buffer. This is specified by providing a special argument value, `MPI_IN_PLACE`, instead of the send buffer or the receive buffer argument, depending on the operation performed.

*Rationale.* The “in place” operations are provided to reduce unnecessary memory motion by both the MPI implementation and by the user. Note that while the simple check of testing whether the send and receive buffers have the same address will work for some cases (e.g., `MPI_ALLREDUCE`), they are inadequate in others (e.g., `MPI_GATHER`, with root not equal to zero). Further, Fortran explicitly prohibits aliasing of arguments; the approach of using a special value to denote “in place” operation eliminates that difficulty. (*End of rationale.*)

*Advice to users.* By allowing the “in place” option, the receive buffer in many of the collective calls becomes a send-and-receive buffer. For this reason, a Fortran binding that includes INTENT must mark these as INOUT, not OUT.

Note that MPI\_IN\_PLACE is a special kind of value; it has the same restrictions on its use that MPI\_BOTTOM has. [ Some intracommunicator collective operations do not support the “in place” option (e.g., MPI\_ALLTOALLV).] (*End of advice to users.*)

## 5.2.2 Applying Collective Operations to Intercommunicators

To understand how collective operations apply to intercommunicators, we can view most MPI intracommunicator collective operations as fitting one of the following categories (see, for instance, [48]):

**All-To-All** All processes contribute to the result. All processes receive the result.

- MPI\_ALLGATHER, MPI\_IALLGATHER, MPI\_ALLGATHERV, MPI\_IALLGATHERV, MPI\_ALLGATHERW, MPI\_IALLGATHERW
- MPI\_ALLTOALL, MPI\_IALLTOALL, MPI\_ALLTOALLV, MPI\_IALLTOALLV, MPI\_ALLTOALLW, MPI\_IALLTOALLW
- MPI\_ALLREDUCE, MPI\_IALLREDUCE, MPI\_REDUCE\_SCATTER\_BLOCK, MPI\_IREDUCE\_SCATTER\_BLOCK, MPI\_REDUCE\_SCATTER, MPI\_IREDUCE\_SCATTER
- MPI\_BARRIER, MPI\_IBARRIER

**All-To-One** All processes contribute to the result. One process receives the result.

- MPI\_GATHER, MPI\_IGATHER, MPI\_GATHERV, MPI\_IGATHERV, MPI\_GATHERW, MPI\_IGATHERW
- MPI\_REDUCE, MPI\_IREDUCE

**One-To-All** One process contributes to the result. All processes receive the result.

- MPI\_BCAST, MPI\_IBCAST
- MPI\_SCATTER, MPI\_ISCATTER, MPI\_SCATTERV, MPI\_ISCATTERV, MPI\_SCATTERW, MPI\_ISCATTERW

**Other** Collective operations that do not fit into one of the above categories.

- MPI\_SCAN, MPI\_ISCAN, MPI\_EXSCAN, MPI\_IEXSCAN

The data movement patterns of MPI\_SCAN, MPI\_ISCAN [and], MPI\_EXSCAN, and MPI\_IEXSCAN do not fit this taxonomy.

The application of collective communication to intercommunicators is best described in terms of two groups. For example, an all-to-all MPI\_ALLGATHER operation can be described as collecting data from all members of one group with the result appearing in all members of the other group (see Figure 5.2). As another example, a one-to-all MPI\_BCAST operation sends data from one member of one group to all members of the other group. Collective computation operations such as MPI\_REDUCE\_SCATTER have a similar interpretation (see Figure 5.3). For intracommunicators, these two groups are the

same. For intercommunicators, these two groups are distinct. For the all-to-all operations, each such operation is described in two phases, so that it has a symmetric, full-duplex behavior.

The following collective operations also apply to intercommunicators:

- MPI\_BARRIER, MPI\_IBARRIER
- MPI\_BCAST, MPI\_IBCAST
- MPI\_GATHER, MPI\_IGATHER, MPI\_GATHERV, MPI\_IGATHERV,
- MPI\_SCATTER, MPI\_ISCATTER, MPI\_SCATTERV, MPI\_ISCATTERV,
- MPI\_ALLGATHER, MPI\_IALLGATHER, MPI\_ALLGATHERV, MPI\_IALLGATHERV,
- MPI\_ALLTOALL, MPI\_IALLTOALL, MPI\_ALLTOALLV, MPI\_IALLTOALLV,
- MPI\_ALLTOALLW, MPI\_IALLTOALLW,
- MPI\_ALLREDUCE, MPI\_IALLREDUCE, MPI\_REDUCE, MPI\_IREDUCE,
- MPI\_REDUCE\_SCATTER\_BLOCK, MPI\_IREDUCE\_SCATTER\_BLOCK,
- MPI\_REDUCE\_SCATTER, MPI\_IREDUCE\_SCATTER.

In C++, the bindings for these functions are in the `MPI::Comm` class. However, since the collective operations do not make sense on a C++ `MPI::Comm` (as it is neither an intercommunicator nor an intracommunicator), the functions are all pure virtual.




Figure 5.2: Intercommunicator allgather. The focus of data to one process is represented, not mandated by the semantics. The two phases do allgathers in both directions.




Figure 5.3: Intercommunicator reduce-scatter. The focus of data to one process is represented, not mandated by the semantics. The two phases do reduce-scatters in both directions.

### 5.2.3 Specifics for Intercommunicator Collective Operations

All processes in both groups identified by the intercommunicator must call the collective routine.

Note that the “in place” option for intracommunicators does not apply to intercommunicators since in the intercommunicator case there is no communication from a process to itself.

For intercommunicator collective communication, if the operation is in the All-To-One or One-To-All categories, then the transfer is unidirectional. The direction of the transfer is indicated by a special value of the root argument. In this case, for the group containing the root process, all processes in the group must call the routine using a special argument for the root. For this, the root process uses the special root value `MPI_ROOT`; all other processes in the same group as the root use `MPI_PROC_NULL`. All processes in the other group (the group that is the remote group relative to the root process) must call the collective routine and provide the rank of the root. If the operation is in the All-To-All category, then the transfer is bidirectional.

*Rationale.* Operations in the All-To-One and One-To-All categories are unidirectional by nature, and there is a clear way of specifying direction. Operations in the All-To-All category will often occur as part of an exchange, where it makes sense to communicate in both directions at once. (*End of rationale.*)

## 5.3 Barrier Synchronization

`MPI_BARRIER(comm)`

IN            `comm`                            communicator (handle)

`int MPI_Barrier(MPI_Comm comm)`

`MPI_BARRIER(COMM, IERROR)`

INTEGER `COMM`, `IERROR`

{`void MPI::Comm::Barrier()` `const = 0`(*binding deprecated, see Section 15.2*) }

If `comm` is an intracommunicator, `MPI_BARRIER` blocks the caller until all group members have called it. The call returns at any process only after all group members have entered the call.

If `comm` is an intercommunicator, `MPI_BARRIER` involves two groups. The call returns at processes in one group (group A) of the intercommunicator only after all members of the other group (group B) have entered the call (and vice versa). A process may return from the call before all processes in its own group have entered the call.

## 5.4 Broadcast

`MPI_BCAST(buffer, count, datatype, root, comm)`

INOUT    `buffer`                            starting address of buffer (choice)

IN        `count`                            number of entries in buffer (non-negative integer)

IN        `datatype`                        data type of buffer (handle)

IN        `root`                            rank of broadcast root (integer)

IN        `comm`                            communicator (handle)

`int MPI_Bcast(void* buffer, int count, MPI_Datatype datatype, int root, MPI_Comm comm)`

`MPI_BCAST(BUFFER, COUNT, DATATYPE, ROOT, COMM, IERROR)`

<type> `BUFFER(*)`

INTEGER `COUNT`, `DATATYPE`, `ROOT`, `COMM`, `IERROR`

{`void MPI::Comm::Bcast(void* buffer, int count, const MPI::Datatype& datatype, int root)` `const = 0`(*binding deprecated, see Section 15.2*) }

If `comm` is an intracommunicator, `MPI_BCAST` broadcasts a message from the process with rank `root` to all processes of the group, itself included. It is called by all members of the group using the same arguments for `comm` and `root`. On return, the content of `root`'s buffer is copied to all other processes.

General, derived datatypes are allowed for `datatype`. The type signature of `count`, `datatype` on any process must be equal to the type signature of `count`, `datatype` at the root. This implies that the amount of data sent must be equal to the amount received, pairwise between each process and the root. `MPI_BCAST` and all other data-movement collective routines make this restriction. Distinct type maps between sender and receiver are still allowed.

The “in place” option is not meaningful here.

If `comm` is an intercommunicator, then the call involves all processes in the intercommunicator, but with one group (group A) defining the root process. All processes in the other group (group B) pass the same value in argument `root`, which is the rank of the root in group A. The root passes the value `MPI_ROOT` in `root`. All other processes in group A pass the value `MPI_PROC_NULL` in `root`. Data is broadcast from the root to all processes in group B. The buffer arguments of the processes in group B must be consistent with the buffer argument of the root.

#### 5.4.1 Example using `MPI_BCAST`

The examples in this section use intracommunicators.

##### Example 5.1

Broadcast 100 ints from process 0 to every process in the group.

```
MPI_Comm comm;
int array[100];
int root=0;
...
MPI_Bcast(array, 100, MPI_INT, root, comm);
```

As in many of our example code fragments, we assume that some of the variables (such as `comm` in the above) have been assigned appropriate values.

## 5.5 Gather

`MPI_GATHER(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, root, comm)`

IN	sendbuf	starting address of send buffer (choice)
IN	sendcount	number of elements in send buffer (non-negative integer)
IN	sendtype	data type of send buffer elements (handle)
OUT	recvbuf	address of receive buffer (choice, significant only at root)
IN	recvcount	number of elements for any single receive (non-negative integer, significant only at root)
IN	recvtype	data type of recv buffer elements (significant only at root) (handle)
IN	root	rank of receiving process (integer)
IN	comm	communicator (handle)

```
int MPI_Gather(void* sendbuf, int sendcount, MPI_Datatype sendtype,
              void* recvbuf, int recvcount, MPI_Datatype recvtype, int root,
              MPI_Comm comm)
```

```
MPI_GATHER(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNT, RECVTYPE,
           ROOT, COMM, IERROR)
<type> SENDBUF(*), RECVBUF(*)
INTEGER SENDCOUNT, SENDTYPE, RECVCOUNT, RECVTYPE, ROOT, COMM, IERROR
```

```
{void MPI::Comm::Gather(const void* sendbuf, int sendcount, const
                        MPI::Datatype& sendtype, void* recvbuf, int recvcount,
                        const MPI::Datatype& recvtype, int root) const = 0(binding
                        deprecated, see Section 15.2) }
```

If `comm` is an intracommunicator, each process (root process included) sends the contents of its send buffer to the root process. The root process receives the messages and stores them in rank order. The outcome is *as if* each of the `n` processes in the group (including the root process) had executed a call to

```
MPI_Send(sendbuf, sendcount, sendtype, root, ...),
```

and the root had executed `n` calls to

```
MPI_Recv(recvbuf + i · recvcount · extent(recvtype), recvcount, recvtype, i, ...),
```

where `extent(recvtype)` is the type extent obtained from a call to `MPI_Type_get_extent()`.

An alternative description is that the `n` messages sent by the processes in the group are concatenated in rank order, and the resulting message is received by the root as if by a call to `MPI_RECV(recvbuf, recvcount·n, recvtype, ...)`.

The receive buffer is ignored for all non-root processes.

General, derived datatypes are allowed for both `sendtype` and `recvtype`. The type signature of `sendcount`, `sendtype` on each process must be equal to the type signature of `recvcount`, `recvtype` at the root. This implies that the amount of data sent must be equal to the amount of data received, pairwise between each process and the root. Distinct type maps between sender and receiver are still allowed.

All arguments to the function are significant on process `root`, while on other processes, only arguments `sendbuf`, `sendcount`, `sendtype`, `root`, and `comm` are significant. The arguments `root` and `comm` must have identical values on all processes.

The specification of counts and types should not cause any location on the root to be written more than once. Such a call is erroneous.

Note that the `recvcount` argument at the root indicates the number of items it receives from *each* process, not the total number of items it receives.

The “in place” option for intracommunicators is specified by passing `MPI_IN_PLACE` as the value of `sendbuf` at the root. In such a case, `sendcount` and `sendtype` are ignored, and the contribution of the root to the gathered vector is assumed to be already in the correct place in the receive buffer.

If `comm` is an intercommunicator, then the call involves all processes in the intercommunicator, but with one group (group A) defining the root process. All processes in the other group (group B) pass the same value in argument `root`, which is the rank of the root in group A. The root passes the value `MPI_ROOT` in `root`. All other processes in group A pass the value `MPI_PROC_NULL` in `root`. Data is gathered from all processes in group B to the root. The send buffer arguments of the processes in group B must be consistent with the receive buffer argument of the root.



1	MPI_GATHERV(sendbuf, sendcount, sendtype, recvbuf, recvcunts, displs, recvtype, root,		
2	comm)		
3	IN	sendbuf	starting address of send buffer (choice)
4	IN	sendcount	number of elements in send buffer (non-negative integer)
5	IN	sendtype	data type of send buffer elements (handle)
6	OUT	recvbuf	address of receive buffer (choice, significant only at root)
7	IN	recvcunts	non-negative integer array (of length group size) containing the number of elements that are received from each process (significant only at root)
8	IN	displs	integer array (of length group size). Entry <i>i</i> specifies the displacement relative to <i>recvbuf</i> at which to place the incoming data from process <i>i</i> (significant only at root)
9	IN	recvtype	data type of recv buffer elements (significant only at root) (handle)
10	IN	root	rank of receiving process (integer)
11	IN	comm	communicator (handle)

```

24 int MPI_Gatherv(void* sendbuf, int sendcount, MPI_Datatype sendtype,
25               void* recvbuf, int *recvcunts, int *displs,
26               MPI_Datatype recvtype, int root, MPI_Comm comm)
27
28 MPI_GATHERV(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNTS, DISPLS,
29             RECVTYPE, ROOT, COMM, IERROR)
30 <type> SENDBUF(*), RECVBUF(*)
31 INTEGER SENDCOUNT, SENDTYPE, RECVCOUNTS(*), DISPLS(*), RECVTYPE, ROOT,
32 COMM, IERROR

```

```

33 {void MPI::Comm::Gatherv(const void* sendbuf, int sendcount, const
34                         MPI::Datatype& sendtype, void* recvbuf,
35                         const int recvcunts[], const int displs[],
36                         const MPI::Datatype& recvtype, int root) const = 0(binding
37                         deprecated, see Section 15.2) }
38

```

MPI\_GATHERV extends the functionality of MPI\_GATHER by allowing a varying count of data from each process, since *recvcunts* is now an array. It also allows more flexibility as to where the data is placed on the root, by providing the new argument, *displs*.

If *comm* is an intracommunicator, the outcome is *as if* each process, including the root process, sends a message to the root,

```
MPI_Send(sendbuf, sendcount, sendtype, root, ...),
```

and the root executes *n* receives,

```
MPI_Recv(recvbuf + displs[j] · extent(recvtype), recvcunts[j], recvtype, i, ...).
```

The data received from process `j` is placed into `recvbuf` of the `root` process beginning at offset `displs[j]` elements (in terms of the `recvtype`).

The receive buffer is ignored for all non-root processes.

The type signature implied by `sendcount`, `sendtype` on process `i` must be equal to the type signature implied by `recvcounts[i]`, `recvtype` at the root. This implies that the amount of data sent must be equal to the amount of data received, pairwise between each process and the root. Distinct type maps between sender and receiver are still allowed, as illustrated in Example 5.6.

All arguments to the function are significant on process `root`, while on other processes, only arguments `sendbuf`, `sendcount`, `sendtype`, `root`, and `comm` are significant. The arguments `root` and `comm` must have identical values on all processes.

The specification of counts, types, and displacements should not cause any location on the root to be written more than once. Such a call is erroneous.

The “in place” option for intracommunicators is specified by passing `MPI_IN_PLACE` as the value of `sendbuf` at the root. In such a case, `sendcount` and `sendtype` are ignored, and the contribution of the root to the gathered vector is assumed to be already in the correct place in the receive buffer

If `comm` is an intercommunicator, then the call involves all processes in the intercommunicator, but with one group (group A) defining the root process. All processes in the other group (group B) pass the same value in argument `root`, which is the rank of the root in group A. The root passes the value `MPI_ROOT` in `root`. All other processes in group A pass the value `MPI_PROC_NULL` in `root`. Data is gathered from all processes in group B to the root. The send buffer arguments of the processes in group B must be consistent with the receive buffer argument of the root.

ticket269.

```

1 MPI_GATHERW(sendbuf, sendcount, sendtype, recvbuf, recvcoun, displs, recvtypes, root,
2             comm)
3
4     IN      sendbuf      starting address of send buffer (choice)
5
6     IN      sendcount    number of elements in send buffer (non-negative integer)
7
8     IN      sendtype     data type of send buffer elements (handle)
9
10    OUT     recvbuf      address of receive buffer (choice, significant only at root)
11
12    IN      recvcoun     non-negative integer array (of length group size) containing the number of elements that are received from each process (significant only at root)
13
14    IN      displs       integer array (of length group size). Entry i specifies the displacement relative to recvbuf at which to place the incoming data from process i (significant only at root)
15
16
17
18    IN      recvtypes     array of datatypes (of length group size). Entry i specifies the type of data received from process i (array of handles)
19
20
21
22    IN      root         rank of receiving process (integer)
23
24    IN      comm         communicator (handle)

```

```

25 int MPI_Gatherw(void* sendbuf, int sendcount, MPI_Datatype sendtype,
26               void* recvbuf, int recvcoun[], int displs[],
27               MPI_Datatype recvtypes[], int root, MPI_Comm comm)
28

```

```

29 MPI_GATHERW(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNTS, DISPLS,
30             RECVTYPES, ROOT, COMM, IERROR)
31 <type> SENDBUF(*), RECVBUF(*)
32 INTEGER SENDCOUNT, SENDTYPE, RECVCOUNTS(*), DISPLS(*), RECVTYPES(*),
33 ROOT, COMM, IERROR

```

MPI\_GATHERW extends the functionality of MPI\_GATHERV by allowing varying datatype specifications for data received from each process. If `comm` is an intracommunicator, the outcome is as if each process, including the root process, sends a message to the root,

```

38 MPI_Send(sendbuf, sendcount, sendtype, root, ...),

```

and the root executes `n` receives,

```

41 MPI_Recv(recvbuf + displs[j] * extent(recvtypes[j]), recvcoun[j], recvtypes[j], i, ...).

```

The data received from process `j` is placed into `recvbuf` of the root process beginning at offset `displs[j]` elements (in terms of `recvtypes[j]`).

The receive buffer is ignored for all non-root processes.

The type signature implied by `sendcount`, `sendtype` on process `i` must be equal to the type signature implied by `recvcoun[i]`, `recvtypes[i]` at the root. This implies that the amount

of data sent must be equal to the amount of data received, pairwise between each process and the root.

All arguments to the function are significant on process root, while on other processes, only arguments `sendbuf`, `sendcount`, `sendtype`, `root`, and `comm` are significant. The arguments `root` and `comm` must have identical values on all processes. The specification of counts, types, and displacements should not cause any location on the root to be written more than once. Such a call is erroneous.

The “in place” option for intracommunicators is specified by passing `MPI_IN_PLACE` as the value of `sendbuf` at the root. In such a case, `sendcount` and `sendtype` are ignored, and the contribution of the root to the gathered vector is assumed to be already in the correct place in the receive buffer.

If `comm` is an intercommunicator, then the call involves all processes in the intercommunicator, but with one group (group A) defining the root process. All processes in the other group (group B) pass the same value in argument `root`, which is the rank of the root in group A. The root passes the value `MPI_ROOT` in `root`. All other processes in group A pass the value `MPI_PROC_NULL` in `root`. Data is gathered from all processes in group B to the root. The send buffer arguments of the processes in group B must be consistent with the receive buffer arguments of the root.

### 5.5.1 Examples using `MPI_GATHER`, `MPI_GATHERV`

The examples in this section use intracommunicators.

#### Example 5.2

Gather 100 ints from every process in group to root. See [\[f\]Figure 5.4](#).

```
MPI_Comm comm;
int gsize, sendarray[100];
int root, *rbuf;
...
MPI_Comm_size(comm, &gsize);
rbuf = (int *)malloc(gsize*100*sizeof(int));
MPI_Gather(sendarray, 100, MPI_INT, rbuf, 100, MPI_INT, root, comm);
```

#### Example 5.3

Previous example modified – only the root allocates memory for the receive buffer.

```
MPI_Comm comm;
int gsize, sendarray[100];
int root, myrank, *rbuf;
...
MPI_Comm_rank(comm, &myrank);
if (myrank == root) {
    MPI_Comm_size(comm, &gsize);
    rbuf = (int *)malloc(gsize*100*sizeof(int));
}
MPI_Gather(sendarray, 100, MPI_INT, rbuf, 100, MPI_INT, root, comm);
```




Figure 5.4: The root process gathers 100 ints from each process in the group.

#### Example 5.4

Do the same as the previous example, but use a derived datatype. Note that the type cannot be the entire set of `gsize*100` ints since type matching is defined pairwise between the root and each process in the gather.

```

16 MPI_Comm comm;
17 int gsize, sendarray[100];
18 int root, *rbuf;
19 MPI_Datatype rtype;
20 ...
21 MPI_Comm_size(comm, &gsize);
22 MPI_Type_contiguous(100, MPI_INT, &rtype);
23 MPI_Type_commit(&rtype);
24 rbuf = (int *)malloc(gsize*100*sizeof(int));
25 MPI_Gather(sendarray, 100, MPI_INT, rbuf, 1, rtype, root, comm);

```

#### Example 5.5

Now have each process send 100 ints to root, but place each set (of 100) `stride` ints apart at receiving end. Use `MPI_GATHERV` and the `displs` argument to achieve this effect. Assume `stride ≥ 100`. See Figure 5.5.

```

32 MPI_Comm comm;
33 int gsize, sendarray[100];
34 int root, *rbuf, stride;
35 int *displs, i, *rcounts;
36 ...
37 ...
38 ...
39 MPI_Comm_size(comm, &gsize);
40 rbuf = (int *)malloc(gsize*stride*sizeof(int));
41 displs = (int *)malloc(gsize*sizeof(int));
42 rcounts = (int *)malloc(gsize*sizeof(int));
43 for (i=0; i<gsize; ++i) {
44     displs[i] = i*stride;
45     rcounts[i] = 100;
46 }
47 MPI_Gatherv(sendarray, 100, MPI_INT, rbuf, rcounts, displs, MPI_INT,
48             root, comm);

```




Figure 5.5: The root process gathers 100 ints from each process in the group, each set is placed `stride` ints apart.

Note that the program is erroneous if `stride < 100`.

### Example 5.6

Same as Example 5.5 on the receiving side, but send the 100 ints from the 0th column of a 100×150 int array, in C. See Figure 5.6.

```

MPI_Comm comm;
int gsize, sendarray[100][150];
int root, *rbuf, stride;
MPI_Datatype stype;
int *displs, i, *rcounts;

...

MPI_Comm_size(comm, &gsize);
rbuf = (int *)malloc(gsize*stride*sizeof(int));
displs = (int *)malloc(gsize*sizeof(int));
rcounts = (int *)malloc(gsize*sizeof(int));
for (i=0; i<gsize; ++i) {
    displs[i] = i*stride;
    rcounts[i] = 100;
}
/* Create datatype for 1 column of array
 */
MPI_Type_vector(100, 1, 150, MPI_INT, &stype);
MPI_Type_commit(&stype);
MPI_Gatherv(sendarray, 1, stype, rbuf, rcounts, displs, MPI_INT,
            root, comm);

```

### Example 5.7

Process `i` sends `(100-i)` ints from the `i`-th column of a 100 × 150 int array, in C. It is received into a buffer with stride, as in the previous two examples. See Figure 5.7.




Figure 5.6: The root process gathers column 0 of a 100×150 C array, and each set is placed `stride` ints apart.

```

MPI_Comm comm;
int gsize, sendarray[100][150], *sptr;
int root, *rbuf, stride, myrank;
MPI_Datatype stype;
int *displs, i, *rcounts;

...

MPI_Comm_size(comm, &gsize);
MPI_Comm_rank(comm, &myrank);
rbuf = (int *)malloc(gsize*stride*sizeof(int));
displs = (int *)malloc(gsize*sizeof(int));
rcounts = (int *)malloc(gsize*sizeof(int));
for (i=0; i<gsize; ++i) {
    displs[i] = i*stride;
    rcounts[i] = 100-i;    /* note change from previous example */
}
/* Create datatype for the column we are sending
 */
MPI_Type_vector(100-myrank, 1, 150, MPI_INT, &stype);
MPI_Type_commit(&stype);
/* sptr is the address of start of "myrank" column
 */
sptr = &sendarray[0][myrank];
MPI_Gatherv(sptr, 1, stype, rbuf, rcounts, displs, MPI_INT,
            root, comm);

```

Note that a different amount of data is received from each process.

### Example 5.8

Same as Example 5.7, but done in a different way at the sending end. We create a datatype that causes the correct striding at the sending end so that we read a column of a C array. A similar thing was done in Example 4.16, Section 4.1.14.

```

MPI_Comm comm;
int gsize, sendarray[100][150], *sptr;

```




Figure 5.7: The root process gathers  $100-i$  ints from column  $i$  of a  $100 \times 150$  C array, and each set is placed `stride` ints apart.

```

int root, *rbuf, stride, myrank, disp[2], blocklen[2];
MPI_Datatype stype, type[2];
int *displs, i, *rcounts;

...

MPI_Comm_size(comm, &gsize);
MPI_Comm_rank(comm, &myrank);
rbuf = (int *)malloc(gsize*stride*sizeof(int));
displs = (int *)malloc(gsize*sizeof(int));
rcounts = (int *)malloc(gsize*sizeof(int));
for (i=0; i<gsize; ++i) {
    displs[i] = i*stride;
    rcounts[i] = 100-i;
}
/* Create datatype for one int, with extent of entire row
 */
disp[0] = 0;      disp[1] = 150*sizeof(int);
type[0] = MPI_INT; type[1] = MPI_UB;
blocklen[0] = 1;  blocklen[1] = 1;
MPI_Type_create_struct(2, blocklen, disp, type, &stype);
MPI_Type_commit(&stype);
sptr = &sendarray[0][myrank];
MPI_Gatherv(sptr, 100-myrank, stype, rbuf, rcounts, displs, MPI_INT,
            root, comm);

```

### Example 5.9

Same as Example 5.7 at sending side, but at receiving side we make the stride between received blocks vary from block to block. See Figure 5.8.

```

MPI_Comm comm;
int gsize, sendarray[100][150], *sptr;
int root, *rbuf, *stride, myrank, bufsize;
MPI_Datatype stype;
int *displs, i, *rcounts, offset;

```






Figure 5.8: The root process gathers  $100-i$  ints from column  $i$  of a  $100 \times 150$  C array, and each set is placed  $\text{stride}[i]$  ints apart (a varying stride).

```

...
MPI_Comm_size(comm, &gsize);
MPI_Comm_rank(comm, &myrank);

stride = (int *)malloc(gsize*sizeof(int));
...
/* stride[i] for i = 0 to gsize-1 is set somehow
 */

/* set up displs and rcounts vectors first
 */
displs = (int *)malloc(gsize*sizeof(int));
rcounts = (int *)malloc(gsize*sizeof(int));
offset = 0;
for (i=0; i<gsize; ++i) {
    displs[i] = offset;
    offset += stride[i];
    rcounts[i] = 100-i;
}
/* the required buffer size for rbuf is now easily obtained
 */
bufsize = displs[gsize-1]+rcounts[gsize-1];
rbuf = (int *)malloc(bufsize*sizeof(int));
/* Create datatype for the column we are sending
 */
MPI_Type_vector(100-myrank, 1, 150, MPI_INT, &styp);
MPI_Type_commit(&styp);
sptr = &sendarray[0][myrank];
MPI_Gatherv(sptr, 1, styp, rbuf, rcounts, displs, MPI_INT,
            root, comm);

```

### Example 5.10

Process *i* sends *num* ints from the *i*-th column of a  $100 \times 150$  int array, in C. The complicating factor is that the various values of *num* are not known to *root*, so a separate gather must first be run to find these out. The data is placed contiguously at the receiving end.

```

MPI_Comm comm;
int gsize, sendarray[100][150], *sptr;
int root, *rbuf, myrank, disp[2], blocklen[2];
MPI_Datatype stype, type[2];
int *displs, i, *rcounts, num;

...

MPI_Comm_size(comm, &gsize);
MPI_Comm_rank(comm, &myrank);

/* First, gather nums to root
 */
rcounts = (int *)malloc(gsize*sizeof(int));
MPI_Gather(&num, 1, MPI_INT, rcounts, 1, MPI_INT, root, comm);
/* root now has correct rcounts, using these we set displs[] so
 * that data is placed contiguously (or concatenated) at receive end
 */
displs = (int *)malloc(gsize*sizeof(int));
displs[0] = 0;
for (i=1; i<gsize; ++i) {
    displs[i] = displs[i-1]+rcounts[i-1];
}
/* And, create receive buffer
 */
rbuf = (int *)malloc(gsize*(displs[gsize-1]+rcounts[gsize-1])
                    *sizeof(int));

/* Create datatype for one int, with extent of entire row
 */
disp[0] = 0;      disp[1] = 150*sizeof(int);
type[0] = MPI_INT; type[1] = MPI_UB;
blocklen[0] = 1;  blocklen[1] = 1;
MPI_Type_create_struct( 2, blocklen, disp, type, &stype );
MPI_Type_commit(&stype);
sptr = &sendarray[0][myrank];
MPI_Gatherv(sptr, num, stype, rbuf, rcounts, displs, MPI_INT,
            root, comm);

```

## 5.6 Scatter

**MPI\_SCATTER**(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, root, comm)

IN	sendbuf	address of send buffer (choice, significant only at root)
IN	sendcount	number of elements sent to each process (non-negative integer, significant only at root)
IN	sendtype	data type of send buffer elements (significant only at root) (handle)
OUT	recvbuf	address of receive buffer (choice)
IN	recvcount	number of elements in receive buffer (non-negative integer)
IN	recvtype	data type of receive buffer elements (handle)
IN	root	rank of sending process (integer)
IN	comm	communicator (handle)

```
int MPI_Scatter(void* sendbuf, int sendcount, MPI_Datatype sendtype,
               void* recvbuf, int recvcount, MPI_Datatype recvtype, int root,
               MPI_Comm comm)
```

```
MPI_SCATTER(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNT, RECVTYPE,
            ROOT, COMM, IERROR)
<type> SENDBUF(*), RECVBUF(*)
INTEGER SENDCOUNT, SENDTYPE, RECVCOUNT, RECVTYPE, ROOT, COMM, IERROR
```

```
{void MPI::Comm::Scatter(const void* sendbuf, int sendcount, const
                        MPI::Datatype& sendtype, void* recvbuf, int recvcount,
                        const MPI::Datatype& recvtype, int root) const = 0(binding
                        deprecated, see Section 15.2) }
```

**MPI\_SCATTER** is the inverse operation to **MPI\_GATHER**.

If *comm* is an intracommunicator, the outcome is *as if* the root executed *n* send operations,

```
MPI_Send(sendbuf + i · sendcount · extent(sendtype), sendcount, sendtype, i, ...),
```

and each process executed a receive,

```
MPI_Recv(recvbuf, recvcount, recvtype, i, ...).
```

An alternative description is that the root sends a message with **MPI\_Send**(sendbuf, sendcount·*n*, sendtype, ...). This message is split into *n* equal segments, the *i*-th segment is sent to the *i*-th process in the group, and each process receives this message as above.

The send buffer is ignored for all non-root processes.

The type signature associated with *sendcount*, *sendtype* at the root must be equal to the type signature associated with *recvcount*, *recvtype* at all processes (however, the type maps may be different). This implies that the amount of data sent must be equal to the

amount of data received, pairwise between each process and the root. Distinct type maps between sender and receiver are still allowed.

All arguments to the function are significant on process `root`, while on other processes, only arguments `recvbuf`, `recvcount`, `recvtype`, `root`, and `comm` are significant. The arguments `root` and `comm` must have identical values on all processes.

The specification of counts and types should not cause any location on the root to be read more than once.

*Rationale.* Though not needed, the last restriction is imposed so as to achieve symmetry with `MPI_GATHER`, where the corresponding restriction (a multiple-write restriction) is necessary. (*End of rationale.*)

The “in place” option for intracommunicators is specified by passing `MPI_IN_PLACE` as the value of `recvbuf` at the root. In such a case, `recvcount` and `recvtype` are ignored, and root “sends” no data to itself. The scattered vector is still assumed to contain  $n$  segments, where  $n$  is the group size; the  $root$ -th segment, which root should “send to itself,” is not moved.

If `comm` is an intercommunicator, then the call involves all processes in the intercommunicator, but with one group (group A) defining the root process. All processes in the other group (group B) pass the same value in argument `root`, which is the rank of the root in group A. The root passes the value `MPI_ROOT` in `root`. All other processes in group A pass the value `MPI_PROC_NULL` in `root`. Data is scattered from the root to all processes in group B. The receive buffer arguments of the processes in group B must be consistent with the send buffer argument of the root.

`MPI_SCATTERV(sendbuf, sendcounts, displs, sendtype, recvbuf, recvcount, recvtype, root, comm)`

IN	<code>sendbuf</code>	address of send buffer (choice, significant only at root)
IN	<code>sendcounts</code>	non-negative integer array (of length group size) specifying the number of elements to send to each processor
IN	<code>displs</code>	integer array (of length group size). Entry $i$ specifies the displacement (relative to <code>sendbuf</code> ) from which to take the outgoing data to process $i$
IN	<code>sendtype</code>	data type of send buffer elements (handle)
OUT	<code>recvbuf</code>	address of receive buffer (choice)
IN	<code>recvcount</code>	number of elements in receive buffer (non-negative integer)
IN	<code>recvtype</code>	data type of receive buffer elements (handle)
IN	<code>root</code>	rank of sending process (integer)
IN	<code>comm</code>	communicator (handle)

```
int MPI_Scatterv(void* sendbuf, int *sendcounts, int *displs,
                MPI_Datatype sendtype, void* recvbuf, int recvcount,
                MPI_Datatype recvtype, int root, MPI_Comm comm)
```

```

1 MPI_SCATTERV(SENDBUF, SENDCOUNTS, DISPLS, SENDTYPE, RECVBUF, REVCOUNT,
2             RECVTYPE, ROOT, COMM, IERROR)
3     <type> SENDBUF(*), RECVBUF(*)
4     INTEGER SENDCOUNTS(*), DISPLS(*), SENDTYPE, REVCOUNT, RECVTYPE, ROOT,
5     COMM, IERROR

```

```

6 {void MPI::Comm::Scatterv(const void* sendbuf, const int sendcounts[],
7     const int displs[], const MPI::Datatype& sendtype,
8     void* recvbuf, int recvcoun, const MPI::Datatype& recvtype,
9     int root) const = 0(binding deprecated, see Section 15.2) }

```

MPI\_SCATTERV is the inverse operation to MPI\_GATHERV.

MPI\_SCATTERV extends the functionality of MPI\_SCATTER by allowing a varying count of data to be sent to each process, since `sendcounts` is now an array. It also allows more flexibility as to where the data is taken from on the root, by providing an additional argument, `displs`.

If `comm` is an intracommunicator, the outcome is as if the root executed `n` send operations,

```
MPI_Send(sendbuf + displs[i] · extent(sendtype), sendcounts[i], sendtype, i, ...),
```

and each process executed a receive,

```
MPI_Recv(recvbuf, recvcoun, recvtype, i, ...).
```

The send buffer is ignored for all non-root processes.

The type signature implied by `sendcount[i]`, `sendtype` at the root must be equal to the type signature implied by `recvcoun`, `recvtype` at process `i` (however, the type maps may be different). This implies that the amount of data sent must be equal to the amount of data received, pairwise between each process and the root. Distinct type maps between sender and receiver are still allowed.

All arguments to the function are significant on process `root`, while on other processes, only arguments `recvbuf`, `recvcoun`, `recvtype`, `root`, and `comm` are significant. The arguments `root` and `comm` must have identical values on all processes.

The specification of counts, types, and displacements should not cause any location on the root to be read more than once.

The “in place” option for intracommunicators is specified by passing `MPI_IN_PLACE` as the value of `recvbuf` at the root. In such a case, `recvcoun` and `recvtype` are ignored, and root “sends” no data to itself. The scattered vector is still assumed to contain  $n$  segments, where  $n$  is the group size; the  $root$ -th segment, which root should “send to itself,” is not moved.

If `comm` is an intercommunicator, then the call involves all processes in the intercommunicator, but with one group (group A) defining the root process. All processes in the other group (group B) pass the same value in argument `root`, which is the rank of the root in group A. The root passes the value `MPI_ROOT` in `root`. All other processes in group A pass the value `MPI_PROC_NULL` in `root`. Data is scattered from the root to all processes in group B. The receive buffer arguments of the processes in group B must be consistent with the send buffer argument of the root.

ticket269.

```

MPI_SCATTERW(sendbuf, sendcount, displs, sendtypes, recvbuf, recvcount, recvtpe, root,
              comm)
    IN      sendbuf      starting address of send buffer (choice, significant only
                        at root)
    IN      sendcount    non-negative integer array (of length group size) speci-
                        fying the number of elements to send to each processor
    IN      displs       integer array (of length group size). Entry i specifies
                        the displacement relative to sendbuf from which to
                        take the outgoing data to process i
    IN      sendtypes     array of datatypes (of length group size). Entry j
                        specifies the type of data to send to process j (array
                        of handles)
    OUT     recvbuf      address of receive buffer (choice)
    IN      recvcount    number of elements in receive buffer (non-negative in-
                        teger)
    IN      recvtpe      data type of receive buffer elements (handle)
    IN      root         rank of sending process (integer)
    IN      comm         communicator (handle)

```

```

int MPI_Scatterw(void* sendbuf, int sendcounts[], int displs[],
                MPI_Datatype sendtypes[], void* recvbuf, int recvcount,
                MPI_Datatype recvtpe, int root, MPI_Comm comm)
MPI_SCATTERW(SENDBUF, SENDCOUNTS, DISPLS, SENDTYPES, RECVBUF, RECVCOUNT,
            RECVTYPE, ROOT, COMM, IERROR)
<type> SENDBUF(*), RECVBUF(*)
INTEGER DISPLS(*), SENDCOUNTS(*), SENDTYPES, RECVCOUNT, RECVTYPE, ROOT,
COMM, IERROR

```

MPI\_SCATTERW is the inverse operation to MPI\_GATHERW.

MPI\_SCATTERW extends the functionality of MPI\_SCATTERV by allowing varying datatype specifications for data sent to each process.

If comm is an intracommunicator, the outcome is as if the root executed n send operations,

```
MPI_Send(sendbuf + displs[i] · extent(sendtypes[i]), sendcounts[i], sendtypes[i], i, ...),
```

and each process executed a receive,

```
MPI_Recv(recvbuf, recvcount, recvtpe, i, ...).
```

The send buffer is ignored for all non-root processes.

The type signature implied by sendcount[i], sendtypes[i] at the root must be equal to the type signature implied by recvcount, recvtpe at process i (however, the type maps may be different). This implies that the amount of data sent must be equal to the amount of data received, pairwise between each process and the root. Distinct type maps between sender and receiver are still allowed.




Figure 5.9: The root process scatters sets of 100 ints to each process in the group.

All arguments to the function are significant on process **root**, while on other processes, only arguments **recvbuf**, **recvcount**, **recvtype**, **root**, and **comm** are significant. The arguments **root** and **comm** must have identical values on all processes.

The specification of counts, types, and displacements should not cause any location on the root to be read more than once.

The “in place” option for intracommunicators is specified by passing **MPI\_IN\_PLACE** as the value of **recvbuf** at the root. In such a case, **recvcount** and **recvtype** are ignored, and root “sends” no data to itself. The scattered vector is still assumed to contain **n** segments, where **n** is the group size; the **root**-th segment, which root should “send to itself”, is not moved.

If **comm** is an intercommunicator, then the call involves all processes in the intercommunicator, but with one group (group A) defining the root process. All processes in the other group (group B) pass the same value in argument **root**, which is the rank of the root in group A. The root passes the value **MPI\_ROOT** in **root**. All other processes in group A pass the value **MPI\_PROC\_NULL** in **root**. Data is scattered from the root to all processes in group B. The receive buffer arguments of the processes in group B must be consistent with the send buffer arguments of the root.

### 5.6.1 Examples using MPI\_SCATTER, MPI\_SCATTERV

The examples in this section use intracommunicators.

#### Example 5.11

The reverse of Example 5.2. Scatter sets of 100 ints from the root to each process in the group. See Figure 5.9.

```
MPI_Comm comm;
int gsize,*sendbuf;
int root, rbuf[100];
...
MPI_Comm_size(comm, &gsize);
sendbuf = (int *)malloc(gsize*100*sizeof(int));
...
MPI_Scatter(sendbuf, 100, MPI_INT, rbuf, 100, MPI_INT, root, comm);
```

#### Example 5.12




Figure 5.10: The root process scatters sets of 100 ints, moving by `stride` ints from send to send in the scatter.

The reverse of Example 5.5. The root process scatters sets of 100 ints to the other processes, but the sets of 100 are *stride* ints apart in the sending buffer. Requires use of `MPI_SCATTERV`. Assume *stride*  $\geq$  100. See Figure 5.10.

```

MPI_Comm comm;
int gsize,*sendbuf;
int root, rbuf[100], i, *displs, *scounts;

...

MPI_Comm_size(comm, &gsize);
sendbuf = (int *)malloc(gsize*stride*sizeof(int));
...
displs = (int *)malloc(gsize*sizeof(int));
scounts = (int *)malloc(gsize*sizeof(int));
for (i=0; i<gsize; ++i) {
    displs[i] = i*stride;
    scounts[i] = 100;
}
MPI_Scatterv(sendbuf, scounts, displs, MPI_INT, rbuf, 100, MPI_INT,
             root, comm);

```

### Example 5.13

The reverse of Example 5.9. We have a varying stride between blocks at sending (root) side, at the receiving side we receive into the *i*-th column of a 100×150 C array. See Figure 5.11.

```

MPI_Comm comm;
int gsize,recvarray[100][150],*rptra;
int root, *sendbuf, myrank, *stride;
MPI_Datatype rtype;
int i, *displs, *scounts, offset;

...
MPI_Comm_size(comm, &gsize);
MPI_Comm_rank(comm, &myrank);

```






Figure 5.11: The root scatters blocks of  $100-i$  ints into column  $i$  of a  $100 \times 150$  C array. At the sending side, the blocks are  $\text{stride}[i]$  ints apart.

```

stride = (int *)malloc(gsize*sizeof(int));
...
/* stride[i] for i = 0 to gsize-1 is set somehow
 * sendbuf comes from elsewhere
 */
...
displs = (int *)malloc(gsize*sizeof(int));
counts = (int *)malloc(gsize*sizeof(int));
offset = 0;
for (i=0; i<gsize; ++i) {
    displs[i] = offset;
    offset += stride[i];
    counts[i] = 100 - i;
}
/* Create datatype for the column we are receiving
 */
MPI_Type_vector(100-myrank, 1, 150, MPI_INT, &rtype);
MPI_Type_commit(&rtype);
rptr = &recvarray[0][myrank];
MPI_Scatterv(sendbuf, counts, displs, MPI_INT, rptr, 1, rtype,
             root, comm);

```

## 5.7 Gather-to-all

<code>MPI_ALLGATHER(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, comm)</code>			1
			2
			3
			4
IN	<code>sendbuf</code>	starting address of send buffer (choice)	5
IN	<code>sendcount</code>	number of elements in send buffer (non-negative integer)	6
			7
IN	<code>sendtype</code>	data type of send buffer elements (handle)	8
OUT	<code>recvbuf</code>	address of receive buffer (choice)	9
IN	<code>recvcount</code>	number of elements received from any process (non-negative integer)	10
			11
IN	<code>recvtype</code>	data type of receive buffer elements (handle)	12
IN	<code>comm</code>	communicator (handle)	13
			14
			15
			16
			17

```

int MPI_Allgather(void* sendbuf, int sendcount, MPI_Datatype sendtype,
                  void* recvbuf, int recvcount, MPI_Datatype recvtype,
                  MPI_Comm comm)
MPI_ALLGATHER(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, REVCOUNT, RECVTYPE,
              COMM, IERROR)
<type> SENDBUF(*), RECVBUF(*)
INTEGER SENDCOUNT, SENDTYPE, REVCOUNT, RECVTYPE, COMM, IERROR
{void MPI::Comm::Allgather(const void* sendbuf, int sendcount, const
                           MPI::Datatype& sendtype, void* recvbuf, int recvcount,
                           const MPI::Datatype& recvtype) const = 0(binding deprecated, see
                           Section 15.2) }
```

MPI\_ALLGATHER can be thought of as MPI\_GATHER, but where all processes receive the result, instead of just the root. The block of data sent from the *j*-th process is received by every process and placed in the *j*-th block of the buffer `recvbuf`.

The type signature associated with `sendcount`, `sendtype`, at a process must be equal to the type signature associated with `recvcount`, `recvtype` at any other process.

If `comm` is an intracommunicator, the outcome of a call to `MPI_ALLGATHER(...)` is as if all processes executed *n* calls to

```

MPI_Gather(sendbuf, sendcount, sendtype, recvbuf, recvcount,
           recvtype, root, comm)
```

for `root = 0, ..., n-1`. The rules for correct usage of `MPI_ALLGATHER` are easily found from the corresponding rules for `MPI_GATHER`.

The “in place” option for intracommunicators is specified by passing the value `MPI_IN_PLACE` to the argument `sendbuf` at all processes. `sendcount` and `sendtype` are ignored. Then the input data of each process is assumed to be in the area where that process would receive its own contribution to the receive buffer.

If `comm` is an intercommunicator, then each process of one group (group A) contributes `sendcount` data items; these data are concatenated and the result is stored at each process

in the other group (group B). Conversely the concatenation of the contributions of the processes in group B is stored at each process in group A. The send buffer arguments in group A must be consistent with the receive buffer arguments in group B, and vice versa.

*Advice to users.* The communication pattern of `MPI_ALLGATHER` executed on an intercommunication domain need not be symmetric. The number of items sent by processes in group A (as specified by the arguments `sendcount`, `sendtype` in group A and the arguments `recvcount`, `recvtype` in group B), need not equal the number of items sent by processes in group B (as specified by the arguments `sendcount`, `sendtype` in group B and the arguments `recvcount`, `recvtype` in group A). In particular, one can move data in only one direction by specifying `sendcount = 0` for the communication in the reverse direction.

*(End of advice to users.)*

`MPI_ALLGATHERV(sendbuf, sendcount, sendtype, recvbuf, recvcounts, displs, recvtype, comm)`

IN	sendbuf	starting address of send buffer (choice)
IN	sendcount	number of elements in send buffer (non-negative integer)
IN	sendtype	data type of send buffer elements (handle)
OUT	recvbuf	address of receive buffer (choice)
IN	recvcounts	non-negative integer array (of length group size) containing the number of elements that are received from each process
IN	displs	integer array (of length group size). Entry <code>i</code> specifies the displacement (relative to <code>recvbuf</code> ) at which to place the incoming data from process <code>i</code>
IN	recvtype	data type of receive buffer elements (handle)
IN	comm	communicator (handle)

```
int MPI_Allgatherv(void* sendbuf, int sendcount, MPI_Datatype sendtype,
                  void* recvbuf, int *recvcounts, int *displs,
                  MPI_Datatype recvtype, MPI_Comm comm)
```

```
MPI_ALLGATHERV(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNTS, DISPLS,
                RECVTYPE, COMM, IERROR)
```

```
<type> SENDBUF(*), RECVBUF(*)
INTEGER SENDCOUNT, SENDTYPE, RECVCOUNTS(*), DISPLS(*), RECVTYPE, COMM,
IERROR
```

```
{void MPI::Comm::Allgatherv(const void* sendbuf, int sendcount, const
                             MPI::Datatype& sendtype, void* recvbuf,
                             const int recvcounts[], const int displs[],
                             const MPI::Datatype& recvtype) const = 0(binding deprecated, see
                             Section 15.2) }
```

MPI\_ALLGATHERV can be thought of as MPI\_GATHERV, but where all processes receive the result, instead of just the root. The block of data sent from the  $j$ -th process is received by every process and placed in the  $j$ -th block of the buffer `recvbuf`. These blocks need not all be the same size.

The type signature associated with `sendcount`, `sendtype`, at process  $j$  must be equal to the type signature associated with `recvcounts[j]`, `recvtype` at any other process.

If `comm` is an intracommunicator, the outcome is as if all processes executed calls to

```
MPI_GATHERV(sendbuf, sendcount, sendtype, recvbuf, recvcounts, displs, recvtype, root, comm),
```

for `root = 0, ..., n-1`. The rules for correct usage of MPI\_ALLGATHERV are easily found from the corresponding rules for MPI\_GATHERV.

The “in place” option for intracommunicators is specified by passing the value `MPI_IN_PLACE` to the argument `sendbuf` at all processes. In such a case, `sendcount` and `sendtype` are ignored, and the input data of each process is assumed to be in the area where that process would receive its own contribution to the receive buffer.

If `comm` is an intercommunicator, then each process of one group (group A) contributes `sendcount` data items; these data are concatenated and the result is stored at each process in the other group (group B). Conversely the concatenation of the contributions of the processes in group B is stored at each process in group A. The send buffer arguments in group A must be consistent with the receive buffer arguments in group B, and vice versa.

ticket269.

```
MPI_ALLGATHERW(sendbuf, sendcount, sendtype, recvbuf, recvcounts, displs, recvtypes, comm)
```

IN	<code>sendbuf</code>	starting address of send buffer (choice)
IN	<code>sendcount</code>	number of elements in send buffer (non-negative integer)
IN	<code>sendtype</code>	data type of send buffer elements (handle)
OUT	<code>recvbuf</code>	address of receive buffer (choice)
IN	<code>recvcounts</code>	non-negative integer array (of length group size) containing the number of elements that are received from each process
IN	<code>displs</code>	integer array (of length group size). Entry $i$ specifies the displacement (relative to <code>recvbuf</code> ) at which to place the incoming data from process $i$
IN	<code>recvtypes</code>	array of datatypes (of length group size). Entry $i$ specifies the type of data received from process $i$ (array of handles)
IN	<code>comm</code>	communicator (handle)

```
int MPI_Allgatherw(void* sendbuf, int sendcount, MPI_Datatype sendtype,
                  void* recvbuf, int *recvcounts, int *displs,
                  MPI_Datatype *recvtypes, MPI_Comm comm)
```

```
MPI_ALLGATHERW(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNTS, DISPLS,
                RECVTYPES, COMM, IERROR)
```

```

1  <type> SENDBUF(*), RECVBUF(*)
2  INTEGER SENDCOUNT, SENDTYPE, RECVCOUNTS(*), DISPLS(*), RECVTYPES(*),
3  COMM, IERROR

```

`MPI_ALLGATHERW` can be thought of as `MPI_GATHERW`, but where all processes receive the result, instead of just the root. The block of data sent from the  $j$ -th process is received by every process and placed in the  $j$ -th block of the buffer `recvbuf`. These blocks need not all be the same size.

The type signature associated with `sendcount`, `sendtype`, at process  $j$  must be equal to the type signature associated with `recvcounts[j]`, `recvtypes[j]` at any other process.

If `comm` is an intracommunicator, the outcome is as if all processes executed calls to

```

12  MPI_GATHERW(sendbuf, sendcount, sendtype, recvbuf, recvcounts, displs, recvtypes, root, comm)

```

for `root = 0, ..., n-1`.

The “in place” option for intracommunicators is specified by passing the value `MPI_IN_PLACE` to the argument `sendbuf` at all processes. In such a case, `sendcount` and `sendtype` are ignored, and the input data of each process is assumed to be in the area where that process would receive its own contribution to the receive buffer.

If `comm` is an intercommunicator, then each process of one group (group A) contributes `sendcount` data items; these data are concatenated and the result is stored at each process in the other group (group B). Conversely the concatenation of the contributions of the processes in group B is stored at each process in group A. The send buffer arguments in group A must be consistent with the receive buffer arguments in group B, and vice versa.

### 5.7.1 Example using `MPI_ALLGATHER`

The example in this section uses intracommunicators.

#### Example 5.14

The all-gather version of Example 5.2. Using `MPI_ALLGATHER`, we will gather 100 `ints` from every process in the group to every process.

```

32  MPI_Comm comm;
33  int gsize, sendarray[100];
34  int *rbuf;
35  ...
36  MPI_Comm_size(comm, &gsize);
37  rbuf = (int *)malloc(gsize*100*sizeof(int));
38  MPI_Allgather(sendarray, 100, MPI_INT, rbuf, 100, MPI_INT, comm);

```

After the call, every process has the group-wide concatenation of the sets of data.

## 5.8 All-to-All Scatter/Gather

`MPI_ALLTOALL(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, comm)`

IN	<code>sendbuf</code>	starting address of send buffer (choice)
IN	<code>sendcount</code>	number of elements sent to each process (non-negative integer)
IN	<code>sendtype</code>	data type of send buffer elements (handle)
OUT	<code>recvbuf</code>	address of receive buffer (choice)
IN	<code>recvcount</code>	number of elements received from any process (non-negative integer)
IN	<code>recvtype</code>	data type of receive buffer elements (handle)
IN	<code>comm</code>	communicator (handle)

```
int MPI_Alltoall(void* sendbuf, int sendcount, MPI_Datatype sendtype,
                void* recvbuf, int recvcount, MPI_Datatype recvtype,
                MPI_Comm comm)
```

```
MPI_ALLTOALL(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNT, RECVTYPE,
              COMM, IERROR)
<type> SENDBUF(*), RECVBUF(*)
INTEGER SENDCOUNT, SENDTYPE, RECVCOUNT, RECVTYPE, COMM, IERROR
```

```
{void MPI::Comm::Alltoall(const void* sendbuf, int sendcount, const
                          MPI::Datatype& sendtype, void* recvbuf, int recvcount,
                          const MPI::Datatype& recvtype) const = 0(binding deprecated, see
                          Section 15.2) }
```

`MPI_ALLTOALL` is an extension of `MPI_ALLGATHER` to the case where each process sends distinct data to each of the receivers. The  $j$ -th block sent from process  $i$  is received by process  $j$  and is placed in the  $i$ -th block of `recvbuf`.

The type signature associated with `sendcount`, `sendtype`, at a process must be equal to the type signature associated with `recvcount`, `recvtype` at any other process. This implies that the amount of data sent must be equal to the amount of data received, pairwise between every pair of processes. As usual, however, the type maps may be different.

If `comm` is an intracommunicator, the outcome is as if each process executed a send to each process (itself included) with a call to,

```
MPI_Send(sendbuf + i · sendcount · extent(sendtype), sendcount, sendtype, i, ...),
```

and a receive from every other process with a call to,

```
MPI_Recv(recvbuf + i · recvcount · extent(recvtype), recvcount, recvtype, i, ...).
```

All arguments on all processes are significant. The argument `comm` must have identical values on all processes.

The “in place” option for intracommunicators is specified by passing `MPI_IN_PLACE` to the argument `sendbuf` at *all* processes. In such a case, `sendcount` and `sendtype` are ignored.

The data to be sent is taken from the `recvbuf` and replaced by the received data. Data sent and received must have the same type map as specified by `recvcount` and `recvtype`.

*Rationale.* For large `MPI_ALLTOALL` instances, allocating both send and receive buffers may consume too much memory. The “in place” option effectively halves the application memory consumption and is useful in situations where the data to be sent will not be used by the sending process after the `MPI_ALLTOALL` exchange (e.g., in parallel Fast Fourier Transforms). (*End of rationale.*)

*Advice to implementors.* Users may opt to use the “in place” option in order to conserve memory. Quality MPI implementations should thus strive to minimize system buffering. (*End of advice to implementors.*)

If `comm` is an intercommunicator, then the outcome is as if each process in group A sends a message to each process in group B, and vice versa. The  $j$ -th send buffer of process  $i$  in group A should be consistent with the  $i$ -th receive buffer of process  $j$  in group B, and vice versa.

*Advice to users.* When a complete exchange is executed on an intercommunication domain, then the number of data items sent from processes in group A to processes in group B need not equal the number of items sent in the reverse direction. In particular, one can have unidirectional communication by specifying `sendcount = 0` in the reverse direction.

(*End of advice to users.*)

`MPI_ALLTOALLV(sendbuf, sendcounts, sdispls, sendtype, recvbuf, recvcounts, rdispls, recvtype, comm)`

IN	<code>sendbuf</code>	starting address of send buffer (choice)
IN	<code>sendcounts</code>	non-negative integer array (of length group size) specifying the number of elements to send to each processor
IN	<code>sdispls</code>	integer array (of length group size). Entry $j$ specifies the displacement (relative to <code>sendbuf</code> ) from which to take the outgoing data destined for process $j$
IN	<code>sendtype</code>	data type of send buffer elements (handle)
OUT	<code>recvbuf</code>	address of receive buffer (choice)
IN	<code>recvcounts</code>	non-negative integer array (of length group size) specifying the number of elements that can be received from each processor
IN	<code>rdispls</code>	integer array (of length group size). Entry $i$ specifies the displacement (relative to <code>recvbuf</code> ) at which to place the incoming data from process $i$
IN	<code>recvtype</code>	data type of receive buffer elements (handle)
IN	<code>comm</code>	communicator (handle)

```
int MPI_Alltoallv(void* sendbuf, int *sendcounts, int *sdispls,
MPI_Datatype sendtype, void* recvbuf, int *recvcounts,
```

```

        int *rdispls, MPI_Datatype recvtype, MPI_Comm comm)
MPI_ALLTOALLV(SENDBUF, SENDCOUNTS, SDISPLS, SENDTYPE, RECVBUF, REVCOUNTS,
              RDISPLS, RECVTYPE, COMM, IERROR)
<type> SENDBUF(*), RECVBUF(*)
INTEGER SENDCOUNTS(*), SDISPLS(*), SENDTYPE, REVCOUNTS(*), RDISPLS(*),
RECVTYPE, COMM, IERROR
{void MPI::Comm::Alltoallv(const void* sendbuf, const int sendcounts[],
    const int sdispls[], const MPI::Datatype& sendtype,
    void* recvbuf, const int recvcounst[], const int rdispls[],
    const MPI::Datatype& recvtype) const = 0(binding deprecated, see
    Section 15.2) }
```

MPI\_ALLTOALLV adds flexibility to MPI\_ALLTOALL in that the location of data for the send is specified by `sdispls` and the location of the placement of the data on the receive side is specified by `rdispls`.

If `comm` is an intracommunicator, then the  $j$ -th block sent from process  $i$  is received by process  $j$  and is placed in the  $i$ -th block of `recvbuf`. These blocks need not all have the same size.

The type signature associated with `sendcounts[j]`, `sendtype` at process  $i$  must be equal to the type signature associated with `recvcounst[i]`, `recvtype` at process  $j$ . This implies that the amount of data sent must be equal to the amount of data received, pairwise between every pair of processes. Distinct type maps between sender and receiver are still allowed.

The outcome is as if each process sent a message to every other process with,

```
MPI_Send(sendbuf + sdispls[i] * extent(sendtype), sendcounts[i], sendtype, i, ...),
```

and received a message from every other process with a call to

```
MPI_Recv(recvbuf + rdispls[i] * extent(recvtype), recvcounst[i], recvtype, i, ...).
```

All arguments on all processes are significant. The argument `comm` must have identical values on all processes.

The “in place” option for intracommunicators is specified by passing `MPI_IN_PLACE` to the argument `sendbuf` at *all* processes. In such a case, `sendcounts`, `sdispls` and `sendtype` are ignored. The data to be sent is taken from the `recvbuf` and replaced by the received data. Data sent and received must have the same type map as specified by the `recvcounst` array and the `recvtype`, and is taken from the locations of the receive buffer specified by `rdispls`.

*Advice to users.* Specifying the “in place” option (which must be given on all processes) implies that the same amount and type of data is sent and received between any two processes in the group of the communicator. Different pairs of processes can exchange different amounts of data. Users must ensure that `recvcounst[j]` and `recvtype` on process  $i$  match `recvcounst[i]` and `recvtype` on process  $j$ . This symmetric exchange can be useful in applications where the data to be sent will not be used by the sending process after the MPI\_ALLTOALLV exchange. (*End of advice to users.*)

If `comm` is an intercommunicator, then the outcome is as if each process in group A sends a message to each process in group B, and vice versa. The  $j$ -th send buffer of process  $i$  in group A should be consistent with the  $i$ -th receive buffer of process  $j$  in group B, and vice versa.



*Rationale.* The definitions of MPI\_ALLTOALL and MPI\_ALLTOALLV give as much flexibility as one would achieve by specifying  $n$  independent, point-to-point communications, with two exceptions: all messages use the same datatype, and messages are scattered from (or gathered to) sequential storage. (*End of rationale.*)

*Advice to implementors.* Although the discussion of collective communication in terms of point-to-point operation implies that each message is transferred directly from sender to receiver, implementations may use a tree communication pattern. Messages can be forwarded by intermediate nodes where they are split (for scatter) or concatenated (for gather), if this is more efficient. (*End of advice to implementors.*)

MPI\_ALLTOALLW(sendbuf, sendcounts, sdispls, sendtypes, recvbuf, recvcoun-  
ts, rdispls, recvtypes, comm)

IN	sendbuf	starting address of send buffer (choice)
IN	sendcounts	non-negative integer array (of length group size) specifying the number of elements to send to each processor
IN	sdispls	integer array (of length group size). Entry $j$ specifies the displacement in bytes (relative to sendbuf) from which to take the outgoing data destined for process $j$ (array of integers)
IN	sendtypes	array of datatypes (of length group size). Entry $j$ specifies the type of data to send to process $j$ (array of handles)
OUT	recvbuf	address of receive buffer (choice)
IN	recvcoun- ts	non-negative integer array (of length group size) specifying the number of elements that can be received from each processor
IN	rdispls	integer array (of length group size). Entry $i$ specifies the displacement in bytes (relative to recvbuf) at which to place the incoming data from process $i$ (array of integers)
IN	recvtypes	array of datatypes (of length group size). Entry $i$ specifies the type of data received from process $i$ (array of handles)
IN	comm	communicator (handle)

```
int MPI_Alltoallw(void* sendbuf, int sendcounts[], int sdispls[],
    MPI_Datatype sendtypes[], void* recvbuf, int recvcoun-
    ts[], int rdispls[], MPI_Datatype recvtypes[], MPI_Comm comm)

MPI_ALLTOALLW(SENDBUF, SENDCOUNTS, SDISPLS, SENDTYPES, RECVBUF, RECVCOUNTS,
    RDISPLS, RECVTYPES, COMM, IERROR)
<type> SENDBUF(*), RECVBUF(*)
```

```

INTEGER SENDCOUNTS(*), SDISPLS(*), SENDTYPES(*), RECVCOUNTS(*),
RDISPLS(*), RECVTYPES(*), COMM, IERROR
{void MPI::Comm::Alltoallw(const void* sendbuf, const int sendcounts[],
    const int sdispls[], const MPI::Datatype sendtypes[], void*
    recvbuf, const int recvcounts[], const int rdispls[], const
    MPI::Datatype recvtypes[]) const = 0(binding deprecated, see
    Section 15.2) }
```

MPI\_ALLTOALLW is the most general form of complete exchange. Like MPI\_TYPE\_CREATE\_STRUCT, the most general type constructor, MPI\_ALLTOALLW allows separate specification of count, displacement and datatype. In addition, to allow maximum flexibility, the displacement of blocks within the send and receive buffers is specified in bytes.

If `comm` is an intracommunicator, then the  $j$ -th block sent from process  $i$  is received by process  $j$  and is placed in the  $i$ -th block of `recvbuf`. These blocks need not all have the same size.

The type signature associated with `sendcounts[j]`, `sendtypes[j]` at process  $i$  must be equal to the type signature associated with `recvcounts[i]`, `recvtypes[i]` at process  $j$ . This implies that the amount of data sent must be equal to the amount of data received, pairwise between every pair of processes. Distinct type maps between sender and receiver are still allowed.

The outcome is as if each process sent a message to every other process with

```
MPI_Send(sendbuf + sdispls[i], sendcounts[i], sendtypes[i], i, ...),
```

and received a message from every other process with a call to

```
MPI_Recv(recvbuf + rdispls[i], recvcounts[i], recvtypes[i], i, ...).
```

All arguments on all processes are significant. The argument `comm` must describe the same communicator on all processes.

Like for MPI\_ALLTOALLV, the “in place” option for intracommunicators is specified by passing MPI\_IN\_PLACE to the argument `sendbuf` at *all* processes. In such a case, `sendcounts`, `sdispls` and `sendtypes` are ignored. The data to be sent is taken from the `recvbuf` and replaced by the received data. Data sent and received must have the same type map as specified by the `recvcounts` and `recvtypes` arrays, and is taken from the locations of the receive buffer specified by `rdispls`.

If `comm` is an intercommunicator, then the outcome is as if each process in group A sends a message to each process in group B, and vice versa. The  $j$ -th send buffer of process  $i$  in group A should be consistent with the  $i$ -th receive buffer of process  $j$  in group B, and vice versa.

*Rationale.* The MPI\_ALLTOALLW function generalizes several MPI functions by carefully selecting the input arguments. For example, by making all but one process have `sendcounts[i] = 0`, this achieves an MPI\_SCATTERW function. (*End of rationale.*)

## 5.9 Global Reduction Operations

The functions in this section perform a global reduce operation (for example sum, maximum, and logical and) across all members of a group. The reduction operation can be either one of

a predefined list of operations, or a user-defined operation. The global reduction functions come in several flavors: a reduce that returns the result of the reduction to one member of a group, an all-reduce that returns this result to all members of a group, and two scan (parallel prefix) operations. In addition, a reduce-scatter operation combines the functionality of a reduce and of a scatter operation.

### 5.9.1 Reduce

`MPI_REDUCE(sendbuf, recvbuf, count, datatype, op, root, comm)`

IN	sendbuf	address of send buffer (choice)
OUT	recvbuf	address of receive buffer (choice, significant only at root)
IN	count	number of elements in send buffer (non-negative integer)
IN	datatype	data type of elements of send buffer (handle)
IN	op	reduce operation (handle)
IN	root	rank of root process (integer)
IN	comm	communicator (handle)

```

int MPI_Reduce(void* sendbuf, void* recvbuf, int count,
               MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm comm)
MPI_REDUCE(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, ROOT, COMM, IERROR)
<type> SENDBUF(*), RECVBUF(*)
INTEGER COUNT, DATATYPE, OP, ROOT, COMM, IERROR
{void MPI::Comm::Reduce(const void* sendbuf, void* recvbuf, int count,
                        const MPI::Datatype& datatype, const MPI::Op& op, int root)
const = 0(binding deprecated, see Section 15.2) }
```

If `comm` is an intracommunicator, `MPI_REDUCE` combines the elements provided in the input buffer of each process in the group, using the operation `op`, and returns the combined value in the output buffer of the process with rank `root`. The input buffer is defined by the arguments `sendbuf`, `count` and `datatype`; the output buffer is defined by the arguments `recvbuf`, `count` and `datatype`; both have the same number of elements, with the same type. The routine is called by all group members using the same arguments for `count`, `datatype`, `op`, `root` and `comm`. Thus, all processes provide input buffers and output buffers of the same length, with elements of the same type. Each process can provide one element, or a sequence of elements, in which case the combine operation is executed element-wise on each entry of the sequence. For example, if the operation is `MPI_MAX` and the send buffer contains two elements that are floating point numbers (`count = 2` and `datatype = MPI_FLOAT`), then `recvbuf(1) = global max(sendbuf(1))` and `recvbuf(2) = global max(sendbuf(2))`.

Section 5.9.2, lists the set of predefined operations provided by MPI. That section also enumerates the datatypes to which each operation can be applied.

In addition, users may define their own operations that can be overloaded to operate on several datatypes, either basic or derived. This is further explained in Section 5.9.5.

The operation `op` is always assumed to be associative. All predefined operations are also assumed to be commutative. Users may define operations that are assumed to be associative, but not commutative. The “canonical” evaluation order of a reduction is determined by the ranks of the processes in the group. However, the implementation can take advantage of associativity, or associativity and commutativity in order to change the order of evaluation. This may change the result of the reduction for operations that are not strictly associative and commutative, such as floating point addition.

*Advice to implementors.* It is strongly recommended that `MPI_REDUCE` be implemented so that the same result be obtained whenever the function is applied on the same arguments, appearing in the same order. Note that this may prevent optimizations that take advantage of the physical location of processors. (*End of advice to implementors.*)

*Advice to users.* Some applications may not be able to ignore the non-associative nature of floating-point operations or may use user-defined operations (see Section 5.9.5) that require a special reduction order and cannot be treated as associative. Such applications should enforce the order of evaluation explicitly. For example, in the case of operations that require a strict left-to-right (or right-to-left) evaluation order, this could be done by gathering all operands at a single process (e.g., with `MPI_GATHER`), applying the reduction operation in the desired order (e.g., with `MPI_REDUCE_LOCAL`), and if needed, broadcast or scatter the result to the other processes (e.g., with `MPI_BCAST`). (*End of advice to users.*)

The `datatype` argument of `MPI_REDUCE` must be compatible with `op`. Predefined operators work only with the MPI types listed in Section 5.9.2 and Section 5.9.4. Furthermore, the `datatype` and `op` given for predefined operators must be the same on all processes.

Note that it is possible for users to supply different user-defined operations to `MPI_REDUCE` in each process. MPI does not define which operations are used on which operands in this case. User-defined operators may operate on general, derived datatypes. In this case, each argument that the reduce operation is applied to is one element described by such a datatype, which may contain several basic values. This is further explained in Section 5.9.5.

*Advice to users.* Users should make no assumptions about how `MPI_REDUCE` is implemented. It is safest to ensure that the same function is passed to `MPI_REDUCE` by each process. (*End of advice to users.*)

Overlapping datatypes are permitted in “send” buffers. Overlapping datatypes in “receive” buffers are erroneous and may give unpredictable results.

The “in place” option for intracommunicators is specified by passing the value `MPI_IN_PLACE` to the argument `sendbuf` at the root. In such a case, the input data is taken at the root from the receive buffer, where it will be replaced by the output data.

If `comm` is an intercommunicator, then the call involves all processes in the intercommunicator, but with one group (group A) defining the root process. All processes in the other group (group B) pass the same value in argument `root`, which is the rank of the root in group A. The root passes the value `MPI_ROOT` in `root`. All other processes in group A pass the value `MPI_PROC_NULL` in `root`. Only send buffer arguments are significant in group B and only receive buffer arguments are significant at the root.

## 5.9.2 Predefined Reduction Operations

The following predefined operations are supplied for `MPI_REDUCE` and related functions `MPI_ALLREDUCE`, `MPI_REDUCE_SCATTER`, `MPI_SCAN`, and `MPI_EXSCAN`. These operations are invoked by placing the following in `op`.

Name	Meaning
<code>MPI_MAX</code>	maximum
<code>MPI_MIN</code>	minimum
<code>MPI_SUM</code>	sum
<code>MPI_PROD</code>	product
<code>MPI_LAND</code>	logical and
<code>MPI_BAND</code>	bit-wise and
<code>MPI_LOR</code>	logical or
<code>MPI BOR</code>	bit-wise or
<code>MPI_LXOR</code>	logical exclusive or (xor)
<code>MPI_BXOR</code>	bit-wise exclusive or (xor)
<code>MPI_MAXLOC</code>	max value and location
<code>MPI_MINLOC</code>	min value and location

The two operations `MPI_MINLOC` and `MPI_MAXLOC` are discussed separately in Section 5.9.4. For the other predefined operations, we enumerate below the allowed combinations of `op` and `datatype` arguments. First, define groups of MPI basic datatypes in the following way.

C integer:	<code>MPI_INT</code> , <code>MPI_LONG</code> , <code>MPI_SHORT</code> , <code>MPI_UNSIGNED_SHORT</code> , <code>MPI_UNSIGNED</code> , <code>MPI_UNSIGNED_LONG</code> , <code>MPI_LONG_LONG_INT</code> , <code>MPI_LONG_LONG</code> (as synonym), <code>MPI_UNSIGNED_LONG_LONG</code> , <code>MPI_SIGNED_CHAR</code> , <code>MPI_UNSIGNED_CHAR</code> , <code>MPI_INT8_T</code> , <code>MPI_INT16_T</code> , <code>MPI_INT32_T</code> , <code>MPI_INT64_T</code> , <code>MPI_UINT8_T</code> , <code>MPI_UINT16_T</code> , <code>MPI_UINT32_T</code> , <code>MPI_UINT64_T</code>
Fortran integer:	<code>MPI_INTEGER</code> , <code>MPI_AINT</code> , <code>MPI_COUNT</code> , <code>MPI_OFFSET</code> , and handles returned from <code>MPI_TYPE_CREATE_F90_INTEGER</code> , and if available: <code>MPI_INTEGER1</code> , <code>MPI_INTEGER2</code> , <code>MPI_INTEGER4</code> , <code>MPI_INTEGER8</code> , <code>MPI_INTEGER16</code>
Floating point:	<code>MPI_FLOAT</code> , <code>MPI_DOUBLE</code> , <code>MPI_REAL</code> , <code>MPI_DOUBLE_PRECISION</code> , <code>MPI_LONG_DOUBLE</code> and handles returned from

	MPI_TYPE_CREATE_F90_REAL,	1
	and if available: MPI_REAL2,	2
	MPI_REAL4, MPI_REAL8, MPI_REAL16	3
Logical:	MPI_LOGICAL, MPI_C_BOOL	4
Complex:	MPI_COMPLEX,	5
	MPI_C_FLOAT_COMPLEX,	6
	MPI_C_DOUBLE_COMPLEX,	7
	MPI_C_LONG_DOUBLE_COMPLEX,	8
	and handles returned from	9
	MPI_TYPE_CREATE_F90_COMPLEX,	10
	and if available: MPI_DOUBLE_COMPLEX,	11
	MPI_COMPLEX4, MPI_COMPLEX8,	12
	MPI_COMPLEX16, MPI_COMPLEX32	13
Byte:	MPI_BYTE	14

Now, the valid datatypes for each option is specified below.

Op	Allowed Types	
MPI_MAX, MPI_MIN	C integer, Fortran integer, Floating point	
MPI_SUM, MPI_PROD	C integer, Fortran integer, Floating point, Complex	
MPI_LAND, MPI_LOR, MPI_LXOR	C integer, Logical	
MPI_BAND, MPI_BOR, MPI_BXOR	C integer, Fortran integer, Byte	

The following examples use intracommunicators.

#### Example 5.15

A routine that computes the dot product of two vectors that are distributed across a group of processes and returns the answer at node zero.

```

SUBROUTINE PAR_BLAS1(m, a, b, c, comm)
REAL a(m), b(m)          ! local slice of array
REAL c                    ! result (at node zero)
REAL sum
INTEGER m, comm, i, ierr

! local sum
sum = 0.0
DO i = 1, m
    sum = sum + a(i)*b(i)
END DO

! global sum
CALL MPI_REDUCE(sum, c, 1, MPI_REAL, MPI_SUM, 0, comm, ierr)
RETURN

```

#### Example 5.16

A routine that computes the product of a vector and an array that are distributed across a group of processes and returns the answer at node zero.

```

1  SUBROUTINE PAR_BLAS2(m, n, a, b, c, comm)
2  REAL a(m), b(m,n)      ! local slice of array
3  REAL c(n)              ! result
4  REAL sum(n)
5  INTEGER n, comm, i, j, ierr
6
7  ! local sum
8  DO j= 1, n
9      sum(j) = 0.0
10     DO i = 1, m
11         sum(j) = sum(j) + a(i)*b(i,j)
12     END DO
13 END DO
14
15 ! global sum
16 CALL MPI_REDUCE(sum, c, n, MPI_REAL, MPI_SUM, 0, comm, ierr)
17
18 ! return result at node zero (and garbage at the other nodes)
19 RETURN
20

```

### 5.9.3 Signed Characters and Reductions

The types `MPI_SIGNED_CHAR` and `MPI_UNSIGNED_CHAR` can be used in reduction operations. `MPI_CHAR`, `MPI_WCHAR`, and `MPI_CHARACTER` (which represent printable characters) cannot be used in reduction operations. In a heterogeneous environment, `MPI_CHAR`, `MPI_WCHAR`, and `MPI_CHARACTER` will be translated so as to preserve the printable character, whereas `MPI_SIGNED_CHAR` and `MPI_UNSIGNED_CHAR` will be translated so as to preserve the integer value.

*Advice to users.* The types `MPI_CHAR`, `MPI_WCHAR`, and `MPI_CHARACTER` are intended for characters, and so will be translated to preserve the printable representation, rather than the integer value, if sent between machines with different character codes. The types `MPI_SIGNED_CHAR` and `MPI_UNSIGNED_CHAR` should be used in C if the integer value should be preserved. (*End of advice to users.*)

### 5.9.4 MINLOC and MAXLOC

The operator `MPI_MINLOC` is used to compute a global minimum and also an index attached to the minimum value. `MPI_MAXLOC` similarly computes a global maximum and index. One application of these is to compute a global minimum (maximum) and the rank of the process containing this value.

The operation that defines `MPI_MAXLOC` is:

$$\begin{pmatrix} u \\ i \end{pmatrix} \circ \begin{pmatrix} v \\ j \end{pmatrix} = \begin{pmatrix} w \\ k \end{pmatrix}$$

where

$$w = \max(u, v)$$

and

$$k = \begin{cases} i & \text{if } u > v \\ \min(i, j) & \text{if } u = v \\ j & \text{if } u < v \end{cases}$$

MPI\_MINLOC is defined similarly:

$$\begin{pmatrix} u \\ i \end{pmatrix} \circ \begin{pmatrix} v \\ j \end{pmatrix} = \begin{pmatrix} w \\ k \end{pmatrix}$$

where

$$w = \min(u, v)$$

and

$$k = \begin{cases} i & \text{if } u < v \\ \min(i, j) & \text{if } u = v \\ j & \text{if } u > v \end{cases}$$

Both operations are associative and commutative. Note that if MPI\_MAXLOC is applied to reduce a sequence of pairs  $(u_0, 0), (u_1, 1), \dots, (u_{n-1}, n-1)$ , then the value returned is  $(u, r)$ , where  $u = \max_i u_i$  and  $r$  is the index of the first global maximum in the sequence. Thus, if each process supplies a value and its rank within the group, then a reduce operation with `op = MPI_MAXLOC` will return the maximum value and the rank of the first process with that value. Similarly, MPI\_MINLOC can be used to return a minimum and its index. More generally, MPI\_MINLOC computes a *lexicographic minimum*, where elements are ordered according to the first component of each pair, and ties are resolved according to the second component.

The reduce operation is defined to operate on arguments that consist of a pair: value and index. For both Fortran and C, types are provided to describe the pair. The potentially mixed-type nature of such arguments is a problem in Fortran. The problem is circumvented, for Fortran, by having the MPI-provided type consist of a pair of the same type as value, and coercing the index to this type also. In C, the MPI-provided pair type has distinct types and the index is an `int`.

In order to use MPI\_MINLOC and MPI\_MAXLOC in a reduce operation, one must provide a `datatype` argument that represents a pair (value and index). MPI provides nine such predefined datatypes. The operations MPI\_MAXLOC and MPI\_MINLOC can be used with each of the following datatypes.

Fortran:

Name	Description
MPI_2REAL	pair of REALs
MPI_2DOUBLE_PRECISION	pair of DOUBLE PRECISION variables
MPI_2INTEGER	pair of INTEGERS

C:

Name	Description
MPI_FLOAT_INT	float and int



```

1      MPI_DOUBLE_INT          double and int
2      MPI_LONG_INT            long and int
3      MPI_2INT                pair of int
4      MPI_SHORT_INT           short and int
5      MPI_LONG_DOUBLE_INT     long double and int

```

6       The datatype MPI\_2REAL is *as if* defined by the following (see Section 4.1).

```

8      MPI_TYPE_CONTIGUOUS(2, MPI_REAL, MPI_2REAL)
9

```

10       Similar statements apply for MPI\_2INTEGER, MPI\_2DOUBLE\_PRECISION, and MPI\_2INT.

11       The datatype MPI\_FLOAT\_INT is *as if* defined by the following sequence of instructions.

```

12
13      type[0] = MPI_FLOAT
14      type[1] = MPI_INT
15      disp[0] = 0
16      disp[1] = sizeof(float)
17      block[0] = 1
18      block[1] = 1
19      MPI_TYPE_CREATE_STRUCT(2, block, disp, type, MPI_FLOAT_INT)
20

```

21       Similar statements apply for MPI\_LONG\_INT and MPI\_DOUBLE\_INT.

22       The following examples use intracommunicators.

### 23   **Example 5.17**

24       Each process has an array of 30 doubles, in C. For each of the 30 locations, compute  
25       the value and rank of the process containing the largest value.

```

26
27      ...
28      /* each process has an array of 30 double: ain[30]
29      */
30      double ain[30], aout[30];
31      int ind[30];
32      struct {
33          double val;
34          int rank;
35      } in[30], out[30];
36      int i, myrank, root;
37
38      MPI_Comm_rank(comm, &myrank);
39      for (i=0; i<30; ++i) {
40          in[i].val = ain[i];
41          in[i].rank = myrank;
42      }
43      MPI_Reduce(in, out, 30, MPI_DOUBLE_INT, MPI_MAXLOC, root, comm);
44      /* At this point, the answer resides on process root
45      */
46      if (myrank == root) {
47          /* read ranks out
48          */

```

```

        for (i=0; i<30; ++i) {
            aout[i] = out[i].val;
            ind[i] = out[i].rank;
        }
    }

```

**Example 5.18**

Same example, in Fortran.

```

...
! each process has an array of 30 double: ain(30)

DOUBLE PRECISION ain(30), aout(30)
INTEGER ind(30)
DOUBLE PRECISION in(2,30), out(2,30)
INTEGER i, myrank, root, ierr

CALL MPI_COMM_RANK(comm, myrank, ierr)
DO I=1, 30
    in(1,i) = ain(i)
    in(2,i) = myrank    ! myrank is coerced to a double
END DO

CALL MPI_REDUCE(in, out, 30, MPI_2DOUBLE_PRECISION, MPI_MAXLOC, root,
               comm, ierr)

! At this point, the answer resides on process root

IF (myrank .EQ. root) THEN
    ! read ranks out
    DO I= 1, 30
        aout(i) = out(1,i)
        ind(i) = out(2,i) ! rank is coerced back to an integer
    END DO
END IF

```

**Example 5.19**

Each process has a non-empty array of values. Find the minimum global value, the rank of the process that holds it and its index on this process.

```

#define LEN 1000

float val[LEN];          /* local array of values */
int count;               /* local number of values */
int myrank, minrank, minindex;
float minval;

struct {

```

```

1      float value;
2      int    index;
3  } in, out;
4
5      /* local minloc */
6  in.value = val[0];
7  in.index = 0;
8  for (i=1; i < count; i++)
9      if (in.value > val[i]) {
10         in.value = val[i];
11         in.index = i;
12     }
13
14     /* global minloc */
15 MPI_Comm_rank(comm, &myrank);
16 in.index = myrank*LEN + in.index;
17 MPI_Reduce( &in, &out, 1, MPI_FLOAT_INT, MPI_MINLOC, root, comm );
18     /* At this point, the answer resides on process root
19        */
20 if (myrank == root) {
21     /* read answer out
22        */
23     minval = out.value;
24     minrank = out.index / LEN;
25     minindex = out.index % LEN;
26 }

```

*Rationale.* The definition of MPI\_MINLOC and MPI\_MAXLOC given here has the advantage that it does not require any special-case handling of these two operations: they are handled like any other reduce operation. A programmer can provide his or her own definition of MPI\_MAXLOC and MPI\_MINLOC, if so desired. The disadvantage is that values and indices have to be first interleaved, and that indices and values have to be coerced to the same type, in Fortran. (*End of rationale.*)

### 5.9.5 User-Defined Reduction Operations

MPI\_OP\_CREATE(function, commute, op)

IN	function	user defined function (function)
IN	commute	true if commutative; false otherwise.
OUT	op	operation (handle)

```
int MPI_Op_create(MPI_User_function *function, int commute, MPI_Op *op)
```

```
MPI_OP_CREATE( FUNCTION, COMMUTE, OP, IERROR)
```

```
EXTERNAL FUNCTION
```

```
LOGICAL COMMUTE
```

INTEGER OP, IERROR

```
{void MPI::Op::Init(MPI::User_function *function, bool commute) (binding
    deprecated, see Section 15.2) }
```

MPI\_OP\_CREATE binds a user-defined reduction operation to an op handle that can subsequently be used in MPI\_REDUCE, MPI\_ALLREDUCE, MPI\_REDUCE\_SCATTER, MPI\_SCAN, and MPI\_EXSCAN. The user-defined operation is assumed to be associative. If `commute = true`, then the operation should be both commutative and associative. If `commute = false`, then the order of operands is fixed and is defined to be in ascending, process rank order, beginning with process zero. The order of evaluation can be changed, taking advantage of the associativity of the operation. If `commute = true` then the order of evaluation can be changed, taking advantage of commutativity and associativity.

The argument `function` is the user-defined function, which must have the following four arguments: `invec`, `inoutvec`, `len` and `datatype`.

The ISO C prototype for the function is the following.

```
typedef void MPI_User_function(void* invec, void* inoutvec, int *len,
    MPI_Datatype *datatype);
```

The Fortran declaration of the user-defined function appears below.

```
SUBROUTINE USER_FUNCTION(INVEC, INOUTVEC, LEN, TYPE)
    <type> INVEC(LEN), INOUTVEC(LEN)
    INTEGER LEN, TYPE
```

The C++ declaration of the user-defined function appears below.

```
{typedef void MPI::User_function(const void* invec, void* inoutvec, int
    len, const Datatype& datatype); (binding deprecated, see
    Section 15.2)}
```

The `datatype` argument is a handle to the data type that was passed into the call to MPI\_REDUCE. The user reduce function should be written such that the following holds: Let  $u[0], \dots, u[\text{len}-1]$  be the `len` elements in the communication buffer described by the arguments `invec`, `len` and `datatype` when the function is invoked; let  $v[0], \dots, v[\text{len}-1]$  be `len` elements in the communication buffer described by the arguments `inoutvec`, `len` and `datatype` when the function is invoked; let  $w[0], \dots, w[\text{len}-1]$  be `len` elements in the communication buffer described by the arguments `inoutvec`, `len` and `datatype` when the function returns; then  $w[i] = u[i] \circ v[i]$ , for  $i=0, \dots, \text{len}-1$ , where  $\circ$  is the reduce operation that the function computes.

Informally, we can think of `invec` and `inoutvec` as arrays of `len` elements that `function` is combining. The result of the reduction over-writes values in `inoutvec`, hence the name. Each invocation of the function results in the pointwise evaluation of the reduce operator on `len` elements: i.e., the function returns in `inoutvec[i]` the value `invec[i]  $\circ$  inoutvec[i]`, for  $i = 0, \dots, \text{count} - 1$ , where  $\circ$  is the combining operation computed by the function.

*Rationale.* The `len` argument allows MPI\_REDUCE to avoid calling the function for each element in the input buffer. Rather, the system can choose to apply the function to chunks of input. In C, it is passed in as a reference for reasons of compatibility with Fortran.

By internally comparing the value of the `datatype` argument to known, global handles, it is possible to overload the use of a single user-defined function for several, different

data types. (*End of rationale.*)

General datatypes may be passed to the user function. However, use of datatypes that are not contiguous is likely to lead to inefficiencies.

No MPI communication function may be called inside the user function. MPI\_ABORT may be called inside the function in case of an error.

*Advice to users.* Suppose one defines a library of user-defined reduce functions that are overloaded: the `datatype` argument is used to select the right execution path at each invocation, according to the types of the operands. The user-defined reduce function cannot “decode” the `datatype` argument that it is passed, and cannot identify, by itself, the correspondence between the datatype handles and the datatype they represent. This correspondence was established when the datatypes were created. Before the library is used, a library initialization preamble must be executed. This preamble code will define the datatypes that are used by the library, and store handles to these datatypes in global, static variables that are shared by the user code and the library code.

The Fortran version of MPI\_REDUCE will invoke a user-defined reduce function using the Fortran calling conventions and will pass a Fortran-type datatype argument; the C version will use C calling convention and the C representation of a datatype handle. Users who plan to mix languages should define their reduction functions accordingly. (*End of advice to users.*)

*Advice to implementors.* We outline below a naive and inefficient implementation of MPI\_REDUCE not supporting the “in place” option.

```

MPI_Comm_size(comm, &groupsize);
MPI_Comm_rank(comm, &rank);
if (rank > 0) {
    MPI_Recv(tempbuf, count, datatype, rank-1,...);
    User_reduce(tempbuf, sendbuf, count, datatype);
}
if (rank < groupsize-1) {
    MPI_Send(sendbuf, count, datatype, rank+1, ...);
}
/* answer now resides in process groupsize-1 ... now send to root
*/
if (rank == root) {
    MPI_Irecv(recvbuf, count, datatype, groupsize-1,..., &req);
}
if (rank == groupsize-1) {
    MPI_Send(sendbuf, count, datatype, root, ...);
}
if (rank == root) {
    MPI_Wait(&req, &status);
}

```

The reduction computation proceeds, sequentially, from process 0 to process `groupsize-1`. This order is chosen so as to respect the order of a possibly non-commutative operator defined by the function `User_reduce()`. A more efficient implementation is achieved by taking advantage of associativity and using a logarithmic tree reduction. Commutativity can be used to advantage, for those cases in which the `commute` argument to `MPI_OP_CREATE` is true. Also, the amount of temporary buffer required can be reduced, and communication can be pipelined with computation, by transferring and reducing the elements in chunks of size `len < count`.

The predefined reduce operations can be implemented as a library of user-defined operations. However, better performance might be achieved if MPI\_REDUCE handles these functions as a special case. (*End of advice to implementors.*)

MPI\_OP\_FREE(op)

INOUT	op	operation (handle)
-------	----	--------------------

```
int MPI_op_free(MPI_Op *op)
```

MPI\_OP\_FREE(OP, IERROR)

INTEGER OP, IERROR

```
{void MPI::Op::Free() (binding deprecated, see Section 15.2)}
```

Marks a user-defined reduction operation for deallocation and sets `op` to `MPI_OP_NULL`.

## Example of User-defined Reduce

It is time for an example of user-defined reduction. The example in this section uses an intracommunicator.

**Example 5.20** Compute the product of an array of complex numbers, in C.

```
typedef struct {
    double real,imag;
} Complex;
```

```
/* the user-defined function
*/
```

```
void myProd(Complex *in, Complex *inout, int *len, MPI_Datatype *dptr)
{
```

```
int i;  
Complex c;
```

```
for (i=0; i< *len; ++i) {
    c.real = inout->real*in->real -
              inout->imag*in->imag;
    c.imag = inout->real*in->imag +
              inout->imag*in->real;
    *inout = c;
}
```

```

1      in++; inout++;
2  }
3  }
4
5  /* and, to call it...
6  */
7  ...
8
9      /* each process has an array of 100 Complexes
10     */
11     Complex a[100], answer[100];
12     MPI_Op myOp;
13     MPI_Datatype ctype;
14
15     /* explain to MPI how type Complex is defined
16     */
17     MPI_Type_contiguous(2, MPI_DOUBLE, &ctype);
18     MPI_Type_commit(&ctype);
19     /* create the complex-product user-op
20     */
21     MPI_Op_create( myProd, 1, &myOp );
22
23     MPI_Reduce(a, answer, 100, ctype, myOp, root, comm);
24
25     /* At this point, the answer, which consists of 100 Complexes,
26     * resides on process root
27     */
28

```

### 5.9.6 All-Reduce

MPI includes a variant of the reduce operations where the result is returned to all processes in a group. MPI requires that all processes from the same group participating in these operations receive identical results.

**MPI\_ALLREDUCE**(sendbuf, recvbuf, count, datatype, op, comm)

IN	sendbuf	starting address of send buffer (choice)
OUT	recvbuf	starting address of receive buffer (choice)
IN	count	number of elements in send buffer (non-negative integer)
IN	datatype	data type of elements of send buffer (handle)
IN	op	operation (handle)
IN	comm	communicator (handle)

```

int MPI_Allreduce(void* sendbuf, void* recvbuf, int count,
                 MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)

```

```

MPI_ALLREDUCE(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, COMM, IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER COUNT, DATATYPE, OP, COMM, IERROR
{void MPI::Comm::Allreduce(const void* sendbuf, void* recvbuf, int count,
    const MPI::Datatype& datatype, const MPI::Op& op)
    const = 0(binding deprecated, see Section 15.2) }

```

If `comm` is an intracommunicator, `MPI_ALLREDUCE` behaves the same as `MPI_REDUCE` except that the result appears in the receive buffer of all the group members.

*Advice to implementors.* The all-reduce operations can be implemented as a reduce, followed by a broadcast. However, a direct implementation can lead to better performance. (*End of advice to implementors.*)

The “in place” option for intracommunicators is specified by passing the value `MPI_IN_PLACE` to the argument `sendbuf` at all processes. In this case, the input data is taken at each process from the receive buffer, where it will be replaced by the output data.

If `comm` is an intercommunicator, then the result of the reduction of the data provided by processes in group A is stored at each process in group B, and vice versa. Both groups should provide `count` and `datatype` arguments that specify the same type signature.

The following example uses an intracommunicator.

#### Example 5.21

A routine that computes the product of a vector and an array that are distributed across a group of processes and returns the answer at all nodes (see also Example 5.16).

```

SUBROUTINE PAR_BLAS2(m, n, a, b, c, comm)
REAL a(m), b(m,n)    ! local slice of array
REAL c(n)            ! result
REAL sum(n)
INTEGER n, comm, i, j, ierr

! local sum
DO j= 1, n
    sum(j) = 0.0
    DO i = 1, m
        sum(j) = sum(j) + a(i)*b(i,j)
    END DO
END DO

! global sum
CALL MPI_ALLREDUCE(sum, c, n, MPI_REAL, MPI_SUM, comm, ierr)

! return result at all nodes
RETURN

```

#### 5.9.7 Process-[l]Local [r]Reduction

The functions in this section are of importance to library implementors who may want to implement special reduction patterns that are otherwise not easily covered by the standard

ticket0.  
ticket0.



MPI operations.

The following function applies a reduction operator to local arguments.

**MPI\_REDUCE\_LOCAL**( inbuf, inoutbuf, count, datatype, op)

IN	inbuf	input buffer (choice)
INOUT	inoutbuf	combined input and output buffer (choice)
IN	count	number of elements in inbuf and inoutbuf buffers (non-negative integer)
IN	datatype	data type of elements of inbuf and inoutbuf buffers (handle)
IN	op	operation (handle)

```
int MPI_Reduce_local(void* inbuf, void* inoutbuf, int count,
                    MPI_Datatype datatype, MPI_Op op)
```

```
MPI_REDUCE_LOCAL(INBUF, INOUTBUF, COUNT, DATATYPE, OP, IERROR)
    <type> INBUF(*), INOUTBUF(*)
    INTEGER COUNT, DATATYPE, OP, IERROR
```

```
{void MPI::Op::Reduce_local(const void* inbuf, void* inoutbuf, int count,
    const MPI::Datatype& datatype) const(binding deprecated, see
    Section 15.2) }
```

The function applies the operation given by **op** element-wise to the elements of **inbuf** and **inoutbuf** with the result stored element-wise in **inoutbuf**, as explained for user-defined operations in Section 5.9.5. Both **inbuf** and **inoutbuf** (input as well as result) have the same number of elements given by **count** and the same datatype given by **datatype**. The **MPI\_IN\_PLACE** option is not allowed.

Reduction operations can be queried for their commutativity.

**MPI\_OP\_COMMUTATIVE**( op, commute)

IN	op	operation (handle)
OUT	commute	true if op is commutative, false otherwise (logical)

```
int MPI_Op_commutative(MPI_Op op, int *commute)
```

```
MPI_OP_COMMUTATIVE(OP, COMMUTE, IERROR)
    LOGICAL COMMUTE
    INTEGER OP, IERROR
```

```
{bool MPI::Op::Is_commutative() const(binding deprecated, see Section 15.2) }
```

## 5.10 Reduce-Scatter

MPI includes variants of the reduce operations where the result is scattered to all processes in a group on return. One variant scatters equal-sized blocks to all processes, while another variant scatters blocks that may vary in size for each process.

### 5.10.1 MPI\_REDUCE\_SCATTER\_BLOCK

```

MPI_REDUCE_SCATTER_BLOCK( sendbuf, recvbuf, recvcnt, datatype, op, comm)

IN      sendbuf      starting address of send buffer (choice)
OUT     recvbuf      starting address of receive buffer (choice)
IN      recvcnt      element count per block (non-negative integer)
IN      datatype     data type of elements of send and receive buffers (handle)
IN      op           operation (handle)
IN      comm         communicator (handle)

int MPI_Reduce_scatter_block(void* sendbuf, void* recvbuf, int recvcnt,
                             MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)

MPI_REDUCE_SCATTER_BLOCK(SENDBUF, RECVBUF, RECVCOUNT, DATATYPE, OP, COMM,
                          IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER RECVCOUNT, DATATYPE, OP, COMM, IERROR

{void MPI::Comm::Reduce_scatter_block(const void* sendbuf, void* recvbuf,
    int recvcnt, const MPI::Datatype& datatype,
    const MPI::Op& op) const = 0(binding deprecated, see Section 15.2) }
```

If `comm` is an intracommunicator, `MPI_REDUCE_SCATTER_BLOCK` first performs a global, element-wise reduction on vectors of `count = n*recvcnt` elements in the send buffers defined by `sendbuf`, `count` and `datatype`, using the operation `op`, where `n` is the number of processes in the group of `comm`. The routine is called by all group members using the same arguments for `recvcnt`, `datatype`, `op` and `comm`. The resulting vector is treated as `n` consecutive blocks of `recvcnt` elements that are scattered to the processes of the group. The `i`-th block is sent to process `i` and stored in the receive buffer defined by `recvbuf`, `recvcnt`, and `datatype`.

*Advice to implementors.* The `MPI_REDUCE_SCATTER_BLOCK` routine is functionally equivalent to: an `MPI_REDUCE` collective operation with `count` equal to `recvcnt*n`, followed by an `MPI_SCATTER` with `sendcount` equal to `recvcnt`. However, a direct implementation may run faster. (*End of advice to implementors.*)

The “in place” option for intracommunicators is specified by passing `MPI_IN_PLACE` in the `sendbuf` argument on *all* processes. In this case, the input data is taken from the receive buffer.

If `comm` is an intercommunicator, then the result of the reduction of the data provided by processes in one group (group A) is scattered among processes in the other group (group B) and vice versa. Within each group, all processes provide the same value for the `recvcount` argument, and provide input vectors of `count = n*recvcount` elements stored in the send buffers, where `n` is the size of the group. The number of elements `count` must be the same for the two groups. The resulting vector from the other group is scattered in blocks of `recvcount` elements among the processes in the group.

*Rationale.* The last restriction is needed so that the length of the send buffer of one group can be determined by the local `recvcount` argument of the other group. Otherwise, a communication is needed to figure out how many elements are reduced.  
(*End of rationale.*)

### 5.10.2 MPI\_REDUCE\_SCATTER

`MPI_REDUCE_SCATTER` extends the functionality of `MPI_REDUCE_SCATTER_BLOCK` such that the scattered blocks can vary in size. Block sizes are determined by the `recvcounts` array, such that the `i`-th block contains `recvcounts[i]` elements.

`MPI_REDUCE_SCATTER( sendbuf, recvbuf, recvcounts, datatype, op, comm)`

IN	<code>sendbuf</code>	starting address of send buffer (choice)
OUT	<code>recvbuf</code>	starting address of receive buffer (choice)
IN	<code>recvcounts</code>	non-negative integer array (of length group size) specifying the number of elements of the result distributed to each process.
IN	<code>datatype</code>	data type of elements of send and receive buffers (handle)
IN	<code>op</code>	operation (handle)
IN	<code>comm</code>	communicator (handle)

```
int MPI_Reduce_scatter(void* sendbuf, void* recvbuf, int *recvcounts,
    MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
```

```
MPI_REDUCE_SCATTER(SENDBUF, RECVBUF, RECVCOUNTS, DATATYPE, OP, COMM,
    IERROR)
```

```
<type> SENDBUF(*), RECVBUF(*)
```

```
INTEGER RECVCOUNTS(*), DATATYPE, OP, COMM, IERROR
```

```
{void MPI::Comm::Reduce_scatter(const void* sendbuf, void* recvbuf,
    int recvcounts[], const MPI::Datatype& datatype,
    const MPI::Op& op) const = 0(binding deprecated, see Section 15.2) }
```

If `comm` is an intracommunicator, `MPI_REDUCE_SCATTER` first performs a global, element-wise reduction on vectors of `count =  $\sum_{i=0}^{n-1} \text{recvcounts}[i]$`  elements in the send buffers defined by `sendbuf`, `count` and `datatype`, using the operation `op`, where `n` is the number of processes in the group of `comm`. The routine is called by all group members using the same arguments for `recvcounts`, `datatype`, `op` and `comm`. The resulting vector is treated as

n consecutive blocks where the number of elements of the i-th block is `recvcounts[i]`. The blocks are scattered to the processes of the group. The i-th block is sent to process i and stored in the receive buffer defined by `recvbuf`, `recvcounts[i]` and `datatype`.

*Advice to implementors.* The `MPI_REDUCE_SCATTER` routine is functionally equivalent to: an `MPI_REDUCE` collective operation with count equal to the sum of `recvcounts[i]` followed by `MPI_SCATTERV` with `sendcounts` equal to `recvcounts`. However, a direct implementation may run faster. (*End of advice to implementors.*)

The “in place” option for intracommunicators is specified by passing `MPI_IN_PLACE` in the `sendbuf` argument. In this case, the input data is taken from the receive buffer. It is not required to specify the “in place” option on all processes, since the processes for which `recvcounts[i]==0` may not have allocated a receive buffer.

If `comm` is an intercommunicator, then the result of the reduction of the data provided by processes in one group (group A) is scattered among processes in the other group (group B), and vice versa. Within each group, all processes provide the same `recvcounts` argument, and provide input vectors of `count =  $\sum_{i=0}^{n-1} \text{recvcounts}[i]$`  elements stored in the send buffers, where `n` is the size of the group. The resulting vector from the other group is scattered in blocks of `recvcounts[i]` elements among the processes in the group. The number of elements `count` must be the same for the two groups.

*Rationale.* The last restriction is needed so that the length of the send buffer can be determined by the sum of the local `recvcounts` entries. Otherwise, a communication is needed to figure out how many elements are reduced. (*End of rationale.*)

## 5.11 Scan

### 5.11.1 Inclusive Scan

`MPI_SCAN(sendbuf, recvbuf, count, datatype, op, comm)`

IN	<code>sendbuf</code>	starting address of send buffer (choice)
OUT	<code>recvbuf</code>	starting address of receive buffer (choice)
IN	<code>count</code>	number of elements in input buffer (non-negative integer)
IN	<code>datatype</code>	data type of elements of input buffer (handle)
IN	<code>op</code>	operation (handle)
IN	<code>comm</code>	communicator (handle)

```
int MPI_Scan(void* sendbuf, void* recvbuf, int count,
             MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
```

```
MPI_SCAN(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, COMM, IERROR)
<type> SENDBUF(*), RECVBUF(*)
INTEGER COUNT, DATATYPE, OP, COMM, IERROR
```

```

1 {void MPI::Intracomm::Scan(const void* sendbuf, void* recvbuf, int count,
2     const MPI::Datatype& datatype, const MPI::Op& op) const(binding
3     deprecated, see Section 15.2) }

```

If `comm` is an intracommunicator, `MPI_SCAN` is used to perform a prefix reduction on data distributed across the group. The operation returns, in the receive buffer of the process with rank `i`, the reduction of the values in the send buffers of processes with ranks `0, ..., i` (inclusive). The type of operations supported, their semantics, and the constraints on send and receive buffers are as for `MPI_REDUCE`.

The “in place” option for intracommunicators is specified by passing `MPI_IN_PLACE` in the `sendbuf` argument. In this case, the input data is taken from the receive buffer, and replaced by the output data.

This operation is invalid for intercommunicators.

### 5.11.2 Exclusive Scan

```

18 MPI_EXSCAN(sendbuf, recvbuf, count, datatype, op, comm)

```

19	IN	sendbuf	starting address of send buffer (choice)
20	OUT	recvbuf	starting address of receive buffer (choice)
21			
22	IN	count	number of elements in input buffer (non-negative integer)
23			
24	IN	datatype	data type of elements of input buffer (handle)
25			
26	IN	op	operation (handle)
27	IN	comm	intracommunicator (handle)

```

28
29 int MPI_Exscan(void* sendbuf, void* recvbuf, int count,
30     MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
31
32 MPI_EXSCAN(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, COMM, IERROR)
33     <type> SENDBUF(*), RECVBUF(*)
34     INTEGER COUNT, DATATYPE, OP, COMM, IERROR

```

```

35 {void MPI::Intracomm::Exscan(const void* sendbuf, void* recvbuf, int count,
36     const MPI::Datatype& datatype, const MPI::Op& op) const(binding
37     deprecated, see Section 15.2) }

```

If `comm` is an intracommunicator, `MPI_EXSCAN` is used to perform a prefix reduction on data distributed across the group. The value in `recvbuf` on the process with rank 0 is undefined, and `recvbuf` is not significant on process 0. The value in `recvbuf` on the process with rank 1 is defined as the value in `sendbuf` on the process with rank 0. For processes with rank `i > 1`, the operation returns, in the receive buffer of the process with rank `i`, the reduction of the values in the send buffers of processes with ranks `0, ..., i - 1` (inclusive). The type of operations supported, their semantics, and the constraints on send and receive buffers, are as for `MPI_REDUCE`.

The “in place” option for intracommunicators is specified by passing `MPI_IN_PLACE` in the `sendbuf` argument. In this case, the input data is taken from the receive buffer, and

replaced by the output data. The receive buffer on rank 0 is not changed by this operation.  
This operation is invalid for intercommunicators.

*Rationale.* The exclusive scan is more general than the inclusive scan. Any inclusive scan operation can be achieved by using the exclusive scan and then locally combining the local contribution. Note that for non-invertable operations such as MPI\_MAX, the exclusive scan cannot be computed with the inclusive scan. (*End of rationale.*)

### 5.11.3 Example using MPI\_SCAN

The example in this section uses an intracommunicator.

#### Example 5.22

This example uses a user-defined operation to produce a *segmented scan*. A segmented scan takes, as input, a set of values and a set of logicals, and the logicals delineate the various segments of the scan. For example:

<i>values</i>	$v_1$	$v_2$	$v_3$	$v_4$	$v_5$	$v_6$	$v_7$	$v_8$
<i>logicals</i>	0	0	1	1	1	0	0	1
<i>result</i>	$v_1$	$v_1 + v_2$	$v_3$	$v_3 + v_4$	$v_3 + v_4 + v_5$	$v_6$	$v_6 + v_7$	$v_8$

The operator that produces this effect is,

$$\begin{pmatrix} u \\ i \end{pmatrix} \circ \begin{pmatrix} v \\ j \end{pmatrix} = \begin{pmatrix} w \\ j \end{pmatrix},$$

where,

$$w = \begin{cases} u + v & \text{if } i = j \\ v & \text{if } i \neq j \end{cases}.$$

Note that this is a non-commutative operator. C code that implements it is given below.

```
typedef struct {
    double val;
    int log;
} SegScanPair;

/* the user-defined function
*/
void segScan(SegScanPair *in, SegScanPair *inout, int *len,
             MPI_Datatype *dptr)
{
    int i;
    SegScanPair c;

    for (i=0; i< *len; ++i) {
        if (in->log == inout->log)
            c.val = in->val + inout->val;
```

```

1      else
2          c.val = inout->val;
3      c.log = inout->log;
4      *inout = c;
5      in++; inout++;
6  }
7  }

```

Note that the `inout` argument to the user-defined function corresponds to the right-hand operand of the operator. When using this operator, we must be careful to specify that it is non-commutative, as in the following.

```

12      int i,base;
13      SegScanPair  a, answer;
14      MPI_Op      myOp;
15      MPI_Datatype type[2] = {MPI_DOUBLE, MPI_INT};
16      MPI_Aint     disp[2];
17      int          blocklen[2] = { 1, 1};
18      MPI_Datatype sspair;
19
20      /* explain to MPI how type SegScanPair is defined
21      */
22      MPI_Get_address( a, disp);
23      MPI_Get_address( a.log, disp+1);
24      base = disp[0];
25      for (i=0; i<2; ++i) disp[i] -= base;
26      MPI_Type_create_struct( 2, blocklen, disp, type, &sspair );
27      MPI_Type_commit( &sspair );
28      /* create the segmented-scan user-op
29      */
30      MPI_Op_create(segScan, 0, &myOp);
31      ...
32      MPI_Scan( &a, &answer, 1, sspair, myOp, comm );

```

ticket109.

## 5.12 Nonblocking Collective Operations

As described in Section 3.7, performance of many applications can be improved by overlapping communication and computation, and many systems enable this. Nonblocking collective operations combine the potential benefits of nonblocking point-to-point operations, to exploit overlap and to avoid synchronization, with the optimized implementation and message scheduling provided by collective operations [27, 30]. One way of doing this would be to perform a blocking collective operation in a separate thread. An alternative mechanism that often leads to better performance (e.g., avoids context switching, scheduler overheads, and thread management) is to use nonblocking collective communication [28].

The nonblocking collective communication model is similar to the model used for non-blocking point-to-point communication. A nonblocking call initiates a collective operation, which must be completed in a separate completion call. Once initiated, the operation

may progress independently of any computation or other communication at participating processes. In this manner, nonblocking collective operations can mitigate possible synchronizing effects of collective operations by running them in the “background.” In addition to enabling communication-computation overlap, nonblocking collective operations can perform collective operations on overlapping communicators, which would lead to deadlocks with blocking operations. Their semantic advantages can also be useful in combination with point-to-point communication.

As in the nonblocking point-to-point case, all calls are local and return immediately, irrespective of the status of other processes. The call initiates the operation, which indicates that the system may start to copy data out of the send buffer and into the receive buffer. Once initiated, all associated send buffers and buffers associated with input arguments (such as arrays of counts, displacements, or datatypes in the vector versions of the collectives) should not be modified, and all associated receive buffers should not be accessed, until the collective operation completes. The call returns a request handle, which must be passed to a completion call.

All completion calls (e.g., `MPI_WAIT`) described in Section 3.7.3 are supported for nonblocking collective operations. Similarly to the blocking case, nonblocking collective operations are considered to be complete when the local part of the operation is finished, i.e., for the caller, the semantics of the operation are guaranteed and all buffers can be safely accessed and modified. Completion does not indicate that other processes have completed or even started the operation (unless otherwise implied by the description of the operation). Completion of a particular nonblocking collective operation also does not indicate completion of any other posted nonblocking collective (or send-receive) operations, whether they are posted before or after the completed operation.

*Advice to users.* Users should be aware that implementations are allowed, but not required (with exception of `MPI_IBARRIER`), to synchronize processes during the completion of a nonblocking collective operation. (*End of advice to users.*)

Upon returning from a completion call in which a nonblocking collective operation completes, the `MPI_ERROR` field in the associated status object is set appropriately, see Section 3.2.5 on page 33. The values of the `MPI_SOURCE` and `MPI_TAG` fields are undefined. It is valid to mix different request types (i.e., any combination of collective requests, I/O requests, generalized requests, or point-to-point requests) in functions that enable multiple completions (e.g., `MPI_WAITALL`). It is erroneous to call `MPI_REQUEST_FREE` or `MPI_CANCEL` for a request associated with a nonblocking collective operation. Nonblocking collective requests are not persistent.

*Rationale.* Freeing an active nonblocking collective request could cause similar problems as discussed for point-to-point requests (see Section 3.7.3). Cancelling a request is not supported because the semantics of this operation are not well-defined. (*End of rationale.*)

Multiple nonblocking collective operations can be outstanding on a single communicator. If the nonblocking call causes some system resource to be exhausted, then it may fail and generate an MPI exception. Quality implementations of MPI should ensure that this happens only in pathological cases. That is, an MPI implementation should be able to support a large number of pending nonblocking operations.



Unlike point-to-point operations, nonblocking collective operations do not match with blocking collective operations, and collective operations do not have a tag argument. All processes must call collective operations (blocking and nonblocking) in the same order per communicator. In particular, once a process calls a collective operation, all other processes in the communicator must eventually call the same collective operation, and no other collective operation with the same communicator in between. This is consistent with the ordering rules for blocking collective operations in threaded environments.

*Rationale.* Matching blocking and nonblocking collective operations is not allowed because the implementation might use different communication algorithms for the two cases. Blocking collective operations may be optimized for minimal time to completion, while nonblocking collective operations may balance time to completion with CPU overhead and asynchronous progression.

The use of tags for collective operations can prevent certain hardware optimizations. (*End of rationale.*)

*Advice to users.* If program semantics require matching blocking and nonblocking collective operations, then a nonblocking collective operation can be initiated and immediately completed with a blocking wait to emulate blocking behavior. (*End of advice to users.*)

In terms of data movements, each nonblocking collective operation has the same effect as its blocking counterpart for intracommunicators and intercommunicators after completion. Likewise, upon completion, nonblocking collective reduction operations have the same effect as their blocking counterparts, and the same restrictions and recommendations on reduction orders apply.

The use of the “in place” option is allowed exactly as described for the corresponding blocking collective operations. When using the “in place” option, message buffers function as both send and receive buffers. Such buffers should not be modified or accessed until the operation completes.

Progression rules for nonblocking collective operations are similar to progression of nonblocking point-to-point operations, refer to Section 3.7.4.

*Advice to implementors.* Nonblocking collective operations can be implemented with local execution schedules [29] using nonblocking point-to-point communication and a reserved tag-space. (*End of advice to implementors.*)

### 5.12.1 Nonblocking Barrier Synchronization

`MPI_IBARRIER(comm , request)`

IN	comm	communicator (handle)
OUT	request	communication request (handle)

`int MPI_Ibarrier(MPI_Comm comm, MPI_Request *request)`

`MPI_IBARRIER(COMM, REQUEST, IERROR)`

INTEGER COMM, REQUEST, IERROR

```
{MPI::Request MPI::Comm::Ibarrier() const = 0(binding deprecated, see  

Section 15.2) }
```

MPI\_IBARRIER is a nonblocking version of MPI\_BARRIER. By calling MPI\_IBARRIER, a process notifies that it has reached the barrier. The call returns immediately, independent of whether other processes have called MPI\_IBARRIER. The usual barrier semantics are enforced at the corresponding completion operation (test or wait), which in the intra-communicator case will complete only after all other processes in the communicator have called MPI\_IBARRIER. In the intercommunicator case, it will complete when all processes in the remote group have called MPI\_IBARRIER.

*Advice to users.* A nonblocking barrier can be used to hide latency. Moving independent computations between the MPI\_IBARRIER and the subsequent completion call can overlap the barrier latency and therefore shorten possible waiting times. The semantic properties are also useful when mixing collective operations and point-to-point messages. (*End of advice to users.*)

### 5.12.2 Nonblocking Broadcast

MPI\_IBCAST(buffer, count, datatype, root, comm, request)

INOUT	buffer	starting address of buffer (choice)
IN	count	number of entries in buffer (non-negative integer)
IN	datatype	data type of buffer (handle)
IN	root	rank of broadcast root (integer)
IN	comm	communicator (handle)
OUT	request	communication request (handle)

```
int MPI_Ibcast(void* buffer, int count, MPI_Datatype datatype, int root,  

MPI_Comm comm, MPI_Request *request)
```

```
MPI_IBCAST(BUFFER, COUNT, DATATYPE, ROOT, COMM, REQUEST, IERROR)
```

```
<type> BUFFER(*)
```

```
INTEGER COUNT, DATATYPE, ROOT, COMM, REQUEST, IERROR
```

```
{MPI::Request MPI::Comm::Ibcast(void* buffer, int count,  

const MPI::Datatype& datatype, int root) const = 0(binding  

deprecated, see Section 15.2) }
```

This call starts a nonblocking variant of MPI\_BCAST (see Section 5.4).

#### Example using MPI\_IBCAST

The example in this section uses an intracommunicator.

#### Example 5.23

1        Start a broadcast of 100 ints from process 0 to every process in the group, perform some  
 2        computation on independent data, and then complete the outstanding broadcast operation.

```
3
4        MPI_Comm comm;
5        int array1[100], array2[100];
6        int root=0;
7        MPI_Request req;
8        ...
9        MPI_Ibcast(array1, 100, MPI_INT, root, comm, &req);
10       compute(array2, 100);
11       MPI_Wait(&req, MPI_STATUS_IGNORE);
```

### 13    5.12.3 Nonblocking Gather

```
14
15
16    MPI_IGATHER(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, root, comm,
17                request)
```

18	IN	sendbuf	starting address of send buffer (choice)
19	IN	sendcount	number of elements in send buffer (non-negative integer)
20			
21			
22	IN	sendtype	data type of send buffer elements (handle)
23	OUT	recvbuf	address of receive buffer (choice, significant only at root)
24			
25			
26	IN	recvcount	number of elements for any single receive (non-negative integer, significant only at root)
27			
28	IN	recvtype	data type of recv buffer elements (significant only at root) (handle)
29			
30			
31	IN	root	rank of receiving process (integer)
32	IN	comm	communicator (handle)
33	OUT	request	communication request (handle)
34			

```
35
36    int MPI_Igather(void* sendbuf, int sendcount, MPI_Datatype sendtype,
37                    void* recvbuf, int recvcount, MPI_Datatype recvtype, int root,
38                    MPI_Comm comm, MPI_Request *request)
```

```
39    MPI_IGATHER(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNT, RECVTYPE,
40                ROOT, COMM, REQUEST, IERROR)
```

```
41       <type> SENDBUF(*), RECVBUF(*)
42       INTEGER SENDCOUNT, SENDTYPE, RECVCOUNT, RECVTYPE, ROOT, COMM, REQUEST,
43       IERROR
```

```
44
45    {MPI::Request MPI::Comm::Igather(const void* sendbuf, int sendcount, const
46                    MPI::Datatype& sendtype, void* recvbuf, int recvcount,
47                    const MPI::Datatype& recvtype, int root) const = 0(binding
48                    deprecated, see Section 15.2) }
```

This call starts a nonblocking variant of `MPI_GATHER` (see Section 5.5).

`MPI_IGATHERV(sendbuf, sendcount, sendtype, recvbuf, recvcunts, displs, recvtype, root, comm, request)`

IN	sendbuf	starting address of send buffer (choice)
IN	sendcount	number of elements in send buffer (non-negative integer)
IN	sendtype	data type of send buffer elements (handle)
OUT	recvbuf	address of receive buffer (choice, significant only at root)
IN	recvcunts	non-negative integer array (of length group size) containing the number of elements that are received from each process (significant only at root)
IN	displs	integer array (of length group size). Entry <i>i</i> specifies the displacement relative to <code>recvbuf</code> at which to place the incoming data from process <i>i</i> (significant only at root)
IN	recvtype	data type of recv buffer elements (significant only at root) (handle)
IN	root	rank of receiving process (integer)
IN	comm	communicator (handle)
OUT	request	communication request (handle)

```
int MPI_Igatherv(void* sendbuf, int sendcount, MPI_Datatype sendtype,
                void* recvbuf, int *recvcunts, int *displs,
                MPI_Datatype recvtype, int root, MPI_Comm comm,
                MPI_Request *request)
```

```
MPI_IGATHERV(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNTS, DISPLS,
             RECVTYPE, ROOT, COMM, REQUEST, IERROR)
```

```
<type> SENDBUF(*), RECVBUF(*)
INTEGER SENDCOUNT, SENDTYPE, RECVCOUNTS(*), DISPLS(*), RECVTYPE, ROOT,
COMM, REQUEST, IERROR
```

```
{MPI::Request MPI::Comm::Igatherv(const void* sendbuf, int sendcount, const
MPI::Datatype& sendtype, void* recvbuf,
const int recvcunts[], const int displs[],
const MPI::Datatype& recvtype, int root) const = 0(binding
deprecated, see Section 15.2) }
```

This call starts a nonblocking variant of `MPI_GATHERV` (see Section 5.5).

ticket269.

```
1 MPI_IGATHERW(sendbuf, sendcount, sendtype, recvbuf, recvcunts, displs, recvtypes, root,
2 comm, request)
```

3	IN	sendbuf	starting address of send buffer (choice)
4	IN	sendcount	number of elements in send buffer (non-negative integer)
5	IN	sendtype	data type of send buffer elements (handle)
6	OUT	recvbuf	address of receive buffer (choice)
7	IN	recvcunts	non-negative integer array (of length group size) containing the number of elements that are received from each process (signifiant only at root)
8	IN	displs	integer array (of length group size). Entry <i>i</i> specifies the displacement relative to <i>recvbuf</i> from which to take the incoming data from process <i>i</i> (significant only at root)
9	IN	recvtypes	array of datatypes (of length group size). Entry <i>i</i> specifies the type of data received from process <i>i</i> (significant only at root) (array of handles)
10	IN	root	rank of sending process (integer)
11	IN	comm	communicator (handle)
12	OUT	request	communication request (handle)

```
25
26 int MPI_Igatherw(void* sendbuf, int sendcount, MPI_Datatype sendtype,
27 void* recvbuf, int recvcunts[], int displs[],
28 MPI_Datatype recvtypes[], int root, MPI_Comm comm,
29 MPI_Request *request)
```

```
30 MPI_IGATHERW(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNTS, DISPLS,
31 RECVTYPES, ROOT, COMM, REQUEST, IERROR)
32 <type> SENDBUF(*), RECVBUF(*)
33 INTEGER SENDCOUNT, SENDTYPE, RECVCOUNTS(*), DISPLS(*), RECVTYPES(*),
34 ROOT, COMM, REQUEST, IERROR
```

ticket109. This call starts a nonblocking variant of MPI\_GATHERW (see Section 5.5).

## 5.12.4 Nonblocking Scatter

`MPI_ISCATTER(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, root, comm, request)`

IN	sendbuf	address of send buffer (choice, significant only at root)
IN	sendcount	number of elements sent to each process (non-negative integer, significant only at root)
IN	sendtype	data type of send buffer elements (significant only at root) (handle)
OUT	recvbuf	address of receive buffer (choice)
IN	recvcount	number of elements in receive buffer (non-negative integer)
IN	recvtype	data type of receive buffer elements (handle)
IN	root	rank of sending process (integer)
IN	comm	communicator (handle)
OUT	request	communication request (handle)

```
int MPI_Iscatter(void* sendbuf, int sendcount, MPI_Datatype sendtype,
               void* recvbuf, int recvcount, MPI_Datatype recvtype, int root,
               MPI_Comm comm, MPI_Request *request)
```

```
MPI_ISCATTER(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNT, RECVTYPE,
             ROOT, COMM, REQUEST, IERROR)
<type> SENDBUF(*), RECVBUF(*)
INTEGER SENDCOUNT, SENDTYPE, RECVCOUNT, RECVTYPE, ROOT, COMM, REQUEST,
IERROR
```

```
{MPI::Request MPI::Comm::Iscatter(const void* sendbuf, int sendcount, const
MPI::Datatype& sendtype, void* recvbuf, int recvcount,
const MPI::Datatype& recvtype, int root) const = 0(binding
deprecated, see Section 15.2) }
```

This call starts a nonblocking variant of `MPI_SCATTER` (see Section 5.6).

```

1  MPI_ISCATTERV(sendbuf, sendcounts, displs, sendtype, recvbuf, recvcount, recvtpe, root,
2                comm, request)
3
4  IN      sendbuf      address of send buffer (choice, significant only at root)
5
6  IN      sendcounts    non-negative integer array (of length group size) speci-
7                      fying the number of elements to send to each processor
8
9  IN      displs        integer array (of length group size). Entry i specifies
10                      the displacement (relative to sendbuf) from which to
11                      take the outgoing data to process i
12
13  IN      sendtype      data type of send buffer elements (handle)
14
15  OUT     recvbuf      address of receive buffer (choice)
16
17  IN      recvcount     number of elements in receive buffer (non-negative in-
18                      teger)
19
20  IN      recvtpe      data type of receive buffer elements (handle)
21
22  IN      root          rank of sending process (integer)
23
24  IN      comm          communicator (handle)
25
26  OUT     request       communication request (handle)
27
28  int MPI_Iscatterv(void* sendbuf, int *sendcounts, int *displs,
29                  MPI_Datatype sendtype, void* recvbuf, int recvcount,
30                  MPI_Datatype recvtpe, int root, MPI_Comm comm,
31                  MPI_Request *request)
32
33  MPI_ISCATTERV(SENDBUF, SENDCOUNTS, DISPLS, SENDTYPE, RECVBUF, RECVCOUNT,
34                RECVTYPE, ROOT, COMM, REQUEST, IERROR)
35
36  <type> SENDBUF(*), RECVBUF(*)
37
38  INTEGER SENDCOUNTS(*), DISPLS(*), SENDTYPE, RECVCOUNT, RECVTYPE, ROOT,
39  COMM, REQUEST, IERROR
40
41  {MPI::Request MPI::Comm::Iscatterv(const void* sendbuf,
42      const int sendcounts[], const int displs[],
43      const MPI::Datatype& sendtype, void* recvbuf, int recvcount,
44      const MPI::Datatype& recvtpe, int root) const = 0(binding
45      deprecated, see Section 15.2) }
46
47
48

```

ticket269. This call starts a nonblocking variant of MPI\_SCATTERV (see Section 5.6).

```

        MPI_ISCATTERW(sendbuf, sendcounts, displs, sendtypes, recvbuf, recvcount,
        recvtype, root, comm, request)
IN      sendbuf      starting address of send buffer (choice, significant only
                        at root)
IN      sendcounts    non-negative integer array (of length group size) speci-
                        fying the number of elements to send to each processor
IN      displs        integer array (of length group size). Entry i specifies
                        the displacement relative to sendbuf from which to
                        take the outgoing data to process i
IN      sendtypes     array of datatypes (of length group size). Entry j
                        specifies the type of data to send to process j (array
                        of handles)
OUT     recvbuf       address of receive buffer (choice)
IN      recvcount     number of elements in receive buffer (non-negative in-
                        teger)
IN      recvtype      data type of receive buffer elements (handle)
IN      root          rank of sending process (integer)
IN      comm          communicator (handle)
OUT     request       communication request (handle)

int MPI_Iscatterw(void* sendbuf, int sendcounts[], int displs[],
        MPI_Datatype sendtypes[], void* recvbuf, int recvcount,
        MPI_Datatype recvtype, int root, MPI_Comm comm,
        MPI_Request *request)

MPI_ISCATTERW(SENDBUF, SENDCOUNTS, DISPLS, SENDTYPES, RECVBUF, RECVCOUNT,
        RECVTYPE, ROOT, COMM, REQUEST, IERROR)
<type> SENDBUF(*), RECVBUF(*)
INTEGER DISPLS(*), SENDCOUNTS(*), SENDTYPES, RECVCOUNT, RECVTYPE, ROOT,
COMM, REQUEST, IERROR

This call starts a nonblocking variant of MPI_SCATTERW (see Section 5.6).

```

ticket109.



### 5.12.5 Nonblocking Gather-to-all

`MPI_IALLGATHER(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, comm, request)`

IN	<code>sendbuf</code>	starting address of send buffer (choice)
IN	<code>sendcount</code>	number of elements in send buffer (non-negative integer)
IN	<code>sendtype</code>	data type of send buffer elements (handle)
OUT	<code>recvbuf</code>	address of receive buffer (choice)
IN	<code>recvcount</code>	number of elements received from any process (non-negative integer)
IN	<code>recvtype</code>	data type of receive buffer elements (handle)
IN	<code>comm</code>	communicator (handle)
OUT	<code>request</code>	communication request (handle)

```
int MPI_Iallgather(void* sendbuf, int sendcount, MPI_Datatype sendtype,
                  void* recvbuf, int recvcount, MPI_Datatype recvtype,
                  MPI_Comm comm, MPI_Request *request)
```

```
MPI_IALLGATHER(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, REVCOUNT, RECVTYPE,
               COMM, REQUEST, IERROR)
<type> SENDBUF(*), RECVBUF(*)
INTEGER SENDCOUNT, SENDTYPE, REVCOUNT, RECVTYPE, COMM, REQUEST, IERROR
```

```
{MPI::Request MPI::Comm::Iallgather(const void* sendbuf, int sendcount,
                                   const MPI::Datatype& sendtype, void* recvbuf, int recvcount,
                                   const MPI::Datatype& recvtype) const = 0 (binding deprecated, see
                                   Section 15.2) }
```

This call starts a nonblocking variant of `MPI_ALLGATHER` (see Section 5.7).

```
MPI_IALLGATHERV(sendbuf, sendcount, sendtype, recvbuf, recvcunts, displs, recvtype, comm,
                 request)
```

IN	sendbuf	starting address of send buffer (choice)
IN	sendcount	number of elements in send buffer (non-negative integer)
IN	sendtype	data type of send buffer elements (handle)
OUT	recvbuf	address of receive buffer (choice)
IN	recvcunts	non-negative integer array (of length group size) containing the number of elements that are received from each process
IN	displs	integer array (of length group size). Entry <i>i</i> specifies the displacement (relative to <b>recvbuf</b> ) at which to place the incoming data from process <i>i</i>
IN	recvtype	data type of receive buffer elements (handle)
IN	comm	communicator (handle)
OUT	request	communication request (handle)

```
int MPI_Iallgatherv(void* sendbuf, int sendcount, MPI_Datatype sendtype,
                   void* recvbuf, int *recvcunts, int *displs,
                   MPI_Datatype recvtype, MPI_Comm comm, MPI_Request* request)
```

```
MPI_IALLGATHERV(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNTS, DISPLS,
                 RECVTYPE, COMM, REQUEST, IERROR)
<type> SENDBUF(*), RECVBUF(*)
INTEGER SENDCOUNT, SENDTYPE, RECVCOUNTS(*), DISPLS(*), RECVTYPE, COMM,
REQUEST, IERROR
```

```
{MPI::Request MPI::Comm::Iallgatherv(const void* sendbuf, int sendcount,
                                     const MPI::Datatype& sendtype, void* recvbuf,
                                     const int recvcunts[], const int displs[],
                                     const MPI::Datatype& recvtype) const = 0 (binding deprecated, see
                                     Section 15.2) }
```

This call starts a nonblocking variant of **MPI\_ALLGATHERV** (see Section 5.7).

ticket269.

```

1  MPI_IALLGATHERW(sendbuf, sendcount, sendtype, recvbuf, recvcoun-
2                      comm, request)
3
4  IN      sendbuf      starting address of send buffer (choice)
5
6  IN      sendcount    number of elements in send buffer (non-negative inte-
7                      ger)
8
9  IN      sendtype     data type of send buffer elements (handle)
10
11 OUT     recvbuf      address of receive buffer (choice)
12
13 IN      recvcoun-    non-negative integer array (of length group size) con-
14                      taining the number of elements that are received from
15                      each process
16
17 IN      displs       integer array (of length group size). Entry i specifies
18                      the displacement (relative to recvbuf) at which to place
19                      the incoming data from process i
20
21 IN      recvtypes     array of datatypes (of length group size). Entry i
22                      specifies the type of data received from process i (ar-
23                      ray of handles)
24
25 IN      comm         communicator (handle)
26
27 OUT     request       communication request (handle)
28
29
30 int MPI_Iallgatherw(void* sendbuf, int sendcount, MPI_Datatype sendtype,
31                     void* recvbuf, int recvcoun-[], int displs[],
32                     MPI_Datatype recvtypes[], MPI_Comm comm, MPI_Request *request)
33
34 MPI_IALLGATHERW(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNTS, DISPLS,
35                 RECVTYPES, COMM, REQUEST, IERROR)
36
37 <type> SENDBUF(*), RECVBUF(*)
38
39 INTEGER SENDCOUNT, SENDTYPE, RECVCOUNTS(*), DISPLS(*), RECVTYPES(*),
40
41 COMM, REQUEST, IERROR

```

ticket109. This call starts a nonblocking variant of `MPI_ALLGATHERW` (see Section 5.7).

## 5.12.6 Nonblocking All-to-All Scatter/Gather

`MPI_IALLTOALL(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, comm, request)`

IN	sendbuf	starting address of send buffer (choice)
IN	sendcount	number of elements sent to each process (non-negative integer)
IN	sendtype	data type of send buffer elements (handle)
OUT	recvbuf	address of receive buffer (choice)
IN	recvcount	number of elements received from any process (non-negative integer)
IN	recvtype	data type of receive buffer elements (handle)
IN	comm	communicator (handle)
OUT	request	communication request (handle)

```

int MPI_Ialltoall(void* sendbuf, int sendcount, MPI_Datatype sendtype,
                 void* recvbuf, int recvcount, MPI_Datatype recvtype,
                 MPI_Comm comm, MPI_Request *request)

MPI_IALLTOALL(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, REVCOUNT, RECVTYPE,
              COMM, REQUEST, IERROR)
<type> SENDBUF(*), RECVBUF(*)
INTEGER SENDCOUNT, SENDTYPE, REVCOUNT, RECVTYPE, COMM, REQUEST, IERROR

{MPI::Request MPI::Comm::Ialltoall(const void* sendbuf, int sendcount,
                                   const MPI::Datatype& sendtype, void* recvbuf, int recvcount,
                                   const MPI::Datatype& recvtype) const = 0 (binding deprecated, see
                                   Section 15.2) }
```

This call starts a nonblocking variant of `MPI_ALLTOALL` (see Section 5.8).

```

1  MPI_IALLTOALLV(sendbuf, sendcounts, sdispls, sendtype, recvbuf, recvcoun-
2      recvtype, comm, request)
3
4      IN      sendbuf      starting address of send buffer (choice)
5
6      IN      sendcounts   non-negative integer array (of length group size) speci-
7                          fying the number of elements to send to each processor
8
9      IN      sdispls      integer array (of length group size). Entry j specifies
10                      the displacement (relative to sendbuf) from which to
11                      take the outgoing data destined for process j
12
13      IN      sendtype     data type of send buffer elements (handle)
14
15      OUT     recvbuf      address of receive buffer (choice)
16
17      IN      recvcoun-    non-negative integer array (of length group size) speci-
18                      fying the number of elements that can be received
19                      from each processor
20
21      IN      rdispls      integer array (of length group size). Entry i specifies
22                      the displacement (relative to recvbuf) at which to place
23                      the incoming data from process i
24
25      IN      recvtype     data type of receive buffer elements (handle)
26
27      IN      comm         communicator (handle)
28
29      OUT     request      communication request (handle)
30
31  int MPI_Ialltoallv(void* sendbuf, int *sendcounts, int *sdispls,
32      MPI_Datatype sendtype, void* recvbuf, int *recvcoun-
33      int *rdispls, MPI_Datatype recvtype, MPI_Comm comm,
34      MPI_Request *request)
35
36  MPI_IALLTOALLV(SENDBUF, SENDCOUNTS, SDISPLS, SENDTYPE, RECVBUF, RECVCOUNTS,
37      RDISPLS, RECVTYPE, COMM, REQUEST, IERROR)
38
39      <type> SENDBUF(*), RECVBUF(*)
40      INTEGER SENDCOUNTS(*), SDISPLS(*), SENDTYPE, REVCOUNTS(*), RDISPLS(*),
41      RECVTYPE, COMM, REQUEST, IERROR
42
43  {MPI::Request MPI::Comm::Ialltoallv(const void* sendbuf,
44      const int sendcounts[], const int sdispls[],
45      const MPI::Datatype& sendtype, void* recvbuf,
46      const int recvcoun-
47      const int rdispls[],
48      const MPI::Datatype& recvtype) const = 0 (binding deprecated, see
49      Section 15.2) }

```

This call starts a nonblocking variant of MPI\_ALLTOALLV (see Section 5.8).

```
MPI_IALLTOALLW(sendbuf, sendcounts, sdispls, sendtypes, recvbuf, recvcounts, rdispls,
                recvtypes, comm, request)
```

IN	sendbuf	starting address of send buffer (choice)
IN	sendcounts	integer array (of length group size) specifying the number of elements to send to each processor (array of non-negative integers)
IN	sdispls	integer array (of length group size). Entry <i>j</i> specifies the displacement in bytes (relative to <code>sendbuf</code> ) from which to take the outgoing data destined for process <i>j</i> (array of integers)
IN	sendtypes	array of datatypes (of length group size). Entry <i>j</i> specifies the type of data to send to process <i>j</i> (array of handles)
OUT	recvbuf	address of receive buffer (choice)
IN	recvcounts	integer array (of length group size) specifying the number of elements that can be received from each processor (array of non-negative integers)
IN	rdispls	integer array (of length group size). Entry <i>i</i> specifies the displacement in bytes (relative to <code>recvbuf</code> ) at which to place the incoming data from process <i>i</i> (array of integers)
IN	recvtypes	array of datatypes (of length group size). Entry <i>i</i> specifies the type of data received from process <i>i</i> (array of handles)
IN	comm	communicator (handle)
OUT	request	communication request (handle)

```
int MPI_Ialltoallw(void* sendbuf, int sendcounts[], int sdispls[],
                  MPI_Datatype sendtypes[], void* recvbuf, int recvcounts[],
                  int rdispls[], MPI_Datatype recvtypes[], MPI_Comm comm,
                  MPI_Request *request)
```

```
MPI_IALLTOALLW(SENDBUF, SENDCOUNTS, SDISPLS, SENDTYPES, RECVBUF,
               RECVCOUNTS, RDISPLS, RECVTYPES, COMM, REQUEST, IERROR)
<type> SENDBUF(*), RECVBUF(*)
INTEGER SENDCOUNTS(*), SDISPLS(*), SENDTYPES(*), RECVCOUNTS(*),
RDISPLS(*), RECVTYPES(*), COMM, REQUEST, IERROR
```

```
{MPI::Request MPI::Comm::Ialltoallw(const void* sendbuf, const int
    sendcounts[], const int sdispls[], const MPI::Datatype
    sendtypes[], void* recvbuf, const int recvcounts[], const int
    rdispls[], const MPI::Datatype recvtypes[]) const = 0(binding
    deprecated, see Section 15.2) }
```

This call starts a nonblocking variant of `MPI_ALLTOALLW` (see Section 5.8).

### 5.12.7 Nonblocking Reduce

`MPI_IREDUCE(sendbuf, recvbuf, count, datatype, op, root, comm, request)`

IN	sendbuf	address of send buffer (choice)
OUT	recvbuf	address of receive buffer (choice, significant only at root)
IN	count	number of elements in send buffer (non-negative integer)
IN	datatype	data type of elements of send buffer (handle)
IN	op	reduce operation (handle)
IN	root	rank of root process (integer)
IN	comm	communicator (handle)
OUT	request	communication request (handle)

```
int MPI_Ireduce(void* sendbuf, void* recvbuf, int count,
               MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm comm,
               MPI_Request *request)
```

```
MPI_IREDUCE(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, ROOT, COMM, REQUEST,
            IERROR)
```

```
<type> SENDBUF(*), RECVBUF(*)
```

```
INTEGER COUNT, DATATYPE, OP, ROOT, COMM, REQUEST, IERROR
```

```
{MPI::Request MPI::Comm::Ireduce(const void* sendbuf, void* recvbuf,
                                int count, const MPI::Datatype& datatype, const MPI::Op& op,
                                int root) const = 0(binding deprecated, see Section 15.2) }
```

This call starts a nonblocking variant of `MPI_REDUCE` (see Section 5.9.1).

*Advice to implementors.* The implementation is explicitly allowed to use different algorithms for blocking and nonblocking reduction operations that might change the order of evaluation of the operations. However, as for `MPI_REDUCE`, it is strongly recommended that `MPI_IREDUCE` be implemented so that the same result be obtained whenever the function is applied on the same arguments, appearing in the same order. Note that this may prevent optimizations that take advantage of the physical location of processes. (*End of advice to implementors.*)

*Advice to users.* For operations which are not truly associative, the result delivered upon completion of the nonblocking reduction may not exactly equal the result delivered by the blocking reduction, even when specifying the same arguments in the same order. (*End of advice to users.*)

## 5.12.8 Nonblocking All-Reduce

`MPI_IALLREDUCE(sendbuf, recvbuf, count, datatype, op, comm, request)`

IN	sendbuf	starting address of send buffer (choice)
OUT	recvbuf	starting address of receive buffer (choice)
IN	count	number of elements in send buffer (non-negative integer)
IN	datatype	data type of elements of send buffer (handle)
IN	op	operation (handle)
IN	comm	communicator (handle)
OUT	request	communication request (handle)

```
int MPI_Iallreduce(void* sendbuf, void* recvbuf, int count,
                  MPI_Datatype datatype, MPI_Op op, MPI_Comm comm,
                  MPI_Request *request)
```

```
MPI_IALLREDUCE(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, COMM, REQUEST,
               IERROR)
```

```
<type> SENDBUF(*), RECVBUF(*)
```

```
INTEGER COUNT, DATATYPE, OP, COMM, REQUEST, IERROR
```

```
{MPI::Request MPI::Comm::Iallreduce(const void* sendbuf, void* recvbuf,
    int count, const MPI::Datatype& datatype, const MPI::Op& op)
    const = 0 (binding deprecated, see Section 15.2) }
```

This call starts a nonblocking variant of `MPI_ALLREDUCE` (see Section 5.9.6).

## 5.12.9 Nonblocking Reduce-Scatter with Equal Blocks

`MPI_IREDUCE_SCATTER_BLOCK(sendbuf, recvbuf, recvcnt, datatype, op, comm, request)`

IN	sendbuf	starting address of send buffer (choice)
OUT	recvbuf	starting address of receive buffer (choice)
IN	recvcnt	element count per block (non-negative integer)
IN	datatype	data type of elements of send and receive buffers (handle)
IN	op	operation (handle)
IN	comm	communicator (handle)
OUT	request	communication request (handle)



```

1  int MPI_Ireduce_scatter_block(void* sendbuf, void* recvbuf, int recvcnt,
2      MPI_Datatype datatype, MPI_Op op, MPI_Comm comm,
3      MPI_Request *request)
4
5  MPI_IREDUCE_SCATTER_BLOCK(SENDBUF, RECVBUF, RECVCOUNT, DATATYPE, OP, COMM,
6      REQUEST, IERROR)
7      <type> SENDBUF(*), RECVBUF(*)
8      INTEGER RECVCOUNT, DATATYPE, OP, COMM, REQUEST, IERROR
9
10 {MPI::Request MPI::Comm::Ireduce_scatter_block(const void* sendbuf,
11     void* recvbuf, int recvcnt, const MPI::Datatype& datatype,
12     const MPI::Op& op) const = 0(binding deprecated, see Section 15.2) }

```

This call starts a nonblocking variant of MPI\_REDUCE\_SCATTER\_BLOCK (see Section 5.10.1).

#### 5.12.10 Nonblocking Reduce-Scatter

```

19 MPI_IREDUCE_SCATTER(sendbuf, recvbuf, recvcnts, datatype, op, comm, request)
20
21 IN      sendbuf      starting address of send buffer (choice)
22 OUT     recvbuf      starting address of receive buffer (choice)
23 IN      recvcnts     non-negative integer array specifying the number of
24                  elements in result distributed to each process. Array
25                  must be identical on all calling processes.
26 IN      datatype     data type of elements of input buffer (handle)
27 IN      op           operation (handle)
28 IN      comm         communicator (handle)
29 OUT     request      communication request (handle)
30
31
32 int MPI_Ireduce_scatter(void* sendbuf, void* recvbuf, int *recvcnts,
33     MPI_Datatype datatype, MPI_Op op, MPI_Comm comm,
34     MPI_Request *request)
35
36 MPI_IREDUCE_SCATTER(SENDBUF, RECVBUF, RECVCOUNTS, DATATYPE, OP, COMM,
37     REQUEST, IERROR)
38     <type> SENDBUF(*), RECVBUF(*)
39     INTEGER RECVCOUNTS(*), DATATYPE, OP, COMM, REQUEST, IERROR
40
41 {MPI::Request MPI::Comm::Ireduce_scatter(const void* sendbuf,
42     void* recvbuf, int recvcnts[],
43     const MPI::Datatype& datatype, const MPI::Op& op)
44     const = 0(binding deprecated, see Section 15.2) }

```

This call starts a nonblocking variant of MPI\_REDUCE\_SCATTER (see Section 5.10.2).

## 5.12.11 Nonblocking Inclusive Scan

`MPI_ISCAN(sendbuf, recvbuf, count, datatype, op, comm, request)`

IN	sendbuf	starting address of send buffer (choice)
OUT	recvbuf	starting address of receive buffer (choice)
IN	count	number of elements in input buffer (non-negative integer)
IN	datatype	data type of elements of input buffer (handle)
IN	op	operation (handle)
IN	comm	communicator (handle)
OUT	request	communication request (handle)

```
int MPI_Iscan(void* sendbuf, void* recvbuf, int count,
              MPI_Datatype datatype, MPI_Op op, MPI_Comm comm,
              MPI_Request *request)
```

```
MPI_ISCAN(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, COMM, REQUEST, IERROR)
<type> SENDBUF(*), RECVBUF(*)
INTEGER COUNT, DATATYPE, OP, COMM, REQUEST, IERROR
```

```
{MPI::Request MPI::Intracomm::Iscan(const void* sendbuf, void* recvbuf,
    int count, const MPI::Datatype& datatype, const MPI::Op& op)
    const(binding deprecated, see Section 15.2) }
```

This call starts a nonblocking variant of `MPI_SCAN` (see Section 5.11).

## 5.12.12 Nonblocking Exclusive Scan

`MPI_IEXSCAN(sendbuf, recvbuf, count, datatype, op, comm, request)`

IN	sendbuf	starting address of send buffer (choice)
OUT	recvbuf	starting address of receive buffer (choice)
IN	count	number of elements in input buffer (non-negative integer)
IN	datatype	data type of elements of input buffer (handle)
IN	op	operation (handle)
IN	comm	intracommunicator (handle)
OUT	request	communication request (handle)

```
int MPI_Iexscan(void* sendbuf, void* recvbuf, int count,
                MPI_Datatype datatype, MPI_Op op, MPI_Comm comm,
                MPI_Request *request)
```

```

1 MPI_IEXSCAN(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, COMM, REQUEST, IERROR)
2     <type> SENDBUF(*), RECVBUF(*)
3     INTEGER COUNT, DATATYPE, OP, COMM, REQUEST, IERROR
4
5 {MPI::Request MPI::Intracomm::Iexscan(const void* sendbuf, void* recvbuf,
6     int count, const MPI::Datatype& datatype, const MPI::Op& op)
7     const(binding deprecated, see Section 15.2) }

```

This call starts a nonblocking variant of MPI\_EXSCAN (see Section 5.11.2).

### 5.13 Correctness

A correct, portable program must invoke collective communications so that deadlock will not occur, whether collective communications are synchronizing or not. The following examples illustrate dangerous use of collective routines on intracommunicators.

#### Example 5.24

The following is erroneous.

```

19 switch(rank) {
20     case 0:
21         MPI_Bcast(buf1, count, type, 0, comm);
22         MPI_Bcast(buf2, count, type, 1, comm);
23         break;
24     case 1:
25         MPI_Bcast(buf2, count, type, 1, comm);
26         MPI_Bcast(buf1, count, type, 0, comm);
27         break;
28 }

```

We assume that the group of `comm` is  $\{0,1\}$ . Two processes execute two broadcast operations in reverse order. If the operation is synchronizing then a deadlock will occur.

Collective operations must be executed in the same order at all members of the communication group.

#### Example 5.25

The following is erroneous.

```

38 switch(rank) {
39     case 0:
40         MPI_Bcast(buf1, count, type, 0, comm0);
41         MPI_Bcast(buf2, count, type, 2, comm2);
42         break;
43     case 1:
44         MPI_Bcast(buf1, count, type, 1, comm1);
45         MPI_Bcast(buf2, count, type, 0, comm0);
46         break;
47     case 2:
48         MPI_Bcast(buf1, count, type, 2, comm2);

```

```

    MPI_Bcast(buf2, count, type, 1, comm1);
    break;
}

```

Assume that the group of `comm0` is  $\{0,1\}$ , of `comm1` is  $\{1, 2\}$  and of `comm2` is  $\{2,0\}$ . If the broadcast is a synchronizing operation, then there is a cyclic dependency: the broadcast in `comm2` completes only after the broadcast in `comm0`; the broadcast in `comm0` completes only after the broadcast in `comm1`; and the broadcast in `comm1` completes only after the broadcast in `comm2`. Thus, the code will deadlock.

Collective operations must be executed in an order so that no cyclic dependencies occur. Nonblocking collective operations can alleviate this issue.

### Example 5.26

The following is erroneous.

```

switch(rank) {
    case 0:
        MPI_Bcast(buf1, count, type, 0, comm);
        MPI_Send(buf2, count, type, 1, tag, comm);
        break;
    case 1:
        MPI_Recv(buf2, count, type, 0, tag, comm, status);
        MPI_Bcast(buf1, count, type, 0, comm);
        break;
}

```

Process zero executes a broadcast, followed by a blocking send operation. Process one first executes a blocking receive that matches the send, followed by broadcast call that matches the broadcast of process zero. This program may deadlock. The broadcast call on process zero *may* block until process one executes the matching broadcast call, so that the send is not executed. Process one will definitely block on the receive and so, in this case, never executes the broadcast.

The relative order of execution of collective operations and point-to-point operations should be such, so that even if the collective operations and the point-to-point operations are synchronizing, no deadlock will occur.

### Example 5.27

An unsafe, non-deterministic program.

```

switch(rank) {
    case 0:
        MPI_Bcast(buf1, count, type, 0, comm);
        MPI_Send(buf2, count, type, 1, tag, comm);
        break;
    case 1:
        MPI_Recv(buf2, count, type, MPI_ANY_SOURCE, tag, comm, status);
        MPI_Bcast(buf1, count, type, 0, comm);
        MPI_Recv(buf2, count, type, MPI_ANY_SOURCE, tag, comm, status);
        break;
}

```

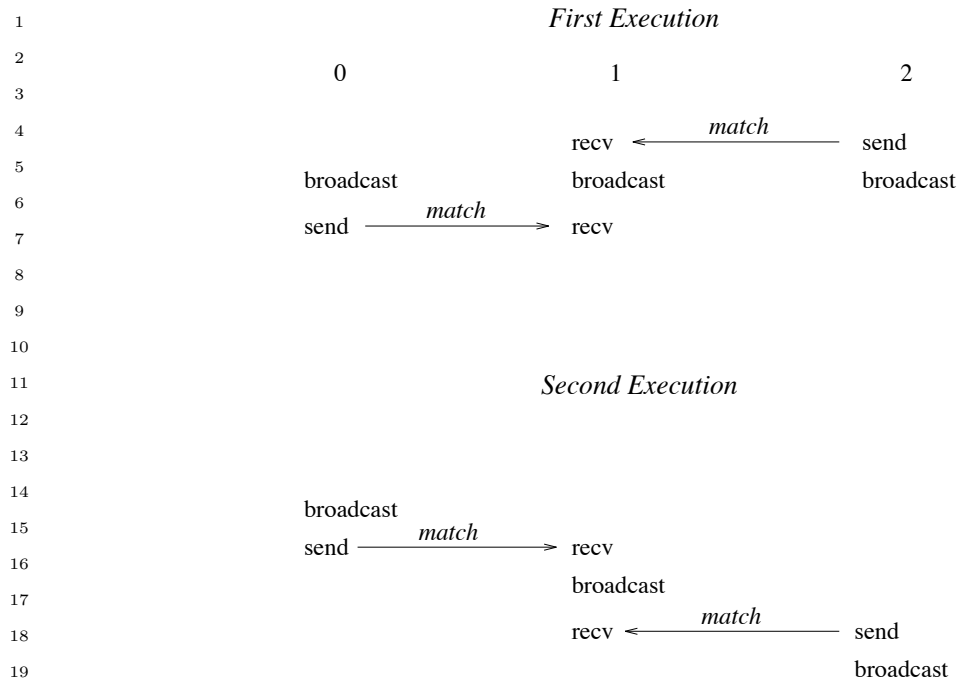


Figure 5.12: A race condition causes non-deterministic matching of sends and receives. One cannot rely on synchronization from a broadcast to make the program deterministic.

```

24 case 2:
25     MPI_Send(buf2, count, type, 1, tag, comm);
26     MPI_Bcast(buf1, count, type, 0, comm);
27     break;
28 }

```

All three processes participate in a broadcast. Process 0 sends a message to process 1 after the broadcast, and process 2 sends a message to process 1 before the broadcast. Process 1 receives before and after the broadcast, with a wildcard source argument.

Two possible executions of this program, with different matchings of sends and receives, are illustrated in Figure 5.12. Note that the second execution has the peculiar effect that a send executed after the broadcast is received at another node before the broadcast. This example illustrates the fact that one should not rely on collective communication functions to have particular synchronization effects. A program that works correctly only when the first execution occurs (only when broadcast is synchronizing) is erroneous.

Finally, in multithreaded implementations, one can have more than one, concurrently executing, collective communication call at a process. In these situations, it is the user's responsibility to ensure that the same communicator is not used concurrently by two different collective communication calls at the same process.

*Advice to implementors.* Assume that broadcast is implemented using point-to-point MPI communication. Suppose the following two rules are followed.

1. All receives specify their source explicitly (no wildcards).
2. Each process sends all messages that pertain to one collective call before sending any message that pertain to a subsequent collective call.

Then, messages belonging to successive broadcasts cannot be confused, as the order of point-to-point messages is preserved.

It is the implementor's responsibility to ensure that point-to-point messages are not confused with collective messages. One way to accomplish this is, whenever a communicator is created, to also create a "hidden communicator" for collective communication. One could achieve a similar effect more cheaply, for example, by using a hidden tag or context bit to indicate whether the communicator is used for point-to-point or collective communication. (*End of advice to implementors.*)

### Example 5.28

Blocking and nonblocking collective operations can be interleaved, i.e., a blocking collective operation can be posted even if there is a nonblocking collective operation outstanding.

```
MPI_Request req;

MPI_Ibarrier(comm, &req);
MPI_Bcast(buf1, count, type, 0, comm);
MPI_Wait(&req, MPI_STATUS_IGNORE);
```

Each process starts a nonblocking barrier operation, participates in a blocking broadcast and then waits until every other process started the barrier operation. This effectively turns the broadcast into a synchronizing broadcast with possible communication/communication overlap (MPI\_Bcast is allowed, but not required to synchronize).

### Example 5.29

The starting order of collective operations on a particular communicator defines their matching. The following example shows an erroneous matching of different collective operations on the same communicator.

```
MPI_Request req;
switch(rank) {
    case 0:
        /* erroneous matching */
        MPI_Ibarrier(comm, &req);
        MPI_Bcast(buf1, count, type, 0, comm);
        MPI_Wait(&req, MPI_STATUS_IGNORE);
        break;
    case 1:
        /* erroneous matching */
        MPI_Bcast(buf1, count, type, 0, comm);
        MPI_Ibarrier(comm, &req);
        MPI_Wait(&req, MPI_STATUS_IGNORE);
        break;
}
```

This ordering would match MPI\_Ibarrier on rank 0 with MPI\_Bcast on rank 1 which is erroneous and the program behavior is undefined. However, if such an order is required, the user must create different duplicate communicators and perform the operations on them. If started with two processes, the following program would be correct:

```

1  MPI_Request req;
2  MPI_Comm dupcomm;
3  MPI_Comm_dup(comm, &dupcomm);
4  switch(rank) {
5      case 0:
6          MPI_Ibarrier(comm, &req);
7          MPI_Bcast(buf1, count, type, 0, dupcomm);
8          MPI_Wait(&req, MPI_STATUS_IGNORE);
9          break;
10     case 1:
11         MPI_Bcast(buf1, count, type, 0, dupcomm);
12         MPI_Ibarrier(comm, &req);
13         MPI_Wait(&req, MPI_STATUS_IGNORE);
14         break;
15 }

```

*Advice to users.* The use of different communicators offers some flexibility regarding the matching of nonblocking collective operations. In this sense, communicators could be used as an equivalent to tags. However, communicator construction might induce overheads so that this should be used carefully. (*End of advice to users.*)

### Example 5.30

Nonblocking collective operations can rely on the same progression rules as nonblocking point-to-point messages. Thus, if started with two processes, the following program is a valid MPI program and is guaranteed to terminate:

```

27 MPI_Request req;
28
29 switch(rank) {
30     case 0:
31         MPI_Ibarrier(comm, &req);
32         MPI_Wait(&req, MPI_STATUS_IGNORE);
33         MPI_Send(buf, count, dtype, 1, tag, comm);
34         break;
35     case 1:
36         MPI_Ibarrier(comm, &req);
37         MPI_Recv(buf, count, dtype, 0, tag, comm, MPI_STATUS_IGNORE);
38         MPI_Wait(&req, MPI_STATUS_IGNORE);
39         break;
40 }

```

The MPI library must progress the barrier in the MPI\_Recv call. Thus, the MPI\_Wait call in rank 0 will eventually complete, which enables the matching MPI\_Send so all calls eventually return.

### Example 5.31

Blocking and nonblocking collective operations do not match. The following example is erroneous.

```

MPI_Request req;

switch(rank) {
    case 0:
        /* erroneous false matching of Alltoall and Ialltoall */
        MPI_Ialltoall(sbuf, scnt, stype, rbuf, rcnt, rtype, comm, &req);
        MPI_Wait(&req, MPI_STATUS_IGNORE);
        break;
    case 1:
        /* erroneous false matching of Alltoall and Ialltoall */
        MPI_Alltoall(sbuf, scnt, stype, rbuf, rcnt, rtype, comm);
        break;
}

```

**Example 5.32**

Collective and point-to-point requests can be mixed in functions that enable multiple completions. If started with two processes, the following program is valid.

```

MPI_Request reqs[2];

switch(rank) {
    case 0:
        MPI_Ibarrier(comm, &reqs[0]);
        MPI_Send(buf, count, dtype, 1, tag, comm);
        MPI_Wait(&reqs[0], MPI_STATUS_IGNORE);
        break;
    case 1:
        MPI_Irecv(buf, count, dtype, 0, tag, comm, &reqs[0]);
        MPI_Ibarrier(comm, &reqs[1]);
        MPI_Waitall(2, reqs, MPI_STATUSES_IGNORE);
        break;
}

```

The Waitall call returns only after the barrier and the receive completed.

**Example 5.33**

Multiple nonblocking collective operations can be outstanding on a single communicator and match in order.

```

MPI_Request reqs[3];

compute(buf1);
MPI_Ibcast(buf1, count, type, 0, comm, &reqs[0]);
compute(buf2);
MPI_Ibcast(buf2, count, type, 0, comm, &reqs[1]);
compute(buf3);
MPI_Ibcast(buf3, count, type, 0, comm, &reqs[2]);
MPI_Waitall(3, reqs, MPI_STATUSES_IGNORE);

```



*Advice to users.* Pipelining and double-buffering techniques can efficiently be used to overlap computation and communication. However, having too many outstanding requests might have a negative impact on performance. (*End of advice to users.*)

*Advice to implementors.* The use of pipelining may generate many outstanding requests. A high-quality hardware-supported implementation with limited resources should be able to fall back to a software implementation if its resources are exhausted. In this way, the implementation could limit the number of outstanding requests only by the available memory. (*End of advice to implementors.*)

### Example 5.34

Nonblocking collective operations can also be used to enable simultaneous collective operations on multiple overlapping communicators (see Figure 5.13). The following example is started with three processes and three communicators. The first communicator `comm1` includes ranks 0 and 1, `comm2` includes ranks 1 and 2 and `comm3` spans ranks 0 and 2. It is not possible to perform a blocking collective operation on all communicators because there exists no deadlock-free order to invoke them. However, nonblocking collective operations can easily be used to achieve this task.

```
MPI_Request reqs[2];

switch(rank) {
    case 0:
        MPI_Iallreduce(sbuf1, rbuf1, count, dtype, MPI_SUM, comm1, &reqs[0]);
        MPI_Iallreduce(sbuf3, rbuf3, count, dtype, MPI_SUM, comm3, &reqs[1]);
        break;
    case 1:
        MPI_Iallreduce(sbuf1, rbuf1, count, dtype, MPI_SUM, comm1, &reqs[0]);
        MPI_Iallreduce(sbuf2, rbuf2, count, dtype, MPI_SUM, comm2, &reqs[1]);
        break;
    case 2:
        MPI_Iallreduce(sbuf2, rbuf2, count, dtype, MPI_SUM, comm2, &reqs[0]);
        MPI_Iallreduce(sbuf3, rbuf3, count, dtype, MPI_SUM, comm3, &reqs[1]);
        break;
}
MPI_Waitall(2, reqs, MPI_STATUSES_IGNORE);
```

*Advice to users.* This method can be useful if overlapping neighboring regions (halo or ghost zones) are used in collective operations. The sequence of the two calls in each process is irrelevant because the two nonblocking operations are performed on different communicators. (*End of advice to users.*)

### Example 5.35

The progress of multiple outstanding nonblocking collective operations is completely independent.




Figure 5.13: Example with overlapping communicators.

```

MPI_Request reqs[2];

compute(buf1);
MPI_Ibcast(buf1, count, type, 0, comm, &reqs[0]);
compute(buf2);
MPI_Ibcast(buf2, count, type, 0, comm, &reqs[1]);
MPI_Wait(&reqs[1], MPI_STATUS_IGNORE);
/* nothing is known about the status of the first bcast here */
MPI_Wait(&reqs[0], MPI_STATUS_IGNORE);

```

Finishing the second `MPI_IBCAST` is completely independent of the first one. This means that it is not guaranteed that the first broadcast operation is finished or even started after the second one is completed via `reqs[1]`.

1  
2  
3  
4  
5  
6  
7  
8  
9  
10  
11  
12  
13  
14  
15  
16  
17  
18  
19  
20  
21  
22  
23  
24  
25  
26  
27  
28  
29  
30  
31  
32  
33  
34  
35  
36  
37  
38  
39  
40  
41  
42  
43  
44  
45  
46  
47  
48

# Chapter 12

## External Interfaces

### 12.1 Introduction

This chapter begins with calls used to create **generalized requests**, which allow users to create new nonblocking operations with an interface similar to what is present in MPI. This can be used to layer new functionality on top of MPI. Next, Section 12.3 deals with setting the information found in `status`. [This is] This functionality is needed for generalized requests. ticket0.

The chapter continues, in Section 12.4, with a discussion of how threads are to be handled in MPI. Although thread compliance is not required, the standard specifies how threads are to work if they are provided.

### 12.2 Generalized Requests

The goal of generalized requests is to allow users to define new nonblocking operations. Such an outstanding nonblocking operation is represented by a (generalized) request. A fundamental property of nonblocking operations is that progress toward the completion of this operation occurs asynchronously, i.e., concurrently with normal program execution. Typically, this requires execution of code concurrently with the execution of the user code, e.g., in a separate thread or in a signal handler. Operating systems provide a variety of mechanisms in support of concurrent execution. MPI does not attempt to standardize or replace these mechanisms: it is assumed programmers who wish to define new asynchronous operations will use the mechanisms provided by the underlying operating system. Thus, the calls in this section only provide a means for defining the effect of MPI calls such as `MPI_WAIT` or `MPI_CANCEL` when they apply to generalized requests, and for signaling to MPI the completion of a generalized operation.

*Rationale.* It is tempting to also define an MPI standard mechanism for achieving concurrent execution of user-defined nonblocking operations. However, it is very difficult to define such a mechanism without consideration of the specific mechanisms used in the operating system. The Forum feels that concurrency mechanisms are a proper part of the underlying operating system and should not be standardized by MPI; the MPI standard should only deal with the interaction of such mechanisms with MPI. (*End of rationale.*)

For a regular request, the operation associated with the request is performed by the MPI implementation, and the operation completes without intervention by the application. For a generalized request, the operation associated with the request is performed by the application; therefore, the application must notify MPI when the operation completes. This is done by making a call to `MPI_GREQUEST_COMPLETE`. MPI maintains the “completion” status of generalized requests. Any other request state has to be maintained by the user.

A new generalized request is started with

```
MPI_GREQUEST_START(query_fn, free_fn, cancel_fn, extra_state, request)
```

IN	query_fn	callback function invoked when request status is queried (function)
IN	free_fn	callback function invoked when request is freed (function)
IN	cancel_fn	callback function invoked when request is cancelled (function)
IN	extra_state	extra state
OUT	request	generalized request (handle)

```
int MPI_Grequest_start(MPI_Grequest_query_function *query_fn,
                      MPI_Grequest_free_function *free_fn,
                      MPI_Grequest_cancel_function *cancel_fn, void *extra_state,
                      MPI_Request *request)

MPI_GREQUEST_START(QUERY_FN, FREE_FN, CANCEL_FN, EXTRA_STATE, REQUEST,
                  IERROR)
INTEGER REQUEST, IERROR
EXTERNAL QUERY_FN, FREE_FN, CANCEL_FN
INTEGER (KIND=MPI_ADDRESS_KIND) EXTRA_STATE

{static MPI::Grequest
    MPI::Grequest::Start(const MPI::Grequest::Query_function*
                        query_fn, const MPI::Grequest::Free_function* free_fn,
                        const MPI::Grequest::Cancel_function* cancel_fn,
                        void *extra_state) (binding deprecated, see Section 15.2) }
```

*Advice to users.* Note that a generalized request belongs, in C++, to the class `MPI::Grequest`, which is a derived class of `MPI::Request`. It is of the same type as regular requests, in C and Fortran. (*End of advice to users.*)

The call starts a generalized request and returns a handle to it in `request`.

The syntax and meaning of the callback functions are listed below. All callback functions are passed the `extra_state` argument that was associated with the request by the starting call `MPI_GREQUEST_START[. This can; extra_state can` be used to maintain user-defined state for the request.

In C, the query function is

```
typedef int MPI_Grequest_query_function(void *extra_state,
                                         MPI_Status *status);
```

in Fortran

```
SUBROUTINE GREQUEST_QUERY_FUNCTION(EXTRA_STATE, STATUS, IERROR)
  INTEGER STATUS(MPI_STATUS_SIZE), IERROR
  INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
```

and in C++

```
{typedef int MPI::Grequest::Query_function(void* extra_state,
                                             MPI::Status& status); (binding deprecated, see Section 15.2)}
```

**[query\_fn]** The **query\_fn** function computes the status that should be returned for the generalized request. The status also includes information about successful/unsuccessful cancellation of the request (result to be returned by `MPI_TEST_CANCELLED`).

**[query\_fn]** The **query\_fn** callback is invoked by the `MPI_{WAIT|TEST}{ANY|SOME|ALL}` call that completed the generalized request associated with this callback. The callback function is also invoked by calls to `MPI_REQUEST_GET_STATUS`, if the request is complete when the call occurs. In both cases, the callback is passed a reference to the corresponding status variable passed by the user to the MPI call; the status set by the callback function is returned by the MPI call. If the user provided `MPI_STATUS_IGNORE` or `MPI_STATUSES_IGNORE` to the MPI function that causes **query\_fn** to be called, then MPI will pass a valid status object to **query\_fn**, and this status will be ignored upon return of the callback function. Note that **query\_fn** is invoked only after `MPI_GREQUEST_COMPLETE` is called on the request; it may be invoked several times for the same generalized request, e.g., if the user calls `MPI_REQUEST_GET_STATUS` several times for this request. Note also that a call to `MPI_{WAIT|TEST}{SOME|ALL}` may cause multiple invocations of **query\_fn** callback functions, one for each generalized request that is completed by the MPI call. The order of these invocations is not specified by MPI.

In C, the free function is

```
typedef int MPI_Grequest_free_function(void *extra_state);
```

and in Fortran

```
SUBROUTINE GREQUEST_FREE_FUNCTION(EXTRA_STATE, IERROR)
  INTEGER IERROR
  INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
```

and in C++

```
{typedef int MPI::Grequest::Free_function(void* extra_state); (binding deprecated, see Section 15.2)}
```

**[free\_fn]** The **free\_fn** function is invoked to clean up user-allocated resources when the generalized request is freed.

**[free\_fn]** The **free\_fn** callback is invoked by the `MPI_{WAIT|TEST}{ANY|SOME|ALL}` call that completed the generalized request associated with this callback. **free\_fn** is invoked after the call to **query\_fn** for the same request. However, if the MPI call completed multiple generalized requests, the order in which **free\_fn** callback functions are invoked is not specified by MPI.

[free\_fn] The `free_fn` callback is also invoked for generalized requests that are freed by a call to `MPI_REQUEST_FREE` (no call to `WAIT_{WAIT|TEST}{ANY|SOME|ALL}` will occur for such a request). In this case, the callback function will be called either in the MPI call `MPI_REQUEST_FREE(request)`, or in the MPI call `MPI_GREQUEST_COMPLETE(request)`, whichever happens last, i.e., in this case the actual freeing code is executed as soon as both calls `MPI_REQUEST_FREE` and `MPI_GREQUEST_COMPLETE` have occurred. The request is not deallocated until after `free_fn` completes. Note that `free_fn` will be invoked only once per request by a correct program.

*Advice to users.* Calling `MPI_REQUEST_FREE(request)` will cause the request handle to be set to `MPI_REQUEST_NULL`. This handle to the generalized request is no longer valid. However, user copies of this handle are valid until after `free_fn` completes since MPI does not deallocate the object until then. Since `free_fn` is not called until after `MPI_GREQUEST_COMPLETE`, the user copy of the handle can be used to make this call. Users should note that MPI will deallocate the object after `free_fn` executes. At this point, user copies of the request handle no longer point to a valid request. MPI will not set user copies to `MPI_REQUEST_NULL` in this case, so it is up to the user to avoid accessing this stale handle. This is a special case [where]in which MPI defers deallocating the object until a later time that is known by the user. (*End of advice to users.*)

In C, the cancel function is  

```
typedef int MPI_Grequest_cancel_function(void *extra_state, int complete);
```

in Fortran

```
SUBROUTINE GREQUEST_CANCEL_FUNCTION(EXTRA_STATE, COMPLETE, IERROR)
  INTEGER IERROR
  INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
  LOGICAL COMPLETE
```

and in C++

```
{typedef int MPI::Grequest::Cancel_function(void* extra_state,
      bool complete); (binding deprecated, see Section 15.2)}
```

[cancel\_fn] The `cancel_fn` function is invoked to start the cancelation of a generalized request. It is called by `MPI_CANCEL(request)`. MPI passes [to the callback function `complete=true`] `complete=true` to the callback function if `MPI_GREQUEST_COMPLETE` was already called on the request, and `complete=false` otherwise.

All callback functions return an error code. The code is passed back and dealt with as appropriate for the error code by the MPI function that invoked the callback function. For example, if error codes are returned then the error code returned by the callback function will be returned by the MPI function that invoked the callback function. In the case of an `MPI_{WAIT|TEST}{ANY}` call that invokes both `query_fn` and `free_fn`, the MPI call will return the error code returned by the last callback, namely `free_fn`. If one or more of the requests in a call to `MPI_{WAIT|TEST}{SOME|ALL}` failed, then the MPI call will return `MPI_ERR_IN_STATUS`. In such a case, if the MPI call was passed an array of statuses, then MPI will return in each of the statuses that correspond to a completed generalized request the error code returned by the corresponding invocation of its `free_fn` callback function.

*Advice to users.* `query_fn` must **not** set the error field of `status` since `query_fn` may be called by `MPI_WAIT` or `MPI_TEST`, in which case the error field of `status` should not change. The MPI library knows the “context” in which `query_fn` is invoked and can decide correctly when to put in the error field of `status` the returned error code. *(End of advice to users.)*

*Advice to implementors.* A call to MPI\_GREQUEST\_COMPLETE may unblock a blocked user process/thread. The MPI library should ensure that the blocked user computation will resume. (*End of advice to implementors.*)

**Example 12.1** This example shows the code for a user-defined reduce operation on an `int` using a binary tree: each non-root node receives two messages, sums them, and sends them up. We assume that no status is returned and that the operation cannot be cancelled.



```

1  typedef struct {
2      MPI_Comm comm;
3      int tag;
4      int root;
5      int valin;
6      int *valout;
7      MPI_Request request;
8      } ARGS;
9
10
11 int myreduce(MPI_Comm comm, int tag, int root,
12             int valin, int *valout, MPI_Request *request)
13 {
14     ARGS *args;
15     pthread_t thread;
16
17     /* start request */
18     MPI_Grequest_start(query_fn, free_fn, cancel_fn, NULL, request);
19
20     args = (ARGS*)malloc(sizeof(ARGS));
21     args->comm = comm;
22     args->tag = tag;
23     args->root = root;
24     args->valin = valin;
25     args->valout = valout;
26     args->request = *request;
27
28     /* spawn thread to handle request */
29     /* The availability of the pthread_create call is system dependent */
30     pthread_create(&thread, NULL, reduce_thread, args);
31
32     return MPI_SUCCESS;
33 }
34
35 /* thread code */
36 void* reduce_thread(void *ptr)
37 {
38     int lchild, rchild, parent, lval, rval, val;
39     MPI_Request req[2];
40     ARGS *args;
41
42     args = (ARGS*)ptr;
43
44     /* compute left,right child and parent in tree; set
45        to MPI_PROC_NULL if does not exist */
46     /* code not shown */
47     ...
48

```

```

MPI_Irecv(&lval, 1, MPI_INT, lchild, args->tag, args->comm, &req[0]); 1
MPI_Irecv(&rval, 1, MPI_INT, rchild, args->tag, args->comm, &req[1]); 2
MPI_Waitall(2, req, MPI_STATUSES_IGNORE); 3
val = lval + args->valin + rval; 4
MPI_Send( &val, 1, MPI_INT, parent, args->tag, args->comm ); 5
if (parent == MPI_PROC_NULL) *(args->valout) = val; 6
MPI_Grequest_complete((args->request)); 7
free(ptr); 8
return(NULL); 9
} 10
11
int query_fn(void *extra_state, MPI_Status *status) 12
{ 13
    /* always send just one int */ 14
    MPI_Status_set_elements(status, MPI_INT, 1); 15
    /* can never cancel so always true */ 16
    MPI_Status_set_cancelled(status, 0); 17
    /* choose not to return a value for this */ 18
    status->MPI_SOURCE = MPI_UNDEFINED; 19
    /* tag has no meaning for this generalized request */ 20
    status->MPI_TAG = MPI_UNDEFINED; 21
    /* this generalized request never fails */ 22
    return MPI_SUCCESS; 23
} 24
25
26
int free_fn(void *extra_state) 27
{ 28
    /* this generalized request does not need to do any freeing */ 29
    /* as a result it never fails here */ 30
    return MPI_SUCCESS; 31
} 32
33
34
int cancel_fn(void *extra_state, int complete) 35
{ 36
    /* This generalized request does not support cancelling. 37
       Abort if not already done. If done then treat as if cancel failed.*/ 38
    if (!complete) { 39
        fprintf(stderr, 40
            "Cannot cancel generalized request - aborting program\n"); 41
        MPI_Abort(MPI_COMM_WORLD, 99); 42
    } 43
    return MPI_SUCCESS; 44
} 45
46
47
48

```

## 12.3 Associating Information with Status

MPI supports several different types of requests besides those for point-to-point operations. These range from MPI calls for I/O to generalized requests. It is desirable to allow these calls *to use* the same request *mechanism. This mechanism, which* allows one to wait or test on different types of requests. However, `MPI_{TEST|WAIT}{ANY|SOME|ALL}` returns a status with information about the request. With the generalization of requests, one needs to define what information will be returned in the status object.

Each MPI call fills in the appropriate fields in the status object. Any unused fields will have undefined values. A call to `MPI_{TEST|WAIT}{ANY|SOME|ALL}` can modify any of the fields in the status object. Specifically, it can modify fields that are undefined. The fields with meaningful *value* *values* for a given request are defined in the sections with the new request.

Generalized requests raise additional considerations. Here, the user provides the functions to deal with the request. Unlike other MPI calls, the user needs to provide the information to be returned in status. The status argument is provided directly to the callback function where the status needs to be set. Users can directly set the values in 3 of the 5 status values. The count and cancel fields are opaque. To overcome this, these calls are provided:

`MPI_STATUS_SET_ELEMENTS(status, datatype, count)`

INOUT	status	status with which to associate count (Status)
IN	datatype	datatype associated with count (handle)
IN	count	number of elements to associate with status (integer)

```
int MPI_Status_set_elements(MPI_Status *status, MPI_Datatype datatype,
                           int count)
```

```
MPI_STATUS_SET_ELEMENTS(STATUS, DATATYPE, COUNT, IERROR)
    INTEGER STATUS(MPI_STATUS_SIZE), DATATYPE, COUNT, IERROR
```

```
{void MPI::Status::Set_elements(const MPI::Datatype& datatype, int
                                count) (binding deprecated, see Section 15.2) }
```

`MPI_STATUS_SET_ELEMENTS_X(status, datatype, count)`

INOUT	status	status with which to associate count (Status)
IN	datatype	datatype associated with count (handle)
IN	count	number of elements to associate with status (integer)

```
int MPI_Status_set_elements_x(MPI_Status *status, MPI_Datatype datatype,
                              MPI_Count count)
```

```
MPI_STATUS_SET_ELEMENTS_X(STATUS, DATATYPE, COUNT, IERROR)
    INTEGER STATUS(MPI_STATUS_SIZE), DATATYPE, IERROR
    INTEGER (KIND=MPI_COUNT_KIND) COUNT
```

This call modifies the opaque part of `status` so that a call to `MPI_GET_ELEMENTS` or `MPI_GET_ELEMENTS_X` will return `count`. `MPI_GET_COUNT` will return a compatible value.

*Rationale.* The number of elements is set instead of the count because the former can deal with a nonintegral number of datatypes. (*End of rationale.*)

A subsequent call to `MPI_GET_COUNT(status, datatype, count)` [ or to ], `MPI_GET_ELEMENTS(status, datatype, count)`, or `MPI_GET_ELEMENTS_X(status, datatype, count)` must use a `datatype` argument that has the same type signature as the `datatype` argument that was used in the call to `MPI_STATUS_SET_ELEMENTS` or `MPI_STATUS_SET_ELEMENTS_X`.

*Rationale.* [This]The requirement of matching type signatures for these calls is similar to the restriction that holds when `count` is set by a receive operation: in that case, the calls to `MPI_GET_COUNT` and `MPI_GET_ELEMENTS` must use a `datatype` with the same signature as the `datatype` used in the receive call. (*End of rationale.*)

`MPI_STATUS_SET_CANCELLED(status, flag)`

INOUT	<code>status</code>	status with which to associate cancel flag (Status)
IN	<code>flag</code>	if true indicates request was cancelled (logical)

`int MPI_Status_set_cancelled(MPI_Status *status, int flag)`

`MPI_STATUS_SET_CANCELLED(STATUS, FLAG, IERROR)`  
`INTEGER STATUS(MPI_STATUS_SIZE), IERROR`  
`LOGICAL FLAG`

`{void MPI::Status::Set_cancelled(bool flag) (binding deprecated, see Section 15.2)`  
`}`

If `flag` is set to true then a subsequent call to `MPI_TEST_CANCELLED(status, flag)` will also return `flag = true`, otherwise it will return false.

*Advice to users.* Users are advised not to reuse the status fields for values other than those for which they were intended. Doing so may lead to unexpected results when using the status object. For example, calling `MPI_GET_ELEMENTS` may cause an error if the value is out of range or it may be impossible to detect such an error. The `extra_state` argument provided with a generalized request can be used to return information that does not logically belong in status. Furthermore, modifying the values in a status set internally by MPI, e.g., `MPI_RECV`, may lead to unpredictable results and is strongly discouraged. (*End of advice to users.*)

## 12.4 MPI and Threads

This section specifies the interaction between MPI calls and threads. The section lists minimal requirements for **thread compliant** MPI implementations and defines functions

that can be used for initializing the thread environment. MPI may be implemented in environments where threads are not supported or perform poorly. Therefore, it is not required that all MPI implementations fulfill all the requirements specified in this section.

This section generally assumes a thread package similar to POSIX threads [33], but the syntax and semantics of thread calls are not specified here — these are beyond the scope of this document.

### 12.4.1 General

In a thread-compliant implementation, an MPI process is a process that may be multi-threaded. Each thread can issue MPI calls; however, threads are not separately addressable: a rank in a send or receive call identifies a process, not a thread. A message sent to a process can be received by any thread in this process.

*Rationale.* This model corresponds to the POSIX model of interprocess communication: the fact that a process is multi-threaded, rather than single-threaded, does not affect the external interface of this process. MPI implementations [where]in which MPI ‘processes’ are POSIX threads inside a single POSIX process are not thread-compliant by this definition (indeed, their “processes” are single-threaded). (*End of rationale.*)

*Advice to users.* It is the user’s responsibility to prevent races when threads within the same application post conflicting communication calls. The user can make sure that two threads in the same process will not issue conflicting communication calls by using distinct communicators at each thread. (*End of advice to users.*)

The two main requirements for a thread-compliant implementation are listed below.

1. All MPI calls are *thread-safe*, i.e., two concurrently running threads may make MPI calls and the outcome will be as if the calls executed in some order, even if their execution is interleaved.
2. Blocking MPI calls will block the calling thread only, allowing another thread to execute, if available. The calling thread will be blocked until the event on which it is waiting occurs. Once the blocked communication is enabled and can proceed, then the call will complete and the thread will be marked runnable, within a finite time. A blocked thread will not prevent progress of other runnable threads on the same process, and will not prevent them from executing MPI calls.

**Example 12.2** Process 0 consists of two threads. The first thread executes a blocking send call `MPI_Send(buff1, count, type, 0, 0, comm)`, whereas the second thread executes a blocking receive call `MPI_Recv(buff2, count, type, 0, 0, comm, &status)`, i.e., the first thread sends a message that is received by the second thread. This communication should always succeed. According to the first requirement, the execution will correspond to some interleaving of the two calls. According to the second requirement, a call can only block the calling thread and cannot prevent progress of the other thread. If the send call went ahead of the receive call, then the sending thread may block, but this will not prevent the receiving thread from executing. Thus, the receive call will occur. Once both calls occur, the communication is enabled and both calls will complete. On the other hand, a single-threaded process that posts a send, followed by a matching receive, may deadlock.

The progress requirement for multithreaded implementations is stronger, as a blocked call cannot prevent progress in other threads.

*Advice to implementors.* MPI calls can be made thread-safe by executing only one at a time, e.g., by protecting MPI code with one process-global lock. However, blocked operations cannot hold the lock, as this would prevent progress of other threads in the process. The lock is held only for the duration of an atomic, locally-completing suboperation such as posting a send or completing a send, and is released in between. Finer locks can provide more concurrency, at the expense of higher locking overheads. Concurrency can also be achieved by having some of the MPI protocol executed by separate server threads. (*End of advice to implementors.*)

### 12.4.2 Clarifications

**Initialization and Completion** The call to MPI\_FINALIZE should occur on the same thread that initialized MPI. We call this thread the **main thread**. The call should occur only after all the process threads have completed their MPI calls, and have no pending communications or I/O operations.

*Rationale.* This constraint simplifies implementation. (*End of rationale.*)

**Multiple threads completing the same request.** A program where two threads block, waiting on the same request, is erroneous. Similarly, the same request cannot appear in the array of requests of two concurrent MPI\_{WAIT|TEST}{ANY|SOME|ALL} calls. In MPI, a request can only be completed once. Any combination of wait or test [which]that violates this rule is erroneous.

*Rationale.* [This]This restriction is consistent with the view that a multithreaded execution corresponds to an interleaving of the MPI calls. In a single threaded implementation, once a wait is posted on a request the request handle will be nullified before it is possible to post a second wait on the same handle. With threads, an MPI\_WAIT{ANY|SOME|ALL} may be blocked without having nullified its request(s) so it becomes the user's responsibility to avoid using the same request in an MPI\_WAIT on another thread. This constraint also simplifies implementation, as only one thread will be blocked on any communication or I/O event. (*End of rationale.*)

**Probe** A receive call that uses source and tag values returned by a preceding call to MPI\_PROBE or MPI\_IProbe will receive the message matched by the probe call only if there was no other matching receive after the probe and before that receive. In a multithreaded environment, it is up to the user to enforce this condition using suitable mutual exclusion logic. This can be enforced by making sure that each communicator is used by only one thread on each process.

**Collective calls** Matching of collective calls on a communicator, window, or file handle is done according to the order in which the calls are issued at each process. If concurrent threads issue such calls on the same communicator, window or file handle, it is up to the user to make sure the calls are correctly ordered, using interthread synchronization.

*Advice to users.* With three concurrent threads in each MPI process of a communicator `comm`, it is allowed that thread A in each MPI process calls a collective operation on `comm`, thread B calls a file operation on an existing filehandle that was formerly opened on `comm`, and thread C invokes one-sided operations on an existing window handle that was also formerly created on `comm`. (*End of advice to users.*)

*Rationale.* As already specified in `MPI_FILE_OPEN` and `MPI_WIN_CREATE`, a file handle and a window handle inherit only the group of processes of the underlying communicator, but not the communicator itself. Accesses to communicators, window handles and file handles cannot affect one another. (*End of rationale.*)

*Advice to implementors.* [Advice to implementors.] If the implementation of file or window operations internally uses MPI communication then a duplicated communicator may be cached on the file or window object. (*End of advice to implementors.*)

**Exception handlers** An exception handler does not necessarily execute in the context of the thread that made the exception-raising MPI call; the exception handler may be executed by a thread that is distinct from the thread that will return the error code.

*Rationale.* The MPI implementation may be multithreaded, so that part of the communication protocol may execute on a thread that is distinct from the thread that made the MPI call. The design allows the exception handler to be executed on the thread where the exception occurred. (*End of rationale.*)

**Interaction with signals and cancellations** The outcome is undefined if a thread that executes an MPI call is cancelled (by another thread), or if a thread catches a signal while executing an MPI call. However, a thread of an MPI process may terminate, and may catch signals or be cancelled by another thread when not executing MPI calls.

*Rationale.* Few C library functions are signal safe, and many have cancellation points — points [where]at which the thread executing them may be cancelled. The above restriction simplifies implementation (no need for the MPI library to be “async-cancel-safe” or [“async-signal-safe.”] “async-signal-safe”). (*End of rationale.*)

*Advice to users.* Users can catch signals in separate, non-MPI threads (e.g., by masking signals on MPI calling threads, and unmasking them in one or more non-MPI threads). A good programming practice is to have a distinct thread blocked in a call to `sigwait` for each user expected signal that may occur. Users must not catch signals used by the MPI implementation; as each MPI implementation is required to document the signals used internally, users can avoid these signals. (*End of advice to users.*)

*Advice to implementors.* The MPI library should not invoke library calls that are not thread safe, if multiple threads execute. (*End of advice to implementors.*)

### 12.4.3 Initialization

The following function may be used to initialize MPI, and initialize the MPI thread environment, instead of `MPI_INIT`.

MPI\_INIT\_THREAD(required, provided)

IN	required	desired level of thread support (integer)
OUT	provided	provided level of thread support (integer)

```
int MPI_Init_thread(int *argc, char *((*argv)[]), int required,
                    int *provided)
```

```
MPI_INIT_THREAD(REQUIRED, PROVIDED, IERROR)
INTEGER REQUIRED, PROVIDED, IERROR
```

```
{int MPI::Init_thread(int& argc, char**& argv, int required) (binding
    deprecated, see Section 15.2) }
```

```
{int MPI::Init_thread(int required) (binding deprecated, see Section 15.2) }
```

*Advice to users.* In C and C++, the passing of `argc` and `argv` is [optional.]optional, as with `MPI_INIT` as discussed in Section 8.7. In C, [this is accomplished by passing the appropriate null pointer.] null pointers may be passed in their place. In C++, [this is accomplished with two separate bindings to cover these two cases. This is as with `MPI_INIT` as discussed in Section 8.7.]two separate bindings support this choice. (End of advice to users.)

This call initializes MPI in the same way that a call to `MPI_INIT` would. In addition, it initializes the thread environment. The argument `required` is used to specify the desired level of thread support. The possible values are listed in increasing order of thread support.

**MPI\_THREAD\_SINGLE** Only one thread will execute.

**MPI\_THREAD\_FUNNELED** The process may be multi-threaded, but the application must ensure that only the main thread makes MPI calls (for the definition of main thread, see `MPI_IS_THREAD_MAIN` on page 419).

**MPI\_THREAD\_SERIALIZED** The process may be multi-threaded, and multiple threads may make MPI calls, but only one at a time: MPI calls are not made concurrently from two distinct threads (all MPI calls are “serialized”).

**MPI\_THREAD\_MULTIPLE** Multiple threads may call MPI, with no restrictions.

These values are monotonic; i.e., `MPI_THREAD_SINGLE < MPI_THREAD_FUNNELED < MPI_THREAD_SERIALIZED < MPI_THREAD_MULTIPLE`.

Different processes in `MPI_COMM_WORLD` may require different levels of thread support.

The call returns in `provided` information about the actual level of thread support that will be provided by MPI. It can be one of the four values listed above.

The level(s) of thread support that can be provided by `MPI_INIT_THREAD` will depend on the implementation, and may depend on information provided by the user before the program started to execute (e.g., with arguments to `mpiexec`). If possible, the call will return `provided = required`. Failing this, the call will return the least supported level such that `provided > required` (thus providing a stronger level of support than required by the



user). Finally, if the user requirement cannot be satisfied, then the call will return in provided the highest supported level.

A **thread compliant** MPI implementation will be able to return provided = MPI\_THREAD\_MULTIPLE. Such an implementation may always return provided = MPI\_THREAD\_MULTIPLE, irrespective of the value of required. At the other extreme, an MPI library that is not thread compliant may always return provided = MPI\_THREAD\_SINGLE, irrespective of the value of required.

A call to MPI\_INIT has the same effect as a call to MPI\_INIT\_THREAD with a required = MPI\_THREAD\_SINGLE.

Vendors may provide (implementation dependent) means to specify the level(s) of thread support available when the MPI program is started, e.g., with arguments to mpiexec. This will affect the outcome of calls to MPI\_INIT and MPI\_INIT\_THREAD. Suppose, for example, that an MPI program has been started so that only MPI\_THREAD\_MULTIPLE is available. Then MPI\_INIT\_THREAD will return provided = MPI\_THREAD\_MULTIPLE, irrespective of the value of required; a call to MPI\_INIT will also initialize the MPI thread support level to MPI\_THREAD\_MULTIPLE. Suppose, on the other hand, that an MPI program has been started so that all four levels of thread support are available. Then, a call to MPI\_INIT\_THREAD will return provided = required; on the other hand, a call to MPI\_INIT will initialize the MPI thread support level to MPI\_THREAD\_SINGLE.

*Rationale.* Various optimizations are possible when MPI code is executed single-threaded, or is executed on multiple threads, but not concurrently: mutual exclusion code may be omitted. Furthermore, if only one thread executes, then the MPI library can use library functions that are not thread safe, without risking conflicts with user threads. Also, the model of one communication thread, multiple computation threads fits many applications well, e.g., if the process code is a sequential Fortran/C/C++ program with MPI calls that has been parallelized by a compiler for execution on an SMP node, in a cluster of SMPs, then the process computation is multi-threaded, but MPI calls will likely execute on a single thread.

The design accommodates a static specification of the thread support level, for environments that require static binding of libraries, and for compatibility for current multi-threaded MPI codes. (*End of rationale.*)

*Advice to implementors.* If provided is not MPI\_THREAD\_SINGLE then the MPI library should not invoke C/ C++/Fortran library calls that are not thread safe, e.g., in an environment where malloc is not thread safe, then malloc should not be used by the MPI library.

Some implementors may want to use different MPI libraries for different levels of thread support. They can do so using dynamic linking and selecting which library will be linked when MPI\_INIT\_THREAD is invoked. If this is not possible, then optimizations for lower levels of thread support will occur only when the level of thread support required is specified at link time. (*End of advice to implementors.*)

The following function can be used to query the current level of thread support.

MPI\_QUERY\_THREAD(provided)

OUT      provided      provided level of thread support (integer)

int MPI\_Query\_thread(int \*provided)

MPI\_QUERY\_THREAD(PROVIDED, IERROR)

INTEGER PROVIDED, IERROR

{int MPI::Query\_thread() *(binding deprecated, see Section 15.2)* }

The call returns in provided the current level of thread [support. This]support, which will be the value returned in provided by MPI\_INIT\_THREAD, if MPI was initialized by a call to MPI\_INIT\_THREAD().

MPI\_IS\_THREAD\_MAIN(flag)

OUT      flag      true if calling thread is main thread, false otherwise  
(logical)

int MPI\_Is\_thread\_main(int \*flag)

MPI\_IS\_THREAD\_MAIN(FLAG, IERROR)

LOGICAL FLAG

INTEGER IERROR

{bool MPI::Is\_thread\_main() *(binding deprecated, see Section 15.2)* }

This function can be called by a thread to [find out whether]determine if it is the main thread (the thread that called MPI\_INIT or MPI\_INIT\_THREAD).

All routines listed in this section must be supported by all MPI implementations.

*Rationale.* MPI libraries are required to provide these calls even if they do not support threads, so that portable code that contains invocations to these functions [be able to]can link correctly. MPI\_INIT continues to be supported so as to provide compatibility with current MPI codes. (*End of rationale.*)

*Advice to users.* It is possible to spawn threads before MPI is initialized, but no MPI call other than MPI\_INITIALIZED should be executed by these threads, until MPI\_INIT\_THREAD is invoked by one thread (which, thereby, becomes the main thread). In particular, it is possible to enter the MPI execution with a multi-threaded process.

The level of thread support provided is a global property of the MPI process that can be specified only once, when MPI is initialized on that process (or before). Portable third party libraries have to be written so as to accommodate any provided level of thread support. Otherwise, their usage will be restricted to specific level(s) of thread support. If such a library can run only with specific level(s) of thread support, e.g., only with MPI\_THREAD\_MULTIPLE, then MPI\_QUERY\_THREAD can be used to check whether the user initialized MPI to the correct level of thread support and, if not, raise an exception. (*End of advice to users.*)

1  
2  
3  
4  
5  
6  
7  
8  
9  
10  
11  
12  
13  
14  
15  
16  
17  
18  
19  
20  
21  
22  
23  
24  
25  
26  
27  
28  
29  
30  
31  
32  
33  
34  
35  
36  
37  
38  
39  
40  
41  
42  
43  
44  
45  
46  
47  
48

# Chapter 13

## I/O

### 13.1 Introduction

POSIX provides a model of a widely portable file system, but the portability and optimization needed for parallel I/O cannot be achieved with the POSIX interface.

The significant optimizations required for efficiency (e.g., grouping [39], collective buffering [6, 13, 40, 44, 51], and disk-directed I/O [35]) can only be implemented if the parallel I/O system provides a high-level interface supporting partitioning of file data among processes and a collective interface supporting complete transfers of global data structures between process memories and files. In addition, further efficiencies can be gained via support for asynchronous I/O, strided accesses, and control over physical file layout on storage devices (disks). The I/O environment described in this chapter provides these facilities.

Instead of defining I/O access modes to express the common patterns for accessing a shared file (broadcast, reduction, scatter, gather), we chose another approach in which data partitioning is expressed using derived datatypes. Compared to a limited set of predefined access patterns, this approach has the advantage of added flexibility and expressiveness.

#### 13.1.1 Definitions

**file** An MPI file is an ordered collection of typed data items. MPI supports random or sequential access to any integral set of these items. A file is opened collectively by a group of processes. All collective I/O calls on a file are collective over this group.

**displacement** A file *displacement* is an absolute byte position relative to the beginning of a file. The displacement defines the location where a *view* begins. Note that a “file displacement” is distinct from a “typemap displacement.”

**etype** An *etype* (*elementary* datatype) is the unit of data access and positioning. It can be any MPI predefined or derived datatype. Derived etypes can be constructed using any of the MPI datatype constructor routines, provided all resulting typemap displacements are non-negative and monotonically nondecreasing. Data access is performed in etype units, reading or writing whole data items of type etype. Offsets are expressed as a count of etypes; file pointers point to the beginning of etypes. Depending on context, the term “etype” is used to describe one of three aspects of an elementary datatype: a particular MPI type, a data item of that type, or the extent of that type.

**filetype** A *filetype* is the basis for partitioning a file among processes and defines a template for accessing the file. A filetype is either a single etype or a derived MPI datatype constructed from multiple instances of the same etype. In addition, the extent of any hole in the filetype must be a multiple of the etype's extent. The displacements in the typemap of the filetype are not required to be distinct, but they must be non-negative and monotonically nondecreasing.

**view** A *view* defines the current set of data visible and accessible from an open file as an ordered set of etypes. Each process has its own view of the file, defined by three quantities: a displacement, an etype, and a filetype. The pattern described by a filetype is repeated, beginning at the displacement, to define the view. The pattern of repetition is defined to be the same pattern that `MPI_TYPE_CONTIGUOUS` would produce if it were passed the filetype and an arbitrarily large count. Figure 13.1 shows how the tiling works; note that the filetype in this example must have explicit lower and upper bounds set in order for the initial and final holes to be repeated in the view. Views can be changed by the user during program execution. The default view is a linear byte stream (displacement is zero, etype and filetype equal to `MPI_BYTE`).

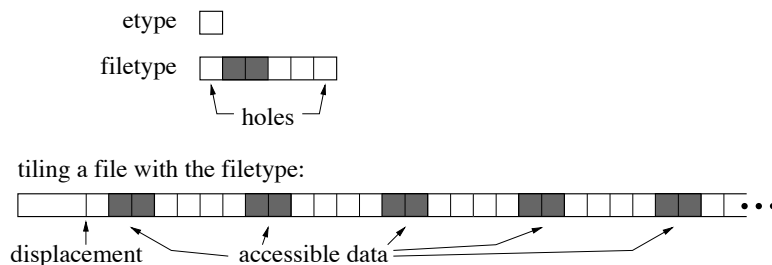


Figure 13.1: Etypes and filetypes

A group of processes can use complementary views to achieve a global data distribution such as a scatter/gather pattern (see Figure 13.2).




Figure 13.2: Partitioning a file among parallel processes

**offset** An *offset* is a position in the file relative to the current view, expressed as a count of etypes. Holes in the view's filetype are skipped when calculating this position. Offset 0 is the location of the first etype visible in the view (after skipping the displacement and any initial holes in the view). For example, an offset of 2 for process 1 in Figure 13.2 is the position of the 8th etype in the file after the displacement. An “explicit offset” is an offset that is used as a formal parameter in explicit data access routines.

**file size and end of file** The *size* of an MPI file is measured in bytes from the beginning of the file. A newly created file has a size of zero bytes. Using the size as an absolute displacement gives the position of the byte immediately following the last byte in the file. For any given view, the *end of file* is the offset of the first etype accessible in the current view starting after the last byte in the file.

**file pointer** A *file pointer* is an implicit offset maintained by MPI. “Individual file pointers” are file pointers that are local to each process that opened the file. A “shared file pointer” is a file pointer that is shared by the group of processes that opened the file.

**file handle** A *file handle* is an opaque object created by `MPI_FILE_OPEN` and freed by `MPI_FILE_CLOSE`. All operations on an open file reference the file through the file handle.

## 13.2 File Manipulation

### 13.2.1 Opening a File

`MPI_FILE_OPEN(comm, filename, amode, info, fh)`

IN	<code>comm</code>	communicator (handle)
IN	<code>filename</code>	name of file to open (string)
IN	<code>amode</code>	file access mode (integer)
IN	<code>info</code>	info object (handle)
OUT	<code>fh</code>	new file handle (handle)

```
int MPI_File_open(MPI_Comm comm, char *filename, int amode, MPI_Info info,
                  MPI_File *fh)
```

```
MPI_FILE_OPEN(COMM, FILENAME, AMODE, INFO, FH, IERROR)
CHARACTER*(*) FILENAME
INTEGER COMM, AMODE, INFO, FH, IERROR
```

```
{static MPI::File MPI::File::Open(const MPI::Intracomm& comm,
    const char* filename, int amode, const MPI::Info& info) (binding
    deprecated, see Section 15.2) }
```

`MPI_FILE_OPEN` opens the file identified by the file name `filename` on all processes in the `comm` communicator group. `MPI_FILE_OPEN` is a collective routine: all processes must provide the same value for `amode`, and all processes must provide `filenames` that reference the same file. (Values for `info` may vary.) `comm` must be an intracommunicator; it is erroneous to pass an intercommunicator to `MPI_FILE_OPEN`. Errors in `MPI_FILE_OPEN` are raised using the default file error handler (see Section 13.7, page 479). A process can open a file independently of other processes by using the `MPI_COMM_SELF` communicator. The file handle returned, `fh`, can be subsequently used to access the file until the file is closed using `MPI_FILE_CLOSE`. Before calling `MPI_FINALIZE`, the user is required to close (via `MPI_FILE_CLOSE`) all files that were opened with `MPI_FILE_OPEN`. Note that the

communicator `comm` is unaffected by `MPI_FILE_OPEN` and continues to be usable in all MPI routines (e.g., `MPI_SEND`). Furthermore, the use of `comm` will not interfere with I/O behavior.

The format for specifying the file name in the `filename` argument is implementation dependent and must be documented by the implementation.

*Advice to implementors.* An implementation may require that `filename` include a string or strings specifying additional information about the file. Examples include the type of filesystem (e.g., a prefix of `ufs:`), a remote hostname (e.g., a prefix of `machine.univ.edu:`), or a file password (e.g., a suffix of `/PASSWORD=SECRET`). (*End of advice to implementors.*)

*Advice to users.* On some implementations of MPI, the file namespace may not be identical from all processes of all applications. For example, `"/tmp/foo"` may denote different files on different processes, or a single file may have many names, dependent on process location. The user is responsible for ensuring that a single file is referenced by the `filename` argument, as it may be impossible for an implementation to detect this type of namespace error. (*End of advice to users.*)

Initially, all processes view the file as a linear byte stream, and each process views data in its own native representation (no data representation conversion is performed). (POSIX files are linear byte streams in the native representation.) The file view can be changed via the `MPI_FILE_SET_VIEW` routine.

The following access modes are supported (specified in `amode`, a bit vector OR of the following integer constants):

- `MPI_MODE_RDONLY` — read only,
- `MPI_MODE_RDWR` — reading and writing,
- `MPI_MODE_WRONLY` — write only,
- `MPI_MODE_CREATE` — create the file if it does not exist,
- `MPI_MODE_EXCL` — error if creating file that already exists,
- `MPI_MODE_DELETE_ON_CLOSE` — delete file on close,
- `MPI_MODE_UNIQUE_OPEN` — file will not be concurrently opened elsewhere,
- `MPI_MODE_SEQUENTIAL` — file will only be accessed sequentially,
- `MPI_MODE_APPEND` — set initial position of all file pointers to end of file.

*Advice to users.* C/C++ users can use bit vector OR (`|`) to combine these constants; Fortran 90 users can use the bit vector `IOR` intrinsic. Fortran 77 users can use (non-portably) bit vector `IOR` on systems that support it. Alternatively, Fortran users can portably use integer addition to OR the constants (each constant should appear at most once in the addition.). (*End of advice to users.*)

*Advice to implementors.* The values of these constants must be defined such that the bitwise OR and the sum of any distinct set of these constants is equivalent. (*End of advice to implementors.*)

The modes `MPI_MODE_RDONLY`, `MPI_MODE_RDWR`, `MPI_MODE_WRONLY`, `MPI_MODE_CREATE`, and `MPI_MODE_EXCL` have identical semantics to their POSIX counterparts [33]. Exactly one of `MPI_MODE_RDONLY`, `MPI_MODE_RDWR`, or `MPI_MODE_WRONLY`, must be specified. It is erroneous to specify `MPI_MODE_CREATE` or `MPI_MODE_EXCL` in conjunction with `MPI_MODE_RDONLY`; it is erroneous to specify `MPI_MODE_SEQUENTIAL` together with `MPI_MODE_RDWR`.

The `MPI_MODE_DELETE_ON_CLOSE` mode causes the file to be deleted (equivalent to performing an `MPI_FILE_DELETE`) when the file is closed.

The `MPI_MODE_UNIQUE_OPEN` mode allows an implementation to optimize access by eliminating the overhead of file locking. It is erroneous to open a file in this mode unless the file will not be concurrently opened elsewhere.

*Advice to users.* For MPI\_MODE\_UNIQUE\_OPEN, *not opened elsewhere* includes both inside and outside the MPI environment. In particular, one needs to be aware of potential external events which may open files (e.g., automated backup facilities). When MPI\_MODE\_UNIQUE\_OPEN is specified, the user is responsible for ensuring that no such external events take place. (*End of advice to users.*)

The `MPI_MODE_SEQUENTIAL` mode allows an implementation to optimize access to some sequential devices (tapes and network streams). It is erroneous to attempt non-sequential access to a file that has been opened in this mode.

Specifying `MPI_MODE_APPEND` only guarantees that all shared and individual file pointers are positioned at the initial end of file when `MPI_FILE_OPEN` returns. Subsequent positioning of file pointers is application dependent. In particular, the implementation does not ensure that all writes are appended.

Errors related to the access mode are raised in the class `MPI_ERR_AMODE`.

The `info` argument is used to provide information regarding file access patterns and file system specifics (see Section 13.2.8, page 430). The constant `MPI_INFO_NULL` can be used when no `info` needs to be specified.

*Advice to users.* Some file attributes are inherently implementation dependent (e.g., file permissions). These attributes must be set using either the `info` argument or facilities outside the scope of MPI. (*End of advice to users.*)

Files are opened by default using nonatomic mode file consistency semantics (see Section 13.6.1, page 469). The more stringent atomic mode consistency semantics, required for atomicity of conflicting accesses, can be set using `MPI_FILE_SET_ATOMICITY`.

### 13.2.2 Closing a File

## MPI\_FILE\_CLOSE(fh)

INOUT	fh	file handle (handle)
-------	----	----------------------

```
int MPI_File_close(MPI_File *fh)
```

MPI\_FILE\_CLOSE(FH, IERROR)

INTEGER FH, IERROR



```
1 {void MPI::File::Close() (binding deprecated, see Section 15.2) }
```

```
2
3 MPI_FILE_CLOSE first synchronizes file state (equivalent to performing an
4 MPI_FILE_SYNC), then closes the file associated with fh. The file is deleted if it was
5 opened with access mode MPI_MODE_DELETE_ON_CLOSE (equivalent to performing an
6 MPI_FILE_DELETE). MPI_FILE_CLOSE is a collective routine.
```

```
7
8 Advice to users. If the file is deleted on close, and there are other processes currently
9 accessing the file, the status of the file and the behavior of future accesses by these
10 processes are implementation dependent. (End of advice to users.)
```

```
11
12 The user is responsible for ensuring that all outstanding nonblocking requests and
13 split collective operations associated with fh made by a process have completed before that
14 process calls MPI_FILE_CLOSE.
```

```
15 The MPI_FILE_CLOSE routine deallocates the file handle object and sets fh to
16 MPI_FILE_NULL.
```

### 17 13.2.3 Deleting a File

```
18
19
20 MPI_FILE_DELETE(filename, info)
```

```
22 IN filename name of file to delete (string)
23 IN info info object (handle)
```

```
25
26 int MPI_File_delete(char *filename, MPI_Info info)
```

```
27 MPI_FILE_DELETE(FILENAME, INFO, IERROR)
```

```
28 CHARACTER*(*) FILENAME
29 INTEGER INFO, IERROR
```

```
30
31 {static void MPI::File::Delete(const char* filename,
32 const MPI::Info& info) (binding deprecated, see Section 15.2) }
```

```
33
34 MPI_FILE_DELETE deletes the file identified by the file name filename. If the file does
35 not exist, MPI_FILE_DELETE raises an error in the class MPI_ERR_NO_SUCH_FILE.
```

```
36 The info argument can be used to provide information regarding file system specifics
37 (see Section 13.2.8, page 430). The constant MPI_INFO_NULL refers to the null info, and
38 can be used when no info needs to be specified.
```

```
39 If a process currently has the file open, the behavior of any access to the file (as well
40 as the behavior of any outstanding accesses) is implementation dependent. In addition,
41 whether an open file is deleted or not is also implementation dependent. If the file is not
42 deleted, an error in the class MPI_ERR_FILE_IN_USE or MPI_ERR_ACCESS will be raised.
43 Errors are raised using the default error handler (see Section 13.7, page 479).
```

## 13.2.4 Resizing a File

`MPI_FILE_SET_SIZE(fh, size)`

INOUT	fh	file handle (handle)
IN	size	size to truncate or expand file (integer)

`int MPI_File_set_size(MPI_File fh, MPI_Offset size)`

`MPI_FILE_SET_SIZE(FH, SIZE, IERROR)`

INTEGER FH, IERROR

INTEGER(KIND=MPI\_OFFSET\_KIND) SIZE

`{void MPI::File::Set_size(MPI::Offset size) (binding deprecated, see Section 15.2)}`

`MPI_FILE_SET_SIZE` resizes the file associated with the file handle `fh`. `size` is measured in bytes from the beginning of the file. `MPI_FILE_SET_SIZE` is collective; all processes in the group must pass identical values for `size`.

If `size` is smaller than the current file size, the file is truncated at the position defined by `size`. The implementation is free to deallocate file blocks located beyond this position.

If `size` is larger than the current file size, the file size becomes `size`. Regions of the file that have been previously written are unaffected. The values of data in the new regions in the file (those locations with displacements between old file size and `size`) are undefined. It is implementation dependent whether the `MPI_FILE_SET_SIZE` routine allocates file space—use `MPI_FILE_PREALLOCATE` to force file space to be reserved.

`MPI_FILE_SET_SIZE` does not affect the individual file pointers or the shared file pointer. If `MPI_MODE_SEQUENTIAL` mode was specified when the file was opened, it is erroneous to call this routine.

*Advice to users.* It is possible for the file pointers to point beyond the end of file after a `MPI_FILE_SET_SIZE` operation truncates a file. This is legal, and equivalent to seeking beyond the current end of file. (*End of advice to users.*)

All nonblocking requests and split collective operations on `fh` must be completed before calling `MPI_FILE_SET_SIZE`. Otherwise, calling `MPI_FILE_SET_SIZE` is erroneous. As far as consistency semantics are concerned, `MPI_FILE_SET_SIZE` is a write operation that conflicts with operations that access bytes at displacements between the old and new file sizes (see Section 13.6.1, page 469).

## 13.2.5 Preallocating Space for a File

`MPI_FILE_PREALLOCATE(fh, size)`

INOUT	fh	file handle (handle)
IN	size	size to preallocate file (integer)

```

1  int MPI_File_preallocate(MPI_File fh, MPI_Offset size)
2
3  MPI_FILE_PREALLOCATE(FH, SIZE, IERROR)
4      INTEGER FH, IERROR
5      INTEGER(KIND=MPI_OFFSET_KIND) SIZE
6
7  {void MPI::File::Preallocate(MPI::Offset size) (binding deprecated, see
8      Section 15.2) }

```

MPI\_FILE\_PREALLOCATE ensures that storage space is allocated for the first `size` bytes of the file associated with `fh`. MPI\_FILE\_PREALLOCATE is collective; all processes in the group must pass identical values for `size`. Regions of the file that have previously been written are unaffected. For newly allocated regions of the file, MPI\_FILE\_PREALLOCATE has the same effect as writing undefined data. If `size` is larger than the current file size, the file size increases to `size`. If `size` is less than or equal to the current file size, the file size is unchanged.

The treatment of file pointers, pending nonblocking accesses, and file consistency is the same as with MPI\_FILE\_SET\_SIZE. If MPI\_MODE\_SEQUENTIAL mode was specified when the file was opened, it is erroneous to call this routine.

*Advice to users.* In some implementations, file preallocation may be expensive. (*End of advice to users.*)

### 13.2.6 Querying the Size of a File

```

26  MPI_FILE_GET_SIZE(fh, size)
27
28      IN          fh          file handle (handle)
29      OUT         size        size of the file in bytes (integer)

```

```

31  int MPI_File_get_size(MPI_File fh, MPI_Offset *size)
32
33  MPI_FILE_GET_SIZE(FH, SIZE, IERROR)
34      INTEGER FH, IERROR
35      INTEGER(KIND=MPI_OFFSET_KIND) SIZE
36
37  {MPI::Offset MPI::File::Get_size() const (binding deprecated, see Section 15.2) }

```

MPI\_FILE\_GET\_SIZE returns, in `size`, the current size in bytes of the file associated with the file handle `fh`. As far as consistency semantics are concerned, MPI\_FILE\_GET\_SIZE is a data access operation (see Section 13.6.1, page 469).

## 13.2.7 Querying File Parameters

`MPI_FILE_GET_GROUP(fh, group)`

IN	fh	file handle (handle)
OUT	group	group which opened the file (handle)

`int MPI_File_get_group(MPI_File fh, MPI_Group *group)`

`MPI_FILE_GET_GROUP(FH, GROUP, IERROR)`  
`INTEGER FH, GROUP, IERROR`

`{MPI::Group MPI::File::Get_group() const(binding deprecated, see Section 15.2) }`

`MPI_FILE_GET_GROUP` returns a duplicate of the group of the communicator used to open the file associated with `fh`. The group is returned in `group`. The user is responsible for freeing `group`.

`MPI_FILE_GET_AMODE(fh, amode)`

IN	fh	file handle (handle)
OUT	amode	file access mode used to open the file (integer)

`int MPI_File_get_amode(MPI_File fh, int *amode)`

`MPI_FILE_GET_AMODE(FH, AMODE, IERROR)`  
`INTEGER FH, AMODE, IERROR`

`{int MPI::File::Get_amode() const(binding deprecated, see Section 15.2) }`

`MPI_FILE_GET_AMODE` returns, in `amode`, the access mode of the file associated with `fh`.

**Example 13.1** In Fortran 77, decoding an `amode` bit vector will require a routine such as the following:

```

      SUBROUTINE BIT_QUERY(TEST_BIT, MAX_BIT, AMODE, BIT_FOUND)
!
!   TEST IF THE INPUT TEST_BIT IS SET IN THE INPUT AMODE
!   IF SET, RETURN 1 IN BIT_FOUND, 0 OTHERWISE
!
      INTEGER TEST_BIT, AMODE, BIT_FOUND, CP_AMODE, HIFOUND
      BIT_FOUND = 0
      CP_AMODE = AMODE
100  CONTINUE
      LBIT = 0
      HIFOUND = 0
      DO 20 L = MAX_BIT, 0, -1
          MATCHER = 2**L

```

```

1      IF (CP_AMODE .GE. MATCHER .AND. HIFOUND .EQ. 0) THEN
2          HIFOUND = 1
3          LBIT = MATCHER
4          CP_AMODE = CP_AMODE - MATCHER
5      END IF
6  20  CONTINUE
7      IF (HIFOUND .EQ. 1 .AND. LBIT .EQ. TEST_BIT) BIT_FOUND = 1
8      IF (BIT_FOUND .EQ. 0 .AND. HIFOUND .EQ. 1 .AND. &
9          CP_AMODE .GT. 0) GO TO 100
10     END

```

This routine could be called successively to decode `amode`, one bit at a time. For example, the following code fragment would check for `MPI_MODE_RDONLY`.

```

14     CALL BIT_QUERY(MPI_MODE_RDONLY, 30, AMODE, BIT_FOUND)
15     IF (BIT_FOUND .EQ. 1) THEN
16         PRINT *, ' FOUND READ-ONLY BIT IN AMODE=', AMODE
17     ELSE
18         PRINT *, ' READ-ONLY BIT NOT FOUND IN AMODE=', AMODE
19     END IF
20
21

```

### 13.2.8 File Info

Hints specified via info (see Section 9, page 331) allow a user to provide information such as file access patterns and file system specifics to direct optimization. Providing hints may enable an implementation to deliver increased I/O performance or minimize the use of system resources. However, hints do not change the semantics of any of the I/O interfaces. In other words, an implementation is free to ignore all hints. Hints are specified on a per file basis, in `MPI_FILE_OPEN`, `MPI_FILE_DELETE`, `MPI_FILE_SET_VIEW`, and `MPI_FILE_SET_INFO`, via the opaque info object. When an info object that specifies a subset of valid hints is passed to `MPI_FILE_SET_VIEW` or `MPI_FILE_SET_INFO`, there will be no effect on previously set or defaulted hints that the info does not specify.

*Advice to implementors.* It may happen that a program is coded with hints for one system, and later executes on another system that does not support these hints. In general, unsupported hints should simply be ignored. Needless to say, no hint can be mandatory. However, for each hint used by a specific implementation, a default value must be provided when the user does not specify a value for this hint. (*End of advice to implementors.*)

`MPI_FILE_SET_INFO(fh, info)`

INOUT	fh	file handle (handle)
IN	info	info object (handle)

`int MPI_File_set_info(MPI_File fh, MPI_Info info)`

`MPI_FILE_SET_INFO(FH, INFO, IERROR)`

INTEGER FH, INFO, IERROR

{void MPI::File::Set\_info(const MPI::Info& info) (*binding deprecated, see Section 15.2*) }

MPI\_FILE\_SET\_INFO sets new values for the hints of the file associated with fh. MPI\_FILE\_SET\_INFO is a collective routine. The info object may be different on each process, but any info entries that an implementation requires to be the same on all processes must appear with the same value in each process's info object.

*Advice to users.* Many info items that an implementation can use when it creates or opens a file cannot easily be changed once the file has been created or opened. Thus, an implementation may ignore hints issued in this call that it would have accepted in an open call. (*End of advice to users.*)

MPI\_FILE\_GET\_INFO(fh, info\_used)

IN	fh	file handle (handle)
OUT	info_used	new info object (handle)

int MPI\_File\_get\_info(MPI\_File fh, MPI\_Info \*info\_used)

MPI\_FILE\_GET\_INFO(FH, INFO\_USED, IERROR)

INTEGER FH, INFO\_USED, IERROR

{MPI::Info MPI::File::Get\_info() const (*binding deprecated, see Section 15.2*) }

MPI\_FILE\_GET\_INFO returns a new info object containing the hints of the file associated with fh. The current setting of all hints actually used by the system related to this open file is returned in info\_used. If no such hints exist, a handle to a newly created info object is returned that contains no key/value pair. The user is responsible for freeing info\_used via MPI\_INFO\_FREE.

*Advice to users.* The info object returned in info\_used will contain all hints currently active for this file. This set of hints may be greater or smaller than the set of hints passed in to MPI\_FILE\_OPEN, MPI\_FILE\_SET\_VIEW, and MPI\_FILE\_SET\_INFO, as the system may not recognize some hints set by the user, and may recognize other hints that the user has not set. (*End of advice to users.*)

### Reserved File Hints

Some potentially useful hints (info key values) are outlined below. The following key values are reserved. An implementation is not required to interpret these key values, but if it does interpret the key value, it must provide the functionality described. (For more details on “info,” see Section 9, page 331.)

These hints mainly affect access patterns and the layout of data on parallel I/O devices. For each hint name introduced, we describe the purpose of the hint, and the type of the hint value. The “[**SAME**]” annotation specifies that the hint values provided by all participating processes must be identical; otherwise the program is erroneous. In addition, some hints are

context dependent, and are only used by an implementation at specific times (e.g., `file_perm` is only useful during file creation).

**access\_style (comma separated list of strings):** This hint specifies the manner in which the file will be accessed until the file is closed or until the `access_style` key value is altered. The hint value is a comma separated list of the following: `read_once`, `write_once`, `read_mostly`, `write_mostly`, `sequential`, `reverse_sequential`, and `random`.

**collective\_buffering (boolean) [SAME]:** This hint specifies whether the application may benefit from collective buffering. Collective buffering is an optimization performed on collective accesses. Accesses to the file are performed on behalf of all processes in the group by a number of target nodes. These target nodes coalesce small requests into large disk accesses. Legal values for this key are `true` and `false`. Collective buffering parameters are further directed via additional hints: `cb_block_size`, `cb_buffer_size`, and `cb_nodes`.

**cb\_block\_size (integer) [SAME]:** This hint specifies the block size to be used for collective buffering file access. *Target nodes* access data in chunks of this size. The chunks are distributed among target nodes in a round-robin (CYCLIC) pattern.

**cb\_buffer\_size (integer) [SAME]:** This hint specifies the total buffer space that can be used for collective buffering on each target node, usually a multiple of `cb_block_size`.

**cb\_nodes (integer) [SAME]:** This hint specifies the number of target nodes to be used for collective buffering.

**chunked (comma separated list of integers) [SAME]:** This hint specifies that the file consists of a multidimensional array that is often accessed by subarrays. The value for this hint is a comma separated list of array dimensions, starting from the most significant one (for an array stored in row-major order, as in C, the most significant dimension is the first one; for an array stored in column-major order, as in Fortran, the most significant dimension is the last one, and array dimensions should be reversed).

**chunked\_item (comma separated list of integers) [SAME]:** This hint specifies the size of each array entry, in bytes.

**chunked\_size (comma separated list of integers) [SAME]:** This hint specifies the dimensions of the subarrays. This is a comma separated list of array dimensions, starting from the most significant one.

**filename (string):** This hint specifies the file name used when the file was opened. If the implementation is capable of returning the file name of an open file, it will be returned using this key by `MPI_FILE_GET_INFO`. This key is ignored when passed to `MPI_FILE_OPEN`, `MPI_FILE_SET_VIEW`, `MPI_FILE_SET_INFO`, and `MPI_FILE_DELETE`.

**file\_perm (string) [SAME]:** This hint specifies the file permissions to use for file creation. Setting this hint is only useful when passed to `MPI_FILE_OPEN` with an `amode` that includes `MPI_MODE_CREATE`. The set of legal values for this key is implementation dependent.

**io\_node\_list (comma separated list of strings) [SAME]:** This hint specifies the list of I/O devices that should be used to store the file. This hint is most relevant when the file is created.

**nb\_proc (integer) [SAME]:** This hint specifies the number of parallel processes that will typically be assigned to run programs that access this file. This hint is most relevant when the file is created.

**num\_io\_nodes (integer) [SAME]:** This hint specifies the number of I/O devices in the system. This hint is most relevant when the file is created.

**striping\_factor (integer) [SAME]:** This hint specifies the number of I/O devices that the file should be striped across, and is relevant only when the file is created.

**striping\_unit (integer) [SAME]:** This hint specifies the suggested striping unit to be used for this file. The striping unit is the amount of consecutive data assigned to one I/O device before progressing to the next device, when striping across a number of devices. It is expressed in bytes. This hint is relevant only when the file is created.

### 13.3 File Views

**MPI\_FILE\_SET\_VIEW(fh, disp, etype, filetype, datarep, info)**

INOUT	fh	file handle (handle)
IN	disp	displacement (integer)
IN	etype	elementary datatype (handle)
IN	filetype	filetype (handle)
IN	datarep	data representation (string)
IN	info	info object (handle)

```
int MPI_File_set_view(MPI_File fh, MPI_Offset disp, MPI_Datatype etype,
                     MPI_Datatype filetype, char *datarep, MPI_Info info)
```

```
MPI_FILE_SET_VIEW(FH, DISP, ETYPE, FILETYPE, DATAREP, INFO, IERROR)
    INTEGER FH, ETYPE, FILETYPE, INFO, IERROR
    CHARACTER*(*) DATAREP
    INTEGER(KIND=MPI_OFFSET_KIND) DISP
```

```
{void MPI::File::Set_view(MPI::Offset disp, const MPI::Datatype& etype,
                          const MPI::Datatype& filetype, const char* datarep,
                          const MPI::Info& info) (binding deprecated, see Section 15.2) }
```

The **MPI\_FILE\_SET\_VIEW** routine changes the process's view of the data in the file. The start of the view is set to **disp**; the type of data is set to **etype**; the distribution of data to processes is set to **filetype**; and the representation of data in the file is set to **datarep**. In addition, **MPI\_FILE\_SET\_VIEW** resets the individual file pointers and the shared file pointer to zero. **MPI\_FILE\_SET\_VIEW** is collective; the values for **datarep** and the extents



of `etype` in the file data representation must be identical on all processes in the group; values for `disp`, `filetype`, and `info` may vary. The datatypes passed in `etype` and `filetype` must be committed.

The `etype` always specifies the data layout in the file. If `etype` is a portable datatype (see Section 2.4, page 11), the extent of `etype` is computed by scaling any displacements in the datatype to match the file data representation. If `etype` is not a portable datatype, no scaling is done when computing the extent of `etype`. The user must be careful when using nonportable `etypes` in heterogeneous environments; see Section 13.5.1, page 462 for further details.

If `MPI_MODE_SEQUENTIAL` mode was specified when the file was opened, the special displacement `MPI_DISPLACEMENT_CURRENT` must be passed in `disp`. This sets the displacement to the current position of the shared file pointer. `MPI_DISPLACEMENT_CURRENT` is invalid unless the `amode` for the file has `MPI_MODE_SEQUENTIAL` set.

*Rationale.* For some sequential files, such as those corresponding to magnetic tapes or streaming network connections, the *displacement* may not be meaningful. `MPI_DISPLACEMENT_CURRENT` allows the view to be changed for these types of files. (*End of rationale.*)

*Advice to implementors.* It is expected that a call to `MPI_FILE_SET_VIEW` will immediately follow `MPI_FILE_OPEN` in numerous instances. A high-quality implementation will ensure that this behavior is efficient. (*End of advice to implementors.*)

The `disp` displacement argument specifies the position (absolute offset in bytes from the beginning of the file) where the view begins.

*Advice to users.* `disp` can be used to skip headers or when the file includes a sequence of data segments that are to be accessed in different patterns (see Figure 13.3). Separate views, each using a different displacement and `filetype`, can be used to access each segment.




Figure 13.3: Displacements

(*End of advice to users.*)

An *etype* (*elementary* datatype) is the unit of data access and positioning. It can be any MPI predefined or derived datatype. Derived `etypes` can be constructed by using any of the MPI datatype constructor routines, provided all resulting `typemap` displacements are non-negative and monotonically nondecreasing. Data access is performed in `etype` units, reading or writing whole data items of type `etype`. Offsets are expressed as a count of `etypes`; file pointers point to the beginning of `etypes`.

*Advice to users.* In order to ensure interoperability in a heterogeneous environment, additional restrictions must be observed when constructing the **etype** (see Section 13.5, page 460). (*End of advice to users.*)

A filetype is either a single etype or a derived MPI datatype constructed from multiple instances of the same etype. In addition, the extent of any hole in the filetype must be a multiple of the etype’s extent. These displacements are not required to be distinct, but they cannot be negative, and they must be monotonically nondecreasing.

If the file is opened for writing, neither the **etype** nor the **filetype** is permitted to contain overlapping regions. This restriction is equivalent to the “datatype used in a receive cannot specify overlapping regions” restriction for communication. Note that **filetypes** from different processes may still overlap each other.

If **filetype** has holes in it, then the data in the holes is inaccessible to the calling process. However, the **disp**, **etype** and **filetype** arguments can be changed via future calls to **MPI\_FILE\_SET\_VIEW** to access a different part of the file.

It is erroneous to use absolute addresses in the construction of the **etype** and **filetype**.

The **info** argument is used to provide information regarding file access patterns and file system specifics to direct optimization (see Section 13.2.8, page 430). The constant **MPI\_INFO\_NULL** refers to the null info and can be used when no info needs to be specified.

The **datarep** argument is a string that specifies the representation of data in the file. See the file interoperability section (Section 13.5, page 460) for details and a discussion of valid values.

The user is responsible for ensuring that all nonblocking requests and split collective operations on **fh** have been completed before calling **MPI\_FILE\_SET\_VIEW**—otherwise, the call to **MPI\_FILE\_SET\_VIEW** is erroneous.

**MPI\_FILE\_GET\_VIEW**(**fh**, **disp**, **etype**, **filetype**, **datarep**)

IN	<b>fh</b>	file handle (handle)
OUT	<b>disp</b>	displacement (integer)
OUT	<b>etype</b>	elementary datatype (handle)
OUT	<b>filetype</b>	filetype (handle)
OUT	<b>datarep</b>	data representation (string)

```
int MPI_File_get_view(MPI_File fh, MPI_Offset *disp, MPI_Datatype *etype,
                     MPI_Datatype *filetype, char *datarep)
```

```
MPI_FILE_GET_VIEW(FH, DISP, ETYPE, FILETYPE, DATAREP, IERROR)
```

```
INTEGER FH, ETYPE, FILETYPE, IERROR
```

```
CHARACTER*(*) DATAREP
```

```
INTEGER(KIND=MPI_OFFSET_KIND) DISP
```

```
{void MPI::File::Get_view(MPI::Offset& disp, MPI::Datatype& etype,
                          MPI::Datatype& filetype, char* datarep) const(binding deprecated,
                          see Section 15.2) }
```

**MPI\_FILE\_GET\_VIEW** returns the process’s view of the data in the file. The current value of the displacement is returned in **disp**. The **etype** and **filetype** are new datatypes with

typemaps equal to the typemaps of the current etype and filetype, respectively.

The data representation is returned in `datarep`. The user is responsible for ensuring that `datarep` is large enough to hold the returned data representation string. The length of a data representation string is limited to the value of `MPI_MAX_DATAREP_STRING`.

In addition, if a portable datatype was used to set the current view, then the corresponding datatype returned by `MPI_FILE_GET_VIEW` is also a portable datatype. If `etype` or `filetype` are derived datatypes, the user is responsible for freeing them. The `etype` and `filetype` returned are both in a committed state.

## 13.4 Data Access

### 13.4.1 Data Access Routines

Data is moved between files and processes by issuing read and write calls. There are three orthogonal aspects to data access: positioning (explicit offset *vs.* implicit file pointer), synchronism (blocking *vs.* nonblocking and split collective), and coordination (noncollective *vs.* collective). The following combinations of these data access routines, including two types of file pointers (individual and shared) are provided in Table 13.1.

positioning	synchronism	coordination	
		<i>noncollective</i>	<i>collective</i>
<i>explicit offsets</i>	<i>blocking</i>	MPI_FILE_READ_AT MPI_FILE_WRITE_AT	MPI_FILE_READ_AT_ALL MPI_FILE_WRITE_AT_ALL
	<i>nonblocking &amp; split collective</i>	MPI_FILE_IREAD_AT  MPI_FILE_IWRITE_AT	MPI_FILE_READ_AT_ALL_BEGIN MPI_FILE_READ_AT_ALL_END MPI_FILE_WRITE_AT_ALL_BEGIN MPI_FILE_WRITE_AT_ALL_END
<i>individual file pointers</i>	<i>blocking</i>	MPI_FILE_READ MPI_FILE_WRITE	MPI_FILE_READ_ALL MPI_FILE_WRITE_ALL
	<i>nonblocking &amp; split collective</i>	MPI_FILE_IREAD  MPI_FILE_IWRITE	MPI_FILE_READ_ALL_BEGIN MPI_FILE_READ_ALL_END MPI_FILE_WRITE_ALL_BEGIN MPI_FILE_WRITE_ALL_END
<i>shared file pointer</i>	<i>blocking</i>	MPI_FILE_READ_SHARED MPI_FILE_WRITE_SHARED	MPI_FILE_READ_ORDERED MPI_FILE_WRITE_ORDERED
	<i>nonblocking &amp; split collective</i>	MPI_FILE_IREAD_SHARED  MPI_FILE_IWRITE_SHARED	MPI_FILE_READ_ORDERED_BEGIN MPI_FILE_READ_ORDERED_END MPI_FILE_WRITE_ORDERED_BEGIN MPI_FILE_WRITE_ORDERED_END

Table 13.1: Data access routines

POSIX `read()/fread()` and `write()/fwrite()` are blocking, noncollective operations and use individual file pointers. The MPI equivalents are `MPI_FILE_READ` and `MPI_FILE_WRITE`.

Implementations of data access routines may buffer data to improve performance. This does not affect reads, as the data is always available in the user's buffer after a read operation completes. For writes, however, the `MPI_FILE_SYNC` routine provides the only guarantee that data has been transferred to the storage device.

## Positioning

MPI provides three types of positioning for data access routines: explicit offsets, individual file pointers, and shared file pointers. The different positioning methods may be mixed within the same program and do not affect each other.

The data access routines that accept explicit offsets contain `_AT` in their name (e.g., `MPI_FILE_WRITE_AT`). Explicit offset operations perform data access at the file position given directly as an argument—no file pointer is used nor updated. Note that this is not equivalent to an atomic seek-and-read or seek-and-write operation, as no “seek” is issued. Operations with explicit offsets are described in Section 13.4.2, page 439.

The names of the individual file pointer routines contain no positional qualifier (e.g., `MPI_FILE_WRITE`). Operations with individual file pointers are described in Section 13.4.3, page 442. The data access routines that use shared file pointers contain `_SHARED` or `_ORDERED` in their name (e.g., `MPI_FILE_WRITE_SHARED`). Operations with shared file pointers are described in Section 13.4.4, page 448.

The main semantic issues with MPI-maintained file pointers are how and when they are updated by I/O operations. In general, each I/O operation leaves the file pointer pointing to the next data item after the last one that is accessed by the operation. In a nonblocking or split collective operation, the pointer is updated by the call that initiates the I/O, possibly before the access completes.

More formally,

$$new\_file\_offset = old\_file\_offset + \frac{elements(datatype)}{elements(etype)} \times count$$

where *count* is the number of *datatype* items to be accessed, *elements(X)* is the number of predefined datatypes in the typemap of *X*, and *old\_file\_offset* is the value of the implicit offset before the call. The file position, *new\_file\_offset*, is in terms of a count of etypes relative to the current view.

## Synchronism

MPI supports blocking and nonblocking I/O routines.

A *blocking* I/O call will not return until the I/O request is completed.

A *nonblocking* I/O call initiates an I/O operation, but does not wait for it to complete. Given suitable hardware, this allows the transfer of data out/in the user’s buffer to proceed concurrently with computation. A separate *request complete* call (`MPI_WAIT`, `MPI_TEST`, or any of their variants) is needed to complete the I/O request, i.e., to confirm that the data has been read or written and that it is safe for the user to reuse the buffer. The nonblocking versions of the routines are named `MPI_FILE_IXXX`, where the *I* stands for immediate.

It is erroneous to access the local buffer of a nonblocking data access operation, or to use that buffer as the source or target of other communications, between the initiation and completion of the operation.

The split collective routines support a restricted form of “nonblocking” operations for collective data access (see Section 13.4.5, page 453).

## Coordination

Every noncollective data access routine `MPI_FILE_XXX` has a collective counterpart. For most routines, this counterpart is `MPI_FILE_XXX_ALL` or a pair of `MPI_FILE_XXX_BEGIN`

and `MPI_FILE_XXX_END`. The counterparts to the `MPI_FILE_XXX_SHARED` routines are `MPI_FILE_XXX_ORDERED`.

The completion of a noncollective call only depends on the activity of the calling process. However, the completion of a collective call (which must be called by all members of the process group) may depend on the activity of the other processes participating in the collective call. See Section 13.6.4, page 473, for rules on semantics of collective calls.

Collective operations may perform much better than their noncollective counterparts, as global data accesses have significant potential for automatic optimization.

## Data Access Conventions

Data is moved between files and processes by calling read and write routines. Read routines move data from a file into memory. Write routines move data from memory into a file. The file is designated by a file handle, `fh`. The location of the file data is specified by an offset into the current view. The data in memory is specified by a triple: `buf`, `count`, and `datatype`. Upon completion, the amount of data accessed by the calling process is returned in a `status`.

An offset designates the starting position in the file for an access. The offset is always in `etype` units relative to the current view. Explicit offset routines pass `offset` as an argument (negative values are erroneous). The file pointer routines use implicit offsets maintained by MPI.

A data access routine attempts to transfer (read or write) `count` data items of type `datatype` between the user's buffer `buf` and the file. The `datatype` passed to the routine must be a committed datatype. The layout of data in memory corresponding to `buf`, `count`, `datatype` is interpreted the same way as in MPI communication functions; see Section 3.2.2 on page 29 and Section 4.1.11 on page 104. The data is accessed from those parts of the file specified by the current view (Section 13.3, page 433). The type signature of `datatype` must match the type signature of some number of contiguous copies of the `etype` of the current view. As in a receive, it is erroneous to specify a `datatype` for reading that contains overlapping regions (areas of memory which would be stored into more than once).

The nonblocking data access routines indicate that MPI can start a data access and associate a request handle, `request`, with the I/O operation. Nonblocking operations are completed via `MPI_TEST`, `MPI_WAIT`, or any of their variants.

Data access operations, when completed, return the amount of data accessed in `status`.

*Advice to users.* To prevent problems with the argument copying and register optimization done by Fortran compilers, please note the hints in subsections “Problems Due to Data Copying and Sequence Association,” and “A Problem with Register Optimization” in Section 16.2.2, pages 514 and 517. (*End of advice to users.*)

For blocking routines, `status` is returned directly. For nonblocking routines and split collective routines, `status` is returned when the operation is completed. The number of `datatype` entries and predefined elements accessed by the calling process can be extracted from `status` by using `MPI_GET_COUNT` and `MPI_GET_ELEMENTS`, respectively. The interpretation of the `MPI_ERROR` field is the same as for other operations — normally undefined, but meaningful if an MPI routine returns `MPI_ERR_IN_STATUS`. The user can pass (in C and Fortran) `MPI_STATUS_IGNORE` in the `status` argument if the return value of this argument is not needed. In C++, the `status` argument is optional. The `status` can be passed to `MPI_TEST_CANCELLED` to determine if the operation was cancelled. All other fields of `status` are undefined.

When reading, a program can detect the end of file by noting that the amount of data read is less than the amount requested. Writing past the end of file increases the file size. The amount of data accessed will be the amount requested, unless an error is raised (or a read reaches the end of file).

### 13.4.2 Data Access with Explicit Offsets

If `MPI_MODE_SEQUENTIAL` mode was specified when the file was opened, it is erroneous to call the routines in this section.

`MPI_FILE_READ_AT(fh, offset, buf, count, datatype, status)`

IN	fh	file handle (handle)
IN	offset	file offset (integer)
OUT	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	status	status object (Status)

```
int MPI_File_read_at(MPI_File fh, MPI_Offset offset, void *buf, int count,
                    MPI_Datatype datatype, MPI_Status *status)
```

```
MPI_FILE_READ_AT(FH, OFFSET, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
```

```
{void MPI::File::Read_at(MPI::Offset offset, void* buf, int count,
    const MPI::Datatype& datatype, MPI::Status& status) (binding
    deprecated, see Section 15.2) }
```

```
{void MPI::File::Read_at(MPI::Offset offset, void* buf, int count,
    const MPI::Datatype& datatype) (binding deprecated, see Section 15.2)
    }
```

`MPI_FILE_READ_AT` reads a file beginning at the position specified by `offset`.

`MPI_FILE_READ_AT_ALL(fh, offset, buf, count, datatype, status)`

IN	fh	file handle (handle)
IN	offset	file offset (integer)
OUT	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	status	status object (Status)

```

1  int MPI_File_read_at_all(MPI_File fh, MPI_Offset offset, void *buf,
2      int count, MPI_Datatype datatype, MPI_Status *status)
3
4  MPI_FILE_READ_AT_ALL(FH, OFFSET, BUF, COUNT, DATATYPE, STATUS, IERROR)
5      <type> BUF(*)
6      INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
7      INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
8
9  {void MPI::File::Read_at_all(MPI::Offset offset, void* buf, int count,
10     const MPI::Datatype& datatype, MPI::Status& status) (binding
11     deprecated, see Section 15.2) }
12
13 {void MPI::File::Read_at_all(MPI::Offset offset, void* buf, int count,
14     const MPI::Datatype& datatype) (binding deprecated, see Section 15.2)
15     }
```

MPI\_FILE\_READ\_AT\_ALL is a collective version of the blocking MPI\_FILE\_READ\_AT interface.

```

16
17
18
19 MPI_FILE_WRITE_AT(fh, offset, buf, count, datatype, status)
20
21     INOUT    fh                file handle (handle)
22
23     IN       offset            file offset (integer)
24
25     IN       buf               initial address of buffer (choice)
26
27     IN       count             number of elements in buffer (integer)
28
29     IN       datatype           datatype of each buffer element (handle)
30
31     OUT      status            status object (Status)
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000
1001
1002
1003
1004
1005
1006
1007
1008
1009
1010
1011
1012
1013
1014
1015
1016
1017
1018
1019
1020
1021
1022
1023
1024
1025
1026
1027
1028
1029
1030
1031
1032
1033
1034
1035
1036
1037
1038
1039
1040
1041
1042
1043
1044
1045
1046
1047
1048
1049
1050
1051
1052
1053
1054
1055
1056
1057
1058
1059
1060
1061
1062
1063
1064
1065
1066
1067
1068
1069
1070
1071
1072
1073
1074
1075
1076
1077
1078
1079
1080
1081
1082
1083
1084
1085
1086
1087
1088
1089
1090
1091
1092
1093
1094
1095
1096
1097
1098
1099
1100
1101
1102
1103
1104
1105
1106
1107
1108
1109
1110
1111
1112
1113
1114
1115
1116
1117
1118
1119
1120
1121
1122
1123
1124
1125
1126
1127
1128
1129
1130
1131
1132
1133
1134
1135
1136
1137
1138
1139
1140
1141
1142
1143
1144
1145
1146
1147
1148
1149
1150
1151
1152
1153
1154
1155
1156
1157
1158
1159
1160
1161
1162
1163
1164
1165
1166
1167
1168
1169
1170
1171
1172
1173
1174
1175
1176
1177
1178
1179
1180
1181
1182
1183
1184
1185
1186
1187
1188
1189
1190
1191
1192
1193
1194
1195
1196
1197
1198
1199
1200
1201
1202
1203
1204
1205
1206
1207
1208
1209
1210
1211
1212
1213
1214
1215
1216
1217
1218
1219
1220
1221
1222
1223
1224
1225
1226
1227
1228
1229
1230
1231
1232
1233
1234
1235
1236
1237
1238
1239
1240
1241
1242
1243
1244
1245
1246
1247
1248
1249
1250
1251
1252
1253
1254
1255
1256
1257
1258
1259
1260
1261
1262
1263
1264
1265
1266
1267
1268
1269
1270
1271
1272
1273
1274
1275
1276
1277
1278
1279
1280
1281
1282
1283
1284
1285
1286
1287
1288
1289
1290
1291
1292
1293
1294
1295
1296
1297
1298
1299
1300
1301
1302
1303
1304
1305
1306
1307
1308
1309
1310
1311
1312
1313
1314
1315
1316
1317
1318
1319
1320
1321
1322
1323
1324
1325
1326
1327
1328
1329
1330
1331
1332
1333
1334
1335
1336
1337
1338
1339
1340
1341
1342
1343
1344
1345
1346
1347
1348
1349
1350
1351
1352
1353
1354
1355
1356
1357
1358
1359
1360
1361
1362
1363
1364
1365
1366
1367
1368
1369
1370
1371
1372
1373
1374
1375
1376
1377
1378
1379
1380
1381
1382
1383
1384
1385
1386
1387
1388
1389
1390
1391
1392
1393
1394
1395
1396
1397
1398
1399
1400
1401
1402
1403
1404
1405
1406
1407
1408
1409
1410
1411
1412
1413
1414
1415
1416
1417
1418
1419
1420
1421
1422
1423
1424
1425
1426
1427
1428
1429
1430
1431
1432
1433
1434
1435
1436
1437
1438
1439
1440
1441
1442
1443
1444
1445
1446
1447
1448
1449
1450
1451
1452
1453
1454
1455
1456
1457
1458
1459
1460
1461
1462
1463
1464
1465
1466
1467
1468
1469
1470
1471
1472
1473
1474
1475
1476
1477
1478
1479
1480
1481
1482
1483
1484
1485
1486
1487
1488
1489
1490
1491
1492
1493
1494
1495
1496
1497
1498
1499
1500
1501
1502
1503
1504
1505
1506
1507
1508
1509
1510
1511
1512
1513
1514
1515
1516
1517
1518
1519
1520
1521
1522
1523
1524
1525
1526
1527
1528
1529
1530
1531
1532
1533
1534
1535
1536
1537
1538
1539
1540
1541
1542
1543
1544
1545
1546
1547
1548
1549
1550
1551
1552
1553
1554
1555
1556
1557
1558
1559
1560
1561
1562
1563
1564
1565
1566
1567
1568
1569
1570
1571
1572
1573
1574
1575
1576
1577
1578
1579
1580
1581
1582
1583
1584
1585
1586
1587
1588
1589
1590
1591
1592
1593
1594
1595
1596
1597
1598
1599
1600
1601
1602
1603
1604
1605
1606
1607
1608
1609
1610
1611
1612
1613
1614
1615
1616
1617
1618
1619
1620
1621
1622
1623
1624
1625
1626
1627
1628
1629
1630
1631
1632
1633
1634
1635
1636
1637
1638
1639
1640
1641
1642
1643
1644
1645
1646
1647
1648
1649
1650
1651
1652
1653
1654
1655
1656
1657
1658
1659
1660
1661
1662
1663
1664
1665
1666
1667
1668
1669
1670
1671
1672
1673
1674
1675
1676
1677
1678
1679
1680
1681
1682
1683
1684
1685
1686
1687
1688
1689
1690
1691
1692
1693
1694
1695
1696
1697
1698
1699
1700
1701
1702
1703
1704
1705
1706
1707
1708
1709
1710
1711
1712
1713
1714
1715
1716
1717
1718
1719
1720
1721
1722
1723
1724
1725
1726
1727
1728
1729
1730
1731
1732
1733
1734
1735
1736
1737
1738
1739
1740
1741
1742
1743
1744
1745
1746
1747
1748
1749
1750
1751
1752
1753
1754
1755
1756
1757
1758
1759
1760
1761
1762
1763
1764
1765
1766
1767
1768
1769
1770
1771
1772
1773
1774
1775
1776
1777
1778
1779
1780
1781
1782
1783
1784
1785
1786
1787
1788
1789
1790
1791
1792
1793
1794
1795
1796
1797
1798
1799
1800
1801
1802
1803
1804
1805
1806
1807
1808
1809
1810
1811
1812
1813
1814
1815
1816
1817
1818
1819
1820
1821
1822
1823
1824
1825
1826
1827
1828
1829
1830
1831
1832
1833
1834
1835
1836
1837
1838
1839
1840
1841
1842
1843
1844
1845
1846
1847
1848
1849
1850
1851
1852
1853
1854
1855
1856
1857
1858
1859
1860
1861
1862
1863
1864
1865
1866
1867
1868
1869
1870
1871
1872
1873
1874
1875
1876
1877
1878
1879
1880
1881
1882
1883
1884
1885
1886
1887
1888
1889
1890
1891
1892
1893
1894
1895
1896
1897
1898
1899
1900
1901
1902
1903
1904
1905
1906
1907
1908
1909
1910
1911
1912
1913
1914
1915
1916
1917
1918
1919
1920
1921
1922
1923
1924
1925
1926
1927
1928
1929
1930
1931
1932
1933
1934
1935
1936
1937
1938
1939
1940
1941
1942
1943
1944
1945
1946
1947
1948
1949
1950
1951
1952
1953
1954
1955
1956
1957
1958
1959
1960
1961
1962
1963
1964
1965
1966
1967
1968
1969
1970
1971
1972
1973
1974
1975
1976
1977
1978
1979
1980
1981
1982
1983
1984
1985
1986
1987
1988
1989
1990
1991
1992
1993
1994
1995
1996
1997
1998
1999
2000
2001
2002
2003
2004
2005
2006
2007
2008
2009
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2059
2060
2061
2062
2063
2064
2065
2066
2067
2068
2069
2070
2071
2072
2073
2074
2075
2076
2077
2078
2079
2080
2081
2082
2083
2084
2085
2086
2087
2088
2089
2090
2091
2092
2093
2094
2095
2096
2097
2098
2099
2100
2101
2102
2103
2104
2105
2106
2107
2108
2109
2110
2111
2112
2113
2114
2115
2116
2117
2118
2119
2120
2121
2122
2123
2124
2125
2126
2127
2128
2129
2130
2131
2132
2133
2134
2135
2136
2137
2138
2139
2140
2141
2142
2143
2144
2145
2146
2147
2148
2149
2150
2151
2152
2153
2154
2155
2156
2157
2158
2159
2160
2161
2162
2163
2164
2165
2166
2167
2168
2169
2170
2171
2172
2173
2174
2175
2176
2177
2178
2179
2180
2181
2182
2183
2184
2185
2186
2187
2188
2189
2190
2191
2192
2193
2194
2195
2196
2197
2198
2199
2200
2201
2202
2203
2204
2205
2206
2207
2208
2209
2210
2211
2212
2213
2214
2215
2216
2217
2218
2219
2220
2221
2222
2223
2224
2225
2226
2227
2228
2229
2230
2231
2232
2233
2234
2235
2236
2237
2238
2239
2240
2241
2242
2243
2244
2245
2246
2247
2248
2249
2250
2251
2252
2253
2254
2255
2256
2257
2258
2259
2260
2261
2262
2263
2264
2265
2266
2267
2268
2269
2270
2271
2272
2273
2274
2275
2276
2277
2278
2279
2280
2281
2282
2283
2284
2285
2286
2287
2288
2289
2290
2291
2292
2293
2294
2295
2296
2297
2298
2299
2300
2301
2302
2303
2304
2305
2306
2307
2308
2309
2310
2311
2312
2313
2314
2315
2316
2317
2318
2319
2320
2321
2322
2323
2324
2325
2326
2327
2328
2329
2330
2331
2332
2333
2334
2335
2336
2337
2338
2339
2340
2341
2342
2343
2344
2345
2346
2347
2348
2349
2350
2351
2352
2353
2354
2355
2356
2357
2358
2359
2360
2361
2362
2363
2364
2365
2366
2367
2368
2369
2370
2371
2372
2373
2374
2375
2376
2377
2378
2379
2380
2381
2382
2383
2384
2385
2386
2387
2388
2389
2390
2391
2392
2393
2394
2395
2396
2397
2398
2399
2400
2401
2402
2403
2404
2405
2406
2407
2408
2409
2410
2411
2412
2413
2414
2415
2416
2417
2418
2419
2420
2421
2422
2423
2424
2425
2426
2427
2428
2429
2430
2431
2432
2433
2434
2435
2436
2437
2438
2439
2440
2441
2442
2443
2444
2445
2446
2447
2448
2449
2450
2451
2452
2453
2454
2455
2456
2457
2458
2459
2460
2461
2462
2463
2464
2465
2466
2467
2468
2469
2470
2471
2472
2473
2474
2475
2476
2477
2478
2479
2480
2481
2482
2483
2484
2485
2486
2487
2488
2489
2490
2491
2492
2493
2494
2495
2496
2497
2498
2499
2500
2501
2502
2503
2504
2505
2506
2507
2508
2509
2510
2511
2512
2513
2514
2515
2516
2517
2518
2519
2520
2521
2522
2523
2524
2525
2526
2527
2528
2529
2530
2531
2532
2533
2534
2535
2536
2537
2538
2539
2540
2541
2542
2543
2544
2545
2546
2547
2548
2549
2550
2551
2552
2553
2554
2555
2556
2557
2558
2559
2560
2561
2562
2563
2564
2565
2566
2567
2568
2569
2570
2571
2572
2573
2574
2575
2576
2577
2578
2579
2580
2581
2582
2583
2584
2585
2586
2587
2588
2589
2590
2591
2592
2593
2594
2595
2596
2597
2598
2599
2600
2601
2602
2603
2604

```



`MPI_FILE_WRITE_AT_ALL(fh, offset, buf, count, datatype, status)`

INOUT	fh	file handle (handle)
IN	offset	file offset (integer)
IN	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	status	status object (Status)

```
int MPI_File_write_at_all(MPI_File fh, MPI_Offset offset, void *buf,
                        int count, MPI_Datatype datatype, MPI_Status *status)
```

```
MPI_FILE_WRITE_AT_ALL(FH, OFFSET, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
```

```
{void MPI::File::Write_at_all(MPI::Offset offset, const void* buf,
    int count, const MPI::Datatype& datatype,
    MPI::Status& status) (binding deprecated, see Section 15.2) }
```

```
{void MPI::File::Write_at_all(MPI::Offset offset, const void* buf,
    int count, const MPI::Datatype& datatype) (binding deprecated, see
    Section 15.2) }
```

`MPI_FILE_WRITE_AT_ALL` is a collective version of the blocking `MPI_FILE_WRITE_AT` interface.

`MPI_FILE_IREAD_AT(fh, offset, buf, count, datatype, request)`

IN	fh	file handle (handle)
IN	offset	file offset (integer)
OUT	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	request	request object (handle)

```
int MPI_File_iread_at(MPI_File fh, MPI_Offset offset, void *buf, int count,
                    MPI_Datatype datatype, MPI_Request *request)
```

```
MPI_FILE_IREAD_AT(FH, OFFSET, BUF, COUNT, DATATYPE, REQUEST, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
```

```
{MPI::Request MPI::File::Iread_at(MPI::Offset offset, void* buf, int count,
    const MPI::Datatype& datatype) (binding deprecated, see Section 15.2)
    }
```



MPI\_FILE\_IREAD\_AT is a nonblocking version of the MPI\_FILE\_READ\_AT interface.

MPI\_FILE\_IWRITE\_AT(fh, offset, buf, count, datatype, request)

INOUT	fh	file handle (handle)
IN	offset	file offset (integer)
IN	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	request	request object (handle)

```
int MPI_File_iwrite_at(MPI_File fh, MPI_Offset offset, void *buf,
                      int count, MPI_Datatype datatype, MPI_Request *request)
```

```
MPI_FILE_IWRITE_AT(FH, OFFSET, BUF, COUNT, DATATYPE, REQUEST, IERROR)
```

```
<type> BUF(*)
```

```
INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR
```

```
INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
```

```
{MPI::Request MPI::File::Iwrite_at(MPI::Offset offset, const void* buf,
                                   int count, const MPI::Datatype& datatype) (binding deprecated, see
                                   Section 15.2) }
```

MPI\_FILE\_IWRITE\_AT is a nonblocking version of the MPI\_FILE\_WRITE\_AT interface.

### 13.4.3 Data Access with Individual File Pointers

MPI maintains one individual file pointer per process per file handle. The current value of this pointer implicitly specifies the offset in the data access routines described in this section. These routines only use and update the individual file pointers maintained by MPI. The shared file pointer is not used nor updated.

The individual file pointer routines have the same semantics as the data access with explicit offset routines described in Section 13.4.2, page 439, with the following modification:

- the `offset` is defined to be the current value of the MPI-maintained individual file pointer.

After an individual file pointer operation is initiated, the individual file pointer is updated to point to the next etype after the last one that will be accessed. The file pointer is updated relative to the current view of the file.

If MPI\_MODE\_SEQUENTIAL mode was specified when the file was opened, it is erroneous to call the routines in this section, with the exception of MPI\_FILE\_GET\_BYTE\_OFFSET.

```

MPI_FILE_READ(fh, buf, count, datatype, status)
    INOUT    fh                file handle (handle)
    OUT      buf                initial address of buffer (choice)
    IN       count              number of elements in buffer (integer)
    IN       datatype            datatype of each buffer element (handle)
    OUT      status              status object (Status)

int MPI_File_read(MPI_File fh, void *buf, int count, MPI_Datatype datatype,
                  MPI_Status *status)

MPI_FILE_READ(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR

{void MPI::File::Read(void* buf, int count, const MPI::Datatype& datatype,
    MPI::Status& status) (binding deprecated, see Section 15.2) }

{void MPI::File::Read(void* buf, int count,
    const MPI::Datatype& datatype) (binding deprecated, see Section 15.2)
    }

```

MPI\_FILE\_READ reads a file using the individual file pointer.

**Example 13.2** The following Fortran code fragment is an example of reading a file until the end of file is reached:

```

!   Read a preexisting input file until all data has been read.
!   Call routine "process_input" if all requested data is read.
!   The Fortran 90 "exit" statement exits the loop.

    integer    bufsize, numread, totprocessed, status(MPI_STATUS_SIZE)
    parameter (bufsize=100)
    real       localbuffer(bufsize)

    call MPI_FILE_OPEN( MPI_COMM_WORLD, 'myoldfile', &
                        MPI_MODE_RDONLY, MPI_INFO_NULL, myfh, ierr )
    call MPI_FILE_SET_VIEW( myfh, 0, MPI_REAL, MPI_REAL, 'native', &
                        MPI_INFO_NULL, ierr )

    totprocessed = 0
    do
        call MPI_FILE_READ( myfh, localbuffer, bufsize, MPI_REAL, &
                            status, ierr )
        call MPI_GET_COUNT( status, MPI_REAL, numread, ierr )
        call process_input( localbuffer, numread )
        totprocessed = totprocessed + numread
        if ( numread < bufsize ) exit
    enddo

```

```

1      write(6,1001) numread, bufsize, totprocessed
2 1001 format( "No more data:  read", I3, "and expected", I3, &
3          "Processed total of", I6, "before terminating job." )
4
5      call MPI_FILE_CLOSE( myfh, ierr )
6
7
8
9  MPI_FILE_READ_ALL(fh, buf, count, datatype, status)
10
11      INOUT    fh                file handle (handle)
12      OUT      buf              initial address of buffer (choice)
13      IN       count            number of elements in buffer (integer)
14      IN       datatype         datatype of each buffer element (handle)
15      OUT      status           status object (Status)
16
17
18  int MPI_File_read_all(MPI_File fh, void *buf, int count,
19                      MPI_Datatype datatype, MPI_Status *status)
20
21  MPI_FILE_READ_ALL(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
22      <type> BUF(*)
23      INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
24
25  {void MPI::File::Read_all(void* buf, int count,
26                          const MPI::Datatype& datatype, MPI::Status& status) (binding
27                          deprecated, see Section 15.2) }
28
29  {void MPI::File::Read_all(void* buf, int count,
30                          const MPI::Datatype& datatype) (binding deprecated, see Section 15.2)
31                          }
32
33  MPI_FILE_READ_ALL is a collective version of the blocking MPI_FILE_READ interface.
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48

```

```

{void MPI::File::Write(const void* buf, int count,
    const MPI::Datatype& datatype, MPI::Status& status) (binding
    deprecated, see Section 15.2) }
{void MPI::File::Write(const void* buf, int count,
    const MPI::Datatype& datatype) (binding deprecated, see Section 15.2)
}
```

MPI\_FILE\_WRITE writes a file using the individual file pointer.

MPI\_FILE\_WRITE\_ALL(fh, buf, count, datatype, status)

INOUT	fh	file handle (handle)
IN	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	status	status object (Status)

```

int MPI_File_write_all(MPI_File fh, void *buf, int count,
    MPI_Datatype datatype, MPI_Status *status)
```

MPI\_FILE\_WRITE\_ALL(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)

<type> BUF(\*)

INTEGER FH, COUNT, DATATYPE, STATUS(MPI\_STATUS\_SIZE), IERROR

```

{void MPI::File::Write_all(const void* buf, int count,
    const MPI::Datatype& datatype, MPI::Status& status) (binding
    deprecated, see Section 15.2) }
```

```

{void MPI::File::Write_all(const void* buf, int count,
    const MPI::Datatype& datatype) (binding deprecated, see Section 15.2)
}
```

MPI\_FILE\_WRITE\_ALL is a collective version of the blocking MPI\_FILE\_WRITE interface.

MPI\_FILE\_IREAD(fh, buf, count, datatype, request)

INOUT	fh	file handle (handle)
OUT	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	request	request object (handle)

```

int MPI_File_iread(MPI_File fh, void *buf, int count,
    MPI_Datatype datatype, MPI_Request *request)
```

MPI\_FILE\_IREAD(FH, BUF, COUNT, DATATYPE, REQUEST, IERROR)

```

1      <type> BUF(*)
2      INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR
3
4      {MPI::Request MPI::File::Iread(void* buf, int count,
5          const MPI::Datatype& datatype) (binding deprecated, see Section 15.2)
6          }

```

MPI\_FILE\_IREAD is a nonblocking version of the MPI\_FILE\_READ interface.

**Example 13.3** The following Fortran code fragment illustrates file pointer update semantics:

```

11
12      ! Read the first twenty real words in a file into two local
13      ! buffers. Note that when the first MPI_FILE_IREAD returns,
14      ! the file pointer has been updated to point to the
15      ! eleventh real word in the file.
16
17      integer    bufsize, req1, req2
18      integer, dimension(MPI_STATUS_SIZE) :: status1, status2
19      parameter (bufsize=10)
20      real       buf1(bufsize), buf2(bufsize)
21
22      call MPI_FILE_OPEN( MPI_COMM_WORLD, 'myoldfile', &
23                          MPI_MODE_RDONLY, MPI_INFO_NULL, myfh, ierr )
24      call MPI_FILE_SET_VIEW( myfh, 0, MPI_REAL, MPI_REAL, 'native', &
25                              MPI_INFO_NULL, ierr )
26      call MPI_FILE_IREAD( myfh, buf1, bufsize, MPI_REAL, &
27                          req1, ierr )
28      call MPI_FILE_IREAD( myfh, buf2, bufsize, MPI_REAL, &
29                          req2, ierr )
30
31      call MPI_WAIT( req1, status1, ierr )
32      call MPI_WAIT( req2, status2, ierr )
33
34      call MPI_FILE_CLOSE( myfh, ierr )
35
36

```

MPI\_FILE\_IWRITE(fh, buf, count, datatype, request)

INOUT	fh	file handle (handle)
IN	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	request	request object (handle)

```

46      int MPI_File_iwrite(MPI_File fh, void *buf, int count,
47                          MPI_Datatype datatype, MPI_Request *request)
48

```

```

MPI_FILE_IWRITE(FH, BUF, COUNT, DATATYPE, REQUEST, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR
{MPI::Request MPI::File::Iwrite(const void* buf, int count,
    const MPI::Datatype& datatype) (binding deprecated, see Section 15.2)
}

```

MPI\_FILE\_IWRITE is a nonblocking version of the MPI\_FILE\_WRITE interface.

```

MPI_FILE_SEEK(fh, offset, whence)

```

INOUT	fh	file handle (handle)
IN	offset	file offset (integer)
IN	whence	update mode (state)

```

int MPI_File_seek(MPI_File fh, MPI_Offset offset, int whence)

```

```

MPI_FILE_SEEK(FH, OFFSET, WHENCE, IERROR)
    INTEGER FH, WHENCE, IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET

```

```

{void MPI::File::Seek(MPI::Offset offset, int whence) (binding deprecated, see
    Section 15.2) }

```

MPI\_FILE\_SEEK updates the individual file pointer according to *whence*, which has the following possible values:

- MPI\_SEEK\_SET: the pointer is set to *offset*
- MPI\_SEEK\_CUR: the pointer is set to the current pointer position plus *offset*
- MPI\_SEEK\_END: the pointer is set to the end of file plus *offset*

The *offset* can be negative, which allows seeking backwards. It is erroneous to seek to a negative position in the view.

```

MPI_FILE_GET_POSITION(fh, offset)

```

IN	fh	file handle (handle)
OUT	offset	offset of individual pointer (integer)

```

int MPI_File_get_position(MPI_File fh, MPI_Offset *offset)

```

```

MPI_FILE_GET_POSITION(FH, OFFSET, IERROR)
    INTEGER FH, IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET

```

```

{MPI::Offset MPI::File::Get_position() const (binding deprecated, see Section 15.2)
}

```

`MPI_FILE_GET_POSITION` returns, in `offset`, the current position of the individual file pointer in etype units relative to the current view.

*Advice to users.* The `offset` can be used in a future call to `MPI_FILE_SEEK` using `whence = MPI_SEEK_SET` to return to the current position. To set the displacement to the current file pointer position, first convert `offset` into an absolute byte position using `MPI_FILE_GET_BYTE_OFFSET`, then call `MPI_FILE_SET_VIEW` with the resulting displacement. (*End of advice to users.*)

`MPI_FILE_GET_BYTE_OFFSET(fh, offset, disp)`

IN	fh	file handle (handle)
IN	offset	offset (integer)
OUT	disp	absolute byte position of offset (integer)

```
int MPI_File_get_byte_offset(MPI_File fh, MPI_Offset offset,
                             MPI_Offset *disp)
```

```
MPI_FILE_GET_BYTE_OFFSET(FH, OFFSET, DISP, IERROR)
```

```
INTEGER FH, IERROR
```

```
INTEGER(KIND=MPI_OFFSET_KIND) OFFSET, DISP
```

```
{MPI::Offset MPI::File::Get_byte_offset(const MPI::Offset disp)
  const(binding deprecated, see Section 15.2) }
```

`MPI_FILE_GET_BYTE_OFFSET` converts a view-relative offset into an absolute byte position. The absolute byte position (from the beginning of the file) of `offset` relative to the current view of `fh` is returned in `disp`.

#### 13.4.4 Data Access with Shared File Pointers

MPI maintains exactly one shared file pointer per collective `MPI_FILE_OPEN` (shared among processes in the communicator group). The current value of this pointer implicitly specifies the offset in the data access routines described in this section. These routines only use and update the shared file pointer maintained by MPI. The individual file pointers are not used nor updated.

The shared file pointer routines have the same semantics as the data access with explicit offset routines described in Section 13.4.2, page 439, with the following modifications:

- the `offset` is defined to be the current value of the MPI-maintained shared file pointer,
- the effect of multiple calls to shared file pointer routines is defined to behave as if the calls were serialized, and
- the use of shared file pointer routines is erroneous unless all processes use the same file view.

For the noncollective shared file pointer routines, the serialization ordering is not deterministic. The user needs to use other synchronization means to enforce a specific order.

After a shared file pointer operation is initiated, the shared file pointer is updated to point to the next etype after the last one that will be accessed. The file pointer is updated relative to the current view of the file.

#### Noncollective Operations

`MPI_FILE_READ_SHARED(fh, buf, count, datatype, status)`

INOUT	fh	file handle (handle)
OUT	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	status	status object (Status)

```
int MPI_File_read_shared(MPI_File fh, void *buf, int count,
                        MPI_Datatype datatype, MPI_Status *status)
```

```
MPI_FILE_READ_SHARED(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
```

```
{void MPI::File::Read_shared(void* buf, int count,
    const MPI::Datatype& datatype, MPI::Status& status) (binding
    deprecated, see Section 15.2) }
```

```
{void MPI::File::Read_shared(void* buf, int count,
    const MPI::Datatype& datatype) (binding deprecated, see Section 15.2)
    }
```

`MPI_FILE_READ_SHARED` reads a file using the shared file pointer.

`MPI_FILE_WRITE_SHARED(fh, buf, count, datatype, status)`

INOUT	fh	file handle (handle)
IN	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	status	status object (Status)

```
int MPI_File_write_shared(MPI_File fh, void *buf, int count,
                        MPI_Datatype datatype, MPI_Status *status)
```

```
MPI_FILE_WRITE_SHARED(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
```



```

1  {void MPI::File::Write_shared(const void* buf, int count,
2      const MPI::Datatype& datatype, MPI::Status& status) (binding
3      deprecated, see Section 15.2) }
4
5  {void MPI::File::Write_shared(const void* buf, int count,
6      const MPI::Datatype& datatype) (binding deprecated, see Section 15.2)
7      }

```

MPI\_FILE\_WRITE\_SHARED writes a file using the shared file pointer.

MPI\_FILE\_IREAD\_SHARED(fh, buf, count, datatype, request)

INOUT	fh	file handle (handle)
OUT	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	request	request object (handle)

```

19
20 int MPI_File_iread_shared(MPI_File fh, void *buf, int count,
21     MPI_Datatype datatype, MPI_Request *request)

```

MPI\_FILE\_IREAD\_SHARED(FH, BUF, COUNT, DATATYPE, REQUEST, IERROR)

```

23     <type> BUF(*)
24     INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR

```

```

26 {MPI::Request MPI::File::Iread_shared(void* buf, int count,
27     const MPI::Datatype& datatype) (binding deprecated, see Section 15.2)
28     }

```

MPI\_FILE\_IREAD\_SHARED is a nonblocking version of the MPI\_FILE\_READ\_SHARED interface.

MPI\_FILE\_IWRITE\_SHARED(fh, buf, count, datatype, request)

INOUT	fh	file handle (handle)
IN	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	request	request object (handle)

```

41
42 int MPI_File_ fwrite_shared(MPI_File fh, void *buf, int count,
43     MPI_Datatype datatype, MPI_Request *request)

```

MPI\_FILE\_IWRITE\_SHARED(FH, BUF, COUNT, DATATYPE, REQUEST, IERROR)

```

45     <type> BUF(*)
46     INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR

```

```
{MPI::Request MPI::File::Iwrite_shared(const void* buf, int count,
    const MPI::Datatype& datatype) (binding deprecated, see Section 15.2)
}
```

MPI\_FILE\_IWRITE\_SHARED is a nonblocking version of the MPI\_FILE\_WRITE\_SHARED interface.

### Collective Operations

The semantics of a collective access using a shared file pointer is that the accesses to the file will be in the order determined by the ranks of the processes within the group. For each process, the location in the file at which data is accessed is the position at which the shared file pointer would be after all processes whose ranks within the group less than that of this process had accessed their data. In addition, in order to prevent subsequent shared offset accesses by the same processes from interfering with this collective access, the call might return only after all the processes within the group have initiated their accesses. When the call returns, the shared file pointer points to the next etype accessible, according to the file view used by all processes, after the last etype requested.

*Advice to users.* There may be some programs in which all processes in the group need to access the file using the shared file pointer, but the program may not *require* that data be accessed in order of process rank. In such programs, using the shared ordered routines (e.g., MPI\_FILE\_WRITE\_ORDERED rather than MPI\_FILE\_WRITE\_SHARED) may enable an implementation to optimize access, improving performance. (*End of advice to users.*)

*Advice to implementors.* Accesses to the data requested by all processes do not have to be serialized. Once all processes have issued their requests, locations within the file for all accesses can be computed, and accesses can proceed independently from each other, possibly in parallel. (*End of advice to implementors.*)

MPI\_FILE\_READ\_ORDERED(fh, buf, count, datatype, status)

INOUT	fh	file handle (handle)
OUT	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	status	status object (Status)

```
int MPI_File_read_ordered(MPI_File fh, void *buf, int count,
    MPI_Datatype datatype, MPI_Status *status)
```

```
MPI_FILE_READ_ORDERED(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
```

```

1  {void MPI::File::Read_ordered(void* buf, int count,
2      const MPI::Datatype& datatype, MPI::Status& status) (binding
3      deprecated, see Section 15.2) }
4
5  {void MPI::File::Read_ordered(void* buf, int count,
6      const MPI::Datatype& datatype) (binding deprecated, see Section 15.2)
7      }

```

MPI\_FILE\_READ\_ORDERED is a collective version of the MPI\_FILE\_READ\_SHARED interface.

MPI\_FILE\_WRITE\_ORDERED(fh, buf, count, datatype, status)

INOUT	fh	file handle (handle)
IN	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	status	status object (Status)

```

20 int MPI_File_write_ordered(MPI_File fh, void *buf, int count,
21     MPI_Datatype datatype, MPI_Status *status)
22

```

```

23 MPI_FILE_WRITE_ORDERED(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
24     <type> BUF(*)
25     INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR

```

```

26 {void MPI::File::Write_ordered(const void* buf, int count,
27     const MPI::Datatype& datatype, MPI::Status& status) (binding
28     deprecated, see Section 15.2) }
29

```

```

30 {void MPI::File::Write_ordered(const void* buf, int count,
31     const MPI::Datatype& datatype) (binding deprecated, see Section 15.2)
32     }
33

```

MPI\_FILE\_WRITE\_ORDERED is a collective version of the MPI\_FILE\_WRITE\_SHARED interface.

### Seek

If MPI\_MODE\_SEQUENTIAL mode was specified when the file was opened, it is erroneous to call the following two routines (MPI\_FILE\_SEEK\_SHARED and MPI\_FILE\_GET\_POSITION\_SHARED).

MPI\_FILE\_SEEK\_SHARED(fh, offset, whence)

INOUT	fh	file handle (handle)
IN	offset	file offset (integer)
IN	whence	update mode (state)

```

int MPI_File_seek_shared(MPI_File fh, MPI_Offset offset, int whence)
MPI_FILE_SEEK_SHARED(FH, OFFSET, WHENCE, IERROR)
    INTEGER FH, WHENCE, IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
{void MPI::File::Seek_shared(MPI::Offset offset, int whence) (binding
    deprecated, see Section 15.2) }

```

MPI\_FILE\_SEEK\_SHARED updates the shared file pointer according to *whence*, which has the following possible values:

- MPI\_SEEK\_SET: the pointer is set to *offset*
- MPI\_SEEK\_CUR: the pointer is set to the current pointer position plus *offset*
- MPI\_SEEK\_END: the pointer is set to the end of file plus *offset*

MPI\_FILE\_SEEK\_SHARED is collective; all the processes in the communicator group associated with the file handle *fh* must call MPI\_FILE\_SEEK\_SHARED with the same values for *offset* and *whence*.

The *offset* can be negative, which allows seeking backwards. It is erroneous to seek to a negative position in the view.

```

MPI_FILE_GET_POSITION_SHARED(fh, offset)
    IN      fh      file handle (handle)
    OUT     offset   offset of shared pointer (integer)

```

```

int MPI_File_get_position_shared(MPI_File fh, MPI_Offset *offset)
MPI_FILE_GET_POSITION_SHARED(FH, OFFSET, IERROR)
    INTEGER FH, IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
{MPI::Offset MPI::File::Get_position_shared() const(binding deprecated, see
    Section 15.2) }

```

MPI\_FILE\_GET\_POSITION\_SHARED returns, in *offset*, the current position of the shared file pointer in byte units relative to the current view.

*Advice to users.* The *offset* can be used in a future call to MPI\_FILE\_SEEK\_SHARED using *whence* = MPI\_SEEK\_SET to return to the current position. To set the displacement to the current file pointer position, first convert *offset* into an absolute byte position using MPI\_FILE\_GET\_BYTE\_OFFSET, then call MPI\_FILE\_SET\_VIEW with the resulting displacement. (*End of advice to users.*)

#### 13.4.5 Split Collective Data Access Routines

MPI provides a restricted form of “nonblocking collective” I/O operations for all data accesses using split collective data access routines. These routines are referred to as “split” collective routines because a single collective operation is split in two: a begin routine and

an end routine. The begin routine begins the operation, much like a nonblocking data access (e.g., `MPI_FILE_IREAD`). The end routine completes the operation, much like the matching test or wait (e.g., `MPI_WAIT`). As with nonblocking data access operations, the user must not use the buffer passed to a begin routine while the routine is outstanding; the operation must be completed with an end routine before it is safe to free buffers, etc.

Split collective data access operations on a file handle `fh` are subject to the semantic rules given below.

- On any MPI process, each file handle may have at most one active split collective operation at any time.
- Begin calls are collective over the group of processes that participated in the collective open and follow the ordering rules for collective calls.
- End calls are collective over the group of processes that participated in the collective open and follow the ordering rules for collective calls. Each end call matches the preceding begin call for the same collective operation. When an “end” call is made, exactly one unmatched “begin” call for the same operation must precede it.
- An implementation is free to implement any split collective data access routine using the corresponding blocking collective routine when either the begin call (e.g., `MPI_FILE_READ_ALL_BEGIN`) or the end call (e.g., `MPI_FILE_READ_ALL_END`) is issued. The begin and end calls are provided to allow the user and MPI implementation to optimize the collective operation.
- Split collective operations do not match the corresponding regular collective operation. For example, in a single collective read operation, an `MPI_FILE_READ_ALL` on one process does not match an `MPI_FILE_READ_ALL_BEGIN`/`MPI_FILE_READ_ALL_END` pair on another process.
- Split collective routines must specify a buffer in both the begin and end routines. By specifying the buffer that receives data in the end routine, we can avoid many (though not all) of the problems described in “A Problem with Register Optimization,” Section 16.2.2, page 517.
- No collective I/O operations are permitted on a file handle concurrently with a split collective access on that file handle (i.e., between the begin and end of the access). That is

```
MPI_File_read_all_begin(fh, ...);
...
MPI_File_read_all(fh, ...);
...
MPI_File_read_all_end(fh, ...);
```

is erroneous.

- In a multithreaded implementation, any split collective begin and end operation called by a process must be called from the same thread. This restriction is made to simplify the implementation in the multithreaded case. (Note that we have already disallowed having two threads begin a split collective operation on the same file handle since only one split collective operation can be active on a file handle at any time.)

The arguments for these routines have the same meaning as for the equivalent collective versions (e.g., the argument definitions for `MPI_FILE_READ_ALL_BEGIN` and `MPI_FILE_READ_ALL_END` are equivalent to the arguments for `MPI_FILE_READ_ALL`). The begin routine (e.g., `MPI_FILE_READ_ALL_BEGIN`) begins a split collective operation that, when completed with the matching end routine (i.e., `MPI_FILE_READ_ALL_END`) produces the result as defined for the equivalent collective routine (i.e., `MPI_FILE_READ_ALL`).

For the purpose of consistency semantics (Section 13.6.1, page 469), a matched pair of split collective data access operations (e.g., `MPI_FILE_READ_ALL_BEGIN` and `MPI_FILE_READ_ALL_END`) compose a single data access.

`MPI_FILE_READ_AT_ALL_BEGIN(fh, offset, buf, count, datatype)`

IN	fh	file handle (handle)
IN	offset	file offset (integer)
OUT	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)

```
int MPI_File_read_at_all_begin(MPI_File fh, MPI_Offset offset, void *buf,
                              int count, MPI_Datatype datatype)
```

```
MPI_FILE_READ_AT_ALL_BEGIN(FH, OFFSET, BUF, COUNT, DATATYPE, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
```

```
{void MPI::File::Read_at_all_begin(MPI::Offset offset, void* buf,
    int count, const MPI::Datatype& datatype) (binding deprecated, see
    Section 15.2) }
```

`MPI_FILE_READ_AT_ALL_END(fh, buf, status)`

IN	fh	file handle (handle)
OUT	buf	initial address of buffer (choice)
OUT	status	status object (Status)

```
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status)
```

```
MPI_FILE_READ_AT_ALL_END(FH, BUF, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR
```

```
{void MPI::File::Read_at_all_end(void* buf, MPI::Status& status) (binding
    deprecated, see Section 15.2) }
```

```
1 {void MPI::File::Read_at_all_end(void* buf) (binding deprecated, see Section 15.2)
2     }
```

```
5 MPI_FILE_WRITE_AT_ALL_BEGIN(fh, offset, buf, count, datatype)
```

```
6     INOUT    fh                file handle (handle)
7
8     IN       offset            file offset (integer)
9
10    IN       buf               initial address of buffer (choice)
11
12    IN       count             number of elements in buffer (integer)
13
14    IN       datatype           datatype of each buffer element (handle)
```

```
14 int MPI_File_write_at_all_begin(MPI_File fh, MPI_Offset offset, void *buf,
15     int count, MPI_Datatype datatype)
```

```
16 MPI_FILE_WRITE_AT_ALL_BEGIN(FH, OFFSET, BUF, COUNT, DATATYPE, IERROR)
17     <type> BUF(*)
18     INTEGER FH, COUNT, DATATYPE, IERROR
19     INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
```

```
21 {void MPI::File::Write_at_all_begin(MPI::Offset offset, const void* buf,
22     int count, const MPI::Datatype& datatype) (binding deprecated, see
23     Section 15.2) }
```

```
26 MPI_FILE_WRITE_AT_ALL_END(fh, buf, status)
```

```
27
28     INOUT    fh                file handle (handle)
29
30     IN       buf               initial address of buffer (choice)
31
32     OUT      status            status object (Status)
```

```
33 int MPI_File_write_at_all_end(MPI_File fh, void *buf, MPI_Status *status)
```

```
34 MPI_FILE_WRITE_AT_ALL_END(FH, BUF, STATUS, IERROR)
35     <type> BUF(*)
36     INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR
```

```
37 {void MPI::File::Write_at_all_end(const void* buf,
38     MPI::Status& status) (binding deprecated, see Section 15.2) }
```

```
40 {void MPI::File::Write_at_all_end(const void* buf) (binding deprecated, see
41     Section 15.2) }
```

```

MPI_FILE_READ_ALL_BEGIN(fh, buf, count, datatype)
    INOUT    fh                file handle (handle)
    OUT      buf                initial address of buffer (choice)
    IN       count              number of elements in buffer (integer)
    IN       datatype           datatype of each buffer element (handle)

int MPI_File_read_all_begin(MPI_File fh, void *buf, int count,
    MPI_Datatype datatype)

MPI_FILE_READ_ALL_BEGIN(FH, BUF, COUNT, DATATYPE, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, IERROR

{void MPI::File::Read_all_begin(void* buf, int count,
    const MPI::Datatype& datatype) (binding deprecated, see Section 15.2)
}

MPI_FILE_READ_ALL_END(fh, buf, status)
    INOUT    fh                file handle (handle)
    OUT      buf                initial address of buffer (choice)
    OUT      status             status object (Status)

int MPI_File_read_all_end(MPI_File fh, void *buf, MPI_Status *status)

MPI_FILE_READ_ALL_END(FH, BUF, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR

{void MPI::File::Read_all_end(void* buf, MPI::Status& status) (binding
    deprecated, see Section 15.2) }

{void MPI::File::Read_all_end(void* buf) (binding deprecated, see Section 15.2) }

MPI_FILE_WRITE_ALL_BEGIN(fh, buf, count, datatype)
    INOUT    fh                file handle (handle)
    IN       buf                initial address of buffer (choice)
    IN       count              number of elements in buffer (integer)
    IN       datatype           datatype of each buffer element (handle)

int MPI_File_write_all_begin(MPI_File fh, void *buf, int count,
    MPI_Datatype datatype)

MPI_FILE_WRITE_ALL_BEGIN(FH, BUF, COUNT, DATATYPE, IERROR)
    <type> BUF(*)

```



```

1      INTEGER FH, COUNT, DATATYPE, IERROR
2
3      {void MPI::File::Write_all_begin(const void* buf, int count,
4          const MPI::Datatype& datatype) (binding deprecated, see Section 15.2)
5          }
6
7
8      MPI_FILE_WRITE_ALL_END(fh, buf, status)
9
10     INOUT    fh                file handle (handle)
11
12     IN       buf                initial address of buffer (choice)
13
14     OUT      status              status object (Status)
15
16
17     int MPI_File_write_all_end(MPI_File fh, void *buf, MPI_Status *status)
18
19     MPI_FILE_WRITE_ALL_END(FH, BUF, STATUS, IERROR)
20     <type> BUF(*)
21     INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR
22
23     {void MPI::File::Write_all_end(const void* buf, MPI::Status& status) (binding
24         deprecated, see Section 15.2) }
25
26     {void MPI::File::Write_all_end(const void* buf) (binding deprecated, see
27         Section 15.2) }
28
29
30     MPI_FILE_READ_ORDERED_BEGIN(fh, buf, count, datatype)
31
32     INOUT    fh                file handle (handle)
33
34     OUT      buf                initial address of buffer (choice)
35
36     IN       count              number of elements in buffer (integer)
37
38     IN       datatype            datatype of each buffer element (handle)
39
40
41     int MPI_File_read_ordered_begin(MPI_File fh, void *buf, int count,
42         MPI_Datatype datatype)
43
44     MPI_FILE_READ_ORDERED_BEGIN(FH, BUF, COUNT, DATATYPE, IERROR)
45     <type> BUF(*)
46     INTEGER FH, COUNT, DATATYPE, IERROR
47
48     {void MPI::File::Read_ordered_begin(void* buf, int count,
49         const MPI::Datatype& datatype) (binding deprecated, see Section 15.2)
50         }

```

```

MPI_FILE_READ_ORDERED_END(fh, buf, status)
    INOUT    fh                file handle (handle)
    OUT      buf                initial address of buffer (choice)
    OUT      status             status object (Status)

int MPI_File_read_ordered_end(MPI_File fh, void *buf, MPI_Status *status)

MPI_FILE_READ_ORDERED_END(FH, BUF, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR

{void MPI::File::Read_ordered_end(void* buf, MPI::Status& status) (binding
    deprecated, see Section 15.2) }

{void MPI::File::Read_ordered_end(void* buf) (binding deprecated, see Section 15.2)
    }

MPI_FILE_WRITE_ORDERED_BEGIN(fh, buf, count, datatype)
    INOUT    fh                file handle (handle)
    IN       buf                initial address of buffer (choice)
    IN       count              number of elements in buffer (integer)
    IN       datatype           datatype of each buffer element (handle)

int MPI_File_write_ordered_begin(MPI_File fh, void *buf, int count,
    MPI_Datatype datatype)

MPI_FILE_WRITE_ORDERED_BEGIN(FH, BUF, COUNT, DATATYPE, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, IERROR

{void MPI::File::Write_ordered_begin(const void* buf, int count,
    const MPI::Datatype& datatype) (binding deprecated, see Section 15.2)
    }

MPI_FILE_WRITE_ORDERED_END(fh, buf, status)
    INOUT    fh                file handle (handle)
    IN       buf                initial address of buffer (choice)
    OUT      status             status object (Status)

int MPI_File_write_ordered_end(MPI_File fh, void *buf, MPI_Status *status)

MPI_FILE_WRITE_ORDERED_END(FH, BUF, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR

```

```

1 {void MPI::File::Write_ordered_end(const void* buf,
2     MPI::Status& status) (binding deprecated, see Section 15.2) }
3
4 {void MPI::File::Write_ordered_end(const void* buf) (binding deprecated, see
5     Section 15.2) }
6
7
8

```

## 13.5 File Interoperability

At the most basic level, file interoperability is the ability to read the information previously written to a file—not just the bits of data, but the actual information the bits represent. MPI guarantees full interoperability within a single MPI environment, and supports increased interoperability outside that environment through the external data representation (Section 13.5.2, page 463) as well as the data conversion functions (Section 13.5.3, page 464).

Interoperability within a single MPI environment (which could be considered “operability”) ensures that file data written by one MPI process can be read by any other MPI process, subject to the consistency constraints (see Section 13.6.1, page 469), provided that it would have been possible to start the two processes simultaneously and have them reside in a single MPI\_COMM\_WORLD. Furthermore, both processes must see the same data values at every absolute byte offset in the file for which data was written.

This single environment file interoperability implies that file data is accessible regardless of the number of processes.

There are three aspects to file interoperability:

- transferring the bits,
- converting between different file structures, and
- converting between different machine representations.

The first two aspects of file interoperability are beyond the scope of this standard, as both are highly machine dependent. However, transferring the bits of a file into and out of the MPI environment (e.g., by writing a file to tape) is required to be supported by all MPI implementations. In particular, an implementation must specify how familiar operations similar to POSIX `cp`, `rm`, and `mv` can be performed on the file. Furthermore, it is expected that the facility provided maintains the correspondence between absolute byte offsets (e.g., after possible file structure conversion, the data bits at byte offset 102 in the MPI environment are at byte offset 102 outside the MPI environment). As an example, a simple off-line conversion utility that transfers and converts files between the native file system and the MPI environment would suffice, provided it maintained the offset coherence mentioned above. In a high-quality implementation of MPI, users will be able to manipulate MPI files using the same or similar tools that the native file system offers for manipulating its files.

The remaining aspect of file interoperability, converting between different machine representations, is supported by the typing information specified in the `etyp` and `filetype`. This facility allows the information in files to be shared between any two applications, regardless of whether they use MPI, and regardless of the machine architectures on which they run.

MPI supports multiple data representations: “native,” “internal,” and “external32.” An implementation may support additional data representations. MPI also supports user-defined data representations (see Section 13.5.3, page 464). The “native” and “internal”

data representations are implementation dependent, while the “external32” representation is common to all MPI implementations and facilitates file interoperability. The data representation is specified in the *datarep* argument to `MPI_FILE_SET_VIEW`.

*Advice to users.* MPI is not guaranteed to retain knowledge of what data representation was used when a file is written. Therefore, to correctly retrieve file data, an MPI application is responsible for specifying the same data representation as was used to create the file. (*End of advice to users.*)

**“native”** Data in this representation is stored in a file exactly as it is in memory. The advantage of this data representation is that data precision and I/O performance are not lost in type conversions with a purely homogeneous environment. The disadvantage is the loss of transparent interoperability within a heterogeneous MPI environment.

*Advice to users.* This data representation should only be used in a homogeneous MPI environment, or when the MPI application is capable of performing the data type conversions itself. (*End of advice to users.*)

*Advice to implementors.* When implementing read and write operations on top of MPI message-passing, the message data should be typed as `MPI_BYTE` to ensure that the message routines do not perform any type conversions on the data. (*End of advice to implementors.*)

**“internal”** This data representation can be used for I/O operations in a homogeneous or heterogeneous environment; the implementation will perform type conversions if necessary. The implementation is free to store data in any format of its choice, with the restriction that it will maintain constant extents for all predefined datatypes in any one file. The environment in which the resulting file can be reused is implementation-defined and must be documented by the implementation.

*Rationale.* This data representation allows the implementation to perform I/O efficiently in a heterogeneous environment, though with implementation-defined restrictions on how the file can be reused. (*End of rationale.*)

*Advice to implementors.* Since “external32” is a superset of the functionality provided by “internal,” an implementation may choose to implement “internal” as “external32.” (*End of advice to implementors.*)

**“external32”** This data representation states that read and write operations convert all data from and to the “external32” representation defined in Section 13.5.2, page 463. The data conversion rules for communication also apply to these conversions (see Section 3.3.2, page 25-27, of the MPI-1 document). The data on the storage medium is always in this canonical representation, and the data in memory is always in the local process’s native representation.

This data representation has several advantages. First, all processes reading the file in a heterogeneous MPI environment will automatically have the data converted to their respective native representations. Second, the file can be exported from one MPI environment and imported into any other MPI environment with the guarantee that the second environment will be able to read all the data in the file.

The disadvantage of this data representation is that data precision and I/O performance may be lost in data type conversions.

*Advice to implementors.* When implementing read and write operations on top of MPI message-passing, the message data should be converted to and from the “external32” representation in the client, and sent as type `MPI_BYTE`. This will avoid possible double data type conversions and the associated further loss of precision and performance. (*End of advice to implementors.*)

### 13.5.1 Datatypes for File Interoperability

If the file data representation is other than “native,” care must be taken in constructing `etypes` and `filetypes`. Any of the datatype constructor functions may be used; however, for those functions that accept displacements in bytes, the displacements must be specified in terms of their values in the file for the file data representation being used. MPI will interpret these byte displacements as is; no scaling will be done. The function `MPI_FILE_GET_TYPE_EXTENT` can be used to calculate the extents of datatypes in the file. For `etypes` and `filetypes` that are portable datatypes (see Section 2.4, page 11), MPI will scale any displacements in the datatypes to match the file data representation. Datatypes passed as arguments to read/write routines specify the data layout in memory; therefore, they must always be constructed using displacements corresponding to displacements in memory.

*Advice to users.* One can logically think of the file as if it were stored in the memory of a file server. The `etype` and `filetype` are interpreted as if they were defined at this file server, by the same sequence of calls used to define them at the calling process. If the data representation is “native”, then this logical file server runs on the same architecture as the calling process, so that these types define the same data layout on the file as they would define in the memory of the calling process. If the `etype` and `filetype` are portable datatypes, then the data layout defined in the file is the same as would be defined in the calling process memory, up to a scaling factor. The routine `MPI_FILE_GET_FILE_EXTENT` can be used to calculate this scaling factor. Thus, two equivalent, portable datatypes will define the same data layout in the file, even in a heterogeneous environment with “internal”, “external32”, or user defined data representations. Otherwise, the `etype` and `filetype` must be constructed so that their `typemap` and `extent` are the same on any architecture. This can be achieved if they have an explicit upper bound and lower bound (defined either using `MPI_LB` and `MPI_UB` markers, or using `MPI_TYPE_CREATE_RESIZED`). This condition must also be fulfilled by any datatype that is used in the construction of the `etype` and `filetype`, if this datatype is replicated contiguously, either explicitly, by a call to `MPI_TYPE_CONTIGUOUS`, or implicitly, by a `blocklength` argument that is greater than one. If an `etype` or `filetype` is not portable, and has a `typemap` or `extent` that is architecture dependent, then the data layout specified by it on a file is implementation dependent.

File data representations other than “native” may be different from corresponding data representations in memory. Therefore, for these file data representations, it is important not to use hardwired byte offsets for file positioning, including the initial displacement that specifies the view. When a portable datatype (see Section 2.4,

page 11) is used in a data access operation, any holes in the datatype are scaled to match the data representation. However, note that this technique only works when all the processes that created the file view build their etypes from the same predefined datatypes. For example, if one process uses an etype built from MPI\_INT and another uses an etype built from MPI\_FLOAT, the resulting views may be nonportable because the relative sizes of these types may differ from one data representation to another. (*End of advice to users.*)

MPI\_FILE\_GET\_TYPE\_EXTENT(fh, datatype, extent)

IN	fh	file handle (handle)
IN	datatype	datatype (handle)
OUT	extent	datatype extent (integer)

```
int MPI_File_get_type_extent(MPI_File fh, MPI_Datatype datatype,
                             MPI_Aint *extent)
```

```
MPI_FILE_GET_TYPE_EXTENT(FH, DATATYPE, EXTENT, IERROR)
    INTEGER FH, DATATYPE, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTENT
```

```
{MPI::Aint MPI::File::Get_type_extent(const MPI::Datatype& datatype)
    const(binding deprecated, see Section 15.2) }
```

Returns the extent of `datatype` in the file `fh`. This extent will be the same for all processes accessing the file `fh`. If the current view uses a user-defined data representation (see Section 13.5.3, page 464), MPI uses the `dtype_file_extent_fn` callback to calculate the extent.

*Advice to implementors.* In the case of user-defined data representations, the extent of a derived datatype can be calculated by first determining the extents of the predefined datatypes in this derived datatype using `dtype_file_extent_fn` (see Section 13.5.3, page 464). (*End of advice to implementors.*)

### 13.5.2 External Data Representation: “external32”

All MPI implementations are required to support the data representation defined in this section. Support of optional datatypes (e.g., MPI\_INTEGER2) is not required.

All floating point values are in big-endian IEEE format [31] of the appropriate size. Floating point values are represented by one of three IEEE formats. These are the IEEE “Single,” “Double,” and “Double Extended” formats, requiring 4, 8 and 16 bytes of storage, respectively. For the IEEE “Double Extended” formats, MPI specifies a Format Width of 16 bytes, with 15 exponent bits, bias = +16383, 112 fraction bits, and an encoding analogous to the “Double” format. All integral values are in two’s complement big-endian format. Big-endian means most significant byte at lowest address byte. For C `_Bool`, Fortran LOGICAL and C++ `bool`, 0 implies false and nonzero implies true. C `float` `_Complex`, `double` `_Complex` and long `double` `_Complex` as well as Fortran COMPLEX and DOUBLE COMPLEX are represented by a pair of floating point format values for the real and imaginary components.

Characters are in ISO 8859-1 format [32]. Wide characters (of type MPI\_WCHAR) are in Unicode format [52].

All signed numerals (e.g., MPI\_INT, MPI\_REAL) have the sign bit at the most significant bit. MPI\_COMPLEX and MPI\_DOUBLE\_COMPLEX have the sign bit of the real and imaginary parts at the most significant bit of each part.

According to IEEE specifications [31], the “NaN” (not a number) is system dependent. It should not be interpreted within MPI as anything other than “NaN.”

*Advice to implementors.* The MPI treatment of “NaN” is similar to the approach used in XDR (see <ftp://ds.internic.net/rfc/rfc1832.txt>). (*End of advice to implementors.*)

All data is byte aligned, regardless of type. All data items are stored contiguously in the file (if the file view is contiguous).

*Advice to implementors.* All bytes of LOGICAL and bool must be checked to determine the value. (*End of advice to implementors.*)

*Advice to users.* The type MPI\_PACKED is treated as bytes and is not converted. The user should be aware that MPI\_PACK has the option of placing a header in the beginning of the pack buffer. (*End of advice to users.*)

The size of the predefined datatypes returned from MPI\_TYPE\_CREATE\_F90\_REAL, MPI\_TYPE\_CREATE\_F90\_COMPLEX, and MPI\_TYPE\_CREATE\_F90\_INTEGER are defined in Section 16.2.5, page 525.

*Advice to implementors.* When converting a larger size integer to a smaller size integer, only the less significant bytes are moved. Care must be taken to preserve the sign bit value. This allows no conversion errors if the data range is within the range of the smaller size integer. (*End of advice to implementors.*)

Table 13.2 specifies the sizes of predefined datatypes in “external32” format.

### 13.5.3 User-Defined Data Representations

There are two situations that cannot be handled by the required representations:

1. a user wants to write a file in a representation unknown to the implementation, and
2. a user wants to read a file written in a representation unknown to the implementation.

User-defined data representations allow the user to insert a third party converter into the I/O stream to do the data representation conversion.

Type	Length	Optional Type	Length
MPI_PACKED	1	MPI_INTEGER1	1
MPI_BYTE	1	MPI_INTEGER2	2
MPI_CHAR	1	MPI_INTEGER4	4
MPI_UNSIGNED_CHAR	1	MPI_INTEGER8	8
MPI_SIGNED_CHAR	1	MPI_INTEGER16	16
MPI_WCHAR	2		
MPI_SHORT	2	MPI_REAL2	2
MPI_UNSIGNED_SHORT	2	MPI_REAL4	4
MPI_INT	4	MPI_REAL8	8
MPI_UNSIGNED	4	MPI_REAL16	16
MPI_LONG	4		
MPI_UNSIGNED_LONG	4	MPI_COMPLEX4	2*2
MPI_LONG_LONG_INT	8	MPI_COMPLEX8	2*4
MPI_UNSIGNED_LONG_LONG	8	MPI_COMPLEX16	2*8
MPI_FLOAT	4	MPI_COMPLEX32	2*16
MPI_DOUBLE	8		
MPI_LONG_DOUBLE	16		
MPI_C_BOOL	4		
MPI_INT8_T	1		
MPI_INT16_T	2		
MPI_INT32_T	4		
MPI_INT64_T	8		
MPI_UINT8_T	1		
MPI_UINT16_T	2		
MPI_UINT32_T	4		
MPI_UINT64_T	8		
MPI_AINT	8		
[ticket265.]MPI_COUNT		8	
MPI_OFFSET	8		
MPI_C_COMPLEX	2*4		
MPI_C_FLOAT_COMPLEX	2*4		
MPI_C_DOUBLE_COMPLEX	2*8		
MPI_C_LONG_DOUBLE_COMPLEX	2*16		
MPI_CHARACTER	1		
MPI_LOGICAL	4		
MPI_INTEGER	4		
MPI_REAL	4		
MPI_DOUBLE_PRECISION	8		
MPI_COMPLEX	2*4		
MPI_DOUBLE_COMPLEX	2*8		

Table 13.2: “external32” sizes of predefined datatypes



```

1  MPI_REGISTER_DATAREP(datarep, read_conversion_fn, write_conversion_fn,
2      dtype_file_extent_fn, extra_state)
3
4  IN      datarep      data representation identifier (string)
5  IN      read_conversion_fn  function invoked to convert from file representation to
6      native representation (function)
7  IN      write_conversion_fn  function invoked to convert from native representation
8      to file representation (function)
9  IN      dtype_file_extent_fn  function invoked to get the extent of a datatype as
10     represented in the file (function)
11
12  IN      extra_state      extra state
13
14  int MPI_Register_datarep(char *datarep,
15      MPI_Datarep_conversion_function *read_conversion_fn,
16      MPI_Datarep_conversion_function *write_conversion_fn,
17      MPI_Datarep_extent_function *dtype_file_extent_fn,
18      void *extra_state)
19
20  MPI_REGISTER_DATAREP(DATAREP, READ_CONVERSION_FN, WRITE_CONVERSION_FN,
21      DTYPE_FILE_EXTENT_FN, EXTRA_STATE, IERROR)
22  CHARACTER*(*) DATAREP
23  EXTERNAL READ_CONVERSION_FN, WRITE_CONVERSION_FN, DTYPE_FILE_EXTENT_FN
24  INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
25  INTEGER IERROR
26
27  {void MPI::Register_datarep(const char* datarep,
28      MPI::Datarep_conversion_function* read_conversion_fn,
29      MPI::Datarep_conversion_function* write_conversion_fn,
30      MPI::Datarep_extent_function* dtype_file_extent_fn,
31      void* extra_state) (binding deprecated, see Section 15.2) }

```

The call associates `read_conversion_fn`, `write_conversion_fn`, and `dtype_file_extent_fn` with the data representation identifier `datarep`. `datarep` can then be used as an argument to `MPI_FILE_SET_VIEW`, causing subsequent data access operations to call the conversion functions to convert all data items accessed between file data representation and native representation. `MPI_REGISTER_DATAREP` is a local operation and only registers the data representation for the calling MPI process. If `datarep` is already defined, an error in the error class `MPI_ERR_DUP_DATAREP` is raised using the default file error handler (see Section 13.7, page 479). The length of a data representation string is limited to the value of `MPI_MAX_DATAREP_STRING`. `MPI_MAX_DATAREP_STRING` must have a value of at least 64. No routines are provided to delete data representations and free the associated resources; it is not expected that an application will generate them in significant numbers.

#### Extent Callback

```

45  typedef int MPI_Datarep_extent_function(MPI_Datatype datatype,
46      MPI_Aint *file_extent, void *extra_state);
47
48  SUBROUTINE DATAREP_EXTENT_FUNCTION(DATATYPE, EXTENT, EXTRA_STATE, IERROR)

```

```

INTEGER DATATYPE, IERROR
INTEGER(KIND=MPI_ADDRESS_KIND) EXTENT, EXTRA_STATE
{typedef void MPI::Datarep_extent_function(const MPI::Datatype& datatype,
MPI::Aint& file_extent, void* extra_state); (binding deprecated,
see Section 15.2)}
```

The function `dtype_file_extent_fn` must return, in `file_extent`, the number of bytes required to store `datatype` in the file representation. The function is passed, in `extra_state`, the argument that was passed to the `MPI_REGISTER_DATAREP` call. MPI will only call this routine with predefined datatypes employed by the user.

#### Datarep Conversion Functions

```

typedef int MPI_Datarep_conversion_function(void *userbuf,
MPI_Datatype datatype, int count, void *filebuf,
MPI_Offset position, void *extra_state);

SUBROUTINE DATAREP_CONVERSION_FUNCTION(USERBUF, DATATYPE, COUNT, FILEBUF,
POSITION, EXTRA_STATE, IERROR)
<TYPE> USERBUF(*), FILEBUF(*)
INTEGER COUNT, DATATYPE, IERROR
INTEGER(KIND=MPI_OFFSET_KIND) POSITION
INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE

{typedef void MPI::Datarep_conversion_function(void* userbuf,
MPI::Datatype& datatype, int count, void* filebuf,
MPI::Offset position, void* extra_state); (binding deprecated, see
Section 15.2)}
```

The function `read_conversion_fn` must convert from file data representation to native representation. Before calling this routine, MPI allocates and fills `filebuf` with `count` contiguous data items. The type of each data item matches the corresponding entry for the predefined datatype in the type signature of `datatype`. The function is passed, in `extra_state`, the argument that was passed to the `MPI_REGISTER_DATAREP` call. The function must copy all `count` data items from `filebuf` to `userbuf` in the distribution described by `datatype`, converting each data item from file representation to native representation. `datatype` will be equivalent to the datatype that the user passed to the read function. If the size of `datatype` is less than the size of the `count` data items, the conversion function must treat `datatype` as being contiguously tiled over the `userbuf`. The conversion function must begin storing converted data at the location in `userbuf` specified by `position` into the (tiled) `datatype`.

*Advice to users.* Although the conversion functions have similarities to `MPI_PACK` and `MPI_UNPACK`, one should note the differences in the use of the arguments `count` and `position`. In the conversion functions, `count` is a count of data items (i.e., count of typemap entries of `datatype`), and `position` is an index into this typemap. In `MPI_PACK`, `incount` refers to the number of whole datatypes, and `position` is a number of bytes. (*End of advice to users.*)

*Advice to implementors.* A converted read operation could be implemented as follows:

1. Get file extent of all data items
2. Allocate a filebuf large enough to hold all count data items
3. Read data from file into filebuf
4. Call `read_conversion_fn` to convert data and place it into userbuf
5. Deallocate filebuf

*(End of advice to implementors.)*

If MPI cannot allocate a buffer large enough to hold all the data to be converted from a read operation, it may call the conversion function repeatedly using the same `datatype` and `userbuf`, and reading successive chunks of data to be converted in `filebuf`. For the first call (and in the case when all the data to be converted fits into `filebuf`), MPI will call the function with `position` set to zero. Data converted during this call will be stored in the `userbuf` according to the first `count` data items in `datatype`. Then in subsequent calls to the conversion function, MPI will increment the value in `position` by the `count` of items converted in the previous call, and the `userbuf` pointer will be unchanged.

*Rationale.* Passing the conversion function a position and one datatype for the transfer allows the conversion function to decode the datatype only once and cache an internal representation of it on the datatype. Then on subsequent calls, the conversion function can use the `position` to quickly find its place in the datatype and continue storing converted data where it left off at the end of the previous call. *(End of rationale.)*

*Advice to users.* Although the conversion function may usefully cache an internal representation on the datatype, it should not cache any state information specific to an ongoing conversion operation, since it is possible for the same datatype to be used concurrently in multiple conversion operations. *(End of advice to users.)*

The function `write_conversion_fn` must convert from native representation to file data representation. Before calling this routine, MPI allocates `filebuf` of a size large enough to hold `count` contiguous data items. The type of each data item matches the corresponding entry for the predefined datatype in the type signature of `datatype`. The function must copy `count` data items from `userbuf` in the distribution described by `datatype`, to a contiguous distribution in `filebuf`, converting each data item from native representation to file representation. If the size of `datatype` is less than the size of `count` data items, the conversion function must treat `datatype` as being contiguously tiled over the `userbuf`.

The function must begin copying at the location in `userbuf` specified by `position` into the (tiled) `datatype`. `datatype` will be equivalent to the datatype that the user passed to the write function. The function is passed, in `extra_state`, the argument that was passed to the `MPI_REGISTER_DATAREP` call.

The predefined constant `MPI_CONVERSION_FN_NULL` may be used as either `write_conversion_fn` or `read_conversion_fn`. In that case, MPI will not attempt to invoke `write_conversion_fn` or `read_conversion_fn`, respectively, but will perform the requested data access using the native data representation.

An MPI implementation must ensure that all data accessed is converted, either by using a `filebuf` large enough to hold all the requested data items or else by making repeated

calls to the conversion function with the same `datatype` argument and appropriate values for `position`.

An implementation will only invoke the callback routines in this section (`read_conversion_fn`, `write_conversion_fn`, and `dtype_file_extent_fn`) when one of the read or write routines in Section 13.4, page 436, or `MPI_FILE_GET_TYPE_EXTENT` is called by the user. `dtype_file_extent_fn` will only be passed predefined datatypes employed by the user. The conversion functions will only be passed datatypes equivalent to those that the user has passed to one of the routines noted above.

The conversion functions must be reentrant. User defined data representations are restricted to use byte alignment for all types. Furthermore, it is erroneous for the conversion functions to call any collective routines or to free `datatype`.

The conversion functions should return an error code. If the returned error code has a value other than `MPI_SUCCESS`, the implementation will raise an error in the class `MPI_ERR_CONVERSION`.

### 13.5.4 Matching Data Representations

It is the user’s responsibility to ensure that the data representation used to read data from a file is *compatible* with the data representation that was used to write that data to the file.

In general, using the same data representation name when writing and reading a file does not guarantee that the representation is compatible. Similarly, using different representation names on two different implementations may yield compatible representations.

Compatibility can be obtained when “external32” representation is used, although precision may be lost and the performance may be less than when “native” representation is used. Compatibility is guaranteed using “external32” provided at least one of the following conditions is met.

- The data access routines directly use types enumerated in Section 13.5.2, page 463, that are supported by all implementations participating in the I/O. The predefined type used to write a data item must also be used to read a data item.
- In the case of Fortran 90 programs, the programs participating in the data accesses obtain compatible datatypes using MPI routines that specify precision and/or range (Section 16.2.5, page 521).
- For any given data item, the programs participating in the data accesses use compatible predefined types to write and read the data item.

User-defined data representations may be used to provide an implementation compatibility with another implementation’s “native” or “internal” representation.

*Advice to users.* Section 16.2.5, page 521, defines routines that support the use of matching datatypes in heterogeneous environments and contains examples illustrating their use. (*End of advice to users.*)

## 13.6 Consistency and Semantics

### 13.6.1 File Consistency

Consistency semantics define the outcome of multiple accesses to a single file. All file accesses in MPI are relative to a specific file handle created from a collective open. MPI

provides three levels of consistency: sequential consistency among all accesses using a single file handle, sequential consistency among all accesses using file handles created from a single collective open with atomic mode enabled, and user-imposed consistency among accesses other than the above. Sequential consistency means the behavior of a set of operations will be as if the operations were performed in some serial order consistent with program order; each access appears atomic, although the exact ordering of accesses is unspecified. User-imposed consistency may be obtained using program order and calls to `MPI_FILE_SYNC`.

Let  $FH_1$  be the set of file handles created from one particular collective open of the file  $FOO$ , and  $FH_2$  be the set of file handles created from a different collective open of  $FOO$ . Note that nothing restrictive is said about  $FH_1$  and  $FH_2$ : the sizes of  $FH_1$  and  $FH_2$  may be different, the groups of processes used for each open may or may not intersect, the file handles in  $FH_1$  may be destroyed before those in  $FH_2$  are created, etc. Consider the following three cases: a single file handle (e.g.,  $fh_1 \in FH_1$ ), two file handles created from a single collective open (e.g.,  $fh_{1a} \in FH_1$  and  $fh_{1b} \in FH_1$ ), and two file handles from different collective opens (e.g.,  $fh_1 \in FH_1$  and  $fh_2 \in FH_2$ ).

For the purpose of consistency semantics, a matched pair (Section 13.4.5, page 453) of split collective data access operations (e.g., `MPI_FILE_READ_ALL_BEGIN` and `MPI_FILE_READ_ALL_END`) compose a single data access operation. Similarly, a non-blocking data access routine (e.g., `MPI_FILE_IREAD`) and the routine which completes the request (e.g., `MPI_WAIT`) also compose a single data access operation. For all cases below, these data access operations are subject to the same constraints as blocking data access operations.

*Advice to users.* For an `MPI_FILE_IREAD` and `MPI_WAIT` pair, the operation begins when `MPI_FILE_IREAD` is called and ends when `MPI_WAIT` returns. (*End of advice to users.*)

Assume that  $A_1$  and  $A_2$  are two data access operations. Let  $D_1$  ( $D_2$ ) be the set of absolute byte displacements of every byte accessed in  $A_1$  ( $A_2$ ). The two data accesses *overlap* if  $D_1 \cap D_2 \neq \emptyset$ . The two data accesses *conflict* if they overlap and at least one is a write access.

Let  $SEQ_{fh}$  be a sequence of file operations on a single file handle, bracketed by `MPI_FILE_SYNC`s on that file handle. (Both opening and closing a file implicitly perform an `MPI_FILE_SYNC`.)  $SEQ_{fh}$  is a “write sequence” if any of the data access operations in the sequence are writes or if any of the file manipulation operations in the sequence change the state of the file (e.g., `MPI_FILE_SET_SIZE` or `MPI_FILE_PREALLOCATE`). Given two sequences,  $SEQ_1$  and  $SEQ_2$ , we say they are not *concurrent* if one sequence is guaranteed to completely precede the other (temporally).

The requirements for guaranteeing sequential consistency among all accesses to a particular file are divided into the three cases given below. If any of these requirements are not met, then the value of all data in that file is implementation dependent.

**Case 1:**  $fh_1 \in FH_1$  All operations on  $fh_1$  are sequentially consistent if atomic mode is set. If nonatomic mode is set, then all operations on  $fh_1$  are sequentially consistent if they are either nonconcurrent, nonconflicting, or both.

**Case 2:**  $fh_{1a} \in FH_1$  and  $fh_{1b} \in FH_1$  Assume  $A_1$  is a data access operation using  $fh_{1a}$ , and  $A_2$  is a data access operation using  $fh_{1b}$ . If for any access  $A_1$ , there is no access  $A_2$  that conflicts with  $A_1$ , then MPI guarantees sequential consistency.

However, unlike POSIX semantics, the default MPI semantics for conflicting accesses do not guarantee sequential consistency. If  $A_1$  and  $A_2$  conflict, sequential consistency can be guaranteed by either enabling atomic mode via the `MPI_FILE_SET_ATOMICITY` routine, or meeting the condition described in Case 3 below.

**Case 3:**  $fh_1 \in FH_1$  and  $fh_2 \in FH_2$  Consider access to a single file using file handles from distinct collective opens. In order to guarantee sequential consistency, `MPI_FILE_SYNC` must be used (both opening and closing a file implicitly perform an `MPI_FILE_SYNC`).

Sequential consistency is guaranteed among accesses to a single file if for any write sequence  $SEQ_1$  to the file, there is no sequence  $SEQ_2$  to the file which is *concurrent* with  $SEQ_1$ . To guarantee sequential consistency when there are write sequences, `MPI_FILE_SYNC` must be used together with a mechanism that guarantees nonconcurrency of the sequences.

See the examples in Section 13.6.10, page 475, for further clarification of some of these consistency semantics.

`MPI_FILE_SET_ATOMICITY(fh, flag)`

INOUT	fh	file handle (handle)
IN	flag	true to set atomic mode, false to set nonatomic mode (logical)

`int MPI_File_set_atomicity(MPI_File fh, int flag)`

`MPI_FILE_SET_ATOMICITY(FH, FLAG, IERROR)`

INTEGER FH, IERROR

LOGICAL FLAG

`{void MPI::File::Set_atomicity(bool flag) (binding deprecated, see Section 15.2) }`

Let  $FH$  be the set of file handles created by one collective open. The consistency semantics for data access operations using  $FH$  is set by collectively calling `MPI_FILE_SET_ATOMICITY` on  $FH$ . `MPI_FILE_SET_ATOMICITY` is collective; all processes in the group must pass identical values for `fh` and `flag`. If `flag` is true, atomic mode is set; if `flag` is false, nonatomic mode is set.

Changing the consistency semantics for an open file only affects new data accesses. All completed data accesses are guaranteed to abide by the consistency semantics in effect during their execution. Nonblocking data accesses and split collective operations that have not completed (e.g., via `MPI_WAIT`) are only guaranteed to abide by nonatomic mode consistency semantics.

*Advice to implementors.* Since the semantics guaranteed by atomic mode are stronger than those guaranteed by nonatomic mode, an implementation is free to adhere to the more stringent atomic mode semantics for outstanding requests. (*End of advice to implementors.*)

1 MPI\_FILE\_GET\_ATOMICITY(fh, flag)

2     IN        fh                           file handle (handle)

3     OUT       flag                        true if atomic mode, false if nonatomic mode (logical)

4     5  
6 int MPI\_File\_get\_atomicity(MPI\_File fh, int \*flag)

7 MPI\_FILE\_GET\_ATOMICITY(FH, FLAG, IERROR)

8     INTEGER FH, IERROR

9     LOGICAL FLAG

10  
11 {bool MPI::File::Get\_atomicity() const(*binding deprecated, see Section 15.2*) }

12  
13 MPI\_FILE\_GET\_ATOMICITY returns the current consistency semantics for data access  
14 operations on the set of file handles created by one collective open. If flag is true, atomic  
15 mode is enabled; if flag is false, nonatomic mode is enabled.

16  
17 MPI\_FILE\_SYNC(fh)

18     INOUT    fh                           file handle (handle)

19  
20  
21 int MPI\_File\_sync(MPI\_File fh)

22 MPI\_FILE\_SYNC(FH, IERROR)

23     INTEGER FH, IERROR

24  
25 {void MPI::File::Sync() (*binding deprecated, see Section 15.2*) }

26  
27 Calling MPI\_FILE\_SYNC with fh causes all previous writes to fh by the calling process  
28 to be transferred to the storage device. If other processes have made updates to the storage  
29 device, then all such updates become visible to subsequent reads of fh by the calling process.  
30 MPI\_FILE\_SYNC may be necessary to ensure sequential consistency in certain cases (see  
31 above).

32 MPI\_FILE\_SYNC is a collective operation.

33 The user is responsible for ensuring that all nonblocking requests and split collective  
34 operations on fh have been completed before calling MPI\_FILE\_SYNC—otherwise, the call  
35 to MPI\_FILE\_SYNC is erroneous.

### 36 13.6.2 Random Access vs. Sequential Files

37  
38 MPI distinguishes ordinary random access files from sequential stream files, such as pipes  
39 and tape files. Sequential stream files must be opened with the MPI\_MODE\_SEQUENTIAL  
40 flag set in the amode. For these files, the only permitted data access operations are shared  
41 file pointer reads and writes. Filetypes and etypes with holes are erroneous. In addition, the  
42 notion of file pointer is not meaningful; therefore, calls to MPI\_FILE\_SEEK\_SHARED and  
43 MPI\_FILE\_GET\_POSITION\_SHARED are erroneous, and the pointer update rules specified  
44 for the data access routines do not apply. The amount of data accessed by a data access  
45 operation will be the amount requested unless the end of file is reached or an error is raised.

46  
47 *Rationale.* This implies that reading on a pipe will always wait until the requested  
48 amount of data is available or until the process writing to the pipe has issued an end



of file. (*End of rationale.*)

Finally, for some sequential files, such as those corresponding to magnetic tapes or streaming network connections, writes to the file may be destructive. In other words, a write may act as a truncate (a `MPI_FILE_SET_SIZE` with `size` set to the current position) followed by the write.

### 13.6.3 Progress

The progress rules of MPI are both a promise to users and a set of constraints on implementors. In cases where the progress rules restrict possible implementation choices more than the interface specification alone, the progress rules take precedence.

All blocking routines must complete in finite time unless an exceptional condition (such as resource exhaustion) causes an error.

Nonblocking data access routines inherit the following progress rule from nonblocking point to point communication: a nonblocking write is equivalent to a nonblocking send for which a receive is eventually posted, and a nonblocking read is equivalent to a nonblocking receive for which a send is eventually posted.

Finally, an implementation is free to delay progress of collective routines until all processes in the group associated with the collective call have invoked the routine. Once all processes in the group have invoked the routine, the progress rule of the equivalent noncollective routine must be followed.

### 13.6.4 Collective File Operations

Collective file operations are subject to the same restrictions as collective communication operations. For a complete discussion, please refer to the semantics set forth in Section 5.13 on page 210.

Collective file operations are collective over a dup of the communicator used to open the file—this duplicate communicator is implicitly specified via the file handle argument. Different processes can pass different values for other arguments of a collective routine unless specified otherwise.

### 13.6.5 Type Matching

The type matching rules for I/O mimic the type matching rules for communication with one exception: if `etype` is `MPI_BYTE`, then this matches any `datatype` in a data access operation. In general, the `etype` of data items written must match the `etype` used to read the items, and for each data access operation, the current `etype` must also match the type declaration of the data access buffer.

*Advice to users.* In most cases, use of `MPI_BYTE` as a wild card will defeat the file interoperability features of MPI. File interoperability can only perform automatic conversion between heterogeneous data representations when the exact `datatypes` accessed are explicitly specified. (*End of advice to users.*)

### 13.6.6 Miscellaneous Clarifications

Once an I/O routine completes, it is safe to free any opaque objects passed as arguments to that routine. For example, the `comm` and `info` used in an `MPI_FILE_OPEN`, or the `etype`



and `filetype` used in an `MPI_FILE_SET_VIEW`, can be freed without affecting access to the file. Note that for nonblocking routines and split collective operations, the operation must be completed before it is safe to reuse data buffers passed as arguments.

As in communication, datatypes must be committed before they can be used in file manipulation or data access operations. For example, the `etype` and `filetype` must be committed before calling `MPI_FILE_SET_VIEW`, and the `datatype` must be committed before calling `MPI_FILE_READ` or `MPI_FILE_WRITE`.

### 13.6.7 MPI\_Offset Type

`MPI_Offset` is an integer type of size sufficient to represent the size (in bytes) of the largest file supported by MPI. Displacements and offsets are always specified as values of type `MPI_Offset`.

In Fortran, the corresponding integer is an integer of kind `MPI_OFFSET_KIND`, defined in `mpif.h` and the `mpi` module.

In Fortran 77 environments that do not support `KIND` parameters, `MPI_Offset` arguments should be declared as an `INTEGER` of suitable size. The language interoperability implications for `MPI_Offset` are similar to those for addresses (see Section 16.3, page 529).

### 13.6.8 Logical vs. Physical File Layout

MPI specifies how the data should be laid out in a virtual file structure (the view), not how that file structure is to be stored on one or more disks. Specification of the physical file structure was avoided because it is expected that the mapping of files to disks will be system specific, and any specific control over file layout would therefore restrict program portability. However, there are still cases where some information may be necessary to optimize file layout. This information can be provided as *hints* specified via *info* when a file is created (see Section 13.2.8, page 430).

### 13.6.9 File Size

The size of a file may be increased by writing to the file after the current end of file. The size may also be changed by calling MPI *size changing* routines, such as `MPI_FILE_SET_SIZE`. A call to a size changing routine does not necessarily change the file size. For example, calling `MPI_FILE_PREALLOCATE` with a size less than the current size does not change the size.

Consider a set of bytes that has been written to a file since the most recent call to a size changing routine, or since `MPI_FILE_OPEN` if no such routine has been called. Let the *high byte* be the byte in that set with the largest displacement. The file size is the larger of

- One plus the displacement of the high byte.
- The size immediately after the size changing routine, or `MPI_FILE_OPEN`, returned.

When applying consistency semantics, calls to `MPI_FILE_SET_SIZE` and `MPI_FILE_PREALLOCATE` are considered writes to the file (which conflict with operations that access bytes at displacements between the old and new file sizes), and `MPI_FILE_GET_SIZE` is considered a read of the file (which overlaps with all accesses to the file).

*Advice to users.* Any sequence of operations containing the collective routines `MPI_FILE_SET_SIZE` and `MPI_FILE_PREALLOCATE` is a write sequence. As such, sequential consistency in nonatomic mode is not guaranteed unless the conditions in Section 13.6.1, page 469, are satisfied. (*End of advice to users.*)

File pointer update semantics (i.e., file pointers are updated by the amount accessed) are only guaranteed if file size changes are sequentially consistent.

*Advice to users.* Consider the following example. Given two operations made by separate processes to a file containing 100 bytes: an `MPI_FILE_READ` of 10 bytes and an `MPI_FILE_SET_SIZE` to 0 bytes. If the user does not enforce sequential consistency between these two operations, the file pointer may be updated by the amount requested (10 bytes) even if the amount accessed is zero bytes. (*End of advice to users.*)

### 13.6.10 Examples

The examples in this section illustrate the application of the MPI consistency and semantics guarantees. These address

- conflicting accesses on file handles obtained from a single collective open, and
- all accesses on file handles obtained from two separate collective opens.

The simplest way to achieve consistency for conflicting accesses is to obtain sequential consistency by setting atomic mode. For the code below, process 1 will read either 0 or 10 integers. If the latter, every element of `b` will be 5. If nonatomic mode is set, the results of the read are undefined.

```
/* Process 0 */
int i, a[10] ;
int TRUE = 1;

for ( i=0;i<10;i++)
    a[i] = 5 ;

MPI_File_open( MPI_COMM_WORLD, "workfile",
               MPI_MODE_RDWR | MPI_MODE_CREATE, MPI_INFO_NULL, &fh0 ) ;
MPI_File_set_view( fh0, 0, MPI_INT, MPI_INT, "native", MPI_INFO_NULL ) ;
MPI_File_set_atomicity( fh0, TRUE ) ;
MPI_File_write_at(fh0, 0, a, 10, MPI_INT, &status) ;
/* MPI_Barrier( MPI_COMM_WORLD ) ; */

/* Process 1 */
int b[10] ;
int TRUE = 1;
MPI_File_open( MPI_COMM_WORLD, "workfile",
               MPI_MODE_RDWR | MPI_MODE_CREATE, MPI_INFO_NULL, &fh1 ) ;
MPI_File_set_view( fh1, 0, MPI_INT, MPI_INT, "native", MPI_INFO_NULL ) ;
MPI_File_set_atomicity( fh1, TRUE ) ;
/* MPI_Barrier( MPI_COMM_WORLD ) ; */
MPI_File_read_at(fh1, 0, b, 10, MPI_INT, &status) ;
```

A user may guarantee that the write on process 0 precedes the read on process 1 by imposing temporal order with, for example, calls to `MPI_BARRIER`.

*Advice to users.* Routines other than `MPI_BARRIER` may be used to impose temporal order. In the example above, process 0 could use `MPI_SEND` to send a 0 byte message, received by process 1 using `MPI_RECV`. (*End of advice to users.*)

Alternatively, a user can impose consistency with nonatomic mode set:

```

/* Process 0 */
int i, a[10] ;
for ( i=0;i<10;i++)
    a[i] = 5 ;

MPI_File_open( MPI_COMM_WORLD, "workfile",
               MPI_MODE_RDWR | MPI_MODE_CREATE, MPI_INFO_NULL, &fh0 ) ;
MPI_File_set_view( fh0, 0, MPI_INT, MPI_INT, "native", MPI_INFO_NULL ) ;
MPI_File_write_at(fh0, 0, a, 10, MPI_INT, &status ) ;
MPI_File_sync( fh0 ) ;
MPI_Barrier( MPI_COMM_WORLD ) ;
MPI_File_sync( fh0 ) ;

/* Process 1 */
int b[10] ;
MPI_File_open( MPI_COMM_WORLD, "workfile",
               MPI_MODE_RDWR | MPI_MODE_CREATE, MPI_INFO_NULL, &fh1 ) ;
MPI_File_set_view( fh1, 0, MPI_INT, MPI_INT, "native", MPI_INFO_NULL ) ;
MPI_File_sync( fh1 ) ;
MPI_Barrier( MPI_COMM_WORLD ) ;
MPI_File_sync( fh1 ) ;
MPI_File_read_at(fh1, 0, b, 10, MPI_INT, &status ) ;

```

The “sync-barrier-sync” construct is required because:

- The barrier ensures that the write on process 0 occurs before the read on process 1.
- The first sync guarantees that the data written by all processes is transferred to the storage device.
- The second sync guarantees that all data which has been transferred to the storage device is visible to all processes. (This does not affect process 0 in this example.)

The following program represents an erroneous attempt to achieve consistency by eliminating the apparently superfluous second “sync” call for each process.

```

/* ----- THIS EXAMPLE IS ERRONEOUS ----- */
/* Process 0 */
int i, a[10] ;
for ( i=0;i<10;i++)
    a[i] = 5 ;

```

```

MPI_File_open( MPI_COMM_WORLD, "workfile",
               MPI_MODE_RDWR | MPI_MODE_CREATE, MPI_INFO_NULL, &fh0 ) ;
MPI_File_set_view( fh0, 0, MPI_INT, MPI_INT, "native", MPI_INFO_NULL ) ;
MPI_File_write_at(fh0, 0, a, 10, MPI_INT, &status ) ;
MPI_File_sync( fh0 ) ;
MPI_Barrier( MPI_COMM_WORLD ) ;

/* Process 1 */
int b[10] ;
MPI_File_open( MPI_COMM_WORLD, "workfile",
               MPI_MODE_RDWR | MPI_MODE_CREATE, MPI_INFO_NULL, &fh1 ) ;
MPI_File_set_view( fh1, 0, MPI_INT, MPI_INT, "native", MPI_INFO_NULL ) ;
MPI_Barrier( MPI_COMM_WORLD ) ;
MPI_File_sync( fh1 ) ;
MPI_File_read_at(fh1, 0, b, 10, MPI_INT, &status ) ;

/* ----- THIS EXAMPLE IS ERRONEOUS ----- */

```

The above program also violates the MPI rule against out-of-order collective operations and will deadlock for implementations in which MPI\_FILE\_SYNC blocks.

*Advice to users.* Some implementations may choose to implement MPI\_FILE\_SYNC as a temporally synchronizing function. When using such an implementation, the “sync-barrier-sync” construct above can be replaced by a single “sync.” The results of using such code with an implementation for which MPI\_FILE\_SYNC is not temporally synchronizing is undefined. (*End of advice to users.*)

### Asynchronous I/O

The behavior of asynchronous I/O operations is determined by applying the rules specified above for synchronous I/O operations.

The following examples all access a preexisting file “myfile.” Word 10 in myfile initially contains the integer 2. Each example writes and reads word 10.

First consider the following code fragment:

```

int a = 4, b, TRUE=1;
MPI_File_open( MPI_COMM_WORLD, "myfile",
               MPI_MODE_RDWR, MPI_INFO_NULL, &fh ) ;
MPI_File_set_view( fh, 0, MPI_INT, MPI_INT, "native", MPI_INFO_NULL ) ;
/* MPI_File_set_atomicity( fh, TRUE ) ; Use this to set atomic mode. */
MPI_File_iwrite_at(fh, 10, &a, 1, MPI_INT, &reqs[0]) ;
MPI_File_iread_at(fh, 10, &b, 1, MPI_INT, &reqs[1]) ;
MPI_Waitall(2, reqs, statuses) ;

```

For asynchronous data access operations, MPI specifies that the access occurs at any time between the call to the asynchronous data access routine and the return from the corresponding request complete routine. Thus, executing either the read before the write, or the write before the read is consistent with program order. If atomic mode is set, then MPI guarantees sequential consistency, and the program will read either 2 or 4 into b. If atomic

mode is not set, then sequential consistency is not guaranteed and the program may read something other than 2 or 4 due to the conflicting data access.

Similarly, the following code fragment does not order file accesses:

```
int a = 4, b;
MPI_File_open( MPI_COMM_WORLD, "myfile",
               MPI_MODE_RDWR, MPI_INFO_NULL, &fh ) ;
MPI_File_set_view( fh, 0, MPI_INT, MPI_INT, "native", MPI_INFO_NULL ) ;
/* MPI_File_set_atomics( fh, TRUE ) ; Use this to set atomic mode. */
MPI_File_iwrite_at(fh, 10, &a, 1, MPI_INT, &reqs[0]) ;
MPI_File_iread_at(fh, 10, &b, 1, MPI_INT, &reqs[1]) ;
MPI_Wait(&reqs[0], &status) ;
MPI_Wait(&reqs[1], &status) ;
```

If atomic mode is set, either 2 or 4 will be read into `b`. Again, MPI does not guarantee sequential consistency in nonatomic mode.

On the other hand, the following code fragment:

```
int a = 4, b;
MPI_File_open( MPI_COMM_WORLD, "myfile",
               MPI_MODE_RDWR, MPI_INFO_NULL, &fh ) ;
MPI_File_set_view( fh, 0, MPI_INT, MPI_INT, "native", MPI_INFO_NULL ) ;
MPI_File_iwrite_at(fh, 10, &a, 1, MPI_INT, &reqs[0]) ;
MPI_Wait(&reqs[0], &status) ;
MPI_File_iread_at(fh, 10, &b, 1, MPI_INT, &reqs[1]) ;
MPI_Wait(&reqs[1], &status) ;
```

defines the same ordering as:

```
int a = 4, b;
MPI_File_open( MPI_COMM_WORLD, "myfile",
               MPI_MODE_RDWR, MPI_INFO_NULL, &fh ) ;
MPI_File_set_view( fh, 0, MPI_INT, MPI_INT, "native", MPI_INFO_NULL ) ;
MPI_File_write_at(fh, 10, &a, 1, MPI_INT, &status) ;
MPI_File_read_at(fh, 10, &b, 1, MPI_INT, &status) ;
```

Since

- nonconcurrent operations on a single file handle are sequentially consistent, and
- the program fragments specify an order for the operations,

MPI guarantees that both program fragments will read the value 4 into `b`. There is no need to set atomic mode for this example.

Similar considerations apply to conflicting accesses of the form:

```
MPI_File_write_all_begin(fh,...) ;
MPI_File_iread(fh,...) ;
MPI_Wait(fh,...) ;
MPI_File_write_all_end(fh,...) ;
```

Recall that constraints governing consistency and semantics are not relevant to the following:

```

MPI_File_write_all_begin(fh,...) ;
MPI_File_read_all_begin(fh,...) ;
MPI_File_read_all_end(fh,...) ;
MPI_File_write_all_end(fh,...) ;

```

since split collective operations on the same file handle may not overlap (see Section 13.4.5, page 453).

## 13.7 I/O Error Handling

By default, communication errors are fatal—`MPI_ERRORS_ARE_FATAL` is the default error handler associated with `MPI_COMM_WORLD`. I/O errors are usually less catastrophic (e.g., “file not found”) than communication errors, and common practice is to catch these errors and continue executing. For this reason, MPI provides additional error facilities for I/O.

*Advice to users.* MPI does not specify the state of a computation after an erroneous MPI call has occurred. A high-quality implementation will support the I/O error handling facilities, allowing users to write programs using common practice for I/O. (*End of advice to users.*)

Like communicators, each file handle has an error handler associated with it. The MPI I/O error handling routines are defined in Section 8.3, page 308.

When MPI calls a user-defined error handler resulting from an error on a particular file handle, the first two arguments passed to the file error handler are the file handle and the error code. For I/O errors that are not associated with a valid file handle (e.g., in `MPI_FILE_OPEN` or `MPI_FILE_DELETE`), the first argument passed to the error handler is `MPI_FILE_NULL`,

I/O error handling differs from communication error handling in another important aspect. By default, the predefined error handler for file handles is `MPI_ERRORS_RETURN`. The default file error handler has two purposes: when a new file handle is created (by `MPI_FILE_OPEN`), the error handler for the new file handle is initially set to the default error handler, and I/O routines that have no valid file handle on which to raise an error (e.g., `MPI_FILE_OPEN` or `MPI_FILE_DELETE`) use the default file error handler. The default file error handler can be changed by specifying `MPI_FILE_NULL` as the `fh` argument to `MPI_FILE_SET_ERRHANDLER`. The current value of the default file error handler can be determined by passing `MPI_FILE_NULL` as the `fh` argument to `MPI_FILE_GET_ERRHANDLER`.

*Rationale.* For communication, the default error handler is inherited from `MPI_COMM_WORLD`. In I/O, there is no analogous “root” file handle from which default properties can be inherited. Rather than invent a new global file handle, the default file error handler is manipulated as if it were attached to `MPI_FILE_NULL`. (*End of rationale.*)

## 13.8 I/O Error Classes

The implementation dependent error codes returned by the I/O routines can be converted into the error classes defined in Table 13.3.

In addition, calls to routines in this chapter may raise errors in other MPI classes, such as `MPI_ERR_TYPE`.

<code>MPI_ERR_FILE</code>	Invalid file handle
<code>MPI_ERR_NOT_SAME</code>	Collective argument not identical on all processes, or collective routines called in a different order by different processes
<code>MPI_ERR_AMODE</code>	Error related to the <code>amode</code> passed to <code>MPI_FILE_OPEN</code>
<code>MPI_ERR_UNSUPPORTED_DATAREP</code>	Unsupported <code>datarep</code> passed to <code>MPI_FILE_SET_VIEW</code>
<code>MPI_ERR_UNSUPPORTED_OPERATION</code>	Unsupported operation, such as seeking on a file which supports sequential access only
<code>MPI_ERR_NO_SUCH_FILE</code>	File does not exist
<code>MPI_ERR_FILE_EXISTS</code>	File exists
<code>MPI_ERR_BAD_FILE</code>	Invalid file name (e.g., path name too long)
<code>MPI_ERR_ACCESS</code>	Permission denied
<code>MPI_ERR_NO_SPACE</code>	Not enough space
<code>MPI_ERR_QUOTA</code>	Quota exceeded
<code>MPI_ERR_READ_ONLY</code>	Read-only file or file system
<code>MPI_ERR_FILE_IN_USE</code>	File operation could not be completed, as the file is currently open by some process
<code>MPI_ERR_DUP_DATAREP</code>	Conversion functions could not be registered because a data representation identifier that was already defined was passed to <code>MPI_REGISTER_DATAREP</code>
<code>MPI_ERR_CONVERSION</code>	An error occurred in a user supplied data conversion function.
<code>MPI_ERR_IO</code>	Other I/O error

Table 13.3: I/O Error Classes

## 13.9 Examples

### 13.9.1 Double Buffering with Split Collective I/O

This example shows how to overlap computation and output. The computation is performed by the function `compute_buffer()`.

```

/*=====
 *
 * Function:          double_buffer
 *
 * Synopsis:
 *   void double_buffer(
 *       MPI_File fh,                ** IN
 *       MPI_Datatype buftype,      ** IN
 *       int bufcount                ** IN

```

```

*      )
*
* Description:
*      Performs the steps to overlap computation with a collective write
*      by using a double-buffering technique.
*
* Parameters:
*      fh              previously opened MPI file handle
*      buftype         MPI datatype for memory layout
*                      (Assumes a compatible view has been set on fh)
*      bufcount        # buftype elements to transfer
*-----*/
1
2
3
4
5
6
7
8
9
10
11
12
13
/* this macro switches which buffer "x" is pointing to */
14
#define TOGGLE_PTR(x) (((x)==(buffer1)) ? (x=buffer2) : (x=buffer1))
15
16
void double_buffer( MPI_File fh, MPI_Datatype buftype, int bufcount)
17
{
18
19
    MPI_Status status;      /* status for MPI calls */
20
    float *buffer1, *buffer2; /* buffers to hold results */
21
    float *compute_buf_ptr;  /* destination buffer */
22
                           /* for computing */
23
    float *write_buf_ptr;    /* source for writing */
24
    int done;               /* determines when to quit */
25
26

    /* buffer initialization */
27
    buffer1 = (float *)
28
                malloc(bufcount*sizeof(float)) ;
29
    buffer2 = (float *)
30
                malloc(bufcount*sizeof(float)) ;
31
    compute_buf_ptr = buffer1 ; /* initially point to buffer1 */
32
    write_buf_ptr   = buffer1 ; /* initially point to buffer1 */
33
34
35

    /* DOUBLE-BUFFER prolog:
36
     *   compute buffer1; then initiate writing buffer1 to disk
37
     */
38
    compute_buffer(compute_buf_ptr, bufcount, &done);
39
    MPI_File_write_all_begin(fh, write_buf_ptr, bufcount, buftype);
40
41

    /* DOUBLE-BUFFER steady state:
42
     *   Overlap writing old results from buffer pointed to by write_buf_ptr
43
     *   with computing new results into buffer pointed to by compute_buf_ptr.
44
     *
45
     *   There is always one write-buffer and one compute-buffer in use
46
     *   during steady state.
47
     */
48

```



```

1  while (!done) {
2      TOGGLE_PTR(compute_buf_ptr);
3      compute_buffer(compute_buf_ptr, bufcount, &done);
4      MPI_File_write_all_end(fh, write_buf_ptr, &status);
5      TOGGLE_PTR(write_buf_ptr);
6      MPI_File_write_all_begin(fh, write_buf_ptr, bufcount, buftype);
7  }
8
9  /* DOUBLE-BUFFER epilog:
10     *   wait for final write to complete.
11     */
12  MPI_File_write_all_end(fh, write_buf_ptr, &status);
13
14
15  /* buffer cleanup */
16  free(buffer1);
17  free(buffer2);
18  }
19

```

### 13.9.2 Subarray Filetype Constructor




Figure 13.4: Example array file layout

Assume we are writing out a 100x100 2D array of double precision floating point numbers that is distributed among 4 processes such that each process has a block of 25 columns (e.g., process 0 has columns 0-24, process 1 has columns 25-49, etc.; see Figure 13.4). To create the filetypes for each process one could use the following C program (see Section 4.1.3 on page 89):

```

43  double subarray[100][25];
44  MPI_Datatype filetype;
45  int sizes[2], subsizes[2], starts[2];
46  int rank;
47

```




Figure 13.5: Example local array filetype for process 1

```

MPI_Comm_rank(MPI_COMM_WORLD, &rank);
sizes[0]=100; sizes[1]=100;
subsizes[0]=100; subsizes[1]=25;
starts[0]=0; starts[1]=rank*subsizes[1];

MPI_Type_create_subarray(2, sizes, subsizes, starts, MPI_ORDER_C,
                        MPI_DOUBLE, &filetype);

```

Or, equivalently in Fortran:

```

double precision subarray(100,25)
integer filetype, rank, ierror
integer sizes(2), subsizes(2), starts(2)

call MPI_COMM_RANK(MPI_COMM_WORLD, rank, ierror)
sizes(1)=100
sizes(2)=100
subsizes(1)=100
subsizes(2)=25
starts(1)=0
starts(2)=rank*subsizes(2)

call MPI_TYPE_CREATE_SUBARRAY(2, sizes, subsizes, starts, &
                             MPI_ORDER_FORTRAN, MPI_DOUBLE_PRECISION, &
                             filetype, ierror)

```

The generated filetype will then describe the portion of the file contained within the process's subarray with holes for the space taken by the other processes. Figure 13.5 shows the filetype created for process 1.

1  
2  
3  
4  
5  
6  
7  
8  
9  
10  
11  
12  
13  
14  
15  
16  
17  
18  
19  
20  
21  
22  
23  
24  
25  
26  
27  
28  
29  
30  
31  
32  
33  
34  
35  
36  
37  
38  
39  
40  
41  
42  
43  
44  
45  
46  
47  
48

# Annex A

## Language Bindings Summary

In this section we summarize the specific bindings for C, Fortran, and C++. First we present the constants, type definitions, info values and keys. Then we present the routine prototypes separately for each binding. Listings are alphabetical within chapter.

### A.1 Defined Values and Handles

#### A.1.1 Defined Constants

The C and Fortran name is listed in the left column and the C++ name is listed in the middle or right column. Constants with the type `const int` may also be implemented as literal integer constants substituted by the preprocessor.

Return Codes	
C type: <code>const int</code> (or unnamed <code>enum</code> )	C++ type: <code>const int</code>
Fortran type: <code>INTEGER</code>	(or unnamed <code>enum</code> )
<code>MPI_SUCCESS</code>	<code>MPI::SUCCESS</code>
<code>MPI_ERR_BUFFER</code>	<code>MPI::ERR_BUFFER</code>
<code>MPI_ERR_COUNT</code>	<code>MPI::ERR_COUNT</code>
<code>MPI_ERR_TYPE</code>	<code>MPI::ERR_TYPE</code>
<code>MPI_ERR_TAG</code>	<code>MPI::ERR_TAG</code>
<code>MPI_ERR_COMM</code>	<code>MPI::ERR_COMM</code>
<code>MPI_ERR_RANK</code>	<code>MPI::ERR_RANK</code>
<code>MPI_ERR_REQUEST</code>	<code>MPI::ERR_REQUEST</code>
<code>MPI_ERR_ROOT</code>	<code>MPI::ERR_ROOT</code>
<code>MPI_ERR_GROUP</code>	<code>MPI::ERR_GROUP</code>
<code>MPI_ERR_OP</code>	<code>MPI::ERR_OP</code>
<code>MPI_ERR_TOPOLOGY</code>	<code>MPI::ERR_TOPOLOGY</code>
<code>MPI_ERR_DIMS</code>	<code>MPI::ERR_DIMS</code>
<code>MPI_ERR_ARG</code>	<code>MPI::ERR_ARG</code>
<code>MPI_ERR_UNKNOWN</code>	<code>MPI::ERR_UNKNOWN</code>
<code>MPI_ERR_TRUNCATE</code>	<code>MPI::ERR_TRUNCATE</code>
<code>MPI_ERR_OTHER</code>	<code>MPI::ERR_OTHER</code>
<code>MPI_ERR_INTERN</code>	<code>MPI::ERR_INTERN</code>
<code>MPI_ERR_PENDING</code>	<code>MPI::ERR_PENDING</code>

(Continued on next page)

**Return Codes (continued)**

MPI_ERR_IN_STATUS	MPI::ERR_IN_STATUS
MPI_ERR_ACCESS	MPI::ERR_ACCESS
MPI_ERR_AMODE	MPI::ERR_AMODE
MPI_ERR_ASSERT	MPI::ERR_ASSERT
MPI_ERR_BAD_FILE	MPI::ERR_BAD_FILE
MPI_ERR_BASE	MPI::ERR_BASE
MPI_ERR_CONVERSION	MPI::ERR_CONVERSION
MPI_ERR_DISP	MPI::ERR_DISP
MPI_ERR_DUP_DATAREP	MPI::ERR_DUP_DATAREP
MPI_ERR_FILE_EXISTS	MPI::ERR_FILE_EXISTS
MPI_ERR_FILE_IN_USE	MPI::ERR_FILE_IN_USE
MPI_ERR_FILE	MPI::ERR_FILE
MPI_ERR_INFO_KEY	MPI::ERR_INFO_VALUE
MPI_ERR_INFO_NOKEY	MPI::ERR_INFO_NOKEY
MPI_ERR_INFO_VALUE	MPI::ERR_INFO_KEY
MPI_ERR_INFO	MPI::ERR_INFO
MPI_ERR_IO	MPI::ERR_IO
MPI_ERR_KEYVAL	MPI::ERR_KEYVAL
MPI_ERR_LOCKTYPE	MPI::ERR_LOCKTYPE
MPI_ERR_NAME	MPI::ERR_NAME
MPI_ERR_NO_MEM	MPI::ERR_NO_MEM
MPI_ERR_NOT_SAME	MPI::ERR_NOT_SAME
MPI_ERR_NO_SPACE	MPI::ERR_NO_SPACE
MPI_ERR_NO_SUCH_FILE	MPI::ERR_NO_SUCH_FILE
MPI_ERR_PORT	MPI::ERR_PORT
MPI_ERR_QUOTA	MPI::ERR_QUOTA
MPI_ERR_READ_ONLY	MPI::ERR_READ_ONLY
MPI_ERR_RMA_CONFLICT	MPI::ERR_RMA_CONFLICT
MPI_ERR_RMA_SYNC	MPI::ERR_RMA_SYNC
MPI_ERR_SERVICE	MPI::ERR_SERVICE
MPI_ERR_SIZE	MPI::ERR_SIZE
MPI_ERR_SPAWN	MPI::ERR_SPAWN
MPI_ERR_UNSUPPORTED_DATAREP	MPI::ERR_UNSUPPORTED_DATAREP
MPI_ERR_UNSUPPORTED_OPERATION	MPI::ERR_UNSUPPORTED_OPERATION
MPI_ERR_WIN	MPI::ERR_WIN
MPI_ERR_LASTCODE	MPI::ERR_LASTCODE

**Buffer Address Constants**

C type: void * const	C++ type:
Fortran type: (predefined memory location)	void * const
MPI_BOTTOM	MPI::BOTTOM
MPI_IN_PLACE	MPI::IN_PLACE

**Assorted Constants**

C type: <code>const int</code> (or unnamed <code>enum</code> )	C++ type:
Fortran type: <code>INTEGER</code>	<code>const int</code> (or unnamed <code>enum</code> )
<code>MPI_PROC_NULL</code>	<code>MPI::PROC_NULL</code>
<code>MPI_ANY_SOURCE</code>	<code>MPI::ANY_SOURCE</code>
<code>MPI_ANY_TAG</code>	<code>MPI::ANY_TAG</code>
<code>MPI_UNDEFINED</code>	<code>MPI::UNDEFINED</code>
<code>MPI_BSEND_OVERHEAD</code>	<code>MPI::BSEND_OVERHEAD</code>
<code>MPI_KEYVAL_INVALID</code>	<code>MPI::KEYVAL_INVALID</code>
<code>MPI_LOCK_EXCLUSIVE</code>	<code>MPI::LOCK_EXCLUSIVE</code>
<code>MPI_LOCK_SHARED</code>	<code>MPI::LOCK_SHARED</code>
<code>MPI_ROOT</code>	<code>MPI::ROOT</code>

**Status size and reserved index values (Fortran only)**

Fortran type: <code>INTEGER</code>	
<code>MPI_STATUS_SIZE</code>	Not defined for C++
<code>MPI_SOURCE</code>	Not defined for C++
<code>MPI_TAG</code>	Not defined for C++
<code>MPI_ERROR</code>	Not defined for C++

**Variable Address Size (Fortran only)**

Fortran type: <code>INTEGER</code>	
<code>MPI_ADDRESS_KIND</code>	Not defined for C++
[ticket265.] <code>MPI_COUNT_KIND</code>	Not defined for C++
<code>MPI_INTEGER_KIND</code>	Not defined for C++
<code>MPI_OFFSET_KIND</code>	Not defined for C++

**Error-handling specifiers**

C type: <code>MPI_Errhandler</code>	C++ type: <code>MPI::Errhandler</code>
Fortran type: <code>INTEGER</code>	
<code>MPI_ERRORS_ARE_FATAL</code>	<code>MPI::ERRORS_ARE_FATAL</code>
<code>MPI_ERRORS_RETURN</code>	<code>MPI::ERRORS_RETURN</code>
	<code>MPI::ERRORS_THROW_EXCEPTIONS</code>

**Maximum Sizes for Strings**

C type: <code>const int</code> (or unnamed <code>enum</code> )	C++ type:
Fortran type: <code>INTEGER</code>	<code>const int</code> (or unnamed <code>enum</code> )
<code>MPI_MAX_PROCESSOR_NAME</code>	<code>MPI::MAX_PROCESSOR_NAME</code>
<code>MPI_MAX_ERROR_STRING</code>	<code>MPI::MAX_ERROR_STRING</code>
<code>MPI_MAX_DATAREP_STRING</code>	<code>MPI::MAX_DATAREP_STRING</code>
<code>MPI_MAX_INFO_KEY</code>	<code>MPI::MAX_INFO_KEY</code>
<code>MPI_MAX_INFO_VAL</code>	<code>MPI::MAX_INFO_VAL</code>
<code>MPI_MAX_OBJECT_NAME</code>	<code>MPI::MAX_OBJECT_NAME</code>
<code>MPI_MAX_PORT_NAME</code>	<code>MPI::MAX_PORT_NAME</code>

Named Predefined Datatypes		C/C++ types
C type: MPI_Datatype	C++ type: MPI::Datatype	
Fortran type: INTEGER		
MPI_CHAR	MPI::CHAR	char (treated as printable character)
MPI_SHORT	MPI::SHORT	signed short int
MPI_INT	MPI::INT	signed int
MPI_LONG	MPI::LONG	signed long
MPI_LONG_LONG_INT	MPI::LONG_LONG_INT	signed long long
MPI_LONG_LONG	MPI::LONG_LONG	long long (synonym)
MPI_SIGNED_CHAR	MPI::SIGNED_CHAR	signed char (treated as integral value)
MPI_UNSIGNED_CHAR	MPI::UNSIGNED_CHAR	unsigned char (treated as integral value)
MPI_UNSIGNED_SHORT	MPI::UNSIGNED_SHORT	unsigned short
MPI_UNSIGNED	MPI::UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	MPI::UNSIGNED_LONG	unsigned long
MPI_UNSIGNED_LONG_LONG	MPI::UNSIGNED_LONG_LONG	unsigned long long
MPI_FLOAT	MPI::FLOAT	float
MPI_DOUBLE	MPI::DOUBLE	double
MPI_LONG_DOUBLE	MPI::LONG_DOUBLE	long double
MPI_WCHAR	MPI::WCHAR	wchar_t (defined in <stddef.h>) (treated as printable character)
MPI_C_BOOL	(use C datatype handle)	_Bool
MPI_INT8_T	(use C datatype handle)	int8_t
MPI_INT16_T	(use C datatype handle)	int16_t
MPI_INT32_T	(use C datatype handle)	int32_t
MPI_INT64_T	(use C datatype handle)	int64_t
MPI_UINT8_T	(use C datatype handle)	uint8_t
MPI_UINT16_T	(use C datatype handle)	uint16_t
MPI_UINT32_T	(use C datatype handle)	uint32_t
MPI_UINT64_T	(use C datatype handle)	uint64_t
MPI_AINT	(use C datatype handle)	MPI_Aint
[ticket265.] MPI_COUNT	(use C datatype handle)	MPI_Count
MPI_OFFSET	(use C datatype handle)	MPI_Offset
MPI_C_COMPLEX	(use C datatype handle)	float _Complex
MPI_C_FLOAT_COMPLEX	(use C datatype handle)	float _Complex
MPI_C_DOUBLE_COMPLEX	(use C datatype handle)	double _Complex
MPI_C_LONG_DOUBLE_COMPLEX	(use C datatype handle)	long double _Complex
MPI_BYTE	MPI::BYTE	(any C/C++ type)
MPI_PACKED	MPI::PACKED	(any C/C++ type)

Named Predefined Datatypes		Fortran types
C type: MPI_Datatype	C++ type: MPI::Datatype	
Fortran type: INTEGER		
MPI_INTEGER	MPI::INTEGER	INTEGER
MPI_REAL	MPI::REAL	REAL
MPI_DOUBLE_PRECISION	MPI::DOUBLE_PRECISION	DOUBLE PRECISION
MPI_COMPLEX	MPI::F_COMPLEX	COMPLEX
MPI_LOGICAL	MPI::LOGICAL	LOGICAL
MPI_CHARACTER	MPI::CHARACTER	CHARACTER(1)
MPI_AINT	(use C datatype handle)	INTEGER (KIND=MPI_ADDRESS_KIND)
MPI_OFFSET	(use C datatype handle)	INTEGER (KIND=MPI_OFFSET_KIND)
[ticket265.] MPI_COUNT	(use C datatype handle)	INTEGER (KIND=MPI_COUNT_KIND)
MPI_BYTE	MPI::BYTE	(any Fortran type)
MPI_PACKED	MPI::PACKED	(any Fortran type)

C++-Only Named Predefined Datatypes	C++ types
C++ type: MPI::Datatype	
MPI::BOOL	bool
MPI::COMPLEX	Complex<float>
MPI::DOUBLE_COMPLEX	Complex<double>
MPI::LONG_DOUBLE_COMPLEX	Complex<long double>

Optional datatypes (Fortran)		Fortran types
C type: MPI_Datatype	C++ type: MPI::Datatype	
Fortran type: INTEGER		
MPI_DOUBLE_COMPLEX	MPI::F_DOUBLE_COMPLEX	DOUBLE COMPLEX
MPI_INTEGER1	MPI::INTEGER1	INTEGER*1
MPI_INTEGER2	MPI::INTEGER2	INTEGER*8
MPI_INTEGER4	MPI::INTEGER4	INTEGER*4
MPI_INTEGER8	MPI::INTEGER8	INTEGER*8
MPI_INTEGER16		INTEGER*16
MPI_REAL2	MPI::REAL2	REAL*2
MPI_REAL4	MPI::REAL4	REAL*4
MPI_REAL8	MPI::REAL8	REAL*8
MPI_REAL16		REAL*16
MPI_COMPLEX4		COMPLEX*4
MPI_COMPLEX8		COMPLEX*8
MPI_COMPLEX16		COMPLEX*16
MPI_COMPLEX32		COMPLEX*32



**Datatypes for reduction functions (C and C++)**

C type: MPI_Datatype	C++ type: MPI::Datatype
Fortran type: INTEGER	

MPI_FLOAT_INT	MPI::FLOAT_INT
MPI_DOUBLE_INT	MPI::DOUBLE_INT
MPI_LONG_INT	MPI::LONG_INT
MPI_2INT	MPI::TWOINT
MPI_SHORT_INT	MPI::SHORT_INT
MPI_LONG_DOUBLE_INT	MPI::LONG_DOUBLE_INT

**Datatypes for reduction functions (Fortran)**

C type: MPI_Datatype	C++ type: MPI::Datatype
Fortran type: INTEGER	

MPI_2REAL	MPI::TWOREAL
MPI_2DOUBLE_PRECISION	MPI::TWODOUBLE_PRECISION
MPI_2INTEGER	MPI::TWOINTEGER

**Special datatypes for constructing derived datatypes**

C type: MPI_Datatype	C++ type: MPI::Datatype
Fortran type: INTEGER	

MPI_UB	MPI::UB
MPI_LB	MPI::LB

**Reserved communicators**

C type: MPI_Comm	C++ type: MPI::Intracomm
Fortran type: INTEGER	

MPI_COMM_WORLD	MPI::COMM_WORLD
MPI_COMM_SELF	MPI::COMM_SELF

**Results of communicator and group comparisons**

C type: const int (or unnamed enum)	C++ type: const int (or unnamed enum)
Fortran type: INTEGER	

MPI_IDENT	MPI::IDENT
MPI_CONGRUENT	MPI::CONGRUENT
MPI_SIMILAR	MPI::SIMILAR
MPI_UNEQUAL	MPI::UNEQUAL

**Environmental inquiry keys**

C type: const int (or unnamed enum)	C++ type: const int (or unnamed enum)
Fortran type: INTEGER	

MPI_TAG_UB	MPI::TAG_UB
MPI_IO	MPI::IO
MPI_HOST	MPI::HOST
MPI_WTIME_IS_GLOBAL	MPI::WTIME_IS_GLOBAL

**Collective Operations**

C type: MPI_Op	C++ type: const MPI::Op
Fortran type: INTEGER	
MPI_MAX	MPI::MAX
MPI_MIN	MPI::MIN
MPI_SUM	MPI::SUM
MPI_PROD	MPI::PROD
MPI_MAXLOC	MPI::MAXLOC
MPI_MINLOC	MPI::MINLOC
MPI_BAND	MPI::BAND
MPI_BOR	MPI::BOR
MPI_BXOR	MPI::BXOR
MPI_LAND	MPI::LAND
MPI_LOR	MPI::LOR
MPI_LXOR	MPI::LXOR
MPI_REPLACE	MPI::REPLACE

**Null Handles**

C/Fortran name	C++ name
C type / Fortran type	C++ type
MPI_GROUP_NULL	MPI::GROUP_NULL
MPI_Group / INTEGER	const MPI::Group
MPI_COMM_NULL	MPI::COMM_NULL
MPI_Comm / INTEGER	<sup>1)</sup>
MPI_DATATYPE_NULL	MPI::DATATYPE_NULL
MPI_Datatype / INTEGER	const MPI::Datatype
MPI_REQUEST_NULL	MPI::REQUEST_NULL
MPI_Request / INTEGER	const MPI::Request
MPI_OP_NULL	MPI::OP_NULL
MPI_Op / INTEGER	const MPI::Op
MPI_ERRHANDLER_NULL	MPI::ERRHANDLER_NULL
MPI_Errhandler / INTEGER	const MPI::Errhandler
MPI_FILE_NULL	MPI::FILE_NULL
MPI_File / INTEGER	
MPI_INFO_NULL	MPI::INFO_NULL
MPI_Info / INTEGER	const MPI::Info
MPI_WIN_NULL	MPI::WIN_NULL
MPI_Win / INTEGER	

<sup>1)</sup> C++ type: See Section 16.1.7 on page 506 regarding class hierarchy and the specific type of MPI::COMM\_NULL

**Empty group**

C type: MPI_Group	C++ type: const MPI::Group
Fortran type: INTEGER	
MPI_GROUP_EMPTY	MPI::GROUP_EMPTY

**Topologies**

C type: <code>const int</code> (or unnamed <code>enum</code> )	C++ type: <code>const int</code>
Fortran type: <code>INTEGER</code>	(or unnamed <code>enum</code> )
<code>MPI_GRAPH</code>	<code>MPI::GRAPH</code>
<code>MPI_CART</code>	<code>MPI::CART</code>
<code>MPI_DIST_GRAPH</code>	<code>MPI::DIST_GRAPH</code>

**Predefined functions**

C/Fortran name	C++ name
C type / Fortran type	C++ type
<code>MPI_COMM_NULL_COPY_FN</code>	<code>MPI_COMM_NULL_COPY_FN</code>
<code>MPI_Comm_copy_attr_function</code>	same as in C <sup>1</sup> )
<code>/ COMM_COPY_ATTR_FN</code>	
<code>MPI_COMM_DUP_FN</code>	<code>MPI_COMM_DUP_FN</code>
<code>MPI_Comm_copy_attr_function</code>	same as in C <sup>1</sup> )
<code>/ COMM_COPY_ATTR_FN</code>	
<code>MPI_COMM_NULL_DELETE_FN</code>	<code>MPI_COMM_NULL_DELETE_FN</code>
<code>MPI_Comm_delete_attr_function</code>	same as in C <sup>1</sup> )
<code>/ COMM_DELETE_ATTR_FN</code>	
<code>MPI_WIN_NULL_COPY_FN</code>	<code>MPI_WIN_NULL_COPY_FN</code>
<code>MPI_Win_copy_attr_function</code>	same as in C <sup>1</sup> )
<code>/ WIN_COPY_ATTR_FN</code>	
<code>MPI_WIN_DUP_FN</code>	<code>MPI_WIN_DUP_FN</code>
<code>MPI_Win_copy_attr_function</code>	same as in C <sup>1</sup> )
<code>/ WIN_COPY_ATTR_FN</code>	
<code>MPI_WIN_NULL_DELETE_FN</code>	<code>MPI_WIN_NULL_DELETE_FN</code>
<code>MPI_Win_delete_attr_function</code>	same as in C <sup>1</sup> )
<code>/ WIN_DELETE_ATTR_FN</code>	
<code>MPI_TYPE_NULL_COPY_FN</code>	<code>MPI_TYPE_NULL_COPY_FN</code>
<code>MPI_Type_copy_attr_function</code>	same as in C <sup>1</sup> )
<code>/ TYPE_COPY_ATTR_FN</code>	
<code>MPI_TYPE_DUP_FN</code>	<code>MPI_TYPE_DUP_FN</code>
<code>MPI_Type_copy_attr_function</code>	same as in C <sup>1</sup> )
<code>/ TYPE_COPY_ATTR_FN</code>	
<code>MPI_TYPE_NULL_DELETE_FN</code>	<code>MPI_TYPE_NULL_DELETE_FN</code>
<code>MPI_Type_delete_attr_function</code>	same as in C <sup>1</sup> )
<code>/ TYPE_DELETE_ATTR_FN</code>	

<sup>1</sup> See the advice to implementors on `MPI_COMM_NULL_COPY_FN`, ... in Section 6.7.2 on page 257

**Deprecated predefined functions**

C/Fortran name	C++ name
C type / Fortran type	C++ type
MPI_NULL_COPY_FN	MPI::NULL_COPY_FN
MPI_Copy_function / COPY_FUNCTION	MPI::Copy_function
MPI_DUP_FN	MPI::DUP_FN
MPI_Copy_function / COPY_FUNCTION	MPI::Copy_function
MPI_NULL_DELETE_FN	MPI::NULL_DELETE_FN
MPI_Delete_function / DELETE_FUNCTION	MPI::Delete_function

**Predefined Attribute Keys**

C type: <code>const int</code> (or unnamed <code>enum</code> )	C++ type:
Fortran type: <code>INTEGER</code>	<code>const int</code> (or unnamed <code>enum</code> )
MPI_APPNUM	MPI::APPNUM
MPI_LASTUSED_CODE	MPI::LASTUSED_CODE
MPI_UNIVERSE_SIZE	MPI::UNIVERSE_SIZE
MPI_WIN_BASE	MPI::WIN_BASE
MPI_WIN_DISP_UNIT	MPI::WIN_DISP_UNIT
MPI_WIN_SIZE	MPI::WIN_SIZE

**Mode Constants**

C type: <code>const int</code> (or unnamed <code>enum</code> )	C++ type:
Fortran type: <code>INTEGER</code>	<code>const int</code> (or unnamed <code>enum</code> )
MPI_MODE_APPEND	MPI::MODE_APPEND
MPI_MODE_CREATE	MPI::MODE_CREATE
MPI_MODE_DELETE_ON_CLOSE	MPI::MODE_DELETE_ON_CLOSE
MPI_MODE_EXCL	MPI::MODE_EXCL
MPI_MODE_NOCHECK	MPI::MODE_NOCHECK
MPI_MODE_NOPRECEDE	MPI::MODE_NOPRECEDE
MPI_MODE_NOPUT	MPI::MODE_NOPUT
MPI_MODE_NOSTORE	MPI::MODE_NOSTORE
MPI_MODE_NOSUCCEED	MPI::MODE_NOSUCCEED
MPI_MODE_RDONLY	MPI::MODE_RDONLY
MPI_MODE_RDWR	MPI::MODE_RDWR
MPI_MODE_SEQUENTIAL	MPI::MODE_SEQUENTIAL
MPI_MODE_UNIQUE_OPEN	MPI::MODE_UNIQUE_OPEN
MPI_MODE_WRONLY	MPI::MODE_WRONLY

**Datatype Decoding Constants**

C type: <code>const int</code> (or unnamed <code>enum</code> )	C++ type:
Fortran type: <code>INTEGER</code>	<code>const int</code> (or unnamed <code>enum</code> )
<code>MPI_COMBINER_CONTIGUOUS</code>	<code>MPI::COMBINER_CONTIGUOUS</code>
<code>MPI_COMBINER_DARRAY</code>	<code>MPI::COMBINER_DARRAY</code>
<code>MPI_COMBINER_DUP</code>	<code>MPI::COMBINER_DUP</code>
<code>MPI_COMBINER_F90_COMPLEX</code>	<code>MPI::COMBINER_F90_COMPLEX</code>
<code>MPI_COMBINER_F90_INTEGER</code>	<code>MPI::COMBINER_F90_INTEGER</code>
<code>MPI_COMBINER_F90_REAL</code>	<code>MPI::COMBINER_F90_REAL</code>
<code>MPI_COMBINER_HINDEXED_INTEGER</code>	<code>MPI::COMBINER_HINDEXED_INTEGER</code>
<code>MPI_COMBINER_HINDEXED</code>	<code>MPI::COMBINER_HINDEXED</code>
<code>MPI_COMBINER_HVECTOR_INTEGER</code>	<code>MPI::COMBINER_HVECTOR_INTEGER</code>
<code>MPI_COMBINER_HVECTOR</code>	<code>MPI::COMBINER_HVECTOR</code>
<code>MPI_COMBINER_INDEXED_BLOCK</code>	<code>MPI::COMBINER_INDEXED_BLOCK</code>
<code>MPI_COMBINER_INDEXED</code>	<code>MPI::COMBINER_INDEXED</code>
<code>MPI_COMBINER_NAMED</code>	<code>MPI::COMBINER_NAMED</code>
<code>MPI_COMBINER_RESIZED</code>	<code>MPI::COMBINER_RESIZED</code>
<code>MPI_COMBINER_STRUCT_INTEGER</code>	<code>MPI::COMBINER_STRUCT_INTEGER</code>
<code>MPI_COMBINER_STRUCT</code>	<code>MPI::COMBINER_STRUCT</code>
<code>MPI_COMBINER_SUBARRAY</code>	<code>MPI::COMBINER_SUBARRAY</code>
<code>MPI_COMBINER_VECTOR</code>	<code>MPI::COMBINER_VECTOR</code>

**Threads Constants**

C type: <code>const int</code> (or unnamed <code>enum</code> )	C++ type:
Fortran type: <code>INTEGER</code>	<code>const int</code> (or unnamed <code>enum</code> )
<code>MPI_THREAD_FUNNELED</code>	<code>MPI::THREAD_FUNNELED</code>
<code>MPI_THREAD_MULTIPLE</code>	<code>MPI::THREAD_MULTIPLE</code>
<code>MPI_THREAD_SERIALIZED</code>	<code>MPI::THREAD_SERIALIZED</code>
<code>MPI_THREAD_SINGLE</code>	<code>MPI::THREAD_SINGLE</code>

**File Operation Constants, Part 1**

C type: <code>const MPI_Offset</code> (or unnamed <code>enum</code> )	C++ type:
Fortran type: <code>INTEGER (KIND=MPI_OFFSET_KIND)</code>	<code>const MPI::Offset</code> (or unnamed <code>enum</code> )
<code>MPI_DISPLACEMENT_CURRENT</code>	<code>MPI::DISPLACEMENT_CURRENT</code>

**File Operation Constants, Part 2**

C type: <code>const int</code> (or unnamed <code>enum</code> )	C++ type:
Fortran type: <code>INTEGER</code>	<code>const int</code> (or unnamed <code>enum</code> )
<code>MPI_DISTRIBUTE_BLOCK</code>	<code>MPI::DISTRIBUTE_BLOCK</code>
<code>MPI_DISTRIBUTE_CYCLIC</code>	<code>MPI::DISTRIBUTE_CYCLIC</code>
<code>MPI_DISTRIBUTE_DFLT_DARG</code>	<code>MPI::DISTRIBUTE_DFLT_DARG</code>
<code>MPI_DISTRIBUTE_NONE</code>	<code>MPI::DISTRIBUTE_NONE</code>
<code>MPI_ORDER_C</code>	<code>MPI::ORDER_C</code>
<code>MPI_ORDER_FORTRAN</code>	<code>MPI::ORDER_FORTRAN</code>
<code>MPI_SEEK_CUR</code>	<code>MPI::SEEK_CUR</code>
<code>MPI_SEEK_END</code>	<code>MPI::SEEK_END</code>
<code>MPI_SEEK_SET</code>	<code>MPI::SEEK_SET</code>

**F90 Datatype Matching Constants**

C type: <code>const int</code> (or unnamed <code>enum</code> )	C++ type:
Fortran type: <code>INTEGER</code>	<code>const int</code> (or unnamed <code>enum</code> )
<code>MPI_TYPECLASS_COMPLEX</code>	<code>MPI::TYPECLASS_COMPLEX</code>
<code>MPI_TYPECLASS_INTEGER</code>	<code>MPI::TYPECLASS_INTEGER</code>
<code>MPI_TYPECLASS_REAL</code>	<code>MPI::TYPECLASS_REAL</code>

**Constants Specifying Empty or Ignored Input**

C/Fortran name	C++ name
C type / Fortran type	C++ type
<code>MPI_ARGVS_NULL</code>	<code>MPI::ARGVS_NULL</code>
<code>char***</code> / 2-dim. array of <code>CHARACTER*(*)</code>	<code>const char ***</code>
<code>MPI_ARGV_NULL</code>	<code>MPI::ARGV_NULL</code>
<code>char**</code> / array of <code>CHARACTER*(*)</code>	<code>const char **</code>
<code>MPI_ERRCODES_IGNORE</code>	Not defined for C++
<code>int*</code> / <code>INTEGER</code> array	
<code>MPI_STATUSES_IGNORE</code>	Not defined for C++
<code>MPI_Status*</code> / <code>INTEGER</code> , <code>DIMENSION(MPI_STATUS_SIZE,*)</code>	
<code>MPI_STATUS_IGNORE</code>	Not defined for C++
<code>MPI_Status*</code> / <code>INTEGER</code> , <code>DIMENSION(MPI_STATUS_SIZE)</code>	
<code>MPI_UNWEIGHTED</code>	Not defined for C++

**C Constants Specifying Ignored Input (no C++ or Fortran)**

C type: <code>MPI_Fint*</code>
<code>MPI_F_STATUSES_IGNORE</code>
<code>MPI_F_STATUS_IGNORE</code>

**C and C++ preprocessor Constants and Fortran Parameters**

C/C++ type: <code>const int</code> (or unnamed <code>enum</code> )
Fortran type: <code>INTEGER</code>
<code>MPI_SUBVERSION</code>
<code>MPI_VERSION</code>

## A.1.2 Types

The following are defined C type definitions, included in the file `mpi.h`.

```
/* C opaque types */
```

```
MPI_Aint
```

```
MPI_Count
```

```
MPI_Fint
```

```
MPI_Offset
```

```
MPI_Status
```

```
/* C handles to assorted structures */
```

```
MPI_Comm
```

```
MPI_Datatype
```

```
MPI_Errhandler
```

```
MPI_File
```

```
MPI_Group
```

```
MPI_Info
```

```
MPI_Op
```

```
MPI_Request
```

```
MPI_Win
```

```
// C++ opaque types (all within the MPI namespace)
```

```
MPI::Aint
```

```
MPI::Offset
```

```
MPI::Status
```

```
// C++ handles to assorted structures (classes,
```

```
// all within the MPI namespace)
```

```
MPI::Comm
```

```
MPI::Intracomm
```

```
MPI::Graphcomm
```

```
MPI::Distgraphcomm
```

```
MPI::Cartcomm
```

```
MPI::Intercomm
```

```
MPI::Datatype
```

```
MPI::Errhandler
```

```
MPI::Exception
```

```
MPI::File
```

```
MPI::Group
```

```
MPI::Info
```

```
MPI::Op
```

```
MPI::Request
```

```
MPI::Prequest
```

```
MPI::Grequest
```

```
MPI::Win
```

## A.1.3 Prototype [d]Definitions

The following are defined C typedefs for user-defined functions, also included in the file `mpi.h`.

```

/* prototypes for user-defined functions */
typedef void MPI_User_function(void *invec, void *inoutvec, int *len,
                               MPI_Datatype *datatype);

typedef int MPI_Comm_copy_attr_function(MPI_Comm oldcomm,
                                         int comm_keyval, void *extra_state, void *attribute_val_in,
                                         void *attribute_val_out, int*flag);
typedef int MPI_Comm_delete_attr_function(MPI_Comm comm,
                                         int comm_keyval, void *attribute_val, void *extra_state);

typedef int MPI_Win_copy_attr_function(MPI_Win oldwin, int win_keyval,
                                         void *extra_state, void *attribute_val_in,
                                         void *attribute_val_out, int *flag);
typedef int MPI_Win_delete_attr_function(MPI_Win win, int win_keyval,
                                         void *attribute_val, void *extra_state);

typedef int MPI_Type_copy_attr_function(MPI_Datatype oldtype,
                                         int type_keyval, void *extra_state,
                                         void *attribute_val_in, void *attribute_val_out, int *flag);
typedef int MPI_Type_delete_attr_function(MPI_Datatype type,
                                         int type_keyval, void *attribute_val, void *extra_state);

typedef void MPI_Comm_errhandler_function(MPI_Comm *, int *, ...);
typedef void MPI_Win_errhandler_function(MPI_Win *, int *, ...);
typedef void MPI_File_errhandler_function(MPI_File *, int *, ...);

typedef int MPI_Grequest_query_function(void *extra_state,
                                         MPI_Status *status);
typedef int MPI_Grequest_free_function(void *extra_state);
typedef int MPI_Grequest_cancel_function(void *extra_state, int complete);

typedef int MPI_Datarep_extent_function(MPI_Datatype datatype,
                                         MPI_Aint *file_extent, void *extra_state);
typedef int MPI_Datarep_conversion_function(void *userbuf,
                                         MPI_Datatype datatype, int count, void *filebuf,
                                         MPI_Offset position, void *extra_state);

```

For Fortran, here are examples of how each of the user-defined subroutines should be declared.

The user-function argument to `MPI_OP_CREATE` should be declared like this:

```

SUBROUTINE USER_FUNCTION(INVEC, INOUTVEC, LEN, TYPE)
  <type> INVEC(LEN), INOUTVEC(LEN)
  INTEGER LEN, TYPE

```



The copy and delete function arguments to MPI\_COMM\_CREATE\_KEYVAL should be declared like these:

```

SUBROUTINE COMM_COPY_ATTR_FN(OLDCOMM, COMM_KEYVAL, EXTRA_STATE,
    ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERROR)
    INTEGER OLDCOMM, COMM_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,
        ATTRIBUTE_VAL_OUT
    LOGICAL FLAG

SUBROUTINE COMM_DELETE_ATTR_FN(COMM, COMM_KEYVAL, ATTRIBUTE_VAL,
    EXTRA_STATE, IERROR)
    INTEGER COMM, COMM_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE

```

The copy and delete function arguments to MPI\_WIN\_CREATE\_KEYVAL should be declared like these:

```

SUBROUTINE WIN_COPY_ATTR_FN(OLDWIN, WIN_KEYVAL, EXTRA_STATE,
    ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERROR)
    INTEGER OLDWIN, WIN_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,
        ATTRIBUTE_VAL_OUT
    LOGICAL FLAG

SUBROUTINE WIN_DELETE_ATTR_FN(WIN, WIN_KEYVAL, ATTRIBUTE_VAL,
    EXTRA_STATE, IERROR)
    INTEGER WIN, WIN_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE

```

The copy and delete function arguments to MPI\_TYPE\_CREATE\_KEYVAL should be declared like these:

```

SUBROUTINE TYPE_COPY_ATTR_FN(OLDTYPE, TYPE_KEYVAL, EXTRA_STATE,
    ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERROR)
    INTEGER OLDTYPE, TYPE_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE,
        ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT
    LOGICAL FLAG

SUBROUTINE TYPE_DELETE_ATTR_FN(TYPE, TYPE_KEYVAL, ATTRIBUTE_VAL,
    EXTRA_STATE, IERROR)
    INTEGER TYPE, TYPE_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE

```

The handler-function argument to MPI\_COMM\_CREATE\_ERRHANDLER should be declared like this:

```

SUBROUTINE COMM_ERRHANDLER_FUNCTION(COMM, ERROR_CODE)
    INTEGER COMM, ERROR_CODE

```

The handler-function argument to MPI\_WIN\_CREATE\_ERRHANDLER should be declared like this:

```
SUBROUTINE WIN_ERRHANDLER_FUNCTION(WIN, ERROR_CODE)
    INTEGER WIN, ERROR_CODE
```

The handler-function argument to MPI\_FILE\_CREATE\_ERRHANDLER should be declared like this:

```
SUBROUTINE FILE_ERRHANDLER_FUNCTION(FILE, ERROR_CODE)
    INTEGER FILE, ERROR_CODE
```

The query, free, and cancel function arguments to MPI\_GREQUEST\_START should be declared like these:

```
SUBROUTINE GREQUEST_QUERY_FUNCTION(EXTRA_STATE, STATUS, IERROR)
    INTEGER STATUS(MPI_STATUS_SIZE), IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
```

```
SUBROUTINE GREQUEST_FREE_FUNCTION(EXTRA_STATE, IERROR)
    INTEGER IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
```

```
SUBROUTINE GREQUEST_CANCEL_FUNCTION(EXTRA_STATE, COMPLETE, IERROR)
    INTEGER IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
    LOGICAL COMPLETE
```

The extend and conversion function arguments to MPI\_REGISTER\_DATAREP should be declared like these:

```
SUBROUTINE DATAREP_EXTENT_FUNCTION(DATATYPE, EXTENT, EXTRA_STATE, IERROR)
    INTEGER DATATYPE, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTENT, EXTRA_STATE
```

```
SUBROUTINE DATAREP_CONVERSION_FUNCTION(USERBUF, DATATYPE, COUNT, FILEBUF,
    POSITION, EXTRA_STATE, IERROR)
    <TYPE> USERBUF(*), FILEBUF(*)
    INTEGER COUNT, DATATYPE, IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) POSITION
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
```

The following are defined C++ typedefs, also included in the file `mpi.h`.

```
namespace MPI {
    typedef void User_function(const void* invec, void *inoutvec,
        int len, const Datatype& datatype);

    typedef int Comm::Copy_attr_function(const Comm& oldcomm,
```

```

1         int comm_keyval, void* extra_state, void* attribute_val_in,
2         void* attribute_val_out, bool& flag);
3     typedef int Comm::Delete_attr_function(Comm& comm, int
4         comm_keyval, void* attribute_val, void* extra_state);
5
6     typedef int Win::Copy_attr_function(const Win& oldwin,
7         int win_keyval, void* extra_state, void* attribute_val_in,
8         void* attribute_val_out, bool& flag);
9     typedef int Win::Delete_attr_function(Win& win, int
10        win_keyval, void* attribute_val, void* extra_state);
11
12    typedef int Datatype::Copy_attr_function(const Datatype& oldtype,
13        int type_keyval, void* extra_state,
14        const void* attribute_val_in, void* attribute_val_out,
15        bool& flag);
16    typedef int Datatype::Delete_attr_function(Datatype& type,
17        int type_keyval, void* attribute_val, void* extra_state);
18
19    typedef void Comm::Errhandler_function(Comm &, int *, ...);
20    typedef void Win::Errhandler_function(Win &, int *, ...);
21    typedef void File::Errhandler_function(File &, int *, ...);
22
23    typedef int Grequest::Query_function(void* extra_state, Status& status);
24    typedef int Grequest::Free_function(void* extra_state);
25    typedef int Grequest::Cancel_function(void* extra_state, bool complete);
26
27    typedef void Datarep_extent_function(const Datatype& datatype,
28        Aint& file_extent, void* extra_state);
29    typedef void Datarep_conversion_function(void* userbuf,
30        Datatype& datatype, int count, void* filebuf,
31        Offset position, void* extra_state);
32 }
33

```

#### ticket0. A.1.4 Deprecated [p]Prototype [d]Definitions

ticket0. The following are defined C typedefs for deprecated user-defined functions, also included in the file `mpi.h`.

```

38 /* prototypes for user-defined functions */
39 typedef int MPI_Copy_function(MPI_Comm oldcomm, int keyval,
40     void *extra_state, void *attribute_val_in,
41     void *attribute_val_out, int *flag);
42 typedef int MPI_Delete_function(MPI_Comm comm, int keyval,
43     void *attribute_val, void *extra_state);
44 typedef void MPI_Handler_function(MPI_Comm *, int *, ...);
45

```

The following are deprecated Fortran user-defined callback subroutine prototypes. The deprecated copy and delete function arguments to `MPI_KEYVAL_CREATE` should be declared like these:

```

SUBROUTINE COPY_FUNCTION(OLDCOMM, KEYVAL, EXTRA_STATE,
                        ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERR)
    INTEGER OLDCOMM, KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,
        ATTRIBUTE_VAL_OUT, IERR
    LOGICAL FLAG

SUBROUTINE DELETE_FUNCTION(COMM, KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE, IERR)
    INTEGER COMM, KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE, IERR

```

The deprecated handler-function for error handlers should be declared like this:

```

SUBROUTINE HANDLER_FUNCTION(COMM, ERROR_CODE)
    INTEGER COMM, ERROR_CODE

```

### A.1.5 Info Keys

access\_style  
 appnum  
 arch  
 cb\_block\_size  
 cb\_buffer\_size  
 cb\_nodes  
 chunked\_item  
 chunked\_size  
 chunked  
 collective\_buffering  
 file\_perm  
 filename  
 file  
 host  
 io\_node\_list  
 ip\_address  
 ip\_port  
 nb\_proc  
 no\_locks  
 num\_io\_nodes  
 path  
 soft  
 striping\_factor  
 striping\_unit  
 wdir

### A.1.6 Info Values

false  
 random  
 read\_mostly  
 read\_once

reverse\_sequential  
sequential  
true  
write\_mostly  
write\_once

## A.2 C Bindings

### A.2.1 Point-to-Point Communication C Bindings

```

int MPI_Bsend(void* buf, int count, MPI_Datatype datatype, int dest,
              int tag, MPI_Comm comm)

int MPI_Bsend_init(void* buf, int count, MPI_Datatype datatype, int dest,
                  int tag, MPI_Comm comm, MPI_Request *request)

int MPI_Buffer_attach(void* buffer, int size)

int MPI_Buffer_detach(void* buffer_addr, int* size)

int MPI_Cancel(MPI_Request *request)

int MPI_Get_count(MPI_Status *status, MPI_Datatype datatype, int *count)

int MPI_Ibsend(void* buf, int count, MPI_Datatype datatype, int dest,
              int tag, MPI_Comm comm, MPI_Request *request)

int MPI_Iprobe(int source, int tag, MPI_Comm comm, int *flag,
              MPI_Status *status)

int MPI_Irecv(void* buf, int count, MPI_Datatype datatype, int source,
              int tag, MPI_Comm comm, MPI_Request *request)

int MPI_Irsend(void* buf, int count, MPI_Datatype datatype, int dest,
              int tag, MPI_Comm comm, MPI_Request *request)

int MPI_Isend(void* buf, int count, MPI_Datatype datatype, int dest,
              int tag, MPI_Comm comm, MPI_Request *request)

int MPI_Issend(void* buf, int count, MPI_Datatype datatype, int dest,
              int tag, MPI_Comm comm, MPI_Request *request)

int MPI_Probe(int source, int tag, MPI_Comm comm, MPI_Status *status)

int MPI_Recv(void* buf, int count, MPI_Datatype datatype, int source,
             int tag, MPI_Comm comm, MPI_Status *status)

int MPI_Recv_init(void* buf, int count, MPI_Datatype datatype, int source,
                 int tag, MPI_Comm comm, MPI_Request *request)

int MPI_Request_free(MPI_Request *request)

int MPI_Request_get_status(MPI_Request request, int *flag,
                          MPI_Status *status)

int MPI_Rsend(void* buf, int count, MPI_Datatype datatype, int dest,
              int tag, MPI_Comm comm)

int MPI_Rsend_init(void* buf, int count, MPI_Datatype datatype, int dest,
                  int tag, MPI_Comm comm, MPI_Request *request)

int MPI_Send(void* buf, int count, MPI_Datatype datatype, int dest,
             int tag, MPI_Comm comm)

```

```

1  int MPI_Send_init(void* buf, int count, MPI_Datatype datatype, int dest,
2      int tag, MPI_Comm comm, MPI_Request *request)
3
4  int MPI_Sendrecv(void *sendbuf, int sendcount, MPI_Datatype sendtype,
5      int dest, int sendtag, void *recvbuf, int recvcount,
6      MPI_Datatype recvttype, int source, int recvttag, MPI_Comm comm,
7      MPI_Status *status)
8
9  int MPI_Sendrecv_replace(void* buf, int count, MPI_Datatype datatype,
10     int dest, int sendtag, int source, int recvttag, MPI_Comm comm,
11     MPI_Status *status)
12
13 int MPI_Ssend(void* buf, int count, MPI_Datatype datatype, int dest,
14     int tag, MPI_Comm comm)
15
16 int MPI_Ssend_init(void* buf, int count, MPI_Datatype datatype, int dest,
17     int tag, MPI_Comm comm, MPI_Request *request)
18
19 int MPI_Start(MPI_Request *request)
20
21 int MPI_Startall(int count, MPI_Request *array_of_requests)
22
23 int MPI_Test(MPI_Request *request, int *flag, MPI_Status *status)
24
25 int MPI_Test_cancelled(MPI_Status *status, int *flag)
26
27 int MPI_Testall(int count, MPI_Request *array_of_requests, int *flag,
28     MPI_Status *array_of_statuses)
29
30 int MPI_Testany(int count, MPI_Request *array_of_requests, int *index,
31     int *flag, MPI_Status *status)
32
33 int MPI_Testsome(int incount, MPI_Request *array_of_requests,
34     int *outcount, int *array_of_indices,
35     MPI_Status *array_of_statuses)
36
37 int MPI_Wait(MPI_Request *request, MPI_Status *status)
38
39 int MPI_Waitall(int count, MPI_Request *array_of_requests,
40     MPI_Status *array_of_statuses)
41
42 int MPI_Waitany(int count, MPI_Request *array_of_requests, int *index,
43     MPI_Status *status)
44
45 int MPI_Waitsome(int incount, MPI_Request *array_of_requests,
46     int *outcount, int *array_of_indices,
47     MPI_Status *array_of_statuses)
48

```

## A.2.2 Datatypes C Bindings

```

44 int MPI_Get_address(void *location, MPI_Aint *address)
45
46 int MPI_Get_elements(MPI_Status *status, MPI_Datatype datatype, int *count)
47
48

```

```

1 int MPI_Get_elements_x(MPI_Status *status, MPI_Datatype datatype,
2     MPI_Count *count)
3
4 int MPI_Pack(void* inbuf, int incount, MPI_Datatype datatype, void *outbuf,
5     int outsize, int *position, MPI_Comm comm)
6
7 int MPI_Pack_external(char *datarep, void *inbuf, int incount,
8     MPI_Datatype datatype, void *outbuf, MPI_Aint outsize,
9     MPI_Aint *position)
10
11 int MPI_Pack_external_size(char *datarep, int incount,
12     MPI_Datatype datatype, MPI_Aint *size)
13
14 int MPI_Pack_size(int incount, MPI_Datatype datatype, MPI_Comm comm,
15     int *size)
16
17 int MPI_Type_commit(MPI_Datatype *datatype)
18
19 int MPI_Type_contiguous(int count, MPI_Datatype oldtype,
20     MPI_Datatype *newtype)
21
22 int MPI_Type_create_darray(int size, int rank, int ndims,
23     int array_of_gsizes[], int array_of_distribs[], int
24     array_of_dargs[], int array_of_psize, int order,
25     MPI_Datatype oldtype, MPI_Datatype *newtype)
26
27 int MPI_Type_create_hindexed(int count, int array_of_blocklengths[],
28     MPI_Aint array_of_displacements[], MPI_Datatype oldtype,
29     MPI_Datatype *newtype)
30
31 int MPI_Type_create_hvector(int count, int blocklength, MPI_Aint stride,
32     MPI_Datatype oldtype, MPI_Datatype *newtype)
33
34 int MPI_Type_create_indexed_block(int count, int blocklength,
35     int array_of_displacements[], MPI_Datatype oldtype,
36     MPI_Datatype *newtype)
37
38 int MPI_Type_create_resized(MPI_Datatype oldtype, MPI_Aint lb, MPI_Aint
39     extent, MPI_Datatype *newtype)
40
41 int MPI_Type_create_struct(int count, int array_of_blocklengths[],
42     MPI_Aint array_of_displacements[],
43     MPI_Datatype array_of_types[], MPI_Datatype *newtype)
44
45 int MPI_Type_create_subarray(int ndims, int array_of_sizes[],
46     int array_of_subsizes[], int array_of_starts[], int order,
47     MPI_Datatype oldtype, MPI_Datatype *newtype)
48
49 int MPI_Type_dup(MPI_Datatype type, MPI_Datatype *newtype)
50
51 int MPI_Type_free(MPI_Datatype *datatype)
52
53 int MPI_Type_get_contents(MPI_Datatype datatype, int max_integers,
54     int max_addresses, int max_datatypes, int array_of_integers[],
55     MPI_Aint array_of_addresses[],
56     MPI_Datatype array_of_datatypes[])
57
58 int MPI_Type_get_enumerator(int count, MPI_Datatype *newtypes)
59
60 int MPI_Type_get_name(MPI_Datatype datatype, char *name, MPI_Aint *length)
61
62 int MPI_Type_get_size(MPI_Datatype datatype, MPI_Aint *size)
63
64 int MPI_Type_is_contiguous(MPI_Datatype datatype)
65
66 int MPI_Type_is_indexed(MPI_Datatype datatype)
67
68 int MPI_Type_is_integrated(MPI_Datatype datatype)
69
70 int MPI_Type_is_lower(MPI_Datatype datatype, MPI_Datatype *newtype)
71
72 int MPI_Type_is_struct(MPI_Datatype datatype)
73
74 int MPI_Type_is_vector(MPI_Datatype datatype)
75
76 int MPI_Type_match(MPI_Datatype type, MPI_Datatype *newtype)
77
78 int MPI_Type_match_multiple(MPI_Datatype *types, MPI_Datatype *newtypes)
79
80 int MPI_Type_pack_size(MPI_Datatype datatype, MPI_Comm comm,
81     MPI_Aint *size)
82
83 int MPI_Type_pack_size_x(MPI_Datatype datatype, MPI_Comm comm,
84     MPI_Aint *size)
85
86 int MPI_Type_rank(MPI_Datatype datatype, int *rank)
87
88 int MPI_Type_size(MPI_Datatype datatype, MPI_Aint *size)
89
90 int MPI_Type_size_x(MPI_Datatype datatype, MPI_Aint *size)
91
92 int MPI_Type_test(MPI_Datatype datatype, MPI_Aint *size)
93
94 int MPI_Type_test_multiple(MPI_Datatype *types, MPI_Aint *size)
95
96 int MPI_Type_test_x(MPI_Datatype datatype, MPI_Aint *size)
97
98 int MPI_Type_unpack_size(MPI_Datatype datatype, MPI_Comm comm,
99     MPI_Aint *size)
100
101 int MPI_Type_unpack_size_x(MPI_Datatype datatype, MPI_Comm comm,
102     MPI_Aint *size)
103
104 int MPI_Type_write_size(MPI_Datatype datatype, MPI_Aint *size)
105
106 int MPI_Type_write_size_x(MPI_Datatype datatype, MPI_Aint *size)
107
108 int MPI_Type_write_x(MPI_Datatype datatype, MPI_Aint *size)
109
110 int MPI_Type_write_x_multiple(MPI_Datatype *types, MPI_Aint *size)
111
112 int MPI_Type_write_x_size(MPI_Datatype datatype, MPI_Aint *size)
113
114 int MPI_Type_write_x_size_x(MPI_Datatype datatype, MPI_Aint *size)
115
116 int MPI_Type_write_x_x(MPI_Datatype datatype, MPI_Aint *size)
117
118 int MPI_Type_write_x_x_multiple(MPI_Datatype *types, MPI_Aint *size)
119
120 int MPI_Type_write_x_x_size(MPI_Datatype datatype, MPI_Aint *size)
121
122 int MPI_Type_write_x_x_size_x(MPI_Datatype datatype, MPI_Aint *size)
123
124 int MPI_Type_write_x_x_x(MPI_Datatype datatype, MPI_Aint *size)
125
126 int MPI_Type_write_x_x_x_multiple(MPI_Datatype *types, MPI_Aint *size)
127
128 int MPI_Type_write_x_x_x_size(MPI_Datatype datatype, MPI_Aint *size)
129
130 int MPI_Type_write_x_x_x_size_x(MPI_Datatype datatype, MPI_Aint *size)
131
132 int MPI_Type_write_x_x_x_x(MPI_Datatype datatype, MPI_Aint *size)
133
134 int MPI_Type_write_x_x_x_x_multiple(MPI_Datatype *types, MPI_Aint *size)
135
136 int MPI_Type_write_x_x_x_x_size(MPI_Datatype datatype, MPI_Aint *size)
137
138 int MPI_Type_write_x_x_x_x_size_x(MPI_Datatype datatype, MPI_Aint *size)
139
140 int MPI_Type_write_x_x_x_x_x(MPI_Datatype datatype, MPI_Aint *size)
141
142 int MPI_Type_write_x_x_x_x_x_multiple(MPI_Datatype *types, MPI_Aint *size)
143
144 int MPI_Type_write_x_x_x_x_x_size(MPI_Datatype datatype, MPI_Aint *size)
145
146 int MPI_Type_write_x_x_x_x_x_size_x(MPI_Datatype datatype, MPI_Aint *size)
147
148 int MPI_Type_write_x_x_x_x_x_x(MPI_Datatype datatype, MPI_Aint *size)
149
150 int MPI_Type_write_x_x_x_x_x_x_multiple(MPI_Datatype *types, MPI_Aint *size)
151
152 int MPI_Type_write_x_x_x_x_x_x_size(MPI_Datatype datatype, MPI_Aint *size)
153
154 int MPI_Type_write_x_x_x_x_x_x_size_x(MPI_Datatype datatype, MPI_Aint *size)
155
156 int MPI_Type_write_x_x_x_x_x_x_x(MPI_Datatype datatype, MPI_Aint *size)
157
158 int MPI_Type_write_x_x_x_x_x_x_x_multiple(MPI_Datatype *types, MPI_Aint *size)
159
160 int MPI_Type_write_x_x_x_x_x_x_x_size(MPI_Datatype datatype, MPI_Aint *size)
161
162 int MPI_Type_write_x_x_x_x_x_x_x_size_x(MPI_Datatype datatype, MPI_Aint *size)
163
164 int MPI_Type_write_x_x_x_x_x_x_x_x(MPI_Datatype datatype, MPI_Aint *size)
165
166 int MPI_Type_write_x_x_x_x_x_x_x_x_multiple(MPI_Datatype *types, MPI_Aint *size)
167
168 int MPI_Type_write_x_x_x_x_x_x_x_x_size(MPI_Datatype datatype, MPI_Aint *size)
169
170 int MPI_Type_write_x_x_x_x_x_x_x_x_size_x(MPI_Datatype datatype, MPI_Aint *size)
171
172 int MPI_Type_write_x_x_x_x_x_x_x_x_x(MPI_Datatype datatype, MPI_Aint *size)
173
174 int MPI_Type_write_x_x_x_x_x_x_x_x_x_multiple(MPI_Datatype *types, MPI_Aint *size)
175
176 int MPI_Type_write_x_x_x_x_x_x_x_x_x_size(MPI_Datatype datatype, MPI_Aint *size)
177
178 int MPI_Type_write_x_x_x_x_x_x_x_x_x_size_x(MPI_Datatype datatype, MPI_Aint *size)
179
180 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x(MPI_Datatype datatype, MPI_Aint *size)
181
182 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_multiple(MPI_Datatype *types, MPI_Aint *size)
183
184 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_size(MPI_Datatype datatype, MPI_Aint *size)
185
186 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_size_x(MPI_Datatype datatype, MPI_Aint *size)
187
188 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x(MPI_Datatype datatype, MPI_Aint *size)
189
190 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_multiple(MPI_Datatype *types, MPI_Aint *size)
191
192 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_size(MPI_Datatype datatype, MPI_Aint *size)
193
194 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_size_x(MPI_Datatype datatype, MPI_Aint *size)
195
196 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x(MPI_Datatype datatype, MPI_Aint *size)
197
198 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_multiple(MPI_Datatype *types, MPI_Aint *size)
199
200 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_size(MPI_Datatype datatype, MPI_Aint *size)
201
202 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_size_x(MPI_Datatype datatype, MPI_Aint *size)
203
204 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x(MPI_Datatype datatype, MPI_Aint *size)
205
206 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_multiple(MPI_Datatype *types, MPI_Aint *size)
207
208 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_size(MPI_Datatype datatype, MPI_Aint *size)
209
210 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_size_x(MPI_Datatype datatype, MPI_Aint *size)
211
212 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x(MPI_Datatype datatype, MPI_Aint *size)
213
214 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_multiple(MPI_Datatype *types, MPI_Aint *size)
215
216 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_size(MPI_Datatype datatype, MPI_Aint *size)
217
218 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_size_x(MPI_Datatype datatype, MPI_Aint *size)
219
220 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x(MPI_Datatype datatype, MPI_Aint *size)
221
222 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_multiple(MPI_Datatype *types, MPI_Aint *size)
223
224 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_size(MPI_Datatype datatype, MPI_Aint *size)
225
226 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_size_x(MPI_Datatype datatype, MPI_Aint *size)
227
228 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x(MPI_Datatype datatype, MPI_Aint *size)
229
230 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_multiple(MPI_Datatype *types, MPI_Aint *size)
231
232 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_size(MPI_Datatype datatype, MPI_Aint *size)
233
234 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_size_x(MPI_Datatype datatype, MPI_Aint *size)
235
236 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x(MPI_Datatype datatype, MPI_Aint *size)
237
238 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_multiple(MPI_Datatype *types, MPI_Aint *size)
239
240 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_size(MPI_Datatype datatype, MPI_Aint *size)
241
242 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_size_x(MPI_Datatype datatype, MPI_Aint *size)
243
244 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x(MPI_Datatype datatype, MPI_Aint *size)
245
246 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_multiple(MPI_Datatype *types, MPI_Aint *size)
247
248 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_size(MPI_Datatype datatype, MPI_Aint *size)
249
250 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_size_x(MPI_Datatype datatype, MPI_Aint *size)
251
252 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x(MPI_Datatype datatype, MPI_Aint *size)
253
254 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_multiple(MPI_Datatype *types, MPI_Aint *size)
255
256 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_size(MPI_Datatype datatype, MPI_Aint *size)
257
258 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_size_x(MPI_Datatype datatype, MPI_Aint *size)
259
260 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x(MPI_Datatype datatype, MPI_Aint *size)
261
262 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_multiple(MPI_Datatype *types, MPI_Aint *size)
263
264 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_size(MPI_Datatype datatype, MPI_Aint *size)
265
266 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_size_x(MPI_Datatype datatype, MPI_Aint *size)
267
268 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x(MPI_Datatype datatype, MPI_Aint *size)
269
270 int MPI_Type_write_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_x_multiple(MPI_Datatype *types, MPI_Aint *size)
271
2
```



```

1      MPI_Datatype array_of_datatypes[])
2
3  int MPI_Type_get_envelope(MPI_Datatype datatype, int *num_integers,
4      int *num_addresses, int *num_datatypes, int *combiner)
5
6  int MPI_Type_get_extent(MPI_Datatype datatype, MPI_Aint *lb,
7      MPI_Aint *extent)
8
9  int MPI_Type_get_extent_x(MPI_Datatype datatype, MPI_Count *lb,
10     MPI_Count *extent)
11
12 int MPI_Type_get_true_extent(MPI_Datatype datatype, MPI_Aint *true_lb,
13     MPI_Aint *true_extent)
14
15 int MPI_Type_get_true_extent_x(MPI_Datatype datatype, MPI_Count *true_lb,
16     MPI_Count *true_extent)
17
18 int MPI_Type_indexed(int count, int *array_of_blocklengths,
19     int *array_of_displacements, MPI_Datatype oldtype,
20     MPI_Datatype *newtype)
21
22 int MPI_Type_size(MPI_Datatype datatype, int *size)
23
24 int MPI_Type_size_x(MPI_Datatype datatype, MPI_Count *size)
25
26 int MPI_Type_vector(int count, int blocklength, int stride,
27     MPI_Datatype oldtype, MPI_Datatype *newtype)
28
29 int MPI_Unpack(void* inbuf, int insize, int *position, void *outbuf,
30     int outcount, MPI_Datatype datatype, MPI_Comm comm)
31
32 int MPI_Unpack_external(char *datarep, void *inbuf, MPI_Aint insize,
33     MPI_Aint *position, void *outbuf, int outcount,
34     MPI_Datatype datatype)

```

### A.2.3 Collective Communication C Bindings

```

33 int MPI_Allgather(void* sendbuf, int sendcount, MPI_Datatype sendtype,
34     void* recvbuf, int recvcount, MPI_Datatype recvtype,
35     MPI_Comm comm)
36
37 int MPI_Allgatherv(void* sendbuf, int sendcount, MPI_Datatype sendtype,
38     void* recvbuf, int *recvcounts, int *displs,
39     MPI_Datatype recvtype, MPI_Comm comm)
40
41 int MPI_Allgatherw(void* sendbuf, int sendcount, MPI_Datatype sendtype,
42     void* recvbuf, int *recvcounts, int *displs,
43     MPI_Datatype *recvtypes, MPI_Comm comm)
44
45 int MPI_Allreduce(void* sendbuf, void* recvbuf, int count,
46     MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
47
48 int MPI_Alltoall(void* sendbuf, int sendcount, MPI_Datatype sendtype,
49     void* recvbuf, int recvcount, MPI_Datatype recvtype,

```

```

        MPI_Comm comm)
1
2
int MPI_Alltoallv(void* sendbuf, int *sendcounts, int *sdispls,
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
        MPI_Datatype sendtype, void* recvbuf, int *recvcounts,
        int *rdispls, MPI_Datatype recvttype, MPI_Comm comm)

int MPI_Alltoallw(void* sendbuf, int sendcounts[], int sdispls[],
        MPI_Datatype sendtypes[], void* recvbuf, int recvcounts[],
        int rdispls[], MPI_Datatype recvtypes[], MPI_Comm comm)

int MPI_Barrier(MPI_Comm comm)

int MPI_Bcast(void* buffer, int count, MPI_Datatype datatype, int root,
        MPI_Comm comm)

int MPI_Exscan(void* sendbuf, void* recvbuf, int count,
        MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)

int MPI_Gather(void* sendbuf, int sendcount, MPI_Datatype sendtype,
        void* recvbuf, int recvcount, MPI_Datatype recvttype, int root,
        MPI_Comm comm)

int MPI_Gatherv(void* sendbuf, int sendcount, MPI_Datatype sendtype,
        void* recvbuf, int *recvcounts, int *displs,
        MPI_Datatype recvttype, int root, MPI_Comm comm)

int MPI_Gatherw(void* sendbuf, int sendcount, MPI_Datatype sendtype,
        void* recvbuf, int recvcounts[], int displs[],
        MPI_Datatype recvtypes[], int root, MPI_Comm comm)

int MPI_Iallgather(void* sendbuf, int sendcount, MPI_Datatype sendtype,
        void* recvbuf, int recvcount, MPI_Datatype recvttype,
        MPI_Comm comm, MPI_Request *request)

int MPI_Iallgatherv(void* sendbuf, int sendcount, MPI_Datatype sendtype,
        void* recvbuf, int *recvcounts, int *displs,
        MPI_Datatype recvttype, MPI_Comm comm, MPI_Request* request)

int MPI_Iallgatherw(void* sendbuf, int sendcount, MPI_Datatype sendtype,
        void* recvbuf, int recvcounts[], int displs[],
        MPI_Datatype recvtypes[], MPI_Comm comm, MPI_Request *request)

int MPI_Iallreduce(void* sendbuf, void* recvbuf, int count,
        MPI_Datatype datatype, MPI_Op op, MPI_Comm comm,
        MPI_Request *request)

int MPI_Ialltoall(void* sendbuf, int sendcount, MPI_Datatype sendtype,
        void* recvbuf, int recvcount, MPI_Datatype recvttype,
        MPI_Comm comm, MPI_Request *request)

int MPI_Ialltoallv(void* sendbuf, int *sendcounts, int *sdispls,
        MPI_Datatype sendtype, void* recvbuf, int *recvcounts,
        int *rdispls, MPI_Datatype recvttype, MPI_Comm comm,

```

```

1         MPI_Request *request)
2
3     int MPI_Ialltoallw(void* sendbuf, int sendcounts[], int sdispls[],
4         MPI_Datatype sendtypes[], void* recvbuf, int recvcounts[],
5         int rdispls[], MPI_Datatype recvtypes[], MPI_Comm comm,
6         MPI_Request *request)
7
8     int MPI_Ibarrier(MPI_Comm comm, MPI_Request *request)
9
10    int MPI_Ibcast(void* buffer, int count, MPI_Datatype datatype, int root,
11        MPI_Comm comm, MPI_Request *request)
12
13    int MPI_Iexscan(void* sendbuf, void* recvbuf, int count,
14        MPI_Datatype datatype, MPI_Op op, MPI_Comm comm,
15        MPI_Request *request)
16
17    int MPI_Igather(void* sendbuf, int sendcount, MPI_Datatype sendtype,
18        void* recvbuf, int recvcount, MPI_Datatype recvttype, int root,
19        MPI_Comm comm, MPI_Request *request)
20
21    int MPI_Igatherv(void* sendbuf, int sendcount, MPI_Datatype sendtype,
22        void* recvbuf, int *recvcounts, int *displs,
23        MPI_Datatype recvttype, int root, MPI_Comm comm,
24        MPI_Request *request)
25
26    int MPI_Igatherw(void* sendbuf, int sendcount, MPI_Datatype sendtype,
27        void* recvbuf, int recvcounts[], int displs[],
28        MPI_Datatype recvtypes[], int root, MPI_Comm comm,
29        MPI_Request *request)
30
31    int MPI_Ireduce(void* sendbuf, void* recvbuf, int count,
32        MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm comm,
33        MPI_Request *request)
34
35    int MPI_Ireduce_scatter(void* sendbuf, void* recvbuf, int *recvcounts,
36        MPI_Datatype datatype, MPI_Op op, MPI_Comm comm,
37        MPI_Request *request)
38
39    int MPI_Ireduce_scatter_block(void* sendbuf, void* recvbuf, int recvcount,
40        MPI_Datatype datatype, MPI_Op op, MPI_Comm comm,
41        MPI_Request *request)
42
43    int MPI_Iscan(void* sendbuf, void* recvbuf, int count,
44        MPI_Datatype datatype, MPI_Op op, MPI_Comm comm,
45        MPI_Request *request)
46
47    int MPI_Iscatter(void* sendbuf, int sendcount, MPI_Datatype sendtype,
48        void* recvbuf, int recvcount, MPI_Datatype recvttype, int root,
49        MPI_Comm comm, MPI_Request *request)
50
51    int MPI_Iscatterv(void* sendbuf, int *sendcounts, int *displs,
52        MPI_Datatype sendtype, void* recvbuf, int recvcount,
53        MPI_Datatype recvttype, int root, MPI_Comm comm,
54        MPI_Request *request)

```

```

int MPI_Iscatterw(void* sendbuf, int sendcounts[], int displs[],
                 MPI_Datatype sendtypes[], void* recvbuf, int recvcnt,
                 MPI_Datatype recvtype, int root, MPI_Comm comm,
                 MPI_Request *request)
int MPI_Op_commutative(MPI_Op op, int *commute)
int MPI_Op_create(MPI_User_function *function, int commute, MPI_Op *op)
int MPI_Reduce(void* sendbuf, void* recvbuf, int count,
              MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm comm)
int MPI_Reduce_local(void* inbuf, void* inoutbuf, int count,
                   MPI_Datatype datatype, MPI_Op op)
int MPI_Reduce_scatter(void* sendbuf, void* recvbuf, int *recvcounts,
                    MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
int MPI_Reduce_scatter_block(void* sendbuf, void* recvbuf, int recvcnt,
                          MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
int MPI_Scan(void* sendbuf, void* recvbuf, int count,
            MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
int MPI_Scatter(void* sendbuf, int sendcount, MPI_Datatype sendtype,
              void* recvbuf, int recvcnt, MPI_Datatype recvtype, int root,
              MPI_Comm comm)
int MPI_Scatterv(void* sendbuf, int *sendcounts, int *displs,
               MPI_Datatype sendtype, void* recvbuf, int recvcnt,
               MPI_Datatype recvtype, int root, MPI_Comm comm)
int MPI_Scatterw(void* sendbuf, int sendcounts[], int displs[],
               MPI_Datatype sendtypes[], void* recvbuf, int recvcnt,
               MPI_Datatype recvtype, int root, MPI_Comm comm)
int MPI_op_free(MPI_Op *op)

```

#### A.2.4 Groups, Contexts, Communicators, and Caching C Bindings

```

int MPI_COMM_DUP_FN(MPI_Comm oldcomm, int comm_keyval, void *extra_state,
                  void *attribute_val_in, void *attribute_val_out, int *flag)
int MPI_COMM_NULL_COPY_FN(MPI_Comm oldcomm, int comm_keyval,
                        void *extra_state, void *attribute_val_in,
                        void *attribute_val_out, int *flag)
int MPI_COMM_NULL_DELETE_FN(MPI_Comm comm, int comm_keyval, void
                          *attribute_val, void *extra_state)
int MPI_Comm_compare(MPI_Comm comm1, MPI_Comm comm2, int *result)
int MPI_Comm_create(MPI_Comm comm, MPI_Group group, MPI_Comm *newcomm)

```

```

1  int MPI_Comm_create_keyval(MPI_Comm_copy_attr_function *comm_copy_attr_fn,
2                             MPI_Comm_delete_attr_function *comm_delete_attr_fn,
3                             int *comm_keyval, void *extra_state)
4
5  int MPI_Comm_delete_attr(MPI_Comm comm, int comm_keyval)
6
7  int MPI_Comm_dup(MPI_Comm comm, MPI_Comm *newcomm)
8
9  int MPI_Comm_free(MPI_Comm *comm)
10
11 int MPI_Comm_free_keyval(int *comm_keyval)
12
13 int MPI_Comm_get_attr(MPI_Comm comm, int comm_keyval, void *attribute_val,
14                       int *flag)
15
16 int MPI_Comm_get_name(MPI_Comm comm, char *comm_name, int *resultlen)
17
18 int MPI_Comm_group(MPI_Comm comm, MPI_Group *group)
19
20 int MPI_Comm_rank(MPI_Comm comm, int *rank)
21
22 int MPI_Comm_remote_group(MPI_Comm comm, MPI_Group *group)
23
24 int MPI_Comm_remote_size(MPI_Comm comm, int *size)
25
26 int MPI_Comm_set_attr(MPI_Comm comm, int comm_keyval, void *attribute_val)
27
28 int MPI_Comm_set_name(MPI_Comm comm, char *comm_name)
29
30 int MPI_Comm_size(MPI_Comm comm, int *size)
31
32 int MPI_Comm_split(MPI_Comm comm, int color, int key, MPI_Comm *newcomm)
33
34 int MPI_Comm_test_inter(MPI_Comm comm, int *flag)
35
36 int MPI_Group_compare(MPI_Group group1, MPI_Group group2, int *result)
37
38 int MPI_Group_difference(MPI_Group group1, MPI_Group group2,
39                          MPI_Group *newgroup)
40
41 int MPI_Group_excl(MPI_Group group, int n, int *ranks, MPI_Group *newgroup)
42
43 int MPI_Group_free(MPI_Group *group)
44
45 int MPI_Group_incl(MPI_Group group, int n, int *ranks, MPI_Group *newgroup)
46
47 int MPI_Group_intersection(MPI_Group group1, MPI_Group group2,
48                            MPI_Group *newgroup)
49
50 int MPI_Group_range_excl(MPI_Group group, int n, int ranges[][3],
51                          MPI_Group *newgroup)
52
53 int MPI_Group_range_incl(MPI_Group group, int n, int ranges[][3],
54                          MPI_Group *newgroup)
55
56 int MPI_Group_rank(MPI_Group group, int *rank)
57
58 int MPI_Group_size(MPI_Group group, int *size)

```

```

int MPI_Group_translate_ranks (MPI_Group group1, int n, int *ranks1,
                               MPI_Group group2, int *ranks2)
int MPI_Group_union(MPI_Group group1, MPI_Group group2,
                    MPI_Group *newgroup)
int MPI_Intercomm_create(MPI_Comm local_comm, int local_leader,
                          MPI_Comm peer_comm, int remote_leader, int tag,
                          MPI_Comm *newintercomm)
int MPI_Intercomm_merge(MPI_Comm intercomm, int high,
                        MPI_Comm *newintracomm)
int MPI_TYPE_DUP_FN(MPI_Datatype oldtype, int type_keyval,
                    void *extra_state, void *attribute_val_in,
                    void *attribute_val_out, int *flag)
int MPI_TYPE_NULL_COPY_FN(MPI_Datatype oldtype, int type_keyval,
                           void *extra_state, void *attribute_val_in,
                           void *attribute_val_out, int *flag)
int MPI_TYPE_NULL_DELETE_FN(MPI_Datatype type, int type_keyval, void
                             *attribute_val, void *extra_state)
int MPI_Type_create_keyval(MPI_Type_copy_attr_function *type_copy_attr_fn,
                           MPI_Type_delete_attr_function *type_delete_attr_fn,
                           int *type_keyval, void *extra_state)
int MPI_Type_delete_attr(MPI_Datatype type, int type_keyval)
int MPI_Type_free_keyval(int *type_keyval)
int MPI_Type_get_attr(MPI_Datatype type, int type_keyval, void
                      *attribute_val, int *flag)
int MPI_Type_get_name(MPI_Datatype type, char *type_name, int *resultlen)
int MPI_Type_set_attr(MPI_Datatype type, int type_keyval,
                      void *attribute_val)
int MPI_Type_set_name(MPI_Datatype type, char *type_name)
int MPI_WIN_DUP_FN(MPI_Win oldwin, int win_keyval, void *extra_state,
                   void *attribute_val_in, void *attribute_val_out, int *flag)
int MPI_WIN_NULL_COPY_FN(MPI_Win oldwin, int win_keyval, void *extra_state,
                           void *attribute_val_in, void *attribute_val_out, int *flag)
int MPI_WIN_NULL_DELETE_FN(MPI_Win win, int win_keyval, void
                             *attribute_val, void *extra_state)
int MPI_Win_create_keyval(MPI_Win_copy_attr_function *win_copy_attr_fn,
                           MPI_Win_delete_attr_function *win_delete_attr_fn,
                           int *win_keyval, void *extra_state)
int MPI_Win_delete_attr(MPI_Win win, int win_keyval)

```

```

1  int MPI_Win_free_keyval(int *win_keyval)
2
3  int MPI_Win_get_attr(MPI_Win win, int win_keyval, void *attribute_val,
4                      int *flag)
5
6  int MPI_Win_get_name(MPI_Win win, char *win_name, int *resultlen)
7
8  int MPI_Win_set_attr(MPI_Win win, int win_keyval, void *attribute_val)
9
10 int MPI_Win_set_name(MPI_Win win, char *win_name)

```

#### A.2.5 Process Topologies C Bindings

```

12 int MPI_Cart_coords(MPI_Comm comm, int rank, int maxdims, int *coords)
13
14 int MPI_Cart_create(MPI_Comm comm_old, int ndims, int *dims, int *periods,
15                   int reorder, MPI_Comm *comm_cart)
16
17 int MPI_Cart_get(MPI_Comm comm, int maxdims, int *dims, int *periods,
18                 int *coords)
19
20 int MPI_Cart_map(MPI_Comm comm, int ndims, int *dims, int *periods,
21                 int *newrank)
22
23 int MPI_Cart_rank(MPI_Comm comm, int *coords, int *rank)
24
25 int MPI_Cart_shift(MPI_Comm comm, int direction, int disp,
26                   int *rank_source, int *rank_dest)
27
28 int MPI_Cart_sub(MPI_Comm comm, int *remain_dims, MPI_Comm *newcomm)
29
30 int MPI_Cartdim_get(MPI_Comm comm, int *ndims)
31
32 int MPI_Dims_create(int nnodes, int ndims, int *dims)
33
34 int MPI_Dist_graph_create(MPI_Comm comm_old, int n, int sources[],
35                          int degrees[], int destinations[], int weights[],
36                          MPI_Info info, int reorder, MPI_Comm *comm_dist_graph)
37
38 int MPI_Dist_graph_create_adjacent(MPI_Comm comm_old, int indegree,
39                                   int sources[], int sourceweights[], int outdegree,
40                                   int destinations[], int destweights[], MPI_Info info,
41                                   int reorder, MPI_Comm *comm_dist_graph)
42
43 int MPI_Dist_graph_neighbors(MPI_Comm comm, int maxindegree, int sources[],
44                             int sourceweights[], int maxoutdegree, int destinations[],
45                             int destweights[])
46
47 int MPI_Dist_graph_neighbors_count(MPI_Comm comm, int *indegree,
48                                   int *outdegree, int *weighted)
49
50 int MPI_Graph_create(MPI_Comm comm_old, int nnodes, int *index, int *edges,
51                    int reorder, MPI_Comm *comm_graph)
52
53 int MPI_Graph_get(MPI_Comm comm, int maxindex, int maxedges, int *index,
54                  int *edges)

```

```

int MPI_Graph_map(MPI_Comm comm, int nnodes, int *index, int *edges,
                  int *newrank)
int MPI_Graph_neighbors(MPI_Comm comm, int rank, int maxneighbors,
                        int *neighbors)
int MPI_Graph_neighbors_count(MPI_Comm comm, int rank, int *nnighbors)
int MPI_Graphdims_get(MPI_Comm comm, int *nnodes, int *nedges)
int MPI_Topo_test(MPI_Comm comm, int *status)

```

### A.2.6 MPI Environmental Management C Bindings

```

double MPI_Wtick(void)
double MPI_Wtime(void)
int MPI_Abort(MPI_Comm comm, int errorcode)
int MPI_Add_error_class(int *errorclass)
int MPI_Add_error_code(int errorclass, int *errorcode)
int MPI_Add_error_string(int errorcode, char *string)
int MPI_Alloc_mem(MPI_Aint size, MPI_Info info, void *baseptr)
int MPI_Comm_call_errhandler(MPI_Comm comm, int errorcode)
int MPI_Comm_create_errhandler(MPI_Comm_errhandler_function *function,
                               MPI_Errhandler *errhandler)
int MPI_Comm_get_errhandler(MPI_Comm comm, MPI_Errhandler *errhandler)
int MPI_Comm_set_errhandler(MPI_Comm comm, MPI_Errhandler errhandler)
int MPI_Errhandler_free(MPI_Errhandler *errhandler)
int MPI_Error_class(int errorcode, int *errorclass)
int MPI_Error_string(int errorcode, char *string, int *resultlen)
int MPI_File_call_errhandler(MPI_File fh, int errorcode)
int MPI_File_create_errhandler(MPI_File_errhandler_function *function,
                               MPI_Errhandler *errhandler)
int MPI_File_get_errhandler(MPI_File file, MPI_Errhandler *errhandler)
int MPI_File_set_errhandler(MPI_File file, MPI_Errhandler errhandler)
int MPI_Finalize(void)
int MPI_Finalized(int *flag)
int MPI_Free_mem(void *base)
int MPI_Get_processor_name(char *name, int *resultlen)

```



```

1  int MPI_Get_version(int *version, int *subversion)
2
3  int MPI_Init(int *argc, char ***argv)
4
5  int MPI_Initialized(int *flag)
6
7  int MPI_Win_call_errhandler(MPI_Win win, int errorcode)
8
9  int MPI_Win_create_errhandler(MPI_Win_errhandler_function *function,
10                               MPI_Errhandler *errhandler)
11
12 int MPI_Win_get_errhandler(MPI_Win win, MPI_Errhandler *errhandler)
13
14 int MPI_Win_set_errhandler(MPI_Win win, MPI_Errhandler errhandler)

```

#### A.2.7 The Info Object C Bindings

```

15 int MPI_Info_create(MPI_Info *info)
16
17 int MPI_Info_delete(MPI_Info info, char *key)
18
19 int MPI_Info_dup(MPI_Info info, MPI_Info *newinfo)
20
21 int MPI_Info_free(MPI_Info *info)
22
23 int MPI_Info_get(MPI_Info info, char *key, int valuelen, char *value,
24                  int *flag)
25
26 int MPI_Info_get_nkeys(MPI_Info info, int *nkeys)
27
28 int MPI_Info_get_nthkey(MPI_Info info, int n, char *key)
29
30 int MPI_Info_get_valuelen(MPI_Info info, char *key, int *valuelen,
31                           int *flag)
32
33 int MPI_Info_set(MPI_Info info, char *key, char *value)

```

#### A.2.8 Process Creation and Management C Bindings

```

34 int MPI_Close_port(char *port_name)
35
36 int MPI_Comm_accept(char *port_name, MPI_Info info, int root,
37                    MPI_Comm comm, MPI_Comm *newcomm)
38
39 int MPI_Comm_connect(char *port_name, MPI_Info info, int root,
40                      MPI_Comm comm, MPI_Comm *newcomm)
41
42 int MPI_Comm_disconnect(MPI_Comm *comm)
43
44 int MPI_Comm_get_parent(MPI_Comm *parent)
45
46 int MPI_Comm_join(int fd, MPI_Comm *intercomm)
47
48 int MPI_Comm_spawn(char *command, char *argv[], int maxprocs, MPI_Info
49                    info, int root, MPI_Comm comm, MPI_Comm *intercomm,
50                    int array_of_errcodes[])

```

```

int MPI_Comm_spawn_multiple(int count, char *array_of_commands[],
                           char **array_of_argv[], int array_of_maxprocs[],
                           MPI_Info array_of_info[], int root, MPI_Comm comm,
                           MPI_Comm *intercomm, int array_of_errcodes[])
int MPI_Lookup_name(char *service_name, MPI_Info info, char *port_name)
int MPI_Open_port(MPI_Info info, char *port_name)
int MPI_Publish_name(char *service_name, MPI_Info info, char *port_name)
int MPI_Unpublish_name(char *service_name, MPI_Info info, char *port_name)

```

### A.2.9 One-Sided Communications C Bindings

```

int MPI_Accumulate(void *origin_addr, int origin_count,
                  MPI_Datatype origin_datatype, int target_rank,
                  MPI_Aint target_disp, int target_count,
                  MPI_Datatype target_datatype, MPI_Op op, MPI_Win win)
int MPI_Get(void *origin_addr, int origin_count, MPI_Datatype
            origin_datatype, int target_rank, MPI_Aint target_disp, int
            target_count, MPI_Datatype target_datatype, MPI_Win win)
int MPI_Put(void *origin_addr, int origin_count, MPI_Datatype
            origin_datatype, int target_rank, MPI_Aint target_disp, int
            target_count, MPI_Datatype target_datatype, MPI_Win win)
int MPI_Win_complete(MPI_Win win)
int MPI_Win_create(void *base, MPI_Aint size, int disp_unit, MPI_Info info,
                  MPI_Comm comm, MPI_Win *win)
int MPI_Win_fence(int assert, MPI_Win win)
int MPI_Win_free(MPI_Win *win)
int MPI_Win_get_group(MPI_Win win, MPI_Group *group)
int MPI_Win_lock(int lock_type, int rank, int assert, MPI_Win win)
int MPI_Win_post(MPI_Group group, int assert, MPI_Win win)
int MPI_Win_start(MPI_Group group, int assert, MPI_Win win)
int MPI_Win_test(MPI_Win win, int *flag)
int MPI_Win_unlock(int rank, MPI_Win win)
int MPI_Win_wait(MPI_Win win)

```

### A.2.10 External Interfaces C Bindings

```

int MPI_Grequest_complete(MPI_Request request)

```

```

1  int MPI_Grequest_start(MPI_Grequest_query_function *query_fn,
2                        MPI_Grequest_free_function *free_fn,
3                        MPI_Grequest_cancel_function *cancel_fn, void *extra_state,
4                        MPI_Request *request)
5
6  int MPI_Init_thread(int argc, char *((*argv)[[]]), int required,
7                    int *provided)
8
9  int MPI_Is_thread_main(int *flag)
10
11 int MPI_Query_thread(int *provided)
12
13 int MPI_Status_set_cancelled(MPI_Status *status, int flag)
14
15 int MPI_Status_set_elements(MPI_Status *status, MPI_Datatype datatype,
16                             int count)
17
18 int MPI_Status_set_elements_x(MPI_Status *status, MPI_Datatype datatype,
19                               MPI_Count count)

```

#### A.2.11 I/O C Bindings

```

20 int MPI_File_close(MPI_File *fh)
21
22 int MPI_File_delete(char *filename, MPI_Info info)
23
24 int MPI_File_get_amode(MPI_File fh, int *amode)
25
26 int MPI_File_get_atomicsity(MPI_File fh, int *flag)
27
28 int MPI_File_get_byte_offset(MPI_File fh, MPI_Offset offset,
29                             MPI_Offset *disp)
30
31 int MPI_File_get_group(MPI_File fh, MPI_Group *group)
32
33 int MPI_File_get_info(MPI_File fh, MPI_Info *info_used)
34
35 int MPI_File_get_position(MPI_File fh, MPI_Offset *offset)
36
37 int MPI_File_get_position_shared(MPI_File fh, MPI_Offset *offset)
38
39 int MPI_File_get_size(MPI_File fh, MPI_Offset *size)
40
41 int MPI_File_get_type_extent(MPI_File fh, MPI_Datatype datatype,
42                             MPI_Aint *extent)
43
44 int MPI_File_get_view(MPI_File fh, MPI_Offset *disp, MPI_Datatype *etype,
45                      MPI_Datatype *filetype, char *datarep)
46
47 int MPI_File_iread(MPI_File fh, void *buf, int count,
48                  MPI_Datatype datatype, MPI_Request *request)
49
50 int MPI_File_iread_at(MPI_File fh, MPI_Offset offset, void *buf, int count,
51                      MPI_Datatype datatype, MPI_Request *request)
52
53 int MPI_File_iread_shared(MPI_File fh, void *buf, int count,
54                          MPI_Datatype datatype, MPI_Request *request)

```

```

int MPI_File_iread(MPI_File fh, void *buf, int count,
    MPI_Datatype datatype, MPI_Request *request)
int MPI_File_iread_at(MPI_File fh, MPI_Offset offset, void *buf,
    int count, MPI_Datatype datatype, MPI_Request *request)
int MPI_File_iread_shared(MPI_File fh, void *buf, int count,
    MPI_Datatype datatype, MPI_Request *request)
int MPI_File_open(MPI_Comm comm, char *filename, int amode, MPI_Info info,
    MPI_File *fh)
int MPI_File_preallocate(MPI_File fh, MPI_Offset size)
int MPI_File_read(MPI_File fh, void *buf, int count, MPI_Datatype datatype,
    MPI_Status *status)
int MPI_File_read_all(MPI_File fh, void *buf, int count,
    MPI_Datatype datatype, MPI_Status *status)
int MPI_File_read_all_begin(MPI_File fh, void *buf, int count,
    MPI_Datatype datatype)
int MPI_File_read_all_end(MPI_File fh, void *buf, MPI_Status *status)
int MPI_File_read_at(MPI_File fh, MPI_Offset offset, void *buf, int count,
    MPI_Datatype datatype, MPI_Status *status)
int MPI_File_read_at_all(MPI_File fh, MPI_Offset offset, void *buf,
    int count, MPI_Datatype datatype, MPI_Status *status)
int MPI_File_read_at_all_begin(MPI_File fh, MPI_Offset offset, void *buf,
    int count, MPI_Datatype datatype)
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status)
int MPI_File_read_ordered(MPI_File fh, void *buf, int count,
    MPI_Datatype datatype, MPI_Status *status)
int MPI_File_read_ordered_begin(MPI_File fh, void *buf, int count,
    MPI_Datatype datatype)
int MPI_File_read_ordered_end(MPI_File fh, void *buf, MPI_Status *status)
int MPI_File_read_shared(MPI_File fh, void *buf, int count,
    MPI_Datatype datatype, MPI_Status *status)
int MPI_File_seek(MPI_File fh, MPI_Offset offset, int whence)
int MPI_File_seek_shared(MPI_File fh, MPI_Offset offset, int whence)
int MPI_File_set_atomicity(MPI_File fh, int flag)
int MPI_File_set_info(MPI_File fh, MPI_Info info)
int MPI_File_set_size(MPI_File fh, MPI_Offset size)

```

```

1  int MPI_File_set_view(MPI_File fh, MPI_Offset disp, MPI_Datatype etype,
2      MPI_Datatype filetype, char *datarep, MPI_Info info)
3
4  int MPI_File_sync(MPI_File fh)
5
6  int MPI_File_write(MPI_File fh, void *buf, int count,
7      MPI_Datatype datatype, MPI_Status *status)
8
9  int MPI_File_write_all(MPI_File fh, void *buf, int count,
10     MPI_Datatype datatype, MPI_Status *status)
11
12 int MPI_File_write_all_begin(MPI_File fh, void *buf, int count,
13     MPI_Datatype datatype)
14
15 int MPI_File_write_all_end(MPI_File fh, void *buf, MPI_Status *status)
16
17 int MPI_File_write_at(MPI_File fh, MPI_Offset offset, void *buf, int count,
18     MPI_Datatype datatype, MPI_Status *status)
19
20 int MPI_File_write_at_all(MPI_File fh, MPI_Offset offset, void *buf,
21     int count, MPI_Datatype datatype, MPI_Status *status)
22
23 int MPI_File_write_at_all_begin(MPI_File fh, MPI_Offset offset, void *buf,
24     int count, MPI_Datatype datatype)
25
26 int MPI_File_write_at_all_end(MPI_File fh, void *buf, MPI_Status *status)
27
28 int MPI_File_write_ordered(MPI_File fh, void *buf, int count,
29     MPI_Datatype datatype, MPI_Status *status)
30
31 int MPI_File_write_ordered_begin(MPI_File fh, void *buf, int count,
32     MPI_Datatype datatype)
33
34 int MPI_File_write_ordered_end(MPI_File fh, void *buf, MPI_Status *status)
35
36 int MPI_File_write_shared(MPI_File fh, void *buf, int count,
37     MPI_Datatype datatype, MPI_Status *status)
38
39 int MPI_Register_datarep(char *datarep,
40     MPI_Datarep_conversion_function *read_conversion_fn,
41     MPI_Datarep_conversion_function *write_conversion_fn,
42     MPI_Datarep_extent_function *dtype_file_extent_fn,
43     void *extra_state)

```

#### A.2.12 Language Bindings C Bindings

```

40 int MPI_Type_create_f90_complex(int p, int r, MPI_Datatype *newtype)
41
42 int MPI_Type_create_f90_integer(int r, MPI_Datatype *newtype)
43
44 int MPI_Type_create_f90_real(int p, int r, MPI_Datatype *newtype)
45
46 int MPI_Type_match_size(int typeclass, int size, MPI_Datatype *type)
47
48 MPI_Fint MPI_Comm_c2f(MPI_Comm comm)

```

```

MPI_Comm MPI_Comm_f2c(MPI_Fint comm)
MPI_Fint MPI_Errhandler_c2f(MPI_Errhandler errhandler)
MPI_Errhandler MPI_Errhandler_f2c(MPI_Fint errhandler)
MPI_Fint MPI_File_c2f(MPI_File file)
MPI_File MPI_File_f2c(MPI_Fint file)
MPI_Fint MPI_Group_c2f(MPI_Group group)
MPI_Group MPI_Group_f2c(MPI_Fint group)
MPI_Fint MPI_Info_c2f(MPI_Info info)
MPI_Info MPI_Info_f2c(MPI_Fint info)
MPI_Fint MPI_Op_c2f(MPI_Op op)
MPI_Op MPI_Op_f2c(MPI_Fint op)
MPI_Fint MPI_Request_c2f(MPI_Request request)
MPI_Request MPI_Request_f2c(MPI_Fint request)
int MPI_Status_c2f(MPI_Status *c_status, MPI_Fint *f_status)
int MPI_Status_f2c(MPI_Fint *f_status, MPI_Status *c_status)
MPI_Fint MPI_Type_c2f(MPI_Datatype datatype)
MPI_Datatype MPI_Type_f2c(MPI_Fint datatype)
MPI_Fint MPI_Win_c2f(MPI_Win win)
MPI_Win MPI_Win_f2c(MPI_Fint win)

```

### A.2.13 Profiling Interface C Bindings

```
int MPI_Pcontrol(const int level, ...)
```

### A.2.14 Deprecated C Bindings

```

int MPI_Address(void* location, MPI_Aint *address)
int MPI_Attr_delete(MPI_Comm comm, int keyval)
int MPI_Attr_get(MPI_Comm comm, int keyval, void *attribute_val, int *flag)
int MPI_Attr_put(MPI_Comm comm, int keyval, void* attribute_val)
int MPI_DUP_FN(MPI_Comm oldcomm, int keyval, void *extra_state,
               void *attribute_val_in, void *attribute_val_out, int *flag)
int MPI_Errhandler_create(MPI_Handler_function *function,
                          MPI_Errhandler *errhandler)

```

```
1  int MPI_Errhandler_get(MPI_Comm comm, MPI_Errhandler *errhandler)
2
3  int MPI_Errhandler_set(MPI_Comm comm, MPI_Errhandler errhandler)
4
5  int MPI_Keyval_create(MPI_Copy_function *copy_fn, MPI_Delete_function
6      *delete_fn, int *keyval, void* extra_state)
7
8  int MPI_Keyval_free(int *keyval)
9
10 int MPI_NULL_COPY_FN(MPI_Comm oldcomm, int keyval, void *extra_state,
11     void *attribute_val_in, void *attribute_val_out, int *flag)
12
13 int MPI_NULL_DELETE_FN(MPI_Comm comm, int keyval, void *attribute_val,
14     void *extra_state)
15
16 int MPI_Type_extent(MPI_Datatype datatype, MPI_Aint *extent)
17
18 int MPI_Type_hindexed(int count, int *array_of_blocklengths,
19     MPI_Aint *array_of_displacements, MPI_Datatype oldtype,
20     MPI_Datatype *newtype)
21
22 int MPI_Type_hvector(int count, int blocklength, MPI_Aint stride,
23     MPI_Datatype oldtype, MPI_Datatype *newtype)
24
25 int MPI_Type_lb(MPI_Datatype datatype, MPI_Aint* displacement)
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
```

## A.3 Fortran Bindings

### A.3.1 Point-to-Point Communication Fortran Bindings

```

MPI_BSEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, IERROR)
    <type> BUF(*)
    INTEGER COUNT, DATATYPE, DEST, TAG, COMM, IERROR

MPI_BSEND_INIT(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
    <type> BUF(*)
    INTEGER REQUEST, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR

MPI_BUFFER_ATTACH(BUFFER, SIZE, IERROR)
    <type> BUFFER(*)
    INTEGER SIZE, IERROR

MPI_BUFFER_DETACH(BUFFER_ADDR, SIZE, IERROR)
    <type> BUFFER_ADDR(*)
    INTEGER SIZE, IERROR

MPI_CANCEL(REQUEST, IERROR)
    INTEGER REQUEST, IERROR

MPI_GET_COUNT(STATUS, DATATYPE, COUNT, IERROR)
    INTEGER STATUS(MPI_STATUS_SIZE), DATATYPE, COUNT, IERROR

MPI_IBSEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
    <type> BUF(*)
    INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR

MPI_IPROBE(SOURCE, TAG, COMM, FLAG, STATUS, IERROR)
    LOGICAL FLAG
    INTEGER SOURCE, TAG, COMM, STATUS(MPI_STATUS_SIZE), IERROR

MPI_IRECV(BUF, COUNT, DATATYPE, SOURCE, TAG, COMM, REQUEST, IERROR)
    <type> BUF(*)
    INTEGER COUNT, DATATYPE, SOURCE, TAG, COMM, REQUEST, IERROR

MPI_IRSEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
    <type> BUF(*)
    INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR

MPI_ISEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
    <type> BUF(*)
    INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR

MPI_ISSEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
    <type> BUF(*)
    INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR

MPI_PROBE(SOURCE, TAG, COMM, STATUS, IERROR)
    INTEGER SOURCE, TAG, COMM, STATUS(MPI_STATUS_SIZE), IERROR

MPI_RECV(BUF, COUNT, DATATYPE, SOURCE, TAG, COMM, STATUS, IERROR)

```



```

1      <type> BUF(*)
2      INTEGER COUNT, DATATYPE, SOURCE, TAG, COMM, STATUS(MPI_STATUS_SIZE),
3      IERROR
4
5      MPI_RECV_INIT(BUF, COUNT, DATATYPE, SOURCE, TAG, COMM, REQUEST, IERROR)
6      <type> BUF(*)
7      INTEGER COUNT, DATATYPE, SOURCE, TAG, COMM, REQUEST, IERROR
8
9      MPI_REQUEST_FREE(REQUEST, IERROR)
10     INTEGER REQUEST, IERROR
11
12     MPI_REQUEST_GET_STATUS( REQUEST, FLAG, STATUS, IERROR)
13     INTEGER REQUEST, STATUS(MPI_STATUS_SIZE), IERROR
14     LOGICAL FLAG
15
16     MPI_RSEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, IERROR)
17     <type> BUF(*)
18     INTEGER COUNT, DATATYPE, DEST, TAG, COMM, IERROR
19
20     MPI_RSEND_INIT(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
21     <type> BUF(*)
22     INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR
23
24     MPI_SEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, IERROR)
25     <type> BUF(*)
26     INTEGER COUNT, DATATYPE, DEST, TAG, COMM, IERROR
27
28     MPI_SENDRECV(SENDBUF, SENDCOUNT, SENDTYPE, DEST, SENDTAG, RECVBUF,
29     RECVCOUNT, RECVMODE, SOURCE, RECVTAG, COMM, STATUS, IERROR)
30     <type> SENDBUF(*), RECVBUF(*)
31     INTEGER SENDCOUNT, SENDTYPE, DEST, SENDTAG, RECVCOUNT, RECVMODE,
32     SOURCE, RECVTAG, COMM, STATUS(MPI_STATUS_SIZE), IERROR
33
34     MPI_SENDRECV_REPLACE(BUF, COUNT, DATATYPE, DEST, SENDTAG, SOURCE, RECVTAG,
35     COMM, STATUS, IERROR)
36     <type> BUF(*)
37     INTEGER COUNT, DATATYPE, DEST, SENDTAG, SOURCE, RECVTAG, COMM,
38     STATUS(MPI_STATUS_SIZE), IERROR
39
40     MPI_SEND_INIT(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
41     <type> BUF(*)
42     INTEGER REQUEST, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR
43
44     MPI_SSEND(BUF, COUNT, DATATYPE, DEST, TAG, COMM, IERROR)
45     <type> BUF(*)
46     INTEGER COUNT, DATATYPE, DEST, TAG, COMM, IERROR
47
48     MPI_SSEND_INIT(BUF, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR)
49     <type> BUF(*)
50     INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR
51
52     MPI_START(REQUEST, IERROR)
53     INTEGER REQUEST, IERROR

```

```

MPI_STARTALL(COUNT, ARRAY_OF_REQUESTS, IERROR)
    INTEGER COUNT, ARRAY_OF_REQUESTS(*), IERROR
MPI_TEST(REQUEST, FLAG, STATUS, IERROR)
    LOGICAL FLAG
    INTEGER REQUEST, STATUS(MPI_STATUS_SIZE), IERROR
MPI_TESTALL(COUNT, ARRAY_OF_REQUESTS, FLAG, ARRAY_OF_STATUSES, IERROR)
    LOGICAL FLAG
    INTEGER COUNT, ARRAY_OF_REQUESTS(*),
    ARRAY_OF_STATUSES(MPI_STATUS_SIZE,*), IERROR
MPI_TESTANY(COUNT, ARRAY_OF_REQUESTS, INDEX, FLAG, STATUS, IERROR)
    LOGICAL FLAG
    INTEGER COUNT, ARRAY_OF_REQUESTS(*), INDEX, STATUS(MPI_STATUS_SIZE),
    IERROR
MPI_TESTSOME(INCOUNT, ARRAY_OF_REQUESTS, OUTCOUNT, ARRAY_OF_INDICES,
    ARRAY_OF_STATUSES, IERROR)
    INTEGER INCOUNT, ARRAY_OF_REQUESTS(*), OUTCOUNT, ARRAY_OF_INDICES(*),
    ARRAY_OF_STATUSES(MPI_STATUS_SIZE,*), IERROR
MPI_TEST_CANCELLED(STATUS, FLAG, IERROR)
    LOGICAL FLAG
    INTEGER STATUS(MPI_STATUS_SIZE), IERROR
MPI_WAIT(REQUEST, STATUS, IERROR)
    INTEGER REQUEST, STATUS(MPI_STATUS_SIZE), IERROR
MPI_WAITALL(COUNT, ARRAY_OF_REQUESTS, ARRAY_OF_STATUSES, IERROR)
    INTEGER COUNT, ARRAY_OF_REQUESTS(*)
    INTEGER ARRAY_OF_STATUSES(MPI_STATUS_SIZE,*), IERROR
MPI_WAITANY(COUNT, ARRAY_OF_REQUESTS, INDEX, STATUS, IERROR)
    INTEGER COUNT, ARRAY_OF_REQUESTS(*), INDEX, STATUS(MPI_STATUS_SIZE),
    IERROR
MPI_WAITSOME(INCOUNT, ARRAY_OF_REQUESTS, OUTCOUNT, ARRAY_OF_INDICES,
    ARRAY_OF_STATUSES, IERROR)
    INTEGER INCOUNT, ARRAY_OF_REQUESTS(*), OUTCOUNT, ARRAY_OF_INDICES(*),
    ARRAY_OF_STATUSES(MPI_STATUS_SIZE,*), IERROR

```

### A.3.2 Datatypes Fortran Bindings

```

MPI_GET_ADDRESS(LOCATION, ADDRESS, IERROR)
    <type> LOCATION(*)
    INTEGER IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) ADDRESS
MPI_GET_ELEMENTS(STATUS, DATATYPE, COUNT, IERROR)
    INTEGER STATUS(MPI_STATUS_SIZE), DATATYPE, COUNT, IERROR
MPI_GET_ELEMENTS_X(STATUS, DATATYPE, COUNT, IERROR)

```

```

1      INTEGER STATUS(MPI_STATUS_SIZE), DATATYPE, IERROR
2      INTEGER (KIND=MPI_COUNT_KIND) COUNT
3
4      MPI_PACK(INBUF, INCOUNT, DATATYPE, OUTBUF, OUTSIZE, POSITION, COMM, IERROR)
5          <type> INBUF(*), OUTBUF(*)
6          INTEGER INCOUNT, DATATYPE, OUTSIZE, POSITION, COMM, IERROR
7
8      MPI_PACK_EXTERNAL(DATAREP, INBUF, INCOUNT, DATATYPE, OUTBUF, OUTSIZE,
9          POSITION, IERROR)
10         INTEGER INCOUNT, DATATYPE, IERROR
11         INTEGER(KIND=MPI_ADDRESS_KIND) OUTSIZE, POSITION
12         CHARACTER*(*) DATAREP
13         <type> INBUF(*), OUTBUF(*)
14
15     MPI_PACK_EXTERNAL_SIZE(DATAREP, INCOUNT, DATATYPE, SIZE, IERROR)
16         INTEGER INCOUNT, DATATYPE, IERROR
17         INTEGER(KIND=MPI_ADDRESS_KIND) SIZE
18         CHARACTER*(*) DATAREP
19
20     MPI_PACK_SIZE(INCOUNT, DATATYPE, COMM, SIZE, IERROR)
21         INTEGER INCOUNT, DATATYPE, COMM, SIZE, IERROR
22
23     MPI_TYPE_COMMIT(DATATYPE, IERROR)
24         INTEGER DATATYPE, IERROR
25
26     MPI_TYPE_CONTIGUOUS(COUNT, OLDTYPE, NEWTYPE, IERROR)
27         INTEGER COUNT, OLDTYPE, NEWTYPE, IERROR
28
29     MPI_TYPE_CREATE_DARRAY(SIZE, RANK, NDIMS, ARRAY_OF_GSIZES,
30         ARRAY_OF_DISTRIBS, ARRAY_OF_DARGS, ARRAY_OF_PSIZEs, ORDER,
31         OLDTYPE, NEWTYPE, IERROR)
32         INTEGER SIZE, RANK, NDIMS, ARRAY_OF_GSIZES(*), ARRAY_OF_DISTRIBS(*),
33         ARRAY_OF_DARGS(*), ARRAY_OF_PSIZEs(*), ORDER, OLDTYPE, NEWTYPE, IERROR
34
35     MPI_TYPE_CREATE_HINDEXED(COUNT, ARRAY_OF_BLOCKLENGTHS,
36         ARRAY_OF_DISPLACEMENTS, OLDTYPE, NEWTYPE, IERROR)
37         INTEGER COUNT, ARRAY_OF_BLOCKLENGTHS(*), OLDTYPE, NEWTYPE, IERROR
38         INTEGER(KIND=MPI_ADDRESS_KIND) ARRAY_OF_DISPLACEMENTS(*)
39
40     MPI_TYPE_CREATE_HVECTOR(COUNT, BLOCKLENGTH, STRIDE, OLDTYPE, NEWTYPE,
41         IERROR)
42         INTEGER COUNT, BLOCKLENGTH, OLDTYPE, NEWTYPE, IERROR
43         INTEGER(KIND=MPI_ADDRESS_KIND) STRIDE
44
45     MPI_TYPE_CREATE_INDEXED_BLOCK(COUNT, BLOCKLENGTH, ARRAY_OF_DISPLACEMENTS,
46         OLDTYPE, NEWTYPE, IERROR)
47         INTEGER COUNT, BLOCKLENGTH, ARRAY_OF_DISPLACEMENTS(*), OLDTYPE,
48         NEWTYPE, IERROR
49
50     MPI_TYPE_CREATE_RESIZED(OLDTYPE, LB, EXTENT, NEWTYPE, IERROR)
51         INTEGER OLDTYPE, NEWTYPE, IERROR
52         INTEGER(KIND=MPI_ADDRESS_KIND) LB, EXTENT

```

```

MPI_TYPE_CREATE_STRUCT(COUNT, ARRAY_OF_BLOCKLENGTHS,
                      ARRAY_OF_DISPLACEMENTS, ARRAY_OF_TYPES, NEWTYPE, IERROR)
INTEGER COUNT, ARRAY_OF_BLOCKLENGTHS(*), ARRAY_OF_TYPES(*), NEWTYPE,
IERROR
INTEGER(KIND=MPI_ADDRESS_KIND) ARRAY_OF_DISPLACEMENTS(*)

MPI_TYPE_CREATE_SUBARRAY(NDIMS, ARRAY_OF_SIZES, ARRAY_OF_SUBSIZES,
                      ARRAY_OF_STARTS, ORDER, OLDTYPE, NEWTYPE, IERROR)
INTEGER NDIMS, ARRAY_OF_SIZES(*), ARRAY_OF_SUBSIZES(*),
ARRAY_OF_STARTS(*), ORDER, OLDTYPE, NEWTYPE, IERROR

MPI_TYPE_DUP(TYPE, NEWTYPE, IERROR)
INTEGER TYPE, NEWTYPE, IERROR

MPI_TYPE_FREE(DATATYPE, IERROR)
INTEGER DATATYPE, IERROR

MPI_TYPE_GET_CONTENTS(DATATYPE, MAX_INTEGERS, MAX_ADDRESSES, MAX_DATATYPES,
                    ARRAY_OF_INTEGERS, ARRAY_OF_ADDRESSES, ARRAY_OF_DATATYPES,
                    IERROR)
INTEGER DATATYPE, MAX_INTEGERS, MAX_ADDRESSES, MAX_DATATYPES,
ARRAY_OF_INTEGERS(*), ARRAY_OF_DATATYPES(*), IERROR
INTEGER(KIND=MPI_ADDRESS_KIND) ARRAY_OF_ADDRESSES(*)

MPI_TYPE_GET_ENVELOPE(DATATYPE, NUM_INTEGERS, NUM_ADDRESSES, NUM_DATATYPES,
                    COMBINER, IERROR)
INTEGER DATATYPE, NUM_INTEGERS, NUM_ADDRESSES, NUM_DATATYPES, COMBINER,
IERROR

MPI_TYPE_GET_EXTENT(DATATYPE, LB, EXTENT, IERROR)
INTEGER DATATYPE, IERROR
INTEGER(KIND = MPI_ADDRESS_KIND) LB, EXTENT

MPI_TYPE_GET_EXTENT_X(DATATYPE, LB, EXTENT, IERROR)
INTEGER DATATYPE, IERROR
INTEGER(KIND = MPI_COUNT_KIND) LB, EXTENT

MPI_TYPE_GET_TRUE_EXTENT(DATATYPE, TRUE_LB, TRUE_EXTENT, IERROR)
INTEGER DATATYPE, IERROR
INTEGER(KIND = MPI_ADDRESS_KIND) TRUE_LB, TRUE_EXTENT

MPI_TYPE_GET_TRUE_EXTENT_X(DATATYPE, TRUE_LB, TRUE_EXTENT, IERROR)
INTEGER DATATYPE, IERROR
INTEGER(KIND = MPI_COUNT_KIND) TRUE_LB, TRUE_EXTENT

MPI_TYPE_INDEXED(COUNT, ARRAY_OF_BLOCKLENGTHS, ARRAY_OF_DISPLACEMENTS,
                OLDTYPE, NEWTYPE, IERROR)
INTEGER COUNT, ARRAY_OF_BLOCKLENGTHS(*), ARRAY_OF_DISPLACEMENTS(*),
OLDTYPE, NEWTYPE, IERROR

MPI_TYPE_SIZE(DATATYPE, SIZE, IERROR)
INTEGER DATATYPE, SIZE, IERROR

```

```

1  MPI_TYPE_SIZE_X(DATATYPE, SIZE, IERROR)
2      INTEGER DATATYPE, IERROR
3      INTEGER (KIND=MPI_COUNT_KIND) SIZE
4
5  MPI_TYPE_VECTOR(COUNT, BLOCKLENGTH, STRIDE, OLDTYPE, NEWTYPE, IERROR)
6      INTEGER COUNT, BLOCKLENGTH, STRIDE, OLDTYPE, NEWTYPE, IERROR
7
8  MPI_UNPACK(INBUF, INSIZE, POSITION, OUTBUF, OUTCOUNT, DATATYPE, COMM,
9      IERROR)
10     <type> INBUF(*), OUTBUF(*)
11     INTEGER INSIZE, POSITION, OUTCOUNT, DATATYPE, COMM, IERROR
12
13 MPI_UNPACK_EXTERNAL(DATAREP, INBUF, INSIZE, POSITION, OUTBUF, OUTCOUNT,
14     DATATYPE, IERROR)
15     INTEGER OUTCOUNT, DATATYPE, IERROR
16     INTEGER(KIND=MPI_ADDRESS_KIND) INSIZE, POSITION
17     CHARACTER*(*) DATAREP
18     <type> INBUF(*), OUTBUF(*)
19

```

### A.3.3 Collective Communication Fortran Bindings

```

20
21 MPI_ALLGATHER(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNT, RECVTYPE,
22     COMM, IERROR)
23     <type> SENDBUF(*), RECVBUF(*)
24     INTEGER SENDCOUNT, SENDTYPE, RECVCOUNT, RECVTYPE, COMM, IERROR
25
26 MPI_ALLGATHERV(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNTS, DISPLS,
27     RECVTYPE, COMM, IERROR)
28     <type> SENDBUF(*), RECVBUF(*)
29     INTEGER SENDCOUNT, SENDTYPE, RECVCOUNTS(*), DISPLS(*), RECVTYPE, COMM,
30     IERROR
31
32 MPI_ALLGATHERW(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNTS, DISPLS,
33     RECVTYPES, COMM, IERROR)
34     <type> SENDBUF(*), RECVBUF(*)
35     INTEGER SENDCOUNT, SENDTYPE, RECVCOUNTS(*), DISPLS(*), RECVTYPES(*),
36     COMM, IERROR
37
38 MPI_ALLREDUCE(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, COMM, IERROR)
39     <type> SENDBUF(*), RECVBUF(*)
40     INTEGER COUNT, DATATYPE, OP, COMM, IERROR
41
42 MPI_ALLTOALL(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNT, RECVTYPE,
43     COMM, IERROR)
44     <type> SENDBUF(*), RECVBUF(*)
45     INTEGER SENDCOUNT, SENDTYPE, RECVCOUNT, RECVTYPE, COMM, IERROR
46
47 MPI_ALLTOALLV(SENDBUF, SENDCOUNTS, SDISPLS, SENDTYPE, RECVBUF, RECVCOUNTS,
48     RDISPLS, RECVTYPE, COMM, IERROR)
49     <type> SENDBUF(*), RECVBUF(*)
50

```

```

    INTEGER SENDCOUNTS(*), SDISPLS(*), SENDTYPE, RECVCOUNTS(*), RDISPLS(*),
    RECVTYPE, COMM, IERROR
MPI_ALLTOALLW(SENDBUF, SENDCOUNTS, SDISPLS, SENDTYPES, RECVBUF, RECVCOUNTS,
    RDISPLS, RECVTYPES, COMM, IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER SENDCOUNTS(*), SDISPLS(*), SENDTYPES(*), RECVCOUNTS(*),
    RDISPLS(*), RECVTYPES(*), COMM, IERROR
MPI_BARRIER(COMM, IERROR)
    INTEGER COMM, IERROR
MPI_BCAST(BUFFER, COUNT, DATATYPE, ROOT, COMM, IERROR)
    <type> BUFFER(*)
    INTEGER COUNT, DATATYPE, ROOT, COMM, IERROR
MPI_EXSCAN(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, COMM, IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER COUNT, DATATYPE, OP, COMM, IERROR
MPI_GATHER(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNT, RECVTYPE,
    ROOT, COMM, IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER SENDCOUNT, SENDTYPE, RECVCOUNT, RECVTYPE, ROOT, COMM, IERROR
MPI_GATHERV(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNTS, DISPLS,
    RECVTYPE, ROOT, COMM, IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER SENDCOUNT, SENDTYPE, RECVCOUNTS(*), DISPLS(*), RECVTYPE, ROOT,
    COMM, IERROR
MPI_GATHERW(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNTS, DISPLS,
    RECVTYPES, ROOT, COMM, IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER SENDCOUNT, SENDTYPE, RECVCOUNTS(*), DISPLS(*), RECVTYPES(*),
    ROOT, COMM, IERROR
MPI_IALLGATHER(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNT, RECVTYPE,
    COMM, REQUEST, IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER SENDCOUNT, SENDTYPE, RECVCOUNT, RECVTYPE, COMM, REQUEST, IERROR
MPI_IALLGATHERV(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNTS, DISPLS,
    RECVTYPE, COMM, REQUEST, IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER SENDCOUNT, SENDTYPE, RECVCOUNTS(*), DISPLS(*), RECVTYPE, COMM,
    REQUEST, IERROR
MPI_IALLGATHERW(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNTS, DISPLS,
    RECVTYPES, COMM, REQUEST, IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER SENDCOUNT, SENDTYPE, RECVCOUNTS(*), DISPLS(*), RECVTYPES(*),

```

```

1      COMM, REQUEST, IERROR
2
3      MPI_IALLREDUCE(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, COMM, REQUEST,
4                    IERROR)
5      <type> SENDBUF(*), RECVBUF(*)
6      INTEGER COUNT, DATATYPE, OP, COMM, REQUEST, IERROR
7
8      MPI_IALLTOALL(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNT, RECVTYPE,
9                    COMM, REQUEST, IERROR)
10     <type> SENDBUF(*), RECVBUF(*)
11     INTEGER SENDCOUNT, SENDTYPE, RECVCOUNT, RECVTYPE, COMM, REQUEST, IERROR
12
13     MPI_IALLTOALLV(SENDBUF, SENDCOUNTS, SDISPLS, SENDTYPE, RECVBUF, RECVCOUNTS,
14                    RDISPLS, RECVTYPE, COMM, REQUEST, IERROR)
15     <type> SENDBUF(*), RECVBUF(*)
16     INTEGER SENDCOUNTS(*), SDISPLS(*), SENDTYPE, RECVCOUNTS(*), RDISPLS(*),
17     RECVTYPE, COMM, REQUEST, IERROR
18
19     MPI_IALLTOALLW(SENDBUF, SENDCOUNTS, SDISPLS, SENDTYPES, RECVBUF,
20                    RECVCOUNTS, RDISPLS, RECVTYPES, COMM, REQUEST, IERROR)
21     <type> SENDBUF(*), RECVBUF(*)
22     INTEGER SENDCOUNTS(*), SDISPLS(*), SENDTYPES(*), RECVCOUNTS(*),
23     RDISPLS(*), RECVTYPES(*), COMM, REQUEST, IERROR
24
25     MPI_IBARRIER(COMM, REQUEST, IERROR)
26     INTEGER COMM, REQUEST, IERROR
27
28     MPI_IBCAST(BUFFER, COUNT, DATATYPE, ROOT, COMM, REQUEST, IERROR)
29     <type> BUFFER(*)
30     INTEGER COUNT, DATATYPE, ROOT, COMM, REQUEST, IERROR
31
32     MPI_IEXSCAN(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, COMM, REQUEST, IERROR)
33     <type> SENDBUF(*), RECVBUF(*)
34     INTEGER COUNT, DATATYPE, OP, COMM, REQUEST, IERROR
35
36     MPI_IGATHER(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNT, RECVTYPE,
37                ROOT, COMM, REQUEST, IERROR)
38     <type> SENDBUF(*), RECVBUF(*)
39     INTEGER SENDCOUNT, SENDTYPE, RECVCOUNT, RECVTYPE, ROOT, COMM, REQUEST,
40     IERROR
41
42     MPI_IGATHERV(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNTS, DISPLS,
43                RECVTYPE, ROOT, COMM, REQUEST, IERROR)
44     <type> SENDBUF(*), RECVBUF(*)
45     INTEGER SENDCOUNT, SENDTYPE, RECVCOUNTS(*), DISPLS(*), RECVTYPE, ROOT,
46     COMM, REQUEST, IERROR
47
48     MPI_IGATHERW(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNTS, DISPLS,
49                RECVTYPES, ROOT, COMM, REQUEST, IERROR)
50     <type> SENDBUF(*), RECVBUF(*)
51     INTEGER SENDCOUNT, SENDTYPE, RECVCOUNTS(*), DISPLS(*), RECVTYPES(*),
52     ROOT, COMM, REQUEST, IERROR

```

```

MPI_IREDUCE(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, ROOT, COMM, REQUEST,
            IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER COUNT, DATATYPE, OP, ROOT, COMM, REQUEST, IERROR
MPI_IREDUCE_SCATTER(SENDBUF, RECVBUF, RECVCOUNTS, DATATYPE, OP, COMM,
                    REQUEST, IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER RECVCOUNTS(*), DATATYPE, OP, COMM, REQUEST, IERROR
MPI_IREDUCE_SCATTER_BLOCK(SENDBUF, RECVBUF, RECVCOUNT, DATATYPE, OP, COMM,
                           REQUEST, IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER RECVCOUNT, DATATYPE, OP, COMM, REQUEST, IERROR
MPI_ISCAN(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, COMM, REQUEST, IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER COUNT, DATATYPE, OP, COMM, REQUEST, IERROR
MPI_ISCATTER(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, REVCOUNT, RECVTYPE,
             ROOT, COMM, REQUEST, IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER SENDCOUNT, SENDTYPE, REVCOUNT, RECVTYPE, ROOT, COMM, REQUEST,
    IERROR
MPI_ISCATTERV(SENDBUF, SENDCOUNTS, DISPLS, SENDTYPE, RECVBUF, REVCOUNT,
              RECVTYPE, ROOT, COMM, REQUEST, IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER SENDCOUNTS(*), DISPLS(*), SENDTYPE, REVCOUNT, RECVTYPE, ROOT,
    COMM, REQUEST, IERROR
MPI_ISCATTERW(SENDBUF, SENDCOUNTS, DISPLS, SENDTYPES, RECVBUF, REVCOUNT,
              RECVTYPE, ROOT, COMM, REQUEST, IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER DISPLS(*), SENDCOUNTS(*), SENDTYPES, REVCOUNT, RECVTYPE, ROOT,
    COMM, REQUEST, IERROR
MPI_OP_COMMUTATIVE(OP, COMMUTE, IERROR)
    LOGICAL COMMUTE
    INTEGER OP, IERROR
MPI_OP_CREATE( FUNCTION, COMMUTE, OP, IERROR)
    EXTERNAL FUNCTION
    LOGICAL COMMUTE
    INTEGER OP, IERROR
MPI_OP_FREE(OP, IERROR)
    INTEGER OP, IERROR
MPI_REDUCE(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, ROOT, COMM, IERROR)
    <type> SENDBUF(*), RECVBUF(*)
    INTEGER COUNT, DATATYPE, OP, ROOT, COMM, IERROR

```



```

1  MPI_REDUCE_LOCAL(INBUF, INOUBUF, COUNT, DATATYPE, OP, IERROR)
2      <type> INBUF(*), INOUBUF(*)
3      INTEGER COUNT, DATATYPE, OP, IERROR
4
5  MPI_REDUCE_SCATTER(SENDBUF, RECVBUF, RECVCOUNTS, DATATYPE, OP, COMM,
6      IERROR)
7      <type> SENDBUF(*), RECVBUF(*)
8      INTEGER RECVCOUNTS(*), DATATYPE, OP, COMM, IERROR
9
10 MPI_REDUCE_SCATTER_BLOCK(SENDBUF, RECVBUF, RECVCOUNT, DATATYPE, OP, COMM,
11 IERROR)
12 <type> SENDBUF(*), RECVBUF(*)
13 INTEGER RECVCOUNT, DATATYPE, OP, COMM, IERROR
14
15 MPI_SCAN(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, COMM, IERROR)
16 <type> SENDBUF(*), RECVBUF(*)
17 INTEGER COUNT, DATATYPE, OP, COMM, IERROR
18
19 MPI_SCATTER(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, REVCOUNT, RECVTYPE,
20 ROOT, COMM, IERROR)
21 <type> SENDBUF(*), RECVBUF(*)
22 INTEGER SENDCOUNT, SENDTYPE, REVCOUNT, RECVTYPE, ROOT, COMM, IERROR
23
24 MPI_SCATTERV(SENDBUF, SENDCOUNTS, DISPLS, SENDTYPE, RECVBUF, REVCOUNT,
25 RECVTYPE, ROOT, COMM, IERROR)
26 <type> SENDBUF(*), RECVBUF(*)
27 INTEGER SENDCOUNTS(*), DISPLS(*), SENDTYPE, REVCOUNT, RECVTYPE, ROOT,
28 COMM, IERROR
29
30 MPI_SCATTERW(SENDBUF, SENDCOUNTS, DISPLS, SENDTYPES, RECVBUF, REVCOUNT,
31 RECVTYPE, ROOT, COMM, IERROR)
32 <type> SENDBUF(*), RECVBUF(*)
33 INTEGER DISPLS(*), SENDCOUNTS(*), SENDTYPES, REVCOUNT, RECVTYPE, ROOT,
34 COMM, IERROR

```

#### A.3.4 Groups, Contexts, Communicators, and Caching Fortran Bindings

```

35 MPI_COMM_COMPARE(COMM1, COMM2, RESULT, IERROR)
36 INTEGER COMM1, COMM2, RESULT, IERROR
37
38 MPI_COMM_CREATE(COMM, GROUP, NEWCOMM, IERROR)
39 INTEGER COMM, GROUP, NEWCOMM, IERROR
40
41 MPI_COMM_CREATE_KEYVAL(COMM_COPY_ATTR_FN, COMM_DELETE_ATTR_FN, COMM_KEYVAL,
42 EXTRA_STATE, IERROR)
43 EXTERNAL COMM_COPY_ATTR_FN, COMM_DELETE_ATTR_FN
44 INTEGER COMM_KEYVAL, IERROR
45 INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
46
47 MPI_COMM_DELETE_ATTR(COMM, COMM_KEYVAL, IERROR)
48 INTEGER COMM, COMM_KEYVAL, IERROR

```

MPI_COMM_DUP(COMM, NEWCOMM, IERROR)	1
INTEGER COMM, NEWCOMM, IERROR	2
	3
MPI_COMM_DUP_FN(OLDCOMM, COMM_KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,	4
ATTRIBUTE_VAL_OUT, FLAG, IERROR)	5
INTEGER OLDCOMM, COMM_KEYVAL, IERROR	6
INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,	7
ATTRIBUTE_VAL_OUT	8
LOGICAL FLAG	9
	10
MPI_COMM_FREE(COMM, IERROR)	11
INTEGER COMM, IERROR	12
	13
MPI_COMM_FREE_KEYVAL(COMM_KEYVAL, IERROR)	14
INTEGER COMM_KEYVAL, IERROR	15
	16
MPI_COMM_GET_ATTR(COMM, COMM_KEYVAL, ATTRIBUTE_VAL, FLAG, IERROR)	17
INTEGER COMM, COMM_KEYVAL, IERROR	18
INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL	19
LOGICAL FLAG	20
	21
MPI_COMM_GET_NAME(COMM, COMM_NAME, RESULTLEN, IERROR)	22
INTEGER COMM, RESULTLEN, IERROR	23
CHARACTER*(*) COMM_NAME	24
	25
MPI_COMM_GROUP(COMM, GROUP, IERROR)	26
INTEGER COMM, GROUP, IERROR	27
	28
MPI_COMM_NULL_COPY_FN(OLDCOMM, COMM_KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,	29
ATTRIBUTE_VAL_OUT, FLAG, IERROR)	30
INTEGER OLDCOMM, COMM_KEYVAL, IERROR	31
INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,	32
ATTRIBUTE_VAL_OUT	33
LOGICAL FLAG	34
	35
MPI_COMM_NULL_DELETE_FN(COMM, COMM_KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE,	36
IERROR)	37
INTEGER COMM, COMM_KEYVAL, IERROR	38
INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE	39
	40
MPI_COMM_RANK(COMM, RANK, IERROR)	41
INTEGER COMM, RANK, IERROR	42
	43
MPI_COMM_REMOTE_GROUP(COMM, GROUP, IERROR)	44
INTEGER COMM, GROUP, IERROR	45
	46
MPI_COMM_REMOTE_SIZE(COMM, SIZE, IERROR)	47
INTEGER COMM, SIZE, IERROR	48
MPI_COMM_SET_ATTR(COMM, COMM_KEYVAL, ATTRIBUTE_VAL, IERROR)	
INTEGER COMM, COMM_KEYVAL, IERROR	
INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL	
MPI_COMM_SET_NAME(COMM, COMM_NAME, IERROR)	

```

1      INTEGER COMM, IERROR
2      CHARACTER*(*) COMM_NAME
3
4      MPI_COMM_SIZE(COMM, SIZE, IERROR)
5          INTEGER COMM, SIZE, IERROR
6
7      MPI_COMM_SPLIT(COMM, COLOR, KEY, NEWCOMM, IERROR)
8          INTEGER COMM, COLOR, KEY, NEWCOMM, IERROR
9
10     MPI_COMM_TEST_INTER(COMM, FLAG, IERROR)
11         INTEGER COMM, IERROR
12         LOGICAL FLAG
13
14     MPI_GROUP_COMPARE(GROUP1, GROUP2, RESULT, IERROR)
15         INTEGER GROUP1, GROUP2, RESULT, IERROR
16
17     MPI_GROUP_DIFFERENCE(GROUP1, GROUP2, NEWGROUP, IERROR)
18         INTEGER GROUP1, GROUP2, NEWGROUP, IERROR
19
20     MPI_GROUP_EXCL(GROUP, N, RANKS, NEWGROUP, IERROR)
21         INTEGER GROUP, N, RANKS(*), NEWGROUP, IERROR
22
23     MPI_GROUP_FREE(GROUP, IERROR)
24         INTEGER GROUP, IERROR
25
26     MPI_GROUP_INCL(GROUP, N, RANKS, NEWGROUP, IERROR)
27         INTEGER GROUP, N, RANKS(*), NEWGROUP, IERROR
28
29     MPI_GROUP_INTERSECTION(GROUP1, GROUP2, NEWGROUP, IERROR)
30         INTEGER GROUP1, GROUP2, NEWGROUP, IERROR
31
32     MPI_GROUP_RANGE_EXCL(GROUP, N, RANGES, NEWGROUP, IERROR)
33         INTEGER GROUP, N, RANGES(3,*), NEWGROUP, IERROR
34
35     MPI_GROUP_RANGE_INCL(GROUP, N, RANGES, NEWGROUP, IERROR)
36         INTEGER GROUP, N, RANGES(3,*), NEWGROUP, IERROR
37
38     MPI_GROUP_RANK(GROUP, RANK, IERROR)
39         INTEGER GROUP, RANK, IERROR
40
41     MPI_GROUP_SIZE(GROUP, SIZE, IERROR)
42         INTEGER GROUP, SIZE, IERROR
43
44     MPI_GROUP_TRANSLATE_RANKS(GROUP1, N, RANKS1, GROUP2, RANKS2, IERROR)
45         INTEGER GROUP1, N, RANKS1(*), GROUP2, RANKS2(*), IERROR
46
47     MPI_GROUP_UNION(GROUP1, GROUP2, NEWGROUP, IERROR)
48         INTEGER GROUP1, GROUP2, NEWGROUP, IERROR
49
50     MPI_INTERCOMM_CREATE(LOCAL_COMM, LOCAL_LEADER, PEER_COMM, REMOTE_LEADER,
51         TAG, NEWINTERCOMM, IERROR)
52         INTEGER LOCAL_COMM, LOCAL_LEADER, PEER_COMM, REMOTE_LEADER, TAG,
53         NEWINTERCOMM, IERROR
54
55     MPI_INTERCOMM_MERGE(INTERCOMM, HIGH, INTRACOMM, IERROR)

```

```

    INTEGER INTERCOMM, INTRACOMM, IERROR
    LOGICAL HIGH
MPI_TYPE_CREATE_KEYVAL(TYPE_COPY_ATTR_FN, TYPE_DELETE_ATTR_FN, TYPE_KEYVAL,
    EXTRA_STATE, IERROR)
    EXTERNAL TYPE_COPY_ATTR_FN, TYPE_DELETE_ATTR_FN
    INTEGER TYPE_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
MPI_TYPE_DELETE_ATTR(TYPE, TYPE_KEYVAL, IERROR)
    INTEGER TYPE, TYPE_KEYVAL, IERROR
MPI_TYPE_DUP_FN(OLDTYPE, TYPE_KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,
    ATTRIBUTE_VAL_OUT, FLAG, IERROR)
    INTEGER OLDTYPE, TYPE_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,
    ATTRIBUTE_VAL_OUT
    LOGICAL FLAG
MPI_TYPE_FREE_KEYVAL(TYPE_KEYVAL, IERROR)
    INTEGER TYPE_KEYVAL, IERROR
MPI_TYPE_GET_ATTR(TYPE, TYPE_KEYVAL, ATTRIBUTE_VAL, FLAG, IERROR)
    INTEGER TYPE, TYPE_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL
    LOGICAL FLAG
MPI_TYPE_GET_NAME(TYPE, TYPE_NAME, RESULTLEN, IERROR)
    INTEGER TYPE, RESULTLEN, IERROR
    CHARACTER*(*) TYPE_NAME
MPI_TYPE_NULL_COPY_FN(OLDTYPE, TYPE_KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,
    ATTRIBUTE_VAL_OUT, FLAG, IERROR)
    INTEGER OLDTYPE, TYPE_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,
    ATTRIBUTE_VAL_OUT
    LOGICAL FLAG
MPI_TYPE_NULL_DELETE_FN(TYPE, TYPE_KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE,
    IERROR)
    INTEGER TYPE, TYPE_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE
MPI_TYPE_SET_ATTR(TYPE, TYPE_KEYVAL, ATTRIBUTE_VAL, IERROR)
    INTEGER TYPE, TYPE_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL
MPI_TYPE_SET_NAME(TYPE, TYPE_NAME, IERROR)
    INTEGER TYPE, IERROR
    CHARACTER*(*) TYPE_NAME
MPI_WIN_CREATE_KEYVAL(WIN_COPY_ATTR_FN, WIN_DELETE_ATTR_FN, WIN_KEYVAL,
    EXTRA_STATE, IERROR)

```

```

1      EXTERNAL WIN_COPY_ATTR_FN, WIN_DELETE_ATTR_FN
2      INTEGER WIN_KEYVAL, IERROR
3      INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
4
5      MPI_WIN_DELETE_ATTR(WIN, WIN_KEYVAL, IERROR)
6      INTEGER WIN, WIN_KEYVAL, IERROR
7
8      MPI_WIN_DUP_FN(OLDWIN, WIN_KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,
9                    ATTRIBUTE_VAL_OUT, FLAG, IERROR)
10     INTEGER OLDWIN, WIN_KEYVAL, IERROR
11     INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,
12     ATTRIBUTE_VAL_OUT
13     LOGICAL FLAG
14
15     MPI_WIN_FREE_KEYVAL(WIN_KEYVAL, IERROR)
16     INTEGER WIN_KEYVAL, IERROR
17
18     MPI_WIN_GET_ATTR(WIN, WIN_KEYVAL, ATTRIBUTE_VAL, FLAG, IERROR)
19     INTEGER WIN, WIN_KEYVAL, IERROR
20     INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL
21     LOGICAL FLAG
22
23     MPI_WIN_GET_NAME(WIN, WIN_NAME, RESULTLEN, IERROR)
24     INTEGER WIN, RESULTLEN, IERROR
25     CHARACTER*(*) WIN_NAME
26
27     MPI_WIN_NULL_COPY_FN(OLDWIN, WIN_KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,
28                          ATTRIBUTE_VAL_OUT, FLAG, IERROR)
29     INTEGER OLDWIN, WIN_KEYVAL, IERROR
30     INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,
31     ATTRIBUTE_VAL_OUT
32     LOGICAL FLAG
33
34     MPI_WIN_NULL_DELETE_FN(WIN, WIN_KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE, IERROR)
35     INTEGER WIN, WIN_KEYVAL, IERROR
36     INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE
37
38     MPI_WIN_SET_ATTR(WIN, WIN_KEYVAL, ATTRIBUTE_VAL, IERROR)
39     INTEGER WIN, WIN_KEYVAL, IERROR
40     INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL
41
42     MPI_WIN_SET_NAME(WIN, WIN_NAME, IERROR)
43     INTEGER WIN, IERROR
44     CHARACTER*(*) WIN_NAME

```

### A.3.5 Process Topologies Fortran Bindings

```

44     MPI_CARTDIM_GET(COMM, NDIMS, IERROR)
45     INTEGER COMM, NDIMS, IERROR
46
47     MPI_CART_COORDS(COMM, RANK, MAXDIMS, COORDS, IERROR)
48     INTEGER COMM, RANK, MAXDIMS, COORDS(*), IERROR

```

```

MPI_CART_CREATE(COMM_OLD, NDIMS, DIMS, PERIODS, REORDER, COMM_CART, IERROR) 1
    INTEGER COMM_OLD, NDIMS, DIMS(*), COMM_CART, IERROR 2
    LOGICAL PERIODS(*), REORDER 3
4
MPI_CART_GET(COMM, MAXDIMS, DIMS, PERIODS, COORDS, IERROR) 5
    INTEGER COMM, MAXDIMS, DIMS(*), COORDS(*), IERROR 6
    LOGICAL PERIODS(*) 7
8
MPI_CART_MAP(COMM, NDIMS, DIMS, PERIODS, NEWRANK, IERROR) 9
    INTEGER COMM, NDIMS, DIMS(*), NEWRANK, IERROR 10
    LOGICAL PERIODS(*) 11
12
MPI_CART_RANK(COMM, COORDS, RANK, IERROR) 13
    INTEGER COMM, COORDS(*), RANK, IERROR 14
15
MPI_CART_SHIFT(COMM, DIRECTION, DISP, RANK_SOURCE, RANK_DEST, IERROR) 16
    INTEGER COMM, DIRECTION, DISP, RANK_SOURCE, RANK_DEST, IERROR 17
18
MPI_CART_SUB(COMM, REMAIN_DIMS, NEWCOMM, IERROR) 19
    INTEGER COMM, NEWCOMM, IERROR 20
    LOGICAL REMAIN_DIMS(*) 21
22
MPI_DIMS_CREATE(NNODES, NDIMS, DIMS, IERROR) 23
    INTEGER NNODES, NDIMS, DIMS(*), IERROR 24
25
MPI_DIST_GRAPH_CREATE(COMM_OLD, N, SOURCES, DEGREES, DESTINATIONS, WEIGHTS, 26
    INFO, REORDER, COMM_DIST_GRAPH, IERROR) 27
    INTEGER COMM_OLD, N, SOURCES(*), DEGREES(*), DESTINATIONS(*), 28
    WEIGHTS(*), INFO, COMM_DIST_GRAPH, IERROR 29
    LOGICAL REORDER 30
31
MPI_DIST_GRAPH_CREATE_ADJACENT(COMM_OLD, INDEGREE, SOURCES, SOURCEWEIGHTS, 32
    OUTDEGREE, DESTINATIONS, DESTWEIGHTS, INFO, REORDER, 33
    COMM_DIST_GRAPH, IERROR) 34
    INTEGER COMM_OLD, INDEGREE, SOURCES(*), SOURCEWEIGHTS(*), OUTDEGREE, 35
    DESTINATIONS(*), DESTWEIGHTS(*), INFO, COMM_DIST_GRAPH, IERROR 36
    LOGICAL REORDER 37
38
MPI_DIST_GRAPH_NEIGHBORS(COMM, MAXINDEGREE, SOURCES, SOURCEWEIGHTS, 39
    MAXOUTDEGREE, DESTINATIONS, DESTWEIGHTS, IERROR) 40
    INTEGER COMM, MAXINDEGREE, SOURCES(*), SOURCEWEIGHTS(*), MAXOUTDEGREE, 41
    DESTINATIONS(*), DESTWEIGHTS(*), IERROR 42
43
MPI_DIST_GRAPH_NEIGHBORS_COUNT(COMM, INDEGREE, OUTDEGREE, WEIGHTED, IERROR) 44
    INTEGER COMM, INDEGREE, OUTDEGREE, IERROR 45
    LOGICAL WEIGHTED 46
47
MPI_GRAPHDIMS_GET(COMM, NNODES, NEDGES, IERROR) 48
    INTEGER COMM, NNODES, NEDGES, IERROR

```

```

1      LOGICAL REORDER
2
3      MPI_GRAPH_GET(COMM, MAXINDEX, MAXEDGES, INDEX, EDGES, IERROR)
4          INTEGER COMM, MAXINDEX, MAXEDGES, INDEX(*), EDGES(*), IERROR
5
6      MPI_GRAPH_MAP(COMM, NNODES, INDEX, EDGES, NEWRANK, IERROR)
7          INTEGER COMM, NNODES, INDEX(*), EDGES(*), NEWRANK, IERROR
8
9      MPI_GRAPH_NEIGHBORS(COMM, RANK, MAXNEIGHBORS, NEIGHBORS, IERROR)
10         INTEGER COMM, RANK, MAXNEIGHBORS, NEIGHBORS(*), IERROR
11
12     MPI_GRAPH_NEIGHBORS_COUNT(COMM, RANK, NNEIGHBORS, IERROR)
13         INTEGER COMM, RANK, NNEIGHBORS, IERROR
14
15     MPI_TOPO_TEST(COMM, STATUS, IERROR)
16         INTEGER COMM, STATUS, IERROR

```

### A.3.6 MPI Environmental Management Fortran Bindings

```

17
18     DOUBLE PRECISION MPI_WTICK()
19
20     DOUBLE PRECISION MPI_WTIME()
21
22     MPI_ABORT(COMM, ERRORCODE, IERROR)
23         INTEGER COMM, ERRORCODE, IERROR
24
25     MPI_ADD_ERROR_CLASS(ERRORCLASS, IERROR)
26         INTEGER ERRORCLASS, IERROR
27
28     MPI_ADD_ERROR_CODE(ERRORCLASS, ERRORCODE, IERROR)
29         INTEGER ERRORCLASS, ERRORCODE, IERROR
30
31     MPI_ADD_ERROR_STRING(ERRORCODE, STRING, IERROR)
32         INTEGER ERRORCODE, IERROR
33         CHARACTER*(*) STRING
34
35     MPI_ALLOC_MEM(SIZE, INFO, BASEPTR, IERROR)
36         INTEGER INFO, IERROR
37         INTEGER(KIND=MPI_ADDRESS_KIND) SIZE, BASEPTR
38
39     MPI_COMM_CALL_ERRHANDLER(COMM, ERRORCODE, IERROR)
40         INTEGER COMM, ERRORCODE, IERROR
41
42     MPI_COMM_CREATE_ERRHANDLER(FUNCTION, ERRHANDLER, IERROR)
43         EXTERNAL FUNCTION
44         INTEGER ERRHANDLER, IERROR
45
46     MPI_COMM_GET_ERRHANDLER(COMM, ERRHANDLER, IERROR)
47         INTEGER COMM, ERRHANDLER, IERROR
48
49     MPI_COMM_SET_ERRHANDLER(COMM, ERRHANDLER, IERROR)
50         INTEGER COMM, ERRHANDLER, IERROR
51
52     MPI_ERRHANDLER_FREE(ERRHANDLER, IERROR)
53         INTEGER ERRHANDLER, IERROR

```

MPI_ERROR_CLASS(ERRORCODE, ERRORCLASS, IERROR)	1
INTEGER ERRORCODE, ERRORCLASS, IERROR	2
	3
MPI_ERROR_STRING(ERRORCODE, STRING, RESULTLEN, IERROR)	4
INTEGER ERRORCODE, RESULTLEN, IERROR	5
CHARACTER*(*) STRING	6
	7
MPI_FILE_CALL_ERRHANDLER(FH, ERRORCODE, IERROR)	8
INTEGER FH, ERRORCODE, IERROR	9
	10
MPI_FILE_CREATE_ERRHANDLER(FUNCTION, ERRHANDLER, IERROR)	11
EXTERNAL FUNCTION	12
INTEGER ERRHANDLER, IERROR	13
	14
MPI_FILE_GET_ERRHANDLER(FILE, ERRHANDLER, IERROR)	15
INTEGER FILE, ERRHANDLER, IERROR	16
	17
MPI_FILE_SET_ERRHANDLER(FILE, ERRHANDLER, IERROR)	18
INTEGER FILE, ERRHANDLER, IERROR	19
	20
MPI_FINALIZE(IERROR)	21
INTEGER IERROR	22
	23
MPI_FINALIZED(FLAG, IERROR)	24
LOGICAL FLAG	25
INTEGER IERROR	26
	27
MPI_FREE_MEM(BASE, IERROR)	28
<type> BASE(*)	29
INTEGER IERROR	30
	31
MPI_GET_PROCESSOR_NAME( NAME, RESULTLEN, IERROR)	32
CHARACTER*(*) NAME	33
INTEGER RESULTLEN, IERROR	34
	35
MPI_GET_VERSION(VERSION, SUBVERSION, IERROR)	36
INTEGER VERSION, SUBVERSION, IERROR	37
	38
MPI_INIT(IERROR)	39
INTEGER IERROR	40
	41
MPI_INITIALIZED(FLAG, IERROR)	42
LOGICAL FLAG	43
INTEGER IERROR	44
	45
MPI_WIN_CALL_ERRHANDLER(WIN, ERRORCODE, IERROR)	46
INTEGER WIN, ERRORCODE, IERROR	47
	48
MPI_WIN_CREATE_ERRHANDLER(FUNCTION, ERRHANDLER, IERROR)	
EXTERNAL FUNCTION	
INTEGER ERRHANDLER, IERROR	
MPI_WIN_GET_ERRHANDLER(WIN, ERRHANDLER, IERROR)	
INTEGER WIN, ERRHANDLER, IERROR	



```

1  MPI_WIN_SET_ERRHANDLER(WIN, ERRHANDLER, IERROR)
2      INTEGER WIN, ERRHANDLER, IERROR
3
4

```

### A.3.7 The Info Object Fortran Bindings

```

6  MPI_INFO_CREATE(INFO, IERROR)
7      INTEGER INFO, IERROR
8
9  MPI_INFO_DELETE(INFO, KEY, IERROR)
10     INTEGER INFO, IERROR
11     CHARACTER*(*) KEY
12
13 MPI_INFO_DUP(INFO, NEWINFO, IERROR)
14     INTEGER INFO, NEWINFO, IERROR
15
16 MPI_INFO_FREE(INFO, IERROR)
17     INTEGER INFO, IERROR
18
19 MPI_INFO_GET(INFO, KEY, VALUELEN, VALUE, FLAG, IERROR)
20     INTEGER INFO, VALUELEN, IERROR
21     CHARACTER*(*) KEY, VALUE
22     LOGICAL FLAG
23
24 MPI_INFO_GET_NKEYS(INFO, NKEYS, IERROR)
25     INTEGER INFO, NKEYS, IERROR
26
27 MPI_INFO_GET_NTHKEY(INFO, N, KEY, IERROR)
28     INTEGER INFO, N, IERROR
29     CHARACTER*(*) KEY
30
31 MPI_INFO_GET_VALUELEN(INFO, KEY, VALUELEN, FLAG, IERROR)
32     INTEGER INFO, VALUELEN, IERROR
33     LOGICAL FLAG
34     CHARACTER*(*) KEY
35
36 MPI_INFO_SET(INFO, KEY, VALUE, IERROR)
37     INTEGER INFO, IERROR
38     CHARACTER*(*) KEY, VALUE
39
40

```

### A.3.8 Process Creation and Management Fortran Bindings

```

37 MPI_CLOSE_PORT(PORT_NAME, IERROR)
38     CHARACTER*(*) PORT_NAME
39     INTEGER IERROR
40
41
42 MPI_COMM_ACCEPT(PORT_NAME, INFO, ROOT, COMM, NEWCOMM, IERROR)
43     CHARACTER*(*) PORT_NAME
44     INTEGER INFO, ROOT, COMM, NEWCOMM, IERROR
45
46 MPI_COMM_CONNECT(PORT_NAME, INFO, ROOT, COMM, NEWCOMM, IERROR)
47     CHARACTER*(*) PORT_NAME
48     INTEGER INFO, ROOT, COMM, NEWCOMM, IERROR

```

```

MPI_COMM_DISCONNECT(COMM, IERROR) 1
    INTEGER COMM, IERROR 2
3
MPI_COMM_GET_PARENT(PARENT, IERROR) 4
    INTEGER PARENT, IERROR 5
6
MPI_COMM_JOIN(FD, INTERCOMM, IERROR) 7
    INTEGER FD, INTERCOMM, IERROR 8
9
MPI_COMM_SPAWN(COMMAND, ARGV, MAXPROCS, INFO, ROOT, COMM, INTERCOMM, 9
    ARRAY_OF_ERRCODES, IERROR) 10
    CHARACTER*(*) COMMAND, ARGV(*) 11
    INTEGER INFO, MAXPROCS, ROOT, COMM, INTERCOMM, ARRAY_OF_ERRCODES(*), 12
    IERROR 13
14
MPI_COMM_SPAWN_MULTIPLE(COUNT, ARRAY_OF_COMMANDS, ARRAY_OF_ARGV, 15
    ARRAY_OF_MAXPROCS, ARRAY_OF_INFO, ROOT, COMM, INTERCOMM, 16
    ARRAY_OF_ERRCODES, IERROR) 17
    INTEGER COUNT, ARRAY_OF_INFO(*), ARRAY_OF_MAXPROCS(*), ROOT, COMM, 18
    INTERCOMM, ARRAY_OF_ERRCODES(*), IERROR 19
    CHARACTER*(*) ARRAY_OF_COMMANDS(*), ARRAY_OF_ARGV(COUNT, *) 20
21
MPI_LOOKUP_NAME(SERVICE_NAME, INFO, PORT_NAME, IERROR) 21
    CHARACTER*(*) SERVICE_NAME, PORT_NAME 22
    INTEGER INFO, IERROR 23
24
MPI_OPEN_PORT(INFO, PORT_NAME, IERROR) 24
    CHARACTER*(*) PORT_NAME 25
    INTEGER INFO, IERROR 26
27
MPI_PUBLISH_NAME(SERVICE_NAME, INFO, PORT_NAME, IERROR) 28
    INTEGER INFO, IERROR 29
    CHARACTER*(*) SERVICE_NAME, PORT_NAME 30
31
MPI_UNPUBLISH_NAME(SERVICE_NAME, INFO, PORT_NAME, IERROR) 31
    INTEGER INFO, IERROR 32
    CHARACTER*(*) SERVICE_NAME, PORT_NAME 33
34
35

```

### A.3.9 One-Sided Communications Fortran Bindings

```

MPI_ACCUMULATE(ORIGIN_ADDR, ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, 37
    TARGET_DISP, TARGET_COUNT, TARGET_DATATYPE, OP, WIN, IERROR) 38
    <type> ORIGIN_ADDR(*) 39
    INTEGER(KIND=MPI_ADDRESS_KIND) TARGET_DISP 40
    INTEGER ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, TARGET_COUNT, 41
    TARGET_DATATYPE, OP, WIN, IERROR 42
43
MPI_GET(ORIGIN_ADDR, ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, 44
    TARGET_DISP, TARGET_COUNT, TARGET_DATATYPE, WIN, IERROR) 45
    <type> ORIGIN_ADDR(*) 46
    INTEGER(KIND=MPI_ADDRESS_KIND) TARGET_DISP 47
48

```

```

1      INTEGER ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, TARGET_COUNT,
2      TARGET_DATATYPE, WIN, IERROR
3
4      MPI_PUT(ORIGIN_ADDR, ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK,
5              TARGET_DISP, TARGET_COUNT, TARGET_DATATYPE, WIN, IERROR)
6      <type> ORIGIN_ADDR(*)
7      INTEGER(KIND=MPI_ADDRESS_KIND) TARGET_DISP
8      INTEGER ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, TARGET_COUNT,
9      TARGET_DATATYPE, WIN, IERROR
10
11     MPI_WIN_COMPLETE(WIN, IERROR)
12     INTEGER WIN, IERROR
13
14     MPI_WIN_CREATE(BASE, SIZE, DISP_UNIT, INFO, COMM, WIN, IERROR)
15     <type> BASE(*)
16     INTEGER(KIND=MPI_ADDRESS_KIND) SIZE
17     INTEGER DISP_UNIT, INFO, COMM, WIN, IERROR
18
19     MPI_WIN_FENCE(ASSERT, WIN, IERROR)
20     INTEGER ASSERT, WIN, IERROR
21
22     MPI_WIN_FREE(WIN, IERROR)
23     INTEGER WIN, IERROR
24
25     MPI_WIN_GET_GROUP(WIN, GROUP, IERROR)
26     INTEGER WIN, GROUP, IERROR
27
28     MPI_WIN_LOCK(LOCK_TYPE, RANK, ASSERT, WIN, IERROR)
29     INTEGER LOCK_TYPE, RANK, ASSERT, WIN, IERROR
30
31     MPI_WIN_POST(GROUP, ASSERT, WIN, IERROR)
32     INTEGER GROUP, ASSERT, WIN, IERROR
33
34     MPI_WIN_START(GROUP, ASSERT, WIN, IERROR)
35     INTEGER GROUP, ASSERT, WIN, IERROR
36
37     MPI_WIN_TEST(WIN, FLAG, IERROR)
38     INTEGER WIN, IERROR
39     LOGICAL FLAG
40
41     MPI_WIN_UNLOCK(RANK, WIN, IERROR)
42     INTEGER RANK, WIN, IERROR
43
44     MPI_WIN_WAIT(WIN, IERROR)
45     INTEGER WIN, IERROR

```

#### A.3.10 External Interfaces Fortran Bindings

```

44     MPI_GREQUEST_COMPLETE(REQUEST, IERROR)
45     INTEGER REQUEST, IERROR
46
47     MPI_GREQUEST_START(QUERY_FN, FREE_FN, CANCEL_FN, EXTRA_STATE, REQUEST,
48                       IERROR)

```

INTEGER REQUEST, IERROR	1
EXTERNAL QUERY_FN, FREE_FN, CANCEL_FN	2
INTEGER (KIND=MPI_ADDRESS_KIND) EXTRA_STATE	3
	4
MPI_INIT_THREAD(REQUIRED, PROVIDED, IERROR)	5
INTEGER REQUIRED, PROVIDED, IERROR	6
	7
MPI_IS_THREAD_MAIN(FLAG, IERROR)	8
LOGICAL FLAG	9
INTEGER IERROR	10
	11
MPI_QUERY_THREAD(PROVIDED, IERROR)	12
INTEGER PROVIDED, IERROR	13
	14
MPI_STATUS_SET_CANCELLED(STATUS, FLAG, IERROR)	15
INTEGER STATUS(MPI_STATUS_SIZE), IERROR	16
LOGICAL FLAG	17
	18
MPI_STATUS_SET_ELEMENTS(STATUS, DATATYPE, COUNT, IERROR)	19
INTEGER STATUS(MPI_STATUS_SIZE), DATATYPE, COUNT, IERROR	20
	21
MPI_STATUS_SET_ELEMENTS_X(STATUS, DATATYPE, COUNT, IERROR)	22
INTEGER STATUS(MPI_STATUS_SIZE), DATATYPE, IERROR	23
INTEGER (KIND=MPI_COUNT_KIND) COUNT	24
	25
A.3.11 I/O Fortran Bindings	26
	27
MPI_FILE_CLOSE(FH, IERROR)	28
INTEGER FH, IERROR	29
	30
MPI_FILE_DELETE(FILENAME, INFO, IERROR)	31
CHARACTER*(*) FILENAME	32
INTEGER INFO, IERROR	33
	34
MPI_FILE_GET_AMODE(FH, AMODE, IERROR)	35
INTEGER FH, AMODE, IERROR	36
	37
MPI_FILE_GET_ATOMICITY(FH, FLAG, IERROR)	38
INTEGER FH, IERROR	39
LOGICAL FLAG	40
	41
MPI_FILE_GET_BYTE_OFFSET(FH, OFFSET, DISP, IERROR)	42
INTEGER FH, IERROR	43
INTEGER(KIND=MPI_OFFSET_KIND) OFFSET, DISP	44
	45
MPI_FILE_GET_GROUP(FH, GROUP, IERROR)	46
INTEGER FH, GROUP, IERROR	47
	48
MPI_FILE_GET_INFO(FH, INFO_USED, IERROR)	
INTEGER FH, INFO_USED, IERROR	
MPI_FILE_GET_POSITION(FH, OFFSET, IERROR)	
INTEGER FH, IERROR	

```

1      INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
2
3      MPI_FILE_GET_POSITION_SHARED(FH, OFFSET, IERROR)
4          INTEGER FH, IERROR
5          INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
6
7      MPI_FILE_GET_SIZE(FH, SIZE, IERROR)
8          INTEGER FH, IERROR
9          INTEGER(KIND=MPI_OFFSET_KIND) SIZE
10
11     MPI_FILE_GET_TYPE_EXTENT(FH, DATATYPE, EXTENT, IERROR)
12         INTEGER FH, DATATYPE, IERROR
13         INTEGER(KIND=MPI_ADDRESS_KIND) EXTENT
14
15     MPI_FILE_GET_VIEW(FH, DISP, ETYPE, FILETYPE, DATAREP, IERROR)
16         INTEGER FH, ETYPE, FILETYPE, IERROR
17         CHARACTER*(*) DATAREP
18         INTEGER(KIND=MPI_OFFSET_KIND) DISP
19
20     MPI_FILE_IREAD(FH, BUF, COUNT, DATATYPE, REQUEST, IERROR)
21         <type> BUF(*)
22         INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR
23
24     MPI_FILE_IREAD_AT(FH, OFFSET, BUF, COUNT, DATATYPE, REQUEST, IERROR)
25         <type> BUF(*)
26         INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR
27         INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
28
29     MPI_FILE_IREAD_SHARED(FH, BUF, COUNT, DATATYPE, REQUEST, IERROR)
30         <type> BUF(*)
31         INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR
32
33     MPI_FILE_IWRITE(FH, BUF, COUNT, DATATYPE, REQUEST, IERROR)
34         <type> BUF(*)
35         INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR
36
37     MPI_FILE_IWRITE_AT(FH, OFFSET, BUF, COUNT, DATATYPE, REQUEST, IERROR)
38         <type> BUF(*)
39         INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR
40         INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
41
42     MPI_FILE_IWRITE_SHARED(FH, BUF, COUNT, DATATYPE, REQUEST, IERROR)
43         <type> BUF(*)
44         INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR
45
46     MPI_FILE_OPEN(COMM, FILENAME, AMODE, INFO, FH, IERROR)
47         CHARACTER*(*) FILENAME
48         INTEGER COMM, AMODE, INFO, FH, IERROR
49
50     MPI_FILE_PREALLOCATE(FH, SIZE, IERROR)
51         INTEGER FH, IERROR
52         INTEGER(KIND=MPI_OFFSET_KIND) SIZE
53
54     MPI_FILE_READ(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)

```

```

    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
MPI_FILE_READ_ALL(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
MPI_FILE_READ_ALL_BEGIN(FH, BUF, COUNT, DATATYPE, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, IERROR
MPI_FILE_READ_ALL_END(FH, BUF, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR
MPI_FILE_READ_AT(FH, OFFSET, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
MPI_FILE_READ_AT_ALL(FH, OFFSET, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
MPI_FILE_READ_AT_ALL_BEGIN(FH, OFFSET, BUF, COUNT, DATATYPE, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
MPI_FILE_READ_AT_ALL_END(FH, BUF, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR
MPI_FILE_READ_ORDERED(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
MPI_FILE_READ_ORDERED_BEGIN(FH, BUF, COUNT, DATATYPE, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, IERROR
MPI_FILE_READ_ORDERED_END(FH, BUF, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR
MPI_FILE_READ_SHARED(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
MPI_FILE_SEEK(FH, OFFSET, WHENCE, IERROR)
    INTEGER FH, WHENCE, IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET

```

```

1  MPI_FILE_SEEK_SHARED(FH, OFFSET, WHENCE, IERROR)
2      INTEGER FH, WHENCE, IERROR
3      INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
4
5  MPI_FILE_SET_ATOMICITY(FH, FLAG, IERROR)
6      INTEGER FH, IERROR
7      LOGICAL FLAG
8
9  MPI_FILE_SET_INFO(FH, INFO, IERROR)
10     INTEGER FH, INFO, IERROR
11
12 MPI_FILE_SET_SIZE(FH, SIZE, IERROR)
13     INTEGER FH, IERROR
14     INTEGER(KIND=MPI_OFFSET_KIND) SIZE
15
16 MPI_FILE_SET_VIEW(FH, DISP, ETYPE, FILETYPE, DATAREP, INFO, IERROR)
17     INTEGER FH, ETYPE, FILETYPE, INFO, IERROR
18     CHARACTER*(*) DATAREP
19     INTEGER(KIND=MPI_OFFSET_KIND) DISP
20
21 MPI_FILE_SYNC(FH, IERROR)
22     INTEGER FH, IERROR
23
24 MPI_FILE_WRITE(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
25     <type> BUF(*)
26     INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
27
28 MPI_FILE_WRITE_ALL(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
29     <type> BUF(*)
30     INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
31
32 MPI_FILE_WRITE_ALL_BEGIN(FH, BUF, COUNT, DATATYPE, IERROR)
33     <type> BUF(*)
34     INTEGER FH, COUNT, DATATYPE, IERROR
35
36 MPI_FILE_WRITE_ALL_END(FH, BUF, STATUS, IERROR)
37     <type> BUF(*)
38     INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR
39
40 MPI_FILE_WRITE_AT(FH, OFFSET, BUF, COUNT, DATATYPE, STATUS, IERROR)
41     <type> BUF(*)
42     INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
43     INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
44
45 MPI_FILE_WRITE_AT_ALL(FH, OFFSET, BUF, COUNT, DATATYPE, STATUS, IERROR)
46     <type> BUF(*)
47     INTEGER FH, COUNT, DATATYPE, IERROR
48     INTEGER(KIND=MPI_OFFSET_KIND) OFFSET

```

```

MPI_FILE_WRITE_AT_ALL_END(FH, BUF, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR
MPI_FILE_WRITE_ORDERED(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
MPI_FILE_WRITE_ORDERED_BEGIN(FH, BUF, COUNT, DATATYPE, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, IERROR
MPI_FILE_WRITE_ORDERED_END(FH, BUF, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR
MPI_FILE_WRITE_SHARED(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
MPI_REGISTER_DATAREP(DATAREP, READ_CONVERSION_FN, WRITE_CONVERSION_FN,
    DTYPE_FILE_EXTENT_FN, EXTRA_STATE, IERROR)
    CHARACTER*(*) DATAREP
    EXTERNAL READ_CONVERSION_FN, WRITE_CONVERSION_FN, DTYPE_FILE_EXTENT_FN
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
    INTEGER IERROR

```

#### A.3.12 Language Bindings Fortran Bindings

```

MPI_SIZEOF(X, SIZE, IERROR)
    <type> X
    INTEGER SIZE, IERROR
MPI_TYPE_CREATE_F90_COMPLEX(P, R, NEWTYPE, IERROR)
    INTEGER P, R, NEWTYPE, IERROR
MPI_TYPE_CREATE_F90_INTEGER(R, NEWTYPE, IERROR)
    INTEGER R, NEWTYPE, IERROR
MPI_TYPE_CREATE_F90_REAL(P, R, NEWTYPE, IERROR)
    INTEGER P, R, NEWTYPE, IERROR
MPI_TYPE_MATCH_SIZE(TYPECLASS, SIZE, TYPE, IERROR)
    INTEGER TYPECLASS, SIZE, TYPE, IERROR

```

#### A.3.13 Profiling Interface Fortran Bindings

```

MPI_PCONTROL(LEVEL)
    INTEGER LEVEL

```



## A.3.14 Deprecated Fortran Bindings

```

MPI_ADDRESS(LOCATION, ADDRESS, IERROR)
    <type> LOCATION(*)
    INTEGER ADDRESS, IERROR

MPI_ATTR_DELETE(COMM, KEYVAL, IERROR)
    INTEGER COMM, KEYVAL, IERROR

MPI_ATTR_GET(COMM, KEYVAL, ATTRIBUTE_VAL, FLAG, IERROR)
    INTEGER COMM, KEYVAL, ATTRIBUTE_VAL, IERROR
    LOGICAL FLAG

MPI_ATTR_PUT(COMM, KEYVAL, ATTRIBUTE_VAL, IERROR)
    INTEGER COMM, KEYVAL, ATTRIBUTE_VAL, IERROR

MPI_DUP_FN(OLDCOMM, KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,
           ATTRIBUTE_VAL_OUT, FLAG, IERR)
    INTEGER OLDCOMM, KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,
    ATTRIBUTE_VAL_OUT, IERR
    LOGICAL FLAG

MPI_ERRHANDLER_CREATE(FUNCTION, ERRHANDLER, IERROR)
    EXTERNAL FUNCTION
    INTEGER ERRHANDLER, IERROR

MPI_ERRHANDLER_GET(COMM, ERRHANDLER, IERROR)
    INTEGER COMM, ERRHANDLER, IERROR

MPI_ERRHANDLER_SET(COMM, ERRHANDLER, IERROR)
    INTEGER COMM, ERRHANDLER, IERROR

MPI_KEYVAL_CREATE(COPY_FN, DELETE_FN, KEYVAL, EXTRA_STATE, IERROR)
    EXTERNAL COPY_FN, DELETE_FN
    INTEGER KEYVAL, EXTRA_STATE, IERROR

MPI_KEYVAL_FREE(KEYVAL, IERROR)
    INTEGER KEYVAL, IERROR

MPI_NULL_COPY_FN(OLDCOMM, KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,
                ATTRIBUTE_VAL_OUT, FLAG, IERR)
    INTEGER OLDCOMM, KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,
    ATTRIBUTE_VAL_OUT, IERR
    LOGICAL FLAG

MPI_NULL_DELETE_FN(COMM, KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE, IERROR)
    INTEGER COMM, KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE, IERROR

MPI_TYPE_EXTENT(DATATYPE, EXTENT, IERROR)
    INTEGER DATATYPE, EXTENT, IERROR

MPI_TYPE_HINDEXED(COUNT, ARRAY_OF_BLOCKLENGTHS, ARRAY_OF_DISPLACEMENTS,
                  OLDTYPE, NEWTYPE, IERROR)

```

```

    INTEGER COUNT, ARRAY_OF_BLOCKLENGTHS(*), ARRAY_OF_DISPLACEMENTS(*),
    OLDTYPE, NEWTYPE, IERROR
MPI_TYPE_HVECTOR(COUNT, BLOCKLENGTH, STRIDE, OLDTYPE, NEWTYPE, IERROR)
    INTEGER COUNT, BLOCKLENGTH, STRIDE, OLDTYPE, NEWTYPE, IERROR
MPI_TYPE_LB( DATATYPE, DISPLACEMENT, IERROR)
    INTEGER DATATYPE, DISPLACEMENT, IERROR
MPI_TYPE_STRUCT(COUNT, ARRAY_OF_BLOCKLENGTHS, ARRAY_OF_DISPLACEMENTS,
    ARRAY_OF_TYPES, NEWTYPE, IERROR)
    INTEGER COUNT, ARRAY_OF_BLOCKLENGTHS(*), ARRAY_OF_DISPLACEMENTS(*),
    ARRAY_OF_TYPES(*), NEWTYPE, IERROR
MPI_TYPE_UB( DATATYPE, DISPLACEMENT, IERROR)
    INTEGER DATATYPE, DISPLACEMENT, IERROR
SUBROUTINE COPY_FUNCTION(OLDCOMM, KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,
    ATTRIBUTE_VAL_OUT, FLAG, IERR)
    INTEGER OLDCOMM, KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,
    ATTRIBUTE_VAL_OUT, IERR
    LOGICAL FLAG
SUBROUTINE DELETE_FUNCTION(COMM, KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE, IERR)
    INTEGER COMM, KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE, IERR

```

## A.4 C++ Bindings (deprecated)

### A.4.1 Point-to-Point Communication C++ Bindings

```

namespace MPI {

    {void Attach_buffer(void* buffer, int size) (binding deprecated, see
        Section 15.2) }

    {void Comm::Bsend(const void* buf, int count, const Datatype& datatype,
        int dest, int tag) const (binding deprecated, see Section 15.2) }

    {Prequest Comm::Bsend_init(const void* buf, int count, const
        Datatype& datatype, int dest, int tag) const (binding deprecated,
        see Section 15.2) }

    {void Request::Cancel() const (binding deprecated, see Section 15.2) }

    {int Detach_buffer(void*& buffer) (binding deprecated, see Section 15.2) }

    {void Request::Free() (binding deprecated, see Section 15.2) }

    {int Status::Get_count(const Datatype& datatype) const (binding deprecated,
        see Section 15.2) }

    {int Status::Get_error() const (binding deprecated, see Section 15.2) }

    {int Status::Get_source() const (binding deprecated, see Section 15.2) }

    {bool Request::Get_status() const (binding deprecated, see Section 15.2) }

    {bool Request::Get_status(Status& status) const (binding deprecated, see
        Section 15.2) }

    {int Status::Get_tag() const (binding deprecated, see Section 15.2) }

    {Request Comm::Ibsend(const void* buf, int count, const
        Datatype& datatype, int dest, int tag) const (binding deprecated,
        see Section 15.2) }

    {bool Comm::Iprobe(int source, int tag) const (binding deprecated, see
        Section 15.2) }

    {bool Comm::Iprobe(int source, int tag, Status& status) const (binding
        deprecated, see Section 15.2) }

    {Request Comm::Irecv(void* buf, int count, const Datatype& datatype,
        int source, int tag) const (binding deprecated, see Section 15.2) }

    {Request Comm::Irsend(const void* buf, int count, const
        Datatype& datatype, int dest, int tag) const (binding deprecated,
        see Section 15.2) }

    {bool Status::Is_cancelled() const (binding deprecated, see Section 15.2) }

    {Request Comm::Isend(const void* buf, int count, const
        Datatype& datatype, int dest, int tag) const (binding deprecated,

```

```

    see Section 15.2) }
1
2
{Request Comm::Issend(const void* buf, int count, const
3
    Datatype& datatype, int dest, int tag) const(binding deprecated,
4
    see Section 15.2) }
5
6
{void Comm::Probe(int source, int tag) const(binding deprecated, see
7
    Section 15.2) }
8
9
{void Comm::Probe(int source, int tag, Status& status) const(binding
10
    deprecated, see Section 15.2) }
11
12
{void Comm::Recv(void* buf, int count, const Datatype& datatype,
13
    int source, int tag) const(binding deprecated, see Section 15.2) }
14
15
{void Comm::Recv(void* buf, int count, const Datatype& datatype,
16
    int source, int tag, Status& status) const(binding deprecated, see
17
    Section 15.2) }
18
19
{Prequest Comm::Recv_init(void* buf, int count, const Datatype& datatype,
20
    int source, int tag) const(binding deprecated, see Section 15.2) }
21
22
{void Comm::Rsend(const void* buf, int count, const Datatype& datatype,
23
    int dest, int tag) const(binding deprecated, see Section 15.2) }
24
25
{Prequest Comm::Rsend_init(const void* buf, int count, const
26
    Datatype& datatype, int dest, int tag) const(binding deprecated,
27
    see Section 15.2) }
28
29
{void Comm::Send(const void* buf, int count, const Datatype& datatype,
30
    int dest, int tag) const(binding deprecated, see Section 15.2) }
31
32
{Prequest Comm::Send_init(const void* buf, int count, const
33
    Datatype& datatype, int dest, int tag) const(binding deprecated,
34
    see Section 15.2) }
35
36
{void Comm::Sendrecv(const void *sendbuf, int sendcount, const
37
    Datatype& sendtype, int dest, int sendtag, void *recvbuf,
38
    int recvcount, const Datatype& recvtype, int source,
39
    int recvtage) const(binding deprecated, see Section 15.2) }
40
41
{void Comm::Sendrecv(const void *sendbuf, int sendcount, const
42
    Datatype& sendtype, int dest, int sendtag, void *recvbuf,
43
    int recvcount, const Datatype& recvtype, int source,
44
    int recvtage, Status& status) const(binding deprecated, see
45
    Section 15.2) }
46
47
{void Comm::Sendrecv_replace(void* buf, int count, const
48
    Datatype& datatype, int dest, int sendtag, int source,
    int recvtage, Status& status) const(binding deprecated, see

```

```

1      Section 15.2) }
2
3      {void Status::Set_error(int error) (binding deprecated, see Section 15.2) }
4
5      {void Status::Set_source(int source) (binding deprecated, see Section 15.2) }
6
7      {void Status::Set_tag(int tag) (binding deprecated, see Section 15.2) }
8
9      {void Comm::Ssend(const void* buf, int count, const Datatype& datatype,
10         int dest, int tag) const (binding deprecated, see Section 15.2) }
11
12      {Prequest Comm::Ssend_init(const void* buf, int count, const
13         Datatype& datatype, int dest, int tag) const (binding deprecated,
14         see Section 15.2) }
15
16      {void Prequest::Start() (binding deprecated, see Section 15.2) }
17
18      {static void Prequest::Startall(int count,
19         Prequest array_of_requests[]) (binding deprecated, see Section 15.2) }
20
21      {bool Request::Test() (binding deprecated, see Section 15.2) }
22
23      {bool Request::Test(Status& status) (binding deprecated, see Section 15.2) }
24
25      {static bool Request::Testall(int count,
26         Request array_of_requests[]) (binding deprecated, see Section 15.2) }
27
28      {static bool Request::Testall(int count, Request array_of_requests[],
29         Status array_of_statuses[]) (binding deprecated, see Section 15.2) }
30
31      {static bool Request::Testany(int count, Request array_of_requests[],
32         int& index) (binding deprecated, see Section 15.2) }
33
34      {static bool Request::Testany(int count, Request array_of_requests[],
35         int& index, Status& status) (binding deprecated, see Section 15.2) }
36
37      {static int Request::Testsome(int incount, Request array_of_requests[],
38         int array_of_indices[]) (binding deprecated, see Section 15.2) }
39
40      {static int Request::Testsome(int incount, Request array_of_requests[],
41         int array_of_indices[], Status array_of_statuses[]) (binding
42         deprecated, see Section 15.2) }
43
44      {void Request::Wait() (binding deprecated, see Section 15.2) }
45
46      {void Request::Wait(Status& status) (binding deprecated, see Section 15.2) }
47
48      {static void Request::Waitall(int count,
49         Request array_of_requests[]) (binding deprecated, see Section 15.2) }
50
51      {static void Request::Waitall(int count, Request array_of_requests[],
52         Status array_of_statuses[]) (binding deprecated, see Section 15.2) }
53
54      {static int Request::Waitany(int count,
55         Request array_of_requests[]) (binding deprecated, see Section 15.2) }

```

```

{static int Request::Waitany(int count, Request array_of_requests[],
    Status& status) (binding deprecated, see Section 15.2) }
{static int Request::Waitsome(int incount, Request array_of_requests[],
    int array_of_indices[]) (binding deprecated, see Section 15.2) }
{static int Request::Waitsome(int incount, Request array_of_requests[],
    int array_of_indices[], Status array_of_statuses[]) (binding
    deprecated, see Section 15.2) }
};

```

#### A.4.2 Datatypes C++ Bindings

```

namespace MPI {
    {void Datatype::Commit() (binding deprecated, see Section 15.2) }
    {Datatype Datatype::Create_contiguous(int count) const (binding deprecated,
        see Section 15.2) }
    {Datatype Datatype::Create_darray(int size, int rank, int ndims,
        const int array_of_gsizes[], const int array_of_distrib[],
        const int array_of_dargs[], const int array_of_psize[],
        int order) const (binding deprecated, see Section 15.2) }
    {Datatype Datatype::Create_hindexed(int count,
        const int array_of_blocklengths[],
        const Aint array_of_displacements[]) const (binding deprecated, see
        Section 15.2) }
    {Datatype Datatype::Create_hvector(int count, int blocklength, Aint
        stride) const (binding deprecated, see Section 15.2) }
    {Datatype Datatype::Create_indexed(int count,
        const int array_of_blocklengths[],
        const int array_of_displacements[]) const (binding deprecated, see
        Section 15.2) }
    {Datatype Datatype::Create_indexed_block(int count, int blocklength,
        const int array_of_displacements[]) const (binding deprecated, see
        Section 15.2) }
    {Datatype Datatype::Create_resized(const Aint lb, const Aint extent)
        const (binding deprecated, see Section 15.2) }
    {static Datatype Datatype::Create_struct(int count,
        const int array_of_blocklengths[], const Aint
        array_of_displacements[],
        const Datatype array_of_types[]) (binding deprecated, see
        Section 15.2) }
}

```

```

1  {Datatype Datatype::Create_subarray(int ndims,
2      const int array_of_sizes[], const int array_of_subsizes[],
3      const int array_of_starts[], int order) const(binding deprecated,
4      see Section 15.2) }
5
6  {Datatype Datatype::Create_vector(int count, int blocklength, int stride)
7      const(binding deprecated, see Section 15.2) }
8
9  {Datatype Datatype::Dup() const(binding deprecated, see Section 15.2) }
10
11 {void Datatype::Free() (binding deprecated, see Section 15.2) }
12
13 {Aint Get_address(void* location) (binding deprecated, see Section 15.2) }
14
15 {void Datatype::Get_contents(int max_integers, int max_addresses,
16     int max_datatypes, int array_of_integers[],
17     Aint array_of_addresses[], Datatype array_of_datatypes[])
18     const(binding deprecated, see Section 15.2) }
19
20 {int Status::Get_elements(const Datatype& datatype) const(binding deprecated,
21     see Section 15.2) }
22
23 {void Datatype::Get_envelope(int& num_integers, int& num_addresses,
24     int& num_datatypes, int& combiner) const(binding deprecated, see
25     Section 15.2) }
26
27 {void Datatype::Get_extent(Aint& lb, Aint& extent) const(binding deprecated,
28     see Section 15.2) }
29
30 {int Datatype::Get_size() const(binding deprecated, see Section 15.2) }
31
32 {void Datatype::Get_true_extent(Aint& true_lb, Aint& true_extent)
33     const(binding deprecated, see Section 15.2) }
34
35 {void Datatype::Pack(const void* inbuf, int incount, void *outbuf,
36     int outsize, int& position, const Comm &comm) const(binding
37     deprecated, see Section 15.2) }
38
39 {void Datatype::Pack_external(const char* datarep, const void* inbuf,
40     int incount, void* outbuf, Aint outsize, Aint& position)
41     const(binding deprecated, see Section 15.2) }
42
43 {Aint Datatype::Pack_external_size(const char* datarep, int incount)
44     const(binding deprecated, see Section 15.2) }
45
46 {int Datatype::Pack_size(int incount, const Comm& comm) const(binding
47     deprecated, see Section 15.2) }
48
49 {void Datatype::Unpack(const void* inbuf, int insize, void *outbuf,
50     int outcount, int& position, const Comm& comm) const(binding
51     deprecated, see Section 15.2) }
52
53 {void Datatype::Unpack_external(const char* datarep, const void* inbuf,
54     Aint insize, Aint& position, void* outbuf, int outcount) const
55     (binding deprecated, see Section 15.2) }

```

};

## A.4.3 Collective Communication C++ Bindings

namespace MPI {

```

{void Comm::Allgather(const void* sendbuf, int sendcount, const
    Datatype& sendtype, void* recvbuf, int recvcount,
    const Datatype& recvttype) const = 0(binding deprecated, see
    Section 15.2) }

{void Comm::Allgatherv(const void* sendbuf, int sendcount, const
    Datatype& sendtype, void* recvbuf, const int recvcounts[],
    const int displs[], const Datatype& recvttype) const = 0(binding
    deprecated, see Section 15.2) }

{void Comm::Allreduce(const void* sendbuf, void* recvbuf, int count,
    const Datatype& datatype, const Op& op) const = 0(binding
    deprecated, see Section 15.2) }

{void Comm::Alltoall(const void* sendbuf, int sendcount, const
    Datatype& sendtype, void* recvbuf, int recvcount,
    const Datatype& recvttype) const = 0(binding deprecated, see
    Section 15.2) }

{void Comm::Alltoallv(const void* sendbuf, const int sendcounts[],
    const int sdispls[], const Datatype& sendtype, void* recvbuf,
    const int recvcounts[], const int rdispls[],
    const Datatype& recvttype) const = 0(binding deprecated, see
    Section 15.2) }

{void Comm::Alltoallw(const void* sendbuf, const int sendcounts[], const
    int sdispls[], const Datatype sendtypes[], void* recvbuf,
    const int recvcounts[], const int rdispls[], const Datatype
    recvtypes[]) const = 0(binding deprecated, see Section 15.2) }

{void Comm::Barrier() const = 0(binding deprecated, see Section 15.2) }

{void Comm::Bcast(void* buffer, int count, const Datatype& datatype,
    int root) const = 0(binding deprecated, see Section 15.2) }

{void Intracomm::Exscan(const void* sendbuf, void* recvbuf, int count,
    const Datatype& datatype, const Op& op) const(binding deprecated,
    see Section 15.2) }

{void Op::Free()(binding deprecated, see Section 15.2) }

{void Comm::Gather(const void* sendbuf, int sendcount, const
    Datatype& sendtype, void* recvbuf, int recvcount,
    const Datatype& recvttype, int root) const = 0(binding deprecated,
    see Section 15.2) }

{void Comm::Gatherv(const void* sendbuf, int sendcount, const
    Datatype& sendtype, void* recvbuf, const int recvcounts[],

```



```

1      const int displs[], const Datatype& recvtype, int root)
2      const = 0(binding deprecated, see Section 15.2) }
3
4      {Request Comm::Iallgather(const void* sendbuf, int sendcount, const
5      Datatype& sendtype, void* recvbuf, int recvcount,
6      const Datatype& recvtype) const = 0(binding deprecated, see
7      Section 15.2) }
8
9      {Request Comm::Iallgatherv(const void* sendbuf, int sendcount, const
10     Datatype& sendtype, void* recvbuf, const int recvcounts[],
11     const int displs[], const Datatype& recvtype) const = 0(binding
12     deprecated, see Section 15.2) }
13
14     {Request Comm::Iallreduce(const void* sendbuf, void* recvbuf, int count,
15     const Datatype& datatype, const Op& op) const = 0(binding
16     deprecated, see Section 15.2) }
17
18     {Request Comm::Ialltoall(const void* sendbuf, int sendcount, const
19     Datatype& sendtype, void* recvbuf, int recvcount,
20     const Datatype& recvtype) const = 0(binding deprecated, see
21     Section 15.2) }
22
23     {Request Comm::Ialltoallv(const void* sendbuf, const int sendcounts[],
24     const int sdispls[], const Datatype& sendtype, void* recvbuf,
25     const int recvcounts[], const int rdispls[],
26     const Datatype& recvtype) const = 0(binding deprecated, see
27     Section 15.2) }
28
29     {Request Comm::Ialltoallw(const void* sendbuf, const int sendcounts[],
30     const int sdispls[], const Datatype sendtypes[], void*
31     recvbuf, const int recvcounts[], const int rdispls[], const
32     Datatype recvtypes[]) const = 0(binding deprecated, see
33     Section 15.2) }
34
35     {Request Comm::Ibarrier() const = 0(binding deprecated, see Section 15.2) }
36
37     {Request Comm::Ibcast(void* buffer, int count, const Datatype& datatype,
38     int root) const = 0(binding deprecated, see Section 15.2) }
39
40     {Request Intracomm::Iexscan(const void* sendbuf, void* recvbuf, int
41     count, const Datatype& datatype, const Op& op) const(binding
42     deprecated, see Section 15.2) }
43
44     {Request Comm::Igatherv(const void* sendbuf, int sendcount, const
45     Datatype& sendtype, void* recvbuf, const int recvcounts[],
46     const int displs[], const Datatype& recvtype, int root)
47     const = 0(binding deprecated, see Section 15.2) }
48

```

```

{void Op::Init(User_function *function, bool commute) (binding deprecated, see
    Section 15.2) }
{Request Comm::Ireduce(const void* sendbuf, void* recvbuf, int count,
    const Datatype& datatype, const Op& op, int root)
    const = 0 (binding deprecated, see Section 15.2) }
{Request Comm::Ireduce_scatter(const void* sendbuf, void* recvbuf,
    int recvcnt, const Datatype& datatype, const Op& op)
    const = 0 (binding deprecated, see Section 15.2) }
{Request Comm::Ireduce_scatter_block(const void* sendbuf, void* recvbuf,
    int recvcnt, const Datatype& datatype, const Op& op)
    const = 0 (binding deprecated, see Section 15.2) }
{bool Op::Is_commutative() const (binding deprecated, see Section 15.2) }
{Request Intracomm::Iscomm(const void* sendbuf, void* recvbuf, int count,
    const Datatype& datatype, const Op& op) const (binding deprecated,
    see Section 15.2) }
{Request Comm::Iscomm(const void* sendbuf, int sendcount, const
    Datatype& sendtype, void* recvbuf, int recvcnt,
    const Datatype& recvttype, int root) const = 0 (binding deprecated,
    see Section 15.2) }
{Request Comm::Iscommerv(const void* sendbuf, const int sendcounts[],
    const int displs[], const Datatype& sendtype, void* recvbuf,
    int recvcnt, const Datatype& recvttype, int root)
    const = 0 (binding deprecated, see Section 15.2) }
{void Comm::Reduce(const void* sendbuf, void* recvbuf, int count,
    const Datatype& datatype, const Op& op, int root)
    const = 0 (binding deprecated, see Section 15.2) }
{void Op::Reduce_local(const void* inbuf, void* inoutbuf, int count,
    const Datatype& datatype) const (binding deprecated, see
    Section 15.2) }
{void Comm::Reduce_scatter(const void* sendbuf, void* recvbuf,
    int recvcnt, const Datatype& datatype, const Op& op)
    const = 0 (binding deprecated, see Section 15.2) }
{void Comm::Reduce_scatter_block(const void* sendbuf, void* recvbuf,
    int recvcnt, const Datatype& datatype, const Op& op)
    const = 0 (binding deprecated, see Section 15.2) }
{void Intracomm::Scan(const void* sendbuf, void* recvbuf, int count,
    const Datatype& datatype, const Op& op) const (binding deprecated,
    see Section 15.2) }
{void Comm::Scatter(const void* sendbuf, int sendcount, const
    Datatype& sendtype, void* recvbuf, int recvcnt,
    const Datatype& recvttype, int root) const = 0 (binding deprecated,

```

```

1         see Section 15.2) }
2
3     {void Comm::Scatterv(const void* sendbuf, const int sendcounts[],
4         const int displs[], const Datatype& sendtype, void* recvbuf,
5         int recvcnt, const Datatype& recvtpe, int root)
6         const = 0(binding deprecated, see Section 15.2) }
7
8 };
9

```

#### A.4.4 Groups, Contexts, Communicators, and Caching C++ Bindings

```

11 namespace MPI {
12
13     {Comm& Comm::Clone() const = 0(binding deprecated, see Section 15.2) }
14
15     {Cartcomm& Cartcomm::Clone() const(binding deprecated, see Section 15.2) }
16
17     {Distgraphcomm& Distgraphcomm::Clone() const(binding deprecated, see
18         Section 15.2) }
19
20     {Graphcomm& Graphcomm::Clone() const(binding deprecated, see Section 15.2) }
21
22     {Intercomm& Intercomm::Clone() const(binding deprecated, see Section 15.2) }
23
24     {Intracomm& Intracomm::Clone() const(binding deprecated, see Section 15.2) }
25
26     {static int Comm::Compare(const Comm& comm1, const Comm& comm2)(binding
27         deprecated, see Section 15.2) }
28
29     {static int Group::Compare(const Group& group1,
30         const Group& group2)(binding deprecated, see Section 15.2) }
31
32     {Intracomm Intracomm::Create(const Group& group) const(binding deprecated,
33         see Section 15.2) }
34
35     {Intercomm Intercomm::Create(const Group& group) const(binding deprecated,
36         see Section 15.2) }
37
38     {Intercomm Intracomm::Create_intercomm(int local_leader, const
39         Comm& peer_comm, int remote_leader, int tag) const(binding
40         deprecated, see Section 15.2) }
41
42     {static int Comm::Create_keyval(Comm::Copy_attr_function*
43         comm_copy_attr_fn,
44         Comm::Delete_attr_function* comm_delete_attr_fn,
45         void* extra_state)(binding deprecated, see Section 15.2) }
46
47     {static int Datatype::Create_keyval(Datatype::Copy_attr_function*
48         type_copy_attr_fn, Datatype::Delete_attr_function*
49         type_delete_attr_fn, void* extra_state)(binding deprecated, see
50         Section 15.2) }
51
52     {static int Win::Create_keyval(Win::Copy_attr_function* win_copy_attr_fn,
53         Win::Delete_attr_function* win_delete_attr_fn,
54         void* extra_state)(binding deprecated, see Section 15.2) }
55

```

```

{void Comm::Delete_attr(int comm_keyval) (binding deprecated, see Section 15.2) }
{void Datatype::Delete_attr(int type_keyval) (binding deprecated, see
    Section 15.2) }
{void Win::Delete_attr(int win_keyval) (binding deprecated, see Section 15.2) }
{static Group Group::Difference(const Group& group1,
    const Group& group2) (binding deprecated, see Section 15.2) }
{Cartcomm Cartcomm::Dup() const (binding deprecated, see Section 15.2) }
{Distgraphcomm Distgraphcomm::Dup() const (binding deprecated, see Section 15.2) }
{Graphcomm Graphcomm::Dup() const (binding deprecated, see Section 15.2) }
{Intercomm Intercomm::Dup() const (binding deprecated, see Section 15.2) }
{Intracomm Intracomm::Dup() const (binding deprecated, see Section 15.2) }
{Group Group::Excl(int n, const int ranks[]) const (binding deprecated, see
    Section 15.2) }
{void Comm::Free() (binding deprecated, see Section 15.2) }
{void Group::Free() (binding deprecated, see Section 15.2) }
{static void Comm::Free_keyval(int& comm_keyval) (binding deprecated, see
    Section 15.2) }
{static void Datatype::Free_keyval(int& type_keyval) (binding deprecated, see
    Section 15.2) }
{static void Win::Free_keyval(int& win_keyval) (binding deprecated, see
    Section 15.2) }
{bool Comm::Get_attr(int comm_keyval, void* attribute_val) const (binding
    deprecated, see Section 15.2) }
{bool Datatype::Get_attr(int type_keyval, void* attribute_val)
    const (binding deprecated, see Section 15.2) }
{bool Win::Get_attr(int win_keyval, void* attribute_val) const (binding
    deprecated, see Section 15.2) }
{Group Comm::Get_group() const (binding deprecated, see Section 15.2) }
{void Comm::Get_name(char* comm_name, int& resultlen) const (binding
    deprecated, see Section 15.2) }
{void Datatype::Get_name(char* type_name, int& resultlen) const (binding
    deprecated, see Section 15.2) }
{void Win::Get_name(char* win_name, int& resultlen) const (binding deprecated,
    see Section 15.2) }
{int Comm::Get_rank() const (binding deprecated, see Section 15.2) }

```

```

1  {int Group::Get_rank() const(binding deprecated, see Section 15.2) }
2
3  {Group Intercomm::Get_remote_group() const(binding deprecated, see Section 15.2)
4      }
5
6  {int Intercomm::Get_remote_size() const(binding deprecated, see Section 15.2) }
7
8  {int Comm::Get_size() const(binding deprecated, see Section 15.2) }
9
10 {int Group::Get_size() const(binding deprecated, see Section 15.2) }
11
12 {Group Group::Incl(int n, const int ranks[]) const(binding deprecated, see
13     Section 15.2) }
14
15 {static Group Group::Intersect(const Group& group1,
16     const Group& group2)(binding deprecated, see Section 15.2) }
17
18 {bool Comm::Is_inter() const(binding deprecated, see Section 15.2) }
19
20 {Intracomm Intercomm::Merge(bool high) const(binding deprecated, see
21     Section 15.2) }
22
23 {Group Group::Range_excl(int n, const int ranges[][3]) const(binding
24     deprecated, see Section 15.2) }
25
26 {Group Group::Range_incl(int n, const int ranges[][3]) const(binding
27     deprecated, see Section 15.2) }
28
29 {void Comm::Set_attr(int comm_keyval, const void* attribute_val)
30     const(binding deprecated, see Section 15.2) }
31
32 {void Datatype::Set_attr(int type_keyval, const void*
33     attribute_val)(binding deprecated, see Section 15.2) }
34
35 {void Win::Set_attr(int win_keyval, const void* attribute_val)(binding
36     deprecated, see Section 15.2) }
37
38 {void Comm::Set_name(const char* comm_name)(binding deprecated, see
39     Section 15.2) }
40
41 {void Datatype::Set_name(const char* type_name)(binding deprecated, see
42     Section 15.2) }
43
44 {void Win::Set_name(const char* win_name)(binding deprecated, see Section 15.2) }
45
46 {Intercomm Intercomm::Split(int color, int key) const(binding deprecated, see
47     Section 15.2) }
48
49 {Intracomm Intracomm::Split(int color, int key) const(binding deprecated, see
50     Section 15.2) }
51
52 {static void Group::Translate_ranks (const Group& group1, int n,
53     const int ranks1[], const Group& group2, int ranks2[])(binding
54     deprecated, see Section 15.2) }
55
56 {static Group Group::Union(const Group& group1,
57     const Group& group2)(binding deprecated, see Section 15.2) }

```

};

## A.4.5 Process Topologies C++ Bindings

namespace MPI {

```
{void Compute_dims(int nnodes, int ndims, int dims[]) (binding deprecated, see
    Section 15.2) }
```

```
{Cartcomm Intracomm::Create_cart(int ndims, const int dims[],
    const bool periods[], bool reorder) const (binding deprecated, see
    Section 15.2) }
```

```
{Graphcomm Intracomm::Create_graph(int nnodes, const int index[],
    const int edges[], bool reorder) const (binding deprecated, see
    Section 15.2) }
```

```
{Distgraphcomm Intracomm::Dist_graph_create(int n, const int sources[],
    const int degrees[], const int destinations[],
    const int weights[], const Info& info, bool reorder)
    const (binding deprecated, see Section 15.2) }
```

```
{Distgraphcomm Intracomm::Dist_graph_create(int n, const int sources[],
    const int degrees[], const int destinations[],
    const Info& info, bool reorder) const (binding deprecated, see
    Section 15.2) }
```

```
{Distgraphcomm Intracomm::Dist_graph_create_adjacent(int indegree,
    const int sources[], const int sourceweights[], int outdegree,
    const int destinations[], const int destweights[],
    const Info& info, bool reorder) const (binding deprecated, see
    Section 15.2) }
```

```
{Distgraphcomm Intracomm::Dist_graph_create_adjacent(int indegree,
    const int sources[], int outdegree, const int destinations[],
    const Info& info, bool reorder) const (binding deprecated, see
    Section 15.2) }
```

```
{int Cartcomm::Get_cart_rank(const int coords[]) const (binding deprecated,
    see Section 15.2) }
```

```
{void Cartcomm::Get_coords(int rank, int maxdims, int coords[])
    const (binding deprecated, see Section 15.2) }
```

```
{int Cartcomm::Get_dim() const (binding deprecated, see Section 15.2) }
```

```
{void Graphcomm::Get_dims(int nnodes[], int nedges[]) const (binding
    deprecated, see Section 15.2) }
```

```
{void Distgraphcomm::Get_dist_neighbors(int maxindegree, int sources[],
    int sourceweights[], int maxoutdegree, int destinations[],
    int destweights[]) (binding deprecated, see Section 15.2) }
```

```

1  {void Distgraphcomm::Get_dist_neighbors_count(int rank, int indegree[],
2      int outdegree[], bool& weighted) const(binding deprecated, see
3      Section 15.2) }
4
5  {void Graphcomm::Get_neighbors(int rank, int maxneighbors, int
6      neighbors[]) const(binding deprecated, see Section 15.2) }
7
8  {int Graphcomm::Get_neighbors_count(int rank) const(binding deprecated, see
9      Section 15.2) }
10
11 {void Cartcomm::Get_topo(int maxdims, int dims[], bool periods[],
12     int coords[]) const(binding deprecated, see Section 15.2) }
13
14 {void Graphcomm::Get_topo(int maxindex, int maxedges, int index[],
15     int edges[]) const(binding deprecated, see Section 15.2) }
16
17 {int Comm::Get_topology() const(binding deprecated, see Section 15.2) }
18
19 {int Cartcomm::Map(int ndims, const int dims[], const bool periods[])
20     const(binding deprecated, see Section 15.2) }
21
22 {int Graphcomm::Map(int nnodes, const int index[], const int edges[])
23     const(binding deprecated, see Section 15.2) }
24
25 {void Cartcomm::Shift(int direction, int disp, int& rank_source,
26     int& rank_dest) const(binding deprecated, see Section 15.2) }
27
28 {Cartcomm Cartcomm::Sub(const bool remain_dims[]) const(binding deprecated,
29     see Section 15.2) }
30
31 };

```

#### A.4.6 MPI Environmental Management C++ Bindings

```

31 namespace MPI {
32
33     {void Comm::Abort(int errorcode)(binding deprecated, see Section 15.2) }
34
35     {int Add_error_class()(binding deprecated, see Section 15.2) }
36
37     {int Add_error_code(int errorclass)(binding deprecated, see Section 15.2) }
38
39     {void Add_error_string(int errorcode, const char* string)(binding deprecated,
40         see Section 15.2) }
41
42     {void* Alloc_mem(Aint size, const Info& info)(binding deprecated, see
43         Section 15.2) }
44
45     {void Comm::Call_errhandler(int errorcode) const(binding deprecated, see
46         Section 15.2) }
47
48     {void File::Call_errhandler(int errorcode) const(binding deprecated, see
49         Section 15.2) }
50
51     {void Win::Call_errhandler(int errorcode) const(binding deprecated, see
52         Section 15.2) }
53
54 }

```



```

{static Errhandler Comm::Create_errhandler(Comm::Errhandler_function*
    function) (binding deprecated, see Section 15.2) }
{static Errhandler File::Create_errhandler(File::Errhandler_function*
    function) (binding deprecated, see Section 15.2) }
{static Errhandler Win::Create_errhandler(Win::Errhandler_function*
    function) (binding deprecated, see Section 15.2) }
{void Finalize() (binding deprecated, see Section 15.2) }
{void Errhandler::Free() (binding deprecated, see Section 15.2) }
{void Free_mem(void *base) (binding deprecated, see Section 15.2) }
{Errhandler Comm::Get_errhandler() const (binding deprecated, see Section 15.2) }
{Errhandler File::Get_errhandler() const (binding deprecated, see Section 15.2) }
{Errhandler Win::Get_errhandler() const (binding deprecated, see Section 15.2) }
{int Get_error_class(int errorcode) (binding deprecated, see Section 15.2) }
{void Get_error_string(int errorcode, char* name, int& resultlen) (binding
    deprecated, see Section 15.2) }
{void Get_processor_name(char* name, int& resultlen) (binding deprecated, see
    Section 15.2) }
{void Get_version(int& version, int& subversion) (binding deprecated, see
    Section 15.2) }
{void Init() (binding deprecated, see Section 15.2) }
{void Init(int& argc, char**& argv) (binding deprecated, see Section 15.2) }
{bool Is_finalized() (binding deprecated, see Section 15.2) }
{bool Is_initialized() (binding deprecated, see Section 15.2) }
{void Comm::Set_errhandler(const Errhandler& errhandler) (binding deprecated,
    see Section 15.2) }
{void File::Set_errhandler(const Errhandler& errhandler) (binding deprecated,
    see Section 15.2) }
{void Win::Set_errhandler(const Errhandler& errhandler) (binding deprecated,
    see Section 15.2) }
{double Wtick() (binding deprecated, see Section 15.2) }
{double Wtime() (binding deprecated, see Section 15.2) }
};

```

#### A.4.7 The Info Object C++ Bindings

```
namespace MPI {
```



```

1  {static Info Info::Create() (binding deprecated, see Section 15.2) }
2
3  {void Info::Delete(const char* key) (binding deprecated, see Section 15.2) }
4
5  {Info Info::Dup() const (binding deprecated, see Section 15.2) }
6
7  {void Info::Free() (binding deprecated, see Section 15.2) }
8
9  {bool Info::Get(const char* key, int valuelen, char* value) const (binding
10     deprecated, see Section 15.2) }
11
12 {int Info::Get_nkeys() const (binding deprecated, see Section 15.2) }
13
14 {void Info::Get_nthkey(int n, char* key) const (binding deprecated, see
15     Section 15.2) }
16
17 {bool Info::Get_valuelen(const char* key, int& valuelen) const (binding
18     deprecated, see Section 15.2) }
19
20 {void Info::Set(const char* key, const char* value) (binding deprecated, see
21     Section 15.2) }
22
23 };

```

#### A.4.8 Process Creation and Management C++ Bindings

```

23 namespace MPI {
24
25     {Intercomm Intracomm::Accept(const char* port_name, const Info& info,
26         int root) const (binding deprecated, see Section 15.2) }
27
28     {void Close_port(const char* port_name) (binding deprecated, see Section 15.2) }
29
30     {Intercomm Intracomm::Connect(const char* port_name, const Info& info,
31         int root) const (binding deprecated, see Section 15.2) }
32
33     {void Comm::Disconnect() (binding deprecated, see Section 15.2) }
34
35     {static Intercomm Comm::Get_parent() (binding deprecated, see Section 15.2) }
36
37     {static Intercomm Comm::Join(const int fd) (binding deprecated, see Section 15.2) }
38
39     {void Lookup_name(const char* service_name, const Info& info,
40         char* port_name) (binding deprecated, see Section 15.2) }
41
42     {void Open_port(const Info& info, char* port_name) (binding deprecated, see
43         Section 15.2) }
44
45     {void Publish_name(const char* service_name, const Info& info,
46         const char* port_name) (binding deprecated, see Section 15.2) }
47
48     {Intercomm Intracomm::Spawn(const char* command, const char* argv[],
49         int maxprocs, const Info& info, int root) const (binding
50         deprecated, see Section 15.2) }

```

```

{Intercomm Intracomm::Spawn(const char* command, const char* argv[],
    int maxprocs, const Info& info, int root,
    int array_of_errcodes[]) const(binding deprecated, see Section 15.2)
    }
}

{Intercomm Intracomm::Spawn_multiple(int count,
    const char* array_of_commands[], const char** array_of_argv[],
    const int array_of_maxprocs[], const Info array_of_info[],
    int root) (binding deprecated, see Section 15.2) }

{Intercomm Intracomm::Spawn_multiple(int count,
    const char* array_of_commands[], const char** array_of_argv[],
    const int array_of_maxprocs[], const Info array_of_info[],
    int root, int array_of_errcodes[]) (binding deprecated, see
    Section 15.2) }

{void Unpublish_name(const char* service_name, const Info& info,
    const char* port_name) (binding deprecated, see Section 15.2) }

};

```

#### A.4.9 One-Sided Communications C++ Bindings

```

namespace MPI {

{void Win::Accumulate(const void* origin_addr, int origin_count, const
    Datatype& origin_datatype, int target_rank, Aint target_disp,
    int target_count, const Datatype& target_datatype, const Op&
    op) const(binding deprecated, see Section 15.2) }

{void Win::Complete() const(binding deprecated, see Section 15.2) }

{static Win Win::Create(const void* base, Aint size, int disp_unit, const
    Info& info, const Intracomm& comm) (binding deprecated, see
    Section 15.2) }

{void Win::Fence(int assert) const(binding deprecated, see Section 15.2) }

{void Win::Free() (binding deprecated, see Section 15.2) }

{void Win::Get(void *origin_addr, int origin_count, const Datatype&
    origin_datatype, int target_rank, Aint target_disp, int
    target_count, const Datatype& target_datatype) const(binding
    deprecated, see Section 15.2) }

{Group Win::Get_group() const(binding deprecated, see Section 15.2) }

{void Win::Lock(int lock_type, int rank, int assert) const(binding
    deprecated, see Section 15.2) }

{void Win::Post(const Group& group, int assert) const(binding deprecated, see
    Section 15.2) }
}

```

```

1  {void Win::Put(const void* origin_addr, int origin_count, const Datatype&
2      origin_datatype, int target_rank, Aint target_disp, int
3      target_count, const Datatype& target_datatype) const(binding
4      deprecated, see Section 15.2) }
5
6  {void Win::Start(const Group& group, int assert) const(binding deprecated,
7      see Section 15.2) }
8
9  {bool Win::Test() const(binding deprecated, see Section 15.2) }
10
11 {void Win::Unlock(int rank) const(binding deprecated, see Section 15.2) }
12
13 {void Win::Wait() const(binding deprecated, see Section 15.2) }
14 };

```

#### A.4.10 External Interfaces C++ Bindings

```

15 namespace MPI {
16
17     {void Grequest::Complete() (binding deprecated, see Section 15.2) }
18
19     {int Init_thread(int required) (binding deprecated, see Section 15.2) }
20
21     {int Init_thread(int& argc, char**& argv, int required) (binding deprecated,
22         see Section 15.2) }
23
24     {bool Is_thread_main() (binding deprecated, see Section 15.2) }
25
26     {int Query_thread() (binding deprecated, see Section 15.2) }
27
28     {void Status::Set_cancelled(bool flag) (binding deprecated, see Section 15.2) }
29
30     {void Status::Set_elements(const Datatype& datatype, int count) (binding
31         deprecated, see Section 15.2) }
32
33     {static Grequest Grequest::Start(const Grequest::Query_function*
34         query_fn, const Grequest::Free_function* free_fn,
35         const Grequest::Cancel_function* cancel_fn,
36         void *extra_state) (binding deprecated, see Section 15.2) }
37
38 };

```

#### A.4.11 I/O C++ Bindings

```

39 namespace MPI {
40
41     {void File::Close() (binding deprecated, see Section 15.2) }
42
43     {static void File::Delete(const char* filename, const Info& info) (binding
44         deprecated, see Section 15.2) }
45
46     {int File::Get_amode() const(binding deprecated, see Section 15.2) }
47
48     {bool File::Get_atomicsity() const(binding deprecated, see Section 15.2) }

```

```

{Offset File::Get_byte_offset(const Offset disp) const(binding deprecated,
    see Section 15.2) }
{Group File::Get_group() const(binding deprecated, see Section 15.2) }
{Info File::Get_info() const(binding deprecated, see Section 15.2) }
{Offset File::Get_position() const(binding deprecated, see Section 15.2) }
{Offset File::Get_position_shared() const(binding deprecated, see Section 15.2) }
{Offset File::Get_size() const(binding deprecated, see Section 15.2) }
{Aint File::Get_type_extent(const Datatype& datatype) const(binding
    deprecated, see Section 15.2) }
{void File::Get_view(Offset& disp, Datatype& etype, Datatype& filetype,
    char* datarep) const(binding deprecated, see Section 15.2) }
{Request File::Iread(void* buf, int count,
    const Datatype& datatype)(binding deprecated, see Section 15.2) }
{Request File::Iread_at(Offset offset, void* buf, int count,
    const Datatype& datatype)(binding deprecated, see Section 15.2) }
{Request File::Iread_shared(void* buf, int count,
    const Datatype& datatype)(binding deprecated, see Section 15.2) }
{Request File::Iwrite(const void* buf, int count,
    const Datatype& datatype)(binding deprecated, see Section 15.2) }
{Request File::Iwrite_at(Offset offset, const void* buf, int count,
    const Datatype& datatype)(binding deprecated, see Section 15.2) }
{Request File::Iwrite_shared(const void* buf, int count,
    const Datatype& datatype)(binding deprecated, see Section 15.2) }
{static File File::Open(const Intracomm& comm, const char* filename,
    int amode, const Info& info)(binding deprecated, see Section 15.2) }
{void File::Preallocate(Offset size)(binding deprecated, see Section 15.2) }
{void File::Read(void* buf, int count, const Datatype& datatype)(binding
    deprecated, see Section 15.2) }
{void File::Read(void* buf, int count, const Datatype& datatype, Status&
    status)(binding deprecated, see Section 15.2) }
{void File::Read_all(void* buf, int count,
    const Datatype& datatype)(binding deprecated, see Section 15.2) }
{void File::Read_all(void* buf, int count, const Datatype& datatype,
    Status& status)(binding deprecated, see Section 15.2) }
{void File::Read_all_begin(void* buf, int count,
    const Datatype& datatype)(binding deprecated, see Section 15.2) }

```

```

1  {void File::Read_all_end(void* buf) (binding deprecated, see Section 15.2) }
2
3  {void File::Read_all_end(void* buf, Status& status) (binding deprecated, see
4      Section 15.2) }
5
6  {void File::Read_at(Offset offset, void* buf, int count,
7      const Datatype& datatype) (binding deprecated, see Section 15.2) }
8
9  {void File::Read_at(Offset offset, void* buf, int count,
10     const Datatype& datatype, Status& status) (binding deprecated, see
11     Section 15.2) }
12
13 {void File::Read_at_all(Offset offset, void* buf, int count,
14     const Datatype& datatype) (binding deprecated, see Section 15.2) }
15
16 {void File::Read_at_all(Offset offset, void* buf, int count,
17     const Datatype& datatype, Status& status) (binding deprecated, see
18     Section 15.2) }
19
20 {void File::Read_at_all_begin(Offset offset, void* buf, int count,
21     const Datatype& datatype) (binding deprecated, see Section 15.2) }
22
23 {void File::Read_at_all_end(void* buf) (binding deprecated, see Section 15.2) }
24
25 {void File::Read_at_all_end(void* buf, Status& status) (binding deprecated,
26     see Section 15.2) }
27
28 {void File::Read_ordered(void* buf, int count,
29     const Datatype& datatype) (binding deprecated, see Section 15.2) }
30
31 {void File::Read_ordered(void* buf, int count, const Datatype& datatype,
32     Status& status) (binding deprecated, see Section 15.2) }
33
34 {void File::Read_ordered_begin(void* buf, int count,
35     const Datatype& datatype) (binding deprecated, see Section 15.2) }
36
37 {void File::Read_ordered_end(void* buf) (binding deprecated, see Section 15.2) }
38
39 {void File::Read_ordered_end(void* buf, Status& status) (binding deprecated,
40     see Section 15.2) }
41
42 {void File::Read_shared(void* buf, int count,
43     const Datatype& datatype) (binding deprecated, see Section 15.2) }
44
45 {void File::Read_shared(void* buf, int count, const Datatype& datatype,
46     Status& status) (binding deprecated, see Section 15.2) }
47
48 {void Register_datarep(const char* datarep,
49     Datarep_conversion_function* read_conversion_fn,
50     Datarep_conversion_function* write_conversion_fn,
51     Datarep_extent_function* dtype_file_extent_fn,
52     void* extra_state) (binding deprecated, see Section 15.2) }
53
54 {void File::Seek(Offset offset, int whence) (binding deprecated, see
55     Section 15.2) }
56

```

```

{void File::Seek_shared(Offset offset, int whence) (binding deprecated, see
    Section 15.2) }
{void File::Set_atomicity(bool flag) (binding deprecated, see Section 15.2) }
{void File::Set_info(const Info& info) (binding deprecated, see Section 15.2) }
{void File::Set_size(Offset size) (binding deprecated, see Section 15.2) }
{void File::Set_view(Offset disp, const Datatype& etype,
    const Datatype& filetype, const char* datarep,
    const Info& info) (binding deprecated, see Section 15.2) }
{void File::Sync() (binding deprecated, see Section 15.2) }
{void File::Write(const void* buf, int count,
    const Datatype& datatype) (binding deprecated, see Section 15.2) }
{void File::Write(const void* buf, int count, const Datatype& datatype,
    Status& status) (binding deprecated, see Section 15.2) }
{void File::Write_all(const void* buf, int count,
    const Datatype& datatype) (binding deprecated, see Section 15.2) }
{void File::Write_all(const void* buf, int count,
    const Datatype& datatype, Status& status) (binding deprecated, see
    Section 15.2) }
{void File::Write_all_begin(const void* buf, int count,
    const Datatype& datatype) (binding deprecated, see Section 15.2) }
{void File::Write_all_end(const void* buf) (binding deprecated, see Section 15.2) }
{void File::Write_all_end(const void* buf, Status& status) (binding
    deprecated, see Section 15.2) }
{void File::Write_at(Offset offset, const void* buf, int count,
    const Datatype& datatype) (binding deprecated, see Section 15.2) }
{void File::Write_at(Offset offset, const void* buf, int count,
    const Datatype& datatype, Status& status) (binding deprecated, see
    Section 15.2) }
{void File::Write_at_all(Offset offset, const void* buf, int count,
    const Datatype& datatype) (binding deprecated, see Section 15.2) }
{void File::Write_at_all(Offset offset, const void* buf, int count,
    const Datatype& datatype, Status& status) (binding deprecated, see
    Section 15.2) }
{void File::Write_at_all_begin(Offset offset, const void* buf, int count,
    const Datatype& datatype) (binding deprecated, see Section 15.2) }
{void File::Write_at_all_end(const void* buf) (binding deprecated, see
    Section 15.2) }

```

```

1  {void File::Write_at_all_end(const void* buf, Status& status) (binding
2      deprecated, see Section 15.2) }
3
4  {void File::Write_ordered(const void* buf, int count,
5      const Datatype& datatype) (binding deprecated, see Section 15.2) }
6
7  {void File::Write_ordered(const void* buf, int count,
8      const Datatype& datatype, Status& status) (binding deprecated, see
9      Section 15.2) }
10
11 {void File::Write_ordered_begin(const void* buf, int count,
12     const Datatype& datatype) (binding deprecated, see Section 15.2) }
13
14 {void File::Write_ordered_end(const void* buf) (binding deprecated, see
15     Section 15.2) }
16
17 {void File::Write_ordered_end(const void* buf, Status& status) (binding
18     deprecated, see Section 15.2) }
19
20 {void File::Write_shared(const void* buf, int count,
21     const Datatype& datatype) (binding deprecated, see Section 15.2) }
22
23 {void File::Write_shared(const void* buf, int count,
24     const Datatype& datatype, Status& status) (binding deprecated, see
25     Section 15.2) }
26
27 };

```

#### A.4.12 Language Bindings C++ Bindings

```

28 namespace MPI {
29
30     {static Datatype Datatype::Create_f90_complex(int p, int r) (binding
31         deprecated, see Section 15.2) }
32
33     {static Datatype Datatype::Create_f90_integer(int r) (binding deprecated, see
34         Section 15.2) }
35
36     {static Datatype Datatype::Create_f90_real(int p, int r) (binding deprecated,
37         see Section 15.2) }
38
39     Exception::Exception(int error_code)
40
41     {int Exception::Get_error_class() const (binding deprecated, see Section 15.2) }
42
43     {int Exception::Get_error_code() const (binding deprecated, see Section 15.2) }
44
45     {const char* Exception::Get_error_string() const (binding deprecated, see
46         Section 15.2) }
47
48     {static Datatype Datatype::Match_size(int typeclass, int size) (binding
49         deprecated, see Section 15.2) }
50
51 };

```

## A.4.13 Profiling Interface C++ Bindings

```
namespace MPI {
    {void Pcontrol(const int level, ...) (binding deprecated, see Section 15.2) }
};
```

[C++ Deprecate Functions section]

## A.4.14 C++ Bindings on all MPI Classes

The C++ language requires all classes to have four special functions: a default constructor, a copy constructor, a destructor, and an assignment operator. The bindings for these functions are listed below; their semantics are discussed in Section 16.1.5. The two constructors are *not virtual*. The bindings prototype functions are using the type `<CLASS>` rather than listing each function for every MPI class. The token `<CLASS>` can be replaced with valid MPI-2 class names, such as `Group`, `Datatype`, etc., except when noted. In addition, bindings are provided for comparison and inter-language operability from Sections 16.1.5 and 16.1.9.

## A.4.15 Construction / Destruction

```
namespace MPI {
    <CLASS>::<CLASS>()
    <CLASS>::~~<CLASS>()
};
```

## A.4.16 Copy / Assignment

```
namespace MPI {
    <CLASS>::<CLASS>(const <CLASS>& data)
    <CLASS>& <CLASS>::operator=(const <CLASS>& data)
};
```

## A.4.17 Comparison

Since `Status` instances are not handles to underlying MPI objects, the `operator==( )` and `operator!=( )` functions are not defined on the `Status` class.

```
namespace MPI {
    bool <CLASS>::operator==(const <CLASS>& data) const
    bool <CLASS>::operator!=(const <CLASS>& data) const
};
```



#### A.4.18 Inter-language Operability

Since there are no C++ `MPI::STATUS_IGNORE` and `MPI::STATUSES_IGNORE` objects, the result of promoting the C or Fortran handles (`MPI_STATUS_IGNORE` and `MPI_STATUSES_IGNORE`) to C++ is undefined.

```
namespace MPI {
    <CLASS>& <CLASS>::operator=(const MPI_<CLASS>& data)
    <CLASS>::<CLASS>(const MPI_<CLASS>& data)
    <CLASS>::operator MPI_<CLASS>() const
};
```

# Annex B

## Change-Log

This annex summarizes changes from the previous version of the MPI standard to the version presented by this document. Only significant changes (i.e., clarifications and new features) that might either require implementation effort in the MPI libraries or change the understanding of MPI from a user's perspective are presented. Editorial modifications, formatting, typo corrections and minor clarifications are not shown.

### B.1 Changes from Version 2.2 to Version 3.0

1. Chapter 5 on page 133 and Section 5.12 on page 190.  
Added nonblocking interfaces to all collective operations.
2. Sections 2.5.8, 3.2.2, 3.3, 5.9.2, on pages 16, 29, 31, 172, Sections 4.1.5, 4.1.7, 4.1.8, 4.1.11, 12.3 on pages 97, 99, 101, 105, 412, Sections 12.3 on pages 412, and Appendix A.1.1 on page 545.  
New inquiry functions, `MPI_TYPE_SIZE_X`, `MPI_TYPE_GET_EXTENT_X`, `MPI_TYPE_GET_TRUE_EXTENT_X`, and `MPI_GET_ELEMENTS_X`, return their results as an `MPI_Count` value, which is a new type large enough to represent element counts in memory, file views, etc. A new function, `MPI_STATUS_SET_ELEMENTS_X`, modifies the opaque part of `MPI_STATUS` so that a call to `MPI_GET_ELEMENTS_X` returns the provided `MPI_Count` value.
3. Sections 3.2.5, 4.1.5, 4.1.11, 4.2 on pages 34, 97, 105, 127.  
The functions `MPI_GET_COUNT`, `MPI_TYPE_SIZE`, and `MPI_GET_ELEMENTS` are now defined to set the count parameter to `MPI_UNDEFINED` when that parameter would overflow. The function `MPI_PACK_SIZE` is now defined to set the size parameter to `MPI_UNDEFINED` when that parameter would overflow. In all other MPI-2.2 routines, the type and semantics of the count arguments are kept unchanged, i.e., `int` or `INTEGER`.

### B.2 Changes from Version 2.1 to Version 2.2

1. Section 2.5.4 on page 14.  
It is now guaranteed that predefined named constant handles (as other constants) can be used in initialization expressions or assignments, i.e., also before the call to `MPI_INIT`.

2. Section 2.6 on page 16, Section 2.6.4 on page 19, and Section 16.1 on page 499.  
The C++ language bindings have been deprecated and may be removed in a future version of the MPI specification.
3. Section 3.2.2 on page 29.  
MPI\_CHAR for printable characters is now defined for C type char (instead of signed char). This change should not have any impact on applications nor on MPI libraries (except some comment lines), because printable characters could and can be stored in any of the C types char, signed char, and unsigned char, and MPI\_CHAR is not allowed for predefined reduction operations.
4. Section 3.2.2 on page 29.  
MPI\_(U)INT{8,16,32,64}\_T, MPI\_AINT, MPI\_OFFSET, MPI\_C\_BOOL, MPI\_C\_COMPLEX, MPI\_C\_FLOAT\_COMPLEX, MPI\_C\_DOUBLE\_COMPLEX, and MPI\_C\_LONG\_DOUBLE\_COMPLEX are now valid predefined MPI datatypes.
5. Section 3.4 on page 41, Section 3.7.2 on page 52, Section 3.9 on page 71, and Section 5.1 on page 133.  
The read access restriction on the send buffer for blocking, non blocking and collective API has been lifted. It is permitted to access for read the send buffer while the operation is in progress.
6. Section 3.7 on page 50.  
The Advice to users for IBSEND and IRSEND was slightly changed.
7. Section 3.7.3 on page 55.  
The advice to free an active request was removed in the Advice to users for MPI\_REQUEST\_FREE.
8. Section 3.7.6 on page 66.  
MPI\_REQUEST\_GET\_STATUS changed to permit inactive or null requests as input.
9. Section 5.8 on page 165.  
“In place” option is added to MPI\_ALLTOALL, MPI\_ALLTOALLV, and MPI\_ALLTOALLW for intracommunicators.
10. Section 5.9.2 on page 172.  
Predefined parameterized datatypes (e.g., returned by MPI\_TYPE\_CREATE\_F90\_REAL) and optional named predefined datatypes (e.g. MPI\_REAL8) have been added to the list of valid datatypes in reduction operations.
11. Section 5.9.2 on page 172.  
MPI\_(U)INT{8,16,32,64}\_T are all considered C integer types for the purposes of the predefined reduction operators. MPI\_AINT and MPI\_OFFSET are considered Fortran integer types. MPI\_C\_BOOL is considered a Logical type. MPI\_C\_COMPLEX, MPI\_C\_FLOAT\_COMPLEX, MPI\_C\_DOUBLE\_COMPLEX, and MPI\_C\_LONG\_DOUBLE\_COMPLEX are considered Complex types.
12. Section 5.9.7 on page 183.  
The local routines MPI\_REDUCE\_LOCAL and MPI\_OP\_COMMUTATIVE have been added.

13. Section 5.10.1 on page 185.  
The collective function `MPI_REDUCE_SCATTER_BLOCK` is added to the MPI standard.
14. Section 5.11.2 on page 188.  
Added in place argument to `MPI_EXSCAN`.
15. Section 6.4.2 on page 232, and Section 6.6 on page 248.  
Implementations that did not implement `MPI_COMM_CREATE` on intercommunicators will need to add that functionality. As the standard described the behavior of this operation on intercommunicators, it is believed that most implementations already provide this functionality. Note also that the C++ binding for both `MPI_COMM_CREATE` and `MPI_COMM_SPLIT` explicitly allow Intercomms.
16. Section 6.4.2 on page 232.  
`MPI_COMM_CREATE` is extended to allow several disjoint subgroups as input if `comm` is an intracommunicator. If `comm` is an intercommunicator it was clarified that all processes in the same local group of `comm` must specify the same value for group.
17. Section 7.5.4 on page 284.  
New functions for a scalable distributed graph topology interface has been added. In this section, the functions `MPI_DIST_GRAPH_CREATE_ADJACENT` and `MPI_DIST_GRAPH_CREATE`, the constants `MPI_UNWEIGHTED`, and the derived C++ class `Distgraphcomm` were added.
18. Section 7.5.5 on page 290.  
For the scalable distributed graph topology interface, the functions `MPI_DIST_NEIGHBORS_COUNT` and `MPI_DIST_NEIGHBORS` and the constant `MPI_DIST_GRAPH` were added.
19. Section 7.5.5 on page 290.  
Remove ambiguity regarding duplicated neighbors with `MPI_GRAPH_NEIGHBORS` and `MPI_GRAPH_NEIGHBORS_COUNT`.
20. Section 8.1.1 on page 303.  
The subversion number changed from 1 to 2.
21. Section 8.3 on page 308, Section 15.2 on page 497, and Annex A.1.3 on page 557.  
Changed function pointer typedef names `MPI_{Comm,File,Win}_errhandler_fn` to `MPI_{Comm,File,Win}_errhandler_function`. Deprecated old “\_fn” names.
22. Section 8.7.1 on page 327.  
Attribute deletion callbacks on `MPI_COMM_SELF` are now called in LIFO order. Implementors must now also register all implementation-internal attribute deletion callbacks on `MPI_COMM_SELF` before returning from `MPI_INIT/MPI_INIT_THREAD`.
23. Section 11.3.4 on page 377.  
The restriction added in MPI 2.1 that the operation `MPI_REPLACE` in `MPI_ACCUMULATE` can be used only with predefined datatypes has been removed. `MPI_REPLACE` can now be used even with derived datatypes, as it was in MPI 2.0. Also, a clarification has been made that `MPI_REPLACE` can be used only in

MPI\_ACCUMULATE, not in collective operations that do reductions, such as MPI\_REDUCE and others.

24. Section 12.2 on page 405.  
Add “\*” to the `query_fn`, `free_fn`, and `cancel_fn` arguments to the C++ binding for `MPI::Grequest::Start()` for consistency with the rest of MPI functions that take function pointer arguments.
25. Section 13.5.2 on page 463, and Table 13.2 on page 465.  
`MPI_(U)INT{8,16,32,64}_T`, `MPI_AINT`, `MPI_OFFSET`, `MPI_C_COMPLEX`, `MPI_C_FLOAT_COMPLEX`, `MPI_C_DOUBLE_COMPLEX`, `MPI_C_LONG_DOUBLE_COMPLEX`, and `MPI_C_BOOL` are added as predefined datatypes in the external32 representation.
26. Section 16.3.7 on page 537.  
The description was modified that it only describes how an MPI implementation behaves, but not how MPI stores attributes internally. The erroneous MPI-2.1 Example 16.17 was replaced with three new examples 16.17, 16.18, and 16.19 on pages 538-540 explicitly detailing cross-language attribute behavior. Implementations that matched the behavior of the old example will need to be updated.
27. Annex A.1.1 on page 545.  
Removed type `MPI::Fint` (compare `MPI_Fint` in Section A.1.2 on page 556).
28. Annex A.1.1 on page 545. Table *Named Predefined Datatypes*.  
Added `MPI_(U)INT{8,16,32,64}_T`, `MPI_AINT`, `MPI_OFFSET`, `MPI_C_BOOL`, `MPI_C_FLOAT_COMPLEX`, `MPI_C_COMPLEX`, `MPI_C_DOUBLE_COMPLEX`, and `MPI_C_LONG_DOUBLE_COMPLEX` are added as predefined datatypes.

### B.3 Changes from Version 2.0 to Version 2.1

1. Section 3.2.2 on page 29, Section 16.1.6 on page 503, and Annex A.1 on page 545.  
In addition, the `MPI_LONG_LONG` should be added as an optional type; it is a synonym for `MPI_LONG_LONG_INT`.
2. Section 3.2.2 on page 29, Section 16.1.6 on page 503, and Annex A.1 on page 545.  
`MPI_LONG_LONG_INT`, `MPI_LONG_LONG` (as synonym), `MPI_UNSIGNED_LONG_LONG`, `MPI_SIGNED_CHAR`, and `MPI_WCHAR` are moved from optional to official and they are therefore defined for all three language bindings.
3. Section 3.2.5 on page 33.  
`MPI_GET_COUNT` with zero-length datatypes: The value returned as the `count` argument of `MPI_GET_COUNT` for a datatype of length zero where zero bytes have been transferred is zero. If the number of bytes transferred is greater than zero, `MPI_UNDEFINED` is returned.
4. Section 4.1 on page 79.  
General rule about derived datatypes: Most datatype constructors have replication count or block length arguments. Allowed values are non-negative integers. If the value is zero, no elements are generated in the type map and there is no effect on datatype bounds or extent.

5. Section 4.3 on page 130.  
MPI\_BYTE should be used to send and receive data that is packed using MPI\_PACK\_EXTERNAL.
6. Section 5.9.6 on page 182.  
If comm is an intercommunicator in MPI\_ALLREDUCE, then both groups should provide count and datatype arguments that specify the same type signature (i.e., it is not necessary that both groups provide the same count value).
7. Section 6.3.1 on page 224.  
MPI\_GROUP\_TRANSLATE\_RANKS and MPI\_PROC\_NULL: MPI\_PROC\_NULL is a valid rank for input to MPI\_GROUP\_TRANSLATE\_RANKS, which returns MPI\_PROC\_NULL as the translated rank.
8. Section 6.7 on page 256.  
About the attribute caching functions:  
  

*Advice to implementors.* High-quality implementations should raise an error when a keyval that was created by a call to MPI\_XXX\_CREATE\_KEYVAL is used with an object of the wrong type with a call to MPI\_YYY\_GET\_ATTR, MPI\_YYY\_SET\_ATTR, MPI\_YYY\_DELETE\_ATTR, or MPI\_YYY\_FREE\_KEYVAL. To do so, it is necessary to maintain, with each keyval, information on the type of the associated user function. (*End of advice to implementors.*)
9. Section 6.8 on page 270.  
In MPI\_COMM\_GET\_NAME: In C, a null character is additionally stored at name[resultlen]. resultlen cannot be larger than MPI\_MAX\_OBJECT\_NAME-1. In Fortran, name is padded on the right with blank characters. resultlen cannot be larger than MPI\_MAX\_OBJECT\_NAME.
10. Section 7.4 on page 278.  
About MPI\_GRAPH\_CREATE and MPI\_CART\_CREATE: All input arguments must have identical values on all processes of the group of comm\_old.
11. Section 7.5.1 on page 280.  
In MPI\_CART\_CREATE: If ndims is zero then a zero-dimensional Cartesian topology is created. The call is erroneous if it specifies a grid that is larger than the group size or if ndims is negative.
12. Section 7.5.3 on page 282.  
In MPI\_GRAPH\_CREATE: If the graph is empty, i.e., nnodes == 0, then MPI\_COMM\_NULL is returned in all processes.
13. Section 7.5.3 on page 282.  
In MPI\_GRAPH\_CREATE: A single process is allowed to be defined multiple times in the list of neighbors of a process (i.e., there may be multiple edges between two processes). A process is also allowed to be a neighbor to itself (i.e., a self loop in the graph). The adjacency matrix is allowed to be non-symmetric.

*Advice to users.* Performance implications of using multiple edges or a non-symmetric adjacency matrix are not defined. The definition of a node-neighbor edge does not imply a direction of the communication. (*End of advice to users.*)

14. Section 7.5.5 on page 290.  
In `MPI_CARTDIM_GET` and `MPI_CART_GET`: If `comm` is associated with a zero-dimensional Cartesian topology, `MPI_CARTDIM_GET` returns `ndims=0` and `MPI_CART_GET` will keep all output arguments unchanged.
15. Section 7.5.5 on page 290.  
In `MPI_CART_RANK`: If `comm` is associated with a zero-dimensional Cartesian topology, `coord` is not significant and 0 is returned in `rank`.
16. Section 7.5.5 on page 290.  
In `MPI_CART_COORDS`: If `comm` is associated with a zero-dimensional Cartesian topology, `coords` will be unchanged.
17. Section 7.5.6 on page 297.  
In `MPI_CART_SHIFT`: It is erroneous to call `MPI_CART_SHIFT` with a direction that is either negative or greater than or equal to the number of dimensions in the Cartesian communicator. This implies that it is erroneous to call `MPI_CART_SHIFT` with a `comm` that is associated with a zero-dimensional Cartesian topology.
18. Section 7.5.7 on page 299.  
In `MPI_CART_SUB`: If all entries in `remain_dims` are false or `comm` is already associated with a zero-dimensional Cartesian topology then `newcomm` is associated with a zero-dimensional Cartesian topology.
- 18.1. Section 8.1.1 on page 303.  
The subversion number changed from 0 to 1.
19. Section 8.1.2 on page 304.  
In `MPI_GET_PROCESSOR_NAME`: In C, a null character is additionally stored at `name[resultlen]`. `resultlen` cannot be larger than `MPI_MAX_PROCESSOR_NAME-1`. In Fortran, `name` is padded on the right with blank characters. `resultlen` cannot be larger than `MPI_MAX_PROCESSOR_NAME`.
20. Section 8.3 on page 308.  
`MPI_{COMM,WIN,FILE}_GET_ERRHANDLER` behave as if a new error handler object is created. That is, once the error handler is no longer needed, `MPI_ERRHANDLER_FREE` should be called with the error handler returned from `MPI_ERRHANDLER_GET` or `MPI_{COMM,WIN,FILE}_GET_ERRHANDLER` to mark the error handler for deallocation. This provides behavior similar to that of `MPI_COMM_GROUP` and `MPI_GROUP_FREE`.
21. Section 8.7 on page 322, see explanations to `MPI_FINALIZE`.  
`MPI_FINALIZE` is collective over all connected processes. If no processes were spawned, accepted or connected then this means over `MPI_COMM_WORLD`; otherwise it is collective over the union of all processes that have been and continue to be connected, as explained in Section 10.5.4 on page 362.

22. Section 8.7 on page 322.

About MPI\_ABORT:

*Advice to users.* Whether the errorcode is returned from the executable or from the MPI process startup mechanism (e.g., mpiexec), is an aspect of quality of the MPI library but not mandatory. (*End of advice to users.*)

*Advice to implementors.* Where possible, a high-quality implementation will try to return the errorcode from the MPI process startup mechanism (e.g. mpiexec or singleton init). (*End of advice to implementors.*)

23. Section 9 on page 331.

An implementation must support info objects as caches for arbitrary (key, value) pairs, regardless of whether it recognizes the key. Each function that takes hints in the form of an MPI\_Info must be prepared to ignore any key it does not recognize. This description of info objects does not attempt to define how a particular function should react if it recognizes a key but not the associated value. MPI\_INFO\_GET\_NKEYS, MPI\_INFO\_GET\_NTHKEY, MPI\_INFO\_GET\_VALUELEN, and MPI\_INFO\_GET must retain all (key,value) pairs so that layered functionality can also use the Info object.

24. Section 11.3 on page 371.

MPI\_PROC\_NULL is a valid target rank in the MPI RMA calls MPI\_ACCUMULATE, MPI\_GET, and MPI\_PUT. The effect is the same as for MPI\_PROC\_NULL in MPI point-to-point communication. See also item 25 in this list.

25. Section 11.3 on page 371.

After any RMA operation with rank MPI\_PROC\_NULL, it is still necessary to finish the RMA epoch with the synchronization method that started the epoch. See also item 24 in this list.

26. Section 11.3.4 on page 377.

MPI\_REPLACE in MPI\_ACCUMULATE, like the other predefined operations, is defined only for the predefined MPI datatypes.

27. Section 13.2.8 on page 430.

About MPI\_FILE\_SET\_VIEW and MPI\_FILE\_SET\_INFO: When an info object that specifies a subset of valid hints is passed to MPI\_FILE\_SET\_VIEW or MPI\_FILE\_SET\_INFO, there will be no effect on previously set or defaulted hints that the info does not specify.

28. Section 13.2.8 on page 430.

About MPI\_FILE\_GET\_INFO: If no hint exists for the file associated with fh, a handle to a newly created info object is returned that contains no key/value pair.

29. Section 13.3 on page 433.

If a file does not have the mode MPI\_MODE\_SEQUENTIAL, then MPI\_DISPLACEMENT\_CURRENT is invalid as disp in MPI\_FILE\_SET\_VIEW.

30. Section 13.5.2 on page 463.

The bias of 16 byte doubles was defined with 10383. The correct value is 16383.



31. Section 16.1.4 on page 500.

In the example in this section, the buffer should be declared as `const void* buf`.

32. Section 16.2.5 on page 521.

About `MPI_TYPE_CREATE_F90_XXXX`:

*Advice to implementors.* An application may often repeat a call to `MPI_TYPE_CREATE_F90_XXXX` with the same combination of `(XXXX,p,r)`. The application is not allowed to free the returned predefined, unnamed datatype handles. To prevent the creation of a potentially huge amount of handles, the MPI implementation should return the same datatype handle for the same `(REAL/COMPLEX/INTEGER,p,r)` combination. Checking for the combination `(p,r)` in the preceding call to `MPI_TYPE_CREATE_F90_XXXX` and using a hash-table to find formerly generated handles should limit the overhead of finding a previously generated datatype with same combination of `(XXXX,p,r)`. (*End of advice to implementors.*)

33. Section A.1.1 on page 545.

`MPI_BOTTOM` is defined as `void * const MPI::BOTTOM`.