

MPI: A Message-Passing Interface Standard

Version 3.0

ticket0.

Message Passing Interface Forum

Draft July 4th, 2011

Contents

1	MPI Environmental Management	1
1.1	Implementation Information	1
1.1.1	Version Inquiries	1
1.1.2	Environmental Inquiries	2
	Tag Values	2
	Host Rank	2
	IO Rank	3
	Clock Synchronization	3
1.2	Memory Allocation	4
1.3	Error Handling	6
1.3.1	Error Handlers for Communicators	8
1.3.2	Error Handlers for Windows	10
1.3.3	Error Handlers for Files	11
1.3.4	Freeing Errorhandlers and Retrieving Error Strings	12
1.4	Error Codes and Classes	13
1.5	Error Classes, Error Codes, and Error Handlers	16
1.6	Timers and Synchronization	19
1.7	Startup	20
1.7.1	Allowing User Functions at Process Termination	25
1.7.2	Determining Whether MPI Has Finished	26
1.8	Portable MPI Process Startup	27
2	Tool Support	30
2.1	Introduction	30
2.2	Profiling Interface	30
2.2.1	Requirements	30
2.2.2	Discussion	31
2.2.3	Logic of the Design	31
2.2.4	Miscellaneous Control of Profiling	32
2.2.5	Profiler Implementation ¶Example	33
2.2.6	MPI Library Implementation ¶Example	33
	Systems with Weak Symbols	33
	Systems Without Weak Symbols	34
2.2.7	Complications	34
	Multiple Counting	34
	Linker Oddities	35
2.2.8	Multiple Levels of Interception	35

ticket266.
ticket266.

2.3	The MPI Tool Information Interface	36	
2.3.1	Verbosity Levels	37	
2.3.2	Binding MPI Tool Information Interface Variables to MPI Objects	37	
2.3.3	Convention for Returning Strings	38	
2.3.4	Initialization and Finalization	38	
2.3.5	Datatype System	40	
2.3.6	Control Variables	42	
	Control Variable Query Functions	42	
	Example: Printing All Control Variables	44	
	Handle Allocation and Deallocation	45	
	Control Variable Access Functions	47	
	Example: Reading the Value of a Control Variable	47	
2.3.7	Performance Variables	48	
	Performance Variable Classes	48	
	Performance Variable Query Functions	51	
	Performance Experiment Sessions	53	
	Handle Allocation and Deallocation	54	
	Starting and Stopping of Performance Variables	55	
	Performance Variable Access Functions	56	
	Example: Tool to Detect Receives with Long Unexpected Message Queues	58	
2.3.8	Variable Categorization	60	
2.3.9	Return Codes for the MPI tool information interface	63	
2.3.10	Profiling Interface	63	
3	Language Bindings Summary	65	
3.1	Defined Values and Handles	65	
3.1.1	Defined Constants	65	
3.1.2	Types	78	
3.1.3	Prototype [d]Definitions	79	ticket0.
3.1.4	Deprecated [p]Prototype [d]Definitions	83	ticket0.
3.1.5	Info Keys	83	ticket0.
3.1.6	Info Values	84	
	Bibliography	85	
	Examples Index	86	
	MPI Constant and Predefined Handle Index	87	
	MPI Declarations Index	94	
	MPI Callback Function Prototype Index	95	
	MPI Function Index	96	

List of Figures

List of Tables

1.1	Error classes (Part 1)	14
1.2	Error classes (Part 2)	15
2.1	MPI tool information interface verbosity levels.	37
2.2	Constants to identify associations of variables.	38
2.3	MPI datatypes that can be used by the MPI tool information interface.	40
2.4	Scopes for control variables.	44
2.5	Return codes used in functions of the MPI tool information interface.	64

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48

Chapter 1

MPI Environmental Management

This chapter discusses routines for getting and, where appropriate, setting various parameters that relate to the MPI implementation and the execution environment (such as error handling). The procedures for entering and leaving the MPI execution environment are also described here.

1.1 Implementation Information

1.1.1 Version Inquiries

In order to cope with changes to the MPI Standard, there are both compile-time and runtime ways to determine which version of the standard is in use in the environment one is using.

The “version” will be represented by two separate integers, for the version and subversion: In C and C++,

```
#define MPI_VERSION    2
#define MPI_SUBVERSION 2
```

in Fortran,

```
INTEGER MPI_VERSION, MPI_SUBVERSION
PARAMETER (MPI_VERSION    = 2)
PARAMETER (MPI_SUBVERSION = 2)
```

For runtime determination,

`MPI_GET_VERSION(version, subversion)`

OUT	version	version number (integer)
OUT	subversion	subversion number (integer)

```
int MPI_Get_version(int *version, int *subversion)
```

```
MPI_GET_VERSION(VERSION, SUBVERSION, IERROR)
  INTEGER VERSION, SUBVERSION, IERROR
```

```
1 {void MPI::Get_version(int& version, int& subversion) (binding deprecated, see  
2 Section ??) }
```

3
4 MPI_GET_VERSION is one of the few functions that can be called before MPI_INIT and
5 after MPI_FINALIZE. Valid (MPI_VERSION, MPI_SUBVERSION) pairs in this and previous
6 versions of the MPI standard are (2,2), (2,1), (2,0), and (1,2).

7 8 1.1.2 Environmental Inquiries

9 A set of attributes that describe the execution environment are attached to the commu-
10 nicator MPI_COMM_WORLD when MPI is initialized. The value of these attributes can be
11 inquired by using the function MPI_COMM_GET_ATTR described in Chapter ?? . It is
12 erroneous to delete these attributes, free their keys, or change their values.

13 The list of predefined attribute keys include

14
15 **MPI_TAG_UB** Upper bound for tag value.

16
17 **MPI_HOST** Host process rank, if such exists, MPI_PROC_NULL, otherwise.

18
19 **MPI_IO** rank of a node that has regular I/O facilities (possibly myrank). Nodes in the same
20 communicator may return different values for this parameter.

21
22 **MPI_WTIME_IS_GLOBAL** Boolean variable that indicates whether clocks are synchronized.

23 Vendors may add implementation specific parameters (such as node number, real mem-
24 ory size, virtual memory size, etc.)

25 These predefined attributes do not change value between MPI initialization (MPI_INIT
26 and MPI completion (MPI_FINALIZE), and cannot be updated or deleted by users.

27
28 *Advice to users.* Note that in the C binding, the value returned by these attributes
29 is a *pointer* to an `int` containing the requested value. (*End of advice to users.*)

30
31 The required parameter values are discussed in more detail below:

32 33 Tag Values

34 Tag values range from 0 to the value returned for MPI_TAG_UB inclusive. These values are
35 guaranteed to be unchanging during the execution of an MPI program. In addition, the tag
36 upper bound value must be *at least* 32767. An MPI implementation is free to make the
37 value of MPI_TAG_UB larger than this; for example, the value $2^{30} - 1$ is also a legal value
38 for MPI_TAG_UB.

39 The attribute MPI_TAG_UB has the same value on all processes of MPI_COMM_WORLD.

41 42 Host Rank

43 The value returned for MPI_HOST gets the rank of the HOST process in the group associated
44 with communicator MPI_COMM_WORLD, if there is such. MPI_PROC_NULL is returned if
45 there is no host. MPI does not specify what it means for a process to be a HOST, nor does
46 it requires that a HOST exists.

47 The attribute MPI_HOST has the same value on all processes of MPI_COMM_WORLD.

IO Rank

The value returned for `MPI_IO` is the rank of a processor that can provide language-standard I/O facilities. For Fortran, this means that all of the Fortran I/O operations are supported (e.g., `OPEN`, `REWIND`, `WRITE`). For C and C++, this means that all of the ISO C and C++, I/O operations are supported (e.g., `fopen`, `fprintf`, `lseek`).

If every process can provide language-standard I/O, then the value `MPI_ANY_SOURCE` will be returned. Otherwise, if the calling process can provide language-standard I/O, then its rank will be returned. Otherwise, if some process can provide language-standard I/O then the rank of one such process will be returned. The same value need not be returned by all processes. If no process can provide language-standard I/O, then the value `MPI_PROC_NULL` will be returned.

Advice to users. Note that input is not collective, and this attribute does *not* indicate which process can or does provide input. (*End of advice to users.*)

Clock Synchronization

The value returned for `MPI_WTIME_IS_GLOBAL` is 1 if clocks at all processes in `MPI_COMM_WORLD` are synchronized, 0 otherwise. A collection of clocks is considered synchronized if explicit effort has been taken to synchronize them. The expectation is that the variation in time, as measured by calls to `MPI_WTIME`, will be less than one half the round-trip time for an MPI message of length zero. If time is measured at a process just before a send and at another process just after a matching receive, the second time should be always higher than the first one.

The attribute `MPI_WTIME_IS_GLOBAL` need not be present when the clocks are not synchronized (however, the attribute key `MPI_WTIME_IS_GLOBAL` is always valid). This attribute may be associated with communicators other than `MPI_COMM_WORLD`.

The attribute `MPI_WTIME_IS_GLOBAL` has the same value on all processes of `MPI_COMM_WORLD`.

`MPI_GET_PROCESSOR_NAME(name, resultlen)`

OUT	name	A unique specifier for the actual (as opposed to virtual) node.
OUT	resultlen	Length (in printable characters) of the result returned in name

```
int MPI_Get_processor_name(char *name, int *resultlen)
```

```
MPI_GET_PROCESSOR_NAME( NAME, RESULTLEN, IERROR)
```

```
CHARACTER*(*) NAME
```

```
INTEGER RESULTLEN, IERROR
```

```
{void MPI::Get_processor_name(char* name, int& resultlen) (binding deprecated,  
see Section ??) }
```

This routine returns the name of the processor on which it was called at the moment of the call. The name is a character string for maximum flexibility. From this value it must be possible to identify a specific piece of hardware; possible values include “processor

9 in rack 4 of mpp.cs.org” and “231” (where 231 is the actual processor number in the running homogeneous system). The argument `name` must represent storage that is at least `MPI_MAX_PROCESSOR_NAME` characters long. `MPI_GET_PROCESSOR_NAME` may write up to this many characters into `name`.

The number of characters actually written is returned in the output argument, `resultlen`. In C, a null character is additionally stored at `name[resultlen]`. The `resultlen` cannot be larger than `MPI_MAX_PROCESSOR_NAME-1`. In Fortran, `name` is padded on the right with blank characters. The `resultlen` cannot be larger than `MPI_MAX_PROCESSOR_NAME`.

Rationale. This function allows MPI implementations that do process migration to return the current processor. Note that nothing in MPI *requires* or defines process migration; this definition of `MPI_GET_PROCESSOR_NAME` simply allows such an implementation. (*End of rationale.*)

Advice to users. The user must provide at least `MPI_MAX_PROCESSOR_NAME` space to write the processor name — processor names can be this long. The user should examine the output argument, `resultlen`, to determine the actual length of the name. (*End of advice to users.*)

The constant `MPI_BSEND_OVERHEAD` provides an upper bound on the fixed overhead per message buffered by a call to `MPI_BSEND` (see Section ??).

1.2 Memory Allocation

In some systems, message-passing and remote-memory-access (RMA) operations run faster when accessing specially allocated memory (e.g., memory that is shared by the other processes in the communicating group on an SMP). MPI provides a mechanism for allocating and freeing such special memory. The use of such memory for message-passing or RMA is not mandatory, and this memory can be used without restrictions as any other dynamically allocated memory. However, implementations may restrict the use of the `MPI_WIN_LOCK` and `MPI_WIN_UNLOCK` functions to windows allocated in such memory (see Section ??).

`MPI_ALLOC_MEM(size, info, baseptr)`

IN	size	size of memory segment in bytes (non-negative integer)
IN	info	info argument (handle)
OUT	baseptr	pointer to beginning of memory segment allocated

`int MPI_Alloc_mem(MPI_Aint size, MPI_Info info, void *baseptr)`

`MPI_ALLOC_MEM(SIZE, INFO, BASEPTR, IERROR)`

INTEGER INFO, IERROR

INTEGER(KIND=MPI_ADDRESS_KIND) SIZE, BASEPTR

{void* MPI::Alloc_mem(MPI::Aint size, const MPI::Info& info) (*binding deprecated, see Section ??*) }

The `info` argument can be used to provide directives that control the desired location of the allocated memory. Such a directive does not affect the semantics of the call. Valid `info` values are implementation-dependent; a null directive value of `info = MPI_INFO_NULL` is always valid.

The function `MPI_ALLOC_MEM` may return an error code of class `MPI_ERR_NO_MEM` to indicate it failed because memory is exhausted.

`MPI_FREE_MEM(base)`

IN	base	initial address of memory segment allocated by <code>MPI_ALLOC_MEM</code> (choice)
----	------	---

`int MPI_Free_mem(void *base)`

`MPI_FREE_MEM(BASE, IERROR)`

`<type> BASE(*)`

`INTEGER IERROR`

`{void MPI::Free_mem(void *base) (binding deprecated, see Section ??) }`

The function `MPI_FREE_MEM` may return an error code of class `MPI_ERR_BASE` to indicate an invalid base argument.

Rationale. The C and C++ bindings of `MPI_ALLOC_MEM` and `MPI_FREE_MEM` are similar to the bindings for the `malloc` and `free` C library calls: a call to `MPI_Alloc_mem(..., &base)` should be paired with a call to `MPI_Free_mem(base)` (one less level of indirection). Both arguments are declared to be of same type `void*` so as to facilitate type casting. The Fortran binding is consistent with the C and C++ bindings: the Fortran `MPI_ALLOC_MEM` call returns in `baseptr` the (integer valued) address of the allocated memory. The `base` argument of `MPI_FREE_MEM` is a choice argument, which passes (a reference to) the variable stored at that location. (*End of rationale.*)

Advice to implementors. If `MPI_ALLOC_MEM` allocates special memory, then a design similar to the design of C `malloc` and `free` functions has to be used, in order to find out the size of a memory segment, when the segment is freed. If no special memory is used, `MPI_ALLOC_MEM` simply invokes `malloc`, and `MPI_FREE_MEM` invokes `free`.

A call to `MPI_ALLOC_MEM` can be used in shared memory systems to allocate memory in a shared memory segment. (*End of advice to implementors.*)

Example 1.1

Example of use of `MPI_ALLOC_MEM`, in Fortran with pointer support. We assume 4-byte REALs, and assume that pointers are address-sized.

```

1  REAL A
2  POINTER (P, A(100,100))    ! no memory is allocated
3  CALL MPI_ALLOC_MEM(4*100*100, MPI_INFO_NULL, P, IERR)
4  ! memory is allocated
5  ...
6  A(3,5) = 2.71;
7  ...
8  CALL MPI_FREE_MEM(A, IERR) ! memory is freed
9

```

Since standard Fortran does not support (C-like) pointers, this code is not Fortran 77 or Fortran 90 code. Some compilers (in particular, at the time of writing, g77 and Fortran compilers for Intel) do not support this code.

Example 1.2 Same example, in C

```

15 float (*f)[100][100] ;
16 /* no memory is allocated */
17 MPI_Alloc_mem(sizeof(float)*100*100, MPI_INFO_NULL, &f);
18 /* memory allocated */
19 ...
20 (*f)[5][3] = 2.71;
21 ...
22 MPI_Free_mem(f);
23

```

1.3 Error Handling

An MPI implementation cannot or may choose not to handle some errors that occur during MPI calls. These can include errors that generate exceptions or traps, such as floating point errors or access violations. The set of errors that are handled by MPI is implementation-dependent. Each such error generates an **MPI exception**.

The above text takes precedence over any text on error handling within this document. Specifically, text that states that errors *will* be handled should be read as *may* be handled.

A user can associate error handlers to three types of objects: communicators, windows, and files. The specified error handling routine will be used for any MPI exception that occurs during a call to MPI for the respective object. MPI calls that are not related to any objects are considered to be attached to the communicator MPI_COMM_WORLD. The attachment of error handlers to objects is purely local: different processes may attach different error handlers to corresponding objects.

Several predefined error handlers are available in MPI:

MPI_ERRORS_ARE_FATAL The handler, when called, causes the program to abort on all executing processes. This has the same effect as if MPI_ABORT was called by the process that invoked the handler.

MPI_ERRORS_RETURN The handler has no effect other than returning the error code to the user.

Implementations may provide additional predefined error handlers and programmers can code their own error handlers.

The error handler `MPI_ERRORS_ARE_FATAL` is associated by default with `MPI_COMM_WORLD` after initialization. Thus, if the user chooses not to control error handling, every error that MPI handles is treated as fatal. Since (almost) all MPI calls return an error code, a user may choose to handle errors in its main code, by testing the return code of MPI calls and executing a suitable recovery code when the call was not successful. In this case, the error handler `MPI_ERRORS_RETURN` will be used. Usually it is more convenient and more efficient not to test for errors after each MPI call, and have such error handled by a non trivial MPI error handler.

After an error is detected, the state of MPI is undefined. That is, using a user-defined error handler, or `MPI_ERRORS_RETURN`, does *not* necessarily allow the user to continue to use MPI after an error is detected. The purpose of these error handlers is to allow a user to issue user-defined error messages and to take actions unrelated to MPI (such as flushing I/O buffers) before a program exits. An MPI implementation is free to allow MPI to continue after an error but is not required to do so.

Advice to implementors. A good quality implementation will, to the greatest possible extent, circumscribe the impact of an error, so that normal processing can continue after an error handler was invoked. The implementation documentation will provide information on the possible effect of each class of errors. (*End of advice to implementors.*)

An MPI error handler is an opaque object, which is accessed by a handle. MPI calls are provided to create new error handlers, to associate error handlers with objects, and to test which error handler is associated with an object. C and C++ have distinct typedefs for user defined error handling callback functions that accept communicator, file, and window arguments. In Fortran there are three user routines.

An error handler object is created by a call to `MPI_XXX_CREATE_ERRHANDLER(function, errhandler)`, where XXX is, respectively, `COMM`, `WIN`, or `FILE`.

An error handler is attached to a communicator, window, or file by a call to `MPI_XXX_SET_ERRHANDLER`. The error handler must be either a predefined error handler, or an error handler that was created by a call to `MPI_XXX_CREATE_ERRHANDLER`, with matching XXX. The predefined error handlers `MPI_ERRORS_RETURN` and `MPI_ERRORS_ARE_FATAL` can be attached to communicators, windows, and files. In C++, the predefined error handler `MPI::ERRORS_THROW_EXCEPTIONS` can also be attached to communicators, windows, and files.

The error handler currently associated with a communicator, window, or file can be retrieved by a call to `MPI_XXX_GET_ERRHANDLER`.

The MPI function `MPI_ERRHANDLER_FREE` can be used to free an error handler that was created by a call to `MPI_XXX_CREATE_ERRHANDLER`.

`MPI_{COMM,WIN,FILE}_GET_ERRHANDLER` behave as if a new error handler object is created. That is, once the error handler is no longer needed, `MPI_ERRHANDLER_FREE` should be called with the error handler returned from `MPI_ERRHANDLER_GET` or `MPI_{COMM,WIN,FILE}_GET_ERRHANDLER` to mark the error handler for deallocation. This provides behavior similar to that of `MPI_COMM_GROUP` and `MPI_GROUP_FREE`.

Advice to implementors. High-quality implementation should raise an error when an error handler that was created by a call to `MPI_XXX_CREATE_ERRHANDLER` is attached to an object of the wrong type with a call to `MPI_YYY_SET_ERRHANDLER`.

To do so, it is necessary to maintain, with each error handler, information on the typedef of the associated user function. (*End of advice to implementors.*)

The syntax for these calls is given below.

1.3.1 Error Handlers for Communicators

MPI_COMM_CREATE_ERRHANDLER(function, errhandler)

IN	function	user defined error handling procedure (function)
OUT	errhandler	MPI error handler (handle)

```
int MPI_Comm_create_errhandler(MPI_Comm_errhandler_function *function,
                               MPI_Errhandler *errhandler)
```

MPI_COMM_CREATE_ERRHANDLER(FUNCTION, ERRHANDLER, IERROR)

EXTERNAL FUNCTION

INTEGER ERRHANDLER, IERROR

```
{static MPI::Errhandler
    MPI::Comm::Create_errhandler(MPI::Comm::Errhandler_function*
    function) (binding deprecated, see Section ??) }
```

Creates an error handler that can be attached to communicators. This function is identical to **MPI_ERRHANDLER_CREATE**, whose use is deprecated.

The user routine should be, in C, a function of type **MPI_Comm_errhandler_function**, which is defined as

```
typedef void MPI_Comm_errhandler_function(MPI_Comm *, int *, ...);
```

The first argument is the communicator in use. The second is the error code to be returned by the MPI routine that raised the error. If the routine would have returned **MPI_ERR_IN_STATUS**, it is the error code returned in the status for the request that caused the error handler to be invoked. The remaining arguments are “stdargs” arguments whose number and meaning is implementation-dependent. An implementation should clearly document these arguments. Addresses are used so that the handler may be written in Fortran. This typedef replaces **MPI_Handler_function**, whose use is deprecated.

In Fortran, the user routine should be of the form:

```
SUBROUTINE COMM_ERRHANDLER_FUNCTION(COMM, ERROR_CODE)
```

```
    INTEGER COMM, ERROR_CODE
```

In C++, the user routine should be of the form:

```
{typedef void MPI::Comm::Errhandler_function(MPI::Comm &, int *, ...);
    (binding deprecated, see Section ??)}
```

Rationale. The variable argument list is provided because it provides an ISO-standard hook for providing additional information to the error handler; without this hook, ISO C prohibits additional arguments. (*End of rationale.*)

Advice to users. A newly created communicator inherits the error handler that is associated with the “parent” communicator. In particular, the user can specify a “global” error handler for all communicators by associating this handler with the communicator MPI_COMM_WORLD immediately after initialization. (*End of advice to users.*)

MPI_COMM_SET_ERRHANDLER(comm, errhandler)

INOUT	comm	communicator (handle)
IN	errhandler	new error handler for communicator (handle)

int MPI_Comm_set_errhandler(MPI_Comm comm, MPI_Errhandler errhandler)

MPI_COMM_SET_ERRHANDLER(COMM, ERRHANDLER, IERROR)

INTEGER COMM, ERRHANDLER, IERROR

{void MPI::Comm::Set_errhandler(const MPI::Errhandler& errhandler) (*binding deprecated, see Section ??*) }

Attaches a new error handler to a communicator. The error handler must be either a predefined error handler, or an error handler created by a call to MPI_COMM_CREATE_ERRHANDLER. This call is identical to MPI_ERRHANDLER_SET, whose use is deprecated.

MPI_COMM_GET_ERRHANDLER(comm, errhandler)

IN	comm	communicator (handle)
OUT	errhandler	error handler currently associated with communicator (handle)

int MPI_Comm_get_errhandler(MPI_Comm comm, MPI_Errhandler *errhandler)

MPI_COMM_GET_ERRHANDLER(COMM, ERRHANDLER, IERROR)

INTEGER COMM, ERRHANDLER, IERROR

{MPI::Errhandler MPI::Comm::Get_errhandler() const (*binding deprecated, see Section ??*) }

Retrieves the error handler currently associated with a communicator. This call is identical to MPI_ERRHANDLER_GET, whose use is deprecated.

Example: A library function may register at its entry point the current error handler for a communicator, set its own private error handler for this communicator, and restore before exiting the previous error handler.

1.3.2 Error Handlers for Windows

MPI_WIN_CREATE_ERRHANDLER(function, errhandler)

IN	function	user defined error handling procedure (function)
OUT	errhandler	MPI error handler (handle)

```
int MPI_Win_create_errhandler(MPI_Win_errhandler_function *function,
                             MPI_Errhandler *errhandler)
```

MPI_WIN_CREATE_ERRHANDLER(FUNCTION, ERRHANDLER, IERROR)

EXTERNAL FUNCTION

INTEGER ERRHANDLER, IERROR

```
{static MPI::Errhandler
    MPI::Win::Create_errhandler(MPI::Win::Errhandler_function*
                                function) (binding deprecated, see Section ??) }
```

Creates an error handler that can be attached to a window object. The user routine should be, in C, a function of type MPI_Win_errhandler_function which is defined as

```
typedef void MPI_Win_errhandler_function(MPI_Win *, int *, ...);
```

The first argument is the window in use, the second is the error code to be returned.

In Fortran, the user routine should be of the form:

```
SUBROUTINE WIN_ERRHANDLER_FUNCTION(WIN, ERROR_CODE)
    INTEGER WIN, ERROR_CODE
```

In C++, the user routine should be of the form:

```
{typedef void MPI::Win::Errhandler_function(MPI::Win &, int *, ...);
    (binding deprecated, see Section ??)}
```

MPI_WIN_SET_ERRHANDLER(win, errhandler)

INOUT	win	window (handle)
IN	errhandler	new error handler for window (handle)

```
int MPI_Win_set_errhandler(MPI_Win win, MPI_Errhandler errhandler)
```

MPI_WIN_SET_ERRHANDLER(WIN, ERRHANDLER, IERROR)

INTEGER WIN, ERRHANDLER, IERROR

```
{void MPI::Win::Set_errhandler(const MPI::Errhandler& errhandler) (binding
    deprecated, see Section ??) }
```

Attaches a new error handler to a window. The error handler must be either a pre-defined error handler, or an error handler created by a call to MPI_WIN_CREATE_ERRHANDLER.

MPI_WIN_GET_ERRHANDLER(win, errhandler)

IN	win	window (handle)
OUT	errhandler	error handler currently associated with window (handle)

int MPI_Win_get_errhandler(MPI_Win win, MPI_Errhandler *errhandler)

MPI_WIN_GET_ERRHANDLER(WIN, ERRHANDLER, IERROR)
 INTEGER WIN, ERRHANDLER, IERROR

{MPI::Errhandler MPI::Win::Get_errhandler() const(*binding deprecated, see Section ??*) }

Retrieves the error handler currently associated with a window.

1.3.3 Error Handlers for Files

MPI_FILE_CREATE_ERRHANDLER(function, errhandler)

IN	function	user defined error handling procedure (function)
OUT	errhandler	MPI error handler (handle)

int MPI_File_create_errhandler(MPI_File_errhandler_function *function,
 MPI_Errhandler *errhandler)

MPI_FILE_CREATE_ERRHANDLER(FUNCTION, ERRHANDLER, IERROR)
 EXTERNAL FUNCTION
 INTEGER ERRHANDLER, IERROR

{static MPI::Errhandler
 MPI::File::Create_errhandler(MPI::File::Errhandler_function*
 function)(*binding deprecated, see Section ??*) }

Creates an error handler that can be attached to a file object. The user routine should be, in C, a function of type MPI_File_errhandler_function, which is defined as

typedef void MPI_File_errhandler_function(MPI_File *, int *, ...);

The first argument is the file in use, the second is the error code to be returned.

In Fortran, the user routine should be of the form:

SUBROUTINE FILE_ERRHANDLER_FUNCTION(FILE, ERROR_CODE)
 INTEGER FILE, ERROR_CODE

In C++, the user routine should be of the form:

{typedef void MPI::File::Errhandler_function(MPI::File &, int *, ...);
 (*binding deprecated, see Section ??*)}

1 MPI_FILE_SET_ERRHANDLER(file, errhandler)

2 INOUT file file (handle)

3 IN errhandler new error handler for file (handle)

4 int MPI_File_set_errhandler(MPI_File file, MPI_Errhandler errhandler)

5 MPI_FILE_SET_ERRHANDLER(FILE, ERRHANDLER, IERROR)

6 INTEGER FILE, ERRHANDLER, IERROR

7 {void MPI::File::Set_errhandler(const MPI::Errhandler& errhandler) (*binding*
8 *deprecated, see Section ??*) }

9 Attaches a new error handler to a file. The error handler must be either a predefined
10 error handler, or an error handler created by a call to MPI_FILE_CREATE_ERRHANDLER.

11 MPI_FILE_GET_ERRHANDLER(file, errhandler)

12 IN file file (handle)

13 OUT errhandler error handler currently associated with file (handle)

14 int MPI_File_get_errhandler(MPI_File file, MPI_Errhandler *errhandler)

15 MPI_FILE_GET_ERRHANDLER(FILE, ERRHANDLER, IERROR)

16 INTEGER FILE, ERRHANDLER, IERROR

17 {MPI::Errhandler MPI::File::Get_errhandler() const (*binding deprecated, see*
18 *Section ??*) }

19 Retrieves the error handler currently associated with a file.

20 1.3.4 Freeing Errorhandlers and Retrieving Error Strings

21 MPI_ERRHANDLER_FREE(errhandler)

22 INOUT errhandler MPI error handler (handle)

23 int MPI_Errhandler_free(MPI_Errhandler *errhandler)

24 MPI_ERRHANDLER_FREE(ERRHANDLER, IERROR)

25 INTEGER ERRHANDLER, IERROR

26 {void MPI::Errhandler::Free() (*binding deprecated, see Section ??*) }

27 Marks the error handler associated with `errhandler` for deallocation and sets `errhandler`
28 to `MPI_ERRHANDLER_NULL`. The error handler will be deallocated after all the objects
29 associated with it (communicator, window, or file) have been deallocated.

MPI_ERROR_STRING(errorcode, string, resultlen)

IN	errorcode	Error code returned by an MPI routine
OUT	string	Text that corresponds to the errorcode
OUT	resultlen	Length (in printable characters) of the result returned in string

int MPI_Error_string(int errorcode, char *string, int *resultlen)

MPI_ERROR_STRING(ERRORCODE, STRING, RESULTLEN, IERROR)

INTEGER ERRORCODE, RESULTLEN, IERROR

CHARACTER*(*) STRING

{void MPI::Get_error_string(int errorcode, char* name,
int& resultlen) (*binding deprecated, see Section ??*) }

Returns the error string associated with an error code or class. The argument `string` must represent storage that is at least `MPI_MAX_ERROR_STRING` characters long.

The number of characters actually written is returned in the output argument, `resultlen`.

Rationale. The form of this function was chosen to make the Fortran and C bindings similar. A version that returns a pointer to a string has two difficulties. First, the return string must be statically allocated and different for each error message (allowing the pointers returned by successive calls to `MPI_ERROR_STRING` to point to the correct message). Second, in Fortran, a function declared as returning `CHARACTER*(*)` can not be referenced in, for example, a `PRINT` statement. (*End of rationale.*)

1.4 Error Codes and Classes

The error codes returned by MPI are left entirely to the implementation (with the exception of `MPI_SUCCESS`). This is done to allow an implementation to provide as much information as possible in the error code (for use with `MPI_ERROR_STRING`).

To make it possible for an application to interpret an error code, the routine `MPI_ERROR_CLASS` converts any error code into one of a small set of standard error codes, called *error classes*. Valid error classes are shown in Table 1.1 and Table 1.2.

The error classes are a subset of the error codes: an MPI function may return an error class number; and the function `MPI_ERROR_STRING` can be used to compute the error string associated with an error class. An MPI error class is a valid MPI error code. Specifically, the values defined for MPI error classes are valid MPI error codes.

The error codes satisfy,

$$0 = \text{MPI_SUCCESS} < \text{MPI_ERR_...} \leq \text{MPI_ERR_LASTCODE}.$$

Rationale. The difference between `MPI_ERR_UNKNOWN` and `MPI_ERR_OTHER` is that `MPI_ERROR_STRING` can return useful information about `MPI_ERR_OTHER`.

Note that `MPI_SUCCESS = 0` is necessary to be consistent with C practice; the separation of error classes and error codes allows us to define the error classes this way. Having a known `LASTCODE` is often a nice sanity check as well. (*End of rationale.*)

MPI_SUCCESS	No error
MPI_ERR_BUFFER	Invalid buffer pointer
MPI_ERR_COUNT	Invalid count argument
MPI_ERR_TYPE	Invalid datatype argument
MPI_ERR_TAG	Invalid tag argument
MPI_ERR_COMM	Invalid communicator
MPI_ERR_RANK	Invalid rank
MPI_ERR_REQUEST	Invalid request (handle)
MPI_ERR_ROOT	Invalid root
MPI_ERR_GROUP	Invalid group
MPI_ERR_OP	Invalid operation
MPI_ERR_TOPOLOGY	Invalid topology
MPI_ERR_DIMS	Invalid dimension argument
MPI_ERR_ARG	Invalid argument of some other kind
MPI_ERR_UNKNOWN	Unknown error
MPI_ERR_TRUNCATE	Message truncated on receive
MPI_ERR_OTHER	Known error not in this list
MPI_ERR_INTERN	Internal MPI (implementation) error
MPI_ERR_IN_STATUS	Error code is in status
MPI_ERR_PENDING	Pending request
MPI_ERR_KEYVAL	Invalid keyval has been passed
MPI_ERR_NO_MEM	MPI_ALLOC_MEM failed because memory is exhausted
MPI_ERR_BASE	Invalid base passed to MPI_FREE_MEM
MPI_ERR_INFO_KEY	Key longer than MPI_MAX_INFO_KEY
MPI_ERR_INFO_VALUE	Value longer than MPI_MAX_INFO_VAL
MPI_ERR_INFO_NOKEY	Invalid key passed to MPI_INFO_DELETE
MPI_ERR_SPAWN	Error in spawning processes
MPI_ERR_PORT	Invalid port name passed to MPI_COMM_CONNECT
MPI_ERR_SERVICE	Invalid service name passed to MPI_UNPUBLISH_NAME
MPI_ERR_NAME	Invalid service name passed to MPI_LOOKUP_NAME
MPI_ERR_WIN	Invalid win argument
MPI_ERR_SIZE	Invalid size argument
MPI_ERR_DISP	Invalid disp argument
MPI_ERR_INFO	Invalid info argument
MPI_ERR_LOCKTYPE	Invalid locktype argument
MPI_ERR_ASSERT	Invalid assert argument
MPI_ERR_RMA_CONFLICT	Conflicting accesses to window
MPI_ERR_RMA_SYNC	Wrong synchronization of RMA calls

Table 1.1: Error classes (Part 1)

MPI_ERR_FILE	Invalid file handle	1
MPI_ERR_NOT_SAME	Collective argument not identical on all processes, or collective routines called in a different order by different processes	2
		3
		4
MPI_ERR_AMODE	Error related to the <code>amode</code> passed to <code>MPI_FILE_OPEN</code>	5
		6
MPI_ERR_UNSUPPORTED_DATAREP	Unsupported <code>datarep</code> passed to <code>MPI_FILE_SET_VIEW</code>	7
		8
MPI_ERR_UNSUPPORTED_OPERATION	Unsupported operation, such as seeking on a file which supports sequential access only	9
		10
MPI_ERR_NO_SUCH_FILE	File does not exist	11
MPI_ERR_FILE_EXISTS	File exists	12
MPI_ERR_BAD_FILE	Invalid file name (e.g., path name too long)	13
MPI_ERR_ACCESS	Permission denied	14
MPI_ERR_NO_SPACE	Not enough space	15
MPI_ERR_QUOTA	Quota exceeded	16
MPI_ERR_READ_ONLY	Read-only file or file system	17
MPI_ERR_FILE_IN_USE	File operation could not be completed, as the file is currently open by some process	18
		19
MPI_ERR_DUP_DATAREP	Conversion functions could not be registered because a data representation identifier that was already defined was passed to <code>MPI_REGISTER_DATAREP</code>	20
		21
		22
MPI_ERR_CONVERSION	An error occurred in a user supplied data conversion function.	23
		24
		25
MPI_ERR_IO	Other I/O error	26
MPI_ERR_LASTCODE	Last error code	27

Table 1.2: Error classes (Part 2)

`MPI_ERROR_CLASS(errorcode, errorclass)`

IN	<code>errorcode</code>	Error code returned by an MPI routine
OUT	<code>errorclass</code>	Error class associated with <code>errorcode</code>

`int MPI_Error_class(int errorcode, int *errorclass)`

`MPI_ERROR_CLASS(ERRORCODE, ERRORCLASS, IERROR)`
`INTEGER ERRORCODE, ERRORCLASS, IERROR`

`{int MPI::Get_error_class(int errorcode) (binding deprecated, see Section ??) }`

The function `MPI_ERROR_CLASS` maps each standard error code (error class) onto itself.

1.5 Error Classes, Error Codes, and Error Handlers

Users may want to write a layered library on top of an existing MPI implementation, and this library may have its own set of error codes and classes. An example of such a library is an I/O library based on MPI, see Chapter ?? on page ?. For this purpose, functions are needed to:

1. add a new error class to the ones an MPI implementation already knows.
2. associate error codes with this error class, so that `MPI_ERROR_CLASS` works.
3. associate strings with these error codes, so that `MPI_ERROR_STRING` works.
4. invoke the error handler associated with a communicator, window, or object.

Several functions are provided to do this. They are all local. No functions are provided to free error classes or codes: it is not expected that an application will generate them in significant numbers.

`MPI_ADD_ERROR_CLASS(errorclass)`

OUT errorclass value for the new error class (integer)

`int MPI_Add_error_class(int *errorclass)`

`MPI_ADD_ERROR_CLASS(ERRORCLASS, IERROR)`

INTEGER ERRORCLASS, IERROR

`{int MPI::Add_error_class() (binding deprecated, see Section ??) }`

Creates a new error class and returns the value for it.

Rationale. To avoid conflicts with existing error codes and classes, the value is set by the implementation and not by the user. (*End of rationale.*)

Advice to implementors. A high-quality implementation will return the value for a new `errorclass` in the same deterministic way on all processes. (*End of advice to implementors.*)

Advice to users. Since a call to `MPI_ADD_ERROR_CLASS` is local, the same `errorclass` may not be returned on all processes that make this call. Thus, it is not safe to assume that registering a new error on a set of processes at the same time will yield the same `errorclass` on all of the processes. However, if an implementation returns the new `errorclass` in a deterministic way, and they are always generated in the same order on the same set of processes (for example, all processes), then the value will be the same. However, even if a deterministic algorithm is used, the value can vary across processes. This can happen, for example, if different but overlapping groups of processes make a series of calls. As a result of these issues, getting the “same” error on multiple processes may not cause the same value of error code to be generated. (*End of advice to users.*)

The value of `MPI_ERR_LASTCODE` is a constant value and is not affected by new user-defined error codes and classes. Instead, a predefined attribute key `MPI_LASTUSED` is associated with `MPI_COMM_WORLD`. The attribute value corresponding to this key is the current maximum error class including the user-defined ones. This is a local value and may be different on different processes. The value returned by this key is always greater than or equal to `MPI_ERR_LASTCODE`.

Advice to users. The value returned by the key `MPI_LASTUSED` will not change unless the user calls a function to explicitly add an error class/code. In a multi-threaded environment, the user must take extra care in assuming this value has not changed. Note that error codes and error classes are not necessarily dense. A user may not assume that each error class below `MPI_LASTUSED` is valid. (*End of advice to users.*)

`MPI_ADD_ERROR_CODE(errorclass, errorcode)`

IN	errorclass	error class (integer)
OUT	errorcode	new error code to associated with errorclass (integer)

`int MPI_Add_error_code(int errorclass, int *errorcode)`

`MPI_ADD_ERROR_CODE(ERRORCLASS, ERRORCODE, IERROR)`
`INTEGER ERRORCLASS, ERRORCODE, IERROR`

`{int MPI::Add_error_code(int errorclass) (binding deprecated, see Section ??) }`

Creates new error code associated with `errorclass` and returns its value in `errorcode`.

Rationale. To avoid conflicts with existing error codes and classes, the value of the new error code is set by the implementation and not by the user. (*End of rationale.*)

Advice to implementors. A high-quality implementation will return the value for a new `errorcode` in the same deterministic way on all processes. (*End of advice to implementors.*)

`MPI_ADD_ERROR_STRING(errorcode, string)`

IN	errorcode	error code or class (integer)
IN	string	text corresponding to errorcode (string)

`int MPI_Add_error_string(int errorcode, char *string)`

`MPI_ADD_ERROR_STRING(ERRORCODE, STRING, IERROR)`
`INTEGER ERRORCODE, IERROR`
`CHARACTER*(*) STRING`

`{void MPI::Add_error_string(int errorcode, const char* string) (binding deprecated, see Section ??) }`

Associates an error string with an error code or class. The string must be no more than `MPI_MAX_ERROR_STRING` characters long. The length of the string is as defined in the calling language. The length of the string does not include the null terminator in C or C++. Trailing blanks will be stripped in Fortran. Calling `MPI_ADD_ERROR_STRING` for an errorcode that already has a string will replace the old string with the new string. It is erroneous to call `MPI_ADD_ERROR_STRING` for an error code or class with a value \leq `MPI_ERR_LASTCODE`.

If `MPI_ERROR_STRING` is called when no string has been set, it will return a empty string (all spaces in Fortran, "" in C and C++).

Section 1.3 on page 6 describes the methods for creating and associating error handlers with communicators, files, and windows.

`MPI_COMM_CALL_ERRHANDLER (comm, errorcode)`

IN comm communicator with error handler (handle)

IN errorcode error code (integer)

int MPI_Comm_call_errhandler(MPI_Comm comm, int errorcode)

MPI_COMM_CALL_ERRHANDLER(COMM, ERRORCODE, IERROR)

INTEGER COMM, ERRORCODE, IERROR

{void MPI::Comm::Call_errhandler(int errorcode) const(*binding deprecated, see Section ??*) }

This function invokes the error handler assigned to the communicator with the error code supplied. This function returns `MPI_SUCCESS` in C and C++ and the same value in `IERROR` if the error handler was successfully called (assuming the process is not aborted and the error handler returns).

Advice to users. Users should note that the default error handler is `MPI_ERRORS_ARE_FATAL`. Thus, calling `MPI_COMM_CALL_ERRHANDLER` will abort the `comm` processes if the default error handler has not been changed for this communicator or on the parent before the communicator was created. (*End of advice to users.*)

`MPI_WIN_CALL_ERRHANDLER (win, errorcode)`

IN win window with error handler (handle)

IN errorcode error code (integer)

int MPI_Win_call_errhandler(MPI_Win win, int errorcode)

MPI_WIN_CALL_ERRHANDLER(WIN, ERRORCODE, IERROR)

INTEGER WIN, ERRORCODE, IERROR

{void MPI::Win::Call_errhandler(int errorcode) const(*binding deprecated, see Section ??*) }

This function invokes the error handler assigned to the window with the error code supplied. This function returns `MPI_SUCCESS` in C and C++ and the same value in `IERROR` if the error handler was successfully called (assuming the process is not aborted and the error handler returns).

Advice to users. As with communicators, the default error handler for windows is `MPI_ERRORS_ARE_FATAL`. (*End of advice to users.*)

`MPI_FILE_CALL_ERRHANDLER` (fh, errorcode)

IN	fh	file with error handler (handle)
IN	errorcode	error code (integer)

`int MPI_File_call_errhandler(MPI_File fh, int errorcode)`

`MPI_FILE_CALL_ERRHANDLER(FH, ERRORCODE, IERROR)`

`INTEGER FH, ERRORCODE, IERROR`

`{void MPI::File::Call_errhandler(int errorcode) const` (*binding deprecated, see Section ??*) `}`

This function invokes the error handler assigned to the file with the error code supplied. This function returns `MPI_SUCCESS` in C and C++ and the same value in `IERROR` if the error handler was successfully called (assuming the process is not aborted and the error handler returns).

Advice to users. Unlike errors on communicators and windows, the default behavior for files is to have `MPI_ERRORS_RETURN`. (*End of advice to users.*)

Advice to users. Users are warned that handlers should not be called recursively with `MPI_COMM_CALL_ERRHANDLER`, `MPI_FILE_CALL_ERRHANDLER`, or `MPI_WIN_CALL_ERRHANDLER`. Doing this can create a situation where an infinite recursion is created. This can occur if `MPI_COMM_CALL_ERRHANDLER`, `MPI_FILE_CALL_ERRHANDLER`, or `MPI_WIN_CALL_ERRHANDLER` is called inside an error handler.

Error codes and classes are associated with a process. As a result, they may be used in any error handler. Error handlers should be prepared to deal with any error code they are given. Furthermore, it is good practice to only call an error handler with the appropriate error codes. For example, file errors would normally be sent to the file error handler. (*End of advice to users.*)

1.6 Timers and Synchronization

MPI defines a timer. A timer is specified even though it is not “message-passing,” because timing parallel programs is important in “performance debugging” and because existing timers (both in POSIX 1003.1-1988 and 1003.4D 14.1 and in Fortran 90) are either inconvenient or do not provide adequate access to high-resolution timers. See also Section ?? on page ??.

1 MPI_WTIME()

2
3 double MPI_Wtime(void)

4 DOUBLE PRECISION MPI_WTIME()

5
6 {double MPI::Wtime() (*binding deprecated, see Section ??*) }

7
8 MPI_WTIME returns a floating-point number of seconds, representing elapsed wall-clock time since some time in the past.

9
10 The “time in the past” is guaranteed not to change during the life of the process. The user is responsible for converting large numbers of seconds to other units if they are preferred.

11
12 This function is portable (it returns seconds, not “ticks”), it allows high-resolution, and carries no unnecessary baggage. One would use it like this:

13
14
15 {
16 double starttime, endtime;
17 starttime = MPI_Wtime();
18 stuff to be timed ...
19 endtime = MPI_Wtime();
20 printf("That took %f seconds\n",endtime-starttime);
21 }
22

23
24 The times returned are local to the node that called them. There is no requirement that different nodes return “the same time.” (But see also the discussion of MPI_WTIME_IS_GLOBAL).

25
26
27 MPI_WTICK()

28
29 double MPI_Wtick(void)

30 DOUBLE PRECISION MPI_WTICK()

31
32 {double MPI::Wtick() (*binding deprecated, see Section ??*) }

33
34 MPI_WTICK returns the resolution of MPI_WTIME in seconds. That is, it returns, as a double precision value, the number of seconds between successive clock ticks. For example, if the clock is implemented by the hardware as a counter that is incremented every millisecond, the value returned by MPI_WTICK should be 10^{-3} .

35 36 37 38 39 1.7 Startup

40
41 One goal of MPI is to achieve *source code portability*. By this we mean that a program written using MPI and complying with the relevant language standards is portable as written, and must not require any source code changes when moved from one system to another. This explicitly does *not* say anything about how an MPI program is started or launched from the command line, nor what the user must do to set up the environment in which an MPI program will run. However, an implementation may require some setup to be performed

before other MPI routines may be called. To provide for this, MPI includes an initialization routine `MPI_INIT`.

`MPI_INIT()`

```
int MPI_Init(int *argc, char ***argv)
```

```
MPI_INIT(IERROR)
```

```
    INTEGER IERROR
```

```
{void MPI::Init(int& argc, char**& argv) (binding deprecated, see Section ??) }
```

```
{void MPI::Init() (binding deprecated, see Section ??) }
```

All MPI programs must contain exactly one call to an MPI initialization routine: `MPI_INIT` or `MPI_INIT_THREAD`. Subsequent calls to any initialization routines are erroneous. The only MPI functions that may be invoked before the MPI initialization routines are called are `MPI_GET_VERSION`, `MPI_INITIALIZED`, [\[and\]](#) `MPI_FINALIZED`[\[\]](#), and any function with the prefix `MPI_T_*` (within the constraints for functions with this prefix listed in [Section 2.3.4](#)). The version for ISO C accepts the `argc` and `argv` that are provided by the arguments to `main` or `NULL`:

```
int main(int argc, char **argv)
```

```
{
```

```
    MPI_Init(&argc, &argv);
```

```
    /* parse arguments */
```

```
    /* main program    */
```

```
    MPI_Finalize();    /* see below */
```

```
}
```

The Fortran version takes only `IERROR`.

Conforming implementations of MPI are required to allow applications to pass `NULL` for both the `argc` and `argv` arguments of `main` in C and C++. In C++, there is an alternative binding for `MPI::Init` that does not have these arguments at all.

Rationale. In some applications, libraries may be making the call to `MPI_Init`, and may not have access to `argc` and `argv` from `main`. It is anticipated that applications requiring special information about the environment or information supplied by `mpiexec` can get that information from environment variables. (*End of rationale.*)

`MPI_FINALIZE()`

```
int MPI_Finalize(void)
```

```
MPI_FINALIZE(IERROR)
```

```
    INTEGER IERROR
```

```
1 {void MPI::Finalize() (binding deprecated, see Section ??) }
```

2
3 This routine cleans up all MPI state. Each process must call MPI_FINALIZE before
4 it exits. Unless there has been a call to MPI_ABORT, each process must ensure that all
5 pending nonblocking communications are (locally) complete before calling MPI_FINALIZE.
6 Further, at the instant at which the last process calls MPI_FINALIZE, all pending sends
7 must be matched by a receive, and all pending receives must be matched by a send.

8 For example, the following program is correct:

9	Process 0	Process 1
10	-----	-----
11	MPI_Init();	MPI_Init();
12	MPI_Send(dest=1);	MPI_Recv(src=0);
13	MPI_Finalize();	MPI_Finalize();

14
15 Without the matching receive, the program is erroneous:

16	Process 0	Process 1
17	-----	-----
18	MPI_Init();	MPI_Init();
19	MPI_Send (dest=1);	
20	MPI_Finalize();	MPI_Finalize();

21
22 A successful return from a blocking communication operation or from MPI_WAIT or
23 MPI_TEST tells the user that the buffer can be reused and means that the communication
24 is completed by the user, but does not guarantee that the local process has no more work
25 to do. A successful return from MPI_REQUEST_FREE with a request handle generated by
26 an MPI_ISEND nullifies the handle but provides no assurance of operation completion. The
27 MPI_ISEND is complete only when it is known by some means that a matching receive has
28 completed. MPI_FINALIZE guarantees that all local actions required by communications
29 the user has completed will, in fact, occur before it returns.

30 MPI_FINALIZE guarantees nothing about pending communications that have not been
31 completed (completion is assured only by MPI_WAIT, MPI_TEST, or MPI_REQUEST_FREE
32 combined with some other verification of completion).
33

34 **Example 1.3** This program is correct:

35	rank 0	rank 1
36	=====	=====
37
38	MPI_Isend();	MPI_Recv();
39	MPI_Request_free();	MPI_Barrier();
40	MPI_Barrier();	MPI_Finalize();
41	MPI_Finalize();	exit();
42	exit();	

43
44
45 **Example 1.4** This program is erroneous and its behavior is undefined:
46
47
48

```

rank 0                                rank 1
=====
...
MPI_Isend();                          MPI_Recv();
MPI_Request_free();                  MPI_Finalize();
MPI_Finalize();                      exit();
exit();

```

If no `MPI_BUFFER_DETACH` occurs between an `MPI_BSEND` (or other buffered send) and `MPI_FINALIZE`, the `MPI_FINALIZE` implicitly supplies the `MPI_BUFFER_DETACH`.

Example 1.5 This program is correct, and after the `MPI_Finalize`, it is as if the buffer had been detached.

```

rank 0                                rank 1
=====
...
buffer = malloc(1000000);             MPI_Recv();
MPI_Buffer_attach();                  MPI_Finalize();
MPI_Bsend();                          exit();
MPI_Finalize();
free(buffer);
exit();

```

Example 1.6 In this example, `MPI_Iprobe()` must return a `FALSE` flag. `MPI_Test_cancelled()` must return a `TRUE` flag, independent of the relative order of execution of `MPI_Cancel()` in process 0 and `MPI_Finalize()` in process 1.

The `MPI_Iprobe()` call is there to make sure the implementation knows that the “tag1” message exists at the destination, without being able to claim that the user knows about it.

```

rank 0                                rank 1
=====
MPI_Init();                          MPI_Init();
MPI_Isend(tag1);                      MPI_Barrier();
MPI_Barrier();                       MPI_Iprobe(tag2);
                                     MPI_Barrier();
MPI_Barrier();                       MPI_Finalize();
                                     exit();

MPI_Cancel();
MPI_Wait();
MPI_Test_cancelled();
MPI_Finalize();
exit();

```

Advice to implementors. An implementation may need to delay the return from `MPI_FINALIZE` until all potential future message cancellations have been processed.

One possible solution is to place a barrier inside MPI_FINALIZE (*End of advice to implementors.*)

Once MPI_FINALIZE returns, no MPI routine (not even MPI_INIT) may be called, except for MPI_GET_VERSION, MPI_INITIALIZED, [and] MPI_FINALIZED[], and any function with the prefix MPI_T_* (within the constraints for functions with this prefix listed in Section 2.3.4). Each process must complete any pending communication it initiated before it calls MPI_FINALIZE. If the call returns, each process may continue local computations, or exit, without participating in further MPI communication with other processes. MPI_FINALIZE is collective over all connected processes. If no processes were spawned, accepted or connected then this means over MPI_COMM_WORLD; otherwise it is collective over the union of all processes that have been and continue to be connected, as explained in Section ?? on page ??.

Advice to implementors. Even though a process has completed all the communication it initiated, such communication may not yet be completed from the viewpoint of the underlying MPI system. E.g., a blocking send may have completed, even though the data is still buffered at the sender. The MPI implementation must ensure that a process has completed any involvement in MPI communication before MPI_FINALIZE returns. Thus, if a process exits after the call to MPI_FINALIZE, this will not cause an ongoing communication to fail. (*End of advice to implementors.*)

Although it is not required that all processes return from MPI_FINALIZE, it is required that at least process 0 in MPI_COMM_WORLD return, so that users can know that the MPI portion of the computation is over. In addition, in a POSIX environment, they may desire to supply an exit code for each process that returns from MPI_FINALIZE.

Example 1.7 The following illustrates the use of requiring that at least one process return and that it be known that process 0 is one of the processes that return. One wants code like the following to work no matter how many processes return.

```
...
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
...
MPI_Finalize();
if (myrank == 0) {
    resultfile = fopen("outfile","w");
    dump_results(resultfile);
    fclose(resultfile);
}
exit(0);
```

MPI_INITIALIZED(flag)

OUT flag

Flag is true if MPI_INIT has been called and false otherwise.

int MPI_Initialized(int *flag)

```
MPI_INITIALIZED(FLAG, IERROR)
```

```
    LOGICAL FLAG
```

```
    INTEGER IERROR
```

```
{bool MPI::Is_initialized() (binding deprecated, see Section ??) }
```

This routine may be used to determine whether MPI_INIT has been called. MPI_INITIALIZED returns true if the calling process has called MPI_INIT. Whether MPI_FINALIZE has been called does not affect the behavior of MPI_INITIALIZED. It is one of the few routines that may be called before MPI_INIT is called.

```
MPI_ABORT( comm, errorcode )
```

```
    IN          comm          communicator of tasks to abort
```

```
    IN          errorcode     error code to return to invoking environment
```

```
int MPI_Abort(MPI_Comm comm, int errorcode)
```

```
MPI_ABORT(COMM, ERRORCODE, IERROR)
```

```
    INTEGER COMM, ERRORCODE, IERROR
```

```
{void MPI::Comm::Abort(int errorcode) (binding deprecated, see Section ??) }
```

This routine makes a “best attempt” to abort all tasks in the group of `comm`. This function does not require that the invoking environment take any action with the error code. However, a Unix or POSIX environment should handle this as a **return errorcode** from the main program.

It may not be possible for an MPI implementation to abort only the processes represented by `comm` if this is a subset of the processes. In this case, the MPI implementation should attempt to abort all the connected processes but should not abort any unconnected processes. If no processes were spawned, accepted or connected then this has the effect of aborting all the processes associated with MPI_COMM_WORLD.

Rationale. The communicator argument is provided to allow for future extensions of MPI to environments with, for example, dynamic process management. In particular, it allows but does not require an MPI implementation to abort a subset of MPI_COMM_WORLD. (*End of rationale.*)

Advice to users. Whether the errorcode is returned from the executable or from the MPI process startup mechanism (e.g., `mpiexec`), is an aspect of quality of the MPI library but not mandatory. (*End of advice to users.*)

Advice to implementors. Where possible, a high-quality implementation will try to return the errorcode from the MPI process startup mechanism (e.g. `mpiexec` or singleton init). (*End of advice to implementors.*)

1.7.1 Allowing User Functions at Process Termination

There are times in which it would be convenient to have actions happen when an MPI process finishes. For example, a routine may do initializations that are useful until the MPI job (or

Advice to implementors. Since attributes can be added from any supported language, the MPI implementation needs to remember the creating language so the correct callback is made. Implementations that use the attribute delete callback on MPI_COMM_SELF internally should register their internal callbacks before returning from MPI_INIT / MPI_INIT_THREAD, so that libraries or applications will not have portions of the MPI implementation shut down before the application-level callbacks are made. (*End of advice to implementors.*)

One of the goals of MPI was to allow for layered libraries. In order for a library to do this cleanly, it needs to know if MPI is active. In MPI the function `MPI_INITIALIZED` was provided to tell if MPI had been initialized. The problem arises in knowing if MPI has been finalized. Once MPI has been finalized it is no longer active and cannot be restarted. A library needs to be able to determine this to act accordingly. To achieve this the following function is needed:

Advice to users. MPI is “active” and it is thus safe to call MPI functions if MPI_INIT has completed and MPI_FINALIZE has not completed. If a library has no other way of knowing whether MPI is active or not, then it can use MPI_INITIALIZED and MPI_FINALIZED to determine this. For example, MPI is “active” in callback functions that are invoked during MPI_FINALIZE. (*End of advice to users.*)

1.8 Portable MPI Process Startup

A number of implementations of MPI provide a startup command for MPI programs that is of the form

```
mpirun <mpirun arguments> <program> <program arguments>
```

Separating the command to start the program from the program itself provides flexibility, particularly for network and heterogeneous implementations. For example, the startup script need not run on one of the machines that will be executing the MPI program itself.

Having a standard startup mechanism also extends the portability of MPI programs one step further, to the command lines and scripts that manage them. For example, a validation suite script that runs hundreds of programs can be a portable script if it is written using such a standard startup mechanism. In order that the “standard” command not be confused with existing practice, which is not standard and not portable among implementations, instead of `mpirun` MPI specifies `mpiexec`.

While a standardized startup mechanism improves the usability of MPI, the range of environments is so diverse (e.g., there may not even be a command line interface) that MPI cannot mandate such a mechanism. Instead, MPI specifies an `mpiexec` startup command and recommends but does not require it, as advice to implementors. However, if an implementation does provide a command called `mpiexec`, it must be of the form described below.

It is suggested that

```
mpiexec -n <numprocs> <program>
```

be at least one way to start `<program>` with an initial `MPI_COMM_WORLD` whose group contains `<numprocs>` processes. Other arguments to `mpiexec` may be implementation-dependent.

Advice to implementors. Implementors, if they do provide a special startup command for MPI programs, are advised to give it the following form. The syntax is chosen in order that `mpiexec` be able to be viewed as a command-line version of `MPI_COMM_SPAWN` (See Section ??).

Analogous to `MPI_COMM_SPAWN`, we have

```
mpiexec -n    <maxprocs>
          -soft <      >
          -host <      >
          -arch <      >
          -wdir <      >
          -path <      >
          -file <      >
          ...
          <command line>
```

for the case where a single command line for the application program and its arguments will suffice. See Section ?? for the meanings of these arguments. For the case corresponding to `MPI_COMM_SPAWN_MULTIPLE` there are two possible formats:

Form A:

```
1      mpiexec { <above arguments> } : { ... } : { ... } : ... : { ... }
```

2
3 As with MPI_COMM_SPAWN, all the arguments are optional. (Even the `-n x` argu-
4 ment is optional; the default is implementation dependent. It might be 1, it might be
5 taken from an environment variable, or it might be specified at compile time.) The
6 names and meanings of the arguments are taken from the keys in the `info` argument
7 to MPI_COMM_SPAWN. There may be other, implementation-dependent arguments
8 as well.

9 Note that Form A, though convenient to type, prevents colons from being program
10 arguments. Therefore an alternate, file-based form is allowed:

11 Form B:

```
12  
13      mpiexec -configfile <filename>
```

14
15 where the lines of `<filename>` are of the form separated by the colons in Form A.
16 Lines beginning with `#` are comments, and lines may be continued by terminating
17 the partial line with `\`.

18
19 **Example 1.8** Start 16 instances of `myprog` on the current or default machine:

```
20      mpiexec -n 16 myprog
```

21
22 **Example 1.9** Start 10 processes on the machine called `ferrari`:

```
23      mpiexec -n 10 -host ferrari myprog
```

24
25 **Example 1.10** Start three copies of the same program with different command-line
26 arguments:

```
27      mpiexec myprog infile1 : myprog infile2 : myprog infile3
```

28
29 **Example 1.11** Start the `ocean` program on five Suns and the `atmos` program on 10
30 RS/6000's:

```
31  
32      mpiexec -n 5 -arch sun ocean : -n 10 -arch rs6000 atmos
```

33
34 It is assumed that the implementation in this case has a method for choosing hosts of
35 the appropriate type. Their ranks are in the order specified.

36
37 **Example 1.12** Start the `ocean` program on five Suns and the `atmos` program on 10
38 RS/6000's (Form B):

```
39      mpiexec -configfile myfile
```

40
41 where `myfile` contains

```
42      -n 5 -arch sun    ocean  
43      -n 10 -arch rs6000 atmos
```

44
45 (*End of advice to implementors.*)

[]

1
2 ticket266.
3 ticket266.
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48

Chapter 2

Tool Support

2.1 Introduction

This chapter discusses a set of interfaces that allows debuggers, performance analyzers, and other tools to extract information about the operation of MPI processes. Specifically, this chapter defines both the MPI profiling interface (Section 2.2), which supports the transparent interception and inspection of MPI calls, and the MPI tool information interface (Section 2.3), which supports the inspection and manipulation of MPI control and performance variables. The interfaces described in this chapter are all defined in the context of an MPI process, i.e., are callable from the same code that invokes other MPI functions.

2.2 Profiling Interface

[WAS: Chapter]

2.2.1 Requirements

[WAS: Section]

To meet [the]the requirements for the MPI profiling interface, an implementation of the MPI functions *must*

1. provide a mechanism through which all of the MPI defined [functions]functions, except those allowed as macros (See Section ??[?]), may be accessed with a name shift. This requires, in C and Fortran, an alternate entry point name, with the prefix `PMPI_` for each MPI function. The profiling interface in C++ is described in Section ??. For routines implemented as macros, it is still required that the `PMPI_` version be supplied and work as expected, but it is not possible to replace at link time the `MPI_` version with a user-defined version.
2. ensure that those MPI functions that are not replaced may still be linked into an executable image without causing name clashes.
3. document the implementation of different language bindings of the MPI interface if they are layered on top of each other, so that the profiler developer knows whether she must implement the profile interface for each binding, or can [economise]economize by implementing it only for the lowest level routines.

4. where the implementation of different language bindings is done through a layered approach ([e.g.]e.g., the Fortran binding is a set of “wrapper” functions that call the C implementation), ensure that these wrapper functions are separable from the rest of the library.

This separability is necessary to allow a separate profiling library to be correctly implemented, since (at least with Unix linker semantics) the profiling library must contain these wrapper functions if it is to perform as expected. This requirement allows the person who builds the profiling library to extract these functions from the original MPI library and add them into the profiling library without bringing along any other unnecessary code.

5. provide a no-op routine `MPI_PCONTROL` in the MPI library.

2.2.2 Discussion

[WAS: Section]

The objective of the MPI profiling interface is to ensure that it is relatively easy for authors of profiling (and other similar) tools to interface their codes to MPI implementations on different machines.

Since MPI is a machine independent standard with many different implementations, it is unreasonable to expect that the authors of profiling tools for MPI will have access to the source code that implements MPI on any particular machine. It is therefore necessary to provide a mechanism by which the implementors of such tools can collect whatever performance information they wish *without* access to the underlying implementation.

We believe that having such an interface is important if MPI is to be attractive to end users, since the availability of many different tools will be a significant factor in attracting users to the MPI standard.

The profiling interface is just that, an interface. It says *nothing* about the way in which it is used. There is therefore no attempt to lay down what information is collected through the interface, or how the collected information is saved, filtered, or displayed.

While the initial impetus for the development of this interface arose from the desire to permit the implementation of profiling tools, it is clear that an interface like that specified may also prove useful for other purposes, such as “internetworking” multiple MPI implementations. Since all that is defined is an interface, there is no objection to its being used wherever it is useful.

As the issues being addressed here are intimately tied up with the way in which executable images are built, which may differ greatly on different machines, the examples given below should be treated solely as one way of implementing the objective of the MPI profiling interface. The actual requirements made of an implementation are those detailed in the Requirements section above, the whole of the rest of this chapter is only present as justification and discussion of the logic for those requirements.

The examples below show one way in which an implementation could be constructed to meet the requirements on a Unix system (there are doubtless others that would be equally valid).

2.2.3 Logic of the Design

[WAS: Section]

Provided that an MPI implementation meets the requirements above, it is possible for the implementor of the profiling system to intercept all of the MPI calls that are made by the user program. She can then collect whatever information she requires before calling the underlying MPI implementation (through its name shifted entry points) to achieve the desired effects.

2.2.4 Miscellaneous Control of Profiling

[WAS: Subsection, Now still a subsection to remove single subsubsection]

There is a clear requirement for the user code to be able to control the profiler dynamically at run time. This is normally used for (at least) the purposes of

- Enabling and disabling profiling depending on the state of the calculation.
- Flushing trace buffers at non-critical points in the [calculation]calculation.
- Adding user events to a trace file.

These requirements are met by use of the MPI_PCONTROL.

MPI_PCONTROL(level, ...)

IN	level	Profiling level
----	-------	-----------------

int MPI_Pcontrol(const int level, ...)

MPI_PCONTROL(LEVEL)

INTEGER LEVEL

{void MPI::Pcontrol(const int level, ...) (binding deprecated, see Section ??) }

MPI libraries themselves make no use of this routine, and simply return immediately to the user code. However the presence of calls to this routine allows a profiling package to be explicitly called by the user.

Since MPI has no control of the implementation of the profiling code, we are unable to specify precisely the semantics that will be provided by calls to MPI_PCONTROL. This vagueness extends to the number of arguments to the function, and their datatypes.

However to provide some level of portability of user codes to different profiling libraries, we request the following meanings for certain values of level.

- level==0 Profiling is disabled.
- level==1 Profiling is enabled at a normal default level of detail.
- level==2 Profile buffers are [flushed. (This may be a no-op in some profilers).]flushed, which may be a no-op in some profilers.
- All other values of level have profile library defined effects and additional arguments.

We also request that the default state after MPI_INIT has been called is for profiling to be enabled at the normal default level. (i.e. as if MPI_PCONTROL had just been called with the argument 1). This allows users to link with a profiling library and obtain profile output without having to modify their source code at all.

ticket0.

The provision of MPI_PCONTROL as a no-op in the standard MPI library [allows them to modify their source code to obtain] supports the collection of more detailed profiling information[, but still be able to link exactly the] with source [same code] code that can still link against the standard MPI library.

[WAS: Subsection Examples]

2.2.5 Profiler Implementation [Example]

[Suppose that the profiler wishes to] A profiler can accumulate the total amount of data sent by the [MPI_SEND] MPI_SEND function, along with the total elapsed time spent in the [function. This could trivially be achieved thus] function, as follows:

```
static int totalBytes = 0;
static double totalTime = 0.0;

int MPI_Send(void* buffer, int count, MPI_Datatype datatype,
             int dest, int tag, MPI_Comm comm)
{
    double tstart = MPI_Wtime();          /* Pass on all the arguments */
    int extent;
    int result = PMPI_Send(buffer, count, datatype, dest, tag, comm);

    MPI_Type_size(datatype, &extent); /* Compute size */
    totalBytes += count*extent;

    totalTime += MPI_Wtime() - tstart;    /* and time */

    return result;
}
```

2.2.6 MPI Library Implementation [Example]

[On a Unix system, in which the MPI library is implemented in C, then] If the MPI library is implemented in C on a Unix system, then there [there are various possible options, of which two of the most obvious] are various options, including the two presented here, for supporting [are presented here. Which is better depends on whether the linker and] the name-shift requirement. The choice between these two options [compiler support weak symbols.] depends partly on whether the linker and compiler support weak symbols.

Systems with Weak Symbols

If the compiler and linker support weak external symbols ([e.g.] e.g., Solaris 2.x, other system V.4 machines), then only a single library is required through the use of `#pragma weak` thus

```
#pragma weak MPI_Example = PMPI_Example

int PMPI_Example(/* appropriate args */)
{
    /* Useful content */
}
```

The effect of this `#pragma` is to define the external symbol `MPI_Example` as a weak definition. This means that the linker will not complain if there is another definition of the symbol (for instance in the profiling library), however if no other definition exists, then the linker will use the weak definition.

Systems Without Weak Symbols

In the absence of weak symbols then one possible solution would be to use the C macro pre-processor thus

```
#ifndef PROFILELIB
#   ifdef __STDC__
#       define FUNCTION(name) P##name
#   else
#       define FUNCTION(name) P/**/name
#   endif
#else
#   define FUNCTION(name) name
#endif
```

Each of the user visible functions in the library would then be declared thus

```
int FUNCTION(MPI_Example)(/* appropriate args */)
{
    /* Useful content */
}
```

The same source file can then be compiled to produce both versions of the library, depending on the state of the `PROFILELIB` macro symbol.

It is required that the standard MPI library be built in such a way that the inclusion of MPI functions can be achieved one at a time. This is a somewhat unpleasant requirement, since it may mean that each external function has to be compiled from a separate file. However this is necessary so that the author of the profiling library need only define those MPI functions that she wishes to intercept, references to any others being fulfilled by the normal MPI library. Therefore the link step can look something like this

```
% cc ... -lmyprof -lpmpi -lmpi
```

Here `libmyprof.a` contains the profiler functions that intercept some of the MPI functions[, `libpmpi.a` contains the “name shifted” MPI functions, and `libmpi.a` contains the normal definitions of the MPI functions.

2.2.7 Complications

Multiple Counting

Since parts of the MPI library may themselves be implemented using more basic MPI functions ([e.g.]e.g., a portable implementation of the collective operations implemented using point to point communications), there is potential for profiling functions to be called from within an MPI function that was called from a profiling function. This could lead to “double

ticket0.

counting” of the time spent in the inner routine. Since this effect could actually be useful under some circumstances ([e.g.]e.g., it might allow one to answer the question “How much time is spent in the point to point routines when they’re called from collective functions?”), we have decided not to enforce any restrictions on the author of the MPI library that would overcome this. Therefore the author of the profiling library should be aware of this problem, and guard against it herself. In a single threaded world this is easily achieved through use of a static variable in the profiling code that remembers if you are already inside a profiling routine. It becomes more complex in a multi-threaded environment (as does the meaning of the times recorded[!]).

Linker Oddities

The Unix linker traditionally operates in one [pass:]pass: the effect of this is that functions from libraries are only included in the image if they are needed at the time the library is scanned. When combined with weak symbols, or multiple definitions of the same function, this can cause odd (and unexpected) effects.

Consider, for instance, an implementation of MPI in which the Fortran binding is achieved by using wrapper functions on top of the C implementation. The author of the profile library then assumes that it is reasonable only to provide profile functions for the C binding, since Fortran will eventually call these, and the cost of the wrappers is assumed to be small. However, if the wrapper functions are not in the profiling library, then none of the profiled entry points will be undefined when the profiling library is called. Therefore none of the profiling code will be included in the image. When the standard MPI library is scanned, the Fortran wrappers will be resolved, and will also pull in the base versions of the MPI functions. The overall effect is that the code will link successfully, but will not be profiled.

To overcome this we must ensure that the Fortran wrapper functions are included in the profiling version of the library. We ensure that this is possible by requiring that these be separable from the rest of the base MPI library. This allows them to be aared out of the base library and into the profiling one.

2.2.8 Multiple Levels of Interception

[WAS: Section] The scheme given here does not directly support the nesting of profiling functions, since it provides only a single alternative name for each MPI function. Consideration was given to an implementation that would allow multiple levels of call interception, however we were unable to construct an implementation of this that did not have the following disadvantages

- assuming a particular implementation language[.],
- imposing a run time cost even when no profiling was taking place.

Since one of the objectives of MPI is to permit efficient, low latency implementations, and it is not the business of a standard to require a particular implementation language, we decided to accept the scheme outlined above.

[Note, however, that it is possible to use the scheme above to implement a multi-level system, since the function called by the user may call many different profiling functions before calling the underlying MPI function.]

[Unfortunately such an implementation may require more cooperation between the different profiling libraries than is required for the single level implementation detailed above.]Note, however, that it is possible to use the scheme above to implement a multi-level system, since the function called by the user may call many different profiling functions before calling the underlying MPI function. This capability has been demonstrated in the P^NMPI tool infrastructure [1].

[]

2.3 The MPI Tool Information Interface

MPI implementations often use internal variables to control their operation and performance. Understanding and manipulating these variables can provide a more efficient execution environment or improve performance for many applications. This section describes the MPI tool information interface, which provides a mechanism for MPI implementors to expose a set of variables, each of which represents a particular property, setting, or performance measurement from within the MPI implementation. The interface is split into two parts: the first part provides information about and supports the setting of control variables through which the MPI implementation tunes its configuration. The second part provides access to performance variables that can provide insight into internal performance information of the MPI implementation.

To avoid restrictions on the MPI implementation, the MPI tool information interface allows the implementation to specify which control and performance variables exist. Additionally, the user of the MPI tool information interface can obtain metadata about each available variable, such as its datatype, and a textual description. The MPI tool information interface provides the necessary routines to find all variables that exist in a particular MPI implementation, to query their properties, to retrieve descriptions about their meaning, and to access and, if appropriate, to alter their values.

The MPI tool information interface can be used independently from the MPI communication functionality. In particular, the routines of this interface can be called before MPI_INIT (or equivalent) and after MPI_FINALIZE. In order to support this behavior cleanly, the MPI tool information interface uses separate initialization and finalization routines. All identifiers used in the MPI tool information interface have the prefix MPI_T_.

On success, all MPI tool information interface routines return MPI_SUCCESS, otherwise they return an appropriate and unique return code indicating the reason why the call was not successfully completed. Details on return codes can be found in Section 2.3.9. However, unsuccessful calls to the MPI tool information interface are not fatal and do not impact the execution of subsequent MPI routines.

Since the MPI tool information interface primarily focuses on tools and support libraries, MPI implementations are only required to provide C bindings for functions introduced in this Section 2.3. Except where otherwise noted, all conventions and principles governing the C bindings of the MPI API also apply to the MPI tool information interface, which is available by including the mpi.h header file. All routines in this interface have local semantics.

Advice to users. The number and type of control variables and performance variables can vary between MPI implementations, platforms and different builds of the same implementation on the same platform as well as between runs. Hence, any application relying on a particular variable will not be portable. Further, there is no guarantee

that number of variables, variable indices, and variable names are the same across processes.

This interface is primarily intended for performance monitoring tools, support tools, and libraries controlling the application's environment. When maximum portability is desired, application programmers should either avoid using the MPI tool information interface or avoid being dependent on the existence of a particular control or performance variable. (*End of advice to users.*)

2.3.1 Verbosity Levels

The MPI tool information interface provides access to internal configuration and performance information through a set of control and performance variables defined by the MPI implementation. Since some implementations may export a large number of variables, variables are classified by a verbosity level that categorizes both their intended audience (end users, performance tuners or MPI implementors) and a relative measure of level of detail (basic, detailed or all). These verbosity levels are described by a single integer. Table 2.1 lists the constants for all possible verbosity levels. The values of the constants are monotonic in the order listed in the table; i.e., `MPI_T_VERBOSITY_USER_BASIC < MPI_T_VERBOSITY_USER_DETAIL < ... < MPI_T_VERBOSITY_MPIDEV_ALL`.

<code>MPI_T_VERBOSITY_USER_BASIC</code>	Basic information of interest to users
<code>MPI_T_VERBOSITY_USER_DETAIL</code>	Detailed information of interest to users
<code>MPI_T_VERBOSITY_USER_ALL</code>	All information of interest to users
<code>MPI_T_VERBOSITY_TUNER_BASIC</code>	Basic information required for tuning
<code>MPI_T_VERBOSITY_TUNER_DETAIL</code>	Detailed information required for tuning
<code>MPI_T_VERBOSITY_TUNER_ALL</code>	All information required for tuning
<code>MPI_T_VERBOSITY_MPIDEV_BASIC</code>	Basic information for MPI implementors
<code>MPI_T_VERBOSITY_MPIDEV_DETAIL</code>	Detailed information for MPI implementors
<code>MPI_T_VERBOSITY_MPIDEV_ALL</code>	All information for MPI implementors

Table 2.1: MPI tool information interface verbosity levels.

2.3.2 Binding MPI Tool Information Interface Variables to MPI Objects

Each MPI tool information interface variable provides access to a particular control setting or performance property of the MPI implementation. A variable may refer to a specific MPI object such as a communicator, datatype, or one-sided communication window, or the variable may refer more generally to the MPI environment of the process. Except for the last case, the variable must be bound to exactly one MPI object before it can be used. Table 2.2 lists all MPI object types to which an MPI tool information interface variable can be bound, together with the matching constant that MPI tool information interface routines return to identify the object type.

Rationale. Some variables have meanings tied to a specific MPI object. Examples include the number of send or receive operations using a particular datatype, the number of times a particular error handler has been called, or the communication protocol and “eager limit” used for a particular communicator. Creating a new MPI

Constant	MPI object
MPI_T_BIND_NO_OBJECT	N/A; applies globally to entire MPI process
MPI_T_BIND_MPI_COMM	MPI communicators
MPI_T_BIND_MPI_DATATYPE	MPI datatypes
MPI_T_BIND_MPI_ERRHANDLER	MPI error handlers
MPI_T_BIND_MPI_FILE	MPI file handles
MPI_T_BIND_MPI_GROUP	MPI groups
MPI_T_BIND_MPI_OP	MPI reduction operators
MPI_T_BIND_MPI_REQUEST	MPI requests
MPI_T_BIND_MPI_WIN	MPI windows for one-sided communication
MPI_T_BIND_MPI_MESSAGE	MPI message object
MPI_T_BIND_MPI_INFO	MPI info object

Table 2.2: Constants to identify associations of variables.

tool information interface variable for each MPI object would cause the number of variables to grow without bounds, since they cannot be reused to avoid naming conflicts. By associating MPI tool information interface variables with a specific MPI object, the MPI implementation only must specify and maintain a single variable, which can then be applied to as many MPI objects of the respective type as created during the program's execution. (*End of rationale.*)

2.3.3 Convention for Returning Strings

Several MPI tool information interface functions return one or more strings. These functions have two arguments for each string to be returned: an OUT parameter that identifies a pointer to the buffer in which the string will be returned, and an IN/OUT parameter to pass the length of the buffer. The user is responsible for the memory allocation of the buffer and must pass the size of the buffer (n) as the length argument. Let n be the length value specified to the function. On return, the function writes at most $n - 1$ of the string's characters into the buffer, followed by a null terminator. If the returned string's length is greater than or equal to n , the string will be truncated to $n - 1$ characters. In this case, the length of the string plus one (for the terminating null character) is returned in the length argument. If the user passes the null pointer as the buffer argument or passes 0 as the length argument, the function does not return the string and only returns the length of the string plus one in the length argument. If the user passes the null pointer as the length argument, the buffer argument is ignored and nothing is returned.

2.3.4 Initialization and Finalization

The MPI tool information interface requires a separate set of initialization and finalization routines.

`MPI_T_INIT_THREAD(required, provided)`

IN	required	desired level of thread support (integer)
OUT	provided	provided level of thread support (integer)

`int MPI_T_init_thread(int required, int *provided)`

All programs or tools that use the MPI tool information interface must initialize the MPI tool information interface in the processes that will use the interface before calling any other of its routines. A user can initialize the MPI tool information interface by calling `MPI_T_INIT_THREAD`, which can be called multiple times. In addition, this routine initializes the thread environment for all routines in the MPI tool information interface. Calling this routine when the MPI tool information interface is already initialized has no effect beyond increasing the reference count of how often the interface has been initialized. The argument `required` is used to specify the desired level of thread support. The possible values and their semantics are identical to the ones that can be used with `MPI_INIT_THREAD` listed in Section ???. The call returns in `provided` information about the actual level of thread support that will be provided by the MPI implementation for calls to MPI tool information interface routines. It can be one of the four values listed in Section ??.

The MPI specification does not require all MPI processes to exist before the call to `MPI_INIT`. If the MPI tool information interface is used before `MPI_INIT` has been called, `MPI_T_INIT_THREAD` must be called on each process that will use the MPI tool information interface. Processes created by the MPI implementation during `MPI_INIT` inherit the status of the MPI tool information interface (whether it is initialized or not as well as all active sessions and handles) from the process from which they are created.

Processes created at runtime as a result of calls MPI's dynamic process management require their own initialization before they can use the MPI tool information interface.

Advice to users. If `MPI_T_INIT_THREAD` is called before `MPI_INIT_THREAD`, the requested and granted thread level for `MPI_T_INIT_THREAD` may influence the behavior and return value of `MPI_INIT_THREAD`. The same is true for the reverse order. (*End of advice to users.*)

Advice to implementors. MPI implementations should strive to make as many control or performance variables available before `MPI_INIT` (instead of adding them within `MPI_INIT`) to allow tools the most flexibility. In particular, control variables should be available before `MPI_INIT` if their value cannot be changed after `MPI_INIT`. (*End of advice to implementors.*)

`MPI_T_FINALIZE()`

`int MPI_T_finalize(void)`

This routine finalizes the use of the MPI tool information interface and may be called as often as the corresponding `MPI_T_INIT_THREAD` routine up to the current point of execution. Calling it more times returns a corresponding error code. As long as the number of calls to `MPI_T_FINALIZE` is smaller than the number of calls to `MPI_T_INIT_THREAD` up to the current point of execution, the MPI tool information interface remains initialized

and calls to its routines are permissible. Further, additional calls to `MPI_T_INIT_THREAD` after one or more calls to `MPI_T_FINALIZE` are permissible.

Once `MPI_T_FINALIZE` is called the same number of times as the routine `MPI_T_INIT_THREAD` up to the current point of execution, the MPI tool information interface is no longer initialized. The interface can be reinitialized by subsequent calls to `MPI_T_INIT_THREAD`.

At the end of the program execution, unless `MPI_ABORT` is called, an application must have called `MPI_T_INIT_THREAD` and `MPI_T_FINALIZE` an equal number of times.

2.3.5 Datatype System

All variables managed through the MPI tool information interface represent their values through typed buffers of a given length and type using an MPI datatype (similar to regular send/receive buffers). Since the initialization of the MPI tool information interface is separate from the initialization of MPI, MPI tool information interface routines can be called before `MPI_INIT`. Consequently, these routines can also use MPI datatypes before `MPI_INIT`. Therefore, within the context of the MPI tool information interface, it is permissible to use a subset of MPI datatypes as specified below before a call to `MPI_INIT` (or equivalent).

```
MPI_INT
MPI_UNSIGNED
MPI_UNSIGNED_LONG
MPI_UNSIGNED_LONG_LONG
MPI_COUNT [ticketcount.][] If the COUNT ticket is passed
MPI_CHAR
MPI_DOUBLE
```

Table 2.3: MPI datatypes that can be used by the MPI tool information interface.

Rationale. The MPI tool information interface relies mainly on unsigned datatypes for integer values since most variables are expected to represent counters or resource sizes. `MPI_INT` is provided for additional flexibility and is expected to be used mainly for control variables and enumeration types (see below).

Providing all basic datatypes, in particular providing all signed and unsigned variants of integer types, would lead to a larger number of types, which tools need to interpret. This would cause unnecessary complexity in the implementation of tools based on the MPI tool information interface. (*End of rationale.*)

The MPI tool information interface only relies on a subset of the basic MPI datatypes and does not use any derived MPI datatypes. Table 2.3 lists all MPI datatypes that can be returned by the MPI tool information interface to represent its variables.

Rationale. The MPI tool information interface requires a significantly simpler type system than MPI itself. Therefore, only its required subset must be present before `MPI_INIT` (or equivalent) and MPI implementations do not need to initialize the complete MPI datatype system. (*End of rationale.*)

For variables of type `MPI_INT`, an MPI implementation can provide additional information by associating names with a fixed number of values. We refer to this information in the following as an enumeration. In this case, the respective calls that provide additional meta-data for each control or performance variable, i.e., `MPI_T_CVAR_GET_INFO` (Section 2.3.6) and `MPI_T_PVAR_GET_INFO` (Section 2.3.7), return a handle of type `MPI_T_enum` that can be passed to the following functions to extract additional information. Thus, the MPI implementation can describe variables with a fixed set of values that each represents a particular state. Each enumeration type can have N different values, with a fixed N that can be queried using `MPI_T_ENUM_GET_INFO`.

`MPI_T_ENUM_GET_INFO(enumtype, num, name, name_len)`

IN	<code>enumtype</code>	enumeration to be queried (handle)
OUT	<code>num</code>	number of discrete values represented by this enumeration (integer)
OUT	<code>name</code>	buffer to return the string containing the name of the enumeration (string)
INOUT	<code>name_len</code>	length of the string and/or buffer for <code>name</code> (integer)

```
int MPI_T_enum_get_info(MPI_T_enum enumtype, int *num, char *name, int
                        *name_len)
```

If `enumtype` is a valid enumeration, this routine returns the number of possible values represented by this enumeration type. `range` and the name of the enumeration. N must be greater than 0, i.e., the enumeration must represent at least one value.

The arguments `name` and `name_len` are used to return the name of the enumerations as described in Section 2.3.3.

The routine is required to return a name of at least length one. This name must be unique with respect to all other names for enumerations that the MPI implementation uses.

Names associated with individual values in each enumeration `enumtype` can be queried using `MPI_T_ENUM_GET_ITEM`.

`MPI_T_ENUM_GET_ITEM(datatype, num, value, name, name_len)`

IN	<code>enumtype</code>	enumeration to be queried (handle)
IN	<code>num</code>	number of the value to be queried in this enumeration (integer)
OUT	<code>value</code>	variable value (integer)
OUT	<code>name</code>	buffer to return the string containing the name of the enumeration item (string)
INOUT	<code>name_len</code>	length of the string and/or buffer for <code>name</code> (integer)

```
int MPI_T_enum_get_item(MPI_T_enum enumtype, int item, char *name, int
                        *name_len)
```

The arguments `name` and `name_len` are used to return the name of the enumeration item as described in Section 2.3.3.

If completed successfully, the routine returns the value/name pair of the *num*-th value describing the enumeration. The call is further required to return a name of at least length one. This name must be unique with respect to all other names of items for the same enumeration.

2.3.6 Control Variables

The routines described in this section of the MPI tool information interface specification focus on the ability to list, query, and possibly set control variables exposed by the MPI implementation. These variables can typically be used by the user to fine tune properties and configuration settings of the MPI implementation. On many systems, such variables can be set using environment variables, although other configuration mechanisms may be available, such as configuration files or central configuration registries. A typical example that is available in several existing MPI implementations is the ability to specify an “eager limit”, i.e., an upper bound on the size of messages sent or received using an eager protocol.

Control Variable Query Functions

An MPI implementation exports a set of N control variables through the MPI tool information interface. If N is zero, then the MPI implementation does not export any control variables, otherwise the provided control variables are indexed from 0 to $N - 1$. This index number is used in subsequent calls to identify the individual variables.

An MPI implementation is allowed to increase the number of control variables during the execution of an MPI application when new variables become available through dynamic loading. However, MPI implementations are not allowed to change the index of a control variable or delete a variable once it has been added to the set. When variables become inactive, e.g., through dynamic unloading, accessing its value should return a corresponding error code.

Advice to users. While the MPI tool information interface guarantees that indices or variable properties do not change during a particular run of an MPI program, it does not provide a similar guarantee between runs. (*End of advice to users.*)

The following function can be used to query the number of control variables, *num_cvar*:

```
MPI_T_CVAR_GET_NUM(num_cvar)
```

```
OUT      num_cvar          returns number of control variables (integer)
```

```
int MPI_T_cvar_get_num(int *num_cvar)
```

The function `MPI_T_CVAR_GET_INFO` provides access to additional information for each variable.


```
MPI_T_CVAR_GET_INFO(cvar_index, name, name_len, verbosity, datatype, enumtype, desc,
                    desc_len, bind, scope)
```

IN	cvar_index	index of the control variable to be queried, value between 0 and <i>num_cvar</i> - 1 (integer)
OUT	name	buffer to return the string containing the name of the control variable (string)
INOUT	name_len	length of the string and/or buffer for <i>name</i> (integer)
OUT	verbosity	verbosity level of this variable (integer)
OUT	datatype	MPI datatype of the information stored in the control variable (handle)
OUT	enumtype	optional descriptor for enumeration information (handle)
OUT	desc	buffer to return the string containing a description of the control variable (string)
INOUT	desc_len	length of the string and/or buffer for <i>desc</i> (integer)
OUT	bind	type of MPI object to which this variable must be bound (integer)
OUT	scope	scope of when changes to this variable are possible (integer)

```
int MPI_T_cvar_get_info(int cvar_index, char *name, int *name_len, int
                        *verbosity, MPI_Datatype *datatype, MPI_T_enum *enumtype, char
                        *desc, int *desc_len, int *bind, int *scope)
```

After a successful call to `MPI_T_CVAR_GET_INFO` for a particular variable, subsequent calls to this routine querying information about the same variable must return the same information. An MPI implementation is not allowed to alter any of the returned values.

The arguments *name* and *name_len* are used to return the name of the control variable as described in Section 2.3.3.

If completed successfully, the routine is required to return a name of at least length one. The name must be unique with respect to all other names for control variables used by the MPI implementation.

The argument *verbosity* returns the verbosity level of the variable (see Section 2.3.1).

The argument *datatype* returns the MPI datatype that is used to represent the control variable.

If the variable is of type `MPI_INT`, MPI can optionally specify an enumeration for the values represented by this variable and return it in *enumtype*. In this case, MPI returns an enumeration identifier, which can then be used to gather more information as described in Section 2.3.5. If the datatype is not `MPI_INT` or the argument *enumtype* is the constant `MPI_T_ENUM_NULL`, no enumeration type is returned.

The arguments *desc* and *desc_len* are used to return a description of the control variable as described in Section 2.3.3.

Returning a description is optional. If an MPI implementation decides not to return a description, the first character for *desc* must be set to the null character and *desc_len* must be set to one at the return of this call.

The parameter `bind` returns the type of the MPI object to which the variable must be bound or the value `MPI_T_BIND_NO_OBJECT` (see Section 2.3.2).

The scope of a variable determines whether changing a variable's value is either local to the process or must be done by the user across multiple processes. The latter is further split into variables that require changes in a group of processes and those that require collective changes among all connected processes. Both cases can require all processes to either be set to consistent (but potentially different) values or to equal values on every participating process. The description provided with the variable must contain an explanation about the requirements and/or restrictions for setting the particular variable.

On successful return from `MPI_T_CVAR_GET_INFO`, the argument `scope` will be set to one of the constants listed in Table 2.4.

Scope Constant	Description
<code>MPI_T_SCOPE_READONLY</code>	read-only, cannot be written
<code>MPI_T_SCOPE_LOCAL</code>	may be writeable, writing is a local operation
<code>MPI_T_SCOPE_GROUP</code>	may be writeable, must be done to a group of processes, all processes in a group must be set to consistent values
<code>MPI_T_SCOPE_GROUP_EQ</code>	may be writeable, must be done to a group of processes, all processes in a group must be set to the same value
<code>MPI_T_SCOPE_ALL</code>	may be writeable, must be done to all processes, all connected processes must be set to consistent values
<code>MPI_T_SCOPE_ALL_EQ</code>	may be writeable, must be done to all processes, all connected processes must be set to the same value

Table 2.4: Scopes for control variables.

Advice to users. The `scope` of a variable only indicates if a variable might be changeable; it is not a guarantee that it can be changed at any time. (*End of advice to users.*)

Example: Printing All Control Variables

Example 2.1

The following example shows how the MPI tool information interface can be used to query and print the names of all available control variables.

```
#include <stdio.h>
#include <stdlibh.h>
#include <mpi.h>

int main(int argc, char **argv) {
    int i, err, num, namelen, bind, verbose, scope;
    int threadsupport;
    char name[100];
    MPI_Datatype datatype;
```

```

err=MPI_T_init_thread(MPI_THREAD_SINGLE,&threadsupport);
if (err!=MPI_SUCCESS)
    return err;

err=MPI_T_cvar_get_num(&num);
if (err!=MPI_SUCCESS)
    return err;

for (i=0; i<num; i++) {
    namelen=100;
    err=MPI_T_cvar_get_info(i, name, &namelen,
        &verbose, &datatype, MPI_T_ENUM_NULL,
        NULL, NULL, /*no description */
        &bind, &scope);
    if (err!=MPI_SUCCESS) return err;
    printf("Var %i: %s\n", i, name);
}

err=MPI_T_finalize();
if (err!=MPI_SUCCESS)
    return 1;
else
    return 0;
}

```

Handle Allocation and Deallocation

Before reading or writing the value of a variable, a user must first allocate a handle of type `MPI_T_cvar_handle` for the variable by binding it to an MPI object (see also Section 2.3.2).

Rationale. Handles used in the MPI tool information interface are distinct from handles used in the remaining parts of the MPI standard because they must be usable before `MPI_INIT` and after `MPI_FINALIZE`. Further, accessing handles, in particular for performance variables, can be time critical and having a separate handle space enables optimizations. (*End of rationale.*)

```
1 MPI_T_CVAR_HANDLE_ALLOC(cvar_index, object, handle, count)
```

2	IN	cvar_index	index of control variable for which handle is to be allocated (index)
3			
4			
5	IN	obj_handle	reference to a handle of the MPI object to which this variable is supposed to be bound (pointer)
6			
7	OUT	handle	allocated handle (handle)
8	OUT	count	number of elements used to represent this variable (integer)
9			
10			

```
11
12 int MPI_T_cvar_handle_alloc(int cvar_index, void *obj_handle,
13 MPI_T_cvar_handle *handle, int *count)
```

14 This routine binds the control variable specified by the argument `index` to an MPI
 15 object. The object is passed in the argument `obj_handle` as an address to a local variable
 16 that stores the object's handle. The handle allocate to reference the variable is returned
 17 in the argument `handle`. Upon successful return, `count` contains the number of elements (of
 18 the datatype returned by a previous `MPI_T_CVAR_GET_INFO` call) used to represent this
 19 variable.
 20

21 *Advice to users.* The `count` can be different based on the MPI object to which is
 22 was bound. For example, variables bound to communicators could have a count that
 23 matches the size of the communicator.

24 It is not portable to pass references to predefined MPI object handles, such as
 25 `MPI_COMM_WORLD` to this routine, since their implementation depends on the MPI
 26 library. Instead, such object handles should be stored in a local variable and the
 27 address of this local variables should be passed into `MPI_T_CVAR_HANDLE_ALLOC`.
 28 (*End of advice to users.*)
 29

30 The value of `cvar_index` should be in the range 0 to `num_cvar - 1`, where `num_cvar`
 31 is the number of available control variables as determined from a prior call to
 32 `MPI_T_CVAR_GET_NUM`. The type of the MPI object it references must be consistent
 33 with the type returned in the `bind` argument in a prior call to `MPI_T_CVAR_GET_INFO`.
 34

35 In the case the `bind` argument returned by `MPI_T_CVAR_GET_INFO` equals
 36 `MPI_T_BIND_NO_OBJECT`, the argument `obj_handle` is ignored.
 37

```
38 MPI_T_CVAR_HANDLE_FREE(handle)
```

39	INOUT	handle	handle to be freed (handle)
40			

```
41
42 int MPI_T_cvar_handle_free(MPI_T_cvar_handle *handle)
```

43 When a handle is no longer needed, a user of the MPI tool information interface should
 44 call `MPI_T_CVAR_HANDLE_FREE` to free the handle and the associated resources in the
 45 MPI implementation. On a successful return, MPI sets the handle to
 46 `MPI_T_CVAR_HANDLE_NULL`.
 47
 48

Control Variable Access Functions

MPI_T_CVAR_READ(handle, buf)

IN	handle	handle to the control variable to be read (handle)
OUT	buf	initial address of storage location for variable value (choice)

```
int MPI_T_cvar_read(MPI_T_cvar_handle handle, void* buf)
```

This routine queries the value of the control variable identified by the argument `handle` and stores the result in the buffer identified by the parameter `buf`. The user must ensure that the buffer is of the appropriate size to hold the entire value of the control variable (based on the returned datatype and count from prior corresponding calls to `MPI_T_CVAR_GET_INFO` and `MPI_T_CVAR_HANDLE_ALLOC`, respectively).

MPI_T_CVAR_WRITE(handle, buf)

IN	handle	handle to the control variable to be written (handle)
IN	buf	initial address of storage location for variable value (choice)

```
int MPI_T_cvar_write(MPI_T_cvar_handle handle, const void* buf)
```

This routine sets the value of the control variable identified by the argument `handle` to the data stored in the buffer identified by the parameter `buf`. The user must ensure that the buffer is of the appropriate size to hold the entire value of the control variable (based on the returned datatype and count from prior corresponding calls to `MPI_T_CVAR_GET_INFO` and `MPI_T_CVAR_HANDLE_ALLOC`, respectively).

If the variable has a global scope (as returned by a prior corresponding `MPI_T_CVAR_GET_INFO` call) any write call to this variable must be issued by the user in all connected (as defined in Section ??) MPI processes. If the variable has a group scope, any write call to this variable must be issued by the user in all MPI processes in the group, which must be described by the MPI implementation in the description by the `MPI_T_CVAR_GET_INFO`.

In both cases, the user must ensure that the writes in all processes are consistent. If the scope is either `MPI_T_SCOPE_GLOBAL_EQ` or `MPI_T_SCOPE_GROUP_EQ` this means that the variable in all processes must be set to the same value.

If it is not possible to change the variable at the time the call is made, the function returns either `MPI_T_ERR_CVAR_SETNOW`, if there may be a later time at which the variable could be set, or `MPI_T_ERR_CVAR_SETNEVER`, if the variable cannot be set for the remainder of the application's execution.

Example: Reading the Value of a Control Variable

Example 2.2

The following example shows a routine that can be used to query the value with a control variable with a given index. The example assumes that the variable is intended to be bound to an MPI communicator.

```

1  int getValue_int_comm(int index, MPI_Comm comm, int *val) {
2      int err,count;
3      MPI_T_cvar_handle handle;
4
5      /* This is example assumes that the variable index */
6      /* can be bound to a communicator */
7
8
9      err=MPI_T_cvar_handle_alloc(index,&comm,&handle,&count);
10     if (err!=MPI_SUCCESS) return err;
11
12     /* The following assumes that the variable is */
13     /* represented by a single integer */
14
15     err=MPI_T_cvar_read(handle,val);
16     if (err!=MPI_SUCCESS) return err;
17
18     err=MPI_T_cvar_handle_free(&handle);
19     return err;
20 }
21
22
23
24

```

2.3.7 Performance Variables

The following section focuses on the ability to list and query performance variables provided by the MPI implementation. Performance variables provide insight into MPI implementation specific internals and can represent information such as the state of the MPI implementation (e.g., waiting blocked, receiving, not active), aggregated timing data for submodules, or queue sizes and lengths.

Rationale. The interface for performance variables is separate from the interface for control variables, since performance variables have different requirements and parameters. By keeping them separate, the interface provides cleaner semantics and allows for more performance optimization opportunities. (*End of rationale.*)

Performance Variable Classes

Each performance variable is associated with a class that describes its basic semantics, possible datatypes, basic behavior, its starting value, whether it can overflow, and when and how an MPI implementation can change the variable's value. The starting value is the value the variable assumes when it is used for the first time or whenever it is reset.

Advice to users. If a performance variable belongs to a class that can overflow, it is up to the user to appropriately protect against this, e.g., by frequently reading and resetting the variable value. (*End of advice to users.*)

Advice to implementors. MPI implementations should use large enough datatypes for each performance variable to avoid overflows under normal circumstances. (*End of advice to implementors.*)

The classes are defined by the following constants:

- **MPI_T_PVAR_CLASS_STATE**

A performance variable in this class represents a set of discrete states. Variables of this class are represented by MPI_INT and can be set by the MPI implementation at any time. Variables of this type should be described further using an enumeration, as discussed in Section 2.3.5. The starting value is the current state of the implementation at the time the starting value is set. MPI implementations must ensure that variables of this class cannot overflow.

- **MPI_T_PVAR_CLASS_LEVEL**

A performance variable in this class represents a value that describes the utilization level of a resource. The value of a variable of this class can change at any time to match the current utilization level of the resource. Values returned from variables in this class are non-negative and represented by one of the following datatypes: MPI_UNSIGNED, MPI_UNSIGNED_LONG, MPI_UNSIGNED_LONG_LONG, MPI_DOUBLE. The starting value is the current utilization level of the resource at the time the starting value is set. MPI implementations must ensure that variables of this class cannot overflow.

- **MPI_T_PVAR_CLASS_SIZE**

A performance variable in this class represents a value that is the fixed size of a resource. Values returned from variables in this class are non-negative and represented by one of the following datatypes: MPI_UNSIGNED, MPI_UNSIGNED_LONG, MPI_UNSIGNED_LONG_LONG, MPI_DOUBLE. The starting value is the current utilization level of the resource at the time the starting value is set. MPI implementations must ensure that variables of this class cannot overflow.

- **MPI_T_PVAR_CLASS_PERCENTAGE**

The value of a performance variable in this class represents the percentage utilization of a finite resource. The value of a variable of this class can change at any time to match the current utilization level of the resource. It will be returned as an MPI_DOUBLE datatype. The value must always be between 0.0 (resource not used at all) and 1.0 (resource completely used). The starting value is the current percentage utilization level of the resource at the time the starting value is set. MPI implementations must ensure that variables of this class cannot overflow.

- **MPI_T_PVAR_CLASS_HIGHWATERMARK**

A performance variable in this class represents a value that describes the high watermark utilization of a resource. The value of a variable of this class is non-negative and grows monotonically from the initialization or reset of the variable. It can be represented by one of the following datatypes: MPI_UNSIGNED, MPI_UNSIGNED_LONG, MPI_UNSIGNED_LONG_LONG, MPI_DOUBLE. The starting value is the current utilization level of the resource at the time the starting value is set. MPI implementations must ensure that variables of this class cannot overflow.

1 • **MPI_T_PVAR_CLASS_LOWWATERMARK**

2 A performance variable in this class represents a value that describes the low wa-
 3 termark utilization of a resource. The value of a variable of this class is non-
 4 negative and decreases monotonically from the initialization or reset of the vari-
 5 able. It can be represented by one of the following datatypes: MPI_UNSIGNED,
 6 MPI_UNSIGNED_LONG, MPI_UNSIGNED_LONG_LONG, MPI_DOUBLE. The
 7 starting value is the current utilization level of the resource at the time the start-
 8 ing value is set. MPI implementations must ensure that variables of this class cannot
 9 overflow.

10 • **MPI_T_PVAR_CLASS_COUNTER**

11 A performance variable in this class counts the number of occurrences of a specific
 12 event (e.g., the number of memory allocations within an MPI library). The value of
 13 a variable of this class increases monotonically from the initialization or reset of the
 14 performance variable by one for each specific event that is observed. Values must be
 15 non-negative and represented by one of the following datatypes: MPI_UNSIGNED,
 16 MPI_UNSIGNED_LONG, MPI_UNSIGNED_LONG_LONG. The starting value for
 17 variables of this class is 0. Variables of this class can overflow.

18 • **MPI_T_PVAR_CLASS_AGGREGATE**

19 The value of a performance variable in this class is an an aggregated value that repre-
 20 sents a sum of arguments processed during a specific event (e.g., the amount of mem-
 21 ory allocated by all memory allocations). This class is similar to the counter class,
 22 but instead of counting individual events, the value can be incremented by arbitrary
 23 amounts. The value of a variable of this class increases monotonically from the initial-
 24 ization or reset of the performance variable. It must be non-negative and represented
 25 by one of the following datatypes: MPI_UNSIGNED, MPI_UNSIGNED_LONG,
 26 MPI_UNSIGNED_LONG_LONG, MPI_DOUBLE. The starting value for variables
 27 of this class is 0. Variables of this class can overflow.

28 • **MPI_T_PVAR_CLASS_TIMER**

29 The value of a performance variable in this class represents the aggregated time
 30 that the MPI implementation spends executing a particular event, type of event,
 31 or section of the MPI library. This class has the same basic semantics as
 32 MPI_T_PVAR_CLASS_AGGREGATE, but explicitly records a timing value. The
 33 value of a variable of this class increases monotonically from the initialization
 34 or reset of the performance variable. It must be non-negative and represented
 35 by one of the following datatypes: MPI_UNSIGNED, MPI_UNSIGNED_LONG,
 36 MPI_UNSIGNED_LONG_LONG, MPI_T_DOUBLE. The starting value for vari-
 37 ables of this class is 0. If the type MPI_DOUBLE is used, the units representing time in
 38 this datatype must match the units used by MPI_WTIME. Otherwise, the time units
 39 should be documented, e.g., in the description returned by MPI_T_PVAR_GET_INFO.
 40 Variables of this class can overflow.

41 • **MPI_T_PVAR_CLASS_GENERIC**

42 This class can be used to describe a variable that does not fit into any of the other
 43 classes. For variables in this class, the starting value is variable specific and imple-
 44 mentation defined.

Performance Variable Query Functions

An MPI implementation exports a set of N performance variables through the MPI tool information interface. If N is zero, then the MPI implementation does not export any performance variables, otherwise the provided performance variables are indexed from 0 to $N - 1$. This index number is used in subsequent calls to identify the individual variables.

An MPI implementation is allowed to increase the number of performance variables during the execution of an MPI application when new variables become available through dynamic loading. However, MPI implementations are not allowed to change the index of a performance variable or delete a variable once it has been added to the set. When variables become inactive, e.g., through dynamic unloading, accessing its value should return a corresponding error code.

The following function can be used to query the number of performance variables, N :

`MPI_T_PVAR_GET_NUM(num_pvar)`

OUT `num_pvar` returns number of performance variables (integer)

`int MPI_T_pvar_get_num(int *num_pvar)`

The function `MPI_T_PVAR_GET_INFO` provides access to additional information for each variable.

```

1 MPI_T_PVAR_GET_INFO(pvar_index, name, name_len, verbosity, varclass, datatype, enum-
2 type, desc, desc_len, bind, readonly, continuous, atomic)
3
4 IN      pvar_index      index of the performance variable to be queried be-
5                          tween 0 and num_pvar - 1 (integer)
6
7 OUT     name            buffer to return the string containing the name of the
8                          performance variable (string)
9
10 INOUT   name_len        length of the string and/or buffer for name (integer)
11
12 OUT     verbosity        verbosity level of this variable (integer)
13
14 OUT     var_class        class of performance variable (integer)
15
16 OUT     datatype          MPI datatype of the information stored in the perfor-
17                          mance variable (handle)
18
19 OUT     enumtype          optional descriptor for enumeration information (han-
20                          dle)
21
22 OUT     desc              buffer to return the string containing a description of
23                          the performance variable (string)
24
25 INOUT   desc_len        length of the string and/or buffer for desc (integer)
26
27 OUT     bind              type of MPI object to which this variable must be
28                          bound (integer)
29
30 OUT     readonly          flag indicating whether the variable can be written/reset
31                          (integer)
32
33 OUT     continuous        flag indicating whether the variable can be started and
34                          stopped or is continuously active (integer)
35
36 OUT     atomic            flag indicating whether the variable can be atomically
37                          read and reset (integer)
38
39 int MPI_T_pvar_get_info(int pvar_index, char *name, int *name_len, int
40                          *verbosity, int *var_class, MPI_Datatype *datatype, MPI_T_enum
41                          *enumtype, char *desc, int *desc_len, int *bind, int
42                          *readonly, int *continuous, int *atomic)

```

After a successful call to `MPI_T_PVAR_GET_INFO` for a particular variable, subsequent calls to this routine querying information about the same variable must return the same information. An MPI implementation is not allowed to alter any of the returned values.

The arguments `name` and `name_len` are used to return the name of the performance variable as described in Section 2.3.3. If completed successfully, the routine is required to return a name of at least length one.

The argument `verbosity` returns the verbosity level of the variable (see Section 2.3.1).

The class of the performance variable is returned in the parameter `var_class`. The class must be one of the constants defined in Section 2.3.7.

The combination of the name and the class of the performance variable must be unique with respect to all other names for performance variables used by the MPI implementation.

Advice to implementors. Groups of variables that belong closely together, but have different classes, can have the same name. This is, e.g., useful to refer to multiple

variables that describe a single resource (like the level, the total size, as well as high and low watermarks). (*End of advice to implementors.*)

The argument `datatype` returns the MPI datatype that is used to represent the performance variable.

If the variable is of type `MPI_INT`, MPI can optionally specify an enumeration for the values represented by this variable and return it in `enumtype`. In this case, MPI returns an enumeration identifier, which can then be used as described in Section 2.3.5 to gather more information. If the datatype is not `MPI_INT` or the argument `enumtype` is the constant `MPI_T_ENUM_NULL`, no enumeration type is returned.

Returning a description is optional. If an MPI implementation decides not to return a description, the first character for `desc` must be set to the null character and `desc_len` must be set to one at the return from this function.

The parameter `bind` returns the type of the MPI object to which the variable must be bound or the value `MPI_T_BIND_NO_OBJECT` (see Section 2.3.2).

Upon return, the argument `readonly` is set to zero if the variable can be written or reset by the user. It is set to one if the variable can only be read.

Upon return, the argument `continuous` is set to zero if the variable can be started and stopped by the user, i.e., it is possible for the user to control if and when the value of a variable is updated. It is set to one if the variable is always active and cannot be controlled by the user.

Upon return, the argument `atomic` is set to zero if the variable cannot be atomically read and reset. Only variables for which the call sets `atomic` to one, can be used in a call to `MPI_T_PVAR_READRESET`.

Performance Experiment Sessions

Within a single program, multiple components can use the MPI tool information interface. To avoid collisions with respect to accesses to performance variables, users of the MPI tool information interface must first create a session. Subsequent calls accessing performance variables can then be made within the context of this session. Any call executed in a session must not influence the results in any other session.

`MPI_T_PVAR_SESSION_CREATE(session)`

OUT `session` identifier of performance session (handle)

```
int MPI_T_pvar_session_create(MPI_T_pvar_session *session)
```

This call creates a new session for accessing performance variables and returns a handle for this session in the argument `session` of type `MPI_T_pvar_session`.

`MPI_T_PVAR_SESSION_FREE(session)`

INOUT `session` identifier of performance experiment session (handle)

```
int MPI_T_pvar_session_free(MPI_T_pvar_session *session)
```

This call frees an existing session. Calls to the MPI tool information interface can no longer be made within the context of a session after it is freed. On a successful return, MPI sets the session identifier to `MPI_T_PVAR_SESSION_NULL`.

Handle Allocation and Deallocation

Before using a performance variable, a user must first allocate a handle of type `MPI_T_pvar_handle` for the variable by binding it to an MPI object (see also Section 2.3.2).

`MPI_T_PVAR_HANDLE_ALLOC(session, pvar_index, obj_handle, handle, count)`

IN	session	identifier of performance experiment session (handle)
IN	pvar_index	index of performance variable for which handle is to be allocated (integer)
IN	obj_handle	reference to a handle of the MPI object to which this variable is supposed to be bound (pointer)
OUT	handle	allocated handle (handle)
OUT	count	number of elements used to represent this variable (integer)

```
int MPI_T_pvar_handle_alloc(MPI_T_pvar_session session, int pvar_index,
                           void *obj_handle, MPI_T_pvar_handle *handle, int *count)
```

This routine binds the performance variable specified by the argument `index` to an MPI object in the session identified by the parameter `session`. The object is passed in the argument `obj_handle` as an address to a local variable that stores the object's handle. The handle allocated to reference the variable is returned in the argument `handle`. Upon successful return, `count` contains the number of elements (of the datatype returned by a previous `MPI_T_PVAR_GET_INFO` call) used to represent this variable.

Advice to users. The count can be different based on the MPI object, to which is was bound. For example, variables bound to communicators could have a count that matches the size of the communicator.

It is not portable to pass references to predefined MPI object handles, such as `MPI_COMM_WORLD`, to this routine, since their implementation depends on the MPI library. Instead, such object handles should be stored in a local variable and the address of this local variables should be passed into `MPI_T_PVAR_HANDLE_ALLOC`. (*End of advice to users.*)

The value of `index` should be in the range 0 to `num_pvar - 1`, where `num_pvar` is the number of available control variables as determined from a prior call to `MPI_T_PVAR_GET_NUM`. The type of the MPI object it references must be consistent with the type returned in the `bind` argument in a prior call to `MPI_T_PVAR_GET_INFO`.

In the case the `bind` argument equals `MPI_T_BIND_NO_OBJECT`, the argument `obj_handle` is ignored.

MPI_T_PVAR_HANDLE_FREE(session, handle)

IN	session	identifier of performance experiment session (handle)
INOUT	handle	handle to be freed (handle)

```
int MPI_T_pvar_handle_free(MPI_T_pvar_session session, MPI_T_pvar_handle
                          *handle)
```

When a handle is no longer needed, a user of the MPI tool information interface should call **MPI_T_PVAR_HANDLE_FREE** to free the handle in the session identified by the parameter **session** and the associated resources in the MPI implementation. On a successful return, MPI sets the handle to **MPI_T_PVAR_HANDLE_NULL**.

Starting and Stopping of Performance Variables

Performance variables that have the continuous flag set during the query operation are continuously operating once a handle has been allocated. Such variables may be queried at any time, but they cannot be started or stopped by the user. All other variables are in a stopped state after their handle has been allocated; their values are not updated until they have been started by the user.

MPI_T_PVAR_START(session, handle)

IN	session	identifier of performance experiment session (handle)
IN	handle	handle of a performance variable (handle)

```
int MPI_T_pvar_start(MPI_T_pvar_session session, MPI_T_pvar_handle handle)
```

This functions starts the performance variable with the handle identified by the parameter **handle** in the session identified by the parameter **session**.

If the constant **MPI_T_PVAR_ALL_HANDLES** is passed in **handle**, the MPI implementation attempts to start all variables within the session identified by the parameter **session** for which handles have been allocated. In this case, the routine returns **MPI_SUCCESS** if all variables are started successfully, otherwise **MPI_T_ERR_PVAR_NOSTARTSTOP** is returned. Continuous variables and variables that are already started are ignored when **MPI_T_PVAR_ALL_HANDLES** is specified.

MPI_T_PVAR_STOP(session, handle)

IN	session	identifier of performance experiment session (handle)
IN	handle	handle of a performance variable (handle)

```
int MPI_T_pvar_stop(MPI_T_pvar_session session, MPI_T_pvar_handle handle)
```

This functions stops the performance variable with the handle identified by the parameter **handle** in the session identified by the parameter **session**.

If the constant **MPI_T_PVAR_ALL_HANDLES** is passed in **handle**, the MPI implementation attempts to stop all variables within the session identified by the parameter **session** for

which handles have been allocated. In this case, the routine returns `MPI_SUCCESS` if all variables are stopped successfully, otherwise `MPI_T_ERR_PVAR_NOSTARTSTOP` is returned. Continuous variables and variables that are already stopped are ignored when `MPI_T_PVAR_ALL_HANDLES` is specified.

Performance Variable Access Functions

`MPI_T_PVAR_READ(session, handle, buf)`

IN	session	identifier of performance experiment session (handle)
IN	handle	handle of a performance variable (handle)
OUT	buf	initial address of storage location for variable value (choice)

```
int MPI_T_pvar_read(MPI_T_pvar_session session, MPI_T_pvar_handle handle,
                    void* buf)
```

The `MPI_T_PVAR_READ` call queries the value of the performance variable with the handle `handle` in the session identified by the parameter `session` and stores the result in the buffer identified by the parameter `buf`. The user is responsible to ensure that the buffer is of the appropriate size to hold the entire value of the performance variable (based on the datatype and count returned by the corresponding previous calls to `MPI_T_PVAR_GET_INFO` and `MPI_T_PVAR_HANDLE_ALLOC`, respectively).

The constant `MPI_T_PVAR_ALL_HANDLES` cannot be used as an argument for the function `MPI_T_PVAR_READ`.

`MPI_T_PVAR_WRITE(session, handle, buf)`

IN	session	identifier of performance experiment session (handle)
IN	handle	handle of a performance variable (handle)
IN	buf	initial address of storage location for variable value (choice)

```
int MPI_T_pvar_write(MPI_T_pvar_session session, MPI_T_pvar_handle handle ,
                     const void* buf)
```

The `MPI_T_PVAR_WRITE` call attempts to write the value of the performance variable with the handle identified by the parameter `handle` in the session identified by the parameter `session`. The value to be written is passed in the buffer identified by the parameter `buf`. The user must ensure that the buffer is of the appropriate size to hold the entire value of the performance variable (based on the datatype and count returned by the corresponding previous calls to `MPI_T_PVAR_GET_INFO` and `MPI_T_PVAR_HANDLE_ALLOC`, respectively).

If it is not possible to change the variable, the function returns `MPI_T_ERR_PVAR_NOWRITE`.

The constant `MPI_T_PVAR_ALL_HANDLES` cannot be used as an argument for the function `MPI_T_PVAR_WRITE`.

MPI_T_PVAR_RESET(session, handle)

IN	session	identifier of performance experiment session (handle)
IN	handle	handle of a performance variable (handle)

int MPI_T_pvar_reset(MPI_T_pvar_session session, MPI_T_pvar_handle handle)

The **MPI_T_PVAR_RESET** call sets the performance variable with the handle identified by the parameter **handle** to its starting value specified in Section 2.3.7. If it is not possible to change the variable, the function returns **MPI_T_ERR_PVAR_NOWRITE**.

If the constant **MPI_T_PVAR_ALL_HANDLES** is passed in **handle**, the MPI implementation attempts to reset all variables within the session identified by the parameter **session** for which handles have been allocated. In this case, the routine returns **MPI_SUCCESS** if all variables are reset successfully, otherwise **MPI_T_ERR_PVAR_NOWRITE** is returned. Read-only variables are ignored when **MPI_T_PVAR_ALL_HANDLES** is specified.

MPI_T_PVAR_READRESET(session, handle, buf)

IN	session	identifier of performance experiment session (handle)
IN	handle	handle of a performance variable (handle)
OUT	buf	initial address of storage location for variable value (choice)

int MPI_T_pvar_readreset(MPI_T_pvar_session session, MPI_T_pvar_handle handle, void* buf)

This call atomically combines the functionality of **MPI_T_PVAR_READ** and **MPI_T_PVAR_RESET** with the same semantics as if these two calls were called separately. If atomic operations on this variable are not supported, this routine returns **MPI_ERR_NOATOMIC**.

The constant **MPI_T_PVAR_ALL_HANDLES** can not be used as an argument for the function **MPI_T_PVAR_READRESET**.

Advice to implementors. Sampling based tools rely on the ability to call the MPI tool information interface, in particular routines to start, stop, read, write and reset performance variables, from any program context, including asynchronous contexts such as signal handlers. MPI implementations should strive, if possible in their particular environment, to enable these usage scenarios for all or a subset of the routines mentioned above. If implementing only a subset, the read, write, and reset routines are typically the most critical for sampling based tools. An MPI implementation should clearly document any restrictions on the program contexts in which the MPI tool information interface can be used. Restrictions might include guaranteeing usage outside of all signals or outside a specific set of signals. Any restrictions could be documented, for example, through the description returned by **MPI_T_PVAR_GET_INFO**. *(End of advice to implementors.)*

Rationale. All routines to read, write or reset performance variables require the session argument. This keeps the interface consistent and allows the use

MPI_T_PVAR_ALL_HANDLES where appropriate. Further, this opens up additional performance optimizations for the implementation of handles. (*End of rationale.*)

Example: Tool to Detect Receives with Long Unexpected Message Queues

Example 2.3

The following example shows a sample tool to identify receive operations that occur during times with long message queues. This examples assumes that the MPI implementation exports a variable with the name "MPI_T_UMQ_LENGTH" to represent the current length of the unexpected message queue. The tool is implemented as a PMPI tool using the MPI profiling interface.

The tool consists of three parts: (1) the initialization (by intercepting the call to MPI_INIT), (2) the test for long unexpected message queues (by intercepting calls to MPI_RECV), and (3) the clean up phase (by intercepting the call to MPI_FINALIZE. To capture all receives, the example would have to be extended to have similar wrappers for all receive operations.

Part 1— Initialization: During initialization, the tool searches for the variable and, once the right index is found, allocates a session and a handle for the variable with the found index, and starts the performance variable.

```
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>
#include <mpi.h>

/* Global variables for the tool */
static MPI_T_pvar_session session;
static MPI_T_pvar_handle handle;

int MPI_Init(int *argc, char ***argv) {
    int err, num, i, index, namelen, verbosity, var_class, bind, threadsup;
    int readonly, continuous, atomic, count;
    char name[17];
    MPI_Comm comm;

    err=PMPI_Init(argc,argv);
    if (err!=MPI_SUCCESS) return err;

    err=PMPI_T_init_thread(MPI_THREAD_SINGLE,&threadsup);
    if (err!=MPI_SUCCESS) return err;

    err=PMPI_T_pvar_get_num(&num);
    if (err!=MPI_SUCCESS) return err;
    index=-1;
    i=0;
    while ((i<num) && (index<0)) {
```



```

        namelen=17;
        err=PMPI_T_pvar_get_info(i, name, namelen, &verbosity, &var_class,
                                NULL, NULL, &bind, &readonly, &continuous, &atomic);
        if (strcmp(name,"MPI_T_UMQ_LENGTH")==0) index=i;
        i++; }

/* this could be handled in a more flexible way for a generic tool */
ASSERT(index>=0);
ASSERT(var_class==MPI_T_PVAR_CLASS_LEVEL);
ASSERT(datatype==MPI_INT);
ASSERT(bind==MPI_T_BIND_MPI_COMM);

/* Create a session */
err=PMPI_T_pvar_session_create(&session);
if (err!=MPI_SUCCESS) return err;

/* Get a handle and bind to MPI_COMM_WORLD */
comm=MPI_COMM_WORLD;
err=PMPI_T_pvar_handle_alloc(session, index, &comm, &handle, &count);
if (err!=MPI_SUCCESS) return err;

/* this could be handled in a more flexible way for a generic tool */
ASSERT(count==1);

/* Start variable */
err=PMPI_T_pvar_start(session, handle);
if (err!=MPI_SUCCESS) return err;

return MPI_SUCCESS;
}

```

Part 2 — Testing the Queue Lengths During Receives: During every receive operation, the tool reads the unexpected queue length through the matching performance variable and compares it against a predefined threshold.

```

#define THRESHOLD 5

int MPI_Recv(void *buf, int count, MPI_Datatype datatype, int source, int tag,
             MPI_Comm comm, MPI_Status *status)
{
    int value, err;

    if (comm==MPI_COMM_WORLD) {
        err=PMPI_T_pvar_read(session, handle, &value);
        if ((err==MPI_SUCCESS) && (value>THRESHOLD))
        {
            /* tool identified receive called with long UMQ */
            /* execute tool functionality, */

```

```

1          /* e.g., gather and print call stack */
2      }
3  }
4
5      return PMPI_Recv(buf, count, dt, source, tag, comm, status);
6  }
7

```

Part 3 — Termination: During MPI_FINALIZE, the MPI tool information interface needs to be finalized as well.

```

11 int MPI_Finalize()
12 {
13     int err;
14     err=MPI_T_handle_free(session, handle);
15     err=MPI_T_session_free(session);
16     err=MPI_T_finalize();
17     return PMPI_Finalize();
18 }
19

```

2.3.8 Variable Categorization

MPI implementations can optionally group performance and control variables into categories to express logical relationships between various variables. For example, an MPI implementation could group all control and performance variables that refer to message transfers in the MPI implementation and thereby distinguish them from variables that refer to local resources such as memory allocations or other interactions with the operating system.

Categories can also contain other categories to form a hierarchical grouping. Categories can never include themselves, either directly or transitively within other included categories. Expanding on the example above, this allows MPI to refine the grouping of variables referring to message transfers into variables to control and monitor message queues, message matching activities and communication protocols. Each of these groups of variables would be represented by a separate category and these categories would then be listed in a single category representing variables for message transfers.

The category information may be queried in a fashion similar to the mechanism for querying variable information. The MPI implementation exports a set of N categories via the MPI tool information interface. If $N = 0$, then the MPI implementation does not export any categories, otherwise the provided categories are indexed from 0 to $N - 1$. This index number is used in subsequent calls to functions of the MPI tool information interface to identify the individual categories.

An MPI implementation is permitted to increase the number of categories during the execution of an MPI program when new categories become available through dynamic loading. However, MPI implementations are not allowed to change the index of a category or delete it once it has been added to the set.

Similarly, MPI implementations are allowed to add variables to categories, but they are not allowed to remove variables from categories or change the order in which they are returned.

The following function can be used to query the number of control variables, N .

```
MPI_T_CATEGORY_GET_NUM(num_cat)
```

```
    OUT      num_cat      current number of categories (integer)
```

```
int MPI_T_category_get_num(int *num_cat)
```

Individual category information can then be queried by calling the following function:

```
MPI_T_CATEGORY_GET_INFO(cat_index, name, name_len, desc, desc_len, num_controlvars,
                        num_perfvars, num_categories)
```

```
    IN      cat_index      index of the category to be queried (integer)
    OUT      name           buffer to return the string containing the name of the
                          category (string)
    INOUT    name_len      length of the string and/or buffer for name (integer)
    OUT      desc           buffer to return the string containing the description
                          of the category (string)
    INOUT    desc_len      length of the string and/or buffer for desc (integer)
    OUT      num_controlvars number of control variables in the category (integer)
    OUT      num_perfvars   number of performance variables in the category (in-
                          te-
                          ger)
    OUT      num_categories number of categories contained in the category (inte-
                          ger)
```

```
int MPI_T_category_get_info(int cat_index, char *name, int *name_len, char
                          *desc, int *desc_len, int *num_controlvars, int *num_perfvars,
                          int *num_categories)
```

The arguments `name` and `name_len` are used to return the name of the category as described in Section 2.3.3.

The routine is required to return a name of at least length one. This name must be unique with respect to all other names for categories used by the MPI implementation.

The arguments `desc` and `desc_len` are used to return the description of the category as described in Section 2.3.3.

Returning a description is optional. If an MPI implementation decides not to return a description, the first character for `desc` must be set to the null character and `desc_len` must be set to one at the return of this call.

The function returns the number of control variables, performance variables and other categories contained in the queried category in the arguments `num_controlvars`, `num_perfvars`, and `num_categories`, respectively.

```
1 MPI_T_CATEGORY_GET_CVARS(cat_index, len, indices)
```

```
2     IN      cat_index      index of the category to be queried, in the range [0, N-
3                               1] (integer)
4
5     IN      len            the length of the indices array (integer)
6
7     OUT     indices        an integer array of size len, indicating control variable
8                               indices (array of integers)
```

```
9 int MPI_T_category_get_cvars(int cat_index, int len, int indices[])
10
```

11 MPI_T_CATEGORY_GET_CVARS can be used to query which control variables are
12 contained in a particular category. A category contains zero or more control variables.

```
14 MPI_T_CATEGORY_GET_PVARS(cat_index,len,indices)
```

```
16     IN      cat_index      index of the category to be queried, in the range [0, N-
17                               1] (integer)
18
19     IN      len            the length of the indices array (integer)
20
21     OUT     indices        an integer array of size len, indicating performance
22                               variable indices (array of integers)
```

```
23 int MPI_T_category_get_pvars(int cat_index, int len, int indices[])
24
```

25 MPI_T_CATEGORY_GET_PVARS can be used to query which performance variables
26 are contained in a particular category. A category contains zero or more performance
27 variables.

```
29 MPI_T_CATEGORY_GET_CATEGORIES(cat_index,len,indices)
```

```
30     IN      cat_index      index of the category to be queried, in the range [0, N-
31                               1] (integer)
32
33     IN      len            the length of the indices array (integer)
34
35     OUT     indices        an integer array of size len, indicating category indices
36                               (array of integers)
```

```
37 int MPI_T_category_get_categories(int cat_index, int len, int indices[])
38
```

39 MPI_T_CATEGORY_GET_CATEGORIES can be used to query which other categories
40 are contained in a particular category. A category contains zero or more other categories.

41 As mentioned above, MPI implementations can grow the number of categories as well
42 as the number of variables or other categories within a category. In order to allow users
43 of the MPI tool information interface to quickly check whether new categories have been
44 added or new variables or categories have been added to a category, MPI maintains a
45 virtual timestamp. This timestamp is monotonically increasing during the execution and is
46 returned by the following function:

MPI_T_CATEGORY_CHANGED(stamp)

OUT **stamp** a virtual time stamp to indicate the last change to the categories (integer)

```
int MPI_T_category_changed(int *stamp)
```

If two subsequent calls to this routine return the same timestamp, it is guaranteed that the category information has not changed between the two calls. If the timestamp retrieved from the second call is higher, then some categories have been added or expanded.

Advice to users. The timestamp value is purely virtual and only intended to check for changes in the category information. It should not be used for any other purpose. (*End of advice to users.*)

The index values returned in indices by `MPI_T_CATEGORY_GET_CVARS`, `MPI_T_CATEGORY_GET_PVARS` and `MPI_T_CATEGORY_GET_CATEGORIES` can be used as input to `MPI_T_CVAR_GET_INFO`, `MPI_T_PVAR_GET_INFO` and `MPI_T_CATEGORY_GET_INFO`, respectively.

The user is responsible for allocating the arrays passed into the functions `MPI_T_CATEGORY_GET_CVARS`, `MPI_T_CATEGORY_GET_PVARS` and `MPI_T_CATEGORY_GET_CATEGORIES`. Starting from array index 0, each function writes up to `len` elements into the array. If the category contains more than `len` elements, the function returns an arbitrary subset of size `len`. Otherwise, the entire set of elements is returned in the beginning entries of the array, and any remaining array entries are not modified.

2.3.9 Return Codes for the MPI tool information interface

All functions defined as part of the MPI tool information interface return an integer return code (see Table 2.5) to indicate whether the function has completed successfully or aborted its execution. In the latter case the return code indicates the reason for not completing the routine. None of the return codes returned by an routine impact the execution of the MPI process and do not invoke MPI error handlers. The execution of the MPI process continues as if the call would have completed. However, the MPI implementation is not required to check all user provided parameters; if a user passes invalid parameter values to any routine the behavior of the implementation is undefined.

All return codes with the prefix `MPI_T_` must be unique values and cannot overlap with any other return values returned by the MPI implementation.

2.3.10 Profiling Interface

All requirements for the profiling interfaces, as described in Section 2.2, also apply to the MPI tool information interface. All rules, guidelines, and recommendations from Section 2.2 apply equally to calls defined as part of the MPI tool information interface.

Return Code	Description
Return Codes for all Functions in the MPI tool information interface	
MPI_SUCCESS	Call completed successfully
MPI_T_ERR_MEMORY	Out of memory
MPI_T_ERR_NOTINITIALIZED	Interface not initialized
MPI_T_ERR_CANTINIT	Interface not in the state to be initialized
Return Codes for Datatype Functions: MPI_T_ENUM_*	
MPI_T_ERR_INVALIDINDEX	The enumeration index is invalid or has been deleted.
MPI_T_ERR_INVALIDITEM	The item index queried is out of range (for MPI_T_ENUMITEM only)
Return Codes for variable and category query functions: MPI_T_*.GET_INFO	
MPI_T_ERR_INVALIDINDEX	The variable or category index is invalid
Return Codes for Handle Functions: MPI_T_*.ALLOCATE,FREE	
MPI_T_ERR_INVALIDINDEX	The variable index is invalid or has been deleted
MPI_T_ERR_INVALIDHANDLE	The handle is invalid
MPI_T_ERR_OUTOFHANDLES	No more handles available
Return Codes for Session Functions: MPI_T_PVAR_SESSION_*	
MPI_T_ERR_OUTOFSESSIONS	No more sessions available
MPI_T_ERR_INVALIDSESSION	Session argument is not a valid session
Return Codes for Control Variable Access Functions:	
MPI_T_CVAR_READ, WRITE	
MPI_T_ERR_CVAR_SETNOTNOW	Variable cannot be set at this moment
MPI_T_ERR_CVAR_SETNEVER	Variable cannot be set until end of execution
MPI_T_ERR_INVALIDHANDLE	The handle is invalid
Return Codes for Performance Variable Access and Control:	
MPI_T_PVAR_START, STOP, READ, WRITE, RESET, READRESET	
MPI_T_ERR_INVALIDHANDLE	The handle is invalid
MPI_T_ERR_INVALIDSESSION	Session argument is not a valid session
MPI_T_ERR_PVAR_NOSTARTSTOP	Variable can not be started or stopped (for MPI_T_PVAR_START and MPI_T_PVAR_STOP)
MPI_T_ERR_PVAR_NOWRITE	Variable can not be written or reset (for MPI_T_PVAR_WRITE and MPI_T_PVAR_RESET)
MPI_T_NOATOMIC	Variable cannot be read and written atomically (for MPI_T_PVAR_READRESET)
Return Codes for Category Functions: MPI_T_CATEGORY_*	
MPI_T_ERR_INVALIDINDEX	The category index is invalid

Table 2.5: Return codes used in functions of the MPI tool information interface.

Chapter 3

Language Bindings Summary

In this section we summarize the specific bindings for C, Fortran, and C++. First we present the constants, type definitions, info values and keys. Then we present the routine prototypes separately for each binding. Listings are alphabetical within chapter.

3.1 Defined Values and Handles

3.1.1 Defined Constants

The C and Fortran name is listed in the left column and the C++ name is listed in the middle or right column. Constants with the type `const int` may also be implemented as literal integer constants substituted by the preprocessor.

Return Codes	
C type: <code>const int</code> (or unnamed <code>enum</code>)	C++ type: <code>const int</code>
Fortran type: <code>INTEGER</code>	(or unnamed <code>enum</code>)
<code>MPI_SUCCESS</code>	<code>MPI::SUCCESS</code>
<code>MPI_ERR_BUFFER</code>	<code>MPI::ERR_BUFFER</code>
<code>MPI_ERR_COUNT</code>	<code>MPI::ERR_COUNT</code>
<code>MPI_ERR_TYPE</code>	<code>MPI::ERR_TYPE</code>
<code>MPI_ERR_TAG</code>	<code>MPI::ERR_TAG</code>
<code>MPI_ERR_COMM</code>	<code>MPI::ERR_COMM</code>
<code>MPI_ERR_RANK</code>	<code>MPI::ERR_RANK</code>
<code>MPI_ERR_REQUEST</code>	<code>MPI::ERR_REQUEST</code>
<code>MPI_ERR_ROOT</code>	<code>MPI::ERR_ROOT</code>
<code>MPI_ERR_GROUP</code>	<code>MPI::ERR_GROUP</code>
<code>MPI_ERR_OP</code>	<code>MPI::ERR_OP</code>
<code>MPI_ERR_TOPOLOGY</code>	<code>MPI::ERR_TOPOLOGY</code>
<code>MPI_ERR_DIMS</code>	<code>MPI::ERR_DIMS</code>
<code>MPI_ERR_ARG</code>	<code>MPI::ERR_ARG</code>
<code>MPI_ERR_UNKNOWN</code>	<code>MPI::ERR_UNKNOWN</code>
<code>MPI_ERR_TRUNCATE</code>	<code>MPI::ERR_TRUNCATE</code>
<code>MPI_ERR_OTHER</code>	<code>MPI::ERR_OTHER</code>
<code>MPI_ERR_INTERN</code>	<code>MPI::ERR_INTERN</code>
<code>MPI_ERR_PENDING</code>	<code>MPI::ERR_PENDING</code>

(Continued on next page)

Return Codes (continued)

MPI_ERR_IN_STATUS	MPI::ERR_IN_STATUS
MPI_ERR_ACCESS	MPI::ERR_ACCESS
MPI_ERR_AMODE	MPI::ERR_AMODE
MPI_ERR_ASSERT	MPI::ERR_ASSERT
MPI_ERR_BAD_FILE	MPI::ERR_BAD_FILE
MPI_ERR_BASE	MPI::ERR_BASE
MPI_ERR_CONVERSION	MPI::ERR_CONVERSION
MPI_ERR_DISP	MPI::ERR_DISP
MPI_ERR_DUP_DATAREP	MPI::ERR_DUP_DATAREP
MPI_ERR_FILE_EXISTS	MPI::ERR_FILE_EXISTS
MPI_ERR_FILE_IN_USE	MPI::ERR_FILE_IN_USE
MPI_ERR_FILE	MPI::ERR_FILE
MPI_ERR_INFO_KEY	MPI::ERR_INFO_VALUE
MPI_ERR_INFO_NOKEY	MPI::ERR_INFO_NOKEY
MPI_ERR_INFO_VALUE	MPI::ERR_INFO_KEY
MPI_ERR_INFO	MPI::ERR_INFO
MPI_ERR_IO	MPI::ERR_IO
MPI_ERR_KEYVAL	MPI::ERR_KEYVAL
MPI_ERR_LOCKTYPE	MPI::ERR_LOCKTYPE
MPI_ERR_NAME	MPI::ERR_NAME
MPI_ERR_NO_MEM	MPI::ERR_NO_MEM
MPI_ERR_NOT_SAME	MPI::ERR_NOT_SAME
MPI_ERR_NO_SPACE	MPI::ERR_NO_SPACE
MPI_ERR_NO_SUCH_FILE	MPI::ERR_NO_SUCH_FILE
MPI_ERR_PORT	MPI::ERR_PORT
MPI_ERR_QUOTA	MPI::ERR_QUOTA
MPI_ERR_READ_ONLY	MPI::ERR_READ_ONLY
MPI_ERR_RMA_CONFLICT	MPI::ERR_RMA_CONFLICT
MPI_ERR_RMA_SYNC	MPI::ERR_RMA_SYNC
MPI_ERR_SERVICE	MPI::ERR_SERVICE
MPI_ERR_SIZE	MPI::ERR_SIZE
MPI_ERR_SPAWN	MPI::ERR_SPAWN
MPI_ERR_UNSUPPORTED_DATAREP	MPI::ERR_UNSUPPORTED_DATAREP
MPI_ERR_UNSUPPORTED_OPERATION	MPI::ERR_UNSUPPORTED_OPERATION
MPI_ERR_WIN	MPI::ERR_WIN
MPI_ERR_LASTCODE	MPI::ERR_LASTCODE

ticket266.

Return Codes for the MPI tool information interface

MPI_T_ERR_CANTINIT
 MPI_T_ERR_NOTINITIALIZED
 MPI_T_ERR_MEMORY
 MPI_T_ERR_INVALIDINDEX
 MPI_T_ERR_INVALIDITEM
 MPI_T_ERR_INVALIDSESSION
 MPI_T_ERR_INVALIDHANDLE
 MPI_T_ERR_OUTOFHANDLES
 MPI_T_ERR_OUTOFSESSIONS
 MPI_T_ERR_CVAR_SETNOTNOW
 MPI_T_ERR_CVAR_SETNEVER
 MPI_T_ERR_PVAR_NOWRITE
 MPI_T_ERR_PVAR_NOSTARTSTOP
 MPI_T_ERR_PVAR_NOATOMIC

Buffer Address Constants

C type: <code>void * const</code>	C++ type:
Fortran type: (predefined memory location)	<code>void * const</code>
MPI_BOTTOM	MPI::BOTTOM
MPI_IN_PLACE	MPI::IN_PLACE

Assorted Constants

C type: <code>const int</code> (or unnamed <code>enum</code>)	C++ type:
Fortran type: <code>INTEGER</code>	<code>const int</code> (or unnamed <code>enum</code>)
MPI_PROC_NULL	MPI::PROC_NULL
MPI_ANY_SOURCE	MPI::ANY_SOURCE
MPI_ANY_TAG	MPI::ANY_TAG
MPI_UNDEFINED	MPI::UNDEFINED
MPI_BSEND_OVERHEAD	MPI::BSEND_OVERHEAD
MPI_KEYVAL_INVALID	MPI::KEYVAL_INVALID
MPI_LOCK_EXCLUSIVE	MPI::LOCK_EXCLUSIVE
MPI_LOCK_SHARED	MPI::LOCK_SHARED
MPI_ROOT	MPI::ROOT

Status size and reserved index values (Fortran only)

Fortran type: <code>INTEGER</code>	
MPI_STATUS_SIZE	Not defined for C++
MPI_SOURCE	Not defined for C++
MPI_TAG	Not defined for C++
MPI_ERROR	Not defined for C++

Variable Address Size (Fortran only)

Fortran type: INTEGER		
MPI_ADDRESS_KIND	Not defined for C++	
MPI_INTEGER_KIND	Not defined for C++	
MPI_OFFSET_KIND	Not defined for C++	

Error-handling specifiers

C type: MPI_Errhandler	C++ type: MPI::Errhandler
Fortran type: INTEGER	
MPI_ERRORS_ARE_FATAL	MPI::ERRORS_ARE_FATAL
MPI_ERRORS_RETURN	MPI::ERRORS_RETURN
	MPI::ERRORS_THROW_EXCEPTIONS

Maximum Sizes for Strings

C type: <code>const int</code> (or unnamed <code>enum</code>)	C++ type:
Fortran type: INTEGER	<code>const int</code> (or unnamed <code>enum</code>)
MPI_MAX_PROCESSOR_NAME	MPI::MAX_PROCESSOR_NAME
MPI_MAX_ERROR_STRING	MPI::MAX_ERROR_STRING
MPI_MAX_DATAREP_STRING	MPI::MAX_DATAREP_STRING
MPI_MAX_INFO_KEY	MPI::MAX_INFO_KEY
MPI_MAX_INFO_VAL	MPI::MAX_INFO_VAL
MPI_MAX_OBJECT_NAME	MPI::MAX_OBJECT_NAME
MPI_MAX_PORT_NAME	MPI::MAX_PORT_NAME

Named Predefined Datatypes		C/C++ types
C type: MPI_Datatype	C++ type: MPI::Datatype	
Fortran type: INTEGER		
MPI_CHAR	MPI::CHAR	char
		(treated as printable character)
MPI_SHORT	MPI::SHORT	signed short int
MPI_INT	MPI::INT	signed int
MPI_LONG	MPI::LONG	signed long
MPI_LONG_LONG_INT	MPI::LONG_LONG_INT	signed long long
MPI_LONG_LONG	MPI::LONG_LONG	long long (synonym)
MPI_SIGNED_CHAR	MPI::SIGNED_CHAR	signed char
		(treated as integral value)
MPI_UNSIGNED_CHAR	MPI::UNSIGNED_CHAR	unsigned char
		(treated as integral value)
MPI_UNSIGNED_SHORT	MPI::UNSIGNED_SHORT	unsigned short
MPI_UNSIGNED	MPI::UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	MPI::UNSIGNED_LONG	unsigned long
MPI_UNSIGNED_LONG_LONG	MPI::UNSIGNED_LONG_LONG	unsigned long long
MPI_FLOAT	MPI::FLOAT	float
MPI_DOUBLE	MPI::DOUBLE	double
MPI_LONG_DOUBLE	MPI::LONG_DOUBLE	long double
MPI_WCHAR	MPI::WCHAR	wchar_t
		(defined in <stddef.h>)
		(treated as printable character)
MPI_C_BOOL	(use C datatype handle)	_Bool
MPI_INT8_T	(use C datatype handle)	int8_t
MPI_INT16_T	(use C datatype handle)	int16_t
MPI_INT32_T	(use C datatype handle)	int32_t
MPI_INT64_T	(use C datatype handle)	int64_t
MPI_UINT8_T	(use C datatype handle)	uint8_t
MPI_UINT16_T	(use C datatype handle)	uint16_t
MPI_UINT32_T	(use C datatype handle)	uint32_t
MPI_UINT64_T	(use C datatype handle)	uint64_t
MPI_AINT	(use C datatype handle)	MPI_Aint
MPI_OFFSET	(use C datatype handle)	MPI_Offset
MPI_C_COMPLEX	(use C datatype handle)	float _Complex
MPI_C_FLOAT_COMPLEX	(use C datatype handle)	float _Complex
MPI_C_DOUBLE_COMPLEX	(use C datatype handle)	double _Complex
MPI_C_LONG_DOUBLE_COMPLEX	(use C datatype handle)	long double _Complex
MPI_BYTE	MPI::BYTE	(any C/C++ type)
MPI_PACKED	MPI::PACKED	(any C/C++ type)

Named Predefined Datatypes		Fortran types
C type: MPI_Datatype	C++ type: MPI::Datatype	
Fortran type: INTEGER		
MPI_INTEGER	MPI::INTEGER	INTEGER
MPI_REAL	MPI::REAL	REAL
MPI_DOUBLE_PRECISION	MPI::DOUBLE_PRECISION	DOUBLE PRECISION
MPI_COMPLEX	MPI::F_COMPLEX	COMPLEX
MPI_LOGICAL	MPI::LOGICAL	LOGICAL
MPI_CHARACTER	MPI::CHARACTER	CHARACTER(1)
MPI_AINT	(use C datatype handle)	INTEGER (KIND=MPI_ADDRESS_KIND)
MPI_OFFSET	(use C datatype handle)	INTEGER (KIND=MPI_OFFSET_KIND)
MPI_BYTE	MPI::BYTE	(any Fortran type)
MPI_PACKED	MPI::PACKED	(any Fortran type)

C++-Only Named Predefined Datatypes	C++ types
C++ type: MPI::Datatype	
MPI::BOOL	bool
MPI::COMPLEX	Complex<float>
MPI::DOUBLE_COMPLEX	Complex<double>
MPI::LONG_DOUBLE_COMPLEX	Complex<long double>

Optional datatypes (Fortran)		Fortran types
C type: MPI_Datatype	C++ type: MPI::Datatype	
Fortran type: INTEGER		
MPI_DOUBLE_COMPLEX	MPI::F_DOUBLE_COMPLEX	DOUBLE COMPLEX
MPI_INTEGER1	MPI::INTEGER1	INTEGER*1
MPI_INTEGER2	MPI::INTEGER2	INTEGER*8
MPI_INTEGER4	MPI::INTEGER4	INTEGER*4
MPI_INTEGER8	MPI::INTEGER8	INTEGER*8
MPI_INTEGER16		INTEGER*16
MPI_REAL2	MPI::REAL2	REAL*2
MPI_REAL4	MPI::REAL4	REAL*4
MPI_REAL8	MPI::REAL8	REAL*8
MPI_REAL16		REAL*16
MPI_COMPLEX4		COMPLEX*4
MPI_COMPLEX8		COMPLEX*8
MPI_COMPLEX16		COMPLEX*16
MPI_COMPLEX32		COMPLEX*32

Datatypes for reduction functions (C and C++)

C type: MPI_Datatype	C++ type: MPI::Datatype
Fortran type: INTEGER	

MPI_FLOAT_INT	MPI::FLOAT_INT
MPI_DOUBLE_INT	MPI::DOUBLE_INT
MPI_LONG_INT	MPI::LONG_INT
MPI_2INT	MPI::TWOINT
MPI_SHORT_INT	MPI::SHORT_INT
MPI_LONG_DOUBLE_INT	MPI::LONG_DOUBLE_INT

Datatypes for reduction functions (Fortran)

C type: MPI_Datatype	C++ type: MPI::Datatype
Fortran type: INTEGER	

MPI_2REAL	MPI::TWOREAL
MPI_2DOUBLE_PRECISION	MPI::TWODOUBLE_PRECISION
MPI_2INTEGER	MPI::TWOINTEGER

Special datatypes for constructing derived datatypes

C type: MPI_Datatype	C++ type: MPI::Datatype
Fortran type: INTEGER	

MPI_UB	MPI::UB
MPI_LB	MPI::LB

Reserved communicators

C type: MPI_Comm	C++ type: MPI::Intracomm
Fortran type: INTEGER	

MPI_COMM_WORLD	MPI::COMM_WORLD
MPI_COMM_SELF	MPI::COMM_SELF

Results of communicator and group comparisons

C type: const int (or unnamed enum)	C++ type: const int
Fortran type: INTEGER	(or unnamed enum)

MPI_IDENT	MPI::IDENT
MPI_CONGRUENT	MPI::CONGRUENT
MPI_SIMILAR	MPI::SIMILAR
MPI_UNEQUAL	MPI::UNEQUAL

Environmental inquiry keys

C type: const int (or unnamed enum)	C++ type: const int
Fortran type: INTEGER	(or unnamed enum)

MPI_TAG_UB	MPI::TAG_UB
MPI_IO	MPI::IO
MPI_HOST	MPI::HOST
MPI_WTIME_IS_GLOBAL	MPI::WTIME_IS_GLOBAL

Collective Operations

C type: MPI_Op	C++ type: const MPI::Op
Fortran type: INTEGER	
MPI_MAX	MPI::MAX
MPI_MIN	MPI::MIN
MPI_SUM	MPI::SUM
MPI_PROD	MPI::PROD
MPI_MAXLOC	MPI::MAXLOC
MPI_MINLOC	MPI::MINLOC
MPI_BAND	MPI::BAND
MPI_BOR	MPI::BOR
MPI_BXOR	MPI::BXOR
MPI_LAND	MPI::LAND
MPI_LOR	MPI::LOR
MPI_LXOR	MPI::LXOR
MPI_REPLACE	MPI::REPLACE

Null Handles

C/Fortran name	C++ name
C type / Fortran type	C++ type
MPI_GROUP_NULL	MPI::GROUP_NULL
MPI_Group / INTEGER	const MPI::Group
MPI_COMM_NULL	MPI::COMM_NULL
MPI_Comm / INTEGER	¹⁾
MPI_DATATYPE_NULL	MPI::DATATYPE_NULL
MPI_Datatype / INTEGER	const MPI::Datatype
MPI_REQUEST_NULL	MPI::REQUEST_NULL
MPI_Request / INTEGER	const MPI::Request
MPI_OP_NULL	MPI::OP_NULL
MPI_Op / INTEGER	const MPI::Op
MPI_ERRHANDLER_NULL	MPI::ERRHANDLER_NULL
MPI_Errhandler / INTEGER	const MPI::Errhandler
MPI_FILE_NULL	MPI::FILE_NULL
MPI_File / INTEGER	
MPI_INFO_NULL	MPI::INFO_NULL
MPI_Info / INTEGER	const MPI::Info
MPI_WIN_NULL	MPI::WIN_NULL
MPI_Win / INTEGER	

¹⁾ C++ type: See Section ?? on page ?? regarding class hierarchy and the specific type of MPI::COMM_NULL

Empty group

C type: MPI_Group	C++ type: const MPI::Group
Fortran type: INTEGER	
MPI_GROUP_EMPTY	MPI::GROUP_EMPTY

Topologies

C type: <code>const int</code> (or unnamed <code>enum</code>)	C++ type: <code>const int</code>
Fortran type: <code>INTEGER</code>	(or unnamed <code>enum</code>)
<code>MPI_GRAPH</code>	<code>MPI::GRAPH</code>
<code>MPI_CART</code>	<code>MPI::CART</code>
<code>MPI_DIST_GRAPH</code>	<code>MPI::DIST_GRAPH</code>

Predefined functions

C/Fortran name	C++ name
C type / Fortran type	C++ type
<code>MPI_COMM_NULL_COPY_FN</code>	<code>MPI_COMM_NULL_COPY_FN</code>
<code>MPI_Comm_copy_attr_function</code>	same as in C ¹)
<code>/ COMM_COPY_ATTR_FN</code>	
<code>MPI_COMM_DUP_FN</code>	<code>MPI_COMM_DUP_FN</code>
<code>MPI_Comm_copy_attr_function</code>	same as in C ¹)
<code>/ COMM_COPY_ATTR_FN</code>	
<code>MPI_COMM_NULL_DELETE_FN</code>	<code>MPI_COMM_NULL_DELETE_FN</code>
<code>MPI_Comm_delete_attr_function</code>	same as in C ¹)
<code>/ COMM_DELETE_ATTR_FN</code>	
<code>MPI_WIN_NULL_COPY_FN</code>	<code>MPI_WIN_NULL_COPY_FN</code>
<code>MPI_Win_copy_attr_function</code>	same as in C ¹)
<code>/ WIN_COPY_ATTR_FN</code>	
<code>MPI_WIN_DUP_FN</code>	<code>MPI_WIN_DUP_FN</code>
<code>MPI_Win_copy_attr_function</code>	same as in C ¹)
<code>/ WIN_COPY_ATTR_FN</code>	
<code>MPI_WIN_NULL_DELETE_FN</code>	<code>MPI_WIN_NULL_DELETE_FN</code>
<code>MPI_Win_delete_attr_function</code>	same as in C ¹)
<code>/ WIN_DELETE_ATTR_FN</code>	
<code>MPI_TYPE_NULL_COPY_FN</code>	<code>MPI_TYPE_NULL_COPY_FN</code>
<code>MPI_Type_copy_attr_function</code>	same as in C ¹)
<code>/ TYPE_COPY_ATTR_FN</code>	
<code>MPI_TYPE_DUP_FN</code>	<code>MPI_TYPE_DUP_FN</code>
<code>MPI_Type_copy_attr_function</code>	same as in C ¹)
<code>/ TYPE_COPY_ATTR_FN</code>	
<code>MPI_TYPE_NULL_DELETE_FN</code>	<code>MPI_TYPE_NULL_DELETE_FN</code>
<code>MPI_Type_delete_attr_function</code>	same as in C ¹)
<code>/ TYPE_DELETE_ATTR_FN</code>	

¹ See the advice to implementors on `MPI_COMM_NULL_COPY_FN`, ... in Section ?? on page ??

Deprecated predefined functions

C/Fortran name	C++ name
C type / Fortran type	C++ type
MPI_NULL_COPY_FN	MPI::NULL_COPY_FN
MPI_Copy_function / COPY_FUNCTION	MPI::Copy_function
MPI_DUP_FN	MPI::DUP_FN
MPI_Copy_function / COPY_FUNCTION	MPI::Copy_function
MPI_NULL_DELETE_FN	MPI::NULL_DELETE_FN
MPI_Delete_function / DELETE_FUNCTION	MPI::Delete_function

Predefined Attribute Keys

C type: <code>const int</code> (or unnamed <code>enum</code>)	C++ type:
Fortran type: <code>INTEGER</code>	<code>const int</code> (or unnamed <code>enum</code>)
MPI_APPNUM	MPI::APPNUM
MPI_LASTUSED_CODE	MPI::LASTUSED_CODE
MPI_UNIVERSE_SIZE	MPI::UNIVERSE_SIZE
MPI_WIN_BASE	MPI::WIN_BASE
MPI_WIN_DISP_UNIT	MPI::WIN_DISP_UNIT
MPI_WIN_SIZE	MPI::WIN_SIZE

Mode Constants

C type: <code>const int</code> (or unnamed <code>enum</code>)	C++ type:
Fortran type: <code>INTEGER</code>	<code>const int</code> (or unnamed <code>enum</code>)
MPI_MODE_APPEND	MPI::MODE_APPEND
MPI_MODE_CREATE	MPI::MODE_CREATE
MPI_MODE_DELETE_ON_CLOSE	MPI::MODE_DELETE_ON_CLOSE
MPI_MODE_EXCL	MPI::MODE_EXCL
MPI_MODE_NOCHECK	MPI::MODE_NOCHECK
MPI_MODE_NOPRECEDE	MPI::MODE_NOPRECEDE
MPI_MODE_NOPUT	MPI::MODE_NOPUT
MPI_MODE_NOSTORE	MPI::MODE_NOSTORE
MPI_MODE_NOSUCCEED	MPI::MODE_NOSUCCEED
MPI_MODE_RDONLY	MPI::MODE_RDONLY
MPI_MODE_RDWR	MPI::MODE_RDWR
MPI_MODE_SEQUENTIAL	MPI::MODE_SEQUENTIAL
MPI_MODE_UNIQUE_OPEN	MPI::MODE_UNIQUE_OPEN
MPI_MODE_WRONLY	MPI::MODE_WRONLY

Datatype Decoding Constants

C type: <code>const int</code> (or unnamed <code>enum</code>)	C++ type:
Fortran type: <code>INTEGER</code>	<code>const int</code> (or unnamed <code>enum</code>)
<code>MPI_COMBINER_CONTIGUOUS</code>	<code>MPI::COMBINER_CONTIGUOUS</code>
<code>MPI_COMBINER_DARRAY</code>	<code>MPI::COMBINER_DARRAY</code>
<code>MPI_COMBINER_DUP</code>	<code>MPI::COMBINER_DUP</code>
<code>MPI_COMBINER_F90_COMPLEX</code>	<code>MPI::COMBINER_F90_COMPLEX</code>
<code>MPI_COMBINER_F90_INTEGER</code>	<code>MPI::COMBINER_F90_INTEGER</code>
<code>MPI_COMBINER_F90_REAL</code>	<code>MPI::COMBINER_F90_REAL</code>
<code>MPI_COMBINER_HINDEXED_INTEGER</code>	<code>MPI::COMBINER_HINDEXED_INTEGER</code>
<code>MPI_COMBINER_HINDEXED</code>	<code>MPI::COMBINER_HINDEXED</code>
<code>MPI_COMBINER_HVECTOR_INTEGER</code>	<code>MPI::COMBINER_HVECTOR_INTEGER</code>
<code>MPI_COMBINER_HVECTOR</code>	<code>MPI::COMBINER_HVECTOR</code>
<code>MPI_COMBINER_INDEXED_BLOCK</code>	<code>MPI::COMBINER_INDEXED_BLOCK</code>
<code>MPI_COMBINER_INDEXED</code>	<code>MPI::COMBINER_INDEXED</code>
<code>MPI_COMBINER_NAMED</code>	<code>MPI::COMBINER_NAMED</code>
<code>MPI_COMBINER_RESIZED</code>	<code>MPI::COMBINER_RESIZED</code>
<code>MPI_COMBINER_STRUCT_INTEGER</code>	<code>MPI::COMBINER_STRUCT_INTEGER</code>
<code>MPI_COMBINER_STRUCT</code>	<code>MPI::COMBINER_STRUCT</code>
<code>MPI_COMBINER_SUBARRAY</code>	<code>MPI::COMBINER_SUBARRAY</code>
<code>MPI_COMBINER_VECTOR</code>	<code>MPI::COMBINER_VECTOR</code>

Threads Constants

C type: <code>const int</code> (or unnamed <code>enum</code>)	C++ type:
Fortran type: <code>INTEGER</code>	<code>const int</code> (or unnamed <code>enum</code>)
<code>MPI_THREAD_FUNNELED</code>	<code>MPI::THREAD_FUNNELED</code>
<code>MPI_THREAD_MULTIPLE</code>	<code>MPI::THREAD_MULTIPLE</code>
<code>MPI_THREAD_SERIALIZED</code>	<code>MPI::THREAD_SERIALIZED</code>
<code>MPI_THREAD_SINGLE</code>	<code>MPI::THREAD_SINGLE</code>

File Operation Constants, Part 1

C type: <code>const MPI_Offset</code> (or unnamed <code>enum</code>)	C++ type:
Fortran type: <code>INTEGER (KIND=MPI_OFFSET_KIND)</code>	<code>const MPI::Offset</code> (or unnamed <code>enum</code>)
<code>MPI_DISPLACEMENT_CURRENT</code>	<code>MPI::DISPLACEMENT_CURRENT</code>

File Operation Constants, Part 2

C type: <code>const int</code> (or unnamed <code>enum</code>)	C++ type:
Fortran type: <code>INTEGER</code>	<code>const int</code> (or unnamed <code>enum</code>)
<code>MPI_DISTRIBUTE_BLOCK</code>	<code>MPI::DISTRIBUTE_BLOCK</code>
<code>MPI_DISTRIBUTE_CYCLIC</code>	<code>MPI::DISTRIBUTE_CYCLIC</code>
<code>MPI_DISTRIBUTE_DFLT_DARG</code>	<code>MPI::DISTRIBUTE_DFLT_DARG</code>
<code>MPI_DISTRIBUTE_NONE</code>	<code>MPI::DISTRIBUTE_NONE</code>
<code>MPI_ORDER_C</code>	<code>MPI::ORDER_C</code>
<code>MPI_ORDER_FORTRAN</code>	<code>MPI::ORDER_FORTRAN</code>
<code>MPI_SEEK_CUR</code>	<code>MPI::SEEK_CUR</code>
<code>MPI_SEEK_END</code>	<code>MPI::SEEK_END</code>
<code>MPI_SEEK_SET</code>	<code>MPI::SEEK_SET</code>

F90 Datatype Matching Constants

C type: <code>const int</code> (or unnamed <code>enum</code>)	C++ type:
Fortran type: <code>INTEGER</code>	<code>const int</code> (or unnamed <code>enum</code>)
<code>MPI_TYPECLASS_COMPLEX</code>	<code>MPI::TYPECLASS_COMPLEX</code>
<code>MPI_TYPECLASS_INTEGER</code>	<code>MPI::TYPECLASS_INTEGER</code>
<code>MPI_TYPECLASS_REAL</code>	<code>MPI::TYPECLASS_REAL</code>

Constants Specifying Empty or Ignored Input

C/Fortran name	C++ name
C type / Fortran type	C++ type
<code>MPI_ARGVS_NULL</code>	<code>MPI::ARGVS_NULL</code>
<code>char***</code> / 2-dim. array of <code>CHARACTER*(*)</code>	<code>const char ***</code>
<code>MPI_ARGV_NULL</code>	<code>MPI::ARGV_NULL</code>
<code>char**</code> / array of <code>CHARACTER*(*)</code>	<code>const char **</code>
<code>MPI_ERRCODES_IGNORE</code>	Not defined for C++
<code>int*</code> / <code>INTEGER</code> array	
<code>MPI_STATUSES_IGNORE</code>	Not defined for C++
<code>MPI_Status*</code> / <code>INTEGER</code> , <code>DIMENSION(MPI_STATUS_SIZE,*)</code>	
<code>MPI_STATUS_IGNORE</code>	Not defined for C++
<code>MPI_Status*</code> / <code>INTEGER</code> , <code>DIMENSION(MPI_STATUS_SIZE)</code>	
<code>MPI_UNWEIGHTED</code>	Not defined for C++

C Constants Specifying Ignored Input (no C++ or Fortran)

C type: <code>MPI_Fint*</code>
<code>MPI_F_STATUSES_IGNORE</code>
<code>MPI_F_STATUS_IGNORE</code>

C and C++ preprocessor Constants and Fortran Parameters

C/C++ type: <code>const int</code> (or unnamed <code>enum</code>)
Fortran type: <code>INTEGER</code>
<code>MPI_SUBVERSION</code>
<code>MPI_VERSION</code>

1

2 ticket266.

Null handles used in the MPI tool information interface

MPI_T_ENUM_NULL

MPI_T_CVAR_HANDLE_NULL

MPI_T_PVAR_HANDLE_NULL

MPI_T_PVAR_SESSION_NULL

Verbosity Levels in the MPI tool information interface

MPI_T_VERBOSITY_USER_BASIC

MPI_T_VERBOSITY_USER_DETAIL

MPI_T_VERBOSITY_USER_ALL

MPI_T_VERBOSITY_TUNER_BASIC

MPI_T_VERBOSITY_TUNER_DETAIL

MPI_T_VERBOSITY_TUNER_ALL

MPI_T_VERBOSITY_MPIDEV_BASIC

MPI_T_VERBOSITY_MPIDEV_DETAIL

MPI_T_VERBOSITY_MPIDEV_ALL

21

ticket0.

Constants to identify associations of variables in the MPI tool information interface

MPI_T_BIND_NO_OBJECT

MPI_T_BIND_MPI_COMM

MPI_T_BIND_MPI_DATATYPE

MPI_T_BIND_MPI_ERRHANDLER

MPI_T_BIND_MPI_FILE

MPI_T_BIND_MPI_GROUP

MPI_T_BIND_MPI_OP

MPI_T_BIND_MPI_REQUEST

MPI_T_BIND_MPI_WIN

MPI_T_BIND_MPI_MESSAGE

MPI_T_BIND_MPI_INFO

Constants describing the scope of a control variable in the MPI tool information interface

MPI_T_SCOPE_READONLY

MPI_T_SCOPE_LOCAL

MPI_T_SCOPE_GROUP

MPI_T_SCOPE_GROUP_EQ

MPI_T_SCOPE_GLOBAL

MPI_T_SCOPE_GLOBAL_EQ

**Additional constants used
by the MPI tool information interface**

MPI_T_PVAR_ALL_HANDLES

**Performance variables classes used by the
MPI tool information interface**

MPI_T_PVAR_CLASS_STATE
MPI_T_PVAR_CLASS_LEVEL
MPI_T_PVAR_CLASS_SIZE
MPI_T_PVAR_CLASS_PERCENTAGE
MPI_T_PVAR_CLASS_HIGHWATERMARK
MPI_T_PVAR_CLASS_LOWWATERMARK
MPI_T_PVAR_CLASS_COUNTER
MPI_T_PVAR_CLASS_AGGREGATE
MPI_T_PVAR_CLASS_TIMER
MPI_T_PVAR_CLASS_GENERIC

3.1.2 Types

The following are defined C type definitions, included in the file `mpi.h`.

```
/* C opaque types */
MPI_Aint
MPI_Fint
MPI_Offset
MPI_Status

/* C handles to assorted structures */
MPI_Comm
MPI_Datatype
MPI_Errhandler
MPI_File
MPI_Group
MPI_Info
MPI_Op
MPI_Request
MPI_Win
```

```
/* Types for the MPI_T interface */
MPI_T_enum
MPI_T_cvar_handle
MPI_T_pvar_handle
MPI_T_pvar_session
```

```
// C++ opaque types (all within the MPI namespace)
MPI::Aint
MPI::Offset
```

```

MPI::Status
// C++ handles to assorted structures (classes,
// all within the MPI namespace)
MPI::Comm
MPI::Intracomm
MPI::Graphcomm
MPI::Distgraphcomm
MPI::Cartcomm
MPI::Intercomm
MPI::Datatype
MPI::Errhandler
MPI::Exception
MPI::File
MPI::Group
MPI::Info
MPI::Op
MPI::Request
MPI::Prequest
MPI::Grequest
MPI::Win

```

3.1.3 Prototype [d]Definitions

The following are defined C typedefs for user-defined functions, also included in the file `mpi.h`.

```

/* prototypes for user-defined functions */
typedef void MPI_User_function(void *invec, void *inoutvec, int *len,
                               MPI_Datatype *datatype);

typedef int MPI_Comm_copy_attr_function(MPI_Comm oldcomm,
                                         int comm_keyval, void *extra_state, void *attribute_val_in,
                                         void *attribute_val_out, int*flag);
typedef int MPI_Comm_delete_attr_function(MPI_Comm comm,
                                         int comm_keyval, void *attribute_val, void *extra_state);

typedef int MPI_Win_copy_attr_function(MPI_Win oldwin, int win_keyval,
                                         void *extra_state, void *attribute_val_in,
                                         void *attribute_val_out, int *flag);
typedef int MPI_Win_delete_attr_function(MPI_Win win, int win_keyval,
                                         void *attribute_val, void *extra_state);

typedef int MPI_Type_copy_attr_function(MPI_Datatype oldtype,
                                         int type_keyval, void *extra_state,
                                         void *attribute_val_in, void *attribute_val_out, int *flag);
typedef int MPI_Type_delete_attr_function(MPI_Datatype type,

```

ticket0.

```

1         int type_keyval, void *attribute_val, void *extra_state);
2
3     typedef void MPI_Comm_errhandler_function(MPI_Comm *, int *, ...);
4     typedef void MPI_Win_errhandler_function(MPI_Win *, int *, ...);
5     typedef void MPI_File_errhandler_function(MPI_File *, int *, ...);
6
7     typedef int MPI_Grequest_query_function(void *extra_state,
8         MPI_Status *status);
9     typedef int MPI_Grequest_free_function(void *extra_state);
10    typedef int MPI_Grequest_cancel_function(void *extra_state, int complete);
11
12    typedef int MPI_Datarep_extent_function(MPI_Datatype datatype,
13        MPI_Aint *file_extent, void *extra_state);
14    typedef int MPI_Datarep_conversion_function(void *userbuf,
15        MPI_Datatype datatype, int count, void *filebuf,
16        MPI_Offset position, void *extra_state);
17

```

For Fortran, here are examples of how each of the user-defined subroutines should be declared.

The user-function argument to MPI_OP_CREATE should be declared like this:

```

21 SUBROUTINE USER_FUNCTION(INVEC, INOUTVEC, LEN, TYPE)
22     <type> INVEC(LEN), INOUTVEC(LEN)
23     INTEGER LEN, TYPE
24

```

The copy and delete function arguments to MPI_COMM_CREATE_KEYVAL should be declared like these:

```

25 SUBROUTINE COMM_COPY_ATTR_FN(OLDCOMM, COMM_KEYVAL, EXTRA_STATE,
26     ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERROR)
27     INTEGER OLDCOMM, COMM_KEYVAL, IERROR
28     INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,
29     ATTRIBUTE_VAL_OUT
30     LOGICAL FLAG
31
32 SUBROUTINE COMM_DELETE_ATTR_FN(COMM, COMM_KEYVAL, ATTRIBUTE_VAL,
33     EXTRA_STATE, IERROR)
34     INTEGER COMM, COMM_KEYVAL, IERROR
35     INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE
36

```

The copy and delete function arguments to MPI_WIN_CREATE_KEYVAL should be declared like these:

```

37 SUBROUTINE WIN_COPY_ATTR_FN(OLDWIN, WIN_KEYVAL, EXTRA_STATE,
38     ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERROR)
39     INTEGER OLDWIN, WIN_KEYVAL, IERROR
40     INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,
41     ATTRIBUTE_VAL_OUT
42     LOGICAL FLAG
43

```

```

SUBROUTINE WIN_DELETE_ATTR_FN(WIN, WIN_KEYVAL, ATTRIBUTE_VAL,
                               EXTRA_STATE, IERROR)
    INTEGER WIN, WIN_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE

```

The copy and delete function arguments to MPI_TYPE_CREATE_KEYVAL should be declared like these:

```

SUBROUTINE TYPE_COPY_ATTR_FN(OLDTYPE, TYPE_KEYVAL, EXTRA_STATE,
                              ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERROR)
    INTEGER OLDTYPE, TYPE_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE,
        ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT
    LOGICAL FLAG

```

```

SUBROUTINE TYPE_DELETE_ATTR_FN(TYPE, TYPE_KEYVAL, ATTRIBUTE_VAL,
                                EXTRA_STATE, IERROR)
    INTEGER TYPE, TYPE_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE

```

The handler-function argument to MPI_COMM_CREATE_ERRHANDLER should be declared like this:

```

SUBROUTINE COMM_ERRHANDLER_FUNCTION(COMM, ERROR_CODE)
    INTEGER COMM, ERROR_CODE

```

The handler-function argument to MPI_WIN_CREATE_ERRHANDLER should be declared like this:

```

SUBROUTINE WIN_ERRHANDLER_FUNCTION(WIN, ERROR_CODE)
    INTEGER WIN, ERROR_CODE

```

The handler-function argument to MPI_FILE_CREATE_ERRHANDLER should be declared like this:

```

SUBROUTINE FILE_ERRHANDLER_FUNCTION(FILE, ERROR_CODE)
    INTEGER FILE, ERROR_CODE

```

The query, free, and cancel function arguments to MPI_GREQUEST_START should be declared like these:

```

SUBROUTINE GREQUEST_QUERY_FUNCTION(EXTRA_STATE, STATUS, IERROR)
    INTEGER STATUS(MPI_STATUS_SIZE), IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE

```

```

SUBROUTINE GREQUEST_FREE_FUNCTION(EXTRA_STATE, IERROR)
    INTEGER IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE

```

```

1
2 SUBROUTINE GREQUEST_CANCEL_FUNCTION(EXTRA_STATE, COMPLETE, IERROR)
3     INTEGER IERROR
4     INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
5     LOGICAL COMPLETE
6

```

The extend and conversion function arguments to MPI_REGISTER_DATAREP should be declared like these:

```

9 SUBROUTINE DATAREP_EXTENT_FUNCTION(DATATYPE, EXTENT, EXTRA_STATE, IERROR)
10     INTEGER DATATYPE, IERROR
11     INTEGER(KIND=MPI_ADDRESS_KIND) EXTENT, EXTRA_STATE
12
13 SUBROUTINE DATAREP_CONVERSION_FUNCTION(USERBUF, DATATYPE, COUNT, FILEBUF,
14     POSITION, EXTRA_STATE, IERROR)
15     <TYPE> USERBUF(*), FILEBUF(*)
16     INTEGER COUNT, DATATYPE, IERROR
17     INTEGER(KIND=MPI_OFFSET_KIND) POSITION
18     INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
19

```

The following are defined C++ typedefs, also included in the file mpi.h.

```

21 namespace MPI {
22     typedef void User_function(const void* invec, void *inoutvec,
23         int len, const Datatype& datatype);
24
25
26     typedef int Comm::Copy_attr_function(const Comm& oldcomm,
27         int comm_keyval, void* extra_state, void* attribute_val_in,
28         void* attribute_val_out, bool& flag);
29     typedef int Comm::Delete_attr_function(Comm& comm, int
30         comm_keyval, void* attribute_val, void* extra_state);
31
32     typedef int Win::Copy_attr_function(const Win& oldwin,
33         int win_keyval, void* extra_state, void* attribute_val_in,
34         void* attribute_val_out, bool& flag);
35     typedef int Win::Delete_attr_function(Win& win, int
36         win_keyval, void* attribute_val, void* extra_state);
37
38     typedef int Datatype::Copy_attr_function(const Datatype& oldtype,
39         int type_keyval, void* extra_state,
40         const void* attribute_val_in, void* attribute_val_out,
41         bool& flag);
42     typedef int Datatype::Delete_attr_function(Datatype& type,
43         int type_keyval, void* attribute_val, void* extra_state);
44
45     typedef void Comm::Errhandler_function(Comm &, int *, ...);
46     typedef void Win::Errhandler_function(Win &, int *, ...);
47     typedef void File::Errhandler_function(File &, int *, ...);
48

```



```

typedef int Grequest::Query_function(void* extra_state, Status& status);
typedef int Grequest::Free_function(void* extra_state);
typedef int Grequest::Cancel_function(void* extra_state, bool complete);

typedef void Datarep_extent_function(const Datatype& datatype,
                                     Aint& file_extent, void* extra_state);
typedef void Datarep_conversion_function(void* userbuf,
                                         Datatype& datatype, int count, void* filebuf,
                                         Offset position, void* extra_state);
}

```

3.1.4 Deprecated [p]Prototype [d]Definitions

The following are defined C typedefs for deprecated user-defined functions, also included in the file `mpi.h`.

```

/* prototypes for user-defined functions */
typedef int MPI_Copy_function(MPI_Comm oldcomm, int keyval,
                             void *extra_state, void *attribute_val_in,
                             void *attribute_val_out, int *flag);
typedef int MPI_Delete_function(MPI_Comm comm, int keyval,
                               void *attribute_val, void *extra_state);
typedef void MPI_Handler_function(MPI_Comm *, int *, ...);

```

The following are deprecated Fortran user-defined callback subroutine prototypes. The deprecated copy and delete function arguments to `MPI_KEYVAL_CREATE` should be declared like these:

```

SUBROUTINE COPY_FUNCTION(OLDCOMM, KEYVAL, EXTRA_STATE,
                        ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERR)
    INTEGER OLDCOMM, KEYVAL, EXTRA_STATE, ATTRIBUTE_VAL_IN,
    ATTRIBUTE_VAL_OUT, IERR
    LOGICAL FLAG

SUBROUTINE DELETE_FUNCTION(COMM, KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE, IERR)
    INTEGER COMM, KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE, IERR

```

The deprecated handler-function for error handlers should be declared like this:

```

SUBROUTINE HANDLER_FUNCTION(COMM, ERROR_CODE)
    INTEGER COMM, ERROR_CODE

```

3.1.5 Info Keys

```

access_style
appnum
arch
cb_block_size
cb_buffer_size

```

1 cb_nodes
2 chunked_item
3 chunked_size
4 chunked
5 collective_buffering
6 file_perm
7 filename
8 file
9 host
10 io_node_list
11 ip_address
12 ip_port
13 nb_proc
14 no_locks
15 num_io_nodes
16 path
17 soft
18 striping_factor
19 striping_unit
20 wdir

21

22

23 3.1.6 Info Values

24

25 false
26 random
27 read_mostly
28 read_once
29 reverse_sequential
30 sequential
31 true
32 write_mostly
33 write_once

34

35

36

37

38

39

40

41

42

43

44

45

46

47

48

Bibliography

- [1] Martin Schulz and Bronis R. de Supinski. P^N MPI Tools: A Whole Lot Greater Than the Sum of Their Parts. In *ACM/IEEE Supercomputing Conference (SC)*, pages 1–10. ACM, 2007. [2.2.8](#)

Examples Index

This index lists code examples throughout the text. Some examples are referred to by content; others are listed by the major MPI function that they are demonstrating. MPI functions listed in all capital letter are Fortran examples; MPI functions listed in mixed case are C/C++ examples.

Basic tool using performance variables in
the MPI tool information interface,
[58](#)

MPI_ALLOC_MEM, [5](#)
MPI_Alloc_mem, [6](#)
MPI_Barrier, [22](#), [23](#)
MPI_Buffer_attach, [23](#)
MPI_Cancel, [23](#)
MPI_Finalize, [22-24](#)
MPI_FREE_MEM, [5](#)
MPI_Iprobe, [23](#)
MPI_Request_free, [22](#)
MPI_Test_cancelled, [23](#)
mpiexec, [28](#)

Profiling interface, [33](#)

Reading the value of a control variable in
the MPI tool information interface,
[47](#)

Using MPI_T_CVAR_GET_INFO to list all
names of control variables., [44](#)

MPI Constant and Predefined Handle Index

This index lists predefined MPI constants and handles.

MPI::ANY_SOURCE, 67	MPI::COMM_NULL, 72
MPI::ANY_TAG, 67	MPI::COMM_SELF, 71
MPI::APPNUM, 74	MPI::COMM_WORLD, 71
MPI::ARGV_NULL, 76	MPI::COMPLEX, 70
MPI::ARGVS_NULL, 76	MPI::CONGRUENT, 71
MPI::BAND, 72	MPI::DATATYPE_NULL, 72
MPI::BOOL, 70	MPI::DISPLACEMENT_CURRENT, 75
MPI::BOR, 72	MPI::DIST_GRAPH, 73
MPI::BOTTOM, 67	MPI::DISTRIBUTE_BLOCK, 76
MPI::BSEND_OVERHEAD, 67	MPI::DISTRIBUTE_CYCLIC, 76
MPI::BXOR, 72	MPI::DISTRIBUTE_DFLT_DARG, 76
MPI::BYTE, 69 , 70	MPI::DISTRIBUTE_NONE, 76
MPI::CART, 73	MPI::DOUBLE, 69
MPI::CHAR, 69	MPI::DOUBLE_COMPLEX, 70
MPI::CHARACTER, 70	MPI::DOUBLE_INT, 71
MPI::COMBINER_CONTIGUOUS, 75	MPI::DOUBLE_PRECISION, 70
MPI::COMBINER_DARRAY, 75	MPI::DUP_FN, 74
MPI::COMBINER_DUP, 75	MPI::ERR_ACCESS, 66
MPI::COMBINER_F90_COMPLEX, 75	MPI::ERR_AMODE, 66
MPI::COMBINER_F90_INTEGER, 75	MPI::ERR_ARG, 65
MPI::COMBINER_F90_REAL, 75	MPI::ERR_ASSERT, 66
MPI::COMBINER_HINDEXED, 75	MPI::ERR_BAD_FILE, 66
MPI::COMBINER_HINDEXED_INTEGER, 75	MPI::ERR_BASE, 66
MPI::COMBINER_HVECTOR, 75	MPI::ERR_BUFFER, 65
MPI::COMBINER_HVECTOR_INTEGER, 75	MPI::ERR_COMM, 65
MPI::COMBINER_INDEXED, 75	MPI::ERR_CONVERSION, 66
MPI::COMBINER_INDEXED_BLOCK, 75	MPI::ERR_COUNT, 65
MPI::COMBINER_NAMED, 75	MPI::ERR_DIMS, 65
MPI::COMBINER_RESIZED, 75	MPI::ERR_DISP, 66
MPI::COMBINER_STRUCT, 75	MPI::ERR_DUP_DATAREP, 66
MPI::COMBINER_STRUCT_INTEGER, 75	MPI::ERR_FILE, 66
MPI::COMBINER_SUBARRAY, 75	MPI::ERR_FILE_EXISTS, 66
MPI::COMBINER_VECTOR, 75	MPI::ERR_FILE_IN_USE, 66
	MPI::ERR_GROUP, 65
	MPI::ERR_IN_STATUS, 66

1	MPI::ERR_INFO, 66	MPI::GROUP_EMPTY, 72
2	MPI::ERR_INFO_KEY, 66	MPI::GROUP_NULL, 72
3	MPI::ERR_INFO_NOKEY, 66	MPI::HOST, 71
4	MPI::ERR_INFO_VALUE, 66	MPI::IDENT, 71
5	MPI::ERR_INTERN, 65	MPI::IN_PLACE, 67
6	MPI::ERR_IO, 66	MPI::INFO_NULL, 72
7	MPI::ERR_KEYVAL, 66	MPI::INT, 69
8	MPI::ERR_LASTCODE, 66	MPI::INTEGER, 70
9	MPI::ERR_LOCKTYPE, 66	MPI::INTEGER1, 70
10	MPI::ERR_NAME, 66	MPI::INTEGER2, 70
11	MPI::ERR_NO_MEM, 66	MPI::INTEGER4, 70
12	MPI::ERR_NO_SPACE, 66	MPI::INTEGER8, 70
13	MPI::ERR_NO_SUCH_FILE, 66	MPI::IO, 71
14	MPI::ERR_NOT_SAME, 66	MPI::KEYVAL_INVALID, 67
15	MPI::ERR_OP, 65	MPI::LAND, 72
16	MPI::ERR_OTHER, 65	MPI::LASTUSED, 74
17	MPI::ERR_PENDING, 65	MPI::LB, 71
18	MPI::ERR_PORT, 66	MPI::LOCK_EXCLUSIVE, 67
19	MPI::ERR_QUOTA, 66	MPI::LOCK_SHARED, 67
20	MPI::ERR_RANK, 65	MPI::LOGICAL, 70
21	MPI::ERR_READ_ONLY, 66	MPI::LONG, 69
22	MPI::ERR_REQUEST, 65	MPI::LONG_DOUBLE, 69
23	MPI::ERR_RMA_CONFLICT, 66	MPI::LONG_DOUBLE_COMPLEX, 70
24	MPI::ERR_RMA_SYNC, 66	MPI::LONG_DOUBLE_INT, 71
25	MPI::ERR_ROOT, 65	MPI::LONG_INT, 71
26	MPI::ERR_SERVICE, 66	MPI::LONG_LONG, 69
27	MPI::ERR_SIZE, 66	MPI::LONG_LONG_INT, 69
28	MPI::ERR_SPAWN, 66	MPI::LOR, 72
29	MPI::ERR_TAG, 65	MPI::LXOR, 72
30	MPI::ERR_TOPOLOGY, 65	MPI::MAX, 72
31	MPI::ERR_TRUNCATE, 65	MPI::MAX_DATAREP_STRING, 68
32	MPI::ERR_TYPE, 65	MPI::MAX_ERROR_STRING, 68
33	MPI::ERR_UNKNOWN, 65	MPI::MAX_INFO_KEY, 68
34	MPI::ERR_UNSUPPORTED_DATAREP, 66	MPI::MAX_INFO_VAL, 68
35	MPI::ERR_UNSUPPORTED_OPERATION,	MPI::MAX_OBJECT_NAME, 68
36	66	MPI::MAX_PORT_NAME, 68
37	MPI::ERR_WIN, 66	MPI::MAX_PROCESSOR_NAME, 68
38	MPI::ERRHANDLER_NULL, 72	MPI::MAXLOC, 72
39	MPI::ERRORS_ARE_FATAL, 68	MPI::MIN, 72
40	MPI::ERRORS_RETURN, 68	MPI::MINLOC, 72
41	MPI::ERRORS_THROW_EXCEPTIONS, 7,	MPI::MODE_APPEND, 74
42	68	MPI::MODE_CREATE, 74
43	MPI::F_COMPLEX, 70	MPI::MODE_DELETE_ON_CLOSE, 74
44	MPI::F_DOUBLE_COMPLEX, 70	MPI::MODE_EXCL, 74
45	MPI::FILE_NULL, 72	MPI::MODE_NOCHECK, 74
46	MPI::FLOAT, 69	MPI::MODE_NOPRECEDE, 74
47	MPI::FLOAT_INT, 71	MPI::MODE_NOPUT, 74
48	MPI::GRAPH, 73	MPI::MODE_NOSTORE, 74

MPI::MODE_NOSUCCEED, 74	MPI::UNSIGNED_LONG, 69	1
MPI::MODE_RDONLY, 74	MPI::UNSIGNED_LONG_LONG, 69	2
MPI::MODE_RDWR, 74	MPI::UNSIGNED_SHORT, 69	3
MPI::MODE_SEQUENTIAL, 74	MPI::WCHAR, 69	4
MPI::MODE_UNIQUE_OPEN, 74	MPI::WIN_BASE, 74	5
MPI::MODE_WRONLY, 74	MPI::WIN_DISP_UNIT, 74	6
MPI::NULL_COPY_FN, 74	MPI::WIN_NULL, 72	7
MPI::NULL_DELETE_FN, 74	MPI::WIN_SIZE, 74	8
MPI::OP_NULL, 72	MPI::WTIME_IS_GLOBAL, 71	9
MPI::ORDER_C, 76	MPI_2DOUBLE_PRECISION, 71	10
MPI::ORDER_FORTRAN, 76	MPI_2INT, 71	11
MPI::PACKED, 69, 70	MPI_2INTEGER, 71	12
MPI::PROC_NULL, 67	MPI_2REAL, 71	13
MPI::PROD, 72	MPI_ADDRESS_KIND, 68	14
MPI::REAL, 70	MPI_AINT, 69, 70	15
MPI::REAL2, 70	MPI_ANY_SOURCE, 3, 67	16
MPI::REAL4, 70	MPI_ANY_TAG, 67	17
MPI::REAL8, 70	MPI_APPNUM, 74	18
MPI::REPLACE, 72	MPI_ARGV_NULL, 76	19
MPI::REQUEST_NULL, 72	MPI_ARGVS_NULL, 76	20
MPI::ROOT, 67	MPI_BAND, 72	21
MPI::SEEK_CUR, 76	MPI_BOR, 72	22
MPI::SEEK_END, 76	MPI_BOTTOM, 67	23
MPI::SEEK_SET, 76	MPI_BSEND_OVERHEAD, 4, 67	24
MPI::SHORT, 69	MPI_BXOR, 72	25
MPI::SHORT_INT, 71	MPI_BYTE, 69, 70	26
MPI::SIGNED_CHAR, 69	MPI_C_BOOL, 69	27
MPI::SIMILAR, 71	MPI_C_COMPLEX, 69	28
MPI::SUCCESS, 65	MPI_C_DOUBLE_COMPLEX, 69	29
MPI::SUM, 72	MPI_C_FLOAT_COMPLEX, 69	30
MPI::TAG_UB, 71	MPI_C_LONG_DOUBLE_COMPLEX, 69	31
MPI::THREAD_FUNNELED, 75	MPI_CART, 73	32
MPI::THREAD_MULTIPLE, 75	MPI_CHAR, 40, 69	33
MPI::THREAD_SERIALIZED, 75	MPI_CHARACTER, 70	34
MPI::THREAD_SINGLE, 75	MPI_COMBINER_CONTIGUOUS, 75	35
MPI::TWODOUBLE_PRECISION, 71	MPI_COMBINER_DARRAY, 75	36
MPI::TWOINT, 71	MPI_COMBINER_DUP, 75	37
MPI::TWOINTEGER, 71	MPI_COMBINER_F90_COMPLEX, 75	38
MPI::TWOREAL, 71	MPI_COMBINER_F90_INTEGER, 75	39
MPI::TYPECLASS_COMPLEX, 76	MPI_COMBINER_F90_REAL, 75	40
MPI::TYPECLASS_INTEGER, 76	MPI_COMBINER_HINDEXED, 75	41
MPI::TYPECLASS_REAL, 76	MPI_COMBINER_HINDEXED_INTEGER,	42
MPI::UB, 71	75	43
MPI::UNDEFINED, 67	MPI_COMBINER_HVECTOR, 75	44
MPI::UNEQUAL, 71	MPI_COMBINER_HVECTOR_INTEGER,	45
MPI::UNIVERSE_SIZE, 74	75	46
MPI::UNSIGNED, 69	MPI_COMBINER_INDEXED, 75	47
MPI::UNSIGNED_CHAR, 69	MPI_COMBINER_INDEXED_BLOCK, 75	48

1	MPI_COMBINER_NAMED, 75	MPI_ERR_INFO_KEY, 14, 66
2	MPI_COMBINER_RESIZED, 75	MPI_ERR_INFO_NOKEY, 14, 66
3	MPI_COMBINER_STRUCT, 75	MPI_ERR_INFO_VALUE, 14, 66
4	MPI_COMBINER_STRUCT_INTEGER, 75	MPI_ERR_INTERN, 14, 65
5	MPI_COMBINER_SUBARRAY, 75	MPI_ERR_IO, 15, 66
6	MPI_COMBINER_VECTOR, 75	MPI_ERR_KEYVAL, 14, 66
7	MPI_COMM_NULL, 72	MPI_ERR_LASTCODE, 13, 15, 17, 18, 66
8	MPI_COMM_SELF, 26, 71	MPI_ERR_LOCKTYPE, 14, 66
9	MPI_COMM_WORLD, 2, 3, 6, 9, 17, 24,	MPI_ERR_NAME, 14, 66
10	25, 27, 46, 54, 71	MPI_ERR_NO_MEM, 5, 14, 66
11	MPI_COMPLEX, 70	MPI_ERR_NO_SPACE, 15, 66
12	MPI_COMPLEX16, 70	MPI_ERR_NO_SUCH_FILE, 15, 66
13	MPI_COMPLEX32, 70	MPI_ERR_NOATOMIC, 57
14	MPI_COMPLEX4, 70	MPI_ERR_NOT_SAME, 15, 66
15	MPI_COMPLEX8, 70	MPI_ERR_OP, 14, 65
16	MPI_CONGRUENT, 71	MPI_ERR_OTHER, 13, 14, 65
17	MPI_COUNT, 40	MPI_ERR_PENDING, 14, 65
18	MPI_DATATYPE_NULL, 72	MPI_ERR_PORT, 14, 66
19	MPI_DISPLACEMENT_CURRENT, 75	MPI_ERR_QUOTA, 15, 66
20	MPI_DIST_GRAPH, 73	MPI_ERR_RANK, 14, 65
21	MPI_DISTRIBUTE_BLOCK, 76	MPI_ERR_READ_ONLY, 15, 66
22	MPI_DISTRIBUTE_CYCLIC, 76	MPI_ERR_REQUEST, 14, 65
23	MPI_DISTRIBUTE_DFLT_DARG, 76	MPI_ERR_RMA_CONFLICT, 14, 66
24	MPI_DISTRIBUTE_NONE, 76	MPI_ERR_RMA_SYNC, 14, 66
25	MPI_DOUBLE, 40, 50, 69	MPI_ERR_ROOT, 14, 65
26	MPI_DOUBLE_COMPLEX, 70	MPI_ERR_SERVICE, 14, 66
27	MPI_DOUBLE_INT, 71	MPI_ERR_SIZE, 14, 66
28	MPI_DOUBLE_PRECISION, 70	MPI_ERR_SPAWN, 14, 66
29	MPI_DUP_FN, 74	MPI_ERR_TAG, 14, 65
30	MPI_ERR_ACCESS, 15, 66	MPI_ERR_TOPOLOGY, 14, 65
31	MPI_ERR_AMODE, 15, 66	MPI_ERR_TRUNCATE, 14, 65
32	MPI_ERR_ARG, 14, 65	MPI_ERR_TYPE, 14, 65
33	MPI_ERR_ASSERT, 14, 66	MPI_ERR_UNKNOWN, 13, 14, 65
34	MPI_ERR_BAD_FILE, 15, 66	MPI_ERR_UNSUPPORTED_DATAREP, 15,
35	MPI_ERR_BASE, 5, 14, 66	66
36	MPI_ERR_BUFFER, 14, 65	MPI_ERR_UNSUPPORTED_OPERATION,
37	MPI_ERR_COMM, 14, 65	15, 66
38	MPI_ERR_CONVERSION, 15, 66	MPI_ERR_WIN, 14, 66
39	MPI_ERR_COUNT, 14, 65	MPI_ERRCODES_IGNORE, 76
40	MPI_ERR_DIMS, 14, 65	MPI_ERRHANDLER_NULL, 12, 72
41	MPI_ERR_DISP, 14, 66	MPI_ERROR, 67
42	MPI_ERR_DUP_DATAREP, 15, 66	MPI_ERROR_STRING, 13
43	MPI_ERR_FILE, 15, 66	MPI_ERRORS_ARE_FATAL, 6, 7, 18, 19,
44	MPI_ERR_FILE_EXISTS, 15, 66	68
45	MPI_ERR_FILE_IN_USE, 15, 66	MPI_ERRORS_RETURN, 6, 7, 19, 68
46	MPI_ERR_GROUP, 14, 65	MPI_F_STATUS_IGNORE, 76
47	MPI_ERR_IN_STATUS, 8, 14, 66	MPI_F_STATUSES_IGNORE, 76
48	MPI_ERR_INFO, 14, 66	MPI_FILE_NULL, 72

MPI_FLOAT, 69	MPI_MODE_APPEND, 74	1
MPI_FLOAT_INT, 71	MPI_MODE_CREATE, 74	2
MPI_GRAPH, 73	MPI_MODE_DELETE_ON_CLOSE, 74	3
MPI_GROUP_EMPTY, 72	MPI_MODE_EXCL, 74	4
MPI_GROUP_NULL, 72	MPI_MODE_NOCHECK, 74	5
MPI_HOST, 2, 71	MPI_MODE_NOPRECEDE, 74	6
MPI_IDENT, 71	MPI_MODE_NOPUT, 74	7
MPI_IN_PLACE, 67	MPI_MODE_NOSTORE, 74	8
MPI_INFO_NULL, 72	MPI_MODE_NOSUCCEED, 74	9
MPI_INT, 40, 41, 43, 49, 53, 69	MPI_MODE_RDONLY, 74	10
MPI_INT16_T, 69	MPI_MODE_RDWR, 74	11
MPI_INT32_T, 69	MPI_MODE_SEQUENTIAL, 74	12
MPI_INT64_T, 69	MPI_MODE_UNIQUE_OPEN, 74	13
MPI_INT8_T, 69	MPI_MODE_WRONLY, 74	14
MPI_INTEGER, 70	MPI_NULL_COPY_FN, 74	15
MPI_INTEGER1, 70	MPI_NULL_DELETE_FN, 74	16
MPI_INTEGER16, 70	MPI_OFFSET, 69, 70	17
MPI_INTEGER2, 70	MPI_OFFSET_KIND, 68	18
MPI_INTEGER4, 70	MPI_OP_NULL, 72	19
MPI_INTEGER8, 70	MPI_ORDER_C, 76	20
MPI_INTEGER_KIND, 68	MPI_ORDER_FORTRAN, 76	21
MPI_IO, 2, 3, 71	MPI_PACKED, 69, 70	22
MPI_KEYVAL_INVALID, 67	MPI_PROC_NULL, 2, 3, 67	23
MPI_LAND, 72	MPI_PROD, 72	24
MPI_LASTUSED_CODE, 17, 74	MPI_REAL, 70	25
MPI_LB, 71	MPI_REAL16, 70	26
MPI_LOCK_EXCLUSIVE, 67	MPI_REAL2, 70	27
MPI_LOCK_SHARED, 67	MPI_REAL4, 70	28
MPI_LOGICAL, 70	MPI_REAL8, 70	29
MPI_LONG, 69	MPI_REPLACE, 72	30
MPI_LONG_DOUBLE, 69	MPI_REQUEST_NULL, 72	31
MPI_LONG_DOUBLE_INT, 71	MPI_ROOT, 67	32
MPI_LONG_INT, 71	MPI_SEEK_CUR, 76	33
MPI_LONG_LONG, 69	MPI_SEEK_END, 76	34
MPI_LONG_LONG_INT, 69	MPI_SEEK_SET, 76	35
MPI_LOR, 72	MPI_SHORT, 69	36
MPI_LXOR, 72	MPI_SHORT_INT, 71	37
MPI_MAX, 72	MPI_SIGNED_CHAR, 69	38
MPI_MAX_DATAREP_STRING, 68	MPI_SIMILAR, 71	39
MPI_MAX_ERROR_STRING, 13, 18, 68	MPI_SOURCE, 67	40
MPI_MAX_INFO_KEY, 14, 68	MPI_STATUS_IGNORE, 76	41
MPI_MAX_INFO_VAL, 14, 68	MPI_STATUS_SIZE, 67	42
MPI_MAX_OBJECT_NAME, 68	MPI_STATUSES_IGNORE, 76	43
MPI_MAX_PORT_NAME, 68	MPI_SUBVERSION, 2, 76	44
MPI_MAX_PROCESSOR_NAME, 4, 68	MPI_SUCCESS, 13, 14, 18, 19, 36, 55–57, 64, 65	45
MPI_MAXLOC, 72	MPI_SUM, 72	46
MPI_MIN, 72	MPI_T_BIND_MPI_COMM, 38, 77	47
MPI_MINLOC, 72		48

1 MPI_T_BIND_MPI_DATATYPE, 38, 77	MPI_T_PVAR_SESSION_NULL, 54, 77
2 MPI_T_BIND_MPI_ERRHANDLER, 38, 77	MPI_T_SCOPE_ALL, 44
3 MPI_T_BIND_MPI_FILE, 38, 77	MPI_T_SCOPE_ALL_EQ, 44
4 MPI_T_BIND_MPI_GROUP, 38, 77	MPI_T_SCOPE_GLOBAL, 77
5 MPI_T_BIND_MPI_INFO, 38, 77	MPI_T_SCOPE_GLOBAL_EQ, 47, 77
6 MPI_T_BIND_MPI_MESSAGE, 38, 77	MPI_T_SCOPE_GROUP, 44, 77
7 MPI_T_BIND_MPI_OP, 38, 77	MPI_T_SCOPE_GROUP_EQ, 44, 47, 77
8 MPI_T_BIND_MPI_REQUEST, 38, 77	MPI_T_SCOPE_LOCAL, 44, 77
9 MPI_T_BIND_MPI_WIN, 38, 77	MPI_T_SCOPE_READONLY, 44, 77
10 MPI_T_BIND_NO_OBJECT, 38, 44, 46, 53,	MPI_T_VERBOSITY_MPIDEV_ALL, 37,
11 54, 77	77
12 MPI_T_CVAR_HANDLE_NULL, 46, 77	MPI_T_VERBOSITY_MPIDEV_BASIC, 37,
13 MPI_T_ENUM_NULL, 43, 53, 77	77
14 MPI_T_ERR_CANTINIT, 64, 67	MPI_T_VERBOSITY_MPIDEV_DETAIL,
15 MPI_T_ERR_CVAR_SETNEVER, 47, 64,	37, 77
16 67	MPI_T_VERBOSITY_TUNER_ALL, 37, 77
17 MPI_T_ERR_CVAR_SETNOTNOW, 47, 64,	MPI_T_VERBOSITY_TUNER_BASIC, 37,
18 67	77
19 MPI_T_ERR_INVALIDHANDLE, 64, 67	MPI_T_VERBOSITY_TUNER_DETAIL, 37,
20 MPI_T_ERR_INVALIDINDEX, 64, 67	77
21 MPI_T_ERR_INVALIDITEM, 64, 67	MPI_T_VERBOSITY_USER_ALL, 37, 77
22 MPI_T_ERR_INVALIDSESSION, 64, 67	MPI_T_VERBOSITY_USER_BASIC, 37, 77
23 MPI_T_ERR_MEMORY, 64, 67	MPI_T_VERBOSITY_USER_DETAIL, 37,
24 MPI_T_ERR_NOTINITIALIZED, 64, 67	77
25 MPI_T_ERR_OUTOFHANDLES, 64, 67	MPI_TAG, 67
26 MPI_T_ERR_OUTOFSESSIONS, 64, 67	MPI_TAG_UB, 2, 71
27 MPI_T_ERR_PVAR_NOATOMIC, 67	MPI_THREAD_FUNNELED, 75
28 MPI_T_ERR_PVAR_NOSTARTSTOP, 55,	MPI_THREAD_MULTIPLE, 75
29 56, 64, 67	MPI_THREAD_SERIALIZED, 75
30 MPI_T_ERR_PVAR_NOWRITE, 56, 57, 64,	MPI_THREAD_SINGLE, 75
31 67	MPI_TYPECLASS_COMPLEX, 76
32 MPI_T_NOATOMIC, 64	MPI_TYPECLASS_INTEGER, 76
33 MPI_T_PVAR_ALL_HANDLES, 55–58, 78	MPI_TYPECLASS_REAL, 76
34 MPI_T_PVAR_CLASS_AGGREGATE, 50,	MPI_UB, 71
35 78	MPI_UINT16_T, 69
36 MPI_T_PVAR_CLASS_COUNTER, 50, 78	MPI_UINT32_T, 69
37 MPI_T_PVAR_CLASS_GENERIC, 50, 78	MPI_UINT64_T, 69
38 MPI_T_PVAR_CLASS_HIGHWATERMARK,	MPI_UINT8_T, 69
39 49, 78	MPI_UNDEFINED, 67
40 MPI_T_PVAR_CLASS_LEVEL, 49, 78	MPI_UNEQUAL, 71
41 MPI_T_PVAR_CLASS_LOWWATERMARK,	MPI_UNIVERSE_SIZE, 74
42 50, 78	MPI_UNSIGNED, 40, 69
43 MPI_T_PVAR_CLASS_PERCENTAGE, 49,	MPI_UNSIGNED_CHAR, 69
44 78	MPI_UNSIGNED_LONG, 40, 69
45 MPI_T_PVAR_CLASS_SIZE, 49, 78	MPI_UNSIGNED_LONG_LONG, 40, 69
46 MPI_T_PVAR_CLASS_STATE, 49, 78	MPI_UNSIGNED_SHORT, 69
47 MPI_T_PVAR_CLASS_TIMER, 50, 78	MPI_UNWEIGHTED, 76
48 MPI_T_PVAR_HANDLE_NULL, 55, 77	MPI_VERSION, 2, 76

MPI_WCHAR, 69	1
MPI_WIN_BASE, 74	2
MPI_WIN_DISP_UNIT, 74	3
MPI_WIN_NULL, 72	4
MPI_WIN_SIZE, 74	5
MPI_WTIME_IS_GLOBAL, 2 , 3 , 20 , 71	6
	7
	8
	9
	10
	11
	12
	13
	14
	15
	16
	17
	18
	19
	20
	21
	22
	23
	24
	25
	26
	27
	28
	29
	30
	31
	32
	33
	34
	35
	36
	37
	38
	39
	40
	41
	42
	43
	44
	45
	46
	47
	48

MPI Declarations Index

This index refers to declarations needed in C/C++, such as address kind integers, handles, etc. The underlined page numbers is the “main” reference (sometimes there are more than one when key concepts are discussed in multiple areas).

MPI::Aint, [78](#)
MPI::Cartcomm, [79](#)
MPI::Comm, [79](#)
MPI::Datatype, [79](#)
MPI::Distgraphcomm, [79](#)
MPI::Errhandler, [8](#), [9–12](#), [79](#)
MPI::Exception, [79](#)
MPI::File, [11](#), [12](#), [19](#), [79](#)
MPI::Graphcomm, [79](#)
MPI::Grequest, [79](#)
MPI::Group, [79](#)
MPI::Info, [4](#), [79](#)
MPI::Intercomm, [79](#)
MPI::Intracomm, [79](#)
MPI::Offset, [78](#)
MPI::Op, [79](#)
MPI::Prequest, [79](#)
MPI::Request, [79](#)
MPI::Status, [79](#)
MPI::Win, [10](#), [11](#), [18](#), [79](#)
MPI_Aint, [78](#)
MPI_Comm, [71](#), [72](#), [78](#)
MPI_Datatype, [69–72](#), [78](#)
MPI_Errhandler, [8](#), [9–12](#), [68](#), [72](#), [78](#)
MPI_File, [11](#), [12](#), [19](#), [72](#), [78](#)
MPI_Fint, [76](#), [78](#)
MPI_Group, [72](#), [78](#)
MPI_Info, [4](#), [72](#), [78](#)
MPI_Offset, [78](#)
MPI_Op, [72](#), [78](#)
MPI_Request, [72](#), [78](#)
MPI_Status, [76](#), [78](#)
MPI_T_cvar_handle, [45](#)
MPI_T_enum, [41](#)
MPI_T_pvar_handle, [54](#)
MPI_T_pvar_session, [53](#)
MPI_Win, [10](#), [11](#), [18](#), [72](#), [78](#)

MPI Callback Function Prototype Index

This index lists the C typedef names for callback routines, such as those used with attribute caching or user-defined reduction operations. C++ names for these typedefs and Fortran example prototypes are given near the text of the C name.

MPI_Comm_copy_attr_function, [73](#), [79](#)
MPI_Comm_delete_attr_function, [73](#), [79](#)
MPI_Comm_errhandler_function, [8](#), [80](#)
MPI_Copy_function, [74](#), [83](#)
MPI_Datarep_conversion_function, [80](#)
MPI_Datarep_extent_function, [80](#)
MPI_Delete_function, [74](#), [83](#)
MPI_File_errhandler_function, [11](#), [80](#)
MPI_Grequest_cancel_function, [80](#)
MPI_Grequest_free_function, [80](#)
MPI_Grequest_query_function, [80](#)
MPI_Handler_function, [83](#)
MPI_Type_copy_attr_function, [73](#), [79](#)
MPI_Type_delete_attr_function, [73](#), [79](#)
MPI_User_function, [79](#)
MPI_Win_copy_attr_function, [73](#), [79](#)
MPI_Win_delete_attr_function, [73](#), [79](#)
MPI_Win_errhandler_function, [10](#), [80](#)

MPI Function Index

The underlined page numbers refer to the function definitions.

- MPI_ABORT, [6](#), [22](#), [25](#), [40](#)
- MPI_ADD_ERROR_CLASS, [16](#), [16](#)
- MPI_ADD_ERROR_CODE, [17](#)
- MPI_ADD_ERROR_STRING, [17](#), [18](#)
- MPI_ALLOC_MEM, [4](#), [5](#), [14](#)
- MPI_BSEND, [4](#), [23](#)
- MPI_BUFFER_DETACH, [23](#)
- MPI_COMM_CALL_ERRHANDLER, [18](#),
[19](#)
- MPI_COMM_CONNECT, [14](#)
- MPI_COMM_CREATE_ERRHANDLER, [7](#),
[8](#), [9](#), [81](#)
- MPI_COMM_CREATE_KEYVAL, [80](#)
- MPI_COMM_DUP_FN, [73](#)
- MPI_COMM_FREE, [26](#)
- MPI_COMM_GET_ATTR, [2](#)
- MPI_COMM_GET_ERRHANDLER, [7](#), [9](#)
- MPI_COMM_GROUP, [7](#)
- MPI_COMM_NULL_COPY_FN, [73](#)
- MPI_COMM_NULL_DELETE_FN, [73](#)
- MPI_COMM_SET_ERRHANDLER, [7](#), [9](#)
- MPI_COMM_SPAWN, [27](#), [28](#)
- MPI_COMM_SPAWN_MULTIPLE, [27](#)
- MPI_ERRHANDLER_CREATE, [8](#)
- MPI_ERRHANDLER_FREE, [7](#), [12](#)
- MPI_ERRHANDLER_GET, [7](#), [9](#)
- MPI_ERRHANDLER_SET, [9](#)
- MPI_ERROR_CLASS, [13](#), [15](#), [15](#), [16](#)
- MPI_ERROR_STRING, [13](#), [13](#), [16](#), [18](#)
- MPI_FILE_CALL_ERRHANDLER, [19](#), [19](#)
- MPI_FILE_CREATE_ERRHANDLER, [7](#), [11](#),
[12](#), [81](#)
- MPI_FILE_GET_ERRHANDLER, [7](#), [12](#)
- MPI_FILE_OPEN, [15](#)
- MPI_FILE_SET_ERRHANDLER, [7](#), [12](#)
- MPI_FILE_SET_VIEW, [15](#)
- MPI_FINALIZE, [2](#), [21](#), [22–26](#), [36](#), [45](#), [58](#),
[60](#)
- MPI_FINALIZED, [21](#), [24](#), [26](#), [26](#)
- MPI_FREE_MEM, [5](#), [5](#), [14](#)
- MPI_GET_PROCESSOR_NAME, [3](#), [4](#)
- MPI_GET_VERSION, [1](#), [2](#), [21](#), [24](#)
- MPI_GREQUEST_START, [81](#)
- MPI_GROUP_FREE, [7](#)
- MPI_INFO_DELETE, [14](#)
- MPI_INIT, [2](#), [21](#), [21](#), [24–26](#), [32](#), [36](#), [39](#), [40](#),
[45](#), [58](#)
- MPI_INIT_THREAD, [21](#), [26](#), [39](#)
- MPI_INITIALIZED, [21](#), [24](#), [24](#), [25](#), [26](#)
- MPI_ISEND, [22](#)
- MPI_KEYVAL_CREATE, [83](#)
- MPI_LOOKUP_NAME, [14](#)
- MPI_OP_CREATE, [80](#)
- MPI_PCONTROL, [31](#), [32](#), [32](#), [33](#)
- MPI_RECV, [58](#)
- MPI_REGISTER_DATAREP, [15](#), [82](#)
- MPI_REQUEST_FREE, [22](#)
- MPI_SEND, [33](#)
- MPI_T_CATEGORY_CHANGED, [63](#)
- MPI_T_CATEGORY_GET_CATEGORIES,
[62](#), [62](#), [63](#)
- MPI_T_CATEGORY_GET_CVARS, [62](#), [62](#),
[63](#)
- MPI_T_CATEGORY_GET_INFO, [61](#), [63](#)
- MPI_T_CATEGORY_GET_NUM, [61](#)
- MPI_T_CATEGORY_GET_PVARS, [62](#), [62](#),
[63](#)
- MPI_T_CVAR_GET_INFO, [41–43](#), [43](#), [44](#),
[46](#), [47](#), [63](#)
- MPI_T_CVAR_GET_NUM, [42](#), [46](#)
- MPI_T_CVAR_HANDLE_ALLOC, [46](#), [46](#),
[47](#)
- MPI_T_CVAR_HANDLE_FREE, [46](#), [46](#)

MPI_T_CVAR_READ, 47	1
MPI_T_CVAR_WRITE, 47	2
MPI_T_ENUM_GET_INFO, 41 , 41	3
MPI_T_ENUM_GET_ITEM, 41 , 41	4
MPI_T_ENUMITEM, 64	5
MPI_T_FINALIZE, 39 , 39 , 40	6
MPI_T_INIT_THREAD, 39 , 39 , 40	7
MPI_T_PVAR_GET_INFO, 41 , 50–52 , 52 , 54 , 56 , 57 , 63	8 9
MPI_T_PVAR_GET_NUM, 51 , 54	10
MPI_T_PVAR_HANDLE_ALLOC, 54 , 54 , 56	11 12
MPI_T_PVAR_HANDLE_FREE, 55 , 55	13
MPI_T_PVAR_READ, 56 , 56 , 57	14
MPI_T_PVAR_READRESET, 53 , 57 , 57 , 64	15 16
MPI_T_PVAR_RESET, 57 , 57 , 64	17
MPI_T_PVAR_SESSION_CREATE, 53	18
MPI_T_PVAR_SESSION_FREE, 53	19
MPI_T_PVAR_START, 55 , 64	20
MPI_T_PVAR_STOP, 55 , 64	21
MPI_T_PVAR_WRITE, 56 , 56 , 64	22
MPI_TEST, 22	23
MPI_TYPE_CREATE_KEYVAL, 81	24
MPI_TYPE_DUP_FN, 73	25
MPI_TYPE_NULL_COPY_FN, 73	26
MPI_TYPE_NULL_DELETE_FN, 73	27
MPI_TYPE_SIZE, 33	28
MPI_UNPUBLISH_NAME, 14	29
MPI_WAIT, 22	30
MPI_WIN_CALL_ERRHANDLER, 18 , 19	31
MPI_WIN_CREATE_ERRHANDLER, 7 , 10 , 10 , 81	32 33
MPI_WIN_CREATE_KEYVAL, 80	34
MPI_WIN_DUP_FN, 73	35
MPI_WIN_GET_ERRHANDLER, 7 , 11	36
MPI_WIN_LOCK, 4	37
MPI_WIN_NULL_COPY_FN, 73	38
MPI_WIN_NULL_DELETE_FN, 73	39
MPI_WIN_SET_ERRHANDLER, 7 , 10	40
MPI_WIN_UNLOCK, 4	41
MPI_WTICK, 20 , 20	42
MPI_WTIME, 3 , 20 , 20 , 33 , 50	43
mpiexec, 21 , 25 , 27 , 27	44
mpirun, 27	45
PMPI_, 30	46 47 48