

```
1 MPI_WTIME()
```

```
2
3 double MPI_Wtime(void)
```

```
4 DOUBLE PRECISION MPI_WTIME()
```

```
5
6 {double MPI::Wtime() (binding deprecated, see Section 15.2) }
```

```
7
8 MPI_WTIME returns a floating-point number of seconds, representing elapsed wall-
9 clock time since some time in the past.
```

```
10 The “time in the past” is guaranteed not to change during the life of the process.
11 The user is responsible for converting large numbers of seconds to other units if they are
12 preferred.
```

```
13 This function is portable (it returns seconds, not “ticks”), it allows high-resolution,
14 and carries no unnecessary baggage. One would use it like this:
```

```
15 {
16     double starttime, endtime;
17     starttime = MPI_Wtime();
18     .... stuff to be timed ...
19     endtime   = MPI_Wtime();
20     printf("That took %f seconds\n",endtime-starttime);
21 }
22
```

```
23 The times returned are local to the node that called them. There is no requirement
24 that different nodes return “the same time.” (But see also the discussion of
25 MPI_WTIME_IS_GLOBAL).
```

```
26
27 MPI_WTICK()
```

```
28
29 double MPI_Wtick(void)
```

```
30 DOUBLE PRECISION MPI_WTICK()
```

```
31
32 {double MPI::Wtick() (binding deprecated, see Section 15.2) }
```

```
33
34 MPI_WTICK returns the resolution of MPI_WTIME in seconds. That is, it returns,
35 as a double precision value, the number of seconds between successive clock ticks. For
36 example, if the clock is implemented by the hardware as a counter that is incremented
37 every millisecond, the value returned by MPI_WTICK should be  $10^{-3}$ .
38
```

39 8.7 Startup

```
40
41 One goal of MPI is to achieve source code portability. By this we mean that a program writ-
42 ten using MPI and complying with the relevant language standards is portable as written,
43 and must not require any source code changes when moved from one system to another.
44 This explicitly does not say anything about how an MPI program is started or launched from
45 the command line, nor what the user must do to set up the environment in which an MPI
46 program will run. However, an implementation may require some setup to be performed
47
48
```

before other MPI routines may be called. To provide for this, MPI includes an initialization routine `MPI_INIT`.

`MPI_INIT()`

```
int MPI_Init(int *argc, char ***argv)
```

```
MPI_INIT(IERROR)
```

```
INTEGER IERROR
```

```
{void MPI::Init(int& argc, char**& argv) (binding deprecated, see Section 15.2) }
```

```
{void MPI::Init() (binding deprecated, see Section 15.2) }
```

[All MPI programs must contain exactly one call to an MPI initialization routine: `MPI_INIT` or `MPI_INIT_THREAD`.] Each MPI process must call an MPI initialization routine, `MPI_INIT` or `MPI_INIT_THREAD`, exactly once. Subsequent calls by that process to any initialization routines are erroneous. The only MPI functions that may be invoked before the MPI initialization routines are called are `MPI_GET_VERSION`, `MPI_INITIALIZED`, and `MPI_FINALIZED`.

The version for ISO C accepts the `argc` and `argv` that are provided by the arguments to `main` or `NULL`:

```
int main(int argc, char **argv)
```

```
{
```

```
    MPI_Init(&argc, &argv);
```

```
    /* parse arguments */
```

```
    /* main program    */
```

```
    MPI_Finalize();    /* see below */
```

```
}
```

The Fortran version takes only `IERROR`.

Conforming implementations of MPI are required to allow applications to pass `NULL` for both the `argc` and `argv` arguments of `main` in C. [and C++. In C++, there is an alternative binding for `MPI::Init` that does not have these arguments at all.

Rationale. In some applications, libraries may be making the call to `MPI_Init`, and may not have access to `argc` and `argv` from `main`. It is anticipated that applications requiring special information about the environment or information supplied by `mpixexec` can get that information from environment variables. (*End of rationale.*)

]

After MPI is initialized, the application can access information about the execution environment by querying the predefined info object `MPI_INFO_GET_ENV`. The following keys are predefined for this object, corresponding to the arguments of `MPI_COMM_SPAWN` or of `mpixexec`:

`command` name of program executed

`argv` (space separated) arguments to command
`maxprocs` Maximum number of MPI processes to start.
`soft` Allowed values for number of processors
`host` Hostname.
`arch` Architecture name.
`wdir` Working directory of the MPI process
`file` Value is the name of a file in which additional information is specified.
`thread_level` Requested level of thread support, if requested before the program started execution.

Note that all values are strings. Thus, the maximum number of processes is represented by a string such as “1024” and the requested level is represented by a string such as “MPI_THREAD_SINGLE”

The info object `MPI_INFO_GET_ENV` need not contain a (key,value) pair for each of these predefined keys; the set of (key,value) pairs provided is implementation-dependent. Implementations may provide additional, implementation specific, (key,value) pairs.

In case where the MPI processes were started with `MPI_COMM_SPAWN_MULTIPLE` or, equivalently, with a startup mechanism that supports multiple process specifications, then the values stored in the info object `MPI_INFO_KEY` at a process are those values that affect the local MPI process.

Example 8.3 If MPI is started with a call to

```
mpiexec -n 5 -arch sun ocean : -n 10 -arch rs6000 atmos
```

Then the first 5 processes will have in their `MPI_INFO_GET_ENV` object the pairs (command, ocean), (maxprocs, 5), and (arch, sun). The next 10 processes will have in `MPI_INFO_KEY` (command, atmos), (maxprocs, 10), and (arch, rs6000)

Advice to users. The values passed in `MPI_INFO_KEY` are the values of the arguments passed to the mechanism that started the MPI execution – not the actual value provided. Thus, the value associated with `maxprocs` is the number of MPI processes requested; it can be larger than the actual number of processes obtained, if the `soft` option was used. (*End of advice to users.*)

Advice to implementors. Good quality implementations will provide a (key,value) pair for each parameter that can be passed to the command that starts an MPI program. (*End of advice to implementors.*)

`MPI_FINALIZE()`

```
int MPI_Finalize(void)
```

```
MPI_FINALIZE(IERROR)
```

INTEGER IERROR

```
{void MPI::Finalize() (binding deprecated, see Section 15.2) }
```

This routine cleans up all MPI state. [

Each process must call `MPI_FINALIZE` before it exits. Unless there has been a call to `MPI_ABORT`, before each process exits process must ensure that all pending nonblocking communications are (locally) complete before calling `MPI_FINALIZE`. Further, at the instant at which the last process calls `MPI_FINALIZE`, all pending sends must be matched by a receive, and all pending receives must be matched by a send.

For example, the following program is correct] If an MPI program terminates normally (i.e., not due to a call to `MPI_ABORT` or an unrecoverable error) then each process must call `MPI_FINALIZE` before it exits.

Before an MPI process invokes `MPI_FINALIZE`, the process must perform all the MPI calls needed to complete its involvement in MPI communications: It must locally complete all the MPI operations it initiated and must execute matching calls needed to complete MPI communications initiated by other processes. For example, if the process executed a nonblocking send, it must follow up with a call to `MPI_WAIT`, `MPI_TEST`, `MPI_REQUEST_FREE` or any derived function; if the process is the target of a send, then it must post the matching receive; if it is part of a group executing a collective operation, then it must have completed its participation in the operations; etc.

The call to `MPI_FINALIZE` does not free objects created by MPI calls; these objects are freed using `MPI_XXX_FREE` calls.

`MPI_FINALIZE` is collective over all connected processes. If no processes were spawned, accepted or connected then this means over `MPI_COMM_WORLD`; otherwise it is collective over the union of all processes that have been and continue to be connected, as explained in Section 10.5.4 on page 356.

The following examples illustrates these rules

Example 8.4 The following code is correct

Process 0	Process 1
-----	-----
<code>MPI_Init();</code>	<code>MPI_Init();</code>
<code>MPI_Send(dest=1);</code>	<code>MPI_Recv(src=0);</code>
<code>MPI_Finalize();</code>	<code>MPI_Finalize();</code>

Example 8.5 Without a matching receive, the program is erroneous

Process 0	Process 1
-----	-----
<code>MPI_Init();</code>	<code>MPI_Init();</code>
<code>MPI_Send (dest=1);</code>	
<code>MPI_Finalize();</code>	<code>MPI_Finalize();</code>

[deleted in April Since `MPI_FINALIZE` is a collective call, a correct MPI program will naturally ensure that all participants in pending collective operations have made the call before calling `MPI_FINALIZE`.

A successful return from a blocking communication operation or from `MPI_WAIT` or `MPI_TEST` tells the user that the buffer can be reused and means that the communication

is completed by the user, but does not guarantee that the local process has no more work to do. A successful return from `MPI_REQUEST_FREE` with a request handle generated by an `MPI_ISEND` nullifies the handle but provides no assurance of operation completion. The `MPI_ISEND` is complete only when it is known by some means that a matching receive has completed. `MPI_FINALIZE` guarantees that all local actions required by communications the user has completed will, in fact, occur before it returns.

`MPI_FINALIZE` guarantees nothing about pending communications that have not been completed (completion is assured only by `MPI_WAIT`, `MPI_TEST`, or `MPI_REQUEST_FREE` combined with some other verification of completion).

Example 8.6 This program is correct HEADER SKIP ENDHEADER

```
rank 0                                rank 1
=====
...
MPI_Isend();                          MPI_Recv();
MPI_Request_free();                   MPI_Barrier();
MPI_Barrier();                       MPI_Finalize();
MPI_Finalize();                      exit();
exit();
```

Example 8.7 This program is erroneous and its behavior is undefined: HEADER SKIP ENDHEADER

```
rank 0                                rank 1
=====
...
MPI_Isend();                          MPI_Recv();
MPI_Request_free();                   MPI_Finalize();
MPI_Finalize();                      exit();
exit();
```

ticket313.

Example 8.8 This program is correct: Process 0 calls `MPI_Finalize` after it has executed the MPI calls that complete the send operation. Likewise, process 1 executes the MPI call that completes the matching receive operation before it calls `MPI_Finalize`.

```
Process 0                            Proces 1
-----
MPI_Init();                          MPI_Init();
MPI_Isend(dest=1);                   MPI_Recv(src=0);
MPI_Request_free();                  MPI_Finalize();
MPI_Finalize();                      exit();
exit();
```

ticket313.

[If no `MPI_BUFFER_DETACH` occurs between an `MPI_BSEND` (or other buffered send) and `MPI_FINALIZE`, the `MPI_FINALIZE` implicitly supplies the `MPI_BUFFER_DETACH`.

Example 8.9 This program is correct, and after the `MPI_Finalize`, it is as if the buffer had been detached. HEADER SKIP ENDHEADER

```

rank 0                                     rank 1
=====
...
buffer = malloc(1000000);
MPI_Buffer_attach();
MPI_Bsend();
MPI_Finalize();
free(buffer);
exit();
]

```

1
2
3
4
5
6
7
8
9
10
11 ticket313.

Example 8.10 This program is correct. The attached buffer is a resource allocated by the user, not by MPI; it is available to the user after MPI is finalized.

```

Process 0                                Process 1
-----
MPI_Init();                               MPI_Init();
buffer = malloc(1000000);                 MPI_Recv(src=0);
MPI_Buffer_attach();                       MPI_Finalize();
MPI_Send(dest=1));                         exit();
MPI_Finalize();
free(buffer);
exit();
[

```

12
13
14
15
16
17
18
19
20
21
22
23 ticket313.

Example 8.11 In this example, `MPI_lprobe()` must return a `FALSE` flag. `MPI_Test_cancelled()` must return a `TRUE` flag, independent of the relative order of execution of `MPI_Cancel()` in process 0 and `MPI_Finalize()` in process 1.

The `MPI_lprobe()` call is there to make sure the implementation knows that the “tag1” message exists at the destination, without being able to claim that the user knows about it.

HEADER SKIP ENDHEADER

```

rank 0                                     rank 1
=====
MPI_Init();                               MPI_Init();
MPI_Isend(tag1);                           MPI_Barrier();
MPI_Barrier();                             MPI_Iprobe(tag2);
MPI_Barrier();                             MPI_Barrier();
                                           MPI_Finalize();
                                           exit();

MPI_Cancel();
MPI_Wait();
MPI_Test_cancelled();
MPI_Finalize();
exit();

```

24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48

ticket313. 1]

2
3 **Example 8.12** This program is correct. The cancel operation must succeed, since the
4 send cannot complete normally. The wait operation, after the call to `MPI_Cancel`, is local
5 – no matching MPI call is required on process 1.

6
7
8 Process 0 Process 1
9 ----- -----
10 MPI_Issend(dest=1); MPI_Finalize();
11 MPI_Cancel();
12 MPI_Wait();
13 MPI_Finalize();

ticket313. 14 [

15
16 *Advice to implementors.* An implementation may need to delay the return from
17 `MPI_FINALIZE` until all potential future message cancellations have been processed.
18 One possible solution is to place a barrier inside `MPI_FINALIZE` (*End of advice to*
19 *implementors.*)

20
21]

22 *Advice to implementors.* Even though a process has [completed all]has executed all
23 the MPI calls needed to complete the communications it [initiated]is involved with,
24 such communication may not yet be completed from the viewpoint of the underlying
25 MPI system. E.g., a blocking send may have returned, even though the data is still
26 buffered at the sender in an MPI buffer; an MPI process may receive a cancel request for
27 a message it has completed receiving; etc. The MPI implementation must ensure that a
28 process has completed any involvement in MPI communication before `MPI_FINALIZE`
29 returns. Thus, if a process exits after the call to `MPI_FINALIZE`, this will not cause
30 an ongoing communication to fail. The MPI implementation should also complete
31 freeing all objects marked for deletion by MPI calls that freed them. (*End of advice*
32 *to implementors.*)

33
34 Once `MPI_FINALIZE` returns, no MPI routine (not even `MPI_INIT`) may be called,
35 except for `MPI_GET_VERSION`, `MPI_INITIALIZED`, and `MPI_FINALIZED`.

36 [Each process must complete any pending communication it initiated before it calls
37 `MPI_FINALIZE`. If the call returns, each process may continue local computations, or exit,
38 without participating in further MPI communication with other processes.

39 `MPI_FINALIZE` is collective on `MPI_COMM_WORLD`. `MPI_FINALIZE` is collective over
40 all connected processes. If no processes were spawned, accepted or connected then this
41 means over `MPI_COMM_WORLD`; otherwise it is collective over the union of all processes
42 that have been and continue to be connected, as explained in Section 10.5.4 on page 356.

43
44 *Advice to implementors.* Even though a process has completed all the communication
45 it initiated, such communication may not yet be completed from the viewpoint of the
46 underlying MPI system. E.g., a blocking send may have completed, even though the
47 data is still buffered at the sender. The MPI implementation must ensure that a
48 process has completed any involvement in MPI communication before `MPI_FINALIZE`

returns. Thus, if a process exits after the call to `MPI_FINALIZE`, this will not cause an ongoing communication to fail. (*End of advice to implementors.*)

]

Although it is not required that all processes return from `MPI_FINALIZE`, it is required that at least process 0 in `MPI_COMM_WORLD` return, so that users can know that the MPI portion of the computation is over. In addition, in a POSIX environment, they may desire to supply an exit code for each process that returns from `MPI_FINALIZE`.

Example 8.13 The following illustrates the use of requiring that at least one process return and that it be known that process 0 is one of the processes that return. One wants code like the following to work no matter how many processes return.

```
...
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
...
MPI_Finalize();
if (myrank == 0) {
    resultfile = fopen("outfile","w");
    dump_results(resultfile);
    fclose(resultfile);
}
exit(0);
```

`MPI_INITIALIZED(flag)`

OUT	flag	Flag is true if <code>MPI_INIT</code> has been called and false otherwise.
-----	------	--

`int MPI_Initialized(int *flag)`

`MPI_INITIALIZED(FLAG, IERROR)`

LOGICAL FLAG

INTEGER IERROR

`{bool MPI::Is_initialized() (binding deprecated, see Section 15.2) }`

This routine may be used to determine whether `MPI_INIT` has been called. `MPI_INITIALIZED` returns true if the calling process has called `MPI_INIT`. Whether `MPI_FINALIZE` has been called does not affect the behavior of `MPI_INITIALIZED`. It is one of the few routines that may be called before `MPI_INIT` is called.

`MPI_ABORT(comm, errorcode)`

IN	comm	communicator of tasks to abort
----	------	--------------------------------

IN	errorcode	error code to return to invoking environment
----	-----------	--

`int MPI_Abort(MPI_Comm comm, int errorcode)`

INTENSIONALLY LEFT EMPTY

The level(s) of thread support that can be provided by `MPI_INIT_THREAD` will depend on the implementation, and may depend on information provided by the user before the program started to execute (e.g., with arguments to `mpiexec`). If possible, the call will return `provided = required`. Failing this, the call will return the least supported level such that `provided > required` (thus providing a stronger level of support than required by the user). Finally, if the user requirement cannot be satisfied, then the call will return in `provided` the highest supported level.

A **thread compliant** MPI implementation will be able to return `provided = MPI_THREAD_MULTIPLE`. Such an implementation may always return `provided = MPI_THREAD_MULTIPLE`, irrespective of the value of `required`. [At the other extreme, an MPI library that is not thread compliant may always return `provided = MPI_THREAD_SINGLE`, irrespective of the value of `required`.]

An MPI library that is not thread compliant must always return `provided=MPI_THREAD_SINGLE`, even if `MPI_INIT_THREAD` is called on a multithreaded process. The library should also return correct values for the MPI calls that can be executed before initialization, even if multiple threads have been spawned.

Rationale. Such code is erroneous, but if the MPI initialization is performed by a library, the error cannot be detected until `MPI_INIT_THREAD` is called. The requirements in the previous paragraph ensure that the error can be properly detected. (*End of rationale.*)

A call to `MPI_INIT` has the same effect as a call to `MPI_INIT_THREAD` with a `required = MPI_THREAD_SINGLE`.

Vendors may provide (implementation dependent) means to specify the level(s) of thread support available when the MPI program is started, e.g., with arguments to `mpiexec`. This will affect the outcome of calls to `MPI_INIT` and `MPI_INIT_THREAD`. Suppose, for example, that an MPI program has been started so that only `MPI_THREAD_MULTIPLE` is available. Then `MPI_INIT_THREAD` will return `provided = MPI_THREAD_MULTIPLE`, irrespective of the value of `required`; a call to `MPI_INIT` will also initialize the MPI thread support level to `MPI_THREAD_MULTIPLE`. Suppose, on the other hand, that an MPI program has been started so that all four levels of thread support are available. Then, a call to `MPI_INIT_THREAD` will return `provided = required`; on the other hand, a call to `MPI_INIT` will initialize the MPI thread support level to `MPI_THREAD_SINGLE`.

Rationale. Various optimizations are possible when MPI code is executed single-threaded, or is executed on multiple threads, but not concurrently: mutual exclusion code may be omitted. Furthermore, if only one thread executes, then the MPI library can use library functions that are not thread safe, without risking conflicts with user threads. Also, the model of one communication thread, multiple computation threads fits many applications well, e.g., if the process code is a sequential Fortran/C/C++ program with MPI calls that has been parallelized by a compiler for execution on an SMP node, in a cluster of SMPs, then the process computation is multi-threaded, but MPI calls will likely execute on a single thread.

The design accommodates a static specification of the thread support level, for environments that require static binding of libraries, and for compatibility for current multi-threaded MPI codes. (*End of rationale.*)

Advice to implementors. If `provided` is not `MPI_THREAD_SINGLE` then the MPI library should not invoke C/ C++/Fortran library calls that are not thread safe, e.g., in an environment where `malloc` is not thread safe, then `malloc` should not be used by the MPI library.

Some implementors may want to use different MPI libraries for different levels of thread support. They can do so using dynamic linking and selecting which library will be linked when `MPI_INIT_THREAD` is invoked. If this is not possible, then optimizations for lower levels of thread support will occur only when the level of thread support required is specified at link time. (*End of advice to implementors.*)

The following function can be used to query the current level of thread support.

`MPI_QUERY_THREAD(provided)`

OUT `provided` provided level of thread support (integer)

`int MPI_Query_thread(int *provided)`

`MPI_QUERY_THREAD(PROVIDED, IERROR)`

INTEGER `PROVIDED`, IERROR

{`int MPI::Query_thread()` (*binding deprecated, see Section 15.2*) }

The call returns in `provided` the current level of thread [support. This]support, which will be the value returned in `provided` by `MPI_INIT_THREAD`, if MPI was initialized by a call to `MPI_INIT_THREAD`.

`MPI_IS_THREAD_MAIN(flag)`

OUT `flag` true if calling thread is main thread, false otherwise
(logical)

`int MPI_Is_thread_main(int *flag)`

`MPI_IS_THREAD_MAIN(FLAG, IERROR)`

LOGICAL `FLAG`

INTEGER IERROR

{`bool MPI::Is_thread_main()` (*binding deprecated, see Section 15.2*) }

This function can be called by a thread to [find out whether]determine if it is the main thread.

All routines listed in this section must be supported by all MPI implementations.

Rationale. MPI libraries are required to provide these calls even if they do not support threads, so that portable code that contains invocations to these functions [be able to]can link correctly. `MPI_INIT` continues to be supported so as to provide compatibility with current MPI codes. (*End of rationale.*)

Advice to users. It is possible to spawn threads before MPI is initialized, but no MPI call other than [`MPI_INITIALIZED`] `MPI_GET_VERSION`, `MPI_INITIALIZED`, or

MPI_FINALIZED should be executed by these threads, until **MPI_INIT_THREAD** is invoked by one thread (which, thereby, becomes the main thread). In particular, it is possible to enter the MPI execution with a multi-threaded process.

The level of thread support provided is a global property of the MPI process that can be specified only once, when MPI is initialized on that process (or before). Portable third party libraries have to be written so as to accommodate any provided level of thread support. Otherwise, their usage will be restricted to specific level(s) of thread support. If such a library can run only with specific level(s) of thread support, e.g., only with **MPI_THREAD_MULTIPLE**, then **MPI_QUERY_THREAD** can be used to check whether the user initialized MPI to the correct level of thread support and, if not, raise an exception. (*End of advice to users.*)