

before other MPI routines may be called. To provide for this, MPI includes an initialization routine `MPI_INIT`.

`MPI_INIT()`

```
int MPI_Init(int *argc, char ***argv)
```

```
MPI_INIT(IERROR)
```

```
INTEGER IERROR
```

```
{void MPI::Init(int& argc, char**& argv) (binding deprecated, see Section 15.2) }
```

```
{void MPI::Init() (binding deprecated, see Section 15.2) }
```

All MPI processes must contain exactly one call to an MPI initialization routine: `MPI_INIT` or `MPI_INIT_THREAD`. Subsequent calls to any initialization routines are erroneous. The only MPI functions that may be invoked before the MPI initialization routines are called are `MPI_GET_VERSION`, `MPI_INITIALIZED`, and `MPI_FINALIZED`.

The version for ISO C accepts the `argc` and `argv` that are provided by the arguments to `main` or `NULL`:

```
int main(int argc, char **argv)
```

```
{
```

```
    MPI_Init(&argc, &argv);
```

```
    /* parse arguments */
```

```
    /* main program    */
```

```
    MPI_Finalize();    /* see below */
```

```
}
```

The Fortran version takes only `IERROR`.

Conforming implementations of MPI are required to allow applications to pass `NULL` for both the `argc` and `argv` arguments of `main` in C. [and C++. In C++, there is an alternative binding for `MPI::Init` that does not have these arguments at all.

Rationale. In some applications, libraries may be making the call to `MPI_Init`, and may not have access to `argc` and `argv` from `main`. It is anticipated that applications requiring special information about the environment or information supplied by `mpiexec` can get that information from environment variables. (*End of rationale.*)

]

After MPI is initialized, the application can access information about the execution environment by querying the predefined info object `MPI_INFO_ENV`. The following keys are predefined for this object, corresponding to the arguments of `MPI_COMM_SPAWN` or of `mpiexec`:

`MPI_ENV_N` Total number of MPI processes.

`MPI_ENV_THREAD_LEVEL` Level of thread support.

MPI_ENV_HOST Hostname.
MPI_ENV_ARCH Architecture name.
MPI_ENV_WDIR Working directory of the MPI process
MPI_ENV_FILE Value is the name of a file in which additional information is specified.
 Implementations may provide additional, implementation specific, keys.

MPI_FINALIZE()

```
int MPI_Finalize(void)
```

```
MPI_FINALIZE(IERROR)
```

```
INTEGER IERROR
```

```
{void MPI::Finalize() (binding deprecated, see Section 15.2) }
```

This routine cleans up all MPI state. [Each process must call **MPI_FINALIZE** before it exits. Unless there has been a call to **MPI_ABORT**, before each process exits process must ensure that all pending nonblocking communications are (locally) complete before calling **MPI_FINALIZE**. Further, at the instant at which the last process calls **MPI_FINALIZE**, all pending sends must be matched by a receive, and all pending receives must be matched by a send.

For example, the following program is correct

] If an MPI program terminates normally (i.e., not due to a call to **MPI_ABORT** or an unrecovered error) then MPI must be finalized at each MPI process. MPI is finalized at a process by a call to **MPI_FINALIZE** on this process.

Before MPI is finalized at an MPI process, the process must locally complete all MPI calls and free all objects created by MPI calls on that process.

When the last process calls **MPI_FINALIZE**, all non-local MPI calls at each process must be matched by MPI calls at the other processes that are needed to complete the relevant operation: For each send, the matching receive has occurred, each collective operation has been invoked at all involved processes, etc.

The following examples illustrates these rules

Example 8.3 The following code is correct

Process 0	Process 1
-----	-----
MPI_Init();	MPI_Init();
MPI_Send(dest=1);	MPI_Recv(src=0);
MPI_Finalize();	MPI_Finalize();

Example 8.4 Without a matching receive, the program is erroneous

Process 0	Process 1
-----	-----
MPI_Init();	MPI_Init();
MPI_Send (dest=1);	
MPI_Finalize();	MPI_Finalize();

A successful return from a blocking communication operation or from `MPI_WAIT` or `MPI_TEST` tells the user that the buffer can be reused and means that the communication is completed by the user, but does not guarantee that the local process has no more work to do. A successful return from `MPI_REQUEST_FREE` with a request handle generated by an `MPI_ISEND` nullifies the handle but provides no assurance of operation completion. The `MPI_ISEND` is complete only when it is known by some means that a matching receive has completed. `MPI_FINALIZE` guarantees that all local actions required by communications the user has completed will, in fact, occur before it returns.

`MPI_FINALIZE` guarantees nothing about pending communications that have not been completed (completion is assured only by `MPI_WAIT`, `MPI_TEST`, or `MPI_REQUEST_FREE` combined with some other verification of completion).

[

Example 8.5 This program is correct `HEADER SKIP ENDHEADER`

```
rank 0                                rank 1
=====
...
MPI_Isend();                          MPI_Recv();
MPI_Request_free();                   MPI_Barrier();
MPI_Barrier();                       MPI_Finalize();
MPI_Finalize();                      exit();
exit();
```

Example 8.6 This program is erroneous and its behavior is undefined: `HEADER SKIP ENDHEADER`

```
rank 0                                rank 1
=====
...
MPI_Isend();                          MPI_Recv();
MPI_Request_free();                   MPI_Finalize();
MPI_Finalize();                      exit();
exit();
```

]

Example 8.7 This program is erroneous: The `MPI_Isend` call is not guaranteed to be complete even if the send request was freed and the matching receive completed; the call `MPI_Finalize` cannot be used to complete the call.

```
Process 0                            Process 1
-----
MPI_Isend();                         MPI_Recv();
MPI_Request_free();                  MPI_Barrier();
MPI_Barrier();                      MPI_Finalize();
MPI_Finalize();
```

[If no `MPI_BUFFER_DETACH` occurs between an `MPI_BSEND` (or other buffered send) and `MPI_FINALIZE`, the `MPI_FINALIZE` implicitly supplies the `MPI_BUFFER_DETACH`.

Example 8.8 This program is correct, and after the `MPI_Finalize`, it is as if the buffer had been detached. HEADER SKIP ENDHEADER

```

rank 0                                rank 1
=====
...
buffer = malloc(1000000);             MPI_Recv();
MPI_Buffer_attach();                  MPI_Finalize();
MPI_Bsend();                          exit();
MPI_Finalize();
free(buffer);
exit();
]
```

Example 8.9 This program is erroneous; the program must call `MPI_Buffer_detach` before the call to `MPI_FINALIZE`.

```

Process 0                             Process 1
-----
buffer = malloc(1000000);             MPI_Recv();
MPI_Buffer_attach();                  MPI_Finalize();
MPI_Bsend();                          exit();
MPI_Finalize();
free(buffer);
exit();
[
```

Example 8.10 In this example, `MPI_lprobe()` must return a `FALSE` flag. `MPI_Test_cancelled()` must return a `TRUE` flag, independent of the relative order of execution of `MPI_Cancel()` in process 0 and `MPI_Finalize()` in process 1.

The `MPI_lprobe()` call is there to make sure the implementation knows that the “tag1” message exists at the destination, without being able to claim that the user knows about it.

HEADER SKIP ENDHEADER

```

rank 0                                rank 1
=====
MPI_Init();                          MPI_Init();
MPI_Isend(tag1);                     MPI_Barrier();
MPI_Barrier();                      MPI_Iprobe(tag2);
MPI_Barrier();                      MPI_Barrier();
MPI_Finalize();                     MPI_Finalize();
exit();
MPI_Cancel();
MPI_Wait();
MPI_Test_cancelled();
```

```
MPI_Finalize();
exit();
```

```
]
```

Example 8.11 This program is correct. The cancel operation must succeed, since the send cannot complete normally.

Process 0

```
MPI_Isend();
MPI_Cancel();
MPI_Wait();
MPI_Finalize();
```

```
[
```

Advice to implementors. An implementation may need to delay the return from MPI_FINALIZE until all potential future message cancellations have been processed. One possible solution is to place a barrier inside MPI_FINALIZE (*End of advice to implementors.*)

```
]
```

Advice to implementors. An implementation may need to delay the return from MPI_FINALIZE on a process even if all communications related to MPI calls by that process have completed; the process may still receive cancel requests for messages it has completed receiving. One possible solution is to place a barrier inside MPI_FINALIZE.

(*End of advice to implementors.*)

Once MPI_FINALIZE returns, no MPI routine (not even MPI_INIT) may be called, except for MPI_GET_VERSION, MPI_INITIALIZED, and MPI_FINALIZED.

[Each process must complete any pending communication it initiated before it calls MPI_FINALIZE. If the call returns, each process may continue local computations, or exit, without participating in further MPI communication with other processes.]

MPI_FINALIZE is collective over all connected processes. If no processes were spawned, accepted or connected then this means over MPI_COMM_WORLD; otherwise it is collective over the union of all processes that have been and continue to be connected, as explained in Section 10.5.4 on page 358.

Advice to implementors. An implementation should check whether MPI is already finalized, before executing the finalization code.

High-quality MPI implementations will free MPI resources not freed by the user before the finalize call, when MPI_FINALIZE executes. (*End of advice to implementors.*)

The level(s) of thread support that can be provided by `MPI_INIT_THREAD` will depend on the implementation, and may depend on information provided by the user before the program started to execute (e.g., with arguments to `mpiexec`). If possible, the call will return `provided = required`. Failing this, the call will return the least supported level such that `provided > required` (thus providing a stronger level of support than required by the user). Finally, if the user requirement cannot be satisfied, then the call will return in `provided` the highest supported level.

A **thread compliant** MPI implementation will be able to return `provided = MPI_THREAD_MULTIPLE`. Such an implementation may always return `provided = MPI_THREAD_MULTIPLE`, irrespective of the value of `required`. [At the other extreme, an MPI library that is not thread compliant may always return `provided = MPI_THREAD_SINGLE`, irrespective of the value of `required`.]

An MPI library that is not thread compliant must always return `provided=MPI_THREAD_SINGLE`, even if `MPI_INIT_THREAD` is called on a multithreaded process.

A call to `MPI_INIT` has the same effect as a call to `MPI_INIT_THREAD` with a `required = MPI_THREAD_SINGLE`.

Vendors may provide (implementation dependent) means to specify the level(s) of thread support available when the MPI program is started, e.g., with arguments to `mpiexec`. This will affect the outcome of calls to `MPI_INIT` and `MPI_INIT_THREAD`. Suppose, for example, that an MPI program has been started so that only `MPI_THREAD_MULTIPLE` is available. Then `MPI_INIT_THREAD` will return `provided = MPI_THREAD_MULTIPLE`, irrespective of the value of `required`; a call to `MPI_INIT` will also initialize the MPI thread support level to `MPI_THREAD_MULTIPLE`. Suppose, on the other hand, that an MPI program has been started so that all four levels of thread support are available. Then, a call to `MPI_INIT_THREAD` will return `provided = required`; on the other hand, a call to `MPI_INIT` will initialize the MPI thread support level to `MPI_THREAD_SINGLE`.

Rationale. Various optimizations are possible when MPI code is executed single-threaded, or is executed on multiple threads, but not concurrently: mutual exclusion code may be omitted. Furthermore, if only one thread executes, then the MPI library can use library functions that are not thread safe, without risking conflicts with user threads. Also, the model of one communication thread, multiple computation threads fits many applications well, e.g., if the process code is a sequential Fortran/C/C++ program with MPI calls that has been parallelized by a compiler for execution on an SMP node, in a cluster of SMPs, then the process computation is multi-threaded, but MPI calls will likely execute on a single thread.

The design accommodates a static specification of the thread support level, for environments that require static binding of libraries, and for compatibility for current multi-threaded MPI codes. (*End of rationale.*)

Advice to implementors. If `provided` is not `MPI_THREAD_SINGLE` then the MPI library should not invoke C/ C++/Fortran library calls that are not thread safe, e.g., in an environment where `malloc` is not thread safe, then `malloc` should not be used by the MPI library.

Some implementors may want to use different MPI libraries for different levels of thread support. They can do so using dynamic linking and selecting which library will be linked when `MPI_INIT_THREAD` is invoked. If this is not possible, then optimizations

for lower levels of thread support will occur only when the level of thread support required is specified at link time. (*End of advice to implementors.*)

The following function can be used to query the current level of thread support.

MPI_QUERY_THREAD(provided)

OUT provided provided level of thread support (integer)

int MPI_Query_thread(int *provided)

MPI_QUERY_THREAD(PROVIDED, IERROR)

INTEGER PROVIDED, IERROR

{int MPI::Query_thread() (*binding deprecated, see Section 15.2*) }

The call returns in provided the current level of thread [support. This]support, which will be the value returned in provided by MPI_INIT_THREAD, if MPI was initialized by a call to MPI_INIT_THREAD().

MPI_IS_THREAD_MAIN(flag)

OUT flag true if calling thread is main thread, false otherwise (logical)

int MPI_Is_thread_main(int *flag)

MPI_IS_THREAD_MAIN(FLAG, IERROR)

LOGICAL FLAG

INTEGER IERROR

{bool MPI::Is_thread_main() (*binding deprecated, see Section 15.2*) }

This function can be called by a thread to [find out whether]determine if it is the main thread (the thread that called MPI_INIT or MPI_INIT_THREAD).

All routines listed in this section must be supported by all MPI implementations.

Rationale. MPI libraries are required to provide these calls even if they do not support threads, so that portable code that contains invocations to these functions [be able to]can link correctly. MPI_INIT continues to be supported so as to provide compatibility with current MPI codes. (*End of rationale.*)

Advice to users. It is possible to spawn threads before MPI is initialized, but no MPI call other than [MPI_INITIALIZED] MPI_GET_VERSION, MPI_INITIALIZED, or MPI_FINALIZED should be executed by these threads, until MPI_INIT_THREAD is invoked by one thread (which, thereby, becomes the main thread). In particular, it is possible to enter the MPI execution with a multi-threaded process.

The level of thread support provided is a global property of the MPI process that can be specified only once, when MPI is initialized on that process (or before). Portable third party libraries have to be written so as to accommodate any provided level of

1 thread support. Otherwise, their usage will be restricted to specific level(s) of thread
2 support. If such a library can run only with specific level(s) of thread support, e.g.,
3 only with `MPI_THREAD_MULTIPLE`, then `MPI_QUERY_THREAD` can be used to check
4 whether the user initialized MPI to the correct level of thread support and, if not,
5 raise an exception. (*End of advice to users.*)
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48

Predefined Attribute Keys

C type: <code>const int</code> (or unnamed <code>enum</code>)	C++ type:
Fortran type: <code>INTEGER</code>	<code>const int</code> (or unnamed <code>enum</code>)
<code>MPI_APPNUM</code>	<code>MPI::APPNUM</code>
<code>MPI_LASTUSED</code>	<code>MPI::LASTUSED</code>
<code>MPI_UNIVERSE_SIZE</code>	<code>MPI::UNIVERSE_SIZE</code>
<code>MPI_WIN_BASE</code>	<code>MPI::WIN_BASE</code>
<code>MPI_WIN_DISP_UNIT</code>	<code>MPI::WIN_DISP_UNIT</code>
<code>MPI_WIN_SIZE</code>	<code>MPI::WIN_SIZE</code>

`MPI_ENV_N`
`MPI_ENV_THREAD_LEVEL`
`MPI_ENV_HOST`
`MPI_ENV_ARCH`
`MPI_ENV_WDIR`
`MPI_ENV_FILE`

Mode Constants

C type: <code>const int</code> (or unnamed <code>enum</code>)	C++ type:
Fortran type: <code>INTEGER</code>	<code>const int</code> (or unnamed <code>enum</code>)
<code>MPI_MODE_APPEND</code>	<code>MPI::MODE_APPEND</code>
<code>MPI_MODE_CREATE</code>	<code>MPI::MODE_CREATE</code>
<code>MPI_MODE_DELETE_ON_CLOSE</code>	<code>MPI::MODE_DELETE_ON_CLOSE</code>
<code>MPI_MODE_EXCL</code>	<code>MPI::MODE_EXCL</code>
<code>MPI_MODE_NOCHECK</code>	<code>MPI::MODE_NOCHECK</code>
<code>MPI_MODE_NOPRECEDE</code>	<code>MPI::MODE_NOPRECEDE</code>
<code>MPI_MODE_NOPUT</code>	<code>MPI::MODE_NOPUT</code>
<code>MPI_MODE_NOSTORE</code>	<code>MPI::MODE_NOSTORE</code>
<code>MPI_MODE_NOSUCCEED</code>	<code>MPI::MODE_NOSUCCEED</code>
<code>MPI_MODE_RDONLY</code>	<code>MPI::MODE_RDONLY</code>
<code>MPI_MODE_RDWR</code>	<code>MPI::MODE_RDWR</code>
<code>MPI_MODE_SEQUENTIAL</code>	<code>MPI::MODE_SEQUENTIAL</code>
<code>MPI_MODE_UNIQUE_OPEN</code>	<code>MPI::MODE_UNIQUE_OPEN</code>
<code>MPI_MODE_WRONLY</code>	<code>MPI::MODE_WRONLY</code>