

D R A F T

Document for a Standard Message-Passing Interface

Message Passing Interface Forum

September 15, 2014

This work was supported in part by NSF and ARPA under NSF contract CDA-9115428 and Esprit under project HPC Standards (21111).

This is the result of a LaTeX run of a draft of a single chapter of the MPIF Final Report document.

Chapter 15

Process Fault Tolerance

15.1 Introduction

In distributed systems with numerous or complex components, a serious risk is that a component fault manifests as a process failure that disrupts the normal execution of a long running application. A process failure is a common outcome for many hardware, network, or software faults that cause a process to crash; it can be more formally defined as a fail-stop failure: the affected process stops communicating permanently. This chapter introduces MPI features that support the development of applications, libraries, and programming languages that can tolerate process failures. The primary goal is to specify error classes and interfaces that permit users to continue simple MPI communication operations after failures have impacted the execution and rebuild MPI objects (communicators, files, etc.) as needed to restore the full capability of MPI to carry out elaborate communication operations (like collective communications). This specification does not include mechanisms to restore the data lost due to process failures. The literature is rich with diverse fault tolerance techniques that the users may employ at their discretion, including checkpoint-restart, algorithmic dataset recovery, and continuation ignoring failed processes. All these fault tolerance approaches benefit from, and often require, the definitions and interfaces specified in this chapter in order to resume communicating after a failure.

The expected behavior of MPI in the case of a process failure is defined by the following statements: any MPI operation that involves a failed process must not block indefinitely but either succeed or raise an MPI exception (see Section 15.2); an MPI operation that does not involve a failed process will complete normally, unless interrupted by the user through provided functionality. Exceptions indicate only the local impact of the failure on an operation, and make no guarantee that other processes have also been notified of the same failure. Asynchronous failure propagation is not guaranteed or required, and users must exercise caution when determining the set of ranks where a failure has been detected and raised an exception. If an application needs global knowledge of failures, it can use the interfaces defined in Section 15.3 to explicitly propagate the notification of locally detected failures.

The typical usage pattern on some reliable machines may not require fault tolerance. An MPI implementation that does not tolerate process failures must never raise a process failure exception (as listed in Section 15.4). Fault-tolerant applications using the interfaces defined in this chapter must be portable across MPI implementations (including these which do not provide resilience, but in this case the interfaces may exhibit undefined be-

havior after a process failure at any rank.) Fault tolerant applications may determine if the implementation supports fault tolerance by querying the predefined attribute `MPI_FT` on `MPI_COMM_WORLD` (see 8.1.2.)

Advice to users. Many of the operations and semantics described in this chapter are applicable only when the MPI application has replaced the default error handler `MPI_ERRORS_ARE_FATAL` on, at least, `MPI_COMM_WORLD`. (*End of advice to users.*)

15.2 Failure Notification

This section specifies the behavior of an MPI communication operation when failures occur on processes involved in the communication. A process is considered involved in a communication (for the purpose of this chapter) if any of the following is true:

- The operation is collective, and the process appears in one of the groups of the associated communication object.
- The process is a specified or matched destination or source in a point-to-point communication.
- The operation is an `MPI_ANY_SOURCE` receive operation and the failed process belongs to the source group.
- The process is a specified target in a remote memory operation.

An operation involving a failed process must always complete in a finite amount of time (possibly by raising one of the process failure exceptions listed in Section 15.4). If an operation does not involve a failed process (such as a point-to-point message between two non-failed processes), it must not raise a process failure exception.

Advice to implementors. A correct MPI implementation may provide failure detection only for processes involved in an ongoing operation and may postpone detection of other failures until necessary. Moreover, as long as an implementation can complete operations, it may choose to delay raising an exception. Another valid implementation might choose to raise an exception as quickly as possible. (*End of advice to implementors.*)

When a communication operation raises a process failure exception, it may not satisfy its specification, (for example, a synchronizing operation may not have synchronized) and the content of the output buffers, targeted memory, or output parameters is *undefined*. Exceptions to this rule are explicitly stated in the remainder of this chapter.

Non-blocking operations must not raise an exception about process failures during initiation. All process failure exception raising is postponed until the corresponding completion function is called.

15.2.1 Startup and Finalize

Initialization does not have any new semantics related to fault tolerance.

Advice to implementors. If a process fails during `MPI_INIT` but its peers are able to complete the `MPI_INIT` successfully, then a high quality implementation will return `MPI_SUCCESS` and delay the reporting of the process failure to a subsequent MPI operation. (*End of advice to implementors.*)

`MPI_FINALIZE` will complete even in the presence of process failures. If process 0 in `MPI_COMM_WORLD` has failed, it is possible that no process returns from `MPI_FINALIZE`.

Advice to users. Applications are encouraged to implement all rank-specific code before the call to `MPI_FINALIZE`. In Example 8.10 in Section 8.7, the process with rank 0 in `MPI_COMM_WORLD` may have failed before, during, or after the call to `MPI_FINALIZE`, possibly leading to this code never being executed. (*End of advice to users.*)

15.2.2 Point-to-Point and Collective Communication

An MPI implementation raises exceptions of the following error classes in order to notify users that a point-to-point communication operation could not complete successfully because of the failure of involved processes:

- `MPI_ERR_PROC_FAILED_PENDING` indicates, for a non-blocking communication, that the communication is a receive operation from `MPI_ANY_SOURCE` and no send operation has matched, yet a potential sending process has failed. Neither the operation nor the request identifying the operation is completed.
- In all other cases, the operation raises an exception of class `MPI_ERR_PROC_FAILED` to indicate that the failure prevents the operation from following its failure-free specification. If there is a request identifying a point-to-point communication, it is completed. Future communication involving the failed process on this communicator must also raise `MPI_ERR_PROC_FAILED`.

When a collective operation cannot be completed because of the failure of an involved process, the collective operation raises an exception of class `MPI_ERR_PROC_FAILED`.

Advice to users.

Depending on how the collective operation is implemented and when a process failure occurs, some participating alive processes may raise an exception while other processes return successfully from the same collective operation. For example, in `MPI_BCAST`, the root process may succeed before a failed process disrupts the operation, resulting in some other processes raising an exception.

Note, however, for some operations' semantics, when a process fails before entering the operation, it forces raising an exception at all ranks. As an example, if an operation on an intracommunicator has raised an exception, the process receiving that exception can then assume that in a subsequent `MPI_BARRIER` on this communicator, all ranks will raise an exception `MPI_ERR_PROC_FAILED` because the participating process is known to have failed before entering the barrier.

(*End of advice to users.*)

Advice to users.

Note that communicator creation functions (e.g., `MPI_COMM_DUP` or `MPI_COMM_SPLIT`) are collective operations. As such, if a failure happened during the call, an exception might be raised at some processes while others succeed and obtain a new communicator. Although it is valid to communicate between processes that succeeded in creating the new communicator, the user is responsible for ensuring a consistent view of the communicator creation, if needed. A conservative solution is to check the global outcome of the communicator creation function with `MPI_COMM_AGREE` (defined in Section 15.3.1), as illustrated in Example 15.1. (*End of advice to users.*)

After a process failure, `MPI_COMM_FREE` (as with all other collective operations) may not complete successfully at all ranks. For any rank that receives the return code `MPI_SUCCESS`, the behavior is defined in Section 6.4.3. If a rank raises a process failure exception (`MPI_ERR_PROC_FAILED` or `MPI_ERR_REVOKED`), the communicator handle `comm` is set to `MPI_COMM_NULL`; however, the implementation makes no guarantee about the success or failure of the `MPI_COMM_FREE` operation, locally or remotely.

Advice to users. Users are encouraged to call `MPI_COMM_FREE` on communicators they do not wish to use anymore, even when they contain failed processes. Although the operation may raise a process failure exception and not synchronize properly, this gives a high quality implementation an opportunity to release local resources and memory consumed by the object. (*End of advice to users.*)

15.2.3 Dynamic Process Management

Dynamic process management functions require some additional semantics from the MPI implementation as detailed below.

- If the MPI implementation raises a process failure exception at the root process of `MPI_COMM_CONNECT` or `MPI_COMM_ACCEPT`, at least the root processes of both intracommunicators must raise the same exception of class `MPI_ERR_PROC_FAILED` (unless required to raise `MPI_ERR_REVOKED` as defined in Section 15.3.1). The same is true if the implementation raises an exception at any process in `MPI_COMM_JOIN`.
- If the MPI implementation raises a process failure exception to the root process of `MPI_COMM_SPAWN` or `MPI_COMM_SPAWN_MULTIPLE`, no spawned processes will be able to communicate on the created intercommunicator.

Advice to users. As with communicator creation functions, if a failure happens during dynamic process management operations, an exception might be raised at some processes while others succeed and obtain a new communicator. (*End of advice to users.*)

15.2.4 One-Sided Communication

When an operation on a window raises an exception related to process failure, the state of all data held in memory exposed by that window becomes undefined.

Advice to users. A high quality implementation may be able to limit the scope of the exposed memory that becomes undefined (as an example, only the memory that has been targeted by remote writes, or origin in remote reads). Note that in passive target mode, when an exception is raised at the origin, the target memory may become undefined before a synchronization raises an exception at the target. Assessing if a particular portion of the exposed memory remains correct is the responsibility of the user. (*End of advice to users.*)

One-sided communication operations whose outputs are undefined due to a process failure (see Section 15.2) are not required to raise a process failure exception. The corresponding epoch closing synchronization must always raise a process failure exception at the ranks where the undefined behavior can be observed.

Advice to users. As with collective operations over MPI communicators, one-sided synchronization operations may raise a process failure exception at some rank while the corresponding operation returned MPI_SUCCESS at some other rank. (*End of advice to users.*)

Passive target synchronization operations may raise a process failure exception when any rank in the window has failed (including non-target ranks).

Rationale. Even when a target argument is specified, a typical implementation of passive target synchronization may need to synchronize with other processes in the window, as an example, a previous owner of an access epoch on the target window. (*End of rationale.*)

After a process failure, MPI_WIN_FREE (as with all other collective operations) may not complete successfully at all ranks. For any rank that receives the return code MPI_SUCCESS, the behavior is defined in Section 11.2.5. If a rank raises a process failure exception (MPI_ERR_PROC_FAILED or MPI_ERR_REVOKED), the window handle win is set to MPI_WIN_NULL; however, the implementation makes no guarantee about the success or failure of the MPI_WIN_FREE operation, locally or remotely.

Advice to users. Users are encouraged to call MPI_WIN_FREE on windows they do not wish to use anymore, even when they contain failed processes. Although the operation may raise a process failure exception and not synchronize properly, this gives a high quality implementation an opportunity to release local resources and memory consumed by the object. Before calling MPI_WIN_FREE, it may be required to call MPI_WIN_REVOKE to close an epoch that couldn't be completed as a consequence of a process failure (see Section 15.3.2). (*End of advice to users.*)

15.2.5 I/O

This section defines the behavior of I/O operations when MPI process failures prevent their successful completion. I/O backend failure error classes and their consequences are defined in Section 13.7.

If a process failure prevents a file operation from completing, an MPI exception of class MPI_ERR_PROC_FAILED is raised. Once an MPI implementation has raised an exception of class MPI_ERR_PROC_FAILED, the state of the file pointers involved in the operation that raised the exception is *undefined*.

Advice to users. Since collective I/O operations may not synchronize with other processes, process failures may not be reported during a collective I/O operation. Users are encouraged to use `MPI_COMM_AGREE` on a communicator containing the same group as the file handle when they need to deduce the completion status of collective operations on file handles and maintain a consistent view of file pointers. The file pointer can be reset by using `MPI_FILE_SEEK` with the `MPI_SEEK_SET` update mode. (*End of advice to users.*)

After a process failure, `MPI_FILE_CLOSE` (as with all other collective operations) may not complete successfully at all ranks. For any rank that receives the return code `MPI_SUCCESS`, the behavior is defined in Section 13.2.2. If a rank raises a process failure exception (`MPI_ERR_PROC_FAILED` or `MPI_ERR_REVOKED`), the file handle `fh` is set to `MPI_FILE_NULL`; however, the implementation makes no guarantee about the success or failure of the `MPI_FILE_CLOSE` operation, locally or remotely.

Advice to users. Users are encouraged to call `MPI_FILE_CLOSE` on files they do not wish to use anymore, even when they contain failed processes. Although the operation may raise a process failure exception and not synchronize properly, this gives a high quality implementation an opportunity to release local resources and memory consumed by the object. (*End of advice to users.*)

15.3 Failure Mitigation Functions

15.3.1 Communicator Functions

MPI provides no guarantee of global knowledge of a process failure. Only processes involved in a communication operation with the failed process are guaranteed to eventually detect its failure (see Section 15.2). If global knowledge is required, MPI provides a function to revoke a communicator at all members.

`MPI_COMM_REVOKE(comm)`

IN `comm` communicator (handle)

`int MPI_Comm_revoke(MPI_Comm comm)`

`MPI_Comm_revoke(comm, ierror)`

TYPE(MPI_Comm), INTENT(IN) :: `comm`

INTEGER, OPTIONAL, INTENT(OUT) :: `ierror`

`MPI_COMM_REVOKE(COMM, IERROR)`

INTEGER `COMM`, `IERROR`

This function notifies all processes in the groups (local and remote) associated with the communicator `comm` that this communicator is now considered revoked. This function is not collective and therefore does not have a matching call on remote processes. All alive processes belonging to `comm` will be notified of the revocation despite failures. The revocation of a communicator completes any non-local MPI operations on `comm` by raising an exception of class `MPI_ERR_REVOKED`, with the exception of `MPI_COMM_SHRINK`,

MPI_COMM_AGREE, and MPI_COMM_IAGREE. A communicator becomes revoked at a process by the following:

- MPI_COMM_REVOKE is called locally,
- or any MPI operation raises an exception of class MPI_ERR_REVOKED because another process in `comm` has called MPI_COMM_REVOKE.

Once a communicator has been revoked, all subsequent non-local operations on that communicator, with the exception of MPI_COMM_SHRINK, MPI_COMM_AGREE, and MPI_COMM_IAGREE are considered local and must complete by raising an exception of class MPI_ERR_REVOKED.

MPI_COMM_SHRINK(`comm`, `newcomm`)

IN	<code>comm</code>	communicator (handle)
OUT	<code>newcomm</code>	communicator (handle)

int MPI_Comm_shrink(MPI_Comm `comm`, MPI_Comm* `newcomm`)

```

MPI_Comm_shrink(comm, newcomm, ierror)
    TYPE(MPI_Comm), INTENT(IN) :: comm
    TYPE(MPI_Comm), INTENT(OUT) :: newcomm
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

```

MPI_COMM_SHRINK(COMM, NEWCOMM, IERROR)
    INTEGER COMM, NEWCOMM, IERROR

```

This collective operation creates a new intra- or intercommunicator `newcomm` from the intra- or intercommunicator `comm`, respectively, by excluding the group of failed processes as agreed upon during the operation. This call is semantically equivalent to an MPI_COMM_SPLIT operation that would succeed despite failures, and where non-failed processes participate with the same color and a key equal to their rank in `comm`, and failed processes implicitly contribute MPI_UNDEFINED. At least every process whose failure raised an MPI exception of class MPI_ERR_PROC_FAILED or MPI_ERR_PROC_FAILED_PENDING on `comm` must be excluded. It is valid to call MPI_COMM_SHRINK on a communicator even after it has been revoked. This function never raises an exception of class MPI_ERR_PROC_FAILED or MPI_ERR_REVOKED.

Advice to users. MPI_COMM_SHRINK maintains its collective behavior even if `comm` is revoked.

This call does not guarantee that all processes in `newcomm` are alive. Any new failure will be detected in subsequent MPI operations. (*End of advice to users.*)

MPI_COMM_FAILURE_ACK(`comm`)

IN	<code>comm</code>	communicator (handle)
----	-------------------	-----------------------

```

1  int MPI_Comm_failure_ack(MPI_Comm comm)
2
3  MPI_Comm_failure_ack(comm, ierror)
4      TYPE(MPI_Comm), INTENT(IN) :: comm
5      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
6
7  MPI_COMM_FAILURE_ACK(COMM, IERROR)
8      INTEGER COMM, IERROR

```

This local operation gives the users a way to *acknowledge* all locally notified failures on `comm`. After the call, unmatched `MPI_ANY_SOURCE` receive operations that would have raised an exception `MPI_ERR_PROC_FAILED_PENDING` due to process failure (see Section 15.2.2) proceed without further raising exceptions due to those acknowledged failures. Also after this call, `MPI_COMM_AGREE` will not raise `MPI_ERR_PROC_FAILED` due to previously acknowledged failures (according to the specification found later in this section).

Advice to users.

Calling `MPI_COMM_FAILURE_ACK` on a communicator with failed processes has no effect on collective operations (except for `MPI_COMM_AGREE`). If a collective operation would raise an exception due to the communicator containing a failed process (as defined in Section 15.2.2), it can continue to raise an exception even after the failure has been acknowledged. In order to use collective operations between processes of a communicator that contains failed processes, users should create a new communicator by calling `MPI_COMM_SHRINK`.

(End of advice to users.)

```

26
27  MPI_COMM_FAILURE_GET_ACKED( comm, failedgrp )
28      IN          comm          communicator (handle)
29      OUT         failedgrp      group of failed processes (handle)
30
31
32  int MPI_Comm_failure_get_acked(MPI_Comm comm, MPI_Group* failedgrp)
33
34  MPI_Comm_failure_get_acked(comm, failedgrp, ierror)
35      TYPE(MPI_Comm), INTENT(IN) :: comm
36      TYPE(MPI_Group), INTENT(OUT) :: failedgrp
37      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
38
39  MPI_COMM_FAILURE_GET_ACKED(COMM, FAILEDGRP, IERROR)
40      INTEGER COMM, FAILEDGRP, IERROR

```

This local operation returns the group `failedgrp` of processes, from the communicator `comm`, that have been locally acknowledged as failed by preceding calls to `MPI_COMM_FAILURE_ACK`. The *failedgrp* can be empty, that is, equal to `MPI_GROUP_EMPTY`.

Advice to users. When they are not separated by a call to `MPI_COMM_FAILURE_ACK`, multiple calls to `MPI_COMM_FAILURE_GET_ACKED` produce similar `failedgrp` groups; that is, the result when providing these groups to `MPI_GROUP_DIFFERENCE` is `MPI_EMPTY`. *(End of advice to users.)*

`MPI_COMM_AGREE(comm, flag)`

IN `comm` communicator (handle)
 INOUT `flag` integer flag

`int MPI_Comm_agree(MPI_Comm comm, int* flag)`

`MPI_Comm_agree(comm, flag, ierror)`
 `TYPE(MPI_Comm), INTENT(IN) :: comm`
 `INTEGER, INTENT(INOUT) :: flag`
 `INTEGER, OPTIONAL, INTENT(OUT) :: ierror`

`MPI_COMM_AGREE(COMM, FLAG, IERROR)`
 `INTEGER COMM, FLAG, IERROR`

This function performs a collective operation on the group of living processes in `comm`. The purpose of this function is to agree on the integer value `flag` and on the group of failed participants in `comm`. When an exception of class `MPI_ERR_PROC_FAILED` is raised, it is consistently raised at all participating ranks; conversely, when `MPI_SUCCESS` is returned, it is consistently returned at all participating ranks.

On completion, all living processes agree to set the output integer value of `flag` to the result of a bitwise ‘AND’ operation over the contributed input values of `flag`. If `comm` is an intercommunicator, the value of `flag` is a bitwise ‘AND’ operation over the values contributed by the remote group.

When a process fails before contributing to the operation, the `flag` is computed ignoring its contribution, and `MPI_COMM_AGREE` raises an exception of class `MPI_ERR_PROC_FAILED`. This exception is raised consistently at all participating ranks (on both the local and remote groups of `comm`). However, if all participants have acknowledged this failure prior to the call to `MPI_COMM_AGREE` (using `MPI_COMM_FAILURE_ACK`), the exception related to this failure is not raised (an exception may still be raised due to other, non-acknowledged failures).

After `MPI_COMM_AGREE` raised an exception of class `MPI_ERR_PROC_FAILED`, a subsequent call to `MPI_COMM_FAILURE_ACK` on `comm` acknowledges (at least) the failure of every process that didn’t contributed to the computation of `flag`.

Rationale. When `MPI_COMM_AGREE` returns `MPI_SUCCESS`, the only ignored contributions are from processes whose failure has been previously acknowledged by `MPI_COMM_FAILURE_ACK` at all ranks.

Using a combination of `MPI_COMM_FAILURE_ACK` and `MPI_COMM_AGREE` as illustrated in Example 15.3, users can propagate and synchronize the knowledge of failures across all ranks in `comm`. (*End of rationale.*)

This function never raises an exception of class `MPI_ERR_REVOKED`.

Advice to users. `MPI_COMM_AGREE` maintains its collective behavior even if the `comm` is revoked. (*End of advice to users.*)

```
1 MPI_COMM_IAGREE( comm, flag, req )
```

```
2     IN          comm          communicator (handle)
```

```
3     INOUT      flag          integer flag
```

```
4     OUT        req          request (handle)
```

```
7 int MPI_Comm_iagree(MPI_Comm comm, int* flag, MPI_Request* req)
```

```
9 MPI_Comm_iagree(comm, flag, req, ierror)
```

```
10     TYPE(MPI_Comm), INTENT(IN) :: comm
```

```
11     INTEGER, INTENT(INOUT) :: flag
```

```
12     TYPE(MPI_Request), INTENT(OUT) :: req
```

```
13     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

```
14 MPI_COMM_IAGREE(COMM, FLAG, REQ, IERROR)
```

```
15     INTEGER COMM, FLAG, REQ, IERROR
```

```
17     This function has the same semantics as MPI_COMM_AGREE except that it is non-
18 blocking.
```

15.3.2 One-Sided Functions

```
23 MPI_WIN_REVOKE( win )
```

```
24     IN          win          window (handle)
```

```
27 int MPI_Win_revoke(MPI_Win win)
```

```
28 MPI_Win_revoke(win, ierror)
```

```
29     TYPE(MPI_Win), INTENT(IN) :: win
```

```
30     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

```
32 MPI_WIN_REVOKE(WIN, IERROR)
```

```
33     INTEGER WIN, IERROR
```

```
34     This function notifies all processes within the window win that this window is now con-
35 sidered revoked. This function is not collective and therefore does not have a matching call
36 on remote processes. All alive processes belonging to win will be notified of the revocation
37 despite failures. The revocation of a window completes any non-local MPI operations on
38 win by raising an exception of class MPI_ERR_REVOKED. Once a window has been revoked,
39 the current epoch is closed (pending operations posted during the epoch are interrupted
40 and completed with undefined outputs), and all subsequent non-local operations on that
41 window are considered local and must raise an exception of class MPI_ERR_REVOKED. A
42 window becomes revoked as soon as either of the following occur:
```

- 44 • MPI_WIN_REVOKE is locally called on it,
- 46 • or any MPI operation raises an exception of class MPI_ERR_REVOKED because another
- 47 process in win has called MPI_WIN_REVOKE.

MPI_WIN_GET_FAILED(win, failedgrp)

IN win window (handle)
OUT failedgrp group of failed processes (handle)

int MPI_Win_get_failed(MPI_Win win, MPI_Group* failedgrp)

MPI_Win_get_failed(win, failedgrp, ierror)
 TYPE(MPI_Win), INTENT(IN) :: win
 TYPE(MPI_Group), INTENT(OUT) :: failedgrp
 INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_WIN_GET_FAILED(WIN, FAILEDGRP, IERROR)
 INTEGER COMM, FAILEDGRP, IERROR

This local operation returns the group `failedgrp` of processes from the window `win` that are locally known to have failed.

Advice to users. MPI makes no assumption about asynchronous progress of the failure detection. A valid MPI implementation may choose to update only the group of locally known failed processes when it enters a synchronization function and must raise a process failure exception. (*End of advice to users.*)

Advice to users. It is possible that only the calling process has detected the reported failure. If global knowledge is necessary, processes detecting failures should use the call `MPI_WIN_REVOKED`. (*End of advice to users.*)

15.3.3 I/O Functions

MPI_FILE_REVOKE(fh)

IN fh file (handle)

int MPI_File_revoke(MPI_File fh)

MPI_File_revoke(fh, ierror)
 TYPE(MPI_File), INTENT(IN) :: fh
 INTEGER, OPTIONAL, INTENT(OUT) :: ierror

MPI_FILE_REVOKE(FH, IERROR)
 INTEGER FH, IERROR

This function notifies all processes within the file handle `fh` that this file handle is now considered revoked. This function is not collective and therefore does not have a matching call on remote processes. All alive processes belonging to the file handle `fh` will be notified of the revocation despite failures. The revocation of a file handle completes any non-local MPI operations on `fh` by raising an exception of class `MPI_ERR_REVOKED`. Once a file handle has been revoked, all subsequent non-local operations on that file handle are considered local and must raise an exception of class `MPI_ERR_REVOKED`. A file handle becomes revoked as soon as either of the following occur:

- `MPI_FILE_REVOKE` is locally called on it,
- or any MPI operation raises an exception of class `MPI_ERR_REVOKED` because another process in `fh` has called `MPI_FILE_REVOKE`.

15.4 Process Failure Error Codes and Classes

The following process failure error classes are added to those defined in Section 8.4:

<code>MPI_ERR_PROC_FAILED</code>	The operation could not complete because of a process failure (a fail-stop failure).
<code>MPI_ERR_PROC_FAILED_PENDING</code>	The operation was interrupted by a process failure (a fail-stop failure). The request is still pending and the operation may be completed later.
<code>MPI_ERR_REVOKED</code>	The communication object used in the operation has been revoked.

Table 15.1: Additional process fault tolerance error classes

15.5 Examples

15.5.1 Safe Communicator Creation

The example below illustrates how a new communicator can be safely created despite disruption by process failures. A child communicator is created with `MPI_COMM_SPLIT`, then the global success of the operation is verified with `MPI_COMM_AGREE`. If any process failed to create the child communicator, all processes are notified by the value of the integer agreed on. Processes that had successfully created the child communicator destroy it, as it cannot be used consistently.

Example 15.1 Fault Tolerant Communicator Split Example

```

int Comm_split_consistent(MPI_Comm parent, int color, int key, MPI_Comm* child)
{
    rc = MPI_Comm_split(parent, color, key, child);
    split_ok = (MPI_SUCCESS == rc);
    rc = MPI_Comm_agree(parent, &split_ok);
    if(split_ok && (MPI_SUCCESS == rc) ) {
        /* All surviving processes have created the "child" comm
         * It may contain supplementary failures and the first
         * operation on it may raise an exception, but it is a
         * workable object that will yield well specified outcomes */
        return MPI_SUCCESS;
    }
    else {
        /* At least one process did not create the child comm properly

```

```

        * if the local rank did succeed in creating it, it disposes
        * of it, as it is a broken, inconsistent object */
    if(MPI_SUCCESS == rc) {
        MPI_Comm_free(child);
    }
    return MPI_ERR_PROC_FAILED;
}
}

```

15.5.2 Obtaining the consistent group of failed processes

Users can invoke `MPI_COMM_FAILURE_ACK`, `MPI_COMM_FAILURE_GET_ACKED`, `MPI_WIN_GET_FAILED`, to obtain the group of failed processes, as detected at the local rank. However, these operations are local, thereby the invocation of the same function at another rank can result in a different group of failed processes being returned.

In the following examples, we illustrate two different approaches that permit obtaining the consistent group of failed processes accross all ranks of a communicator. The first one employs `MPI_COMM_SHRINK` to create a temporary communicator were all alive processes are agreed on. The second one employs `MPI_COMM_AGREE` to synchronize the set of acknowledged failures.

Example 15.2 Fault-Tolerant Consistent Group of Failures Example (Shrink variant)

```

Comm_failure_allget(MPI_Comm c, MPI_Group * g) {
    MPI_Comm s; MPI_Group c_grp, s_grp;

    /* Using shrink to create a new communicator, the underlying
     * group is necessarily consistent across all ranks, and excludes
     * all processes detected to have failed before the call */
    MPI_Comm_shrink(c, &s);
    /* Extracting the groups from the communicators */
    MPI_Comm_group(c, &c_grp);
    MPI_Comm_group(s, &s_grp);
    /* s_grp is the group of still alive processes, we want to
     * return the group of failed processes. */
    MPI_Group_diff(c_grp, s_grp, g);

    MPI_Group_free(&c_grp); MPI_Group_free(&s_grp);
    MPI_Comm_free(&s);
}

```

Example 15.3 Fault-Tolerant Consistent Group of Failures Example (Agree variant)

```

Comm_failure_allget2(MPI_Comm c, MPI_Group * g) {
    int rc; int T=1;

    do {
        /* this routine is not pure: calling MPI_Comm_failure_ack
         * affects the state of the communicator c */

```

```

1      MPI_Comm_failure_ack(comm);
2      /* we simply ignore the value in this example */
3      rc = MPI_Comm_agree(comm, &T);
4  } while( rc != MPI_SUCCESS );
5  /* after this loop, MPI_Comm_agree has returned MPI_SUCCESS at
6   * all ranks, so all ranks have Acknowledged the same set of
7   * failures. Let's get that set of failures in the g group. */
8  MPI_Comm_failure_get_acked(comm, g);
9  }

```

15.5.3 Fault-Tolerant Master/Worker

The example below presents a master code that handles worker failures by discarding failed worker processes and resubmitting the work to the remaining workers. It demonstrates the different failure cases that may occur when posting receptions from `MPI_ANY_SOURCE` as discussed in the advice to users in Section 15.2.2.

Example 15.4 Fault-Tolerant Master Example

```

19 int master(void)
20 {
21     MPI_Comm_set_errhandler(comm, MPI_ERRORS_RETURN);
22     MPI_Comm_size(comm, &size);
23
24     /* ... submit the initial work requests ... */
25
26     /* Progress engine: Get answers, send new requests,
27      and handle process failures */
28     MPI_Irecv( buffer, 1, MPI_INT, MPI_ANY_SOURCE, tag, comm, &req );
29     while( (active_workers > 0) && work_available ) {
30         rc = MPI_Wait( &req, &status );
31         if( MPI_SUCCESS == rc ) {
32             /* ... process the answer and update work_available ... */
33         }
34         else {
35             MPI_Error_class(rc, &ec);
36             if( (MPI_ERR_PROC_FAILED == ec) ||
37                 (MPI_ERR_PROC_FAILED_PENDING == ec) ) {
38                 MPI_Comm_failure_ack(comm);
39                 MPI_Comm_failure_get_acked(comm, &g);
40                 MPI_Group_size(g, &gsize);
41
42                 /* ... find the lost work and requeue it ... */
43
44                 active_workers = size - gsize - 1;
45                 MPI_Group_free(&g);
46
47                 /* no need to repost when the request is still pending */
48                 if( ec == MPI_ERR_PROC_FAILED_PENDING )

```



```

        continue;
    }
}
/* get ready to recieve more notifications from workers */
MPI_Irecv( buffer, 1, MPI_INT, MPI_ANY_SOURCE, tag, comm, &req );
}
/* ... cancel request and cleanup ... */
}

```

15.5.4 Fault-Tolerant Iterative Refinement

The example below demonstrates a method of fault tolerance for detecting and handling failures. At each iteration, the algorithm checks the return code of the `MPI_ALLREDUCE`. If the return code indicates a process failure for at least one process, the algorithm revokes the communicator, agrees on the presence of failures, and shrinks it to create a new communicator. By calling `MPI_COMM_REVOKE`, the algorithm ensures that all processes will be notified of process failure and enter the `MPI_COMM_AGREE`. If a process fails, the algorithm must complete at least one more iteration to ensure a correct answer.

Example 15.5 Fault-tolerant iterative refinement with shrink and agreement

```

while( gnorm > epsilon ) {
    /* Add a computation iteration to converge and
       compute local norm in lnorm */
    rc = MPI_Allreduce(&lnorm, &gnorm, 1, MPI_DOUBLE, MPI_MAX, comm);
    ec = MPI_Error_class(rc, &ec);

    if( (MPI_ERR_PROC_FAILED == ec) ||
        (MPI_ERR_REVOKED == ec) ||
        (gnorm <= epsilon) ) {

        /* This rank detected a failure, but other ranks may have
           * proceeded into the next MPI_Allreduce. Since this rank
           * will not match that following MPI_Allreduce, these other
           * ranks would be at risk of deadlocking. This process thus
           * calls MPI_Comm_revoke to interrupt other ranks and notify
           * them that it has detected a failure and is leaving the
           * failure free execution path to go into recovery. */
        if( MPI_ERR_PROC_FAILED == ec )
            MPI_Comm_revoke(comm);

        /* About to leave: let's be sure that everybody
           received the same information */
        allsucceeded = (rc == MPI_SUCCESS);
        rc = MPI_Comm_agree(comm, &allsucceeded);
        MPI_Error_class(rc, &ec);
        if( ec == MPI_ERR_PROC_FAILED || !allsucceeded ) {
            MPI_Comm_shrink(comm, &comm2);
        }
    }
}

```

```
1      MPI_Comm_free(comm); /* Release the revoked communicator */
2      comm = comm2;
3      gnorm = epsilon + 1.0; /* Force one more iteration */
4  }
5      }
6  }
```

7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48

Index

CONST:MPI_ANY_SOURCE, [2](#), [3](#), [8](#), [14](#)
 CONST:MPI_Comm, [6–10](#)
 CONST:MPI_COMM_NULL, [4](#)
 CONST:MPI_COMM_WORLD, [2](#), [3](#)
 CONST:MPI_EMPTY, [8](#)
 CONST:MPI_ERR_PROC_FAILED, [3–9](#), [12](#)
 CONST:MPI_ERR_PROC_FAILED_PENDING, [3](#), [7](#), [8](#), [12](#)
 CONST:MPI_ERR_REVOKED, [4–7](#), [9–12](#)
 CONST:MPI_ERRORS_ARE_FATAL, [2](#)
 CONST:MPI_File, [11](#)
 CONST:MPI_FILE_NULL, [6](#)
 CONST:MPI_FT, [2](#)
 CONST:MPI_Group, [8](#), [11](#)
 CONST:MPI_GROUP_EMPTY, [8](#)
 CONST:MPI_Request, [10](#)
 CONST:MPI_SEEK_SET, [6](#)
 CONST:MPI_SUCCESS, [3–6](#), [9](#)
 CONST:MPI_UNDEFINED, [7](#)
 CONST:MPI_Win, [10](#), [11](#)
 CONST:MPI_WIN_NULL, [5](#)

 EXAMPLES:Comm_failure_allget example, [13](#)
 EXAMPLES:Comm_failure_allget2 example, [13](#)
 EXAMPLES:Fault-tolerant iterative refinement with shrink and agreement, [15](#)
 EXAMPLES:Master example, [14](#)
 EXAMPLES:MPI_COMM_AGREE, [12](#), [13](#), [15](#)
 EXAMPLES:MPI_COMM_FAILURE_ACK, [13](#), [14](#)
 EXAMPLES:MPI_COMM_FAILURE_GET_ACKED, [13](#), [14](#)
 EXAMPLES:MPI_COMM_FREE, [12](#), [13](#), [15](#)
 EXAMPLES:MPI_COMM_GROUP, [13](#)
 EXAMPLES:MPI_COMM_REVOKE, [15](#)
 EXAMPLES:MPI_COMM_SHRINK, [13](#), [15](#)
 EXAMPLES:MPI_COMM_SPLIT, [12](#)
 EXAMPLES:MPI_GROUP_DIFF, [13](#)
 EXAMPLES:MPI_GROUP_FREE, [13](#)

 MPI_ALLREDUCE, [15](#)
 MPI_BARRIER, [3](#)
 MPI_BCAST, [3](#)
 MPI_COMM_ACCEPT, [4](#)
 MPI_COMM_AGREE, [4](#), [6–10](#), [12](#), [13](#), [15](#)
 MPI_COMM_AGREE(comm, flag), [9](#)
 MPI_COMM_CONNECT, [4](#)
 MPI_COMM_DUP, [4](#)
 MPI_COMM_FAILURE_ACK, [8](#), [9](#), [13](#)
 MPI_COMM_FAILURE_ACK(comm), [7](#)
 MPI_COMM_FAILURE_GET_ACKED, [8](#), [13](#)
 MPI_COMM_FAILURE_GET_ACKED(comm, failedgrp), [8](#)
 MPI_COMM_FREE, [4](#)
 MPI_COMM_IAGREE, [7](#)
 MPI_COMM_IAGREE(comm, flag, req), [10](#)
 MPI_COMM_JOIN, [4](#)
 MPI_COMM_REVOKE, [7](#), [15](#)
 MPI_COMM_REVOKE(comm), [6](#)
 MPI_COMM_SHRINK, [6–8](#), [13](#)
 MPI_COMM_SHRINK(comm, newcomm), [7](#)
 MPI_COMM_SPAWN, [4](#)
 MPI_COMM_SPAWN_MULTIPLE, [4](#)
 MPI_COMM_SPLIT, [4](#), [7](#), [12](#)
 MPI_FILE_CLOSE, [6](#)
 MPI_FILE_REVOKE, [12](#)
 MPI_FILE_REVOKE(fh), [11](#)
 MPI_FILE_SEEK, [6](#)
 MPI_FILE_FINALIZE, [3](#)
 MPI_GROUP_DIFFERENCE, [8](#)
 MPI_INIT, [3](#)
 MPI_WIN_FREE, [5](#)
 MPI_WIN_GET_FAILED, [13](#)
 MPI_WIN_GET_FAILED(win, failedgrp), [11](#)

MPI_WIN_REVOKE, [5](#), [10](#)
MPI_WIN_REVOKE(win), [10](#)
MPI_WIN_REVOKED, [11](#)